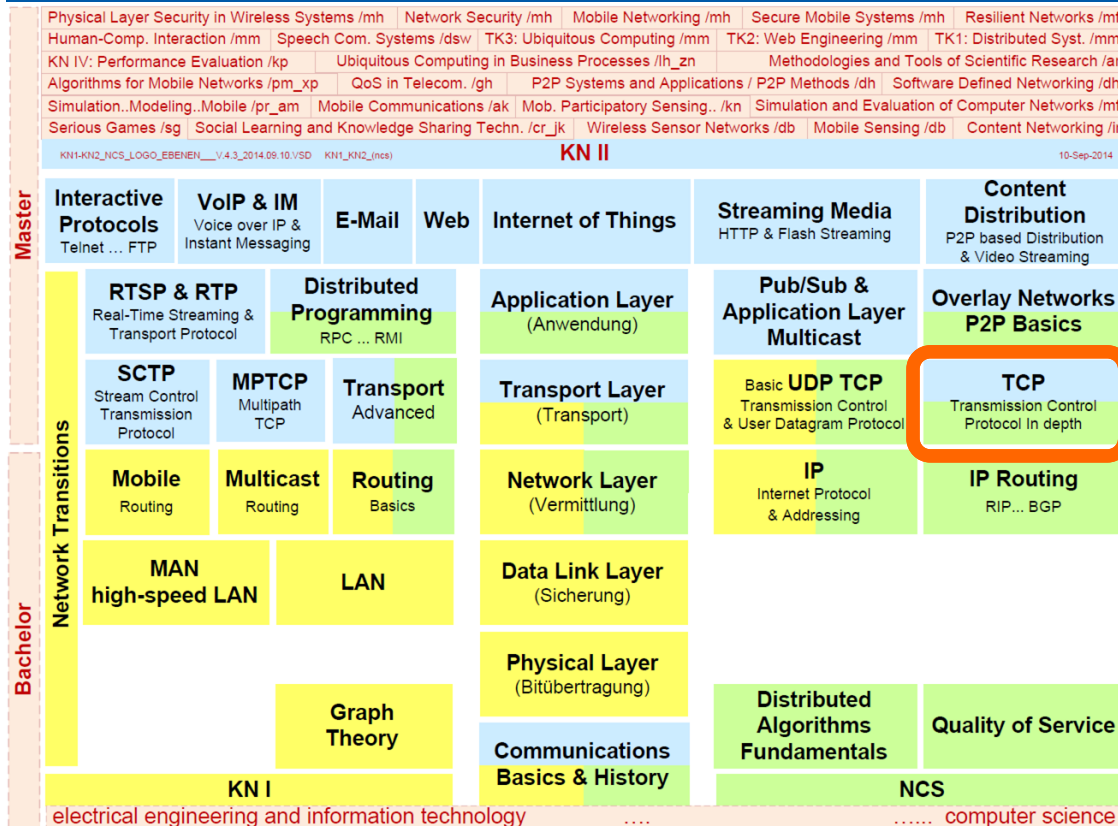


# Communication Networks II

## Transmission Control Protocol - TCP



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



10. September 2014

- 1 TCP – Transmission Control Protocol - Basics**
- 2 TCP Message Format**
- 3 Segments & Fragments**
- 4 TCP Connection**
  - 4.1 Connection Management**
  - 4.2 Connection - Setup**
  - 4.3 Connection – Tear Down**
  - 4.4 Sequence Number Space**
  - 4.5 Error Control**
- 5 Flow Control – in General**
  - 5.1 TCP Flow Control – in General**
  - 5.2 TCP Window Flow Control - Special Cases**
  - 5.3 Ongoing Communication**
- 6 Round Trip Time**
  - 6.1 Setting Retransmission Timeout (RTO)**
  - 6.2 Original TCP Round-trip Estimator (Jacobson, 1988)**
  - 6.3 RTT Sample Ambiguity**
- 7 Congestion Control**
  - 7.1 Congestion at Transport Layer - in General**
  - 7.2 Limits of TCP Window Flow Control**
  - 7.3 Three Congestion Control Problems**
  - 7.4 General Approaches**
  - 7.5 TCP Congestion Control**
  - 7.6 Flow Control & Congestion Control**
  - 7.7 Congestion Window (cwnd)**
  - 7.8 Basic Component 1: Detecting Congestion**
  - 7.9 Basic Component 2: Rate Adjustment**
- 8 Fairness of TCP Congestion Control**
  - 8.1 Additive Increase Multiplicative Decrease (AIMD)**
  - 8.2 Additive Increase Additive Decrease (AIAD)**
- 9 Phases of TCP Congestion Control**
  - 9.1 Initialization**
  - 9.2 Phase 1: TCP Slow Start**
  - 9.3 Phase 2: Congestion Avoidance (After Slow Start)**
  - 9.4 Example of Slow Start + Congestion Avoidance**
  - 9.5 TCP illustrated**

Physical Layer Security in Wireless Systems /mh	Network Security /mh	Mobile Networking /mh	Secure Mobile Systems /mh	Resilient Networks /mf
Human-Comp. Interaction /mm	Speech Com. Systems /dsw	TK3: Ubiquitous Computing /mm	TK2: Web Engineering /mm	TK1: Distributed Syst. /mm
KN IV: Performance Evaluation /kp	Ubiquitous Computing in Business Processes /lh_zn		Methodologies and Tools of Scientific Research /ar	
Algorithms for Mobile Networks /pm_xp	QoS in Telecom. /gh	P2P Systems and Applications / P2P Methods /dh	Software Defined Networking /dh	
Simulation..Modeling..Mobile /pr_am	Mobile Communications /ak	Mob. Participatory Sensing.. /kn	Simulation and Evaluation of Computer Networks /mf	
Serious Games /sg	Social Learning and Knowledge Sharing Techn. /cr_jk	Wireless Sensor Networks /db	Mobile Sensing /db	Content Networking /ir

KN1-KN2\_NCS\_LOGO\_EBENEN\_\_\_V.4.3\_2014.09.10.VSD KN1\_KN2\_(ncs)

## KN II

10-Sep-2014

Master

<b>Interactive Protocols</b> Telnet ... FTP	<b>VoIP &amp; IM</b> Voice over IP & Instant Messaging	<b>E-Mail</b>	<b>Web</b>	<b>Internet of Things</b>	<b>Streaming Media</b> HTTP & Flash Streaming	<b>Content Distribution</b> P2P based Distribution & Video Streaming
--	---	---------------	------------	---------------------------	--	---

<b>RTSP &amp; RTP</b> Real-Time Streaming & Transport Protocol	<b>Distributed Programming</b> RPC ... RMI		<b>Application Layer</b> (Anwendung)	<b>Pub/Sub &amp; Application Layer Multicast</b>	<b>Overlay Networks</b> <b>P2P Basics</b>
---	---	--	---	--	--

<b>SCTP</b> Stream Control Transmission Protocol	<b>MPTCP</b> Multipath TCP	<b>Transport Advanced</b>	<b>Transport Layer</b> (Transport)	<b>Basic UDP TCP</b> Transmission Control & User Datagram Protocol	<b>TCP</b> Transmission Control Protocol In depth
---	-------------------------------	---------------------------	---------------------------------------	---	--

<b>Mobile Routing</b>	<b>Multicast Routing</b>	<b>Routing Basics</b>	<b>Network Layer</b> (Vermittlung)	<b>IP</b> Internet Protocol & Addressing	<b>IP Routing</b> RIP... BGP
-----------------------	--------------------------	-----------------------	---------------------------------------	---	---------------------------------

<b>MAN high-speed LAN</b>	<b>LAN</b>	<b>Data Link Layer</b> (Sicherung)
---------------------------	------------	---------------------------------------

<b>Graph Theory</b>	<b>Physical Layer</b> (Bitübertragung)	<b>Distributed Algorithms Fundamentals</b>	<b>Quality of Service</b>
---------------------	---	--	---------------------------


<b>KN I</b>	<b>Communications Basics &amp; History</b>	<b>NCS</b>
-------------	--	------------

electrical engineering and information technology ..... computer science

# 1 TCP – Transmission Control Protocol - Basics


## Internet layer offers best effort packet delivery

- From host to host



## TCP offers reliable byte stream


- From application to application



# UDP – User Datagram Protocol (vs. TCP)


## Internet layer offers best effort packet delivery

- From host to host



## UDP offers best effort message delivery

- From application to application



# TCP – Transmission Control Protocol - Basics

## **Motivation: network layer provides unreliable connectionless service**

- packets and messages may be
  - duplicated, delivered in wrong order, faulty
- given such an unreliable service
  - each application would have to implement error detection and correction separately
- network or service can
  - impose packet length
  - define additional requirements to optimize data transmission
  - i.e. each application would have to be adapted separately
- → do not reinvent the wheel for every application

## **→ TCP is the Internet transport protocol providing**

- reliable end-to end byte stream over an unreliable internetwork

## **Specification:**

- RFC 793 - Transmission Control Protocol: originally
- RFC 1122 and RFC 1323: errors corrected, enhancements implemented

# TCP in Use & Application Areas

## Each machine supporting TCP has a TCP transport entity composed of

- library procedure
- user process
- part of kernel

## TCP transport entity manages

- TCP streams
- interfaces to IP layer

## TCP transport entity at sending side

- accepts user data streams from local processes
- splits them into pieces  $\leq 64$  KB
  - typically 1460 bytes  
(to fit into single Ethernet frame with IP and TCP headers)
- sends each piece as separate IP datagram

## TCP transport entity at receiving side

- gets TCP data from datagram received at host
- reconstructs original byte streams

# TCP in Use & Application Areas

## Two-way communications (fully duplex)

- data may be transmitted simultaneously in both directions over a TCP connection

## Point-to-point

- each connection has exactly two endpoints

## TCP must ensure reliability

- IP layer doesn't guarantee that datagram will be delivered properly / in order
  - TCP must handle this, e.g. timeout and retransmit / reorder
    - i.e. reliable
- fully ordered, fully reliable
  - sequence maintained
  - no data loss, no duplicates, no modified data



# TCP in Use & Application Areas

## Benefits of TCP

- reliable data transmission
- efficient data transmission despite complexity
  - (up to 8 Mbps on 10 Mbps Ethernet)
- can be used with LAN and WAN for
  - low data rates  
(e.g. interactive terminal)
  - high data rates  
(e.g. file transfer)

## Disadvantages when compared with UDP

- higher resource requirements
  - buffering, status information, timer usage
  - connection set-up and disconnect necessary
    - even in case of short data transmissions

## Applications

- file transfer (FTP)
- interactive terminal (Telnet)
- e-mail (SMTP)
- X-Windows

## Some Missing Characteristics

### **no broadcast**

- no possibility to address all applications at the same time with a single message

### **no multicasting**

- group addressing not possible

### **no QoS parameters**

- not suited for different media characteristics

### **no real-time support**

- no correct treatment/communications of audio or video possible
- e.g. no Forward Error Correction (FEC)

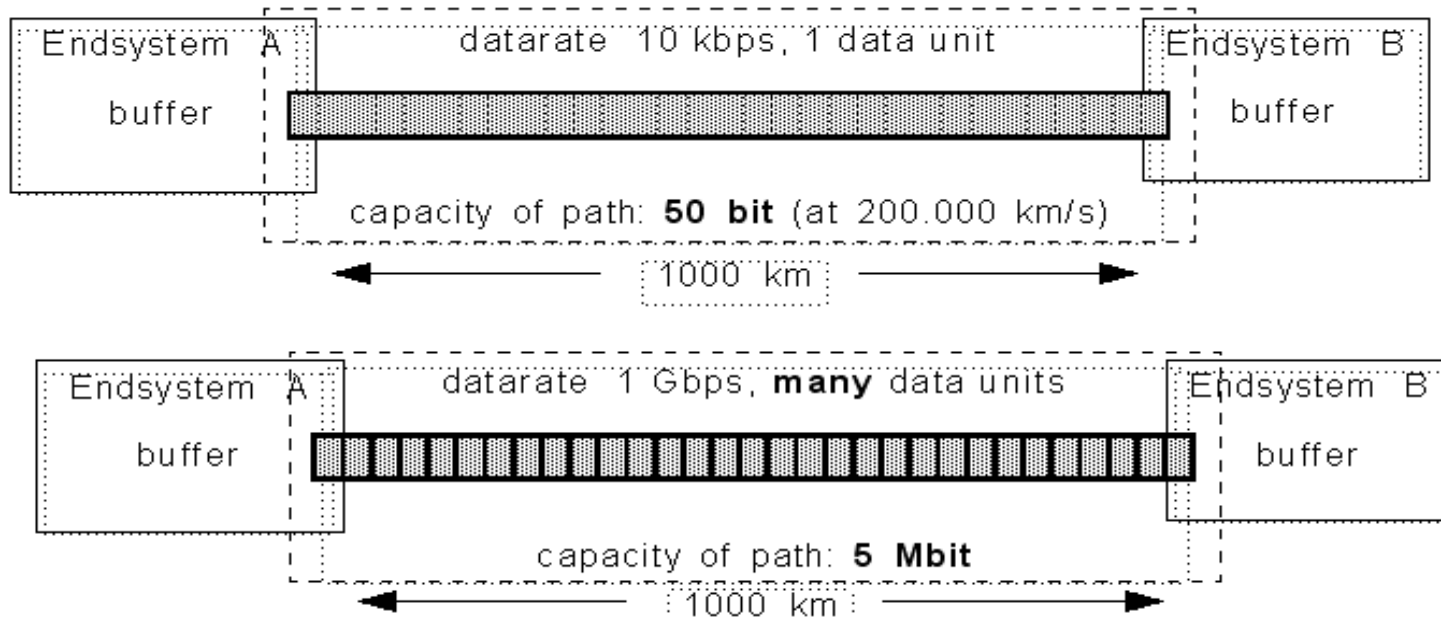
## Motivation

- networks and applications have changed

## Networks

- higher data rates
- also farther distances (e.g. also via satellite)
- networks

$$\text{Data amount} = \frac{\text{Data rate} \times \text{Distance}}{\text{Velocity of Propagation}}$$



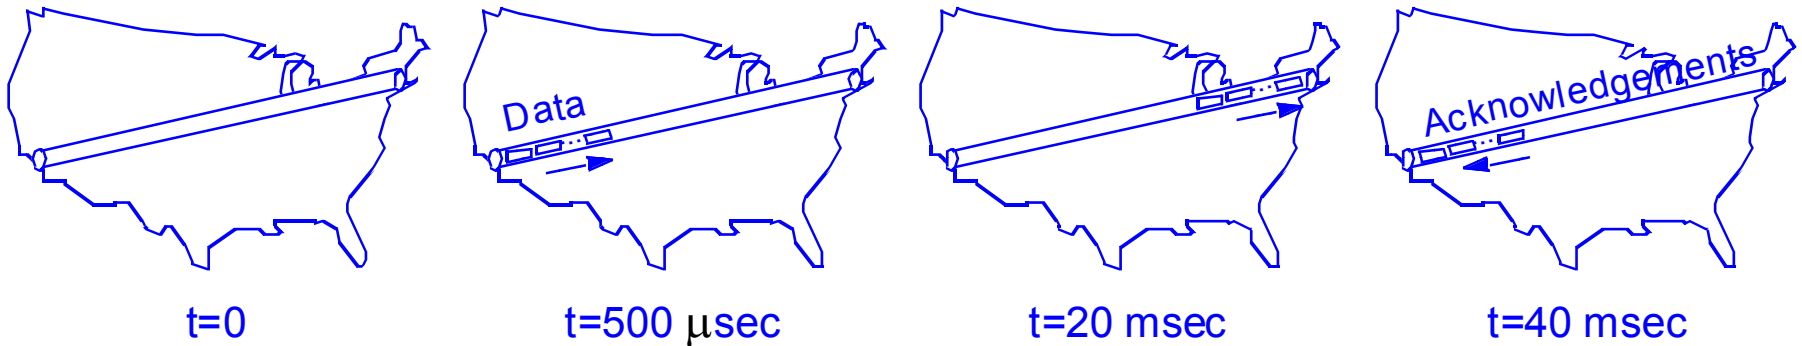
# Further Development of Transport Protocols

## Bandwidth-Delay Product increases

- $\text{bandwidth [bits/sec]} * \text{round-trip delay [sec]}$

## Useful parameter for network performance analysis

- Capacity of pipe from sender to receiver and back (in bits)



## Example:

- Transmission from San Diego to Boston
  - sending 64 KB burst (receiver buffer 64 KB), link: 1 Gbps
  - one-way propagation delay (speed-of-light in fiber): 20 msec
- Bandwidth-delay product: 40 million bit
- i.e.: sender would have to transmit burst of 40 million bits to keep pipe busy till ACK

## Receiver window must be $\geq$ bandwidth-delay product

- for good performance

## 2 TCP Message Format



### Reliable bidirectional in-order byte stream

- Socket: SOCK\_STREAM

### Connections established & torn down

### Multiplexing/ demultiplexing

- ports at both ends

### Error control

- users see correct, ordered byte sequences

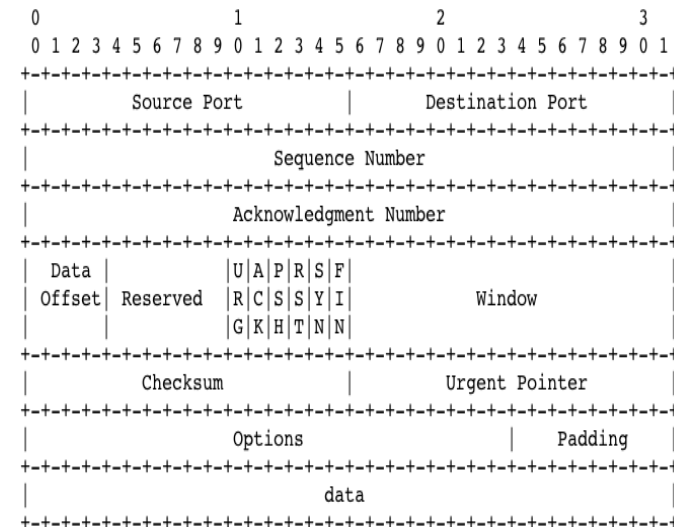
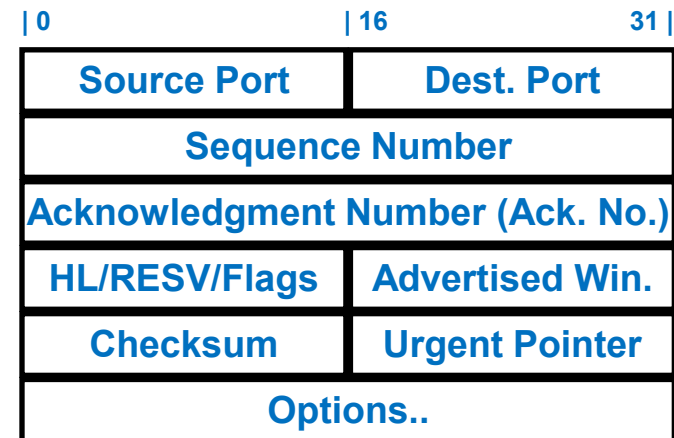
### End-to-end flow control

- avoid overwhelming the machines at either end

### Congestion avoidance

- avoid creating traffic jams within network

### TCP Header:



# TCP Message Format



## Variable length header

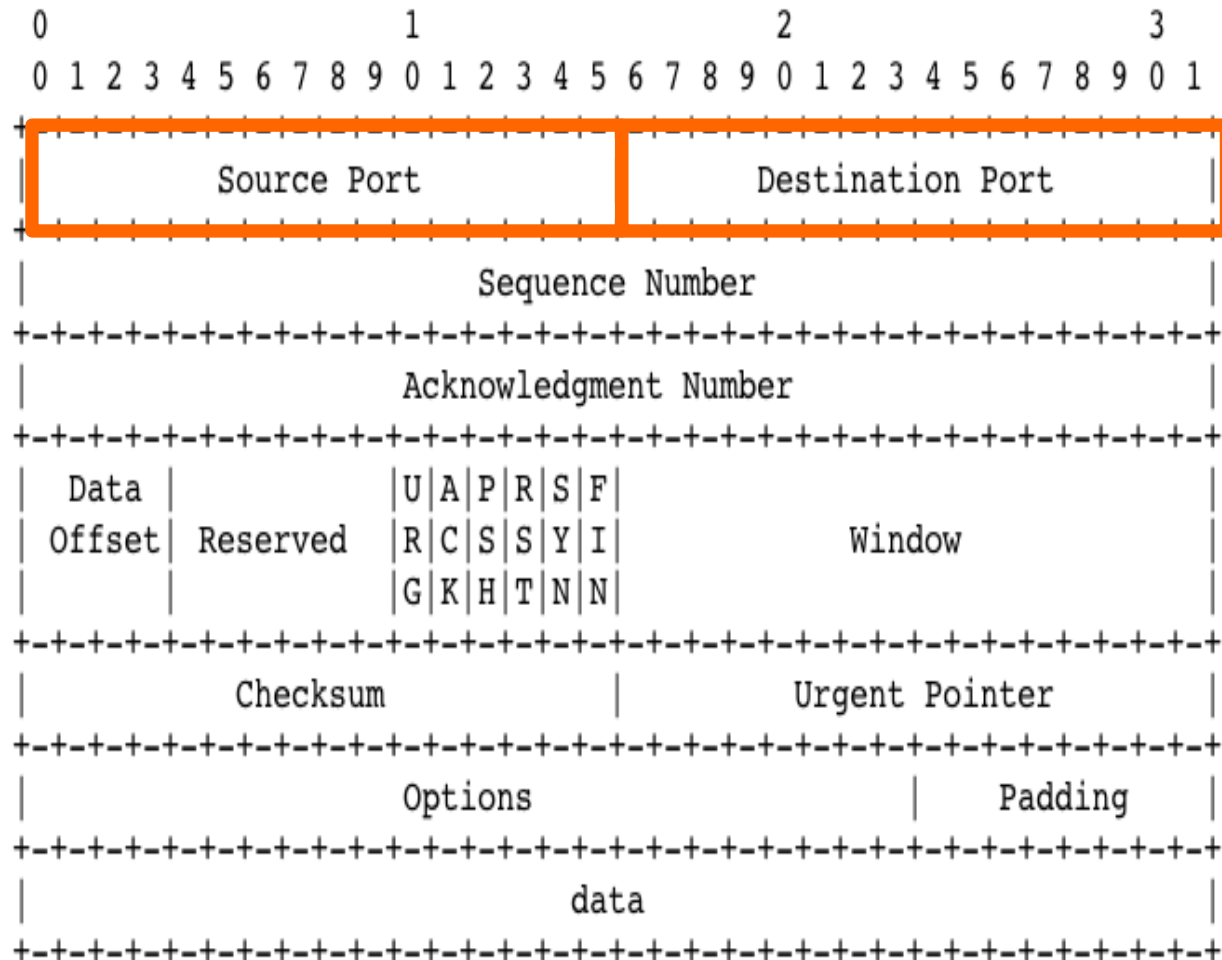
- Min. 20 byte
- Variable length options
- Multiple of 4 byte

## Source port

- 16 bit identifier of sending application

## Destination port

- 16 bit identifier of receiving application



TCP Header Format

# Message Format - Ports

## Port numbers identify sending/receiving application/process

- In analogy to UDP ports
  - Ports 0 – 40151 assigned by Internet Assigned Numbers Authority (IANA)
- System ports (or: well known ports), 0 – 1023
  - E.g. ports 20 and 21: File Transfer Protocol (FTP) data and control
  - E.g. port 22: Secure Shell (SSH)
  - E.g. port 25: Simple Mail Transfer Protocol (SMTP)
  - E.g. port 80: Hypertext Transfer Protocol (HTTP)
- User ports (or: registered ports), 1024 – 40151
  - E.g. port 1194: OpenVPN
  - E.g. port 3689: Digital Audio Access Protocol (DAAP)
  - E.g. port 17500: Dropbox LANsync data
    - Compare UDP port 17500 Dropbox LANsync discovery → conflict?
- Dynamic ports (or: private/ephemeral ports)
  - 40152 – 65535 for dynamic use

# Message Format

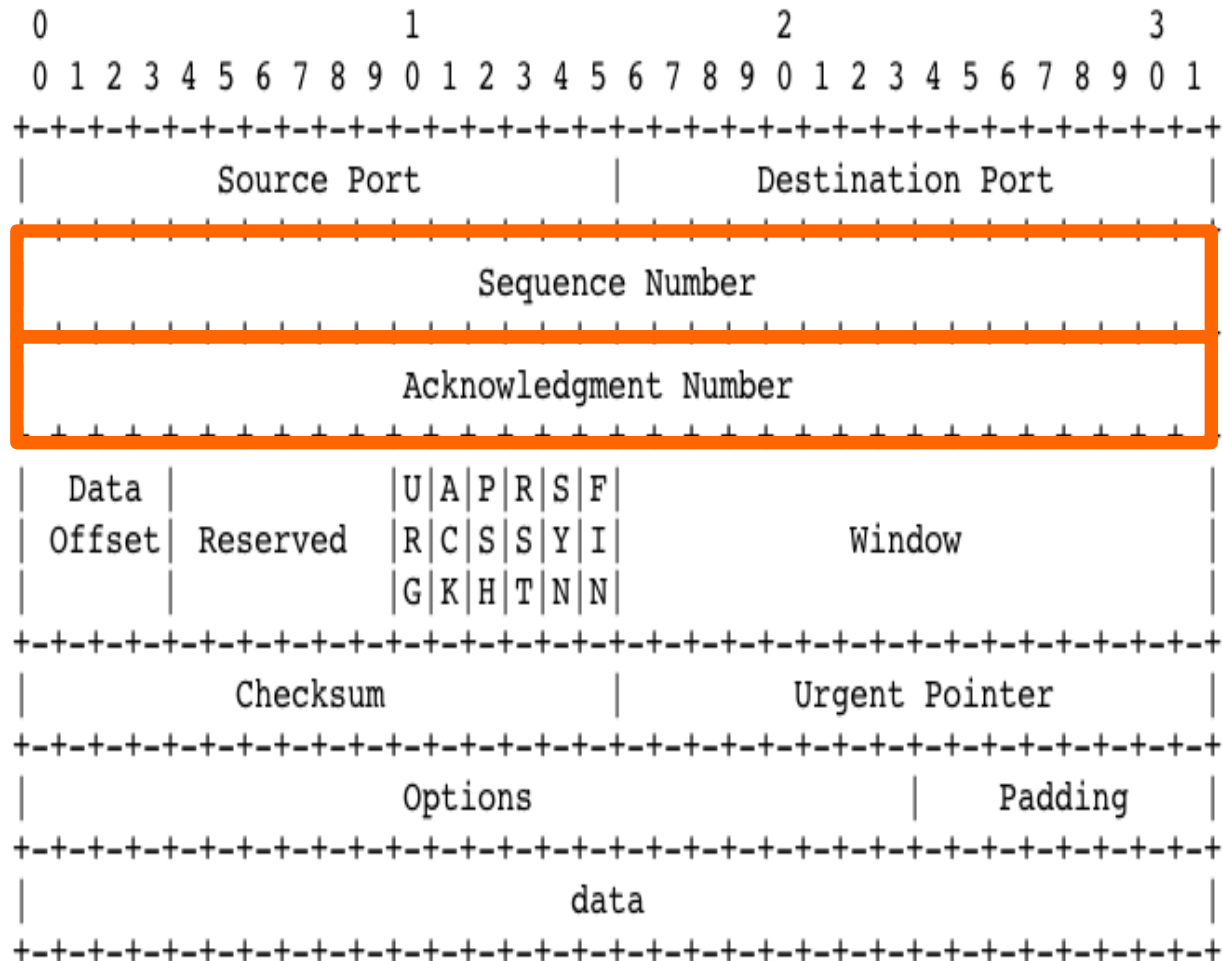


## Sequence number

- Connection setup:  
Initial sequence  
number negotiation
- Data transfer:  
Number of first byte  
in data field

## Acknowledgment number

- Connection setup:  
Initial sequence  
number negotiation
- Data transfer:  
Next expected byte
  - Cumulative  
acknowledgment



TCP Header Format



# Message Format

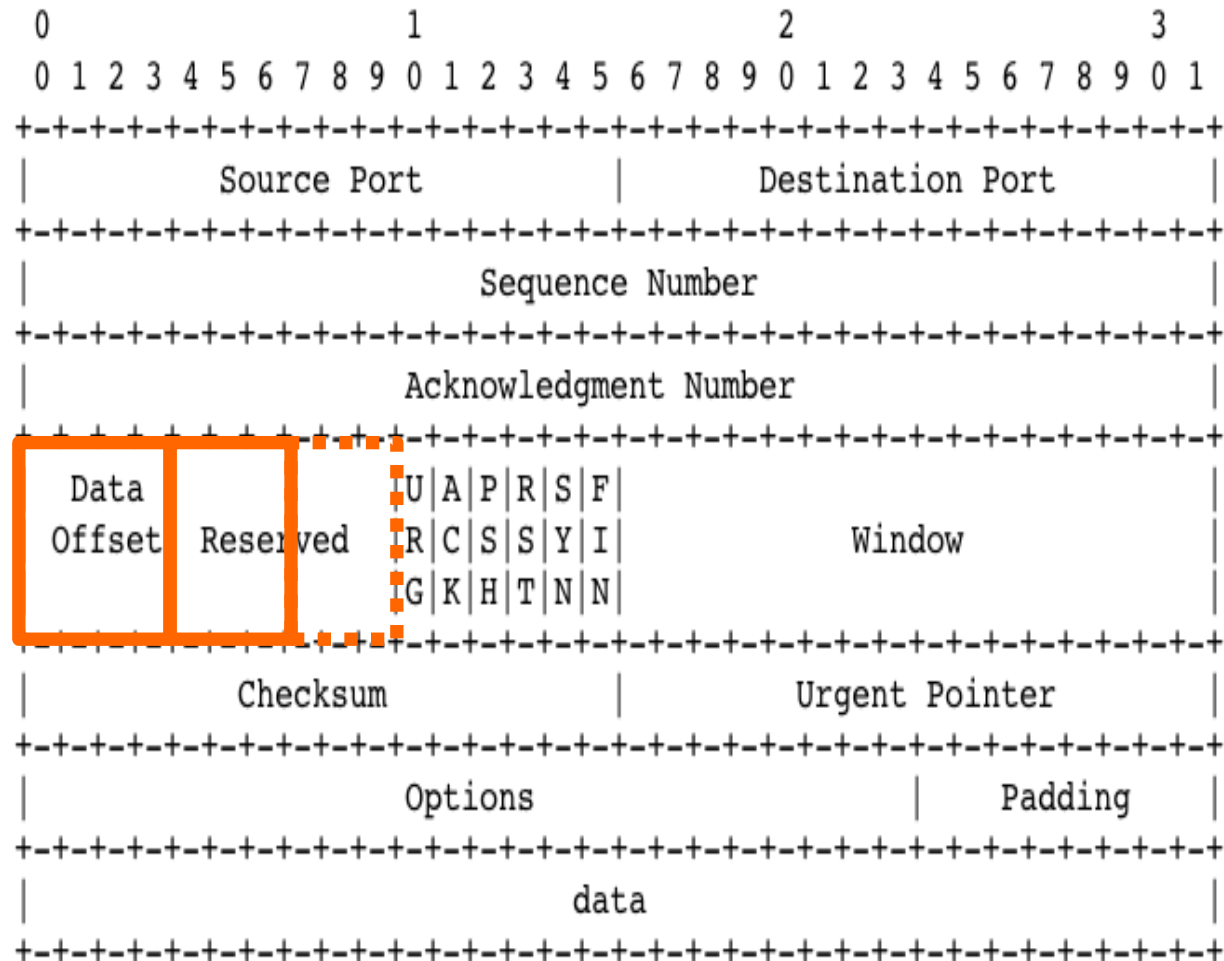


## Data offset

- Header length in words (4 byte / 32 bit)
- Indicates beginning of data field
- Remember: variable length options
- Also (later RFCs) called header length

## Reserved


- Originally 6 bits
  - RFC 793
- Now 3 bits
  - After RFCs 3168 and 3540
  - .. see next slide
- Must be zero



TCP Header Format

## Further Flags

- NS: nonce sum
  - Used for Explicit Congestion Notification (ECN) congestion control
  - Added June 2003 in RFC 3540
- CWR: congestion window reduced
  - Used for ECN
  - Added September 2001 in RFC 3168
- ECE: ECN Echo
  - Used for ECN
  - Added September 2001 in RFC 3168

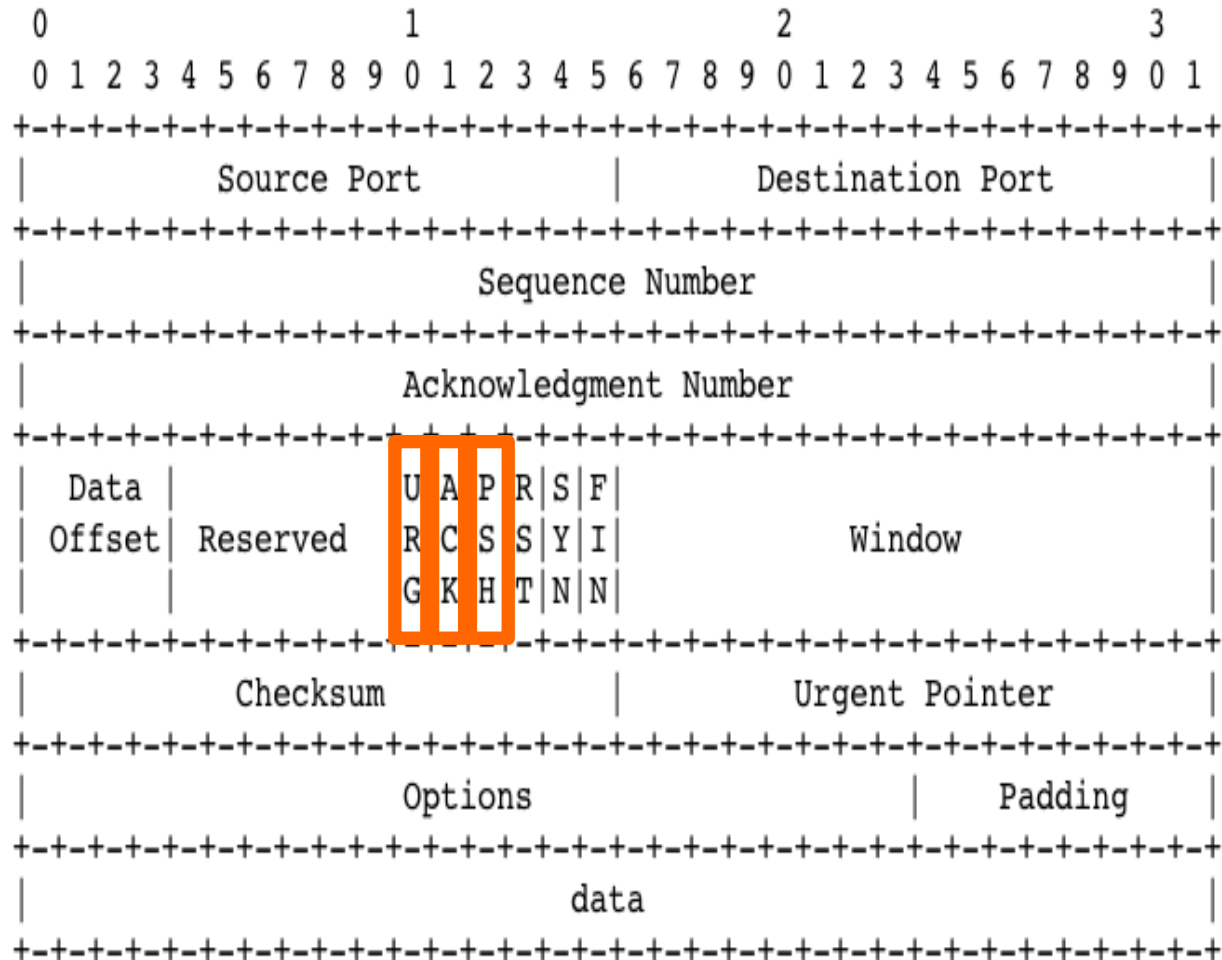


# Message Format



## Original flags

- URG: data field contains urgent data
  - Application layer to be notified immediately
  - Not used in practice
  
- ACK: message contains acknowledgment
  - Acknowledgment number significant
  
- PSH: push function
  - Data to be delivered to application layer immediately



TCP Header Format

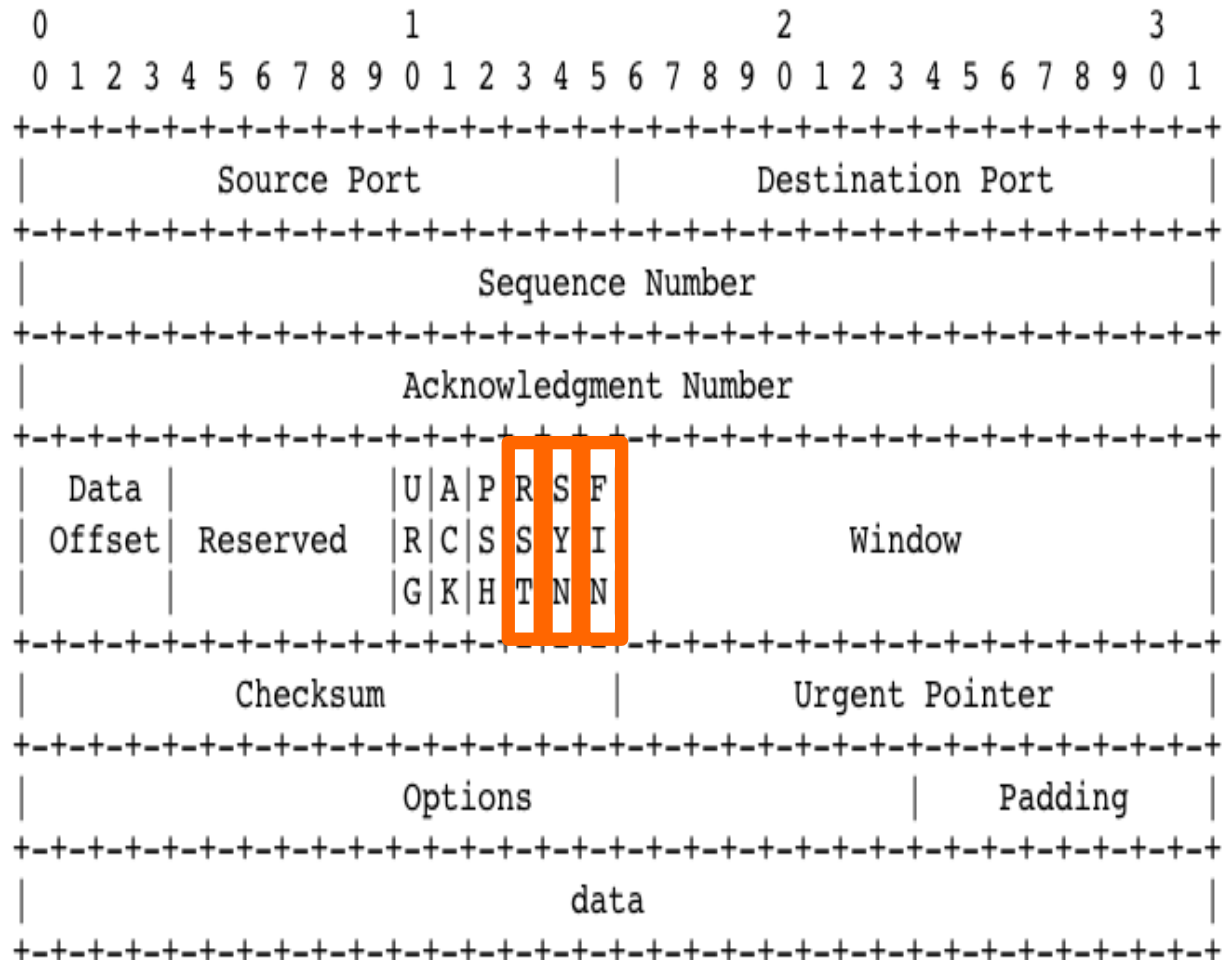
# Message Format



## Original flags

### Used for connection management

- RST: reset connection
- SYN: synchronize sequence numbers
- FIN: data transfer finished



TCP Header Format

# Message Format

## Window

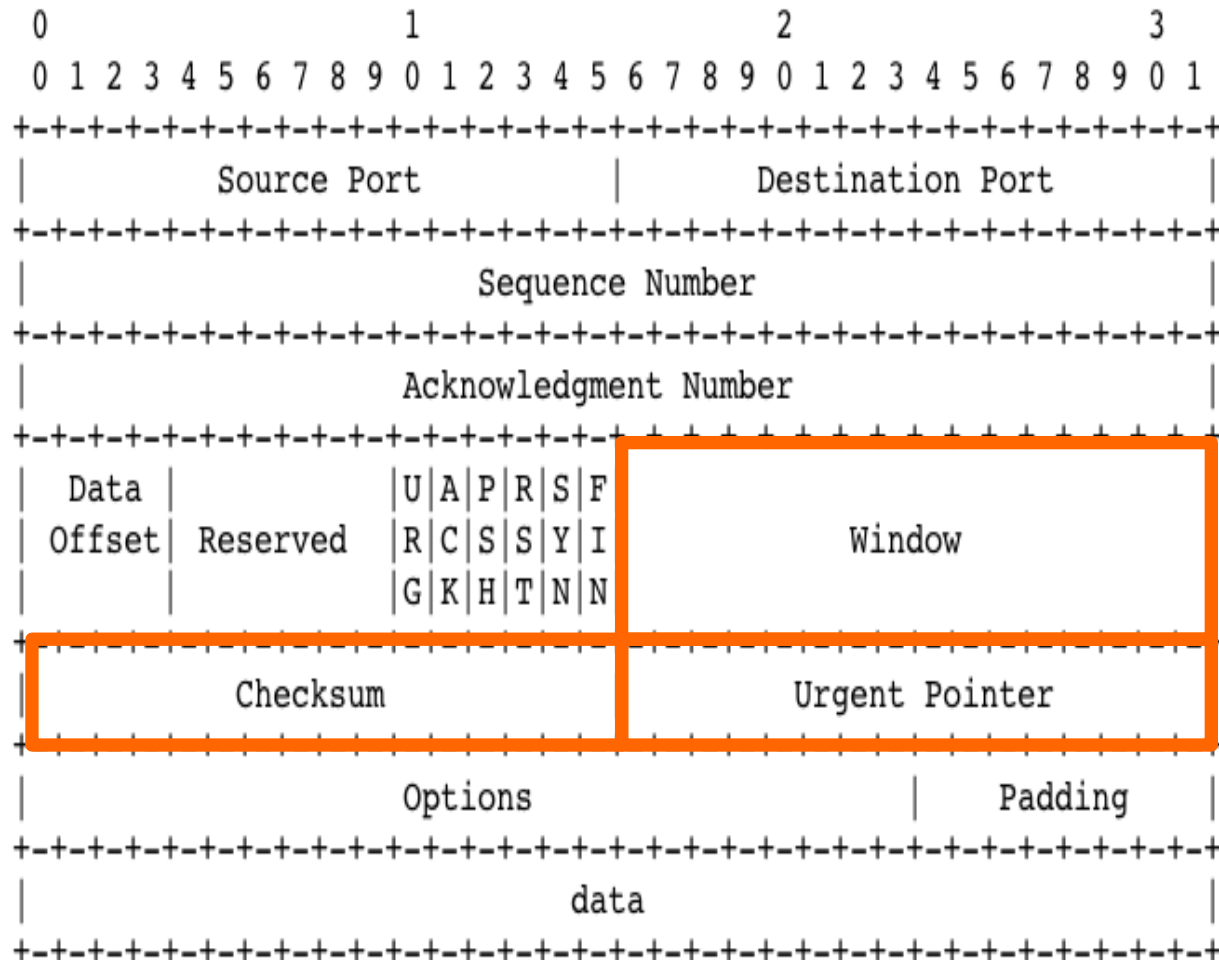
- Number of bytes sender of message can accept
- Indicates available buffer space of sender
- Used for flow control

## Checksum

- Used for error detection
- Same recipe as UDP

## Urgent pointer

- Number of urgent bytes in data field
- Not used in practice



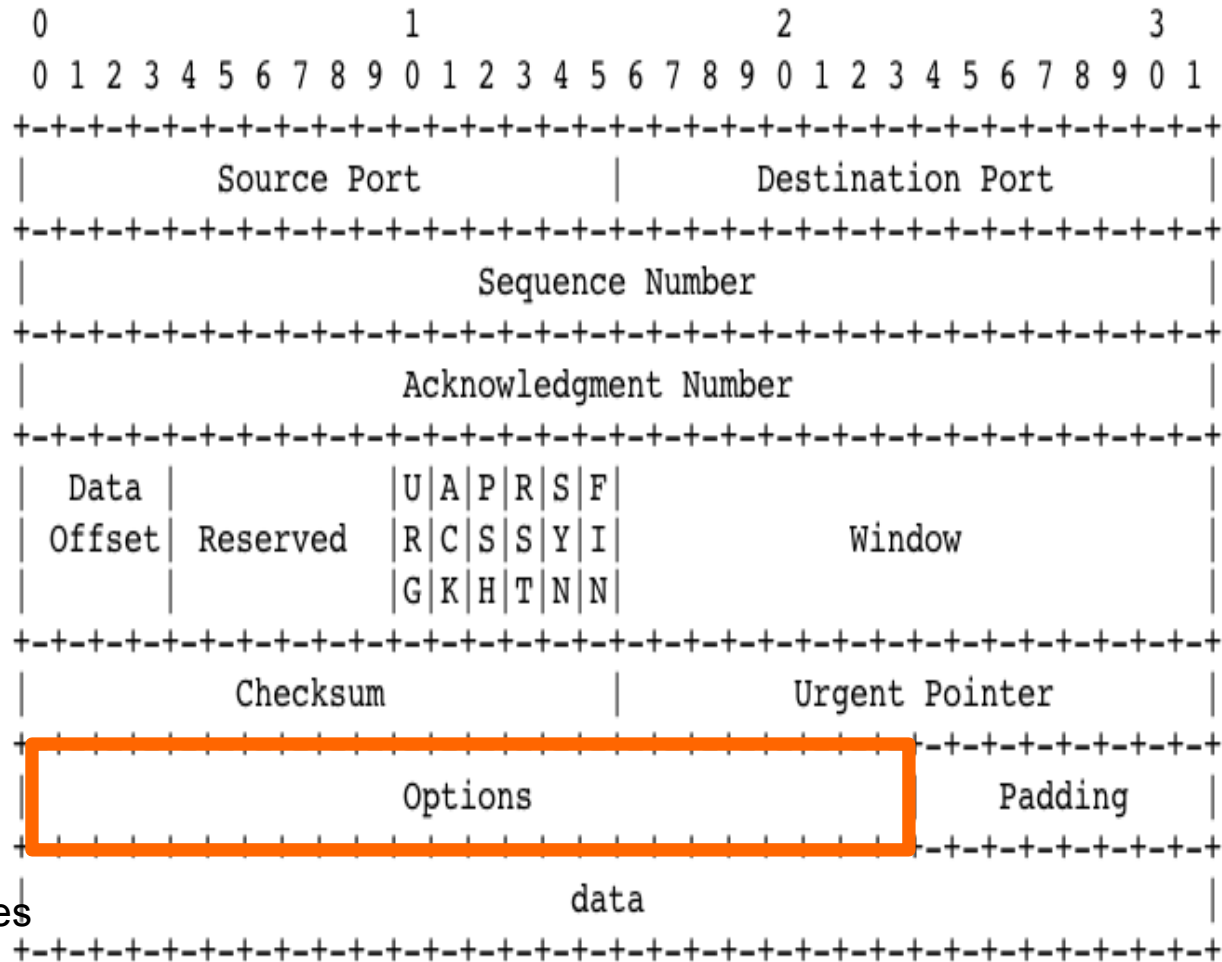
TCP Header Format

# Message Format



## Options

- Definition of protocol options
- Valid for connection
  
- E.g.
  - Definition of max. segment size
    - TCP message is called segment
  - Selective acknowledgment
    - Useful in certain cases of packet loss
  - Window scaling
    - Original window size limited to  $2^{16} = 65535$  bytes
    - Scaling factor used to optimize TCP for high bandwidth connections



TCP Header Format

# Message Format

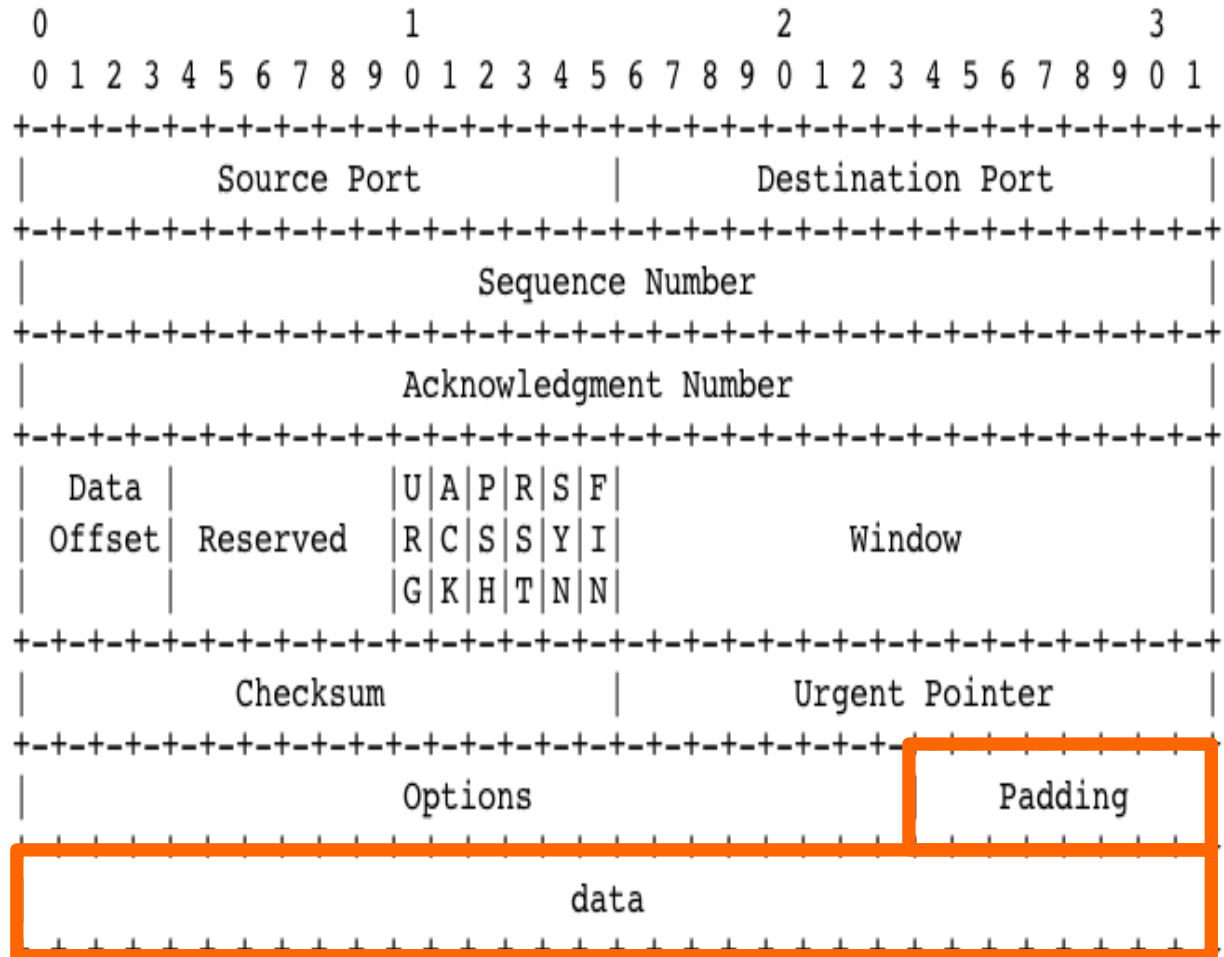


## Padding

- Zeros to align options to 4 byte boundary

## Data

- Application data
  - Application header
  - User data



TCP Header Format

### Segments

#### Challenge:

- Buffered data transmission
  - byte stream – is not a message stream

➔ **Segments are introduced**

#### Transport layer

- reassembles segments

### Fragments

#### Challenge:

- size of packets of underlying network are smaller than size of IP packet

➔ **Fragments are introduced**

#### Network layer

- reassembles fragments at final destination



# Fragments

## Challenge:

**size of packets of underlying network  
are smaller than size of IP packet**


- → Fragments are introduced

## Fragments

- IP packets are split (if necessary) into FRAGMENTS
  - in order to adapt them to underlying networks
- max. IP packet size is limited by MTU (maximum transfer unit) of underlying network,
  - e.g. for Ethernet MTU=1500 byte

## IP layer

- reassembles fragments at final destination



## Challenge: Buffered data transmission

- byte stream – is not a message stream:
  - message boundaries are not preserved
  - no way for receiver to detect the unit(s) in which data were written

## ➔Segments are introduced

## Segments

- TCP DATA STREAM is split into segments
  - SEGMENTS sent as payload of IP packets
  - max. segment size (mss) limits the size of a segment
  - mss is negotiated at connection setup
    - using TCP options (as discussed previously)

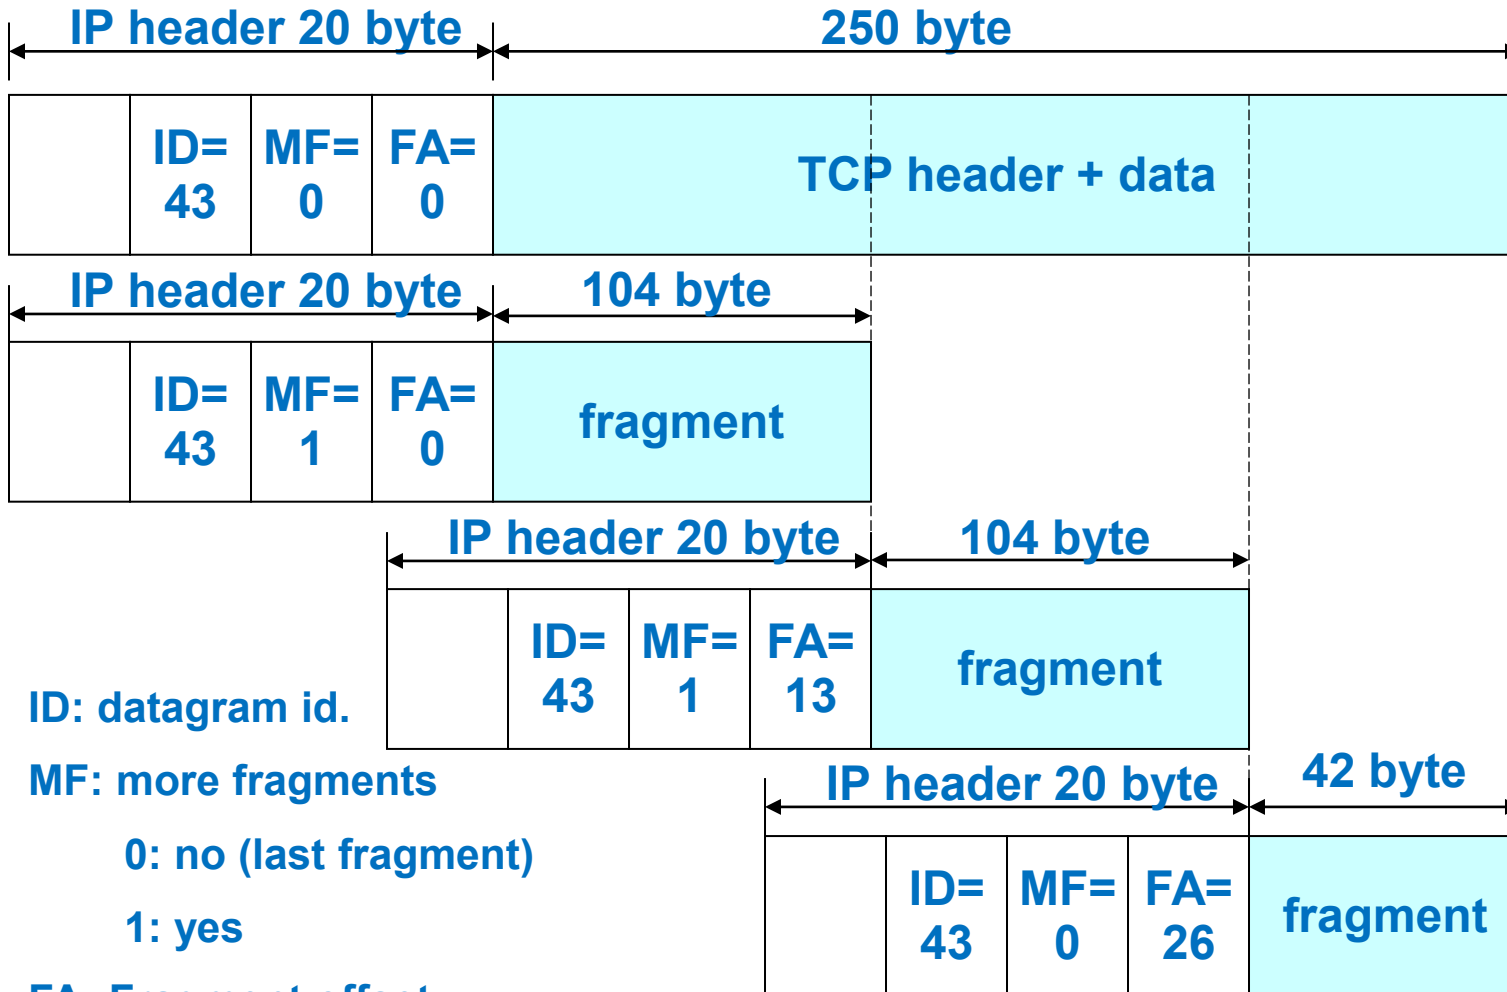
## Transport layer

- reassembles segments

# Segments & Fragments



## Fragmentation of segments e.g. case MTU = 128 Bytes



ID: datagram id.

MF: more fragments

0: no (last fragment)

1: yes

FA: Fragment offset

$n \rightarrow n * 8$  byte

### Why is there a need for connection setup?

#### Mainly to agree on starting sequence numbers

- starting sequence number is randomly chosen
- reason:
  - to reduce the chance that sequence numbers of old and new connections overlap

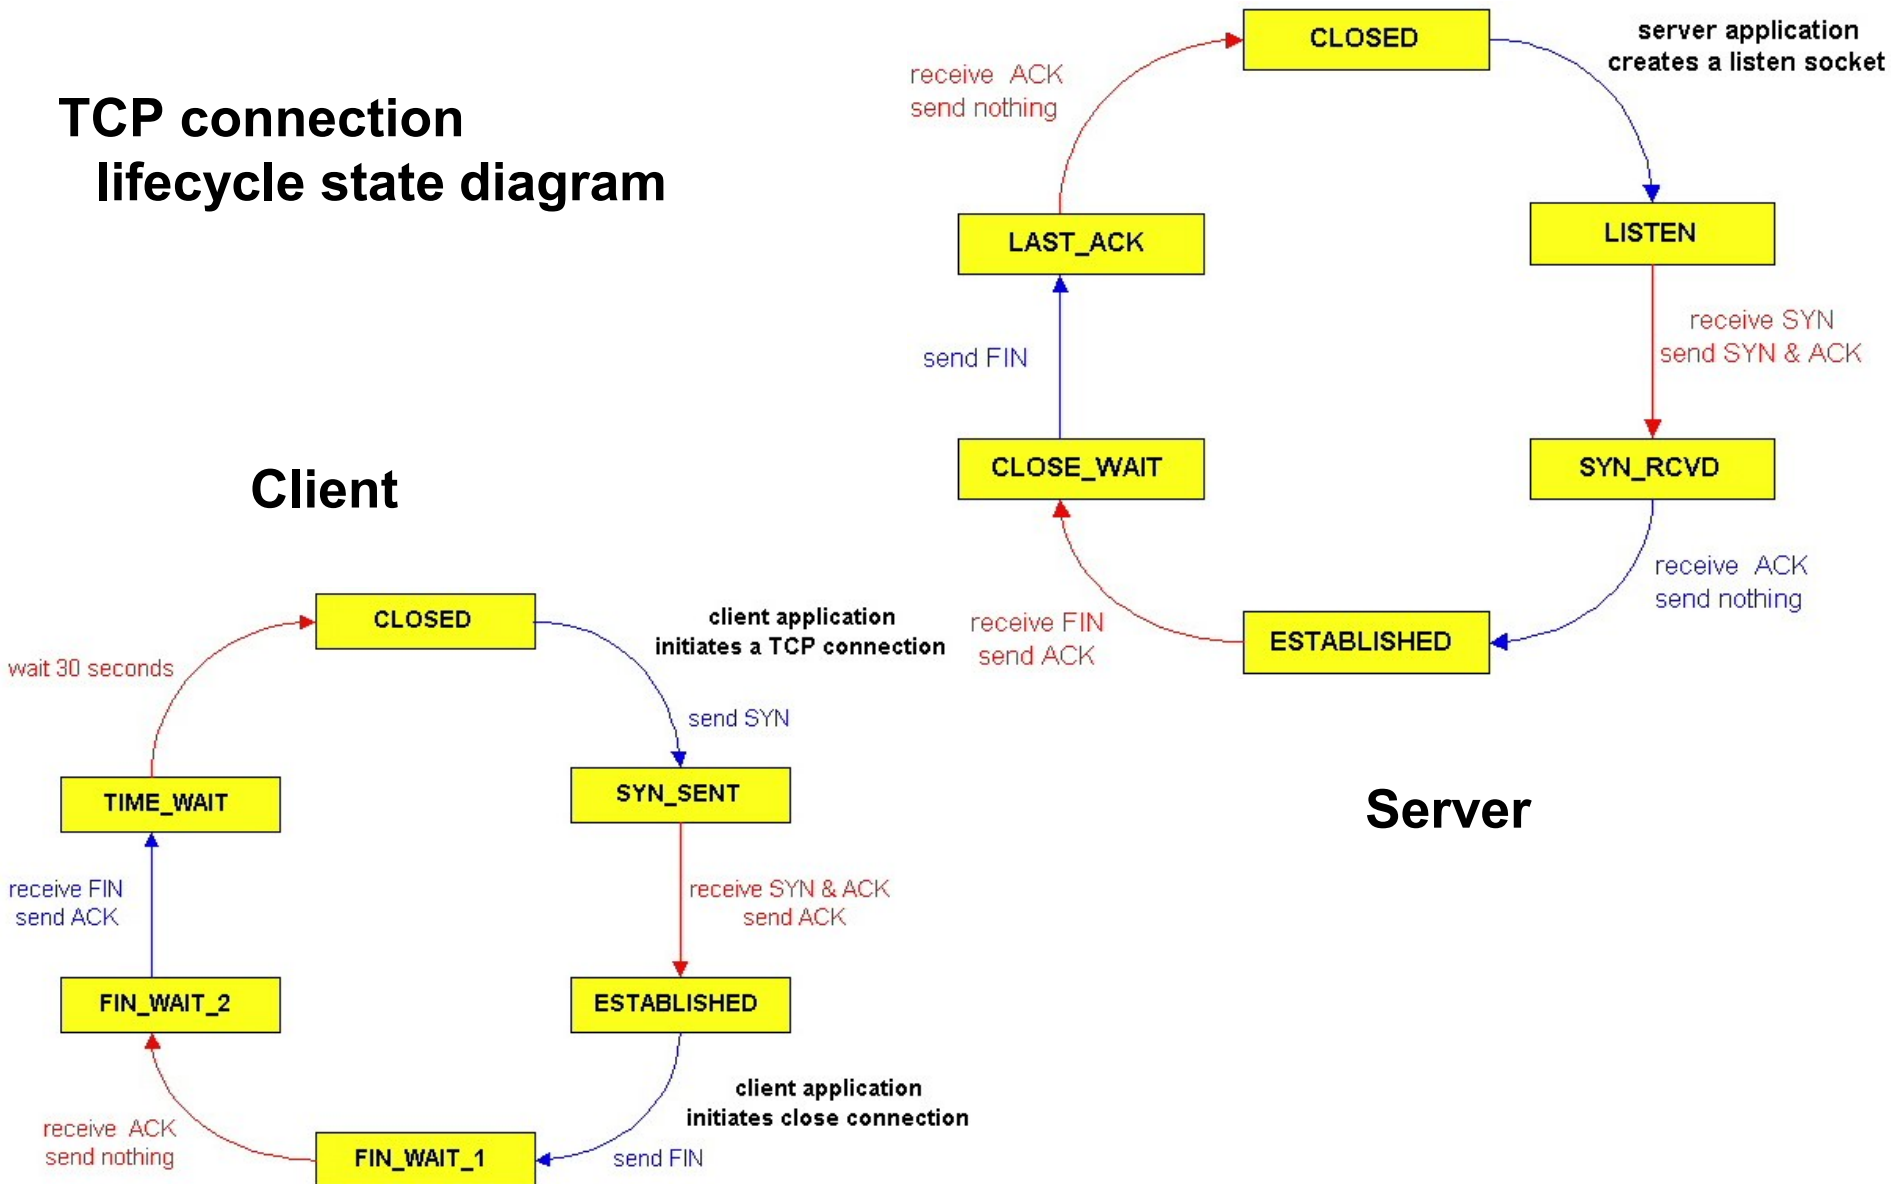
## 4.1 Connection Management



### TCP connection lifecycle state diagram

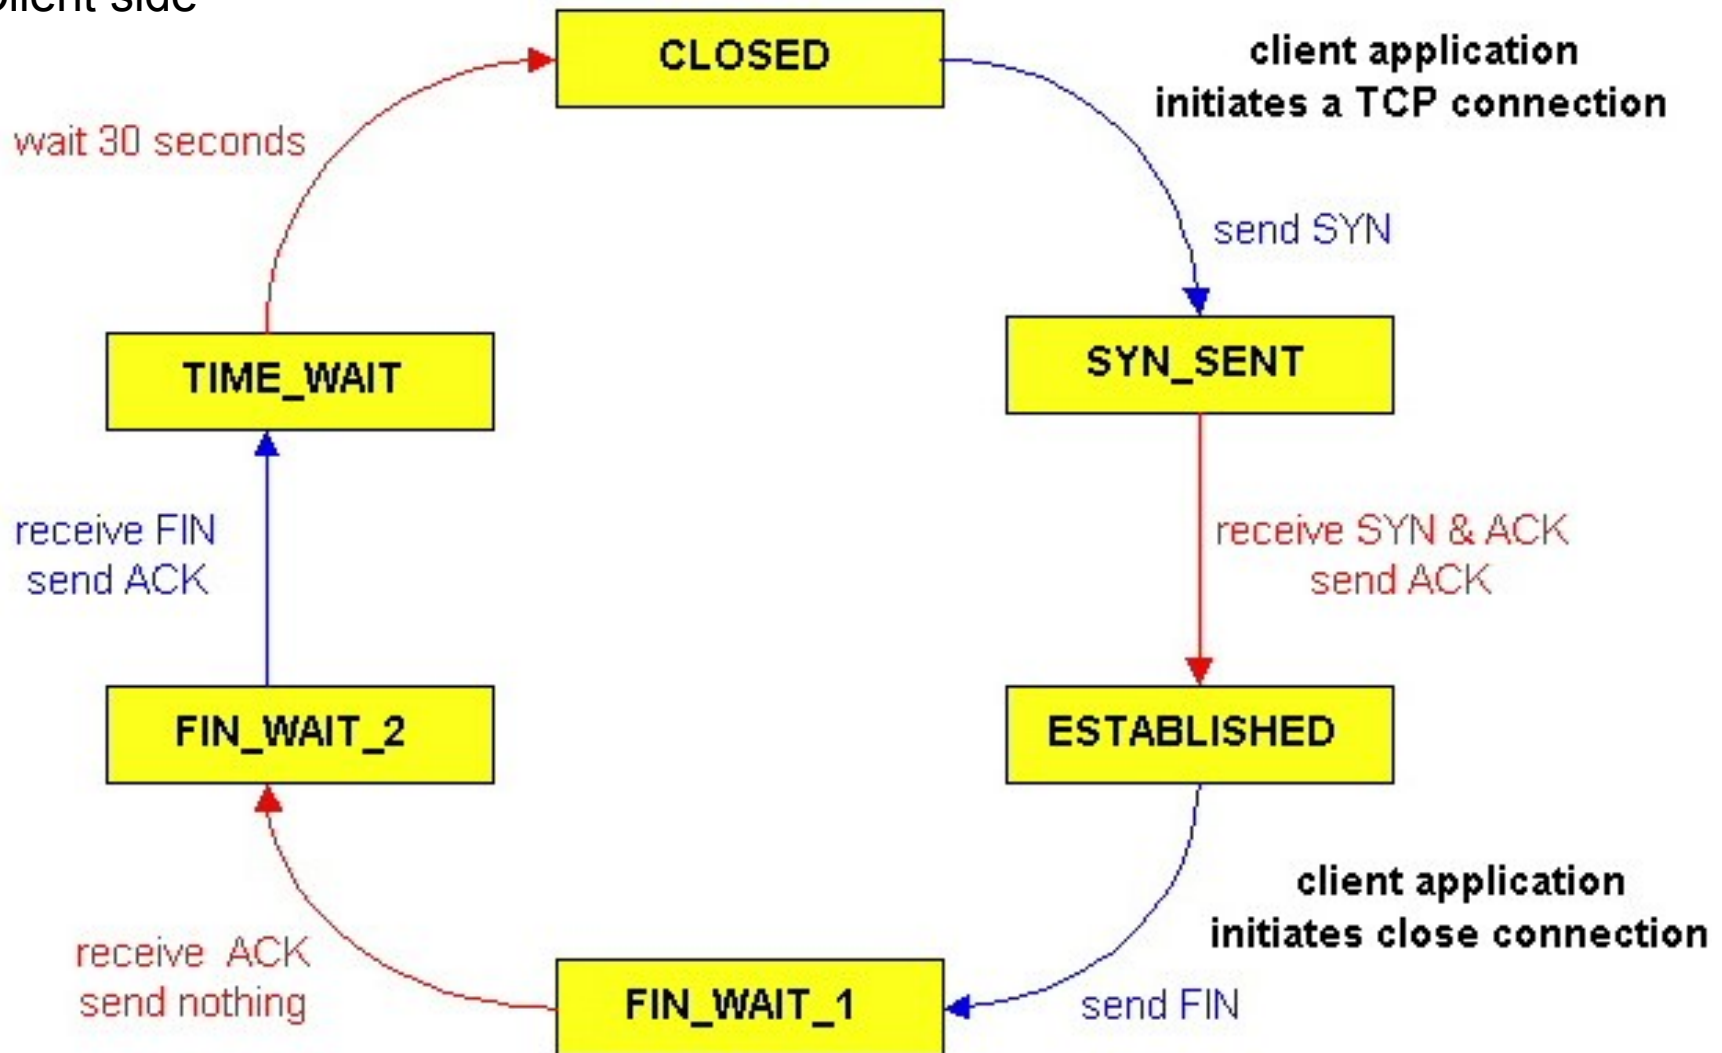
#### Client

#### Server



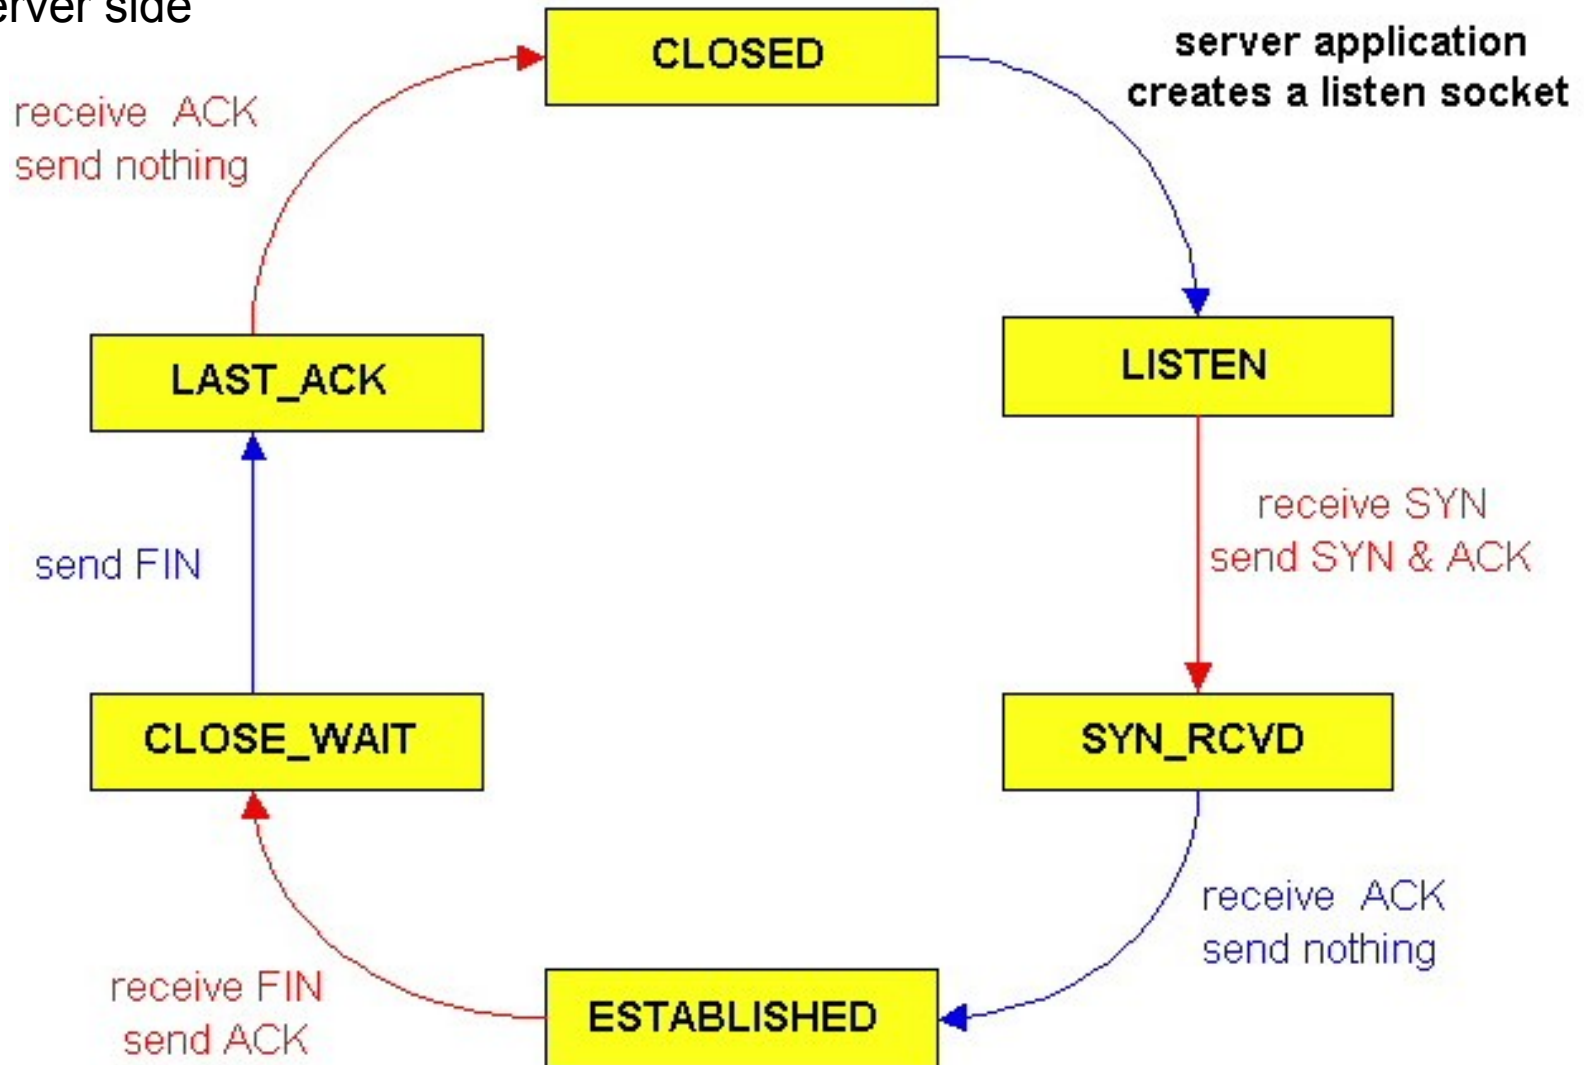
## TCP connection lifecycle state diagram

- Client side



## TCP connection lifecycle state diagram

- Server side




## 4.2 Connection - Setup



### TCP connection setup: 3-way handshake

- Step 1:  
client sends message with
  - SYN flag set
  - Sequence number (SN) field containing Client Initial Sequence Number ( $ISN_C$ )
- Step 2:  
server sends message with
  - SYN and ACK flags set
  - Acknowledgment field containing  $ISN_C+1$
  - Sequence number field containing Server Initial Sequence Number ( $ISN_S$ )
- Step 3:  
client send message with
  - ACK flag set
  - Acknowledgment field containing  $ISN_S+1$
  - Sequence number field containing  $ISN_C+1$






# Connection Setup

## TCP connection setup – numerical example

- Step 1:  
client sends message with
  - SYN flag set
  - Client Initial Sequence Number=100
- Step 2:  
server sends message with
  - SYN and ACK flags set
  - Client's Initial Sequence Number acknowledged
    - By setting acknowledgment field to 101
    - Next expected sequence number
  - Server Initial Sequence Number=300
- Step 3:  
client send message with
  - ACK flag set
  - Server's Initial Sequence Number acknowledged
    - By setting acknowledgment field to 301
    - Next expected sequence number
  - Sequence number field set to 101



## 4.3 Connection – Tear Down

### Either side can initiate tear down

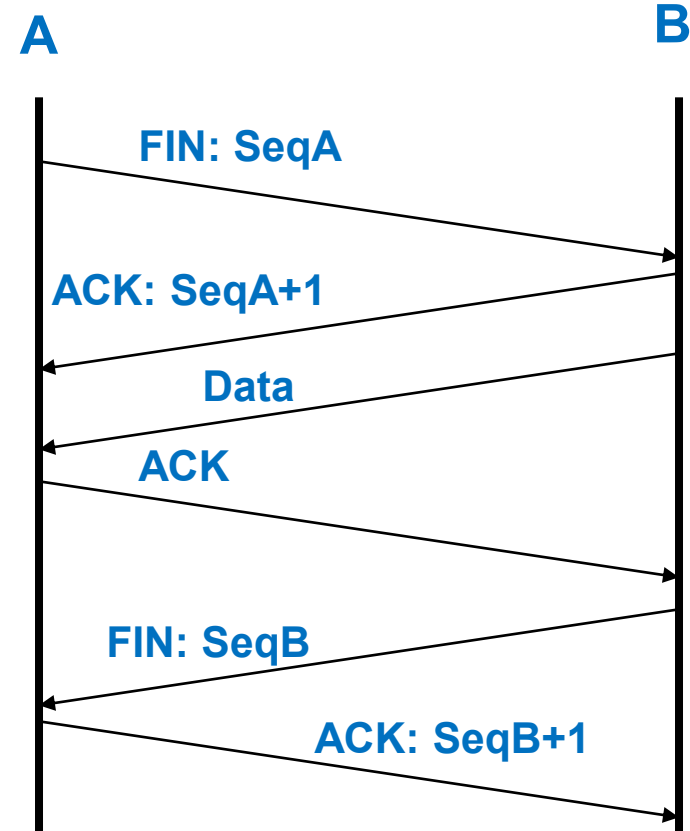
- send FIN signal
- i.e. “I’m not going to send any more data”

### “other side” can continue sending Data

- half open connection
- Has to continue to acknowledge

### Acknowledge with FIN

- acknowledge  
last sequence number + 1



# Connection – Tear Down

## TCP connection teardown: 4-way handshake

- Can be initiated by both sides
- E.g. client

### Step 1: Client sends FIN message

### Step 2: Server sends ACK message

- Half-duplex data transfer may continue
  - From Server to Client
  - Client still acknowledges data

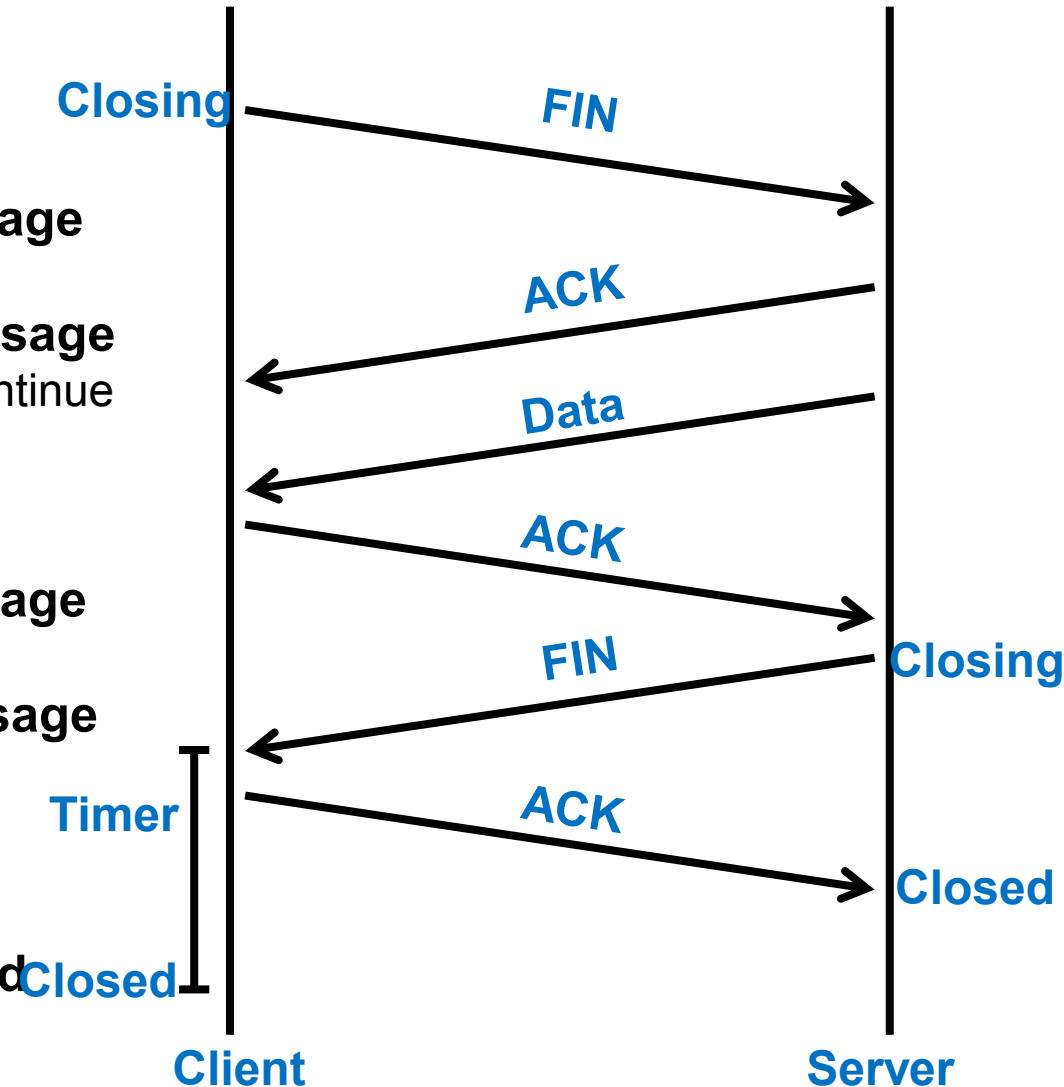
### Step 3: Server sends FIN message

### Step 4: Client sends ACK message

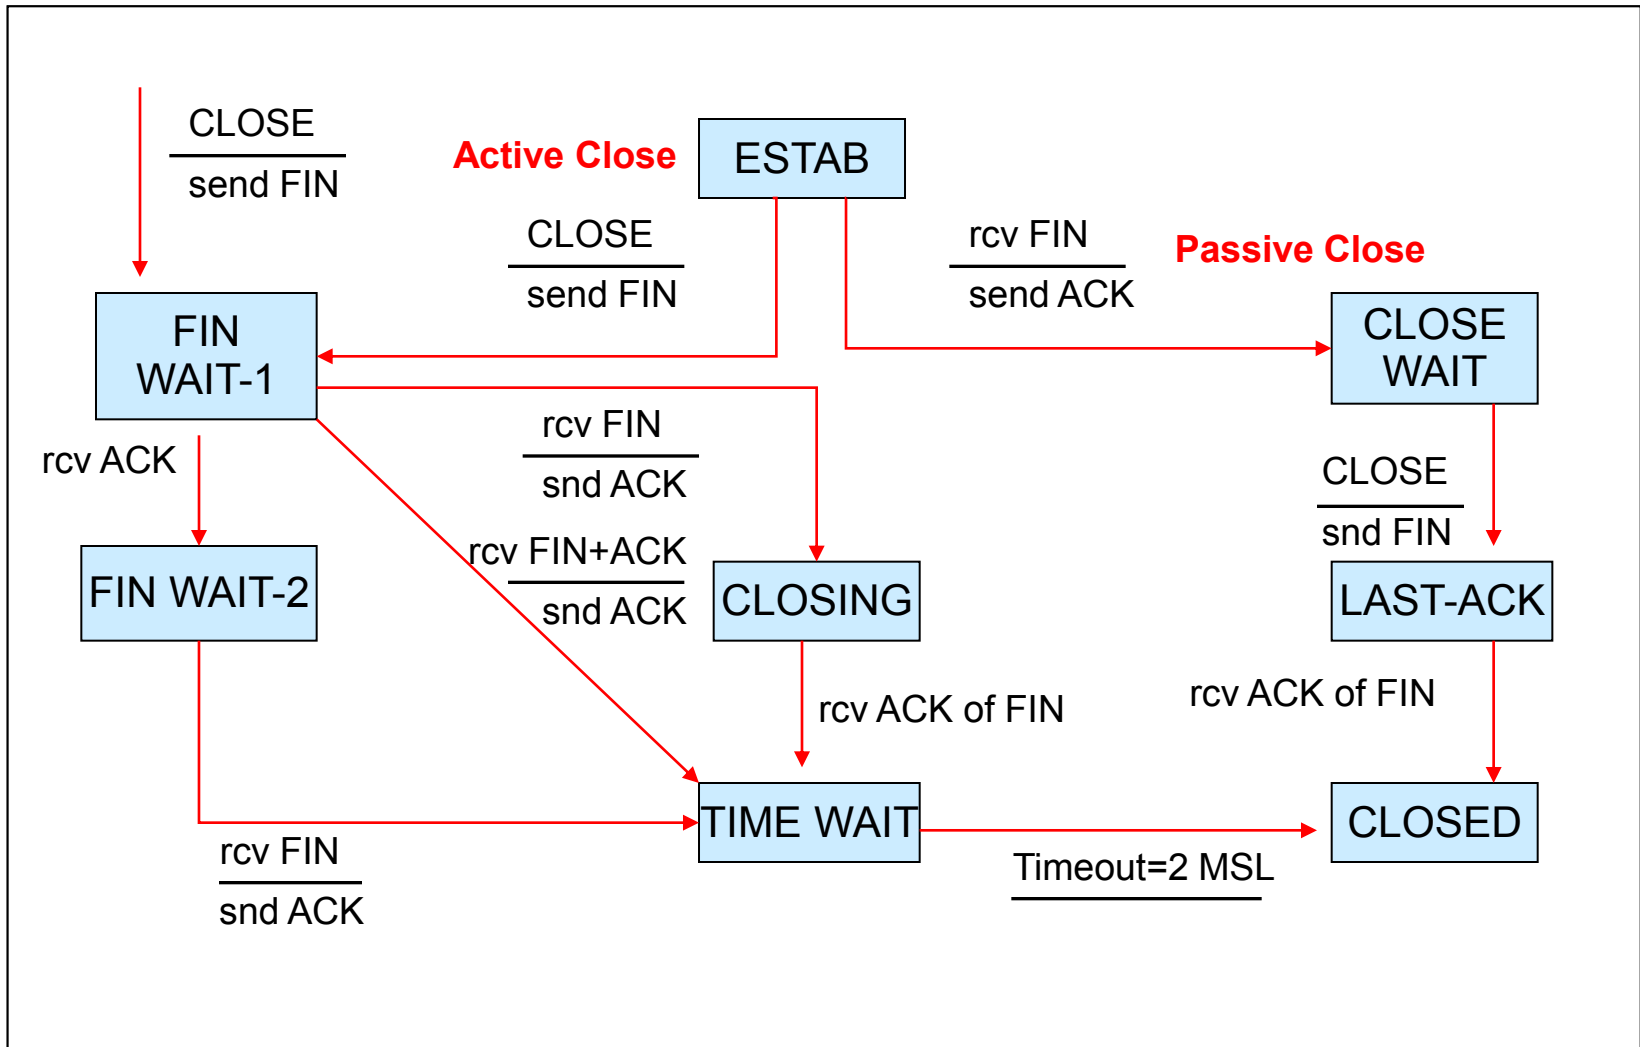
- Starts closing timer
- Closes connection at timeout
- May resend ACK message if lost

### Steps 2 and 3 may be combined

- If server has no data to send



# State Diagram: Connection Tear Down



## 4.4 Sequence Number Space

### Each byte in the byte stream is numbered


- 32 bit value
- wraps around
- initial values selected at start up time

### TCP breaks up the byte stream in packets (“segments”)

- packet size is limited to the Maximum Segment Size
- set to prevent packet fragmentation

### Each segment has a sequence number

- indicates where it fits in the byte stream



# Sequence Numbers

## 32 Bits, Unsigned

### Why so big?

- for sliding window, must have
  - $|\text{Sequence Space}| > |\text{Sending Window}| + |\text{Receiving Window}|$
  - $2^{32} > 2 * 2^{16}$ .
  - No problem
- also, want to guard against stray packets  
(stray packets – “vagabundierende Pakete”)
  - with IP, assume packets have maximum segment lifetime (MSL) of 120 sec
    - i.e. can linger in network for up to 120sec
  - sequence number would wrap around in this time at 286Mbps

### Additional reading:

- RFC 1323 / PAWS
  - protect against wrapped sequence numbers
  - TCP extension for high-speed paths

**Checksum (mostly) guarantees end-to-end data integrity**

**Sequence numbers detect packet sequencing problems:**

- duplicate: → ignore
- reordered: → reorder or drop
- lost: → retransmit


**Lost segments detected by sender**

- use time out to detect lack of acknowledgment
- need estimate of the roundtrip time to set timeout

**Retransmission requires that sender keeps a copy of the data**

- copy is discarded when ACK is received

## Bidirectional Communication (Duplex)




**Each side of a connection can send and receive (duplex)**

**i.e.**


- to maintain different sequence numbers for each direction
- single segment can contain new data for one direction, plus acknowledgement for the other
  - but some contain only data & others only acknowledgement



## 5 Flow Control – in General



**Controlled by  
Window: advertised window awnd**



**Controlled by  
Window: congestion window cwnd**

## 5.1 TCP Flow Control – in General

### Sliding window protocol

- for window size  $n$ 
  - sender can send up to  $n$  bytes without receiving an acknowledgement
- when the data is acknowledged then the window slides forward

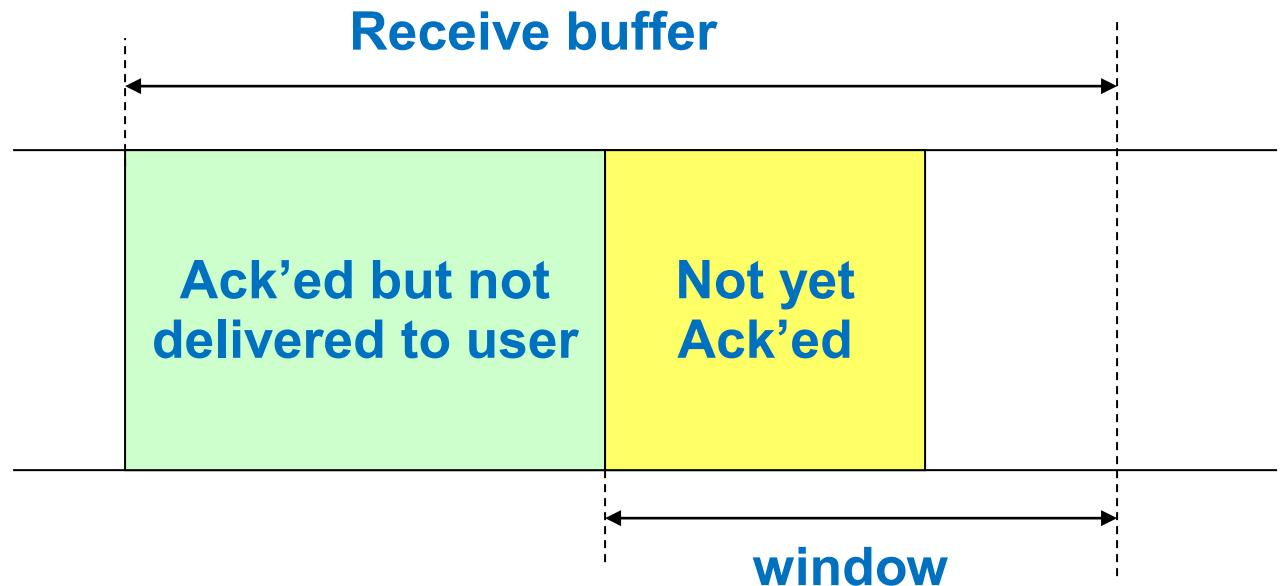
### Window size determines

- how much unacknowledged data the sender can send

**But there is one more detail ...**

## Complication

- TCP receiver can delete acknowledged data
  - only after the data has been delivered to the application
- I.e. depending on how fast the application is reading the data,
  - the receiver's window size may change!



## Receiver informs the sender

- about the current window size
  - in every packet it transmits to the sender

## Sender uses

- this current window size
- instead of a fixed value

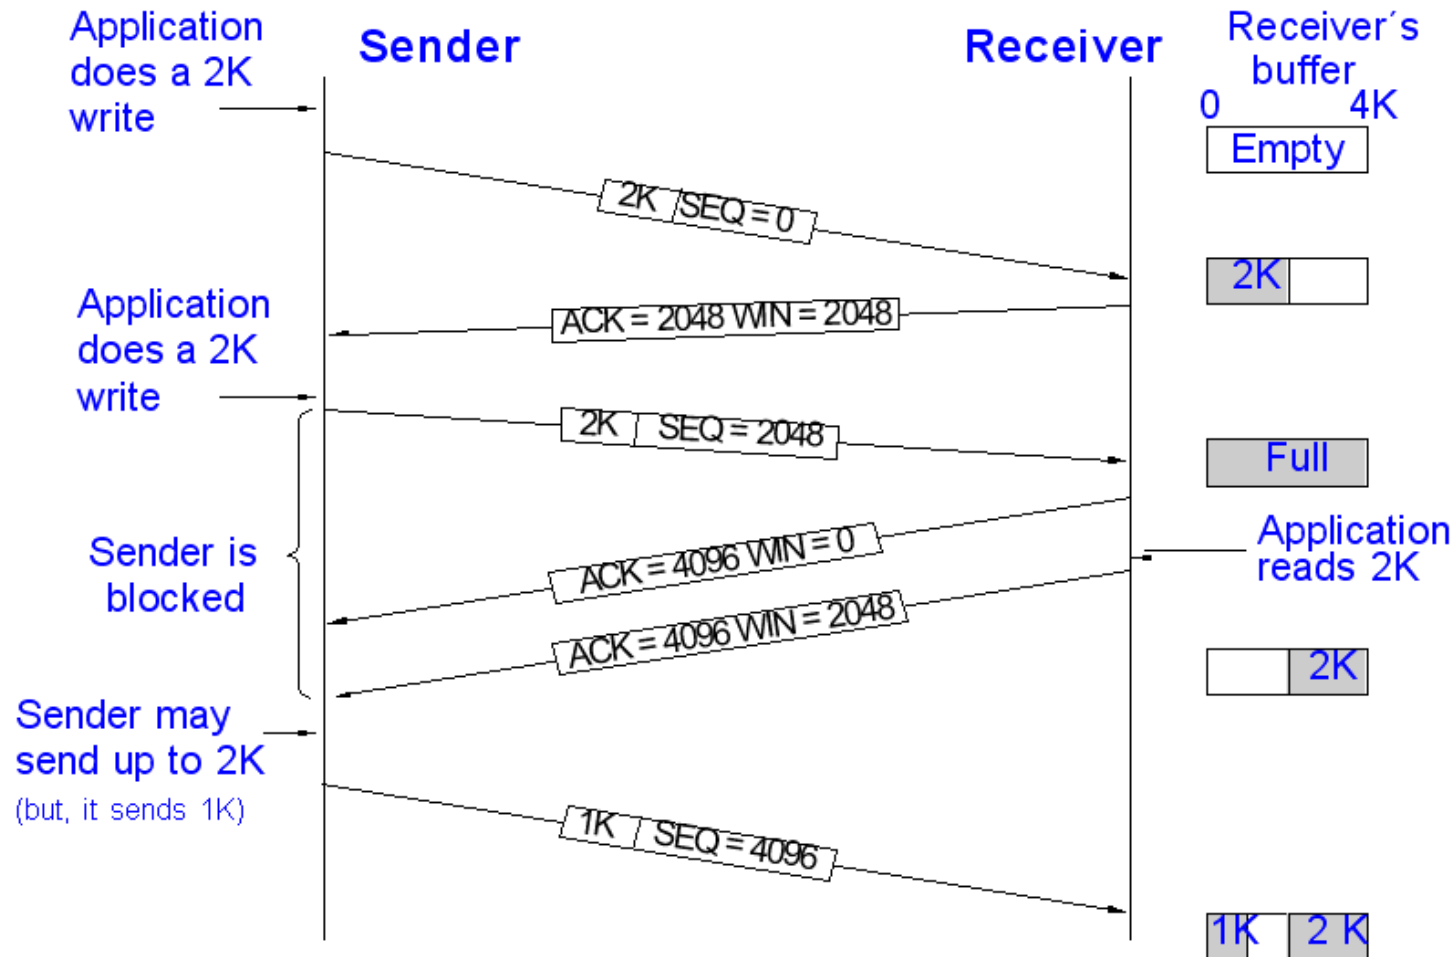
## Window size (also called ADVERTISED WINDOW)

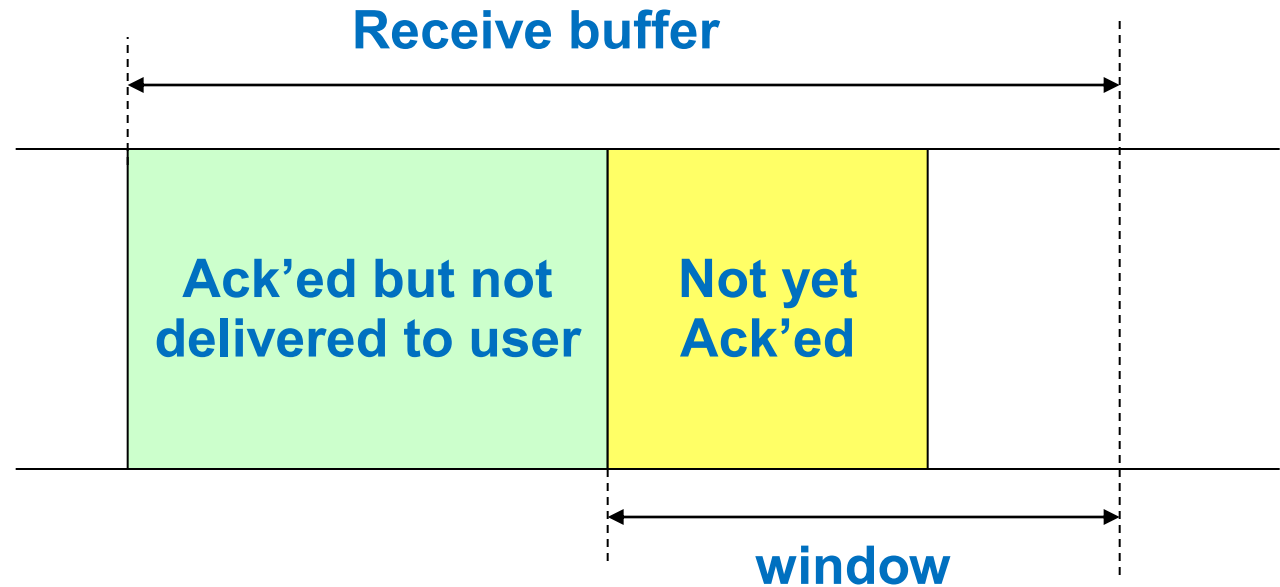
- is continuously changing

## May go to zero!

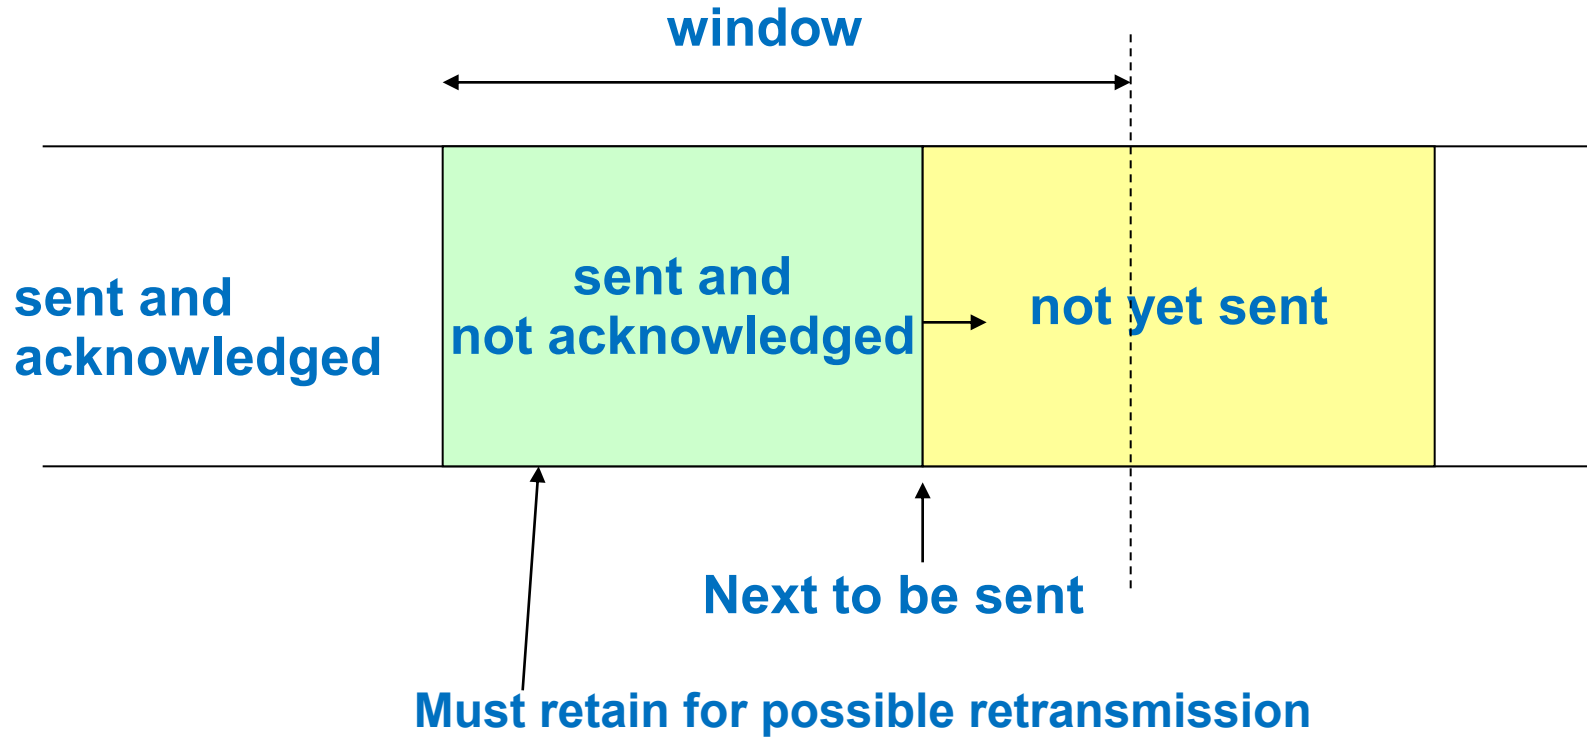
- sender not allowed to send anything!

## Solution, e.g.





# Window Flow Control: Senders Side

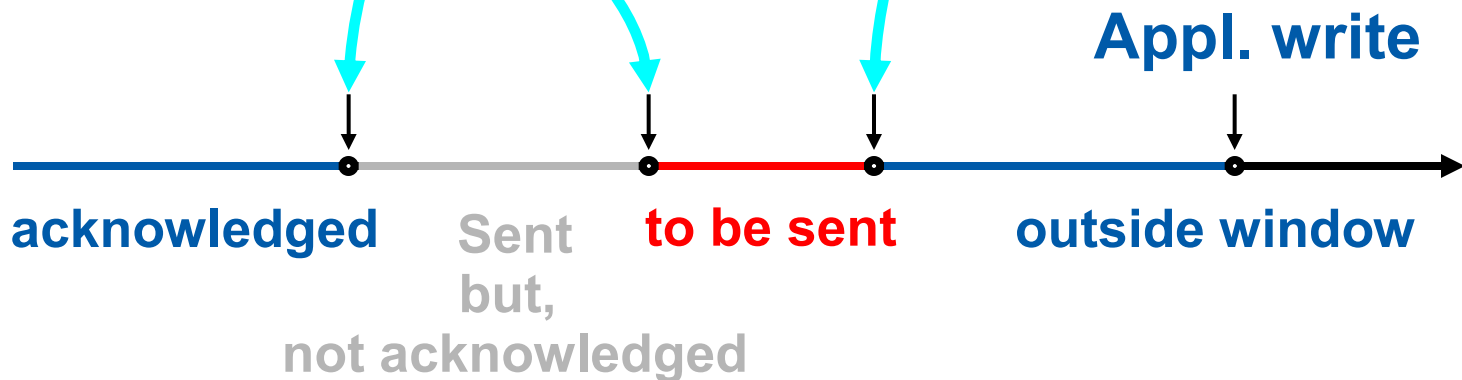


## Packet Sent

Source Port	Dest. Port
Sequence Number	
Acknowledgment	
HL/Flags	Window
Checksum	Urgent Pointer
Options..	

## Packet Received

Source Port	Dest. Port
Sequence Number	
Acknowledgment	
HL/Flags	Window
Checksum	Urgent Pointer
Options..	





## 5.2 TCP Window Flow Control - Special Cases

### Optimization for low throughput rate

#### Problem:

**Telnet (and ssh) connection to interactive editor reacting on every keystroke**

- 1 character typed requires up to 162 bytes
  - data:
    - 20 bytes TCP header,
    - 20 bytes IP header,
    - 1 byte payload
  - ACK:
    - 20 bytes TCP header, 20 bytes IP header
  - editor echoes character:
    - 20 bytes TCP header, 20 bytes IP header, 1 byte payload
  - ACK:
    - 20 bytes TCP header, 20 bytes IP header

#### commonly used solution

- delay acknowledgements and window update by 500 ms
- hoping for more data to come

#### Nagle's algorithm, 1984

- send first byte immediately
- keep on buffering bytes until first byte has been acknowledged
- then send all buffered characters in one TCP segment and start buffering again

#### comment

- effect e.g. X-Windows: jerky pointer movements

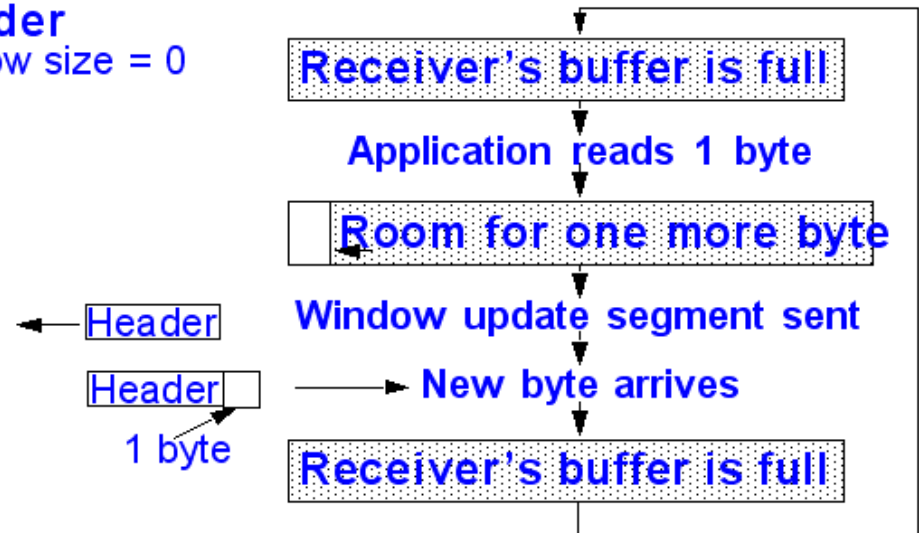
# TCP Window Flow Control: Special Cases

## Silly window syndrome (Clark, 1982)

### Problem

- data on sending side arrives in large blocks
- but receiving side reads data only one byte at a time

Sender  
window size = 0



### Clark's solution:

- prevent receiver from sending window update for 1 byte
- certain amount of space must be available in order to send window update
- $\min(X, Y)$ :
  - $X = \text{max. segment size (MSS)}$ ,
  - $Y = \text{buffer}/2$

### Bidirectional Communication

- each side acts as sender & receiver
- every message
  - contains acknowledgement of received sequence
    - even if no new data has been received
  - advertises window size
    - size of its receiving window
  - contains sent sequence number
    - even if no new data is being sent


## When does a sender actually sends a message?

- when sending buffer contains at least max. segment size (- header sizes) bytes
- when application tells it to
  - set PUSH flag for last segment sent
- when timer expires


## 6 Round Trip Time



**Ideal situation:**



**Example of real situation**





### TCP Must Operate Over Any Internet Path

- Retransmission time-out should be set based on round-trip delay
- But round-trip delay different for each path!

➔ **Must estimate RTT dynamically**

## 6.1 Setting Retransmission Timeout (RTO)



### Retransmission Timeout (RTO)


- time between sending & resending segment

### Challenge

- too long:
  - Add latency to communication when packets dropped
- too short:
  - Send too many duplicate packets
- general principle:
  - $RTO > 1 \text{ Round Trip Time (RTT)}$

# Round-trip Time Estimation

Every Data/Ack pair gives new RTT estimate



Can get lots of short-term fluctuations

## 6.2 Original TCP Round-trip Estimator (Jacobson, 1988)

### Round trip times estimated as a moving average:

- $\text{new\_RTT} = \alpha * (\text{old\_RTT}) + (1 - \alpha) * (\text{new\_sample})$

### Smoothing factor

- recommended value for  $\alpha$ :
  - 7/8 (0.8 - 0.9)
  - 0.875 for most TCP's

### Retransmit timer set to

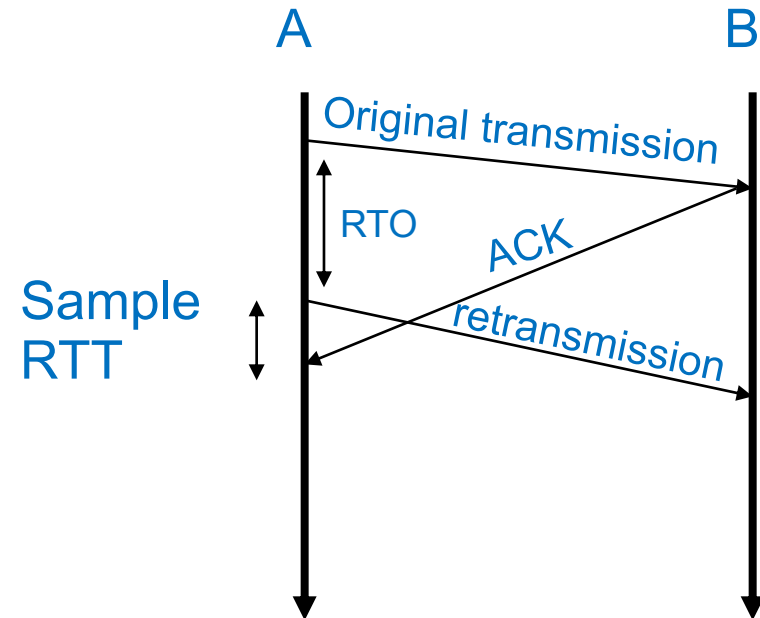

- $\text{RTO} = \beta \text{ new\_RTT}$ ,
  - where  $\beta = 2$
  - want to be somewhat conservative about retransmitting

### Problem

- static  $\beta$  not able to adapt to high variation in observed RTTs during high load conditions
- solution used today
  - estimate both RTT and variance in RTT
  - use the estimated variance in place of constant  $\beta$




## 6.3 RTT Sample Ambiguity



### Solution

- Karn's algorithm
  - ignore sample for segment that has been retransmitted
  - timer backoff strategy
    - retransmission timer ( $\text{new\_RTO} = \gamma \cdot \text{old\_RTO}$ )
    - typical  $\gamma = 2$
    - resume normal computation when ACK received for non-retransmitted segment

## 7 Congestion Control



**Controlled by  
Window: advertised window  $awnd$**

**Controlled by  
Window: congestion window  $cwnd$**


## 7.1 Congestion at Transport Layer - in General

**Load placed on the network  
is higher than  
the capacity of the network**

- not surprising: independent senders place load on network

**Results in packet loss: routers have no choice**

- can only buffer finite amount of data
- end-to-end protocol will typically react, e.g. TCP





# Congestion – the Problem

**Wasted bandwidth: retransmission of dropped packets**


**Poor user service: unpredictable delay, low user goodput**

**Increased load can even result in lower network goodput**

- switched nets:
  - packet losses create lots of retransmissions
- broadcast Ethernet:
  - high demand → many collisions



# Sending Rate of Sliding Window Protocol



## Suppose

- Sender A uses a sliding window protocol to transmit a large data file to B
- Window size = 64KB (i.e.  $2^{16}$  Bytes)
- Network round-trip delay is 1 second

## Which is the expected sending rate?

- $\rightarrow 64\text{KB in } 1 \text{ second } (2^{16} \text{ Bytes} / 1 \text{ second} = 65536\text{Byte/second}$   
 $= 524288\text{bits/second} \approx 524.3\text{kbps})$

## What if a

- network link is (only) 524.3kbps
- but there are 1000 people who are transferring files over that link using the sliding window protocol?

**$\rightarrow$  Packet losses, timeouts, retransmissions, more packet losses...**

- nothing useful gets through, collapse due to congestion!

## 7.2 Limits of TCP Window Flow Control




### Packet Sent

Source Port	Dest. Port
Sequence Number	
Acknowledgment	
HL/Flags	Window
Checksum	Urgent Pointer
Options..	

### Packet Received

Source Port	Dest. Port
Sequence Number	
Acknowledgment	
HL/Flags	Window
Checksum	Urgent Pointer
Options..	



# **TCP Flow Control alone is not enough**

**We have talked about how TCP's advertised window  
which is used for flow control**

- to prevent the sender sending faster than the receiver can handle

**If the receiver is sufficiently fast,**

**→ the advertised window will be maximized at all time**

**→ leads to collapse due to congestion**

- as in the previous example if
  - there are too many senders or
  - the network is too slow

**Key 1: Window size determines sending rate**

**Key 2: Window size must be dynamically adjusted  
to prevent collapse due to congestion**


# How Fast to Send?

## Send too slow: link sits idle

- wastes time

## Send too fast: link is kept busy but....

- queue builds up in router buffer (delay)
- overflow buffers in routers (loss)
- many retransmissions, many losses
- network goodput goes down







## We ignore

- internal structure of the network and
- model network as having a single bottleneck link


## 7.3 Three Congestion Control Problems

---

**Adjusting to bottleneck bandwidth**


**Adjusting to variations in bandwidth**

**Sharing bandwidth between flows**




## Adjust rate to match bottleneck bandwidth

- without any a priori knowledge
- could be gigabit link, could be a modem



**Adjust rate to match instantaneous bandwidth**

**Bottleneck can change because of a routing change**




## (two) issues:

- Adjust total sending rate to match bottleneck bandwidth
- Allocation of bandwidth between flows


# Multiple Flows and how TCP shares capacity



individual  
flow rates

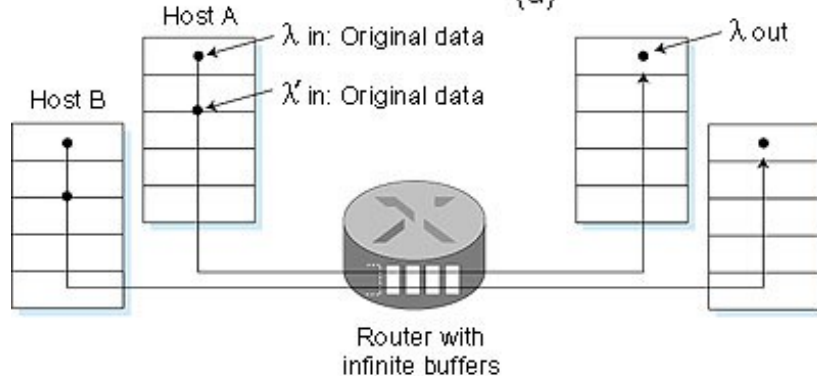
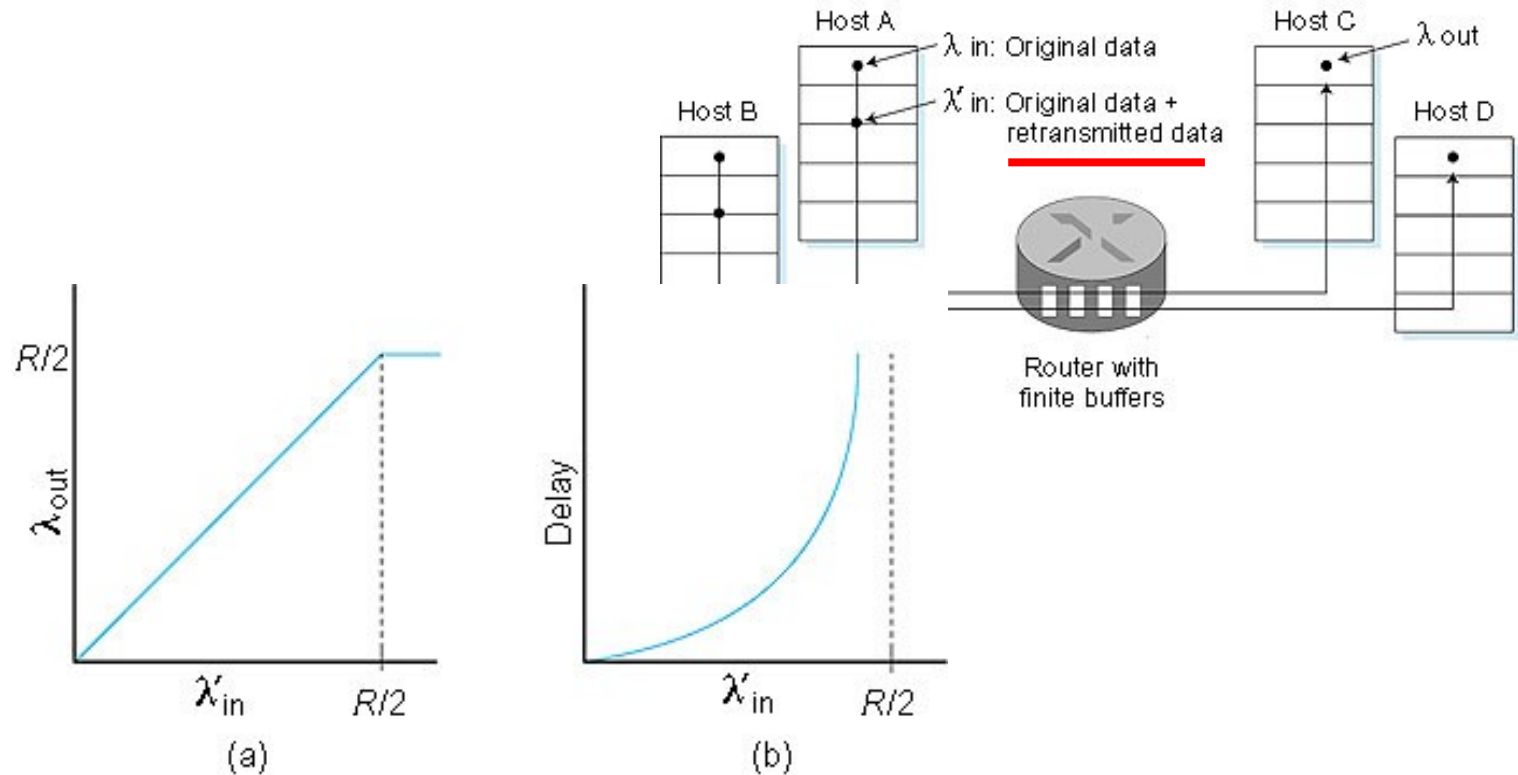


aggregate  
flow rate



available  
capacity

# Why is Congestion Bad - Revised?



### Send without care

- many packet drops
- could cause collapse due to congestion

### Reservations

- pre-arrange bandwidth allocations
- requires negotiation before sending packets

### Pricing


- don't drop packets for the high-bidders
- requires payment model

### Dynamic Adjustment (TCP)

- technique
  - every sender probes network to test level of congestion
  - speed up when no congestion
  - slow down when congestion
- evaluation
  - suboptimal, messy dynamics, simple to implement
  - distributed coordination problem



## 7.5 TCP Congestion Control



**Controlled by  
Window: advertised window  $awnd$**

**Controlled by  
Window: congestion window  $cwnd$**

## **TCP connection has window**

- controls number of unacknowledged packets
- Sending rate:  $\sim \text{Window} / \text{RTT}$
- Vary window size to control sending rate

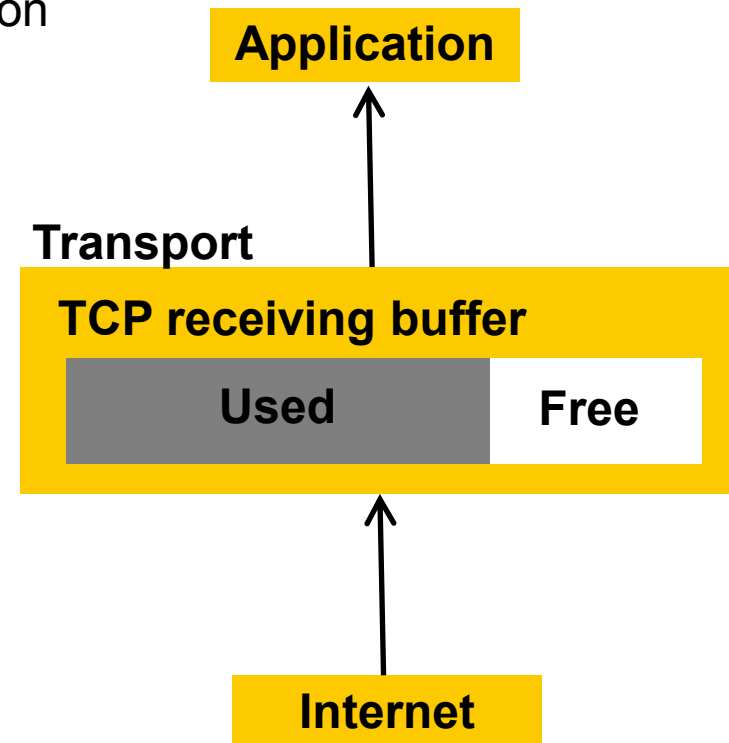
## **Introduce a new parameter called congestion window (cwnd) at the sender**

- congestion control is mainly a sender-side operation

## 7.6 Flow Control & Congestion Control

### Flow control required to avoid flooding receiver

- Receiver allocates buffer space for receiving messages
  - Buffer becomes available when
    - Data is acknowledged and
    - Data is read from buffer by receiving application
  - But: application may be slower than network
    - E.g. high load situations
- Receiving buffer may become full
- How to react?



## 7.7 Congestion Window (cwnd)

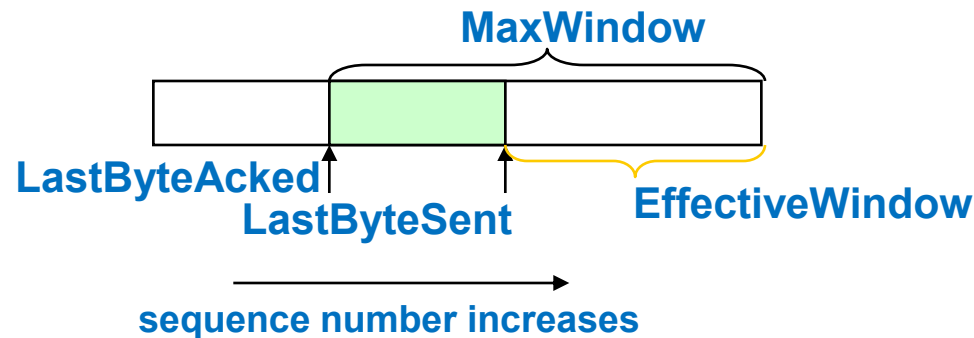
Limits how much data can be in transit

Implemented as # of bytes

Described as # packets in this lecture

$$\text{MaxWindow} = \min(\text{cwnd}, \text{awnd})$$

$$\text{EffectiveWindow} = \text{MaxWindow} - (\text{LastByteSent} - \text{LastByteAcked})$$



# Two Basic Components of Congestion Window (cwnd)



## Detecting congestion

## Rate adjustment algorithm (change cwnd size)

- depends on congestion or not

## 7.8 Basic Component 1: Detecting Congestion

### Packet dropping is best indication for congestion

- delay-based methods are hard and risky

### How do you detect packet drops? ... ACKs

- TCP uses ACKs to signal receipt of data
- ACK denotes last contiguous byte received
  - actually, ACKs indicate next segment expected

### Two signs of packet drops

- no ACK after certain time interval: time-out
- several duplicate ACKs (ignore for now)

### May not work well for wireless networks, why?





- Fading, echoes, .. → sporadic losses ... → underutilized links with TCP

# Sliding (Congestion) Window

## Sliding window:

**each ACK = permission to send a new packet**


- example cwnd = 3



# Self-clocking

If we have a large window,  
ACKs “self-clock” the data to the rate of the bottleneck link

Observe: received ACK spacing  $\cong$  bottleneck bandwidth





## 7.9 Basic Component 2: Rate Adjustment

### Algorithm

- upon receipt of ACK (of new data):
  - ➔ increase rate
- data successfully delivered, perhaps can send faster
- upon detection of loss:
  - ➔ decrease rate

### But which increase/decrease functions should we use?


- .... depends on what problem we are solving

## 8 Fairness of TCP Congestion Control





### Two competing sessions:

- Additive increase (AI) gives slope of 1,
  - as throughput increases
- multiplicative decrease (MD) decreases throughput proportionally



## 8.1 Additive Increase Multiplicative Decrease (AIMD)



### 1: x and y

- have certain throughput
- No loss  $\rightarrow$  x, y cwnd by +1 every RTT

### 2: x and y



- loss  $\rightarrow$  decrease by factor 2

### 3: x and y

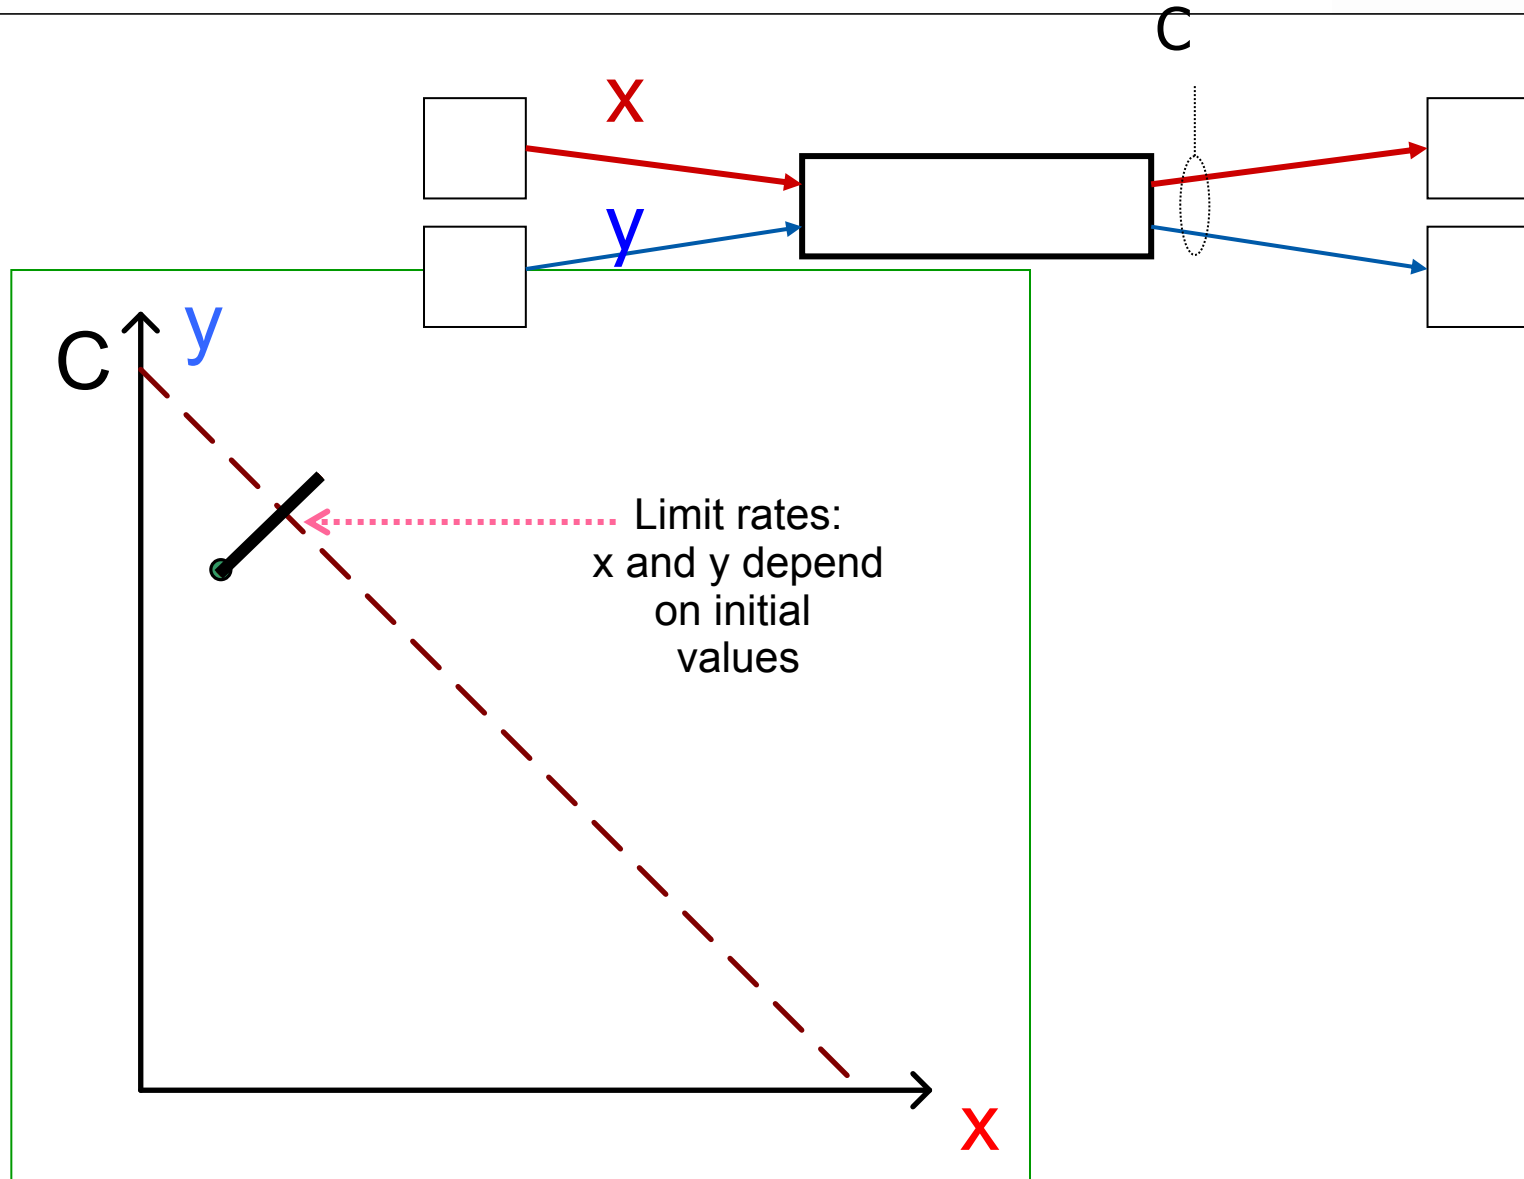
- ...

Limit rates:  
 $x = y$



# Additive Increase Multiplicative Decrease: Sharing Dynamics



## 8.2 Additive Increase Additive Decrease (AIAD)



# Additive Increase Additive Decrease: Sharing Dynamics



## 9 Phases of TCP Congestion Control

### So far

- sliding window + self-clocking of ACKs

### How to know

- the best congestion window  $cwnd$  ?
- and best transmission rate?

### → Adapting Congestion Window $cwnd$

# Phases of Congestion Control

## Phase 1: Slow start (getting to equilibrium)

- want to find this extremely fast and wasting time

## Phase 2: Congestion Avoidance

- additive increase
  - gradually probing for additional bandwidth
- multiplicative decrease
  - decreasing cwnd upon loss/timeout



## 9.1 Initialization

### Congestion Window (cwnd)

- Initial value is 1 MSS (=maximum segment size)
  - counted as bytes

### Slow-start threshold Value (ss\_thresh)

- Initial value is advertised window size

#### i.e. phase 1:

- slow start ( $\text{cwnd} < \text{ss\_thresh}$ )

#### i.e. phase 2:

- congestion avoidance ( $\text{cwnd} \geq \text{ss\_thresh}$ )

## 9.2 Phase 1: TCP Slow Start

### Goal:

- to discover roughly the proper sending rate quickly

### Whenever

- starting traffic on a new connection, or whenever
- increasing traffic after congestion was experienced:

➔ **initialize cwnd =1**

**each time a segment is acknowledged,**

➔ **increment cwnd by one (cwnd++)**

### Continue until

- reach ss\_thresh
- packet loss

# Slow Start Illustration

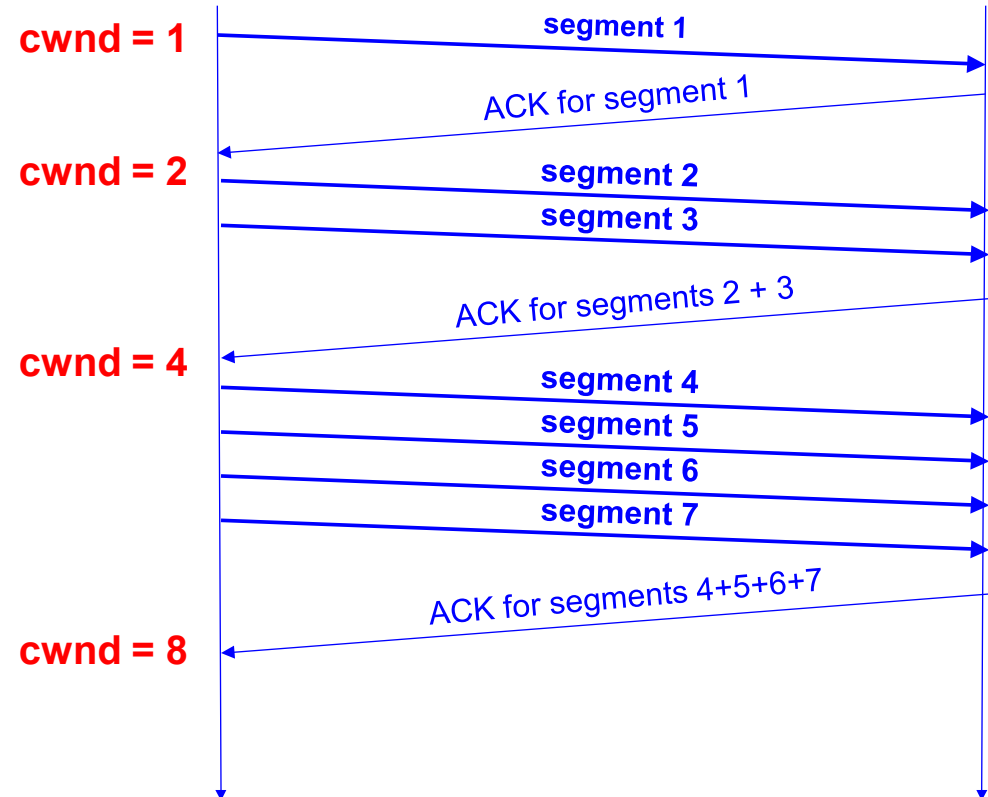
congestion window size grows very rapidly

TCP slows down the increase of cwnd

- when  $\text{cwnd} \geq \text{ss\_thresh}$

Observe:

- each ACK generates 2 packets
- slow start increases rate exponentially
  - (doubled every RTT)



## 9.3 Phase 2: Congestion Avoidance (After Slow Start)



### Slow Start

- roughly figures out rate at which the network starts to get congested

### Congestion Avoidance

- continues to react to network condition
  - probes for more bandwidth
    - increase cwnd
      - if more bandwidth available
  - if congestion detected,
    - aggressive cut back cwnd

## Phase 2: Congestion Avoidance (Additive Increase)

**After exiting slow start**

**→ slowly increase cwnd  
to probe for additional available bandwidth**

- competing flows may end transmission
- may have been “unlucky” with an early drop

**If  $cwnd > ss\_thresh$  then**

- each time a segment is acknowledged
  - increment cwnd by  $1/cwnd$  ( $cwnd += 1/cwnd$ ).

**i.e. cwnd is increased by 1**



- only if all segments have been acknowledged

**→ increases by 1 per RTT, vs. doubling per RTT**

## 9.4 Example of Slow Start + Congestion Avoidance



Assume that  $ss\_thresh = 8$



# Detecting Congestion via Timeout

---

## **If there is a packet loss**

- the ACK for that packet will not be received

## **packet will eventually time out**

- no ACK is seen as a sign of congestion

# Congestion Avoidance: Multiplicative Decrease

**Timeout = congestion**

**Each time when congestion occurs,**

- `ss_thresh` is set to 50%  
of the current size of the congestion window:
  - $\text{ss\_thresh} = \text{cwnd} / 2$
- `cwnd` is reset to one:
  - $\text{cwnd} = 1$
- and  
slow-start is entered



## 9.5 TCP illustrated

