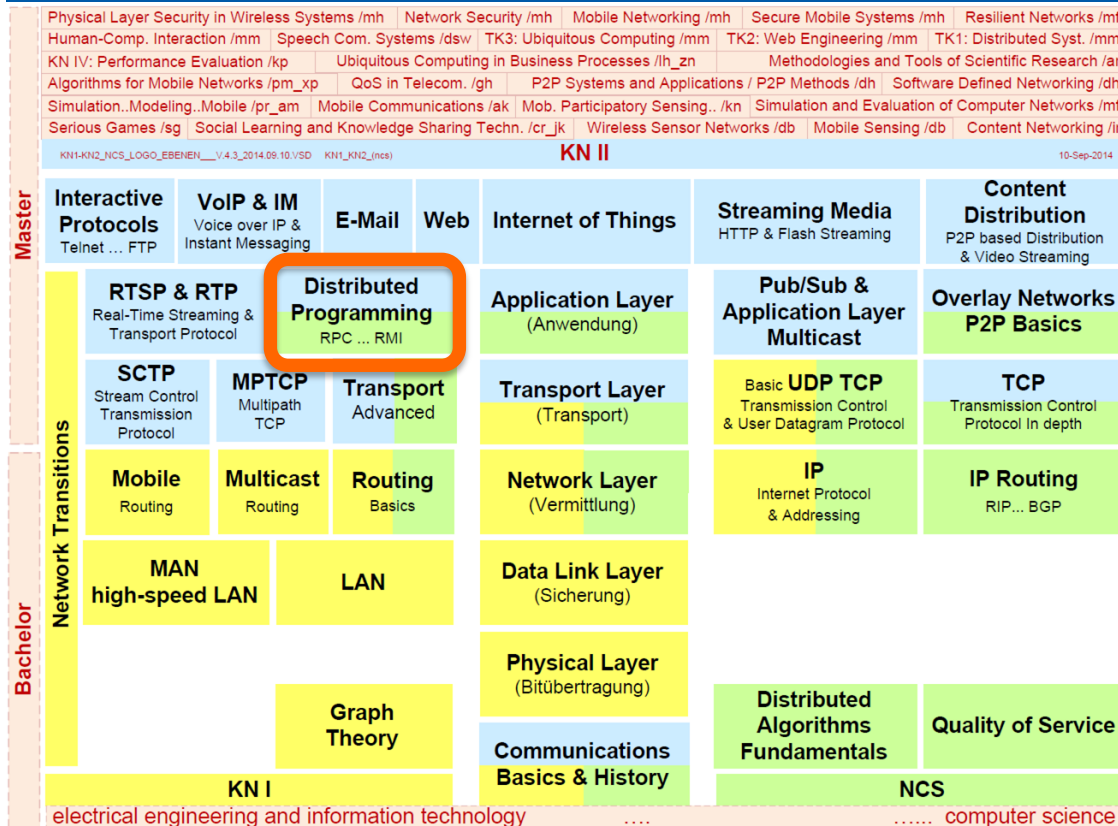


Communication Networks II

Distributed Programming



TECHNISCHE
UNIVERSITÄT
DARMSTADT



10. September 2014

Prof. Dr.-Ing. Ralf Steinmetz
KOM - Multimedia Communications Lab

1 Distributed Systems

- 1.1 Distributed Systems: Introduction**
- 1.2 Transparency Concepts**
- 1.3 Programming Abstractions**

2 System Architectures

- 2.1 Client/Server Model**
- 2.2 Object-oriented Model**
- 2.3 Component-based Model**
- 2.4 Service-oriented Model**
- 2.5 Multi-Tier Architectures**
- 2.6 Further System Architectures**

3 Interprocess Communication

- 3.1 Message Destinations**
- 3.2 Sockets**
- 3.3 UDP/IP Sockets**
- 3.4 TCP/IP Sockets**

4 Remote Procedure Call (RPC)

- 4.1 RPC: Introduction**
- 4.2 RPC: Principle**
- 4.3 RPC: Marshalling and Presentation**
- 4.4 RPC: Binding**
- 4.5 RPC: Protocols**
- 4.6 RPC: Errors**
- 4.7 RPC: Failure Semantics**

5 Web Services

- 5.1 Web Service Model**
- 5.2 SOAP – Simple Object Access Protocol**
- 5.3 WSDL – Web Service Description Language**
- 5.4 UDDI – Universal Description, Discovery and Integration**
- 5.5 Web 2.0 – AJAX**

6 Event Based Systems (Event Routing Systems)

- 6.1 Basic Principles**
- 6.2 Classification**
- 6.3 Subscription Semantics**

7 Further Technologies

- 7.1 Remote Method Invocation - RMI**
- 7.2 Common Object Request Broker Architecture - CORBA**
- 7.3 Enterprise Java Beans - EJB**

Physical Layer Security in Wireless Systems /mh	Network Security /mh	Mobile Networking /mh	Secure Mobile Systems /mh	Resilient Networks /mf
Human-Comp. Interaction /mm	Speech Com. Systems /dsw	TK3: Ubiquitous Computing /mm	TK2: Web Engineering /mm	TK1: Distributed Syst. /mm
KN IV: Performance Evaluation /kp	Ubiquitous Computing in Business Processes /lh_zn		Methodologies and Tools of Scientific Research /ar	
Algorithms for Mobile Networks /pm_xp	QoS in Telecom. /gh	P2P Systems and Applications / P2P Methods /dh	Software Defined Networking /dh	
Simulation..Modeling..Mobile /pr_am	Mobile Communications /ak	Mob. Participatory Sensing.. /kn	Simulation and Evaluation of Computer Networks /mf	
Serious Games /sg	Social Learning and Knowledge Sharing Techn. /cr_jk	Wireless Sensor Networks /db	Mobile Sensing /db	Content Networking /ir

KN1-KN2_NCS_LOGO_EBENEN___V.4.3_2014.09.10.VSD KN1_KN2_(ncs)

KN II

10-Sep-2014

Master

Interactive Protocols Telnet ... FTP	VoIP & IM Voice over IP & Instant Messaging	E-Mail	Web	Internet of Things	Streaming Media HTTP & Flash Streaming	Content Distribution P2P based Distribution & Video Streaming
--	---	---------------	------------	---------------------------	--	---

RTSP & RTP Real-Time Streaming & Transport Protocol	Distributed Programming RPC ... RMI	Application Layer (Anwendung)	Pub/Sub & Application Layer Multicast	Overlay Networks P2P Basics
---	---	---	--	--

SCTP Stream Control Transmission Protocol	MPTCP Multipath TCP	Transport Advanced	Transport Layer (Transport)	Basic UDP TCP Transmission Control & User Datagram Protocol	TCP Transmission Control Protocol In depth
---	-------------------------------	---------------------------	---------------------------------------	---	--

Mobile Routing	Multicast Routing	Routing Basics	Network Layer (Vermittlung)	IP Internet Protocol & Addressing	IP Routing RIP... BGP
-----------------------	--------------------------	-----------------------	---------------------------------------	---	---------------------------------

MAN high-speed LAN	LAN	Data Link Layer (Sicherung)
---------------------------	------------	---------------------------------------

Graph Theory	Physical Layer (Bitübertragung)
---------------------	---

Communications Basics & History	Distributed Algorithms Fundamentals	Quality of Service
--	--	---------------------------

KN I	NCS
-------------	------------

electrical engineering and information technology

..... computer science

Motivation for building distributed systems:

Scalability

- A single host can handle only a limited load
- If the load gets higher → split the system in components and distribute them

Openness

- Communication to 3rd party systems is required
- E.g. online shop with Paypal integration

Heterogeneity

- Systems grow over time
- Old hardware and software components get unavailable and must be replaced with new components

Fault tolerance

- The question is not if failures happen – the question is when...
- The system should be available even if some components fail

1.1 Distributed Systems: Introduction

Our Understanding of “Distributed System”

**“A distributed system is one in which the failure of a computer you didn’t even know existed can render your own computer unusable”.
(Leslie Lamport, 1987)**

Properties

- Composed of several distributed, autonomous components
- Components are linked by a computer network
- Typically, components do not use same storage space
- Components have to communicate by exchanging messages
- Components cooperate to achieve a common goal
- Largest distributed system: World Wide Web

Goals

- Parallelization (e.g., parallel processing of customer data)
- Load-balancing (e.g., subdivide workload into different tasks)
- Fault tolerance, resilience, availability (e.g., replicate critical components)
- Scalability (e.g., increase capacity by adding additional components)

Basic Concepts

- Name and Directory Services
- Distributed Transactions
- Security (Authentication, Authorization, Encryption)
- Transparency (specified in reference models: ANSA, ISO)
 - Location
 - Access
 - Failure
 - Concurrency
 - Replication
 - Mobility [ANSA] / Migration [ISO]
 - Scalability
 - Performance

Typical Approach

- Middleware solutions, which implement the concepts above

1.2 Transparency Concepts

Access Transparency

- Remote methods should look like local methods from the clients point of view
- The interface for local and remote method calls should be equal

Location Transparency

- A location transparent name contains no information about the named object's physical location
- E.g. visiting www.tu-darmstadt.de via Web-Browser is location transparent
 - (human) visitor does not need to know the IP address or physical location of the hosting server

Mobility / Migration Transparency

- An object could be moved from one host to another without affecting other components

Replication Transparency

- Other components(=users) do not need to care whether an accessed object is the original or a replica

Transparency Concepts

Concurrency Transparency

- Users and Applications should be able to access shared data or objects without interference between each other

Failure Transparency

- Enables the concealment of faults, allowing user and application programs to complete their tasks despite the failure of hardware or software components
- This rather abstract criterion is supported by the replication transparency and by the location transparency

Performance Transparency

- Allows the system to be reconfigured to improve the performance as the load varies
- E.g. any component should respond as fast as possible from the clients point of view

Scaling Transparency

- A system should be able to grow without affecting application algorithms
- E.g. the Internet scales (nearly) transparent because the addition of new hosts does not affect existing nodes

1.3 Programming Abstractions

Distributed operating system

- Support for distributed programming is part of operating system
 - Pro: Quite general solution
 - Con: Needs wide-scale adoption of the same system
Large systems always heterogeneous

Distributed database approach

- Same as above, except OS replaced by a database system
 - Pro: Allows for all features of databases (semantics, etc.)
 - Con: Independent applications with shared database
 - Con: Many distributed algorithms hard to realize in this case

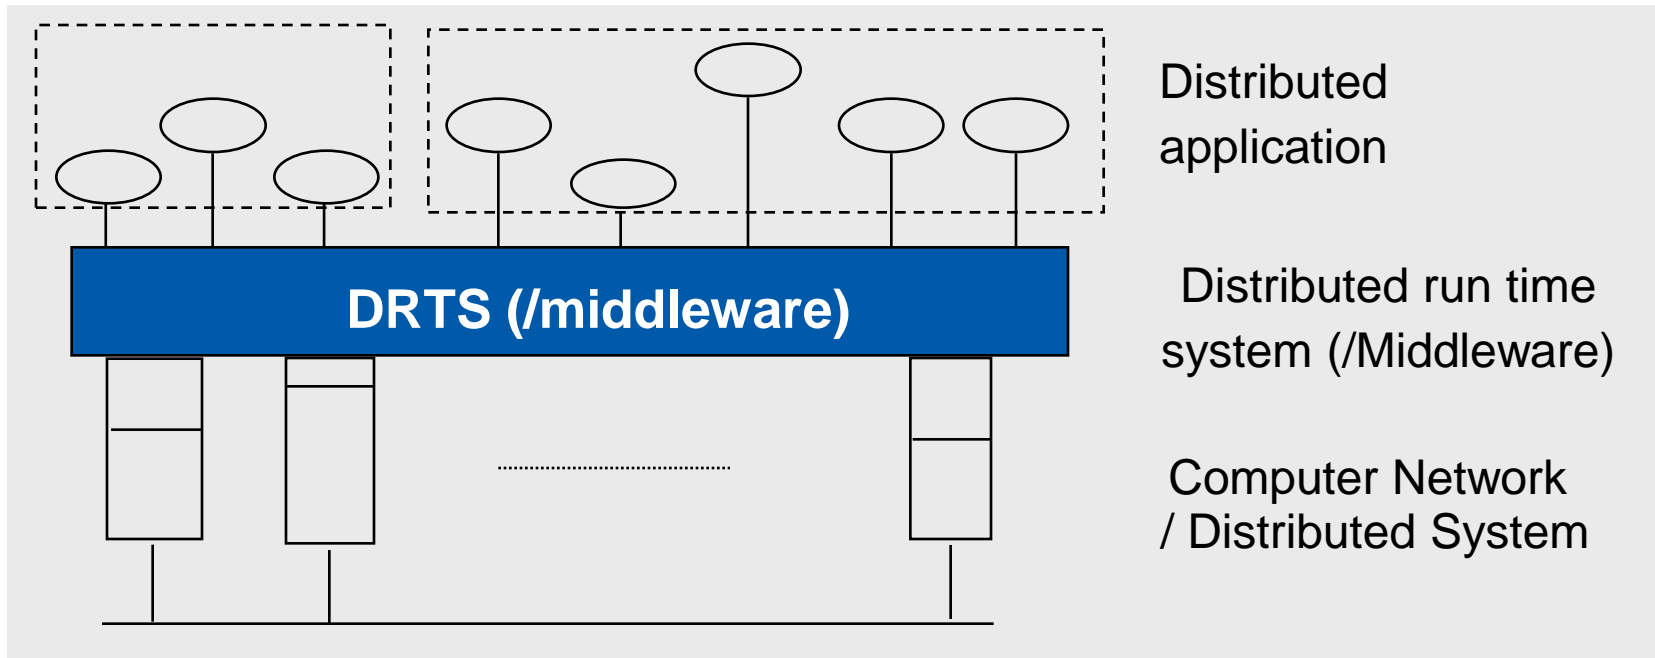
Protocols

- Standardized APIs for connecting to servers (e.g., HTTP)
 - Pro: Open, global
 - Con: Limited to standard functionalities

Our Model: Programming Language

Arbitrary, heterogeneous OSES (and databases)

Distributed programming language



Modern approach:

- Sequential programming language + extensions + middleware
- Compiler will not see the distributed program

Architectural Model defines

- Placement of components across a computer network
- Way of interaction between components

Classification of Components

- Necessary to
 - Define responsibilities
 - Assess workloads and impact of failures
- Component types: server, client, peer

Interaction Models

- Synchronous communication
 - Transmitted messages are received within known period of time
 - Requires synchronization between sender and receiver
- Asynchronous communication
 - Requires no assumptions of execution time intervals

2.1 Client/Server Model

Components

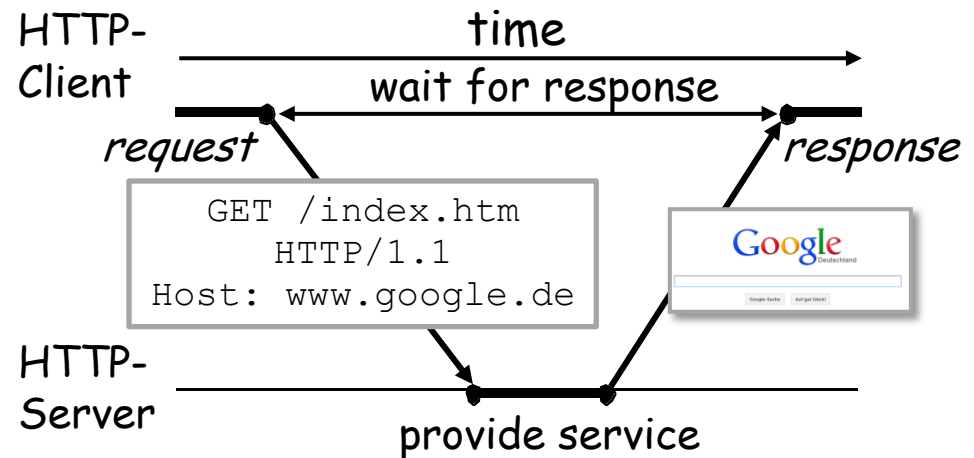
- Client: requests a service
- Server: provides a service

Typical Interaction

- Synchronous communication
- Client:
 - Sends request
 - Waits for server response
- Server:
 - Receives request
 - Processes request
 - Sends response

Examples:

- Web Client/Server
- Remote Procedure Call (RPC)




2.2 Object-oriented Model

Components

- Based on client/server model
- But objects provide/request services

Typical Interaction

- Objects exchange messages
- Local and remote object calls
- Objects as input parameters
 - By Value
 - By Reference



Examples:

- Remote Method Invocation (RMI)
- Common Object Request Broker Architecture (CORBA)

2.3 Component-based Model

Components


- Idea: separation of concerns
- Decompose application functionality into reusable, functional components
- Add properties for distributed systems at deployment time, such as
 - Persistence
 - Security
- Specific runtime environment realizes deployment properties

Typical Interaction

- Components expose well-defined interfaces
- Depend on a specific runtime environment

Examples:

- Enterprise JavaBeans
- CORBA Component Model (CCM)



2.4 Service-oriented Model

Components


- Coarse-granular services
- Encapsulate a specific functionality

Typical Interaction

- Services expose well-defined interfaces
- Loose coupling of services to build complex workflows
- Interoperability across platforms and organizations

Examples:


- Web Services
- Service-oriented Architectures (SOA)



2.5 Multi-Tier Architectures

Two-Tier Architecture

- Client/server-based
- Three layers mapped on two components
- Example:
 - “Fat” client (presentation, application logic)
 - Server (persistence/data, e.g. DBMS)
 - → DBMS could not be easily replaced



Three-Tier Architecture

- Each layer mapped onto separate component
- Different layers expose well-defined interfaces
- Functionality of each layer can be modified without affecting another layer


Real World: N-Tier Architectures

- Typically you will find more tiers (e.g., 2nd layer split into multiple tiers)
- Systems are realized using specific application servers/middleware

2.6 Further System Architectures


Peer-to-Peer Computing

- Combined client and server functionality
- Direct interaction between peers
- No centralized usage/provisioning of a service



Grid Computing


- Aggregation and shared usage of heterogeneous, interconnected resources
- High-end computers with standard OS
- Usually applied to solve large problems
- Multiple owners/decentralized administration



<http://dame.ds.f.unina.it/images/grid01.jpg>

Cloud Computing

- On-demand service provisioning (“computing as a utility”)
- Shared pool of configurable resources using virtual machines
- Rapidly provisioned with minimal management effort
- Single owners/centralized administration



www.definethecloud.net

3 Interprocess Communication

Characteristics

- Two message types for communication between processes: send and receive
- A queue is associated with each message destination
 - Send causes message to be added to remote queue
 - Receive causes message to be removed from local queue
- Communication may be either synchronous or asynchronous

Synchronous

- Sender and receiver synchronize at every message
- Send and receive are blocking operations
 - Sending process is blocked until corresponding receive is issued
 - Receiving process does not proceed until remote data is arrived

Asynchronous

- Send is a non-blocking operation, i.e., transmission in parallel with sending
- Receive is blocking or non-blocking (process proceeds after issuing a receive)

3.1 Message Destinations

Messages are sent to (Internet address, local port) pairs

Local port

- Is message destination within a computer (specified by int number)
- Has exactly one receiver (except for multicast ports)
- Can have many senders

Processes may use multiple ports to receive messages


Servers publish their port number for use by clients

Some port numbers officially assigned to different services (IANA)

Port	Description
20/21	File Transfer Protocol (FTP)
25	Simple Mail Transfer Protocol (SMTP)
53	Domain Name System (DNS)
80	Hypertext Transfer Protocol (HTTP)

3.2 Sockets

- Endpoints of a bidirectional communication channel between client and server
- Different socket types for different underlying protocols (e.g., TCP/IP)
- Foundation for more complex mechanisms (e.g., RPC, RMI, CORBA)



Typical Interaction

- (1) Server opens a server socket to provide a specific service
- (2) Client sends a connection request to the server socket
- (3) Server accepts the client request, opens a new socket, and sends a response
- (4) Client receives and processes the server response
- (5) Client opens a socket for communication
- (6) Client and server communicate via sockets

3.3 UDP/IP Sockets

User Datagram Protocol (UDP)

Transmits datagrams without acknowledgments or retries

Sockets API for UDP

- Receiver creates socket and binds to it (local address, local port)
- Sender sends from any socket, specifies destination in datagram
- Synchronization
 - Send is non-blocking
 - Receive is blocking, gets messages from local queue

Reliability of UDP

- Integrity: checksums to detect corrupt packets
- No guaranteed delivery, ordering or duplicate detection

Usage examples:

- DNS, VoIP, RTP

3.4 TCP/IP Sockets

Transmission Control Protocol (TCP)

TCP abstraction: stream of bytes („pipe“);

- To which data may be written and
- From which data may be read

Stream abstraction hides TCP/IP internals such as

- Lost messages, duplicates, ordering, flow control
- Establishing a connection between communicating processes before they can communicate over a stream

Reliability of TCP

- Integrity: checksums/seq. numbers to detect corrupt/duplicate packets
- Validity: timeouts and retransmissions to handle lost packets
- Limitation: if a connection breaks before an explicit close operation

Usage Examples: HTTP, FTP, Telnet, SMTP

Simple Example: TCP/IP Echo Server

Client

```
Socket sock = new Socket(192.168.2.135, 1234);

DataInputStream in = new DataInputStream(sock.getInputStream());
DataOutputStream out = new DataOutputStream(sock.getOutputStream());

out.writeUTF("Hello");
String data = in.readUTF();
System.out.println("Received: " + data);
```

Server

```
ServerSocket listenSock = new ServerSocket(1234);

for(;;){
    Socket sock = listenSock.accept();

    DataInputStream in = new DataInputStream(sock.getInputStream());
    DataOutputStream out = new DataOutputStream(sock.getOutputStream());

    String data = in.readUTF();
    out.writeUTF(data);
    sock.close();
}
```

4 Remote Procedure Call (RPC)

One example of above is RPC

Fundamental idea behind RPC:

Processes can call procedures on other computers

Goal: Syntactic and semantic uniformity of local and remote calls in terms of

- Call mechanism
- “Expressive power“ of language
- Error handling

→ Goal cannot be fully achieved...

Regardless of problems, RPC is widely used in the computation world

Other goals

- Distribution transparency
- Simplicity (no message conversion/serialization, no packing, no ACKs, ...)
- Type safety (type checking for parameters at client, server)

Definition (Nelson, 1984):

- Synchronous control flow handshake
- Appears to be at level of programming language
- Separate address spaces
- Connected via “narrow-band” channel
- Data passing (message exchange): call parameters results

RPC: “Political” Issues

RPC has been successful

RPC presented/marketed in a good way

- “As easy as what-you-know” (procedural programming)
- Virtually not different from what-you-know
- But: Many subtle differences make it much different and rather difficult


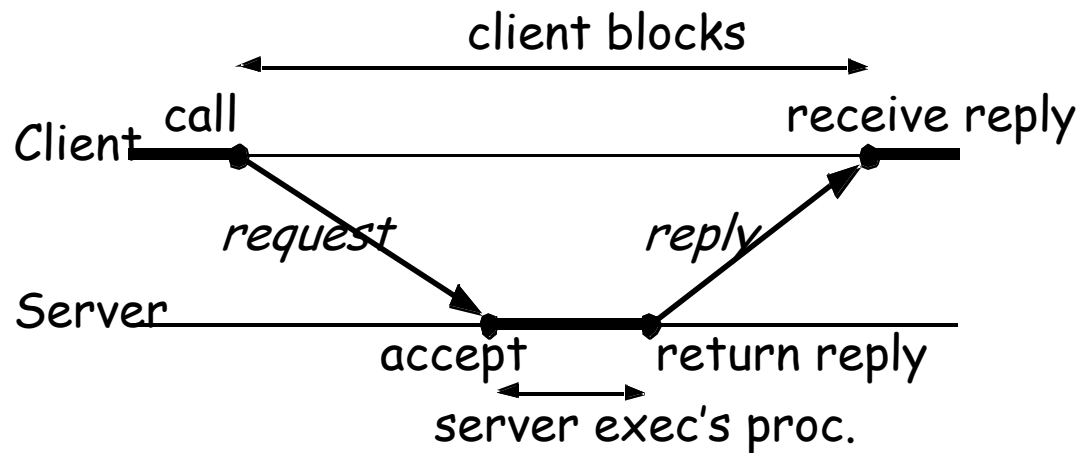
Basis for client/server computing

- Popular in commercial context
- Easy to model and understand, suitable for mid-size problems
- But: in many applications today, client/server model is reaching its limits!
- But: Does not (easily) scale

Sold at right time, together with platforms

- True, but that has nothing to do with RPC in particular

4.2 RPC: Principle



ndr: network data representation
(not a common term)

rpc: rpc protocol

Code in client:

```
Set out-parameters
Call X(out-parameters,result)
Use result
```

Code on server:

```
Proc X(parameters)
Do stuff
Return (result)
```

Client

```
int i = 1, j = 2;  
int res;  
  
res = RPC_add(i, j);  
  
printf("Result %d\n", res);
```

Server

```
int RPC_add(int a, int b) {  
    int res;  
  
    res = a + b;  
  
    return res;  
}
```

Client calls `RPC_add` like any other function

- Function `RPC_add` is implemented in server like any other function
- Client does not see any difference between RPC and local functions (at least in principle)
- Code for communication is hidden in middleware

Problems:

- How does the call to `RPC_add` find the right server?
- How do we know everything works as planned?

RPC: Basic Properties

Synchronous communication

Only 1 call needed to access remote procedure

- Other approaches (e.g., IPC) require more

System takes care of all “small details”

- Message assembly and disassembly, etc.

Complexity same as normal procedure call

- Only one call in progress at a time

Transparent to distribution

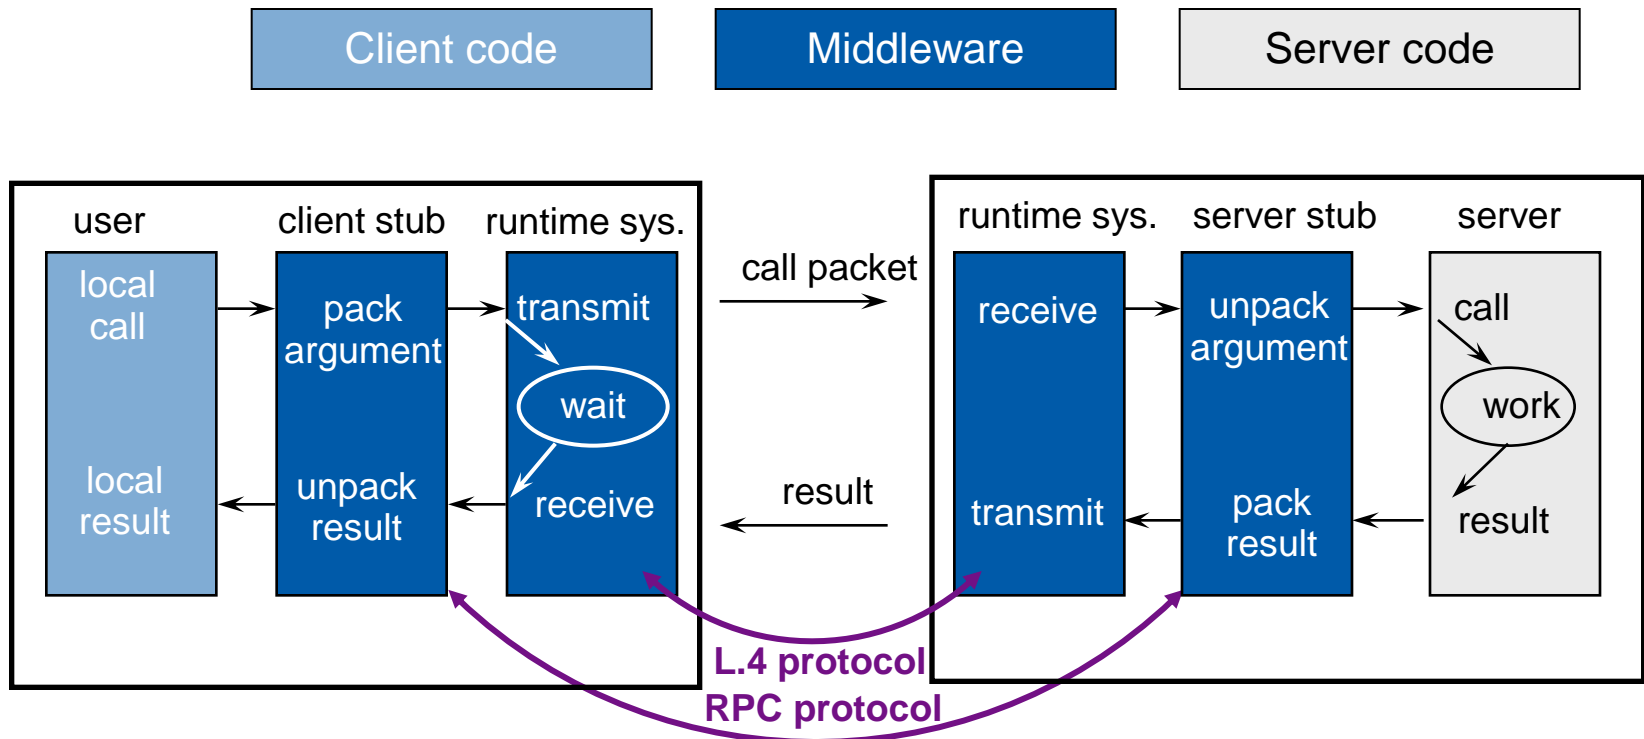
- As long as client can find server, it does not matter where it is

RPC: Implementation / Steps

Server must implement “dispatcher”:

- Dispatcher waits for incoming requests
- When request comes, invokes/dispatches the correct procedure

Flow of control:



Stubs:

- Mimic local procedure call, hide “networking”
 - Call remoteserver.X (out: a; in: b); → extended to client-stub code
- Pack / unpack messages (call, reply, ...):

→ Marshalling

- May convert to / from network data representation
- Support mapping from client to server:

→ Binding

- Carry out RPC protocol

Client stub is proxy for server at client side

Server stub(skeleton) is proxy for client at server side

RPC: Interface definition language (IDL)

Automatic generation of stub code: stub compiler

Basis:

→ **IDL: Interface definition language**

One middleware-specific extension to many programming languages

- Corba-IDL, DCE-IDL, SUN's XDR (external data rep.), Mach's Matchmaker

Internet-wide: XML

Stub/IDL compiler knows types

- used in host programming language
(if strongly typed language)

May use libraries

- in order to insert code
for format conversion, marshalling, transmission control, ...

4.3 RPC: Marshalling and Presentation

Marshalling challenge:

Flatten (serialize) complex data structures

- Basic data types *plus* structural info

Accomodate presentation

in different OSes, languages, and hardware architectures:

- Integer (size?, 1's / 2's complement?) / float/real (size?, IEEE?, ...)
- Character (ASCII, unicode?) / string (`\0` end-flag or byte-count?)
- Arrays (row- /column-based), struct/union/set... (organization?)
- Little-endian, big-endian, bit order (MSB→LSB or inverse)

Worst-case: n systems \rightarrow need $\sim n^2$ conversions

Reality: Often 2 possibilities (see below)

OSI model devoted layer 6 to this problem

- Never widely used, only ASN.1 and BER existed

Two possibilities from above:

“Receiver makes it right”

- Mark representation type
- Between client and server with same representation, no need for translation (80% of the cases?)

Abstract syntax (IDL, or ISO ASN.1) plus standardized network data representation

- ISO-ASN.1 (abstract syntax notation #1): called BER (basic encoding rules)
- SUN-XDR (same name as for IDL), adopted by DCE
- Corba: CDR (common data representation)
- Java-RMI: “Java serialized form”
- XML: SOAP, XMLP

Corba IDL & CDR example:

```
struct Person {  
    string name;  
    string place;  
    long year;  
};
```

index in sequence of bytes ← **4 bytes** → *notes on representation*

0–3	5	<i>length of string 'Smith'</i>
4–7	"Smit"	
8–11	"h__"	
12–15	6	<i>length of string 'London'</i>
16–19	"Lond"	
20–23	"on__"	
24–27	1934	<i>unsigned long</i>


struct with value: {'Smith', 'London', 1934}

```
...  
void CreateNewEntry  
    (in Person newMember,  
     out int MemberNo);  
...
```

Note: BER (for ASN.1) more sophisticated:

- Codes basic, constructed types via tags
- Codes user-defined types (cf. Person)
- Considered unnecessary in IDLs since: both sides know IDL → know “what comes next” in marshalled messages
- Corba IDL uses type tags where needed, e.g., for “union”

4.4 RPC: Binding



Binding matches Clients and Servers

- Clients, servers developed independently

Binding may distinguish just name of server or up to program, version, protocol

Three possibilities

- Static (at compilation time): fast, no middleware overhead
- Semi-dynamic (at startup time): logical name/ DB/ multicast/ service
- Dynamic (= per call), as semi-dynamic plus: fault tolerance, load balancing

Binding via intermediate service (“trader”, “broker”):

- Can become a bottleneck
- Expensive (execution time, triangle communication)
- “Yellow pages” (search via attributes, description)

RPC: Binding Examples

Binding for Internet-RPC (e.g., based on SUN-RPC / XDR):

- Port mapper
rpcbind: name service for RPC servers on server node: port #111
 - Server registers program & version numbers with local port mapper
 - Check out all registered RPC servers with: rpcinfo -p
 - When client calls clnt_create, RPC request is sent to server's port mapper asking info for given program, version, & protocol
 - RPC request returns server's port number to client

For Java: “you’re supposed to know your URLs”

For Corba:

- May use naming service to translate logical name
- May use trader service for yellow pages

4.5 RPC: Protocols


Underlying Layer 4 protocol:

- Connectionless (e.g., UDP): (normally used in Internet)
 - Pro: Need no L.4 acks since we have reply anyway
 - Con: Max. packet size not good for large objects
- RPC-optimized L.4 protocols may provide **implicit** (dis)connect

RPC protocol:

- Distinguish two kinds of RPC protocols
 - **RR**: request-reply
 - **RRA**: request-reply-ack (ACK by client)
- Must realize three functions:
 - client: **do_operation** (callee, op_id, *in_args)
 - server: **getRequest**(...) and **sendReply** (caller, *results)
- Possible errors:
 - Request omission, reply omission (both mean lost messages)
 - Server crash, client crash

Crashes



Server crash: **How far did it get?**

- How to distinguish from omissions (n timeouts?)
- If other server takes over, how should it know?
- For restart, how to cope with locks, dirty states, info recovery?

Client crash: Server executes **“orphan”**


- Procedure execution maybe erroneous, “costly”, ...
- Restarted client should not be puzzled by orphan-results
 - Might even tell server to stop them?
- Note: For lengthy procedures, server may poll client

RPC Errors: Request Omission


Basic countermeasure: Client sets timeout

If no answer within timeout, re-send request

Four possible error cases in this case



- Server must recognize duplicate requests --> Unique ID for requests
- Server must keep state, length of time determines residual error probability



Same countermeasure as for request omission (timeout)

For client, indistinguishable from request omission & delays

For server, means to memorize results for potential 2nd reply

- More states, more memory, especially if parallel calls per client
- Also scalability problem: What if server has thousands of clients?

Now we see why RRA makes sense

- If client ACKs reception of reply, server can throw away stored result

4.7 RPC: Failure Semantics

Maybe-Semantics: No repeated requests (replies, ...)

- Simple, fast, efficient but often not sufficient
- Idea: User will try again in case of failure (check eMail, ...)

At-Least-Once-Semantics: Infinite retry

- Repeated requests, but stateless servers (no duplicates recognized)
- Restricted to idempotent operations (basically, “read”-operations)

At-Most-Once-Semantics: Tolerate omissions

- Repeated requests & replies, duplicates recognized → exec only once
- Server crash → no result (no reply), server may have executed or not

Exactly-Once: Tolerates crashes

- For normal commercial RPC systems, this remains a dream
- Transactional systems, fail-safe solutions needed

Solutions in order or increasing effort

- Offer the choice to programmer
- Commercial systems usually offer some choice of 1/2/3

In summary: Forget transparency, local call \neq RPC

Error Semantic	Absence of errors	In case of omissions	In case of server crash
Maybe	1 proc-exec. 1 result returned	0 1 proc-exec. 0 results returned	0 1 proc-exec. 0 results returned
At-least-once	1 proc-exec. 1 result returned	≥ 1 proc-exec. ≥ 1 result returned	≥ 0 proc-exec. ≥ 0 result returned
At-most-once	1 proc-exec. 1 result returned	1 proc-exec. 1 result returned	0 1 proc-exec. 0 results returned
Exactly-once	1 proc-exec. 1 result returned	1 proc-exec. 1 result returned	1 proc-exec. 1 result returned

Exactly-once remains a dream

- Can approach via redundancy, but this is expensive

Definition:

- A Web Service is a software application identified by an URI, whose interfaces and bindings are capable of being defined, described, and discovered as XML artifacts. A Web service supports direct interactions with other software agents using XML-based messages exchanged via internet-based protocols (W3C).


Web Services

- Are a technology to realize distributed systems.
- Functionalities offered by Web services are similar to traditional Remote Procedure Calls (RPC) apart from the message encoding.


Web Services use XML

- For data exchange and Remote Procedure Calls:
 - XML allows platform independent description of data
 - But performance drawback because of conversion from/into XML

5.1 Web Service Model



Publish-Find-Bind-Execute Paradigm



Roles in a Web Service-based architecture:

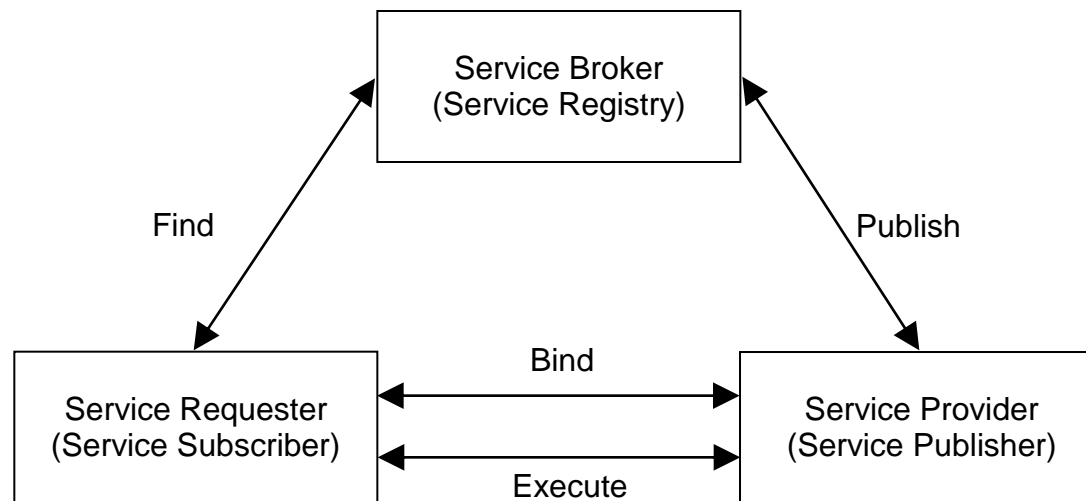
- Service Requestor
- Service Provider
- Service Registry (optional)

Interaction between roles (including Service Registry):

- Service Provider registers Web service with Service Broker
- Service Requestor searches for Web service in registry and receives information about where to find and how to use Web service
- Service Requestor is able to bind service based on the information given by the Service Broker
- Service Requestor is able to execute services provided by Service Provider

Publish-Find-Bind-Execute Paradigm

.. mentioning UDDI, SOAP, WSDL ..



Roles in a Web Service-based architecture:

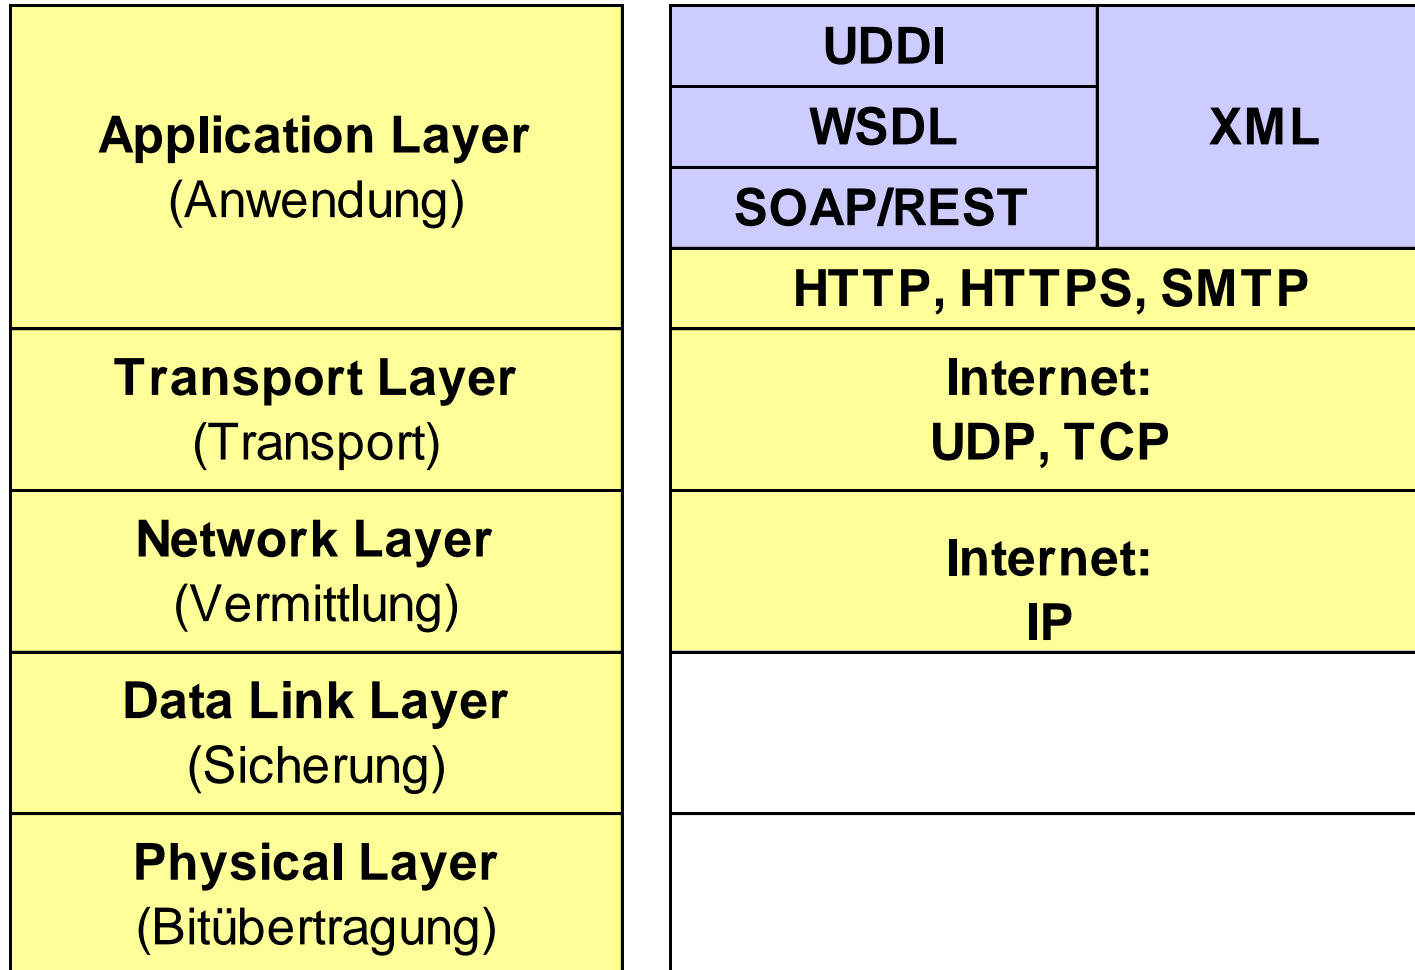
- Service Requestor
- Service Provider
- Service Registry (optional)

Interaction between roles (including Service Registry):

- Service Provider registers Web service with Service Broker (using UDDI).
- Service Requestor searches for Web service in UDDI registry and receives information about where to find and how to use Web service (via UDDI API).
- Service Requestor is able to bind service based on the information given by the Service Broker (based on WSDL).
- Service Requestor is able to execute services provided by Service Provider (using SOAP).

Features of Web Services technology:

- Reusable logic is divided into services
- Services abstract underlying logic
- Services are composable
- Services are autonomous
- Services share a formal contract
- Services are loosely coupled
- Services are stateless
- Services are discoverable



5.2 SOAP – Simple Object Access Protocol

SOAP is an XML-based specification of messages used for communication with Web Services

Content of the SOAP standard

- Contains a syntax for the definition of XML-based messages
- SOAP is a communication model specifying message exchange
- Defines rules for possible content of corresponding messages
- Defines rules for message transport using various L5-protocols (HTTP, SMTP)
- Conventions for Remote Procedure Calls

Messages




- Are always exchanged between two participants
- Can be processed by different intermediaries on their way to the final receiver (so called endpoint)

Every receiver (including intermediaries)

- Opens the messages and processes the part intended for him
- Is able to be sender again; so message passing is possible

Possible communication models are:

- Request-response: sender sends message, endpoint responds to message
- Broadcast: sender generates messages concurrently transmitted to various receivers
- Workflow: chain of senders and receivers processing a single task forming a circle



5.3 WSDL – Web Service Description Language

WSDL

- Is a specification to describe interfaces of Web Services using XML documents
- Defines rules for method invocation (so called contracts)

Contents of a WSDL document describing a Web Service are:

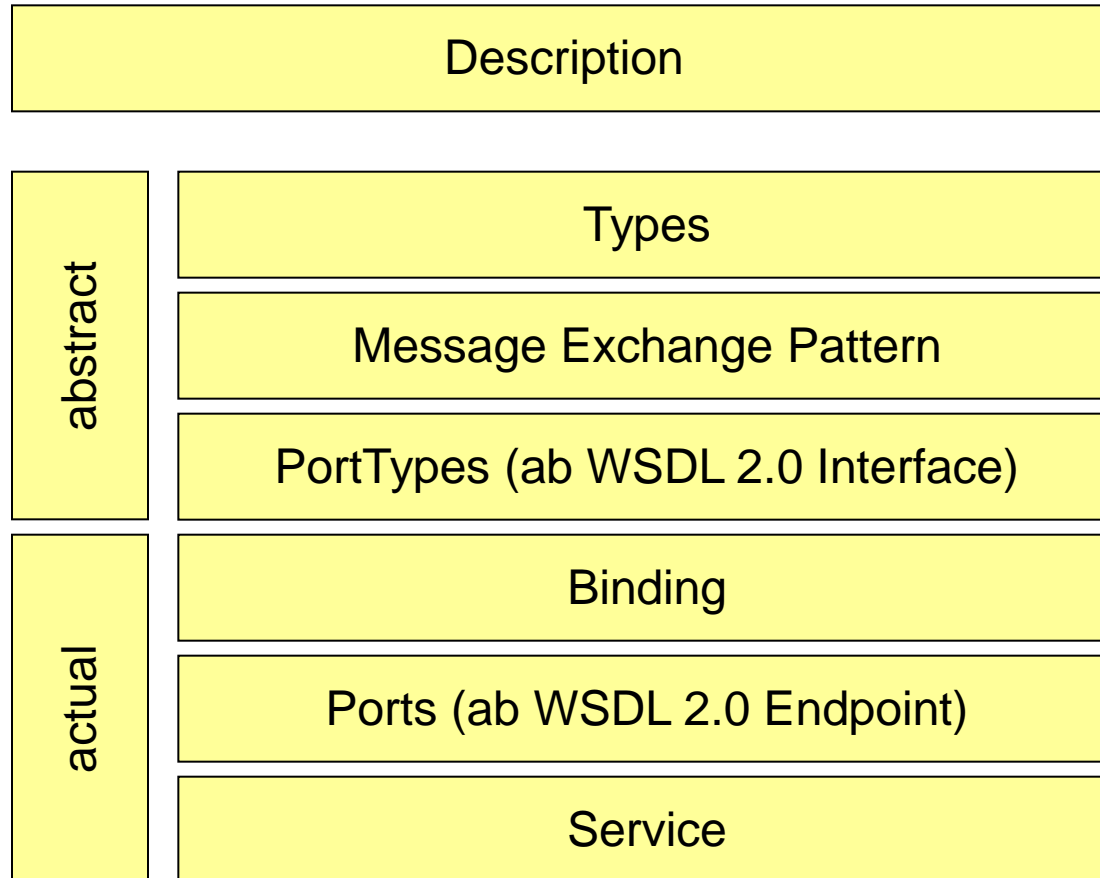
- Definition of data types
- Declaration of methods
- Combination of methods to Web Services
- Mapping of protocols for method invocation (binding)

WSDL allows the following bindings to be used:

- SOAP
- HTTP GET/POST
- MIME

Possible types of communication defined by WSDL:

- One-Way: clients sends message to server, no response by server
- Request-Response: client sends message to server, server sends response
- Solicit-Response: server sends message to client, response by client
- Notification: server sends message to client, no response by client



```
<?xml version="1.0" ?>
  <definitions name="StockQuote"
    targetNamespace="http://example.com/stockquote.wsdl"
    xmlns:tns="http://example.com/stockquote.wsdl"
    xmlns:xsd1="http://example.com/stockquote.xsd"
    xmlns:soap="http://schemas.xmlsoap.org/wsdl/soap/"
    xmlns="http://schemas.xmlsoap.org/wsdl/">
    <types>
      <schema targetNamespace="http://example.com/stockquote.xsd"
        xmlns="http://www.w3.org/2001/XMLSchema">
        <element name="TradePriceRequest">
          <complexType>
            <all>
              <element name="tickerSymbol" type="string"/>
            </all>
          </complexType>
        </element>
        <element name="TradePrice">
          <complexType>
            <all>
              <element name="price" type="float"/>
            </all>
          </complexType>
        </element>
      </schema>
    </types>
```

WSDL – Example

```
...
<message name="GetLastTradePriceInput">
  <part name="body" element="xsd1:TradePriceRequest"/>
</message>
<message name="GetLastTradePriceOutput">
  <part name="body" element="xsd1:TradePrice"/>
</message>
<portType name="StockQuotePortType">
  <operation name="GetLastTradePrice">
    <input message="tns:GetLastTradePriceInput"/>
    <output message="tns:GetLastTradePriceOutput"/>
  </operation>
</portType>
<binding name="StockQuoteSoapBinding" type="tns:StockQuotePortType">
  <soap:binding style="document" transport="http://schemas.xmlsoap.org/soap/http"/>
  <operation name="GetLastTradePrice">
    <soap:operation soapAction="http://example.com/GetLastTradePrice"/>
    <input>
      <soap:body use="literal"/>
    </input>
    <output>
      <soap:body use="literal"/>
    </output>
  </operation>
</binding>
...
```

```
...  
<service name="StockQuoteService">  
  <documentation>My first service</documentation>  
  <port name="StockQuotePort" binding="tns:StockQuoteSoapBinding">  
    <soap:address location="http://example.com/stockquote"/>  
  </port>  
</service>  
</definitions>
```


SOAP – Example Request

POST /InStock HTTP/1.1

Host: www.example.org

Content-Type: application/soap+xml; charset=utf-8

Content-Length: nnn

```
<?xml version="1.0"?>
```

```
<soap:Envelope
```

```
  xmlns:soap="http://www.w3.org/2001/12/soap-envelope"
```

```
  soap:encodingStyle="http://www.w3.org/2001/12/soap-encoding">
```

```
  <soap:Body xmlns:m="http://www.example.org/stock">
```

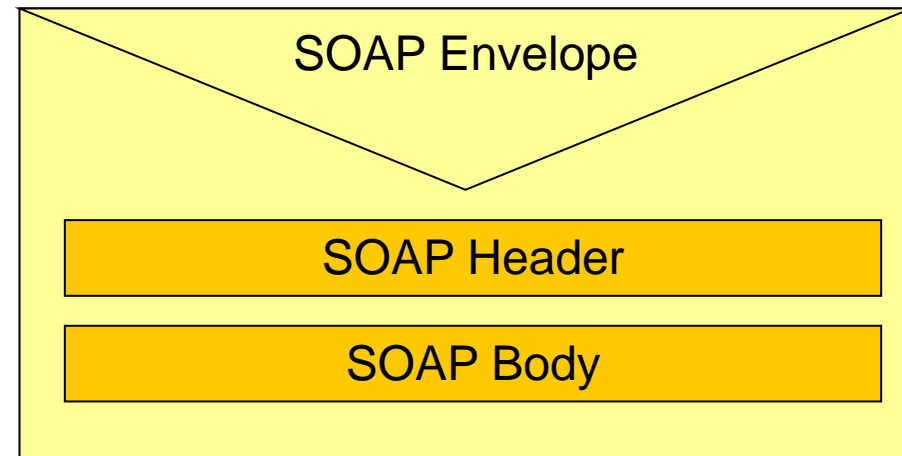
```
    <m:GetStockPrice>
```

```
      <m:StockName>IBM</m:StockName>
```

```
    </m:GetStockPrice>
```

```
  </soap:Body>
```

```
</soap:Envelope>
```



SOAP – Example Response

HTTP/1.1 200 OK

Content-Type: application/soap+xml; charset=utf-8

Content-Length: nnn

```
<?xml version="1.0"?>
```

```
<soap:Envelope
```

```
  xmlns:soap="http://www.w3.org/2001/12/soap-envelope"
```

```
  soap:encodingStyle="http://www.w3.org/2001/12/soap-encoding">
```

```
  <soap:Body xmlns:m="http://www.example.org/stock">
```

```
    <m:GetStockPriceResponse>
```

```
      <m:Price>34.5</m:Price>
```

```
    </m:GetStockPriceResponse>
```

```
  </soap:Body>
```

```
</soap:Envelope>
```

UDDI contains various aspects needed to support dynamic access and usage of Web Services:

- A logically unique but physically distributed directory service architecture in which Web Services can be published and searched
- Requirements on providers of those directory services
 - A description of an API enabling the publishing and searching of Web Services
 - An XML-based data model to describe companies offering Web Services and the Web Services itself

Contents of the UDDI data model are (not limited to Web Services):

- White Pages: information about the company offering the Web Service
- Yellow Pages: classification of the company offering the Web Service
- Green Pages: description of the offered Web Service, containing technical information how to find and use the offered Web Service

5.5 Web 2.0 – AJAX

AJAX – Asynchronous JavaScript and XML

Idea:

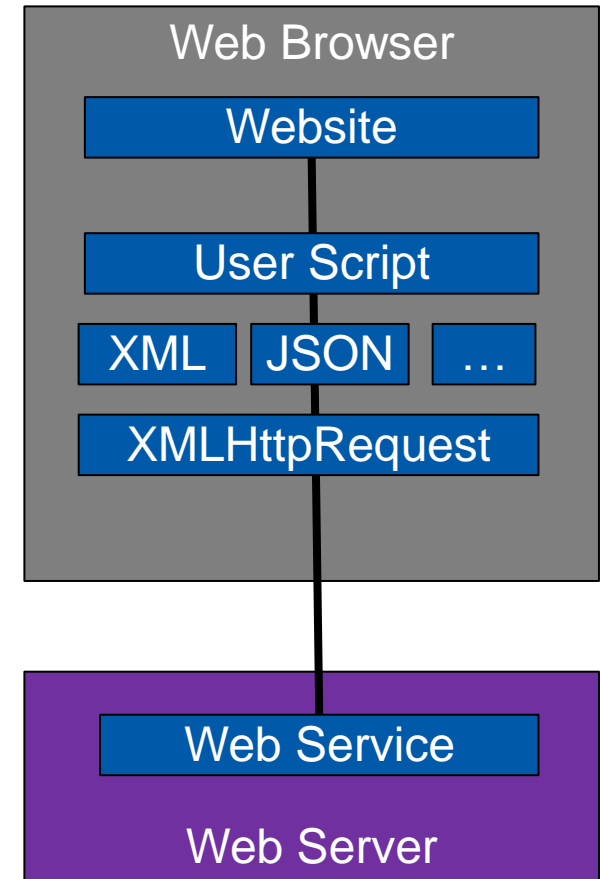
- To build Rich Internet Applications (RIA), i.e., Web applications behaving and looking like “normal” applications

Core Element is XMLHttpRequest

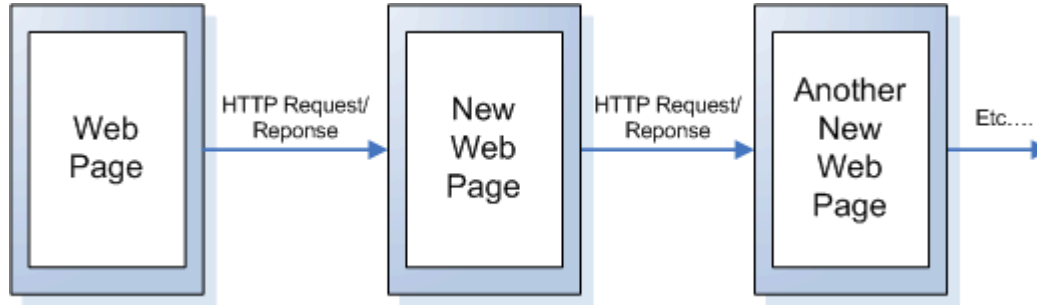
- Retrieval of data without loading complete website
- A Web service invocation
- But: Not limited to transmission of XML (JSON etc. more often used)

Sites using AJAX and therefore WS:


- Gmail
- Facebook
- Yahoo
- ...



Traditional Web interaction



AJAX based style



Basics

- Introduced by Roy Thomas Fielding in his doctoral dissertation in 2000
- Is an architectural style for distributed Web applications, i.e.,
 - REST is not a protocol (in contrast to SOAP)
 - REST is not a standard (in the sense of the W3C)
- Leverages the existing principles and standards of the WWW (e.g. HTTP, URI)
- An application which adheres to REST architectural style is called „RESTful“
- RESTful applications manage an arbitrary number of resources
- When a client follows a URL within an application, the resulting page represents the next state of the application (→ REpresentational State Transfer)

Resource

- Can be everything that can be referenced by URIs (acc. to W3C)
- Can be identified uniquely
- Has one or more representations (i.e., current state, file format)
- Should be a noun, not a verb
- Example: a song (e.g., three representations: audio file, video file, text file)

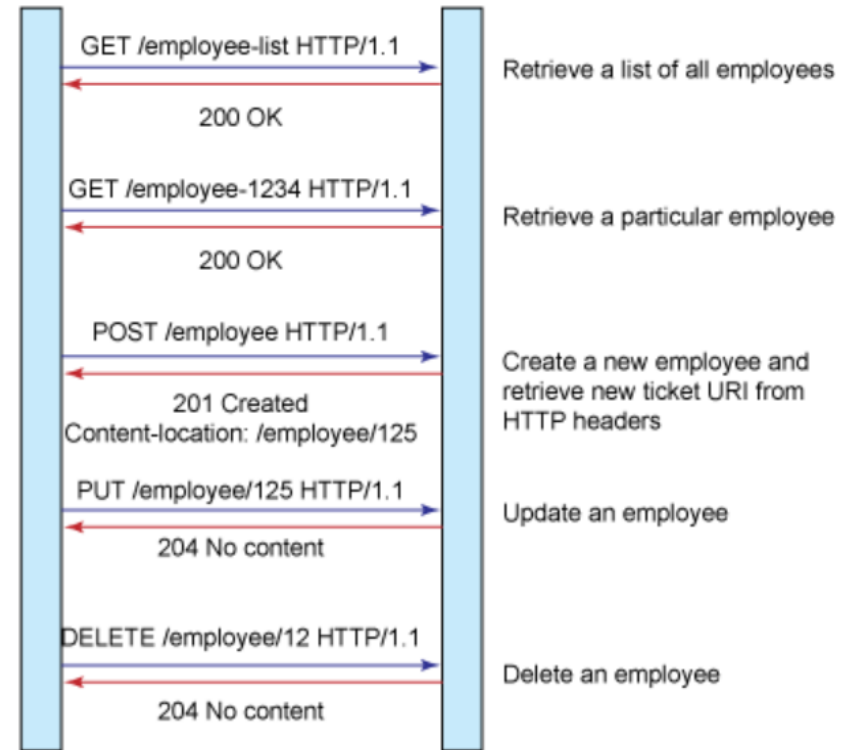
REST – Representational State Transfer (cont'd)

Communication

- Relies on simple point-to-point communication over HTTP
- HTTP methods (generic interface):
 - GET: get representation of a resource
 - POST: update a resource
 - PUT: create or overwrite a resource
 - DELETE: remove a resource
- All interactions are stateless
- Typically XML or HTML payload

RESTful Web Service Example

- Manage your employees
- <http://www.employee-details.com/employee-list...>



Source: <http://www.ibm.com/developerworks/webservices/library/ws-restajax/>

AJAX and RESTful Web Services

- XMLHttpRequest object allows for GET, POST, PUT, DELETE to server
- AJAX-based HTML page can act as client to the above RESTful Web service

Benefits

- HTTP widespread
→ standardized implementation;
easy to build
- Generic, i.e., uniform interface
→ clients can be reused
- HTTP Caching
→ faster response times
- Stateless
→ improved scalability and load-balancing
- Not a lot of additional markup
information
→ lightweight messages

**Does not support the whole
Webservice-* (e.g., Webservice-
security) advanced protocol
stack**

Limitations

- Relies on HTTP protocol
→ not transport-agnostic
- Point-to-point communication
→ not usable for more complex
communication models in
distributed environments (cf.
SOAP communication models)
- No formal interface description
language like WSDL
- Only synchronous message
exchange (cf. callbacks using
SOAP and SMTP)
- No support for features such as
security, reliable messaging, etc.
→ no support for complex user
requirements (cf. specific SOAP
header)

What's “wrong” with RPC and other similar mechanisms?

They are based on Request/Reply approach:

- Client
 - has initiative, client “pulls” data from server
- Peer
 - “locked” in handshake
- Bad if:
 - Lots of data
 - many receivers which come and go
 - information not generated very often

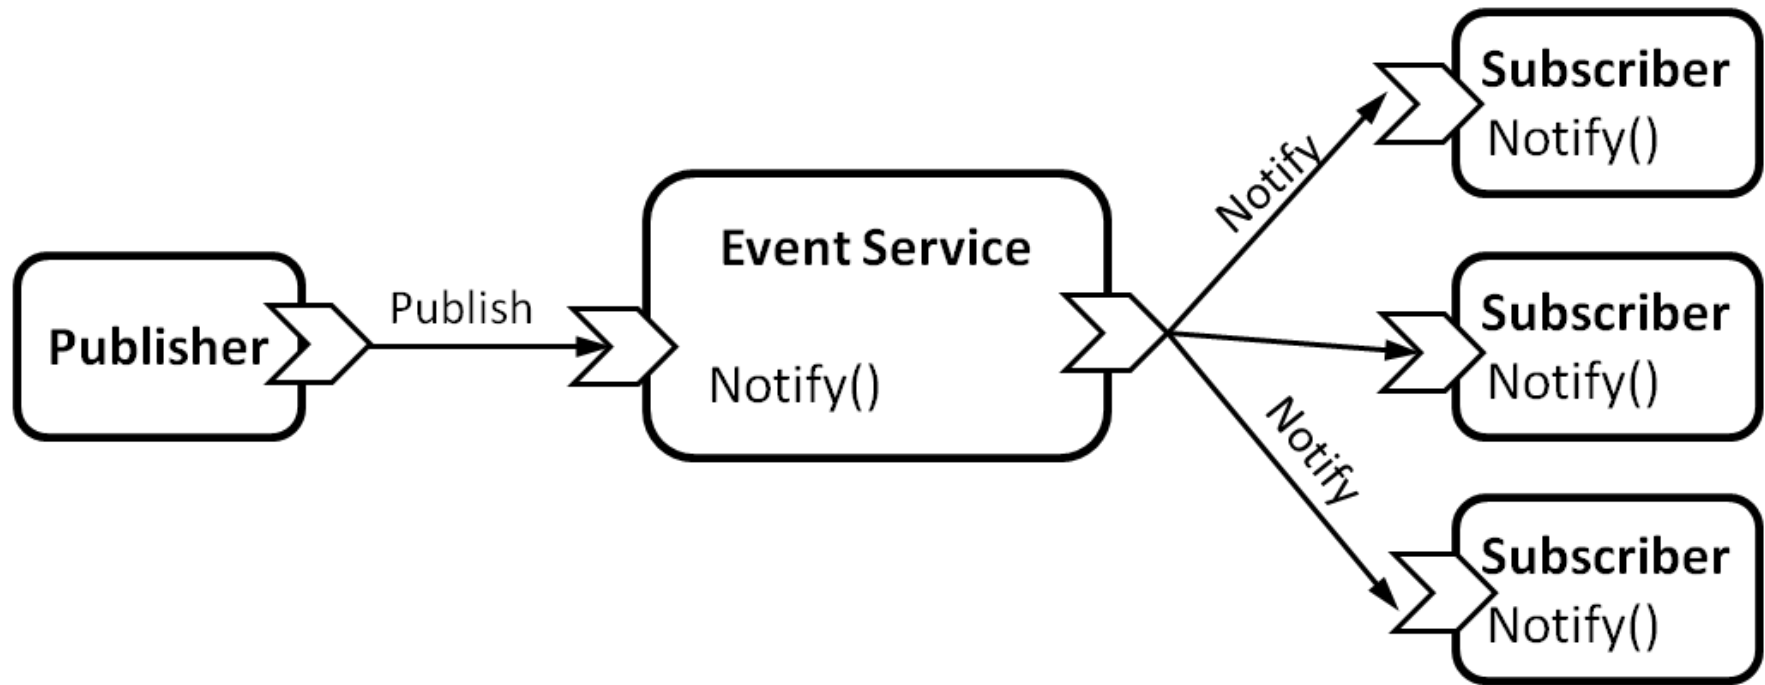
Push approach

- Idea:
 - Peer with data “pushes” it to interested parties
- Producer:
 - Immediate information delivery (“publish”)
- Consumer:
 - Initial “subscribe” for event-type/channel (= info-category)
- Consumer receives events
 - that match the subscription asynchronously as they are generated by producer

Such systems are also called publish/subscribe

Note: Event systems widely used in many fields, e.g., GUI

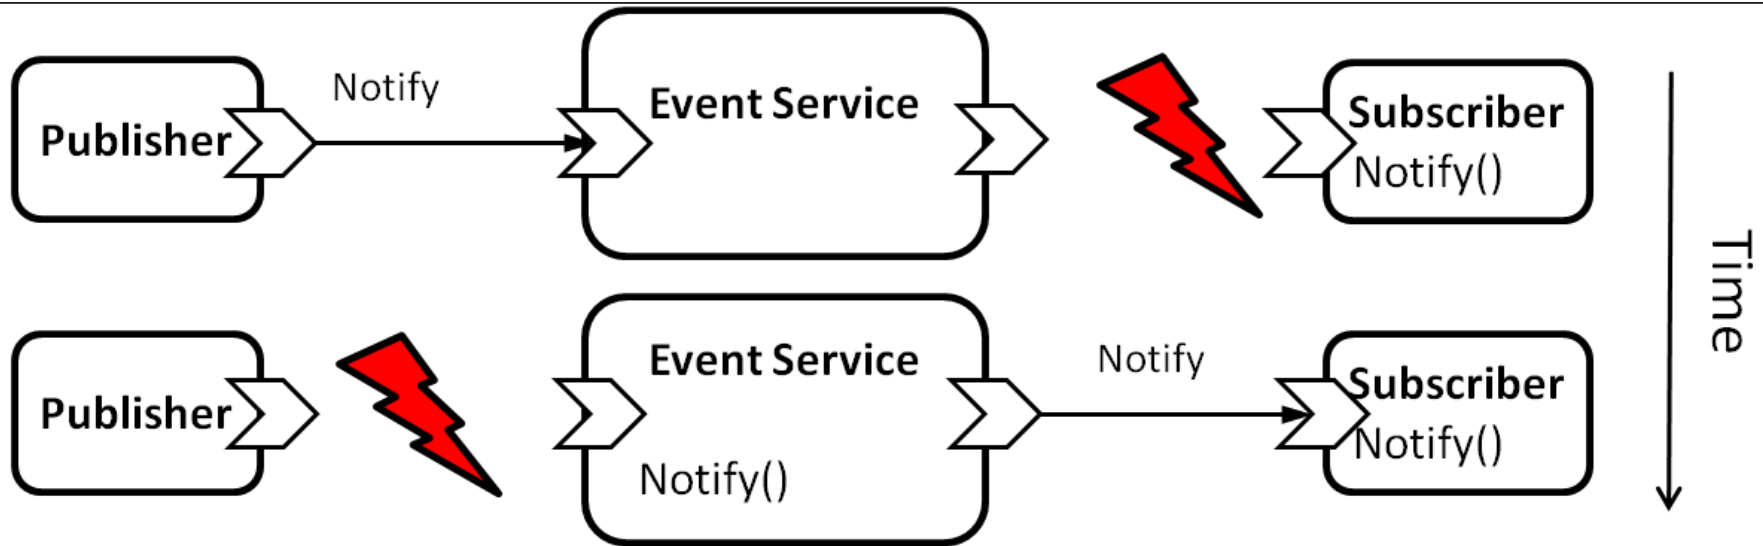
6.1 Basic Principles



Space decoupling

- The interacting parties do not need to know each other
- Producers publish messages through an event service
- Subscribers indirectly receive messages from event service
- One-to-many communication patterns possible


Basic Principles: Time Decoupling



Time decoupling

- The interacting parties do not need to be actively participating in the interaction at the same time


Reducing space and time dependencies greatly reduces the need for coordination between systems



Permit
loosely coupled,
asynchronous point-to-multipoint communication patterns

Application independent infrastructures

Clients communicating via a logically centralized component



Logically centralized component

- Single server /
Network of event routers
- Transparent for application
(=Client)
- Router network can be reconfigured
 - independently and
 - without changes to the application

--> Scalability

Messaging Domain

- Point-to-Point (Producer -> Consumer)
- Publish/Subscribe
 - Subscription-based
 - Advertisement-based

Subscription Mechanism

- Channel-based (=Topic-based) Subscription
- Content-based Subscription
- Subject-based Subscription (limited form of Content-based sub.)

Classification

Event Data Model


- Un-typed
- Typed
- Object-oriented

Event Filters

- Expressiveness and flexibility of subscription language
 - Simple Expressions
 - SQL-like Query Language
 - (Mobile) Code
- Pattern Monitoring: Temporal sequence of events
- Evaluated in router network

Scalability <-> Expressiveness Tradeoff

- Simple Expressions permit Filter Merging -> better scalability




Each message has only one consumer

Receiver acknowledges successful processing of message


No timing dependencies between sender and receiver

Queue stores message (persistent), until

- It is read by a receiver
- The message expires (Leases)



- Interested parties can subscribe to a channel (topic)
- Applications post messages explicitly to specific channels
- Each message may have multiple receivers
- Timing dependency between publishers and subscribers



- Publisher advertises topic before publishing
- Subscribers can get a list of advertised topics
- Avoids problem of subscriber having to figure out which topics are available for subscription

Topic-Based vs. Content-Based

Topic-based subscription

- Messages sent to a well-known topic
- Subscribers subscribe to topics
- Topics are typically expressed as strings

Content-based subscription

- Subscriptions are matched against the content of the message
- Subscribers describe their interest as filter expressions

Special case: Subject-based subscription

- Special case of content-based subscription
- Well-known subject in messages
- Subscriptions matched against the subject
- Subject typically strings or key-value-pairs

6.3 Subscription Semantics



Covering Relations

- Attribute Filter: Filter covers Notification (= message)

$$\phi \subset_f^n \alpha :\Leftrightarrow \phi \text{ covers } \alpha$$

$$\phi \subset_f^n \alpha :\Leftrightarrow \phi.name = \alpha.name \wedge \phi.type = \alpha.type \wedge \phi.match(\alpha.value, \phi.value)$$

- Subscription: Subscription covers Notification

$$s \subset_s^N n :\Leftrightarrow s \text{ covers } n$$

$$N_s(s) \subseteq N; n \in N_s(s) :\Leftrightarrow s \subset_s^N n$$

$$N_s(s) = \{n \in N : \forall \phi \in s : \exists \alpha \in n : \phi \subset_f^n \alpha\}$$

- Examples:

$$\begin{array}{l} \text{String event=alarm} \\ \text{Time date=02:40:03} \end{array} \subset_s^N \begin{array}{l} \text{String event=alarm} \\ \text{Time date=02:40:03} \end{array}$$

$$\begin{array}{l} \text{String event=alarm} \\ \text{Integer level>3} \end{array} \not\subset_s^N \begin{array}{l} \text{String event=alarm} \\ \text{Time date=02:40:03} \end{array}$$

Covering Relations

- Advertisement: Advertisement covers Notification

$$a \subset_A^N n :\Leftrightarrow a \text{ covers } n$$

$$N_A(s) \subseteq N; \quad n \in N_A(a) :\Leftrightarrow a \subset_A^N n$$

$$N_A(a) = \{n \in N : (\forall \alpha \in n : \exists \phi \in a : \phi.name = \alpha.name) \\ \wedge (\forall \alpha \in n : \forall \phi \in a : \phi.name = \alpha.name \Rightarrow \phi \subset_f^n \alpha)\}$$

- Advertisement covers Subscription

$$a \subset_A^S s :\Leftrightarrow N_A(a) \cap N_S(s) \neq \emptyset$$

- “a is relevant for s”

String event=alarm
Time date any
Integer level>0

\subset_A^S

String event=alarm
Integer level>3

String event=alarm
Time date any
Integer level>0

$\not\subset_A^S$

String event=alarm
Integer level>3
String user any

String event=alarm
Time date any
Integer level>0

\subset_A^S

Integer level>5

Subscription-based event service

- Service delivers notification n to party X iff
 - X subscribes s
 - $s \subset_S^N n$

Advertisement-based event service

- Service delivers notification n posted by object Y to party X iff
 - Y advertises a
 - X subscribes s
 - $a \subset_A^S s$
 - $s \subset_S^N n$


Remote Method Invocation (RMI)

Common Object Request Broker Architecture (CORBA)

Enterprise Java Beans (EJB)

7.1 Remote Method Invocation - RMI

**Developed by SUN
for the Java Platform**



Server Object Interface

- Plain Old Java Interface
- Describes Operations provided by the server object
- Shared between the client and the server

Remote Object = Distributed Object

RMI Components

Server Stub (automatically generated)

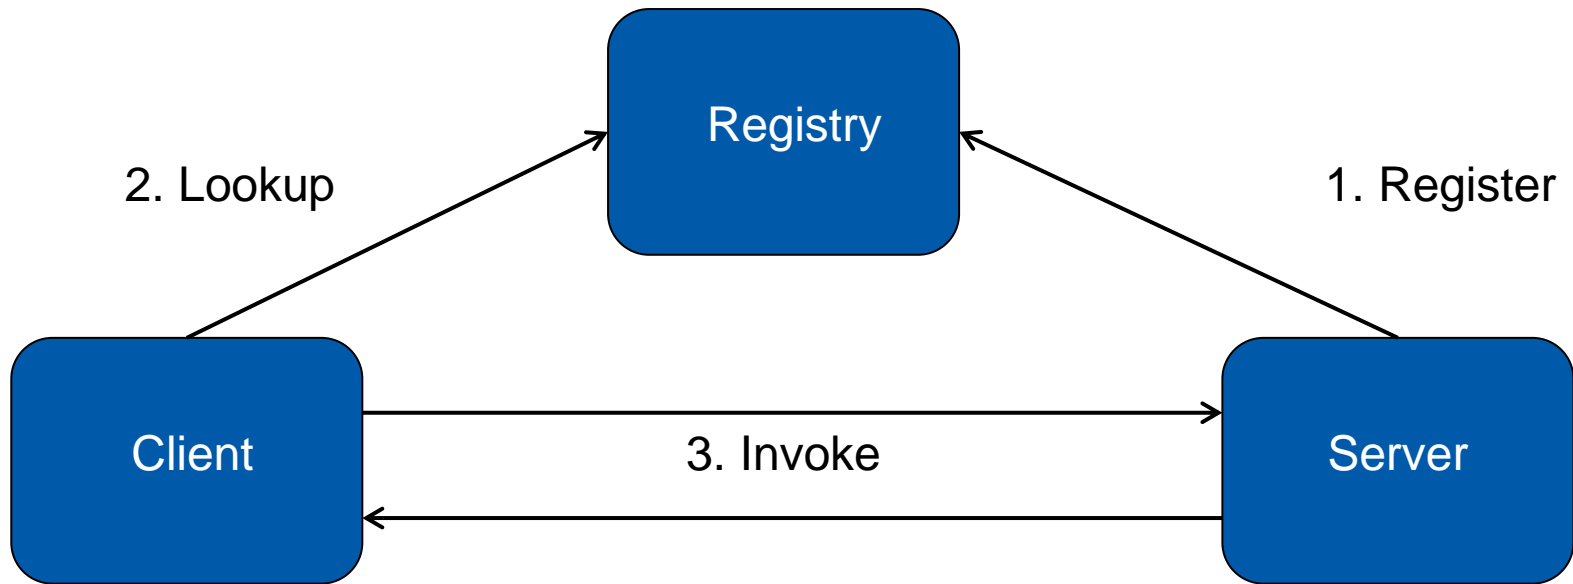
- Provide an interface to call the remote method(access transparently)
- Pack the arguments and transfer them to the server host
- Unpack the result and return it to the caller

Server Skeleton (automatically generated)

- Unpack remote method call parameters
- Invoke the method on the server object
- Pack the result of the invocation
- Transfer the result to the client host

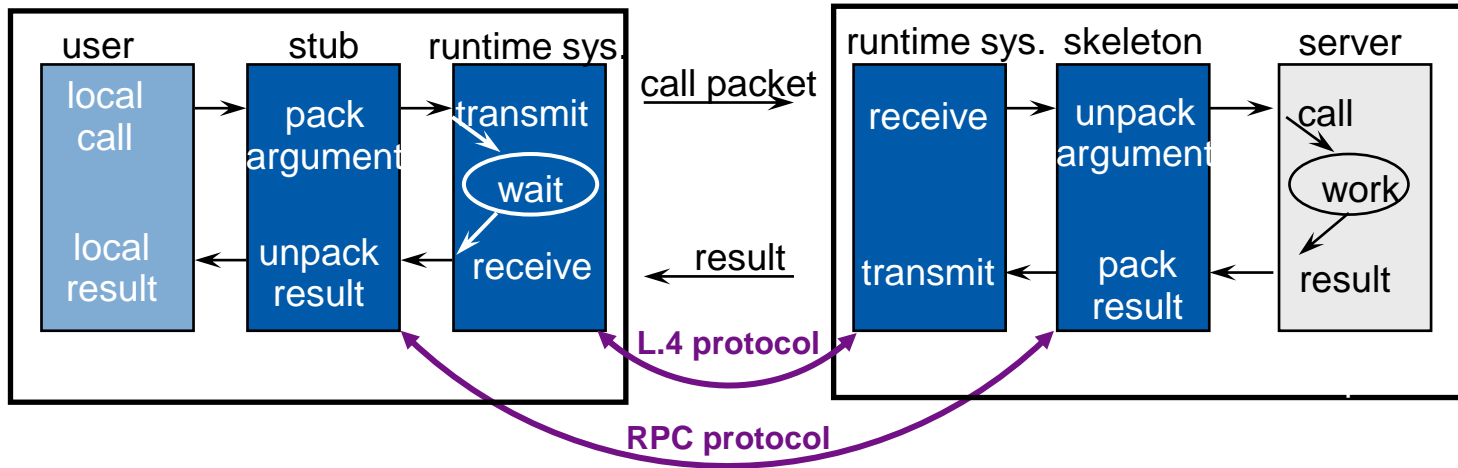
RMI Registry

- Provides a name service for server objects
- The Server Host register server objects
- The Client Host lookup server objects by their name
- → The client is supposed to know the name of the server object



Workflow (On Bootstrap)

- Server Object Interface is shared between server and client
- 1. Server registers the Server object at the RMI Registry
- 2A. Client looks up the required server object
- 2B. RMI Registry returns address of the server object
- 3. client could invoke the server object



Workflow (Remote Method Call)

- Equal to RPC (see drawing)
- 1. Client calls the Server stub(which is a local method)
- 2. server stub serializes the arguments and writes them to the network stream
- 3. server skeleton reads the network stream, unpacks the arguments and calls the server object method
- 4. server object method performs the computation
- 5. server skeleton packs the result and writes them to the network stream
- 6. server stub reads the response from the network stream unpacks it and returns it to the caller

Remote Object Interface

```
public interface HelloServer extends Remote {  
    public String sayHello() throws RemoteException;  
}
```

Server Side(Remote Object)

```
public class HelloImpl extends UnicastRemoteObject  
    implements HelloServer {  
    public String sayHello() {  
        return "Hello from the remote server!";  
    }  
}
```

Registration

```
Naming.rebind("server", new HelloImpl() );
```

Client (Usage of remote object)

```
HelloServer server(HelloServer)Naming.lookup("//localhost/server");  
String result = server.sayHello();  
System.out.println(result);
```

Local vs. Distributed Objects

Differences between local and distributed objects

Distributed objects have more meta data

- (object location, unique id, type meta data)
- Depending on the middleware

Remote method invocations have much higher latency

- three to four orders of magnitude

Creation:

- Objects are either created on another host
- or migrated to another host

Cleanup:

- Network communication necessary to decrease reference counter
- For each reference to an object

Referencing:

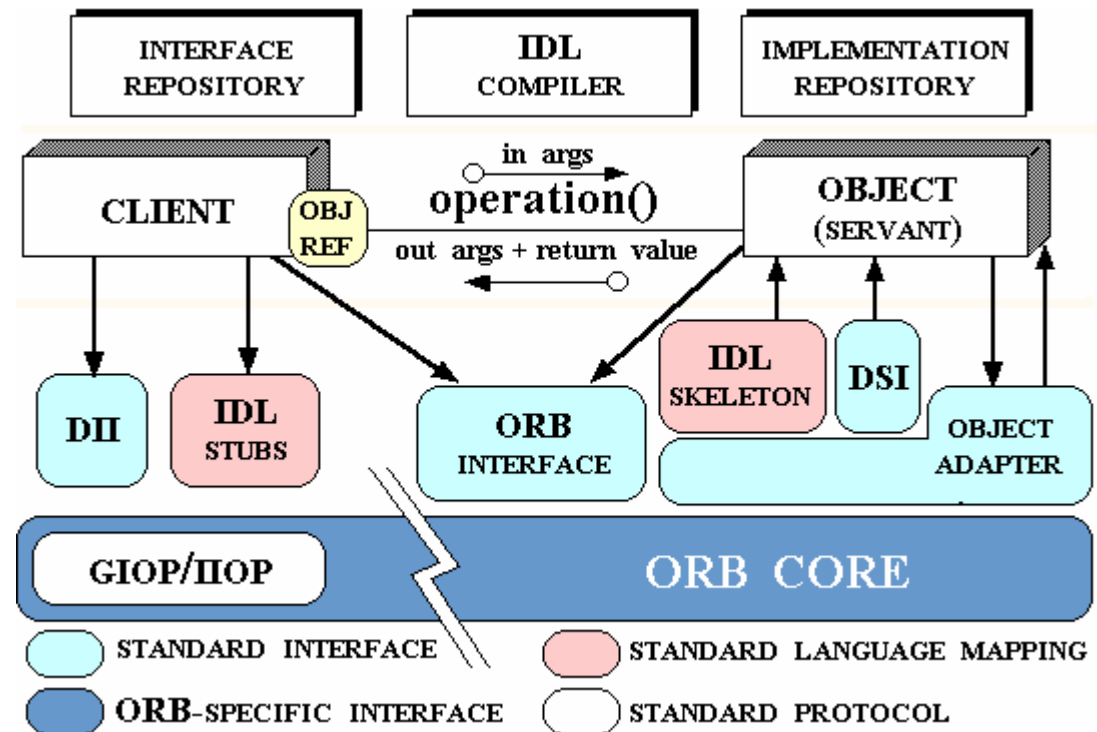
- network communication necessary to increase reference counters

7.2 Common Object Request Broker Architecture - CORBA

Common Object Request Broker Architecture

Specified by the Object Management Group

- Defines an interface definition language (IDL)
- Defines a set of language bindings (C, C++, Java, Smalltalk, Ada95, ...)
- Stubs and Skeletons are automatically from the IDL



DII = Dynamic Invocation Interface
DSI = Dynamic skeleton interface
ORB = Object request broker
IDL = Interface Definition Language
GIOP = General Inter-ORB Protocol
IIOP = Internet Inter-ORB Protocol

The Rise and Fall of CORBA

Building dist. Systems in the early 90s was a nightmare

- because of many different
 - hardware,
 - operating systems and
 - programming languages
- E.g. different byte order, floating point representation or character encoding
- So Programmers wrote an entire protocol stack themselves

CORBA was designated to solve these problems

- By specification of a platform independent middleware

**The C++ and Java binding gave developers a tool to build
heterogeneous distributed systems**

The Rise and Fall of CORBA

But technical problems stopped the show:

- CORBA API was too complex and error prone
- Programming Component models (COM, EJB) were a lot simpler
- During CORBA's growth phase in the mid 90s, Java and the web changed the computing landscape
- CORBA did not cooperate with the rapidly expanding web

Because:

The standardization committee consisted of too many members

- The committee is based on consensus
- Consensus by saying yes is easier than by saying no
- So most request for proposals were added
- Not even a prototypical implementation was required in the request for proposal

→ **The resulting standard was too complex**

7.3 Enterprise Java Beans - EJB

Enterprise Java Beans (EJB)

**Managed, server-side COMPONENT architecture
for modular construction of distributed ENTERPRISE ARCHITECTURE**

A BEAN is a component

- Entity bean
 - Contains the data model of the business process
 - Persistence is managed either manually or automatically
 - E.g. all articles in an online store are modeled as Entity bean
- Session Bean
 - Encapsulates operations to transform entities
 - Distinguish between stateless and stateful session beans
 - Stateless session beans could have a web service interface
 - E.g. adding an article to the customers shopping cart of an online shop
- Message Driven Bean
 - Covers asynchronously executed operations
 - E.g. sending confirmation emails to customers

Enterprise Java Beans: Architecture

