
Exercise Session 2

Page Addressing, Update Strategies and Shadow Paging

Note: For simplicity, we will assume in exercises 1 and 2 below that database page size is equal to disk block size.

2.1 Segments and Pages, Addressing and Update Strategies

- a) Explain the role of **segments** and **pages** in a DBMS
- b) Explain the differences between:
 - Direct and Indirect Page Addressing
 - Direct and Indirect Update Strategy

2.2 Shadow Paging

- a) Describe the Shadow Paging Method. What data structures are used and what happens when pages of a segment are updated?
- b) Which data structures need to be written atomically to disk when a segment is closed? Why?
- c) Given is a segment S1 with 8 pages, stored in a 16 block file F.

Page	1	2	3	4	5	6	7	8
Block	2	3	4	0	0	0	0	0

Lets assume that the first 8 blocks of F are used (blocks 1 thru 8), while the rest are free. Go step-by-step through the execution of the following operations on S1 and show how the auxiliary data structures (CMAP, MAP0, MAP1, V10, V11) are used:

- OpenSegment(S1)
- Read(Page 1)
- Update(Page 1)

-
- Read(Page 2)
 - Update(Page 2)
 - Update(Page 1)
 - CloseSegment(S1) // Savepoint / Checkpoint

2.3 Costs of Shadow Paging

Given is a segment S of size 128 MB stored in a file F:

- Database page size = 4 KB
 - Disk block size = 512 Byte
 - Size of F = 2 (Size of S) = 256 MB
- a) Assuming that S is full and that there are no other segments stored in F, what is the maximum number of pages that can be updated with a single opening/closing of S?
- b) Calculate the sizes (in blocks) of the free space bitmap MAP0 (assuming that 1 bit is used per page) and the page table V of S (assuming that addressing is performed at the page level).
- c) Calculate the number of block accesses at save/checkpoint if:
- only 1 page was updated
 - all pages were updated

Assume that master record size is 2 blocks and that all updated pages need to be flushed at savepoint - i.e. none of them has already been flushed.

- d) What is the overhead in c)?

2.4 Shadow Paging and Transactions

- a) To what extent can the Shadow Paging Algorithm be used as a basis for implementing a Recovery Mechanism in a Transaction Processing System? What problems need to be addressed when concurrent transactions update the same segment?
- b) What are the advantages/disadvantages of recovery systems based on shadowing in comparison to conventional recovery systems based on in-place updating?