

---

## Exercise Session 7


# Further Concurrency Control Techniques

### 7.1 Data Contention, Resource Contention and Thrashing

- Explain the difference between **data contention** and **resource contention** and the way they could lead to data thrashing.
- What could be done to reduce data contention and achieve throughput beyond the thrashing point?

### 7.2 Tree Locking (TL) Protocols

- Explain briefly the main principles of TL protocols.
- Explain how TL protocols can be used to synchronize access to B+-Trees.
- Given is the following B+-Tree:



- Which nodes are safe (unsafe) for the insert/delete operation?

- Show how the operations Insert(271) and Insert(272) would be executed if the Bayer/Schkolnick algorithm is used.
- To what extent are exclusive locks unnecessarily set in the Bayer/Schkolnick algorithm? How can this be avoided?

## 7.3 Timestamp Ordering Protocols

- Explain why the TO Rule guarantees serializability.
- Describe a possible implementation of a TO-based Scheduler.
  - What data structures are used?
  - How are the Basic TO Rules enforced?
- Given are the transactions T1 (timestamp=1), T5 (timestamp=5) and T10 (timestamp=10). T1, T2 and T3 send the following operations to the scheduler:

| Point in Time | T1    | T2    | T3     |
|---------------|-------|-------|--------|
| 1             | BOT   |       |        |
| 5             |       | BOT   |        |
| 10            | w1[x] |       | BOT    |
| 11            |       | r5[x] |        |
| 12            |       |       | r10[x] |
| 13            |       | w5[x] |        |
| 14            |       | c5    |        |
| 15            |       |       | w10[x] |
| 16            | c1    |       |        |
| 17            |       |       | c10    |

We assume that if an operation is sent to the Data Manager at point in time  $t$ , the scheduler receives confirmation from the Data Manager at point in time  $t+1$  for read operations and  $t+2$  for write operations. Parallel read operations are possible. Under these assumptions, describe in detail the sequence of actions performed by the scheduler. Show how the TO data structures are used.

| Time | BTO Test | w-max-sched[x] | r-max-sched[x] | w-in-transit[x] | r-in-transit[x] | queue[x] |
|------|----------|----------------|----------------|-----------------|-----------------|----------|
| ...  |          |                |                |                 |                 |          |

If a transaction  $T$  sets  $o\_max\_scheduled[x]$  to  $TS(T)$  and is subsequently aborted can  $o\_max\_scheduled[x]$  be restored to its before image with respect to  $T$ ?

- Discuss the following issues in the context of the Basic TO Method:
  - Deadlocks
  - Livelocks, Starvation
  - Strictness

## 7.4 Optimistic Concurrency Control

- Explain the different paradigms of pessimistic and optimistic CC methods.

---

b) Describe the typical phases in the execution of transactions by an optimistic CC scheduler.

c) Describe a possible implementation of a locking-based optimistic CC protocol.

- What criteria are used for validation?
- Which of these criteria for backwards validation would be sufficient, if additional information about the relative order of transaction's phases is available? When should transaction  $T_n$  be backwards validated with respect to  $T_m$ , if:
  - 1.)  $\text{end\_write\_phase}(T_m) < \text{start\_read\_phase}(T_n)$
  - 2.)  $\text{end\_write\_phase}(T_m) < \text{start\_write\_phase}(T_n)$
  - 3.)  $\text{end\_read\_phase}(T_m) < \text{end\_read\_phase}(T_n)$
- Which transactions must the scheduler backwards-validate against? When can the protocol data (read/write sets) for a transaction be deleted?