

Administrative Info

— OS1-Labs —

Administrative Info: Labs

□ General info

- Not mandatory, but you can gain bonus points (cf. Intro Lecture). **We strongly recommend to participate in the labs!**
- Programming tasks, C language
- Overall, there will be **3 labs** that offer hands-on experience with the Linux (kernel) programming environment.
- Labs get published one after another via TUCaN.
The next lab starts after the testing of the previous one.
 - Lab 1: Monday, Nov 3
 - Lab 2: Monday, Nov 24
 - Lab 3: Monday, Dec 15
- **Work in groups of 2-3 people!**
 - No individual submissions!
- Check our website regularly for updates!

Administrative Info: Labs

❑ Lab submissions

- Submit your group solutions via email
 - os-lab@deeds.informatik.tu-darmstadt.de
- List the **names, matriculation numbers and email addresses** of all group members in your submission email.
- Package your solution as tarball or zip archive
 - Include all files needed to build and execute your solution
- **Indicate time collisions** with other courses for testing in your email!
(see next slide)

Administrative Info: Labs

□ Lab testing

- For each lab, you have to pass the testing to get the bonus points.
- **Attendance is mandatory** for all group members.
- We build and execute your submitted solutions; you have to explain it and answer some questions.
- Testing takes place on the Friday after the submission deadline for each lab.
- We assign testing time slots for each group.
 - **Indicate collisions with other courses in your submission email!**
Otherwise, we cannot take collisions into account.
 - The **time slots are fixed** once they are assigned! We won't accommodate change requests afterwards!

Administrative Info: Labs

❑ Recommended setup and advice for all labs

- Use a **recent Linux** distribution with up-to-date tools and kernel
 - Linux kernel 3.12 or newer should work
 - We use kernel 3.16 (or newer) for testing
 - Do not use unstable or outdated kernel versions
 - GNU make 3.8+ and GCC 4.6+
 - GCC 4.8 or 4.9 recommended
- Use a **virtual machine** for your test/development environment
 - Later labs deal with kernel modules and may crash your system!
- Avoid hardcoded paths in your Makefiles and shell scripts, especially if they include kernel specific components.
 - Paths and kernel versions on your test system may differ from ours.

Administrative Info: Lab 1

- ❑ Lab 1: Processes and IPC
 - Starts next week
 - Covers contents from this and last lecture

- ❑ Schedule: Lab 1
 - **Monday, Nov 3**
 - Publication of lab description via TUCaN
 - **Tuesday, Nov 18**
 - Deadline for solution submissions
 - **Friday, Nov 21**
 - Testing
 - Takes place in the morning in room E302



Processes (and Threads)

- Process Management
- Thread Models
- Inter-process Communication (IPC)
- Hierarchical Microkernel and IPC

Cooperating Processes

❑ Processes can be

- **Independent**

- A process is independent if it cannot affect or be affected by other processes executing in the system.
- Independent processes do not share data with other processes.

- **Cooperating**

- A process is cooperating if it can affect or be affected by other processes executing in the system.
- A process that shares data with other processes is a cooperating process.

Cooperating Processes

❑ Why do we want cooperating processes?


- Information sharing
- Speedup
- Modularity
- Convenience
- Privilege separation

❑ Process cooperation issues

- data sharing
- need a communication channel/medium
- coordination & synchronization (later lectures)
 - race conditions, deadlocks
 - critical sections, locking, semaphores


Inter-process Communication (IPC) Models

© Silberschatz et al.: Operating System Concepts, 8th ed., Wiley



(a)

Message Passing




(b)

Shared Memory

Shared-Memory: Producer-Consumer Model

❑ Cooperating processes

- **Producer** process produces information/data
- **Consumer** process consumes information/data
- Both share a common memory region (buffer)



❑ How large should the buffer be?

- bounded → assumes fixed buffer size
 - full buffer?
- unbounded → no practical limit on buffer size
 - realistic?

Shared-Memory: Bounded Buffer Producer-Consumer

- ❑ Variables shared between producer and consumer processes

```
#define BUFFER_SIZE 10

typedef struct {
    ...
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

buffer bound: 10 items

shared circular buffer

next free position, shared

first full position, shared

- ❑ How to recognize that buffer is

- empty?

```
in == out
```

- full?

```
(in + 1) % BUFFER_SIZE == out
```

Shared-Memory: Bounded Buffer Producer-Consumer

```
while (true) {  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; // do nothing, buffer full  
  
    // produce an item  
    buffer[in] = new_item;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

❑ Producer

- waits if buffer is full
- inserts new items

```
while (true) {  
    while (in == out)  
        ; // do nothing, buffer empty  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    process(item);  
}
```

❑ Consumer

- waits if buffer is empty
- removes items

explicit synchronization needed!
covered in later lectures

Concurrent access to shared variables?

Message Based Communication

- ❑ **Shared memory** IPC – processes use shared memory region as communication medium

- ❑ **Message passing** IPC – processes use (abstract) communication link
 - two primitive operations, message size may be fixed or variable
 - **send**(message)
 - **receive**(message)
 - If P and Q wish to communicate, they need to:
 - establish **a communication link** between them
 - exchange messages via send/receive
 - Implementation/design of the communication link
 - physical (shared memory, hardware bus)
 - logical (logical properties)

Link Implementation Issues

- ❑ How to establish the links?
 - ❑ Can a link be associated with more than two processes?
 - ❑ How many links can exist between every pair of communicating processes?
 - ❑ What is the capacity of a link?
 - ❑ Is the size of a message that a link can accommodate fixed or variable?
 - ❑ Is a link unidirectional or bi-directional?
-
- Addressing, naming (direct vs. indirect communication)
 - Synchronization Aspects (blocking vs. nonblocking)
 - Buffering (link capacity)

Addressing: Direct Communication

❑ Processes must name each other explicitly:

- P: **send**(Q, *msg*)
 - process P sends a message to process Q
- Q: **receive**(P, *msg*)
 - process Q receives a message from process P

- Symmetric addressing
 - send(Q, *msg*); receive(P, *msg*)
- Asymmetric addressing
 - send(Q, *msg*); receive(&pid, *msg*)

receive msg from any process, receiver also gets sender's process id

❑ Properties of communication link

- Links are established automatically
- Link associated with exactly one pair of communicating processes
- Between each pair there exists exactly one link
- The link may be unidirectional, but is usually bi-directional
- Various queue models (zero, bounded, unbounded)

Addressing: Indirect Communication

- ❑ Messages are directed to and received from **mailboxes or ports**
 - Communicating processes do not have to know each others names
 - Each mailbox has a unique id X
 - operations on mailboxes
 - `send(X , msg)`, `receive(X , msg)`
 - Processes can communicate only if they share a mailbox
 - Example Unix pipes: `cat xyz | lpr`; implicit FIFO mailbox

- ❑ Properties of communication link
 - Link is established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional
 - Various queue models (zero, bounded, unbounded)

Addressing: Indirect Communication

❑ Mailbox operations

- create a new mailbox, once per mailbox
- send and receive messages through mailbox
- destroy a mailbox, once per mailbox


❑ Who owns a mailbox?

- Process
 - resides in process address space
 - mailbox only exists as long as the process lives
 - owner process receives, others send
- Operating system
 - mailbox exists on its own, not coupled to some process
 - possibly multiple receivers

Addressing: Indirect Communication

❑ Mailbox sharing

- P_1 , P_2 , and P_3 share mailbox A
- P_1 sends; P_2 and P_3 execute receive
- Who gets the message?



❑ Solution choices based on whether we

- allow a link to be associated with at most two processes
- allow only one process at a time to execute a receive operation
- allow the system to select arbitrarily the receiver; sender is notified who the receiver was.

Synchronization Issues

- ❑ Message passing may be either **blocking** or **non-blocking**
 - different combinations for operations possible

- ❑ **Blocking** (B) is considered **synchronous**
 - **B.Send**: the sender blocks until the message is received by the receiving process or mailbox
 - **B.Receive**: the receiver blocks until a message is available
 - **Rendezvous**: both send and receiver are blocking; no buffers needed

- ❑ **Non-blocking** (NB) is considered **asynchronous**
 - **NB.Send**: the sender sends the message and continues
 - **NB.Receive**: the receiver receives a valid message or null

Message Buffering (Link Capacity)

- ❑ Messages of communicating processes reside in temporary queues; implemented in one of three ways:
 - **Zero capacity** – 0 messages
 - Sender must wait for receiver for each message (rendezvous)
 - **Bounded capacity** – finite length of n messages
 - Sender must only wait if link is full
 - **Unbounded capacity** – infinite length
 - Sender never waits (no blocking)

What about receiver? When does it wait?
Are there more possibilities in terms of blocking vs. non-blocking?

IPC Example: (A)LPC in Windows


□ (Advanced) Local Procedure Call in Windows

- Message-passing IPC facility
- Supports blocking and non-blocking semantics
- Arbitrary sized messages
- Only works between processes on the same system
- Uses ports (like mailboxes) to establish and maintain communication channels
- Client-Server style setup
- Windows internal mechanisms, not directly accessible to 3rd parties
- further details in
 - Russinovich, Solomon & Ionescu, *Windows Internals*, 5th or 6th edition

IPC Example: (A)LPC in Windows

❑ Sketch of communication:

- The client opens a handle to the subsystem's connection port object.
- The client sends a connection request.
- The server creates two private communication ports and returns the handle to one of them to the client.
- The client and server use the corresponding port handle to send messages (or callbacks and to listen for replies).



Client-Server Communication

- ☐ Shared Memory Communication
- ☐ Message Passing

- ☐ Additional strategies suitable for client-server systems
 - Higher level abstractions?

- ☐ Sockets
 - Stream/packets of data
- ☐ Remote Procedure Calls (RPC)
 - Abstract from procedure calls
- ☐ Pipes
 - Stream of data

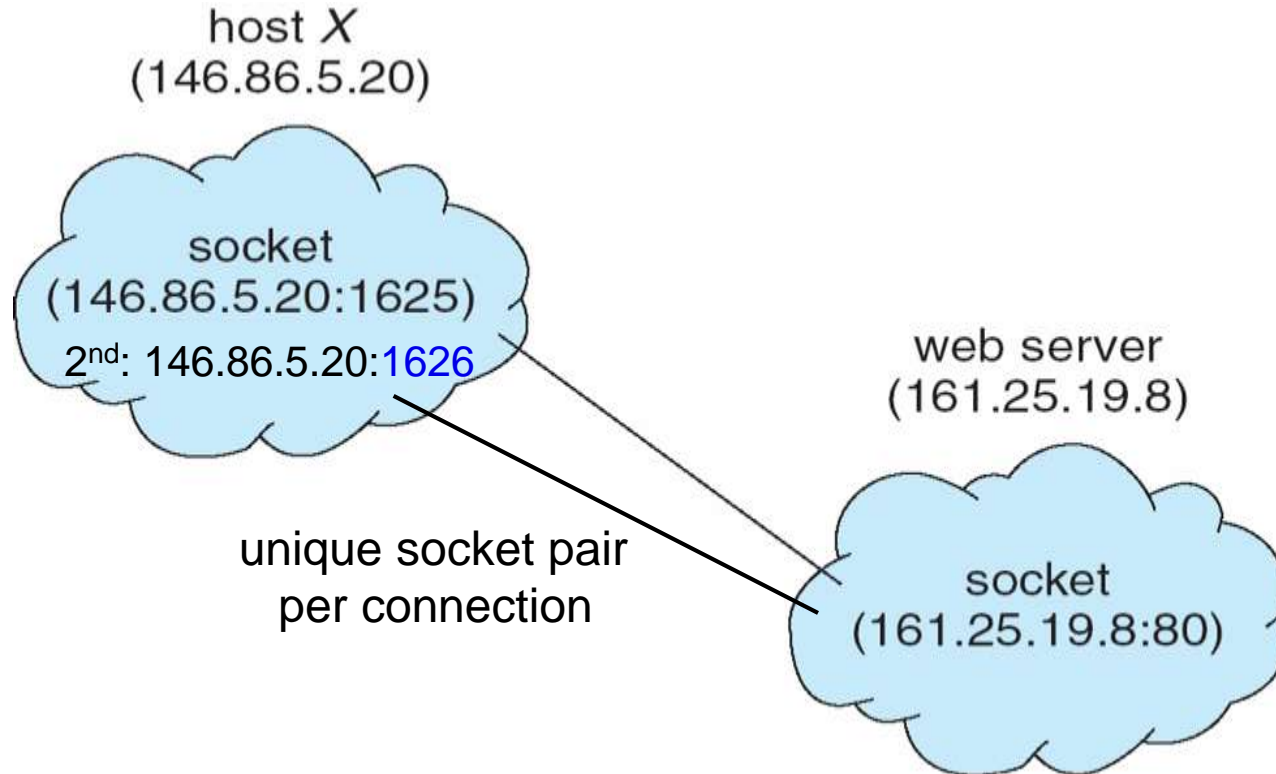
Sockets

- ❑ **Socket:** an endpoint for communication
 - A pair of communicating processes employ sockets at each end (1 per process)
 - Server listens (waits) for incoming client requests
 - different kinds of sockets → here: network (Internet) sockets

- ❑ Concatenation of IP address and port
 - IP identifies host
 - port identifies process on host
 - Example: the socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8

- ❑ Communication occurs between socket pairs using some communication protocol
 - Connection-oriented (e.g. TCP)
 - Connection-less (e.g. UDP)

Socket Communication Example



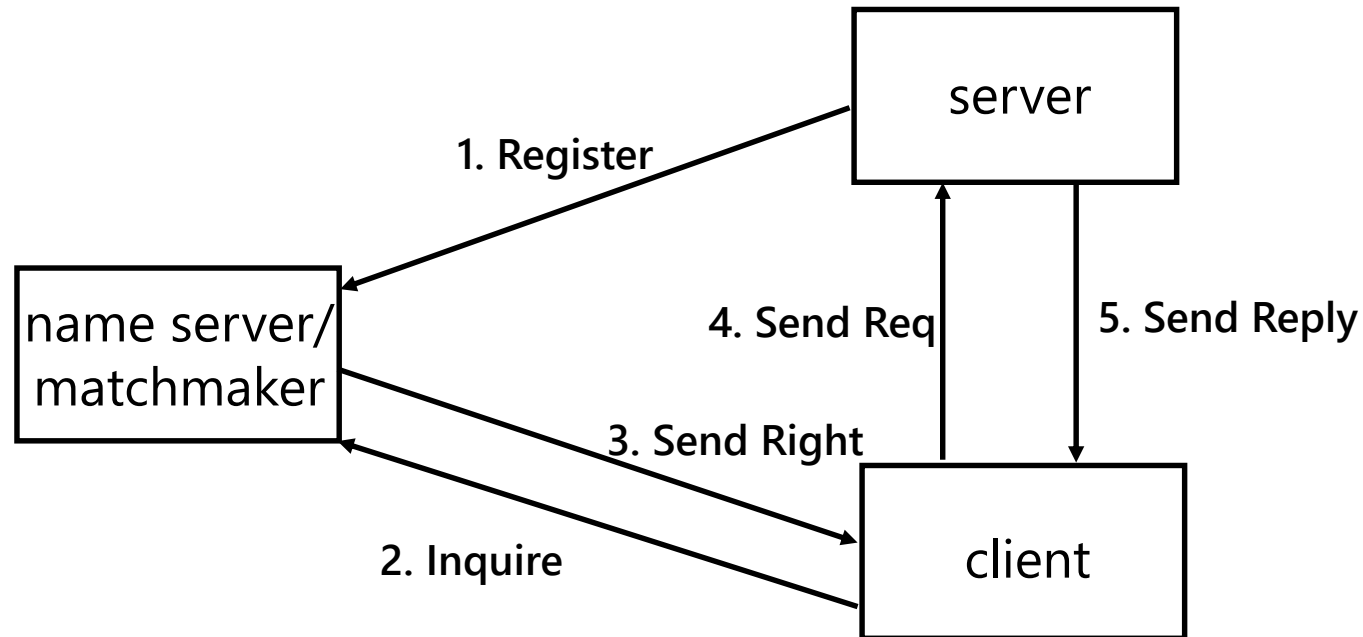
© Silberschatz et al.: Operating System Concepts, 8th ed., Wiley

- host X requests webpages
 - 1st request → socket pair: 146.86.2.20:1625, 161.25.19.8:80
 - 2nd request → socket pair: 146.86.2.20:1626, 161.25.19.8:80

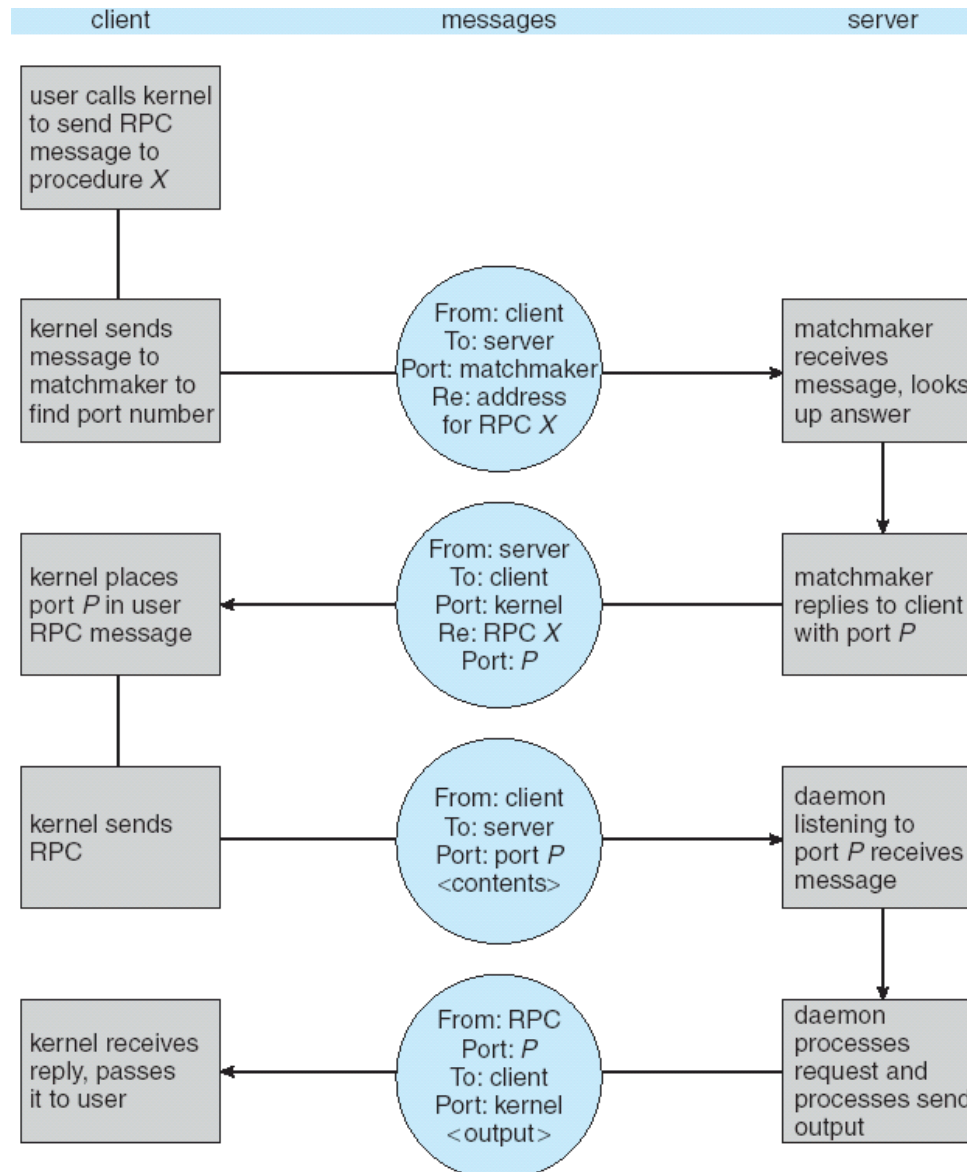
Remote Procedure Calls (RPC)

- ❑ **Remote procedure call:** abstracts procedure calls between processes on networked systems.
- ❑ **Stubs** – client-side proxy for the actual procedure on the server.
- ❑ The client-side stub locates the server and marshals the parameters → message.
- ❑ The server-side stub receives this message, unpacks the marshaled parameters, and performs the procedure on the server.
- ❑ Susceptible to communication faults
 - exactly once and at most once semantics
- ❑ How does a client know the server address/socket?

Client-Server Models Behind RPC



Execution of RPC



© Silberschatz et al.:
Operating System Concepts,
8th ed., Wiley

Pipes


❑ **Pipe:** conduit, allowing two processes to communicate

❑ Implementation issues:

- Unidirectional or bi-directional communication?
- Half-duplex or full-duplex (for bi-directional mode)?
- Relationship between communicating processes (e.g. parent/child)?
- Communication via network or only locally on same machine?

❑ Ordinary pipe (anonymous)

- Unidirectional (one-way)
- Created using `pipe()` system call
- Two file descriptors:
 - `fd[0]`: read-end
 - `fd[1]`: write-end
- Not accessible outside the creating process



© Silberschatz et al.: Operating System Concepts, 8th ed., Wiley

Summary

❑ Processes

- Definition, PCB
- States (& changing states)
- Creation, switching, termination

❑ Threads

❑ Inter-process communication

- Shared memory
- Message passing
 - Links
 - Addressing, synchronization, buffering
- Client-Server
 - Sockets
 - RPC
 - Pipes

Hierarchical Microkernel

Hierarchical Microkernel: Introduction

❑ Goals:

- Make the system flexibly adjustable (for students to play)
- Make the system robust (for students to play & other reasons)

❑ Flexible adjustability:


Easily exchange core OS functionality

- At compile time
- At runtime
- Examples: scheduler, memory manager, process loader, ...

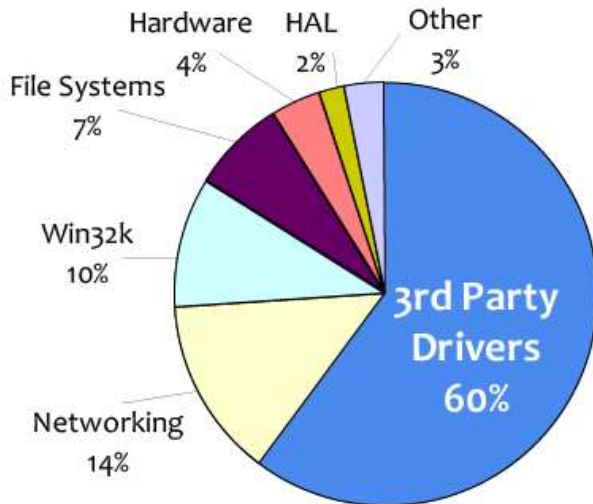
❑ Robustness:

- Easier debugging
- Prevent system failure caused by single OS component

The Driver Problem: Windows/Linux Example



Causes of WinXP outages



Causes of Win2k outages

❑ Device drivers


- **Numerous:** 250 installed (100 active) drivers in XP/Vista
- **Large & complex:** 70% of Linux code base
- **Immature:** every day 25 new / 100 revised versions Vista drivers
- **Access Rights:** kernel mode operation in monolithic OSs

❑ Device drivers are the **dominant cause** of OS failures despite sustained testing efforts

❑ Situation similar in Linux and Windows

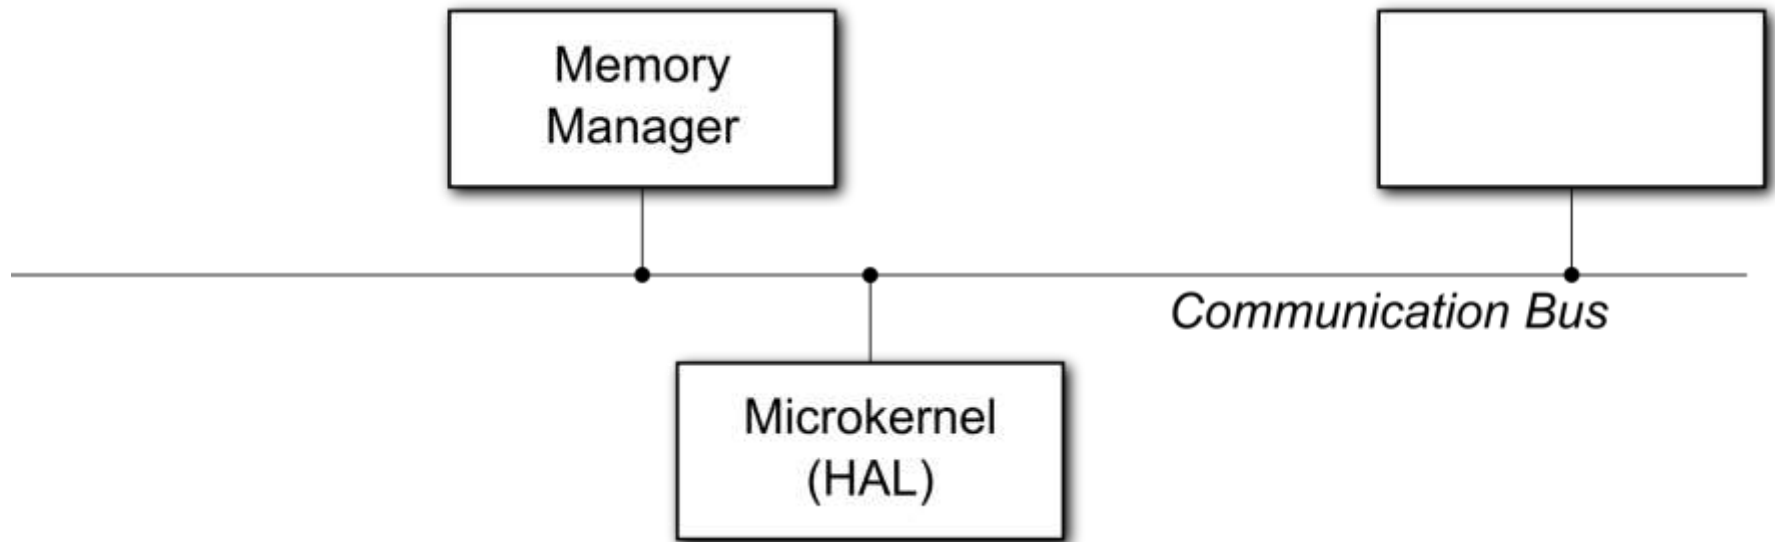
- ❑ Graph illustrates Windows crash reasons

Hierarchical Microkernel: First Sketch



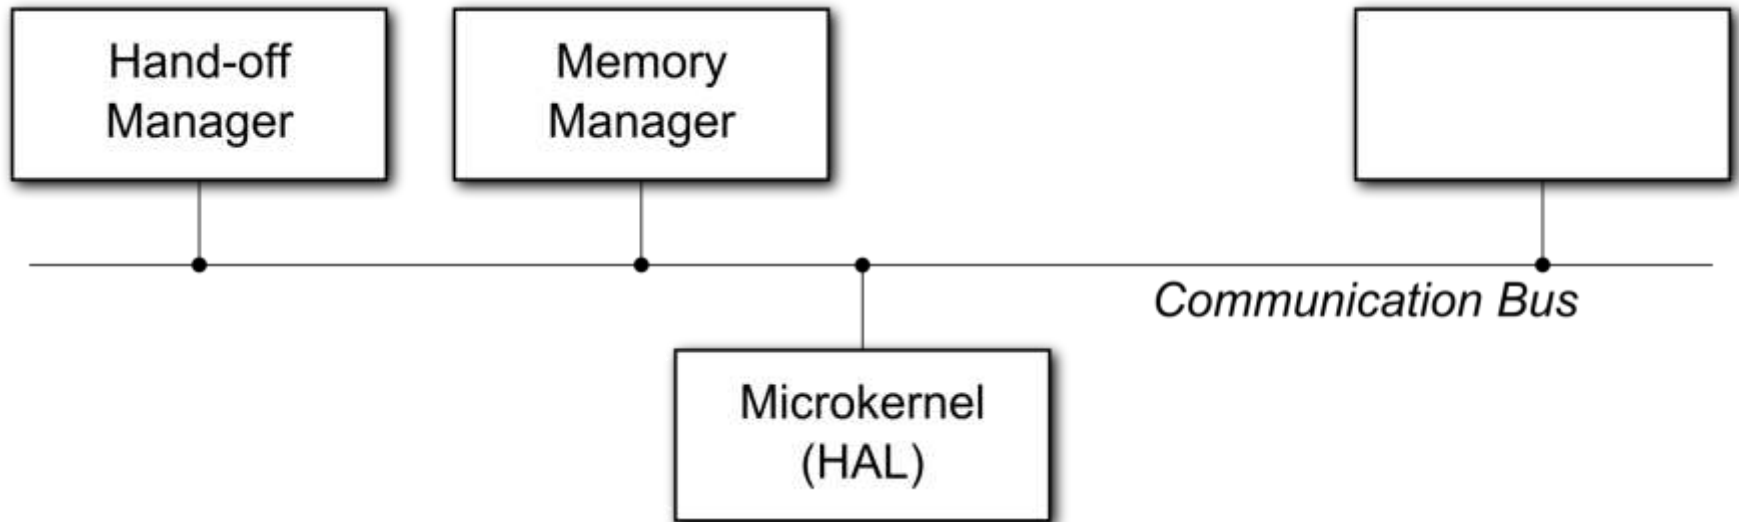
- ❑ “Ultra micro” microkernel
 - provides only hardware abstraction layer (HAL)
 - executes in supervisor mode
- ❑ Isolated “modules”
 - all user-specific and system functions
- ❑ Communication bus
 - simple send/receive interface
 - broadcast communication

Hierarchical microkernel: Flexible Module Exchange



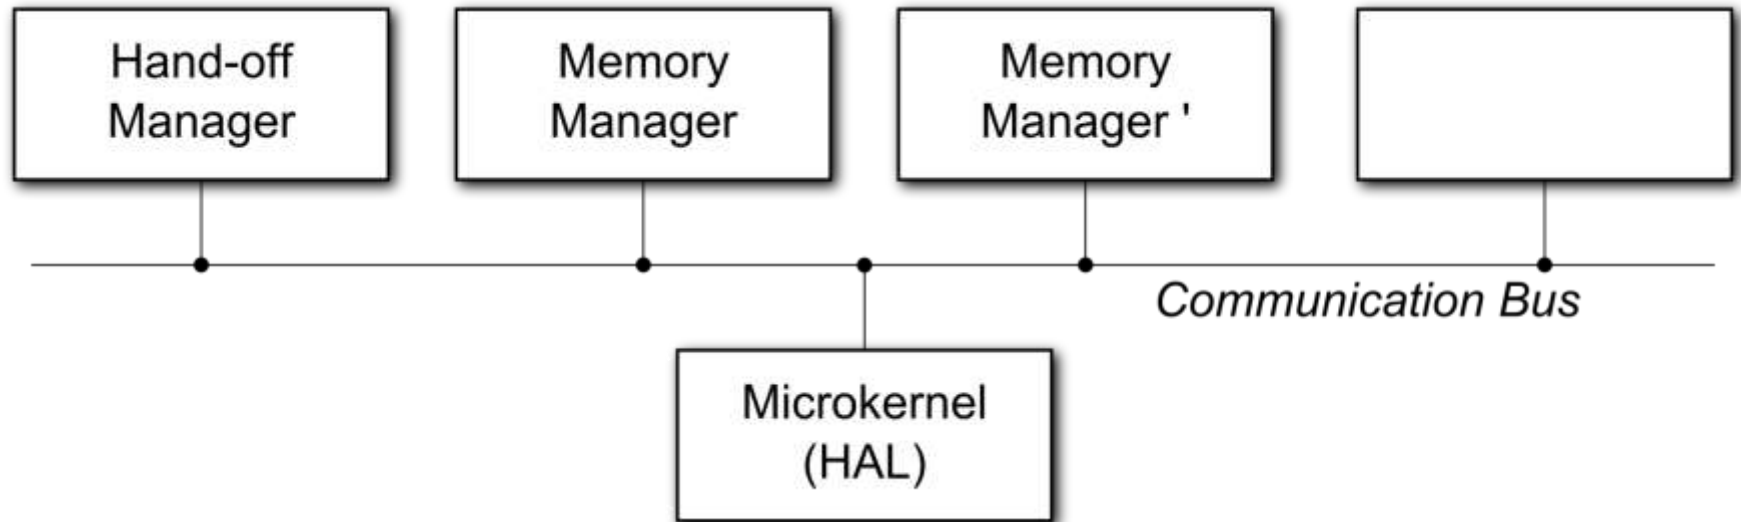
- ❑ Example: memory manager exchange
 - microkernel receives memory manager exchange request

Hierarchical microkernel: Flexible Module Exchange



- ❑ Example: memory manager exchange
 - microkernel receives memory manager exchange request
 - microkernel loads hand-off manager (via separate loader module)

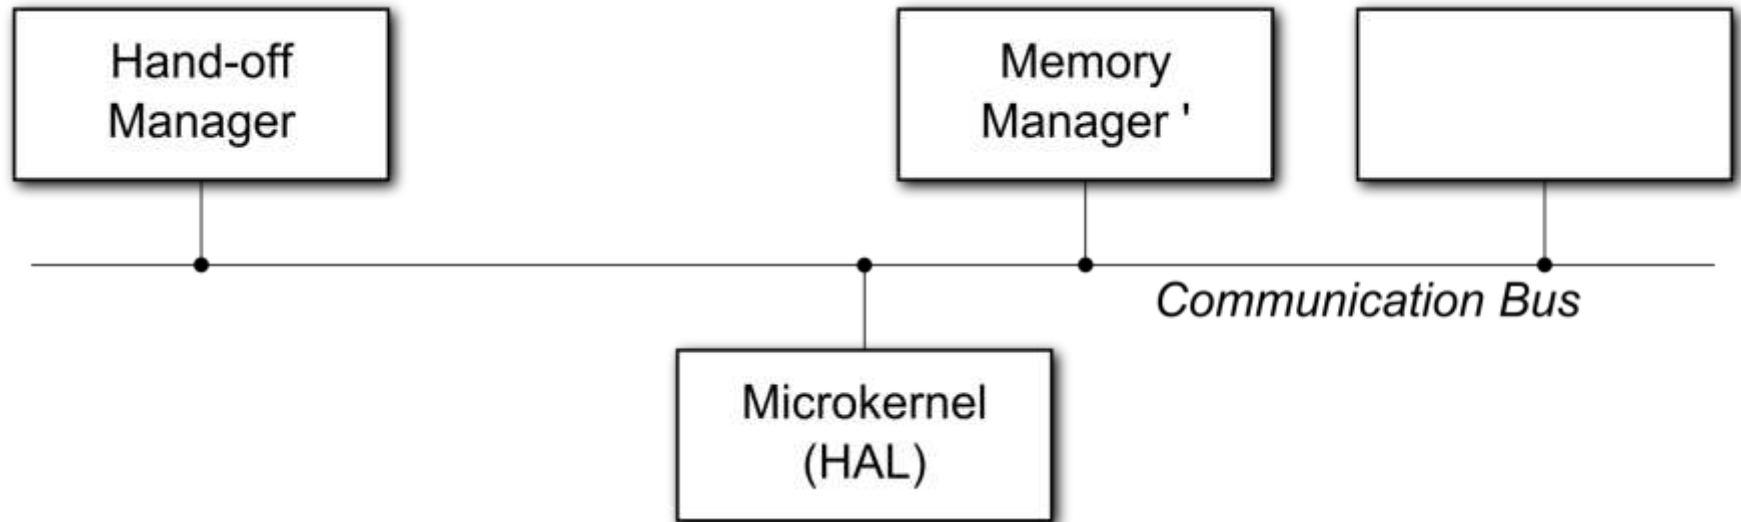
Hierarchical microkernel: Flexible Module Exchange



❑ Example: memory manager exchange

- microkernel receives memory manager (mm) exchange request
- microkernel loads hand-off manager (hom) (via separate loader module)
- hom (indirectly) loads mm'
- hom (indirectly) stops mm and queues all subsequent requests to mm that traverse the bus
- hom initiates state transfer from mm to mm'


Hierarchical microkernel: Flexible Module Exchange



❑ Example: memory manager exchange

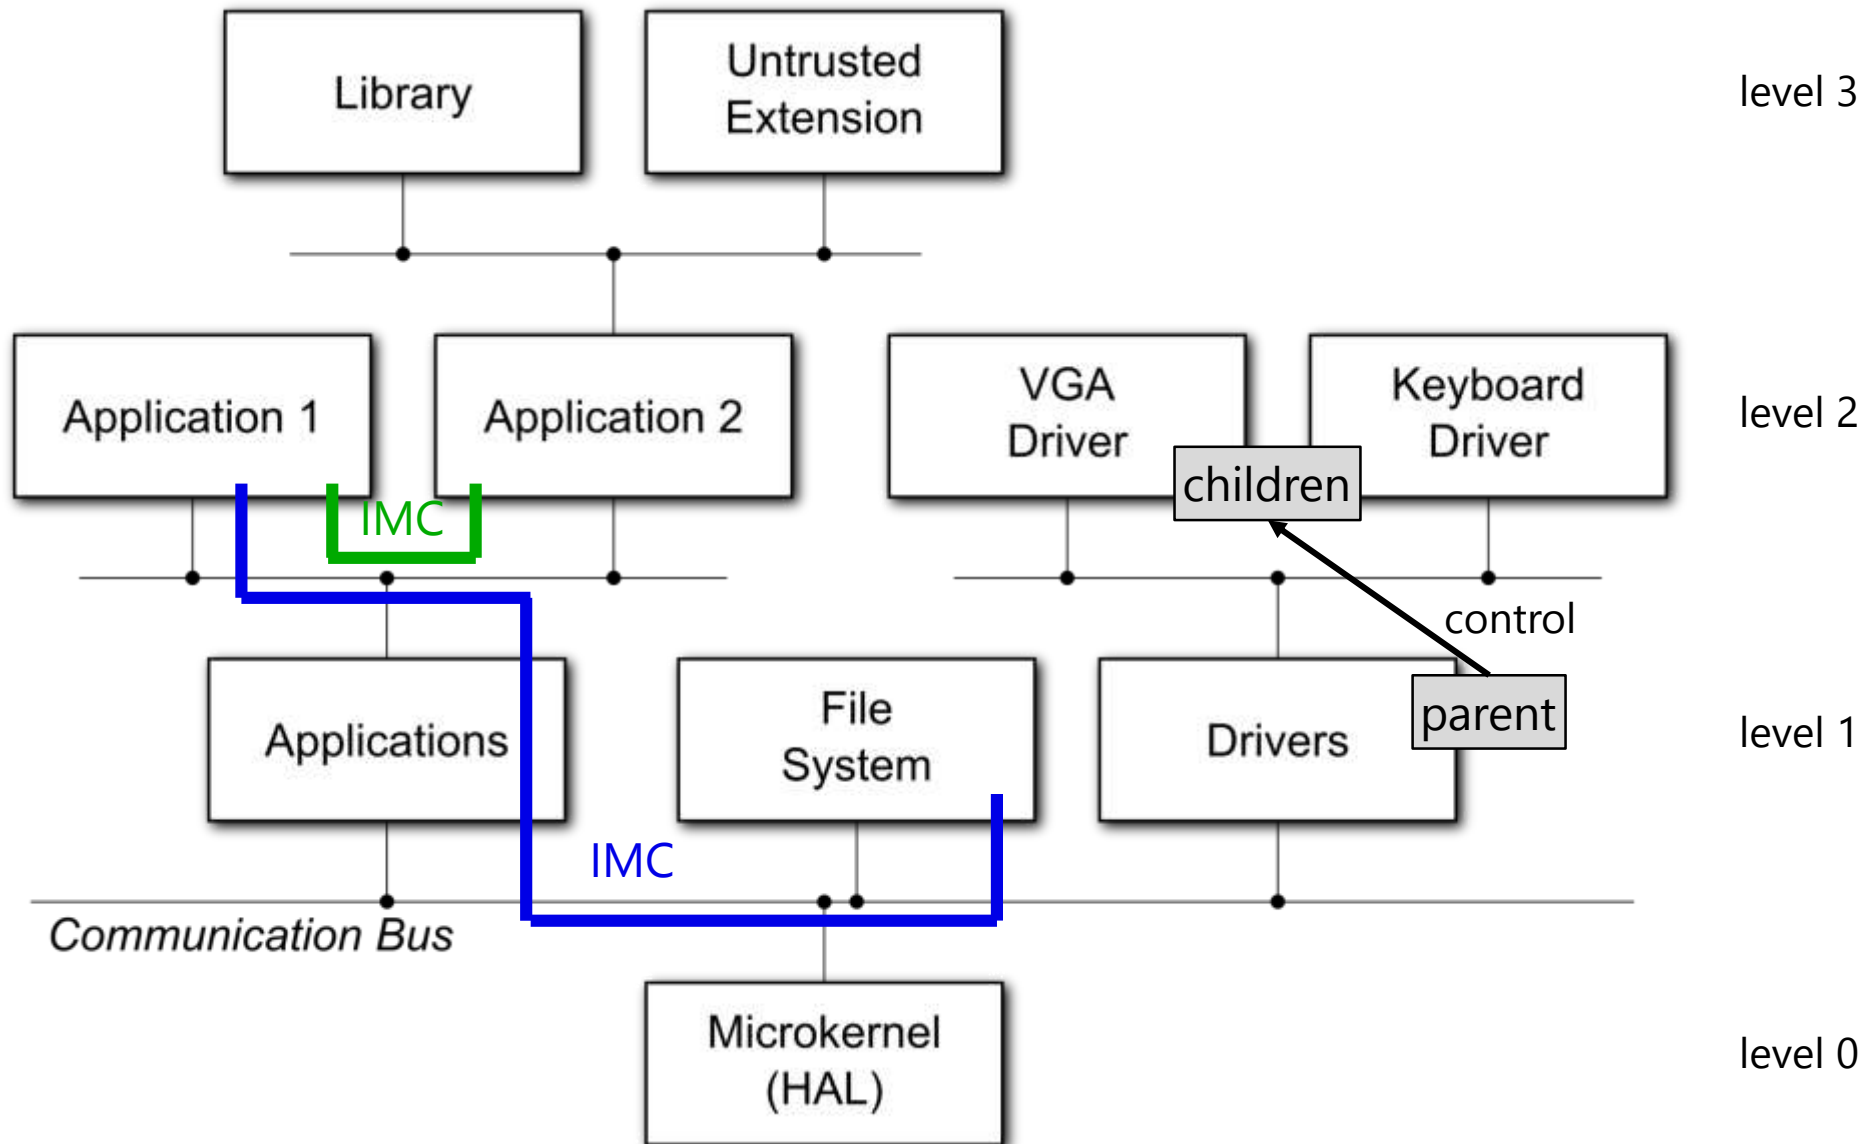
- microkernel receives memory manager (mm) exchange request
- microkernel loads hand-off manager (hom) (via separate loader module)
- hom (indirectly) loads mm'
- hom (indirectly) stops mm and queues all subsequent requests to mm that traverse the bus
- hom initiates state transfer from mm to mm'
- hom removes mm, starts mm', replays all queued requests for mm

Hierarchical Microkernel: Bus Downsides



- ☐ Flexibility (and simple implementation of fault-tolerance mechanisms) through “public” communication
- ☐ Congestion?
- ☐ Single point of failure (SPOF)?
- ☐ Denial of service (DOS) attacks?

Hierarchical microkernel: Hierarchy



Hierarchical Microkernel Wrap-up

❑ Modules instead of processes

⇒ inter-module communication (IMC) instead of IPC ☺

⇒ Is there a difference?

❑ Strictly message-passing IMC

- also for “system calls”
- multicast (bus-local broadcast)
- bus implementation customizable!
 - synchronous vs asynchronous
 - buffering vs rendezvous
 - message ordering characteristics/guarantees
- modules can implement custom protocols on top of the busses they are attached to