


# Resource Sharing & Inter-Process Interactions

- **The Issues:** Critical Sections, Mutual Exclusion, Deadlocks...
- **Task Orderings**: Scheduling issues and solutions
- **Algorithmic Solutions:** Races, Ordering, Alternations...
- **Program Level Solutions:** Semaphores, Monitors

# Alternatives?

## ❑ Algorithms to allocate a resource

- E.g., let's give resources to shortest job first
- Works great for multiple short jobs in a system
- May cause indefinite postponement of long jobs even though not blocked (starvation?)



## ❑ Other Solutions?

- First-come, first-served (FIFO) policy
- Shortest job first, highest priority, deadline first etc...

## Scheduling Policies!

... efficient, fair, deadlock & race-free


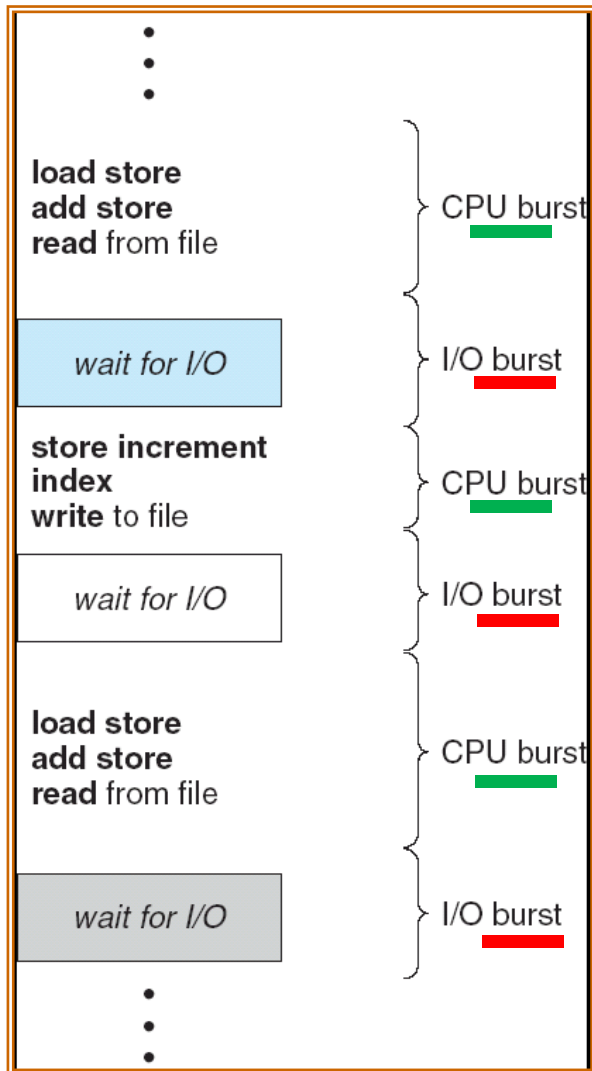
# OS Scheduling Criteria (Optimizations)

- **CPU utilization:** keep CPU as busy as possible (maximize) ↑
    - Nice target but the CPU better be doing something useful!  
Remember livelock?
  - **Throughput:** # of processes completed/time unit (maximize) ↑
  - **Turnaround time:** average amount of time to execute a particular process (minimize) ↓
  - **Waiting time:** process wait time in ready queue (minimize) ↓
  - **Response time (minimize):** amount of time it takes from when a request was submitted until the first response (**not** output) is produced. [result output time = RT deadline] ↓
- Efficiency → Optimize all 😊 ...plus want Fairness!  
...and across very diverse applications!

# Scheduling Variants

- A. First Come, First Served (**FCFS**)
- B. Shortest Job First (**SJF**)
- C. Round Robin (**RR**)
- D. Priority Based (**PB**)

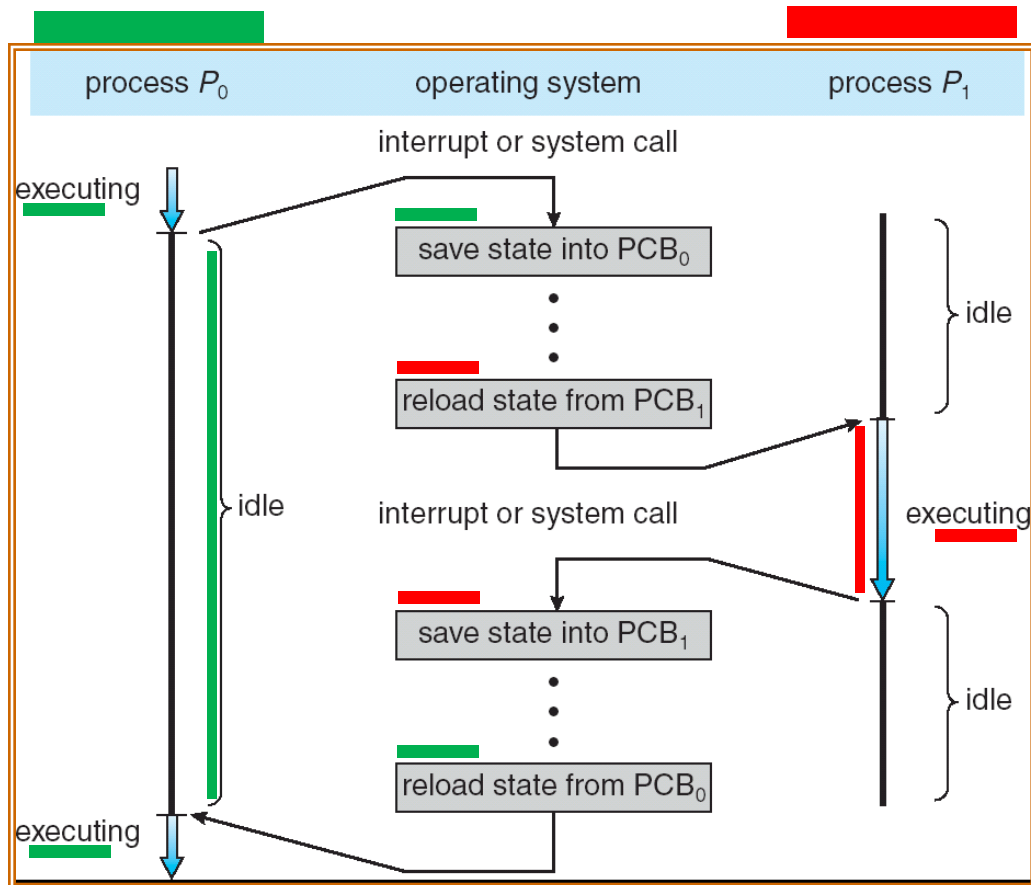
# Issue 1. No Single Process/CPU Usage Model



- Apps “typically” stay CPU or I/O-bound
- Hybrids “typically” for interactive apps

# Issue 2. Process/Thread Preemptability


- ❑ **Non preemptive (NP):** An ongoing task **cannot** be displaced
- ❑ **Preemptive:** Ongoing tasks **can** be switched in/out as needed



# A: First-Come, First-Served (FCFS) Scheduling (Non-Preemptive)

<u>Process</u>	<u>Execution Length (CPU Burst Time)</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Processes arrive (& get executed) in their arrival order:  $P_1, P_2, P_3$   
Single queue of “ready” non-preemptible processes




- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$

# FCFS Scheduling (NP: Non-Preemptive)

Suppose that the processes arrive in a different order as:

$P_2, P_3, P_1$  [3,3,24]


- Process Order/Schedule:




- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Variability of wait time (from 17 to 3!) – control? response prediction?



# Scheduling in Batch Systems where we can actually see and re-order FCFS?



(a)



(b)

Shortest job first (SJF) scheduling?

# B: Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of expected CPU burst
- Use these lengths to schedule the process with the shortest time

## Options:


- **non-preemptive** – once CPU given to the process it cannot be preempted until it completes its CPU burst
- **preemptive** – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt.

Variation: Shortest-Remaining-Time-First (SRTF)

# Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>		<u>Burst Time</u>
$P_1$	0.00	0.0	7
$P_2$	0.0+	2.0	4
$P_3$	0.0+	4.0	1
$P_4$	0.0+	5.0	4

- SJF (non-preemptive)  $P_1$ , then  $P_3$  then  $P_2$ ,  $P_4$  (FCFS across equals)



- Average waiting time =  $(0 + 8 + 7 + 12)/4 = 6.75$  (fixed arrival)
- Av. waiting =  $(0 + [8-2] + [7-4] + [12-5])/4 = 4.00$  (staggered arrival)

# Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
- Use these lengths to schedule the process with the shortest time

## Options:

– **non-preemptive** – once CPU given to the process it cannot be preempted until it completes its CPU burst

➤ **preemptive** – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt.


Variation: Shortest-Remaining-Time-First (SRTF)

# Preemptive SJF

(Shortest Remaining Time First: Re-assess SRT on arrivals)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (preemptive)



- Average waiting time =  $P_1:[0, (11-2)]$ ;  $P_2:[0, (5-4)]$ ;  $P_3: 0$ ;  $P_4: (7-5)$
- Average waiting time =  $(9 + 1 + 0 + 2)/4 = 3 \dots [6.75; 4.00]$
- Dynamic scheduling (SRT) calculations + context switches → **cost !!!**

# Determining Length of Next CPU Burst

- Can only estimate the task length (simple for batch jobs)
- Use the length of previous CPU bursts (history logs & ageing)

1.  $t_n$  = actual length of  $n^{th}$  CPU burst

2.  $\tau_{n+1}$  = predicted value for the next CPU burst

3.  $\alpha, 0 \leq \alpha \leq 1$ ;  $\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$

4. Define :  $\tau_{n+1} = \tau_n$  ( $\alpha = 0$ ) direct fn. of prior prediction

$\tau_{n+1} = t_n$  ( $\alpha = 1$ ) direct fn. of most recent run

$\alpha = 1/2$

$\tau_1 = 1/2 t_0 + 1/2 \tau_0$  add + simple right shift ( $\rightarrow$  divide by 2)

$\tau_2 = 1/2 t_1 + 1/2 \tau_1 = 1/2 t_1 + 1/2 (1/2 t_0 + 1/2 \tau_0) = 1/2 t_1 + 1/4 t_0 + 1/4 \tau_0$

$\tau_3 = 1/2 t_2 + 1/2 \tau_2 = 1/2 t_2 + 1/4 t_1 + 1/8 t_0 + 1/8 \tau_0$

# Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
- Use these lengths to schedule the process with the shortest time

Options:

- non-preemptive – once CPU given to the process it cannot be preempted until completes its CPU burst
- preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. Shortest-Remaining-Time-First (SRTF)

- ❑ SJF is optimal – gives minimum average waiting time for a given set of processes
- ❑ SJF can still lead to starvation: if new short jobs keep coming in and getting considered in the “shortest” ranking, a “long” job can potentially starve! [Convoy Effect]

# Variations

✓ SJF → EDF (Earliest Deadline First)

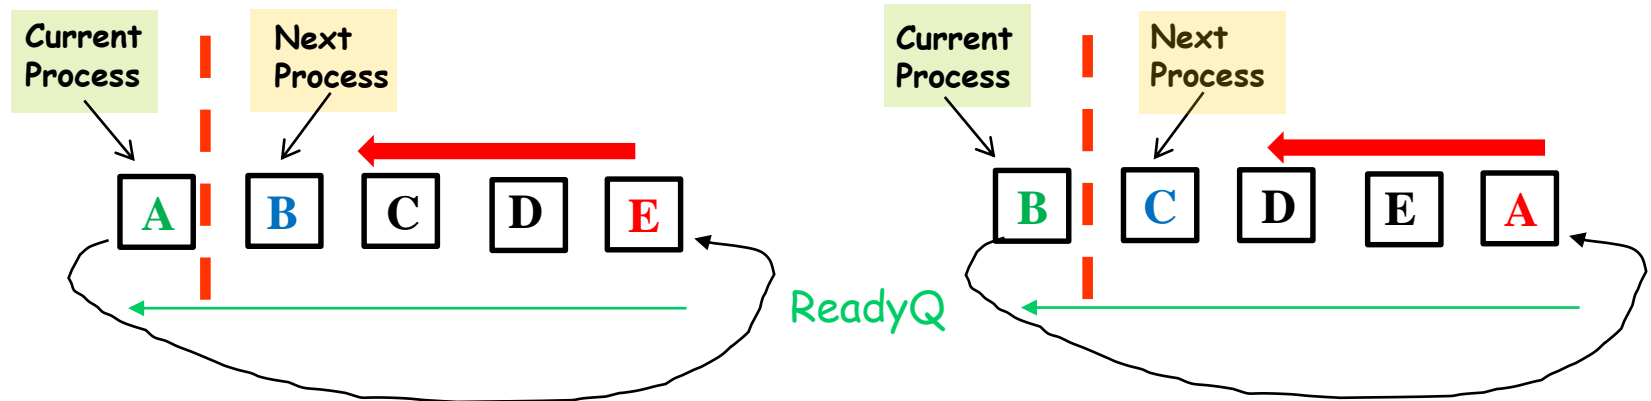


# Onwards to Preemptive Scheduling

## Preemptive FCFS → Round Robin


# C: \*Round Robin (RR/TDMA): Fixed Slots

- Serial ReadyQ of processes (Circular Queue)
- Each process gets a fixed slice of CPU time (*time quantum:  $q$* ), 10-100ms
  - Slot finishes  $\rightarrow$  process “**preempted**”, added to the end of Q: **Sliding Window**
  - $n$  processes in ReadyQ and time quantum is  $q \rightarrow$  each process gets  $1/n$  of CPU time in chunks of  $q$  time units. No process waits more than  $(n-1)q$  time units.



- FIFO Performance?
  - $q$  too large  $\Rightarrow$  poor response time & poor CPU utilization – wasted time
  - $q$  too small  $\Rightarrow$  context switching.  $q$  must be large wrt context switch overhead, else context switching overhead reduces efficiency


# RR with Fixed/Static time Quantum = 20



# Better: Dynamic RR with Quantum = 20

Dynamic Q: if a process finishes in less than 20, then start the next slot earlier!

<u>Process</u>	<u>Burst Time</u>
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24



- Typically, higher average turnaround than SJF, but better *response*
  - Turnaround time** – amount of time to execute a particular process (*minimize*)
  - Response time** (*min*) – amount of time it takes from request submission until first response
- Scheduling is static though calculating/estimating “right” length of quantum is crucial!!!
- Alternate: Have quanta of varying lengths!!! (dynamic scheduling costs ~ SJF)

# Input Queue Policy?

- **Static** – Preemptive FCFS scheduling with predictable wait time
- **Dynamic?**
  - If Q keeps adding new jobs → slot waiting times can vary (extend!) for existing jobs leading to potential starvation (Admission Control)

# D: \*Priority Scheduling

- ❑ A priority number (integer) is associated with each process
- ❑ The CPU is allocated to the process with the highest priority

Generally, **Highest = 's Biggest #** ... BUT in Unix: smallest integer  $\equiv$  highest priority

- Preemptive (priority and order changes possible at runtime)
- Non-preemptive (initial execution order chosen by static priorities)

- ❖ FCFS is priority scheduling where arrival order determines priority
- ❖ SJF is priority scheduling where CPU burst time determines priority
- ❖ RR is equal priority scheduling (albeit in FCFS order)

- ❖ Problem  $\equiv$  Starvation: low priority processes may never execute
- ❖ Solution  $\equiv$  **Ageing**: as time progresses increase priority of process

# Ageing and Priority

- Priority can be static or dynamic (as a function of wait-time & usage-time for fairness etc: Unix BSD, Linux)

**Note: In BSD → Large # = 's Low Priority!**

$$P_{\text{user-pri}} = P_{\text{USER}} + P_{\text{cpu}}/4 + (2 * P_{\text{nice}})$$

- ❑ Process **using** CPU: **P\_cpu increases** with time → **P\_user-pri goes up** (ie lower priority)
- ❑ Process **waiting**: **P\_cpu keeps decreasing** → **P\_user-pri keeps going down** (ie higher priority)

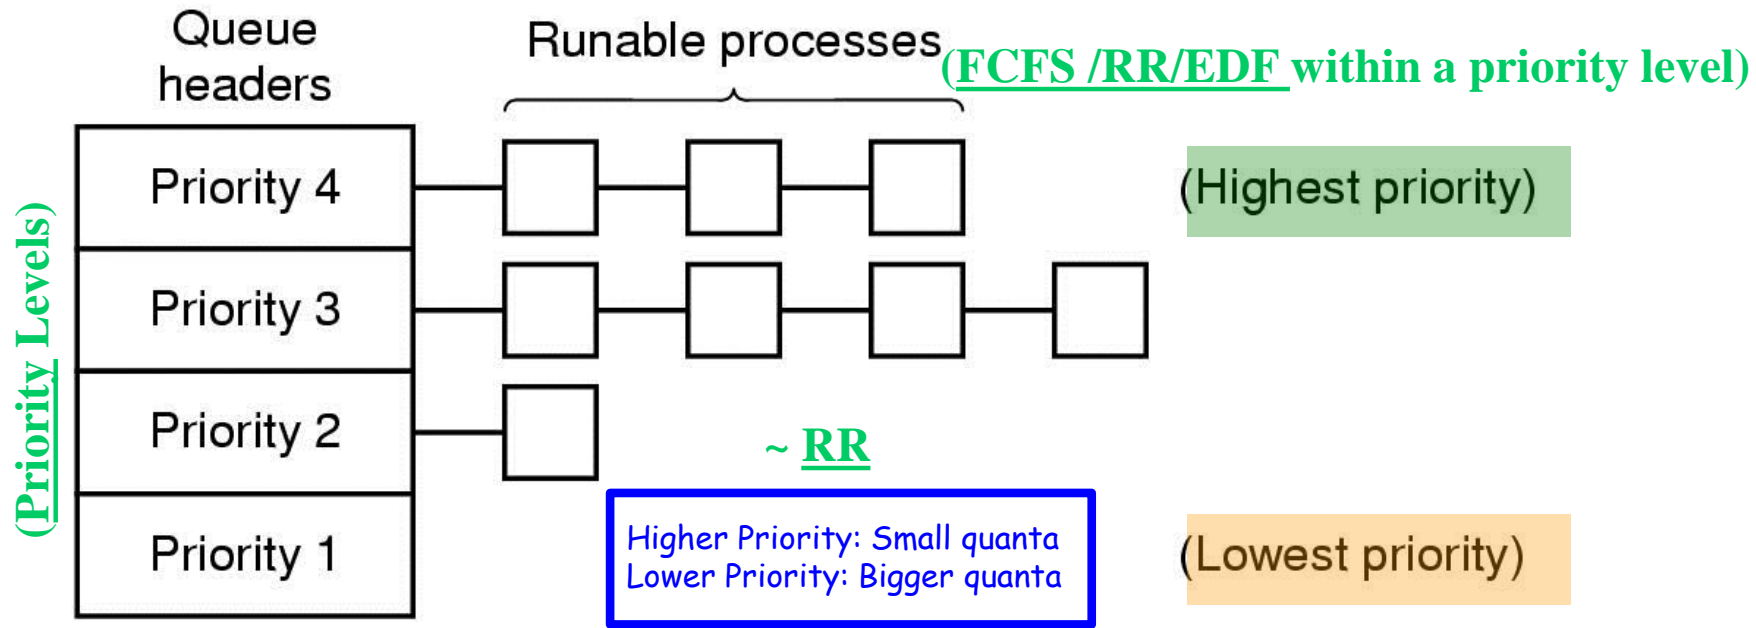
# Which Scheduling Policy to Use?

- First Come, First Served (**FCFS**)
- Shortest Job First (**SJF**)
- Round Robin (**RR**)
- Priority Based (**PB**)

- OS schedulers implement \*all\* scheduling policies
- Choice of policy is based on the nature of specific applications (computing, I/O, multimedia etc)
- Mixture of applications (at kernel- level, at user-level) determines the specific scheduling policy or groups of policies to use!



# Scheduling Options (Solaris, Windows, Linux)




- ❑ Kernel process: “generally” non-preemptive
- ❑ User processes: “mostly” preemptive

# Where Should the Scheduler be?

- ☐ Kernel level?
- ☐ User level?

# Thread Scheduling (Kernel Scheduler)



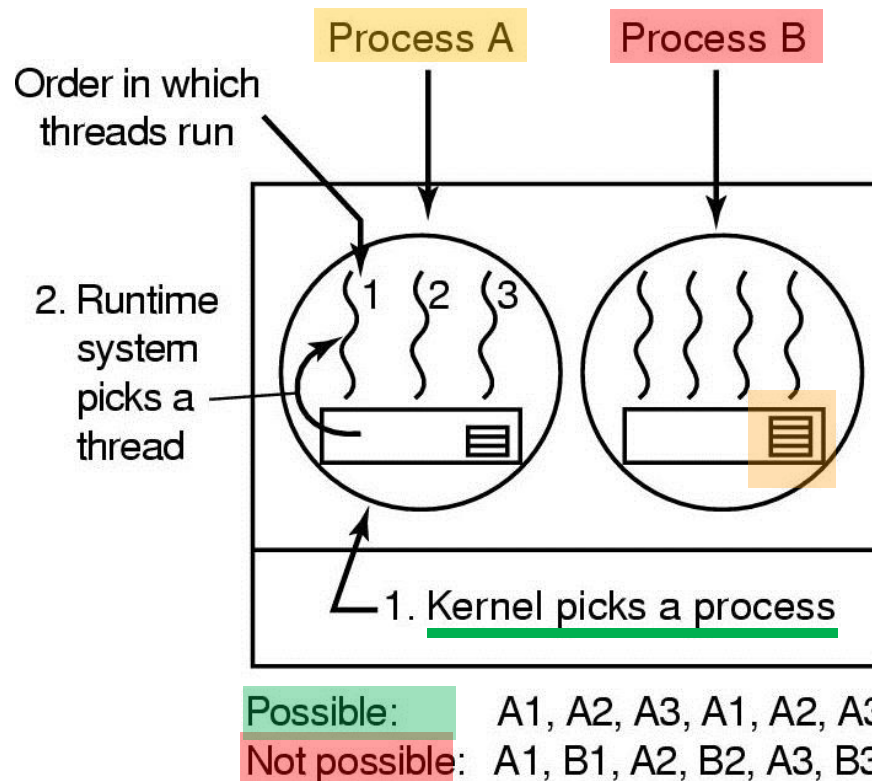
Possible: A1, A2, A3, A1, A2, A3

Also possible: A1, B1, A2, B2, A3, B3

## Possible scheduling of kernel-level threads

- Flexible & efficient for quanta usage...but costly full context switch for threads
- Monolithic kernel scheduler also avoids thread blocks (on I/O etc)

# Thread Scheduling (User Scheduler)




## Possible scheduling of user-level threads

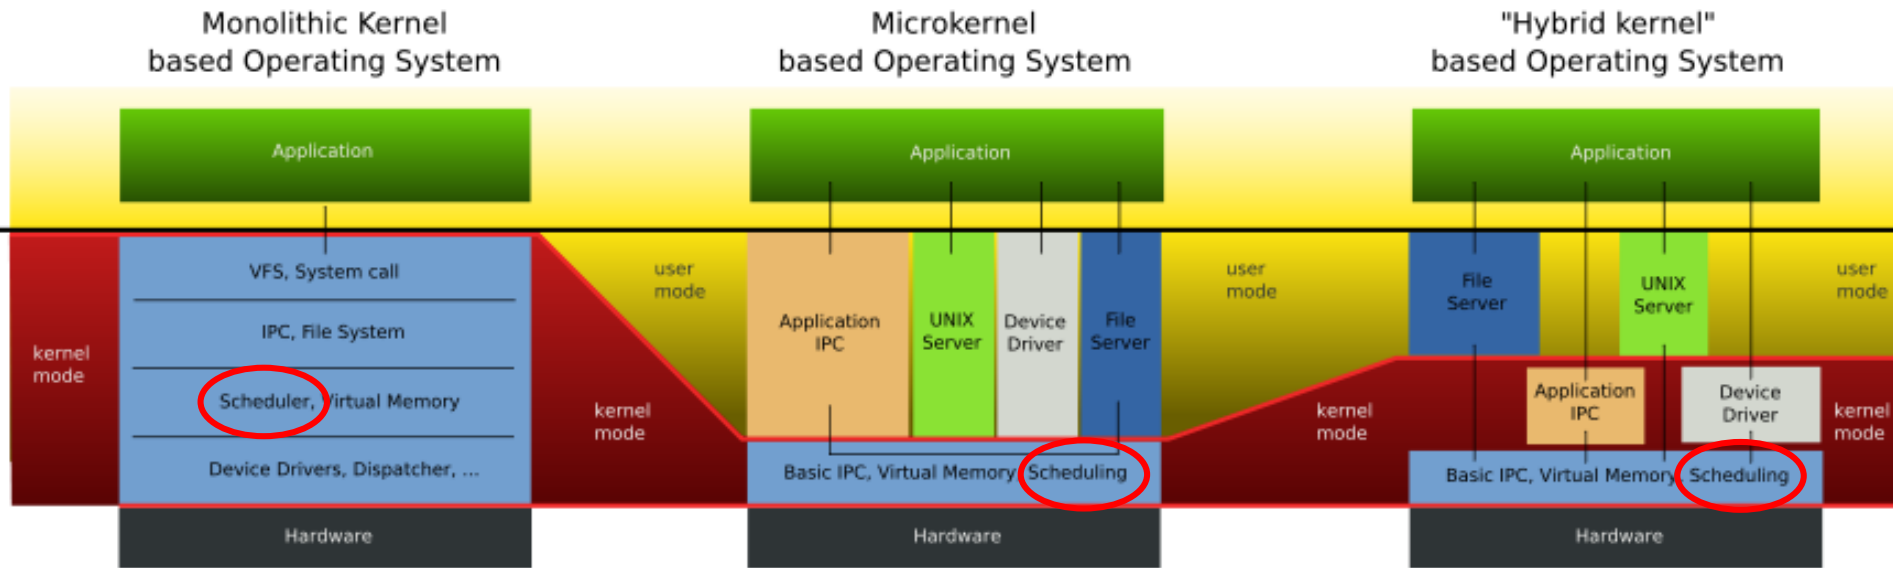
- Simple process level context switch; thread block → process blocks
- Split level scheduling (kernel + dedicated application/thread type specific)

# Thread Scheduling

- ❑ Local Scheduling: Threads library decides which thread to put onto an available LWP
- ❑ Global Scheduling: Kernel decides which kernel thread to run next



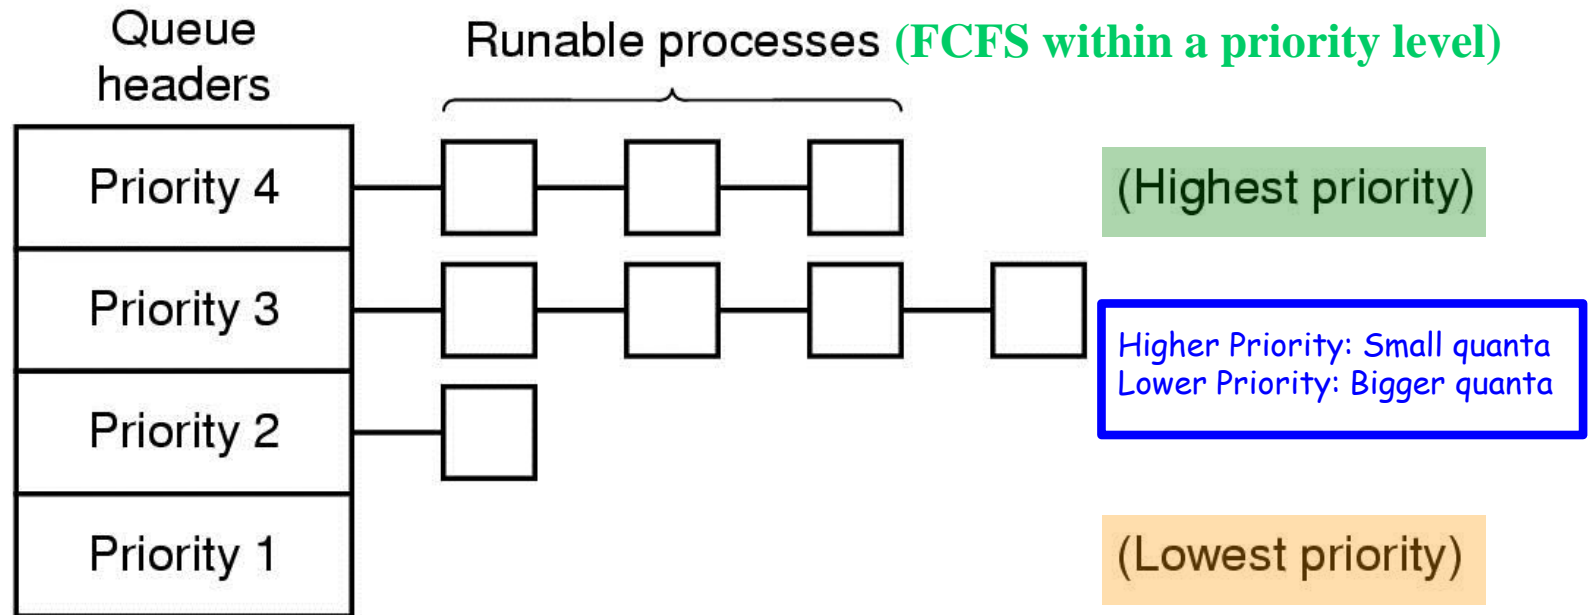
# Scheduler Locations?



# OS Examples

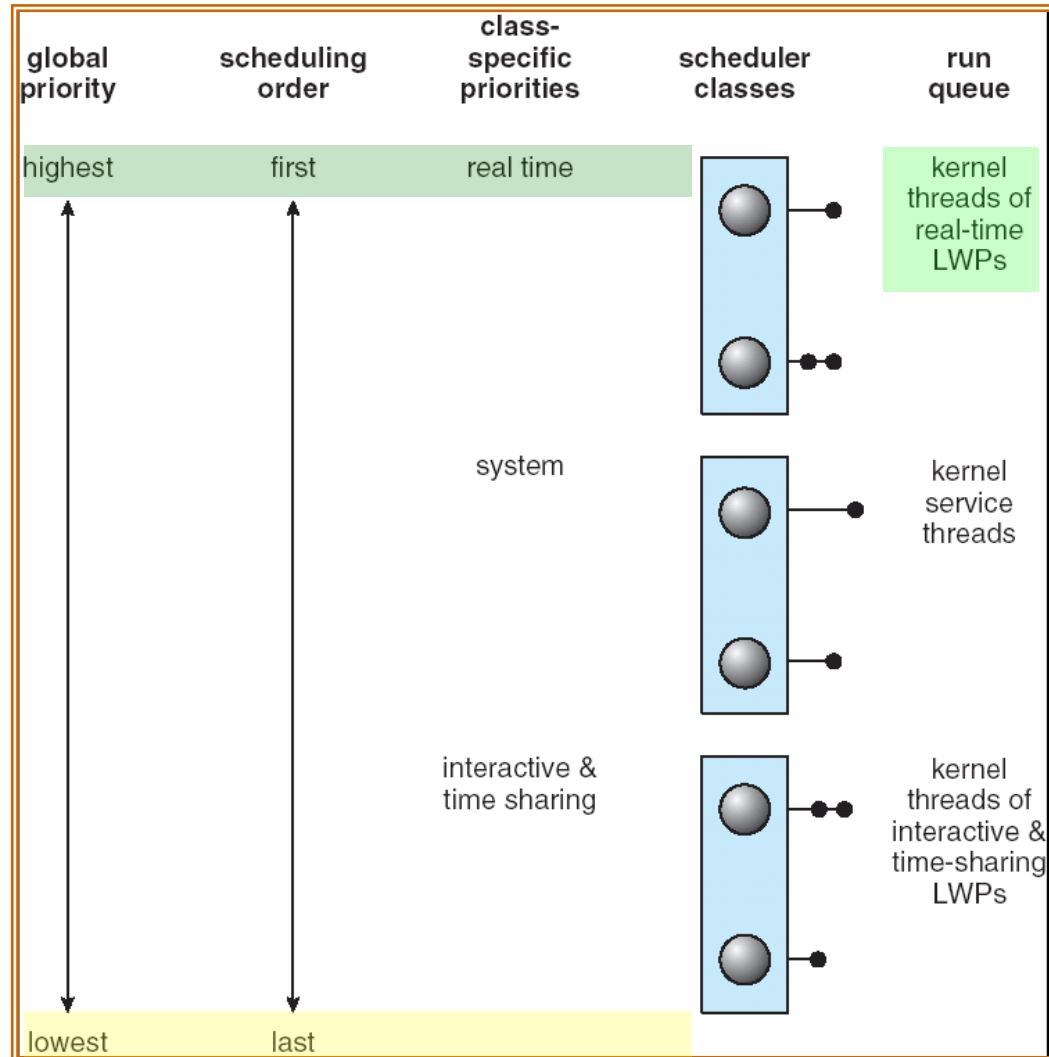
- ❑ Solaris scheduling
- ❑ Windows scheduling
- ❑ Linux scheduling

- All work with dynamically altering priorities & time slices
- All have preemptive tasking



# Linux/Solaris Scheduling

Higher the priority,  
shorter the time slice





# Solaris Dispatch Table

		TQE	RFS
priority	time quantum(msec)	time quantum expired	return from sleep
0	200	0	50
5	200	0	50
10	160	0	51
15	160	5	51
20	120	10	52
25	120	15	52
30	80	20	53
35	80	25	54
40	40	30	55
45	40	35	56

← New dynamically altered priorities  
...and time slots

CPU-tasks mid-priority,  
Interactive: higher priority

\* Lowest priority 0 (diff. from BSD) is given largest quantum

**TQE**: amount of quanta used without blocking

(CPU Intensive Tasks) ; Solaris 9+ : No TQE (fixed priority & CPU shares)

**RFS**: (raised) priority of awakened tasks (responsiveness for I/O Tasks)

# Windows: Priority Based, Preemptive

(Fixed & Variable Components)

## priority classes

	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

relative priority within a priority class (High # = 's High Priority; diff. from BSD)

- Fixed: Priority Class (eg RT) will always run when invoked
- Variable: Thread runs till quanta finished; priority lowered; task coming in from "wait" gets priority raised – but only within its priority class!

# Linux Scheduling (0=highest)

## ❖ Two algorithms: time-sharing and real-time

### ✓ Time-sharing

- Prioritized credit-based – process with most credits is scheduled next
- Credit subtracted when timer interrupt occurs
- When credit = 0, another process chosen
- When all processes have credit = 0, re-crediting occurs
  - Based on factors including priority and history

### ✓ Real-time

- Soft real-time
- Two classes
  - FCFS and RR
  - Highest priority process always runs first