

# ME Solutions: Today

- Memory-sharing
  - All previous primitives/mechanisms are memory-sharing
  - Semaphores as *mutex* and as *synchronization primitive*
  - Monitors
- Message-passing solutions
- Barriers
- Classical ME / concurrency problems
  - Access problems (Readers/Writers Problem)
  - Synchronization problems (Dining Philosophers Problem)
  - Scheduling (Sleeping Barber Problem)

# Semaphores?

- SW mechanism for synchronization on higher abstraction than TSL assembly
- Semaphores (S): Integer Variables [System Calls]
  - Two standard *atomic* operations:
    - wait()
    - signal()

```
□ wait(S) {  
    while S  $\leq$  0  
        ; // no-op  
    S--;  
}  
□ signal(S) {  
    S++;  
}
```

S can be an integer resource counter;  
If S is binary, it is called a “**mutex**”

wait(1) → progress & decrement

wait(0) → block

signal(0) → increment and unblock

# ME using Mutexs

```
do {  
    wait (mutex);          mutex initialized to 1  
        // critical section  
    signal (mutex);  
        // non critical section  
} while (TRUE);
```

# Binary Semaphores with TSL: Mutexs

mutex\_lock:

TSL REGISTER,MUTEX

| copy mutex to register and set mutex to 1

CMP REGISTER,#0

| was mutex zero?

JZE ok

| if it was zero, mutex was unlocked, so return

CALL thread\_yield

| mutex is busy; schedule another thread

JMP mutex\_lock

| try again later

ok: RET | return to caller; critical region entered

mutex\_unlock:

MOVE MUTEX,#0

| store a 0 in mutex

RET | return to caller

# Ordering (Sync.) with Semaphores

- Consider 2 concurrent processes P1/P2 with statements S1/S2
- Would like S2 to be executed only after S1 completed
- Let P1 and P2 share a common semaphore “**sync**” set to 0

– [in P1] **S1;**                      *Statement S1 executes in Process P1*  
   **signal(sync);**

– [in P2] **wait(sync);**  
   **S2;**                      *Statement S2 executes in Process P2*

[sync = 0 ➔ P2 executes S2 only after P1 has invoked signal(sync);  
which is only after S1 has been executed]

# Semaphores

mutual exclusion →  
synchronization →

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

```
void producer(void)
{
    int item;
```

```
    while (TRUE) {
        item = produce_item();
        down(&empty);
        CS → down(&mutex);
            insert_item(item);
            up(&mutex);
            up(&full);
        }
    }
```

```
void consumer(void)
{
    int item;
```

```
    while (TRUE) {
        down(&full);
        CS → down(&mutex);
            item = remove_item();
            up(&mutex);
            up(&empty);
            consume_item(item);
        }
    }
```

```
/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */
```

```
/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */
```

```
/* infinite loop */
/* decrement full count */
/* enter critical region */
/* take item from buffer */
/* leave critical region */
/* increment count of empty slots */
/* do something with the item */
```

## The producer-consumer problem using semaphores

# Semaphore Constructs in Java

- Implemented by **java.util.concurrent.Semaphore** class
  - The class uses the wait() and signal() system calls
  - `public Semaphore available = new Semaphore(100);`
  - `available.acquire();`               //decreases semaphore value,  
  uses wait() syscall;
  - `available.release();`            //increases semaphore value,  
  uses signal() syscall;
- ...and other available methods, as `acquire(int n);` `release (int n);` `acquireUninterruptibly()` etc.

# Semaphore Problems?

- Semaphore → effective SW level synchronization
- System calls (simple!)
- But, synchronization errors are still possible through misuse of wait/signal ☹
  - programmer interchanges order of wait() and signal() ops on the semaphore
    - wait(mutex) CS signal(mutex) → signal(mutex) ... CS ... wait(mutex)
    - several processes may end up executing their CS concurrently
  - Suppose a user replaces signal(mutex) with wait(mutex)
    - wait(mutex) ... CS ... wait(mutex)
    - deadlock!!!
  - Suppose the process omits wait(mutex) or signal(mutex) or both
    - ME violated or deadlock



# Monitors: Language Constructs not System Calls

```
monitor example  
    integer i;           // shared variable declarations  
    condition c;  
  
    procedure producer( );  
        .  
        .  
        .  
    end;  
  
    procedure consumer( );  
        .  
        .  
        .  
    end;  
  
end monitor;
```

# Monitors

```
monitor ProducerConsumer
  condition full, empty;
  integer count;
  procedure insert(item: integer);
  begin
    if count = N then wait(full);
    insert_item(item);
    count := count + 1;
    if count = 1 then signal(empty)
  end;
  function remove: integer;
  begin
    if count = 0 then wait(empty);
    remove = remove_item;
    count := count - 1;
    if count = N - 1 then signal(full)
  end;
  count := 0;
end monitor;
```

```
procedure producer;
begin
  while true do
    begin
      item = produce_item;
      ProducerConsumer.insert(item)
    end
  end;
procedure consumer;
begin
  while true do
    begin
      item = ProducerConsumer.remove;
      consume_item(item)
    end
  end;
end;
```

## Outline of producer-consumer problem with monitors

- only one monitor procedure active at one time
- buffer has  $N$  slots

# Monitors in Java

```
static class our_monitor {           // this is a monitor
    private int buffer[ ] = new int[N];
    private int count = 0, lo = 0, hi = 0; // counters and indices
    public synchronized void insert(int val) {
        if (count == N) go_to_sleep(); // if the buffer is full, go to sleep
        buffer [hi] = val;             // insert an item into the buffer
        hi = (hi + 1) % N;             // slot to place next item in
        count = count + 1;             // one more item in the buffer now
        if (count == 1) notify();      // if consumer was sleeping, wake it up
    }
    public synchronized int remove( ) {
        int val;
        if (count == 0) go_to_sleep(); // if the buffer is empty, go to sleep
        val = buffer [lo];             // fetch an item from the buffer
        lo = (lo + 1) % N;             // slot to fetch next item from
        count = count - 1;             // one few items in the buffer
        if (count == N - 1) notify();  // if producer was sleeping, wake it up
        return val;
    }
    private void go_to_sleep( ) { try{wait( );} catch(InterruptedException exc) {};}
}
```

Solution to producer-consumer problem in Java

# Monitors in Java

```
public class ProducerConsumer {
    static final int N = 100;           // constant giving the buffer size
    static producer p = new producer(); // instantiate a new producer thread
    static consumer c = new consumer(); // instantiate a new consumer thread
    static our_monitor mon = new our_monitor(); // instantiate a new monitor
    public static void main(String args[]) {
        p.start();                      // start the producer thread
        c.start();                      // start the consumer thread
    }
    static class producer extends Thread {
        public void run() {              // run method contains the thread code
            int item;
            while (true) {               // producer loop
                item = produce_item();
                mon.insert(item);
            }
        }
        private int produce_item() { ... } // actually produce
    }
    static class consumer extends Thread {
        public void run() {              // run method contains the thread code
            int item;
            while (true) {              // consumer loop
                item = mon.remove();
                consume_item (item);
            }
        }
        private void consume_item(int item) { ... } // actually consume
    }
}
```

Solution to producer-consumer problem in Java

# Problems?

- **TSL**: lowest level (HW),
- **Semaphores**: low-level (kernel), depend too much on programmer's skills
- **Monitors**: need language support (C/Pascal?)
- **All**: **memory sharing** solutions, work only on the same machine but not if the processes sit in different machines (LAN etc.)
- Let's look at **message passing** solutions (send/receive)

# Producer-Consumer with Message Passing

```
#define N 100                                /* number of slots in the buffer */

void producer(void)
{
    int item;
    message m;                               /* message buffer */


    while (TRUE) {
        item = produce_item();               /* generate something to put in buffer */
        receive(consumer, &m);               /* wait for an empty to arrive */
        build_message(&m, item);             /* construct a message to send */
        send(consumer, &m);                   /* send item to consumer */
    }
}

void consumer(void)
{
    int item, i;
    message m;

    for (i = 0; i < N; i++) send(producer, &m); /* send N empties */
    while (TRUE) {
        receive(producer, &m);               /* get message containing item */
        item = extract_item(&m);             /* extract item from message */
        send(producer, &m);                  /* send back empty reply */
        consume_item(item);                  /* do something with the item */
    }
}
```

Ques: what happens if the producer (or the consumer) is much faster at processing messages than the consumer (or producer)?

# Barriers (primitives) for Synchronization



## Use of a barrier (~ AND operation)

- a) processes approaching a barrier
- b) all processes blocked at barrier, waiting for C
- c) last process (C) arrives, all are let through

# Synchronization Implementations

- Solaris
  - adaptive mutex, semaphores, RW locks, threads blocked by locks etc
- Windows XP
  - interrupt masking, busy-waiting spin locks (for short code segments), mutex, semaphores, monitors, msg. passing
- Linux
  - pre v2.6 (non-preemptible); post v2.6 preemptible: interrupts
  - semaphores, spin-locks (for short CS's in kernel only)




# Classical ME/Concurrency Problems

- Synchron. problems (Dining Philosophers Problem)
- Access problems (Readers/Writers Problem)
- Scheduling (Sleeping Barber Problem)

# Dining Philosophers

- Philosophers eat/think
- Eating needs 2 chopsticks (forks)
- Pick one instrument at a time
- Solutions (with semaphores?)...



# Dining Philosophers – Obvious Solution

```
#define N 5                                /* number of philosophers */

void philosopher(int i)                    /* i: philosopher number, from 0 to 4 */
{
    while (TRUE) {
        think( );                          /* philosopher is thinking */
        wait  take_fork(i);                 /* take left fork */
        wait  take_fork((i+1) % N);         /* take right fork; % is modulo operator */
        eat( );                             /* yum-yum, spaghetti */
        signal put_fork(i);                 /* put left fork back on the table */
        signal put_fork((i+1) % N);         /* put right fork back on the table */
    }
}
```

Deadlocks? – all pick the left fork at the same time  
→ add a check if the fork is available → **Livelock!**

# Dining-Philosophers Problem

- Philosopher  $i$ :

```
while (true) {  
  
    ...think...  
  
    wait ( chopstick[i] );  
    wait ( chopstick[ (i + 1) % N] )  
    signal (...) ?  
    ...eat...  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % N] );  
}
```

Needs:

- No deadlock
- No starvation for anyone
- Maximum parallelism

# Dining Philosophers – No deadlocks - Max Parallelism Solution

```
#define N          5          /* number of philosophers */
#define LEFT      (i+N-1)%N   /* number of i's left neighbor */
#define RIGHT     (i+1)%N     /* number of i's right neighbor */
#define THINKING  0          /* philosopher is thinking */
#define HUNGRY    1          /* philosopher is trying to get forks */
#define EATING    2          /* philosopher is eating */
typedef int semaphore;       /* semaphores are a special kind of int */
int state[N];               /* array to keep track of everyone's state */
semaphore mutex = 1;        /* mutual exclusion for critical regions */
semaphore s[N];             /* one semaphore per philosopher */

void philosopher(int i)      /* i: philosopher number, from 0 to N-1 */
{
    while (TRUE) {           /* repeat forever */
        think( );           /* philosopher is thinking */
        take_forks(i);       /* acquire two forks or block */
        eat( );              /* yum-yum, spaghetti */
        put_forks(i);        /* put both forks back on table */
    }
}
```

# Continuation...

```
void take_forks(int i)                /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);                      /* enter critical region */
    state[i] = HUNGRY;                 /* record fact that philosopher i is hungry */
    test(i);                          /* try to acquire 2 forks */
    up(&mutex);                        /* exit critical region */
    down(&s[i]);                       /* block if forks were not acquired */
}

void put_forks(i)                    /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);                      /* enter critical region */
    state[i] = THINKING;              /* philosopher has finished eating */
    test(LEFT);                      /* see if left neighbor can now eat */
    test(RIGHT);                    /* see if right neighbor can now eat */
    up(&mutex);                       /* exit critical region */
}

void test(i)                        /* i: philosopher number, from 0 to N-1 */
{
    if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
        state[i] = EATING;
        up(&s[i]);    // acquire forks
    }
}
```

Deadlock free + max. parallelism

# Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - Readers – only read the database; they do **not** perform any updates
  - Writers – can both read and write.
- Problem – allow multiple readers to queue to read at the same time. Only one single writer can access the shared data at a time.
- Shared Data
  - Database
  - Integer **readcount** initialized to 0 (# of processes currently reading object)
  - Semaphore **mutex** initialized to 1; controls the access to **readcount**
  - Semaphore **db** initialized to 1; controls access to database;

# The Readers and Writers Problem

```
typedef int semaphore;
semaphore mutex = 1;
semaphore db = 1;
int rc = 0;

/* use your imagination */
/* controls access to 'rc' */
/* controls access to the database */
/* # of processes reading or wanting to */

void reader(void)
{
    while (TRUE) {
        down(&mutex);
        rc = rc + 1;
        if (rc == 1) down(&db);
        up(&mutex);
        read_data_base();
        down(&mutex);
        rc = rc - 1;
        if (rc == 0) up(&db);
        up(&mutex);
        use_data_read();
    }
}

/* repeat forever */
/* get exclusive access to 'rc' */
/* one reader more now */
/* if this is the first reader ... */
/* release exclusive access to 'rc' */
/* access the data */
/* get exclusive access to 'rc' */
/* one reader fewer now */
/* if this is the last reader ... */
/* release exclusive access to 'rc' */
/* noncritical region */

void writer(void)
{
    while (TRUE) {
        think_up_data();
        down(&db);
        write_data_base();
        up(&db);
    }
}
```



# The Sleeping Barber Problem

- 1 Barber
- 1 Barber Chair
- $N$  Customer Chairs



# The Sleeping Barber Problem

```
#define CHAIRS 5                /* # chairs for waiting customers */

typedef int semaphore;          /* use your imagination */

semaphore customers = 0;        /* # of customers waiting for service */
semaphore barbers = 0;         /* # of barbers waiting for customers */
semaphore mutex = 1;           /* for mutual exclusion */
int waiting = 0;               /* customers are waiting (not being cut) */

void barber(void)
{
    while (TRUE) {
        down(&customers);      /* go to sleep if # of customers is 0 */
        down(&mutex);           /* acquire access to 'waiting' */
        waiting = waiting - 1;  /* decrement count of waiting customers */
        up(&barbers);           /* one barber is now ready to cut hair */
        up(&mutex);             /* release 'waiting' */
        cut_hair();             /* cut hair (outside critical region) */
    }
}

void customer(void)
{
    down(&mutex);               /* enter critical region */
    if (waiting < CHAIRS) {     /* if there are no free chairs, leave */
        waiting = waiting + 1; /* increment count of waiting customers */
        up(&customers);         /* wake up barber if necessary */
        up(&mutex);             /* release access to 'waiting' */
        down(&barbers);         /* go to sleep if # of free barbers is 0 */
        get_haircut();          /* be seated and be serviced */
    } else {
        up(&mutex);             /* shop is full; do not wait */
    }
}
```