

Embedded Operating Systems

- Requirements
- Design
- Implementation



Embedded Systems

❑ Computer system embedded within larger system

- Often not perceived as computer system
- Handles particular task
- Interacts with the physical world
- Large variety



❑ Application:

- Automotive
- Avionic
- Consumer electronics
- Mobile devices
- ... and many more



Requirements of Embedded Systems

Fundamental differences to desktop (operating) systems!

❑ Costs

- Hardware (CPU, RAM, Flash, etc)
- Licenses
- Software ?

❑ Real-time constraints

- Correctness of system also depends on **timely delivery**
- Soft real-time constraints
 - Slight violations can be tolerated (e.g. multimedia)
- Hard real-time constraints
 - Violations must be avoided (life can be at risk)
 - Potentially excessive costs (not just monetary)

Requirements of Embedded Systems

□ Safety

- "Absence of catastrophic consequences on the user(s) and the environment" (Avizienis et al., 2004)
- Safety standards and specifications, e.g. IEC 61508 or ISO 26262

□ Robustness of hardware and software

- Environmental influences
- Software failures

□ Security

- Protection from unauthorized access, use, modification
- Internal / external attackers
- Attack surface:
 - Direct connection (e.g., debug ports)
 - Short- & Long-range wireless communication
- System lifetime

How to Fullfill These Requirements?

□ Costs

- Minimize per-piece costs
 - Low cost components
 - Microcontrollers (8, 16, 32bit)
 - Low amount of memory:
1MB Flash / 128KB RAM are "high-end"
- Development costs are a minor issue

□ Real-time constraints

- Provide guarantees on worst-case execution time (WCET)
 - Static analysis, Simulation, Dynamic analysis
- Eliminate non-determinism
 - Static scheduling
 - Static memory layout and allocation

How to Fullfill These Requirements?

❑ Robustness

- “Ability of a system to handle errors during execution”
- Redundancy
- Fault tolerance

❑ Safety & Security

- Rigorous software design processes
- Following established standards / best practices
 - Static analysis
 - Dynamic testing
 - Fault injection
 - Run-time monitoring
 - Code reviews
 - Penetration testing
 - ... and more

Implications on Embedded OS


❑ OS still is **resource allocator** and **control program**

❑ Key differences

- Resources are (highly) constrained
- No multi-user
- Kernel-/user-mode less common

❑ What about


- Process management
- Memory management
- I/O
- Scheduling
- Overall architecture



Example: AUTOSAR OS

Modern automobile

- Up to 80 electronic control units
 - Engine
 - Transmission
 - Brakes
 - Airbag
 - Infotainment
- Distributed system
 - CAN, FlexRay, LIN, MOST
- Safety-critical
- Increasingly complex



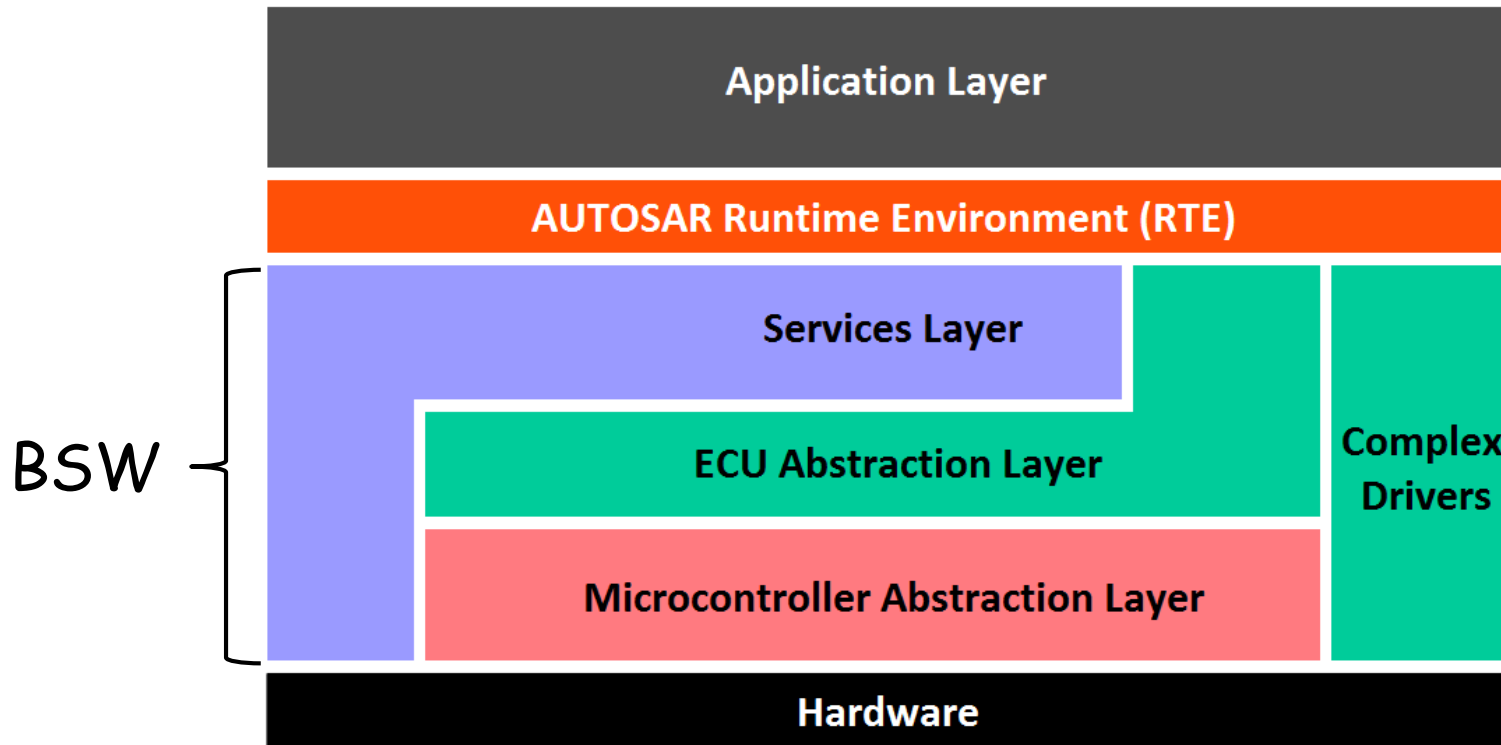
AUTOSAR standard

- Open automotive system architecture
- Partners: BMW, Daimler, Ford, GM, VW and many more
- "Cooperate on Standards, Compete on Implementations"

AUTOSAR Architecture

Layered architecture

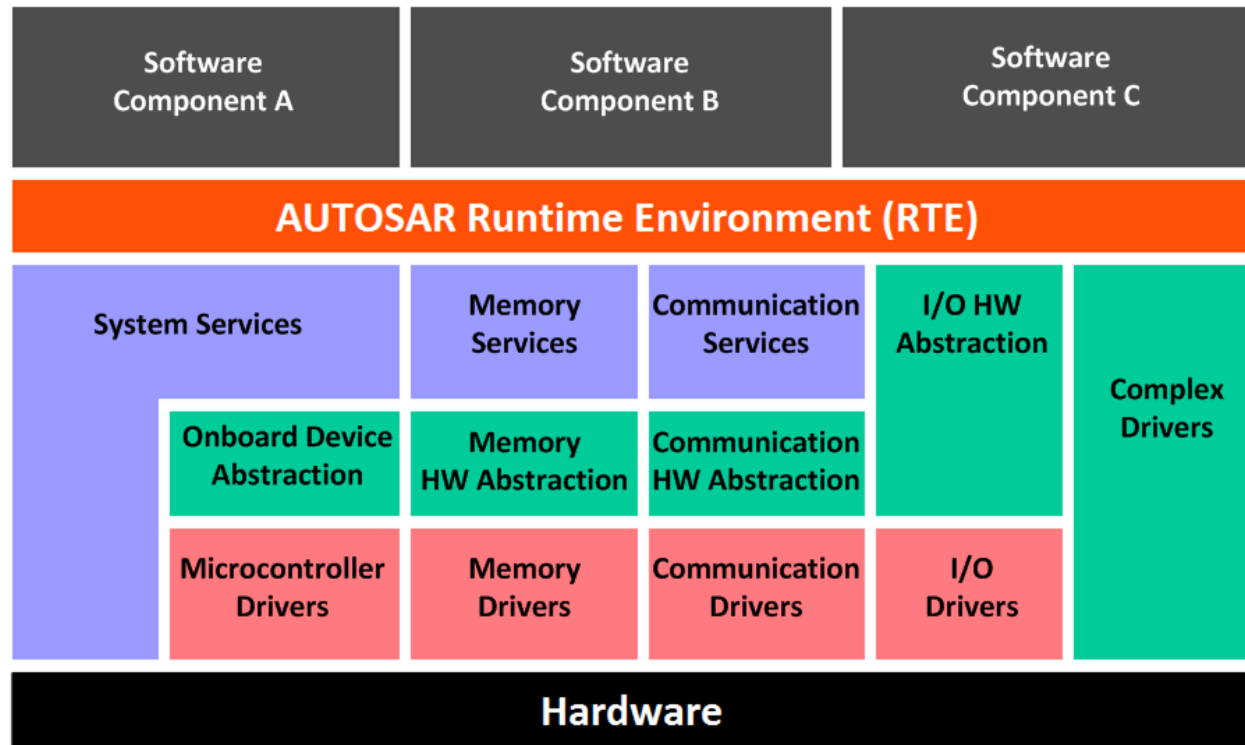
- Application: Software components implementing runnables
- RTE: Communication and comm. abstraction services
- BSW: Operating system, services and hardware abstraction



AUTOSAR Architecture

Component-based architecture

- Application Software Components (SW-Cs)
- SW-Cs contain **runnables** as functional blocks
- Runnables are grouped to **tasks**
- Tasks are scheduled by the OS



AUTOSAR OS

- ❑ Based on the widely used OSEK/VDX
 - ❑ Event triggered OS
 - Timers
 - Signals (e.g., message arrival)
→ Both events drive task scheduling
 - ❑ Basic features
 - Configured and scaled statically
 - Fixed, priority-based scheduling
 - Hostable on low end controllers
 - No external resources
- Goal: Achieve **determinism** & **predictability** !

❑ Process Management

- No programs or processes → instead: tasks
- Tasks contain runnables (functional entities)
- Task types: non-preemptive or preemptive
- Static and “monolithic” system (but modular)

❑ Memory Management

- Static memory layout
- Eliminates non-determinism
- Functions such as `malloc()` not implemented


❑ Scheduling

- Statically configured via timers / alarms
- Worst-case execution time analysis / estimation
- Different scheduling planning algorithms

Problem: OS does not provide schedulability analysis!

Timing Violation Example

Task ID	Priority	Worst-case Execution Time	Deadline / Period
A	High	1	5
B	Medium	3	10
C	Low	5	15




- Ready time of A, B, C is 0
- All tasks keep their execution time limits

→ Everything fine here

Timing Violation Example

Task ID	Priority	Worst-case Execution Time	Deadline / Period
A	High	1	5
B	Medium	3	10
C	Low	5	15



- B and A both violate their execution time budget
- C behaves correctly, still misses deadline

→ Error propagation

Real Time and Scheduling Analysis

Goal: Ensure that all tasks meet their deadlines

☐ Testing is inappropriate

- Only verifies if deadlines are usually met, not always
- Constructing the worst case is often difficult

→ Formulate assumptions that **constrain** & **ease** the problem

☐ Ensure testing for the worst case gets easier

- Static scheduling

☐ Ensure worst case is mathematically understood


- Rate monotonic scheduling

Real Time and Scheduling Analysis: Assumptions

- ❑ All tasks $\{T_0, T_1, \dots, T_{N-1}\}$ are perfectly periodic
- ❑ All tasks are completely independent
 - In practice difficult to realize
 - Shared resources, variables, multiplexed A/D, etc...
- ❑ Worst case execution time (WCET) is known
 - Longest possible path through branches
 - Worst combination of cache misses
 - Computation / analysis tools (e.g. AbsInt)
 - Measurement
 - Add margins
- ❑ Zero time task switching
 - Add to task execution time
 - Leave some % of CPU capacity as margin

Pure Static Scheduling (non-preemptive)

- ❑ All tasks run in a big loop
- ❑ No interrupts, no preemption
- ❑ Pros:
 - Simplicity (no scheduling SW, no OS interaction needed)
 - Easy to test / validate
 - Deterministic
- ❑ Cons:
 - Unflexible
 - No prioritization
 - I/O polling



© P. Koopman – Better Embedded Software, 2010

Pure Static Scheduling (non-preemptive)

Scheduling analysis terminology

N Number of tasks

C_i Computation time of task T_i (worst case!)

D_i Deadline of task T_i


P_{main} Period of main loop

$$\sum_{i=0}^{N-1} C_i \leq \min_{i=0}^{N-1} (D_i)$$

$$\sum_{i=0}^{N-1} C_i \leq P_{main} \leq \min_{i=0}^{N-1} (D_i)$$

□ Improvements / Variations

- Multiple runs of a task (ensure spacing!)
- Run one background interrupt



© P. Koopman – Better Embedded Software, 2010

Static Scheduling with Small Helper Interrupts

❑ Helper interrupts

- Strict limitation of functionality
- I/O servicing (read/write device)
- Service counter/timers
- Small task with very fast deadline


→ Ensure ISRs are fast w/o data sharing problems

❑ Pros:

- Interrupt usage
- No polled I/O
- Tight deadlines handled in ISR

❑ Cons:

- ISR functionality limited to few dozen instructions



© P. Koopman – Better Embedded Software, 2010

Static Scheduling with Small Helper Interrupts

Scheduling analysis - Easy case

Assumptions:

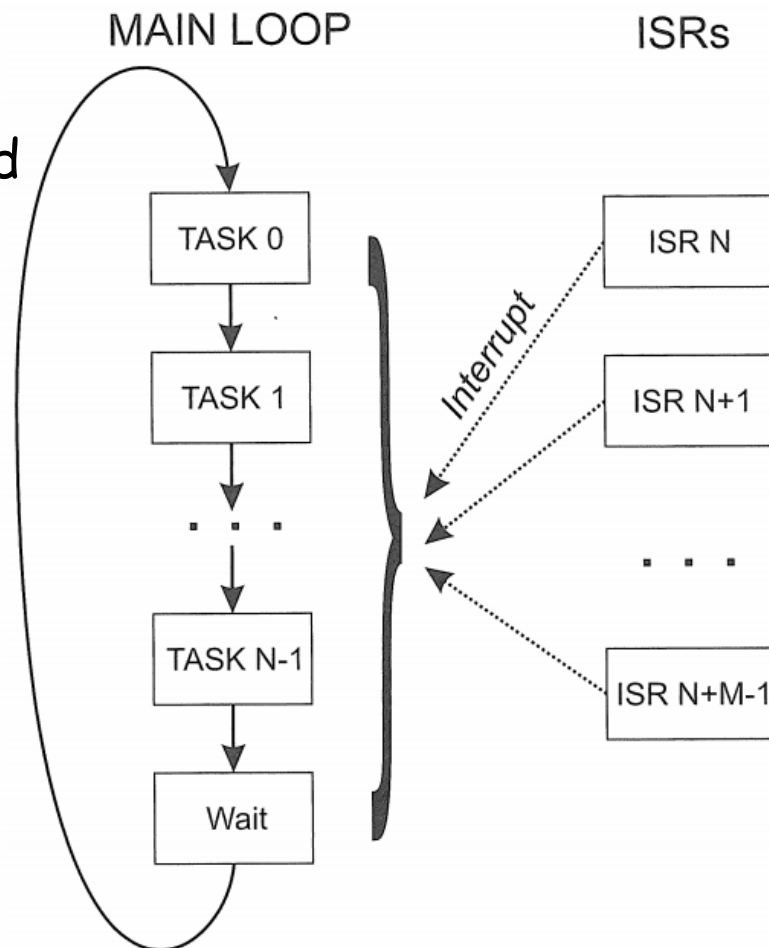
- Every ISR activates once per period
- ISR deadline not shorter than mainloop

→ Limited practical use - helps understanding difficult case!

M Number of ISRs

$$\sum_{i=0}^{N+M-1} C_i \leq P_{main} \leq \min_{i=0}^{N+M-1} (D_i)$$

→ ISRs treated as task!



© P. Koopman – Better Embedded Software, 2010

Static Scheduling with Small Helper Interrupts

Scheduling analysis - Difficult case

Assumption:

- ISR may execute more than once per main loop:

$$\rightarrow D_{ISR} \leq P_{main}$$

Problems:

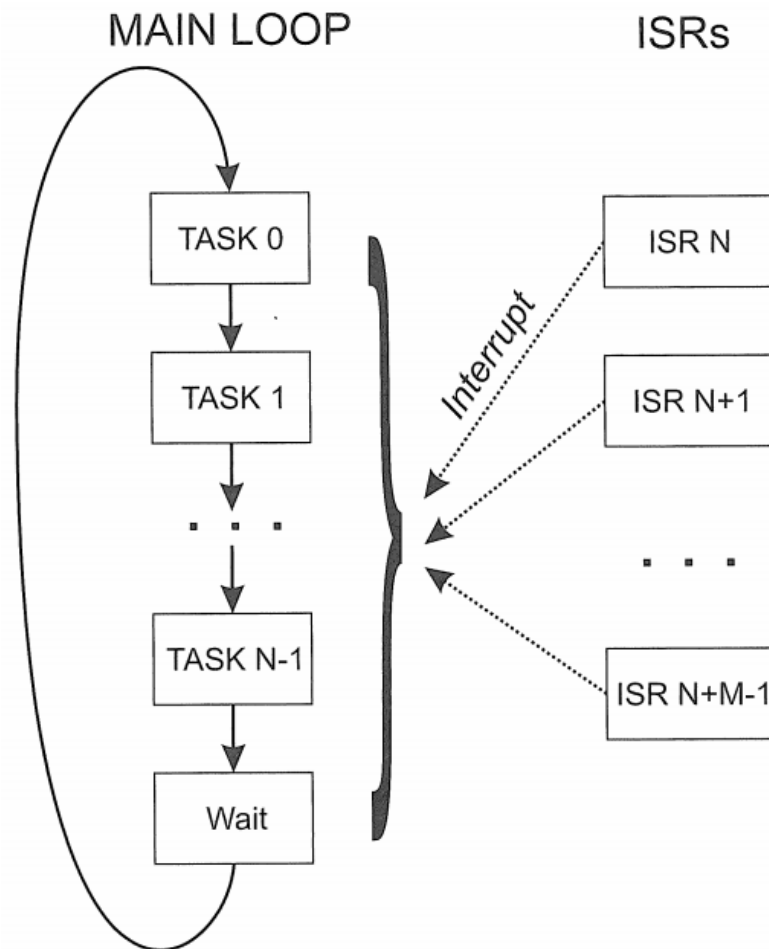
- Starvation
- Priority inversion (ISR)

Solution:

- Each ISR can execute only once within smallest ISR deadline

$$\sum_{i=N}^{N+M-1} C_i \leq \min_{i=N}^{N+M-1} (D_i)$$

Limitation: No long running ISRs!



© P. Koopman – Better Embedded Software, 2010

Static Scheduling with Small Helper Interrupts

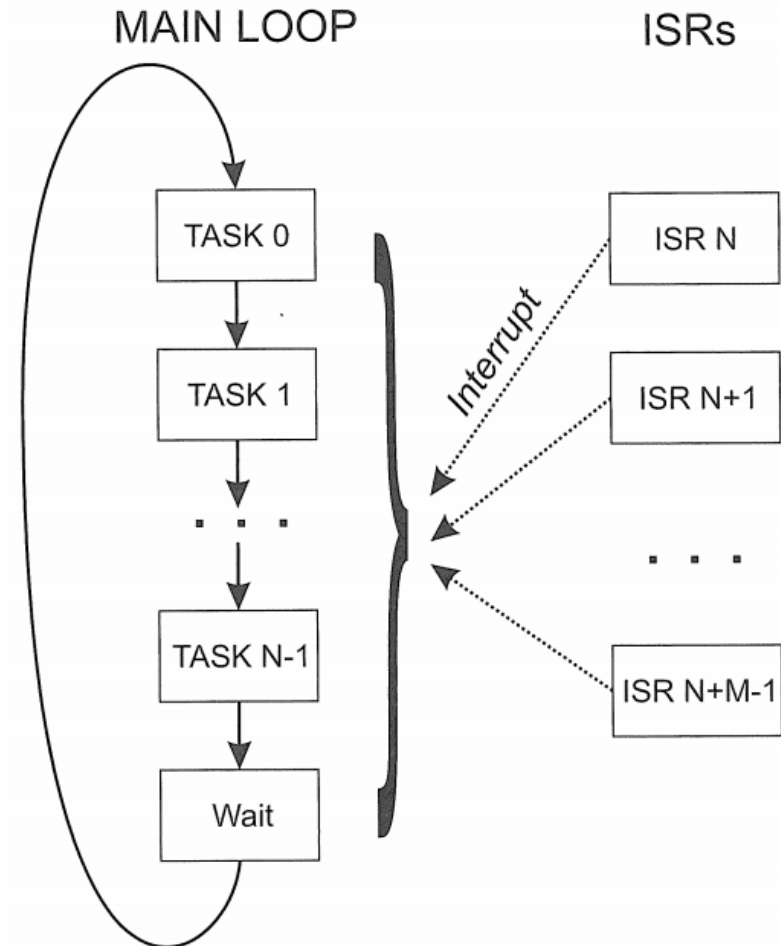
Scheduling analysis - Difficult case

Maximum computation time of all interrupts in P_{main} :

$$C_{ISR_Total} \leq \sum_{i=N}^{i=N+M-1} \left(\left\lceil \frac{P_{main}}{P_i} \right\rceil \cdot C_i \right)$$

Schedulability condition:

$$\sum_{i=0}^{N-1} C_i + C_{ISR_Total} \leq P_{main}$$



© P. Koopman – Better Embedded Software, 2010

Rate Monotonic Scheduling

□ Properties

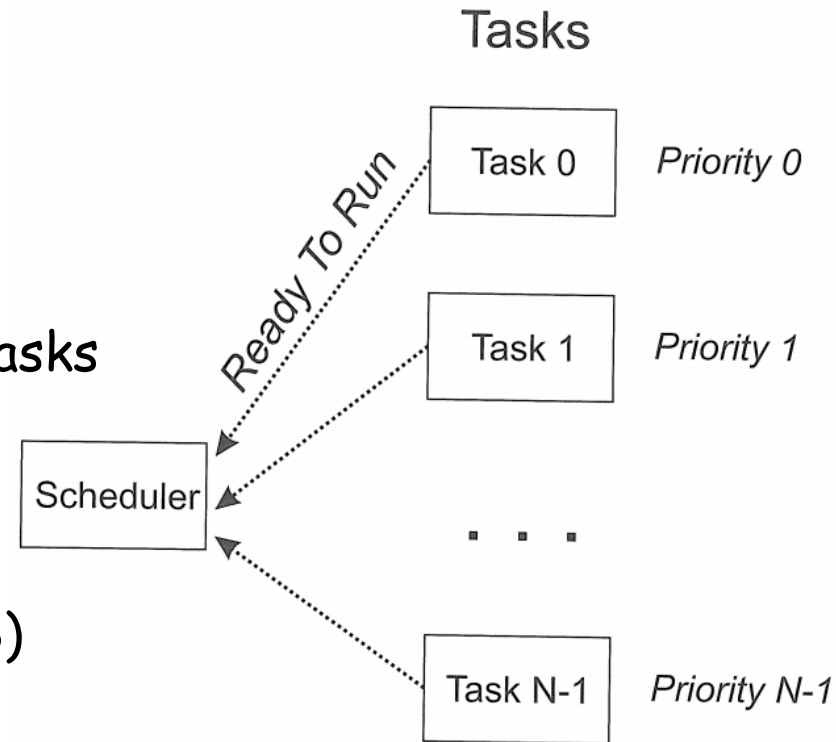
- Most flexible/complex
- Supports task preemption
- Scheduler instead of main loop
- Priorities statically assigned

□ Pros:

- Supports tasks with long periods
- Mix non-realtime with realtime tasks
- Simple scheduling system

□ Cons:

- Requires a scheduler (with RTOS)
- Assumes no interrupts
- Assumes zero latency preemption



© P. Koopman – Better Embedded Software, 2010

Rate Monotonic Scheduling

Scheduling analysis

❑ Complex!

- "Scheduling algorithms for multiprogramming in a hard real-time environment", C. L. Liu and J. Layland, Journal of the ACM 20 (1): 46–61, 1973.

doi:10.1145/321738.321743 (over 9500 citations)

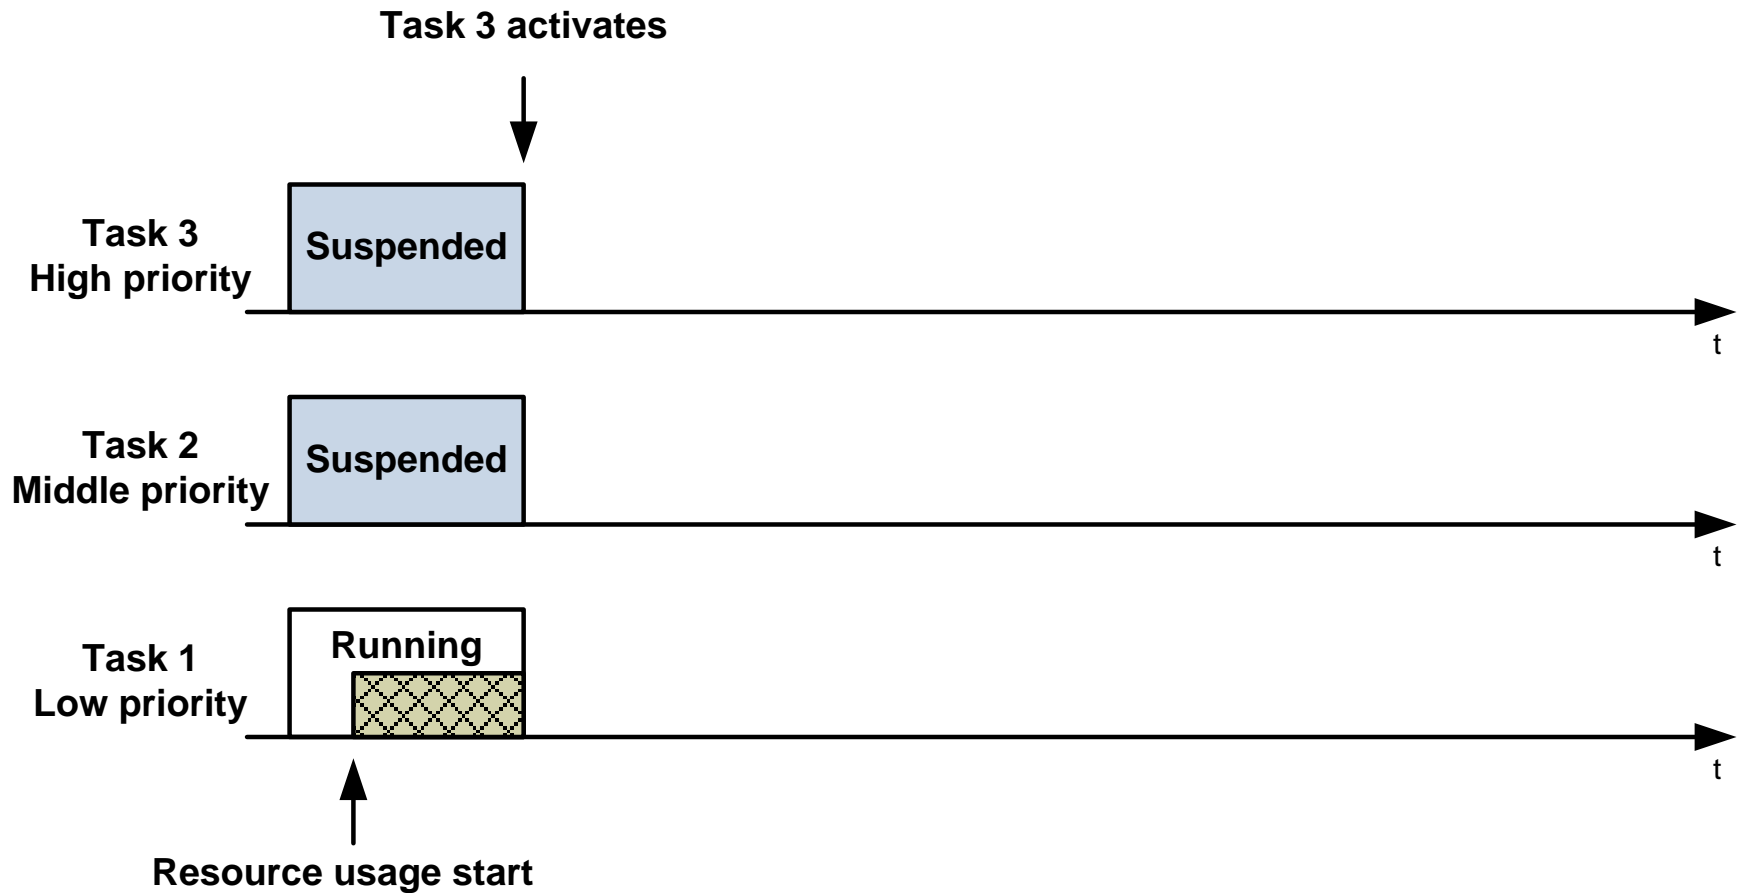
❑ Special case:

- Schedulability for 100% CPU usage only guaranteed for harmonic task periods
- Example: {2ms, 10ms, 20ms} vs. {2ms, 10ms, 25ms}

❑ Pitfalls:

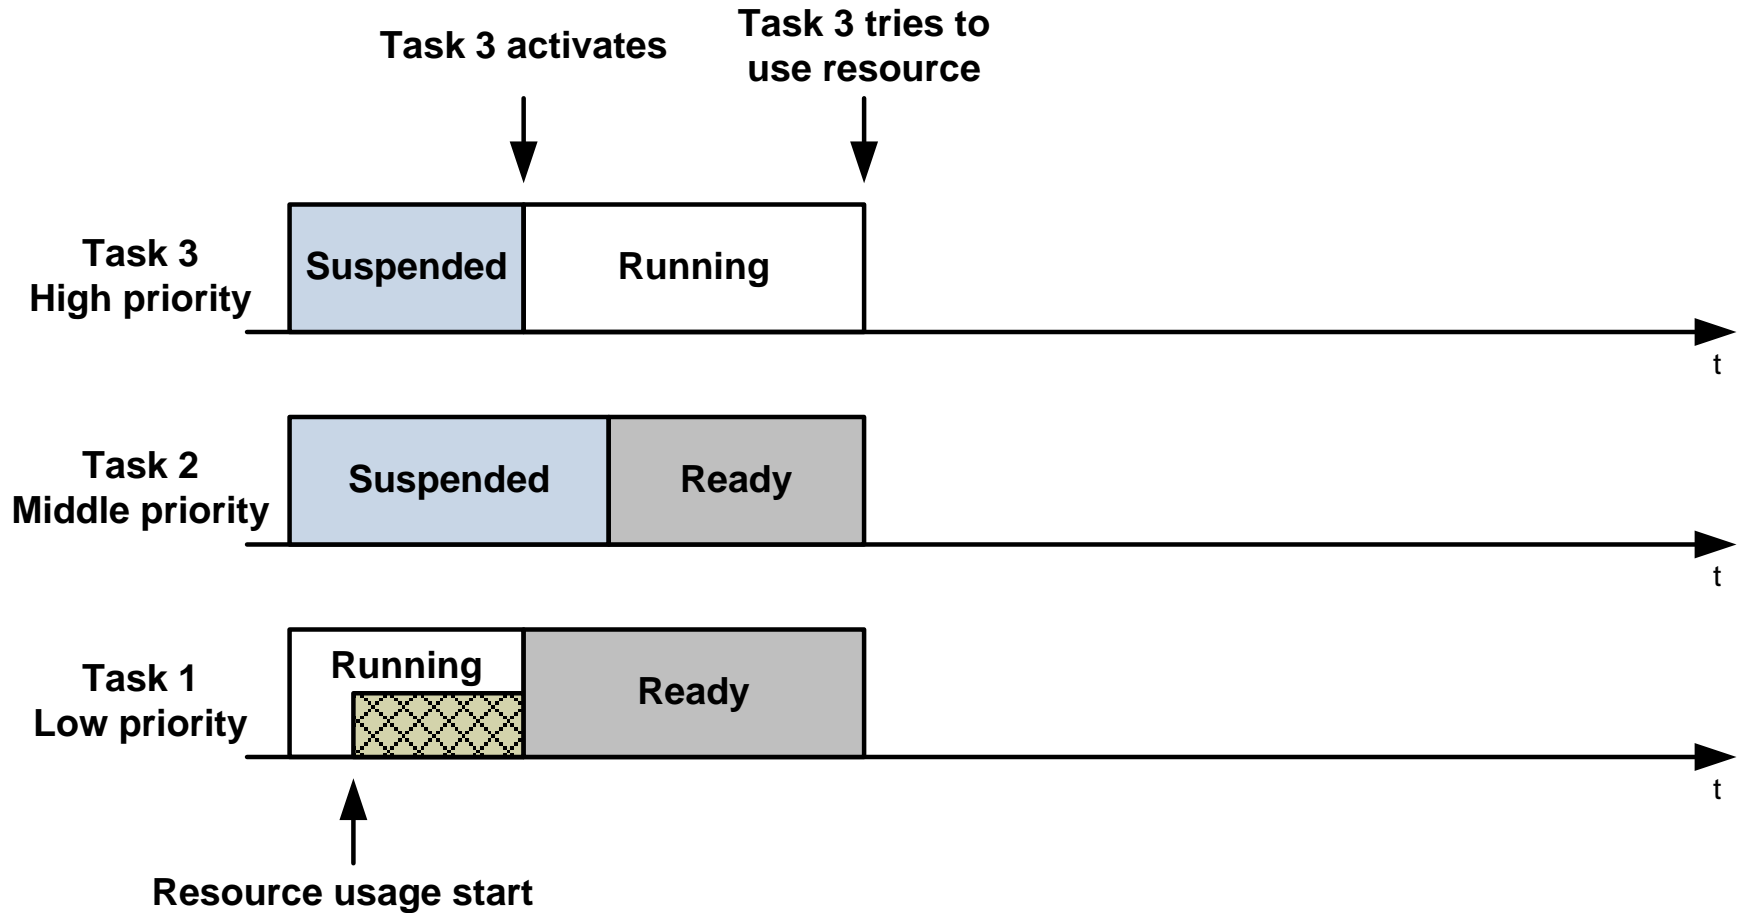
- No interaction between tasks
- Prone to priority inversion!!

Priority Inversion Problem



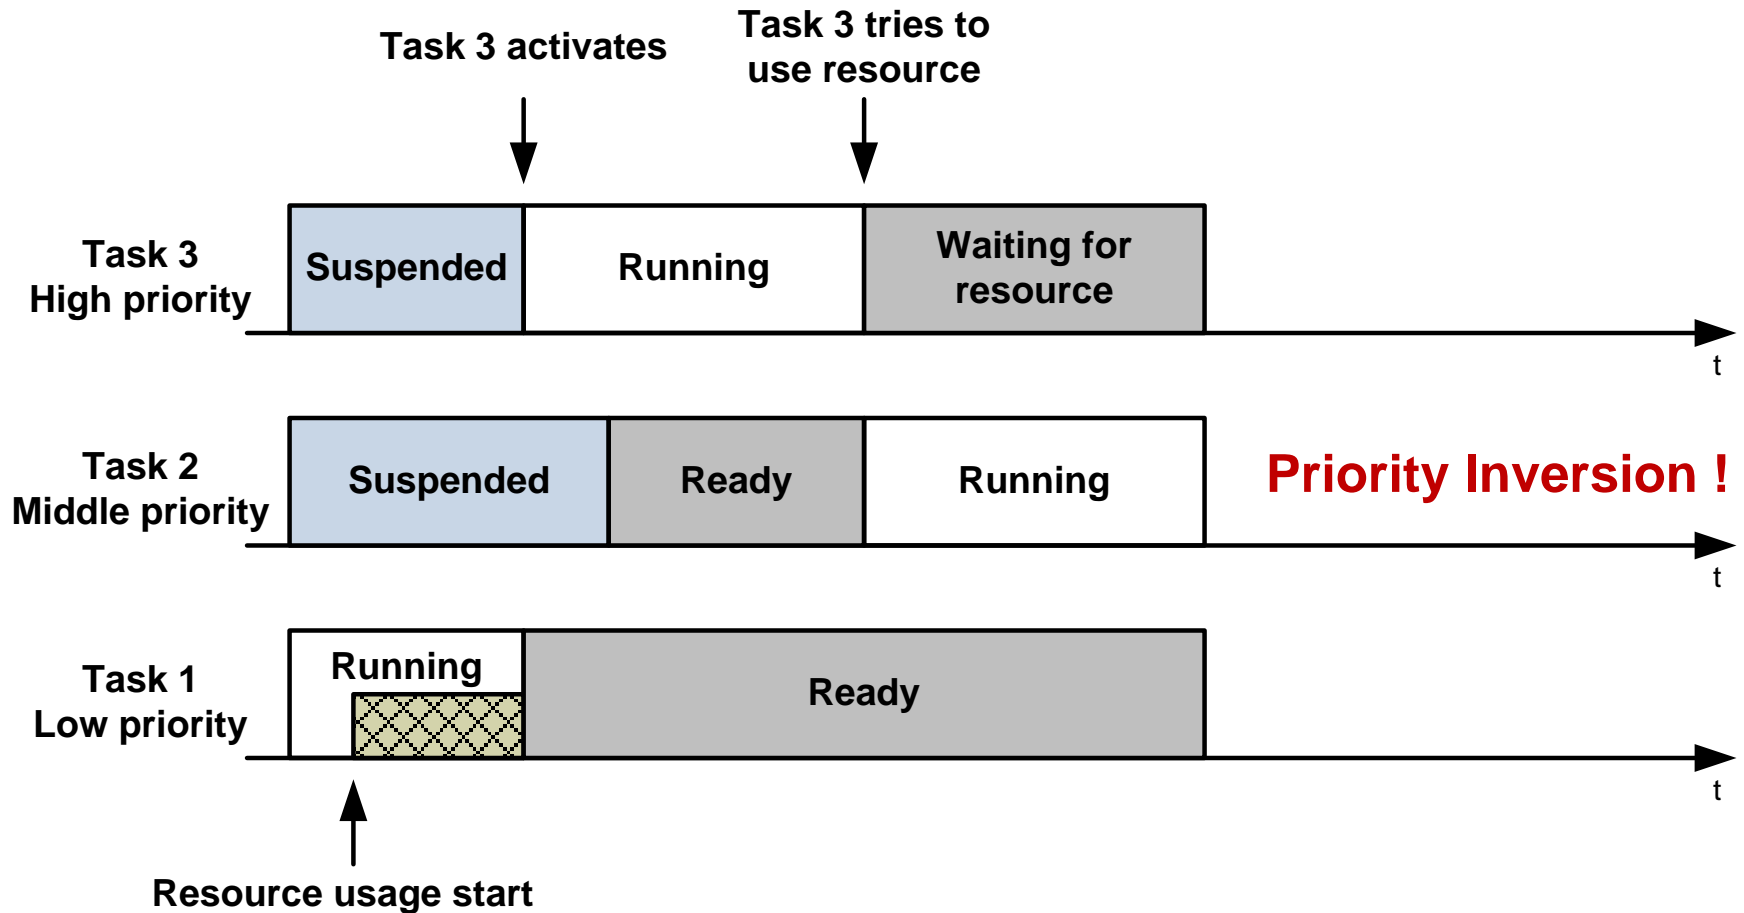
© W. Zimmermann, et al. – Bussysteme in der Fahrzeugtechnik

Priority Inversion Problem



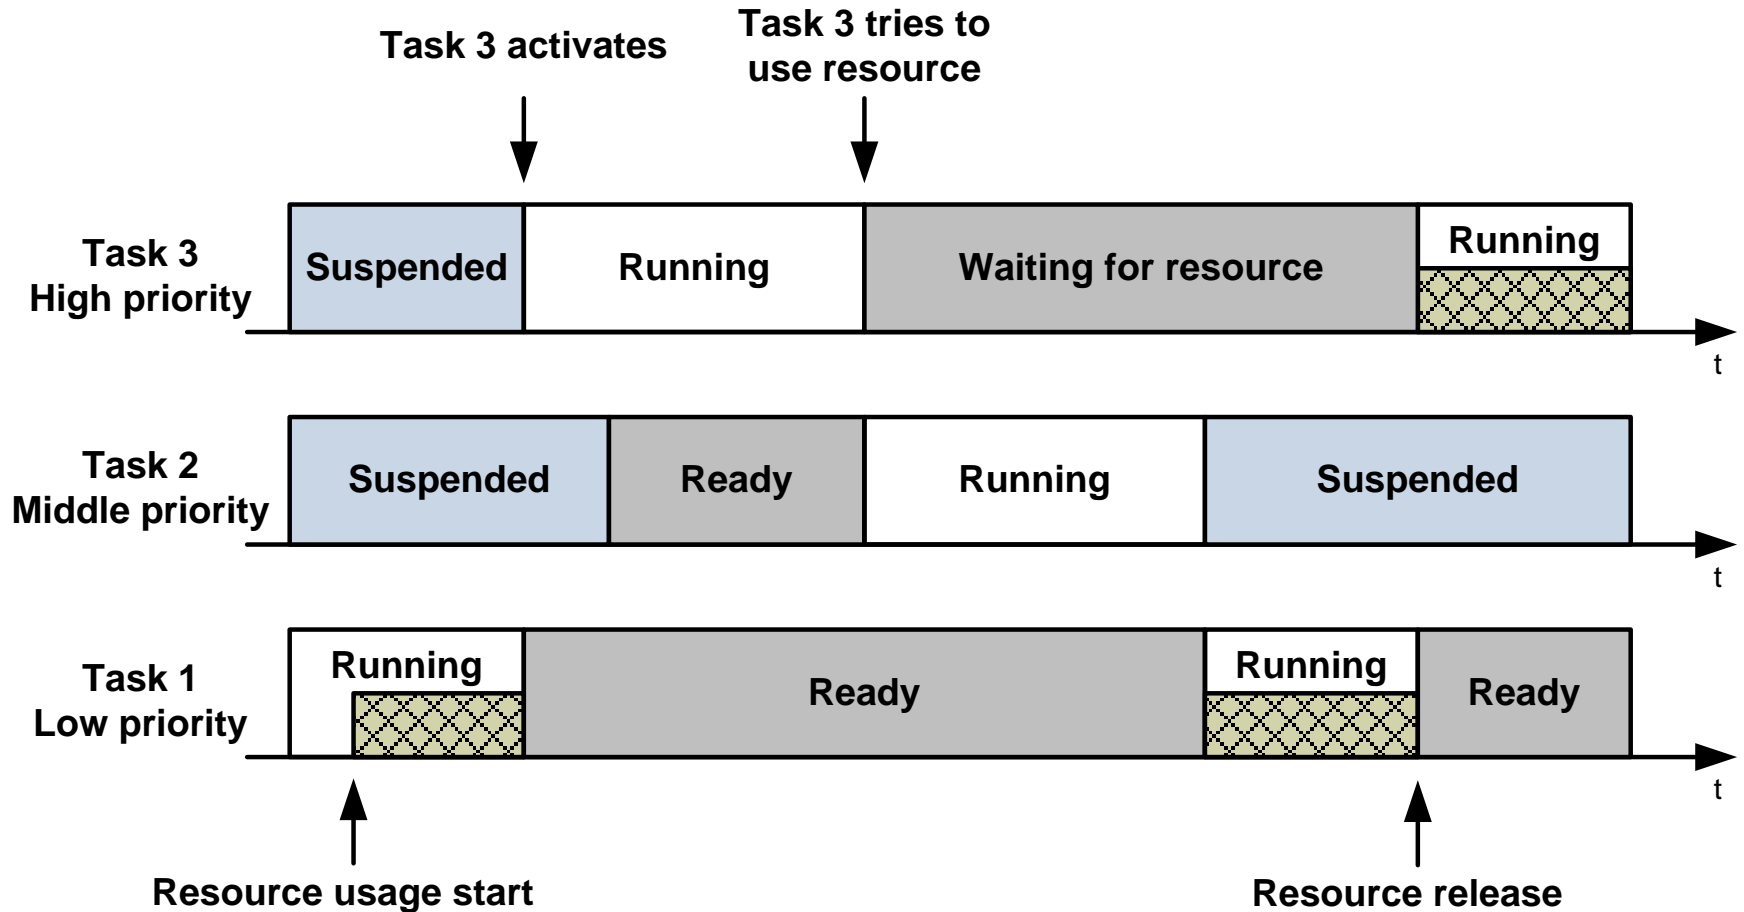
© W. Zimmermann, et al. – Bussysteme in der Fahrzeugtechnik

Priority Inversion Problem



© W. Zimmermann, et al. – Bussysteme in der Fahrzeugtechnik

Priority Inversion Problem



© W. Zimmermann, et al. – Bussysteme in der Fahrzeugtechnik

Famous example: NASA Mars Pathfinder Mission

http://research.microsoft.com/en-us/um/people/mbj/mars_pathfinder/mars_pathfinder.html

Priority Ceiling / Inheritance

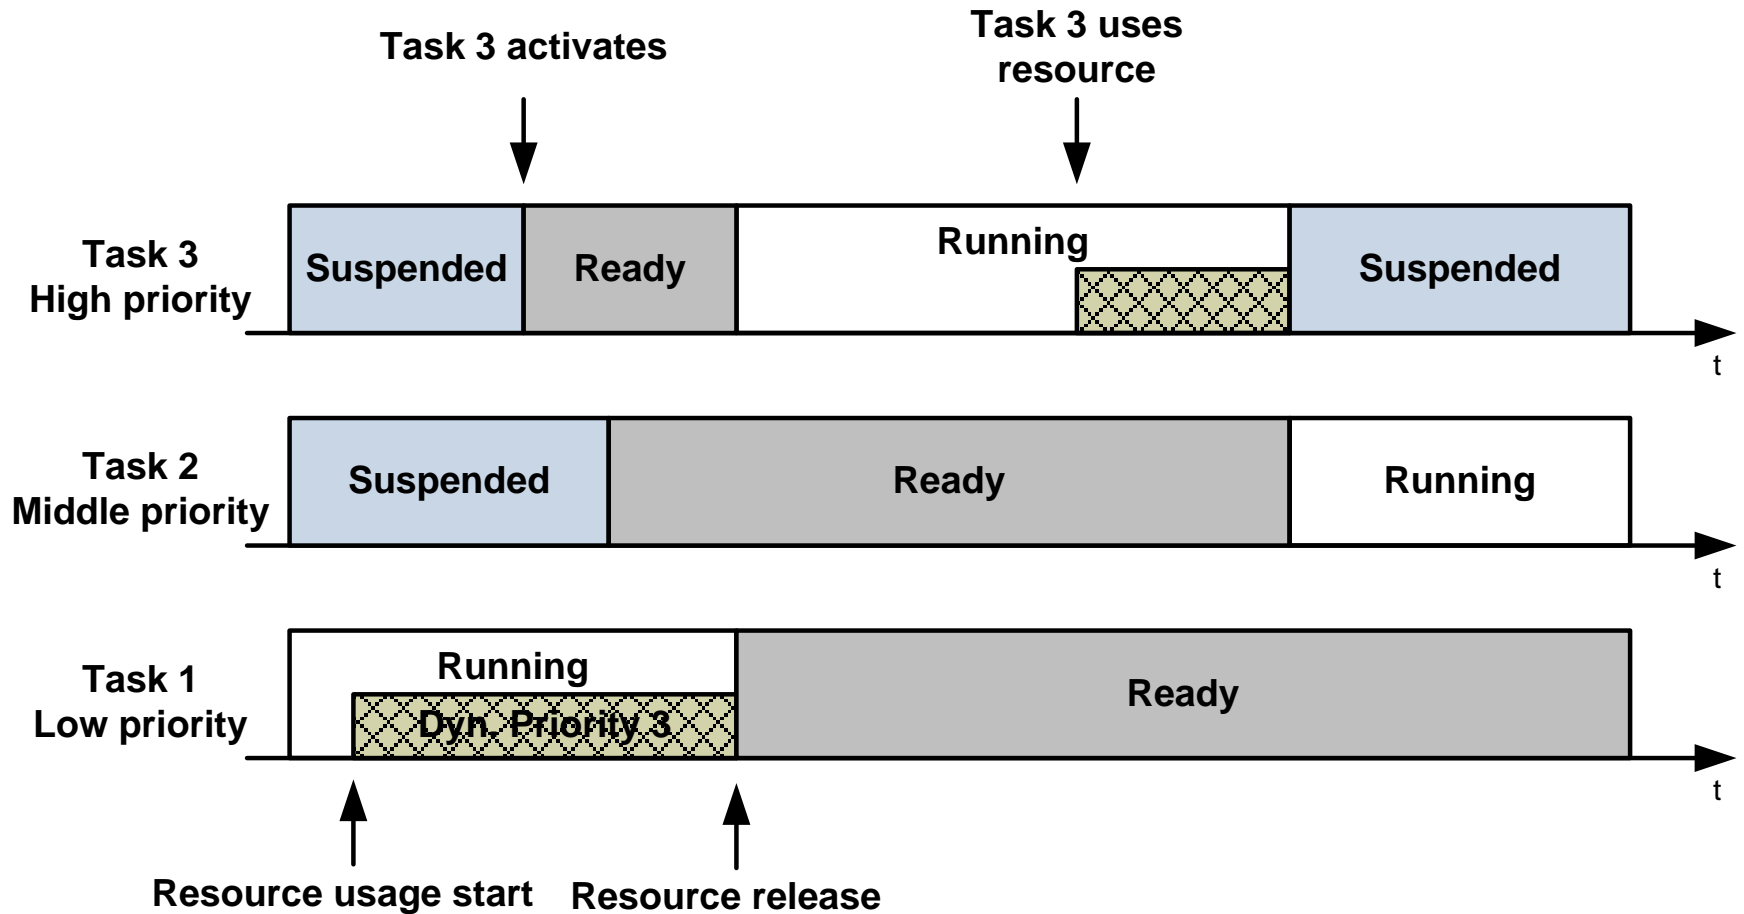
□ Priority ceiling

- Each resource is assigned a **priority ceiling**:
Priority ceiling of resource is the highest priority of any task that may request the resource
- Task's priority is raised to priority ceiling of a resource when they acquire it
- Effective from the moment of resource acquisition

□ Priority inheritance

- Scheduler dynamically changes priority of task, when a higher prioritized task is waiting for the resource
- Effective from the moment of higher prioritized resource acquisition attempt
- Low priority tasks gets the priority of the waiting task

Priority Ceiling



© W. Zimmermann, et al. – Bussysteme in der Fahrzeugtechnik

Further Schedule Planning Algorithms

- ❑ Brute force - planning by searching
 - Simple and naïve approach
 - Exhaustive search, number of combinations for n tasks is $n!$
 - NP-complete problem when disregarding ready times + deadlines

- ❑ Earliest deadline first
 - Schedules tasks by their deadlines
 - Optimal, if tasks have same ready times
 - If not, not feasible in non-preemptive case

- ❑ Other examples:
 - Least laxity first / least slack time
 - For multi-processor systems
 - NP-complete in non-preemptive case

→ Will be covered in exercise

AUTOSAR Watchdog Manager

❑ Watchdog timer

- Hardware device (internal/external)
- Expects a signal within a configurable time
- Triggers reaction on timeout

❑ Alive supervision

- Checkpoints in supervised code sections
- Watchdog checks periodically if checkpoints have been reached
- Not suitable for aperiodical events

❑ Deadline supervision

- Checkpoints in supervised code sections (e.g. beginning/end)
- Watchdog checks **timing** of checkpoint **transitions** (aperiodic/episodical)


See: http://www.autosar.org/download/R4.0/AUTOSAR_SWS_WatchdogManager.pdf

Robustness Mechanisms

AUTOSAR Watchdog Manager

□ Logical supervision

- Checkpoints in supervised code sections
- Program flow monitoring
- Watchdog checks **legality** of transitions
- Example: Transition to bb5 only valid if it originates from bb4 or bb0



See: http://www.autosar.org/download/R4.0/AUTOSAR_SWS_WatchdogManager.pdf

AUTOSAR Operating System

☐ Memory protection

- Hardware support required! (MPU/MMU)
- Protect code / data / stack segments from unauthorized writes and reads

☐ Timing protection

- Handle missed deadlines at runtime (timing fault)
- Timing fault types:
 - Execution time protection
 - Resource locking time protection
 - Inter-arrival time protection

☐ Hardware protection

- Privileged and non-privileged mode (kernel/user modes)

See: http://www.autosar.org/download/R4.0/AUTOSAR_SWS_OS.pdf