



Multiprocessing Support


- Parallel computer architectures
- OS design issues

Parallel computer architectures

- ❑ Parallel computing motivation
- ❑ Parallel computing architectures
- ❑ Consequences for OS implementations


Parallel Processing Motivation

- ❑ Single Processor is a Single Point Of Failure (SPOF)
- ❑ Not tolerable in critical applications
- ❑ Redundancy helps, e.g. Triple Modular Redundancy (TMR)



(Why not just double?)

Parallel Processing Motivation



Performance gain through clock rate
Power consumption? Implications?

Parallel Processing Overview

□ Flynn's Taxonomy

- Single instruction single data (SISD) stream
no parallel operation
- Single instruction multiple data (SIMD) stream
vector and array processors, MMX & SSE instructions
- Multiple instruction single data (MISD) stream
never implemented
- Multiple instruction multiple data (MIMD) stream
multicore, multiprocessor

Parallel Processing Overview

❑ Past: Multithreaded Single-Core Processors

- No real parallelism, just improved thread switching
- Choose two threads
- If one blocks, the other one can run w/o loading its data into register set

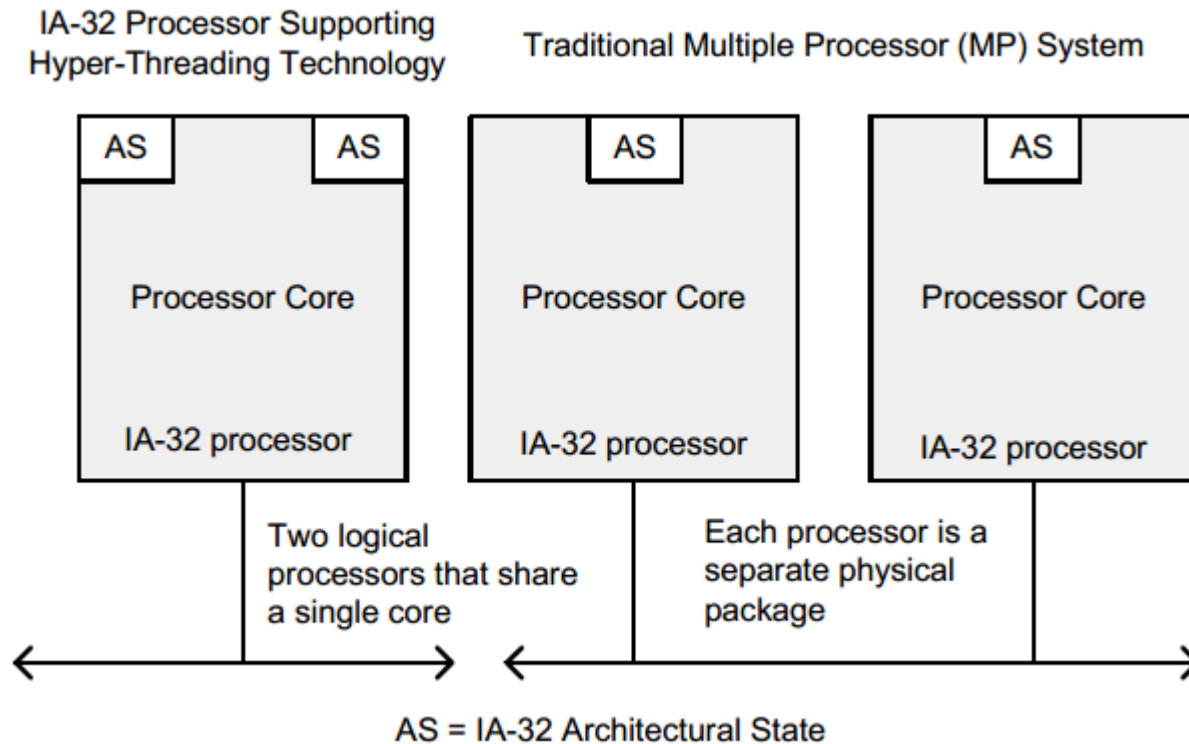
❑ Present: Homogeneous Multi-Processors (*plus* multithreading)

- Nonuniform Memory Access (NUMA)
- Uniform Memory Access (UMA)

❑ Future: Heterogeneous MPs

- General purpose CPUs
- General purpose GPUs
- Specific coprocessors
- Specific accelerators
- Reconfigurable HW

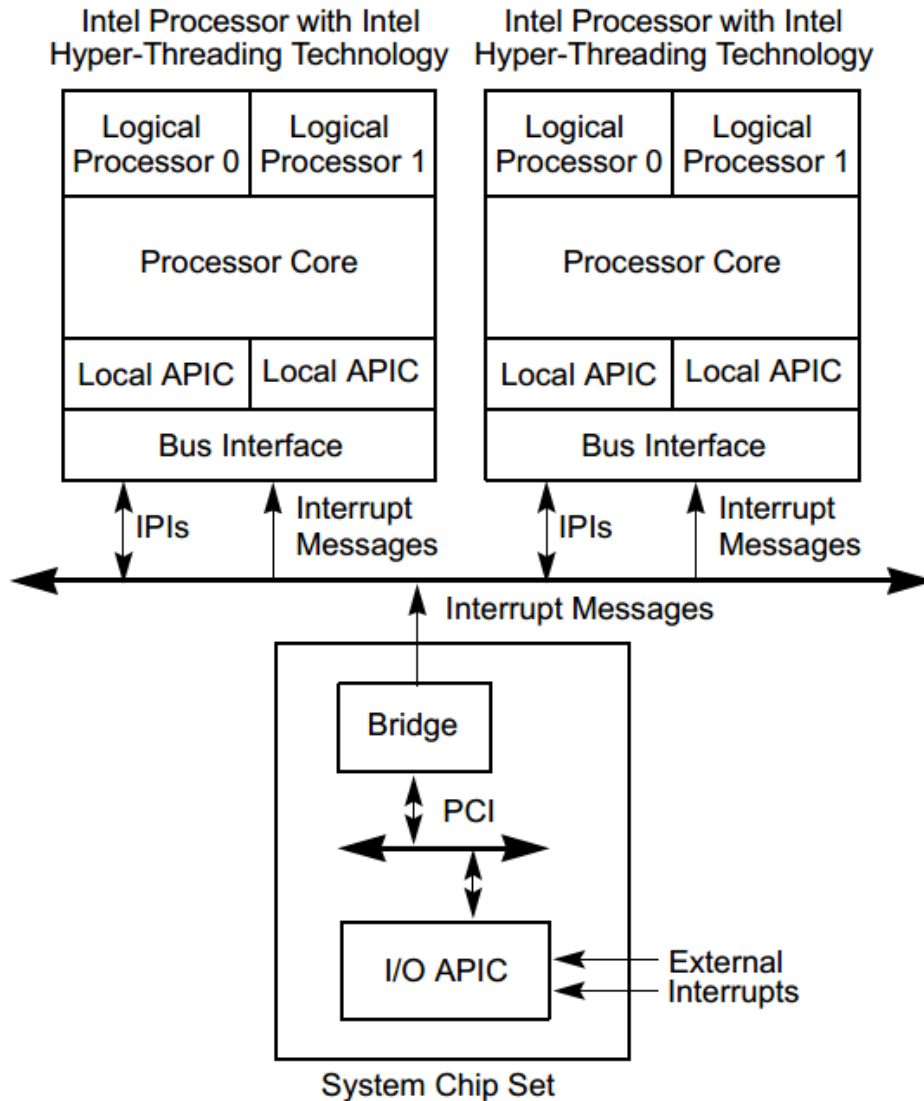
Past: Multithreaded Processors



❑ Intel "Hyperthreading" vs AMD full replication policy

- replicated register set
- save state of two threads that share address space
- speed up context switching
- BUT: shared execution engine

(Logical) Processor/Core Interaction



- ❑ Local APIC per (logical) core: Advanced Programmable Interrupt Controller
- ❑ Inter-Processor Interrupts to synchronize execution
 - start/halt processors
 - flush TLBs
 - ...
- ❑ I/O interrupts via separate APIC

Present: Homogeneous Multi-Processors


❑ Nonuniform Memory Access (NUMA)

- Memory access differs for different memory areas

❑ Uniform Memory Access (UMA)


- Memory access is identical for all memory areas

Multiprocessing Hardware Architectures




(a)

Shared-memory
multiprocessor



(b)

Message-passing
multicomputer




(c)

Distributed
System

© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall


Non-Uniform Memory Access (NUMA)



© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

- ❑ Single address space
- ❑ Load/Store interface
- ❑ Access times differ!

UMA 1: Crossbar-switched UMA Architecture




© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

Scalability?

UMA 2: Multistage Network UMA Architecture



(a)



(b)

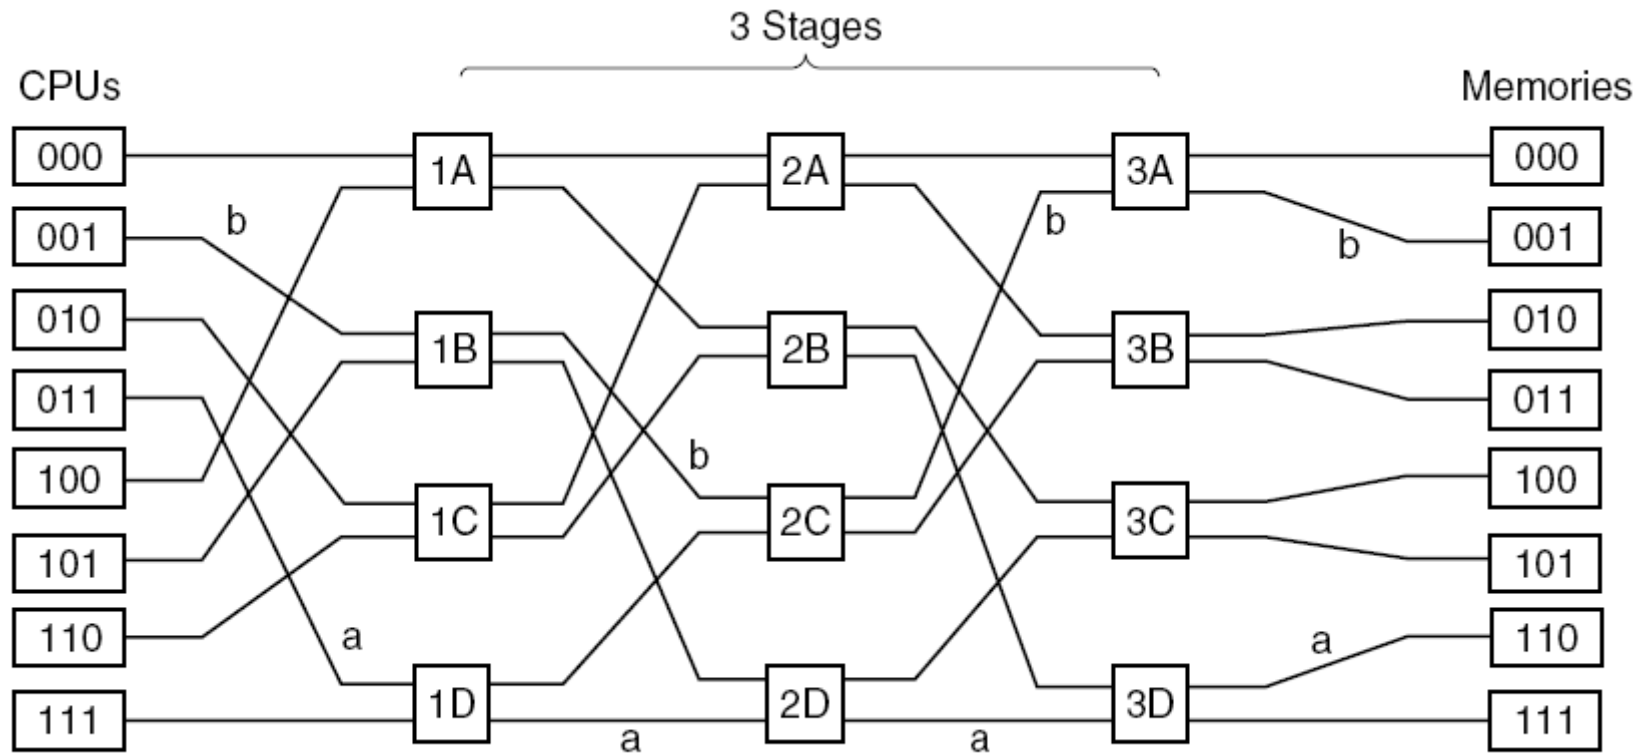
© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

Network of 2x2 switches

$A \rightarrow X, X \rightarrow A$
 $A \rightarrow Y, Y \rightarrow A$
 $B \rightarrow X, X \rightarrow B$
 $B \rightarrow Y, Y \rightarrow B$

Memory addressing scheme

Multistage Omega Network UMA Architecture



© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

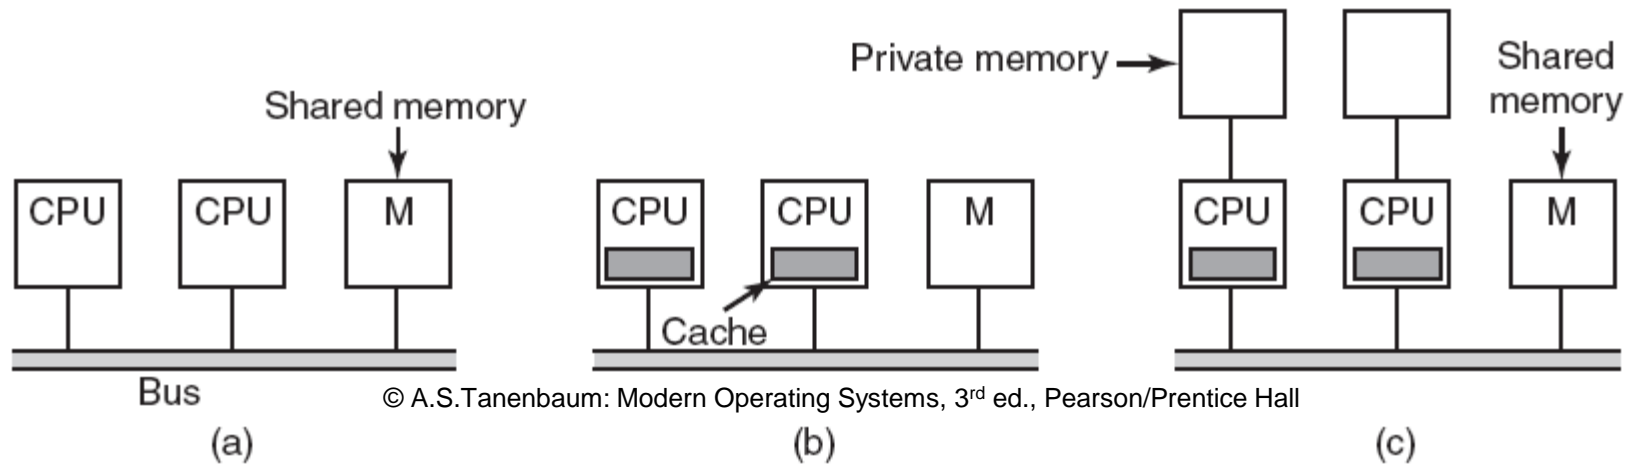
011 sends READ (a) to module 110
001 sends READ (b) to module 001

Scalability?

n CPUs \times n memory units: $n/2(\log_2 n)$ switches

Other problems?

UMA 3: Bus-based UMA Architectures



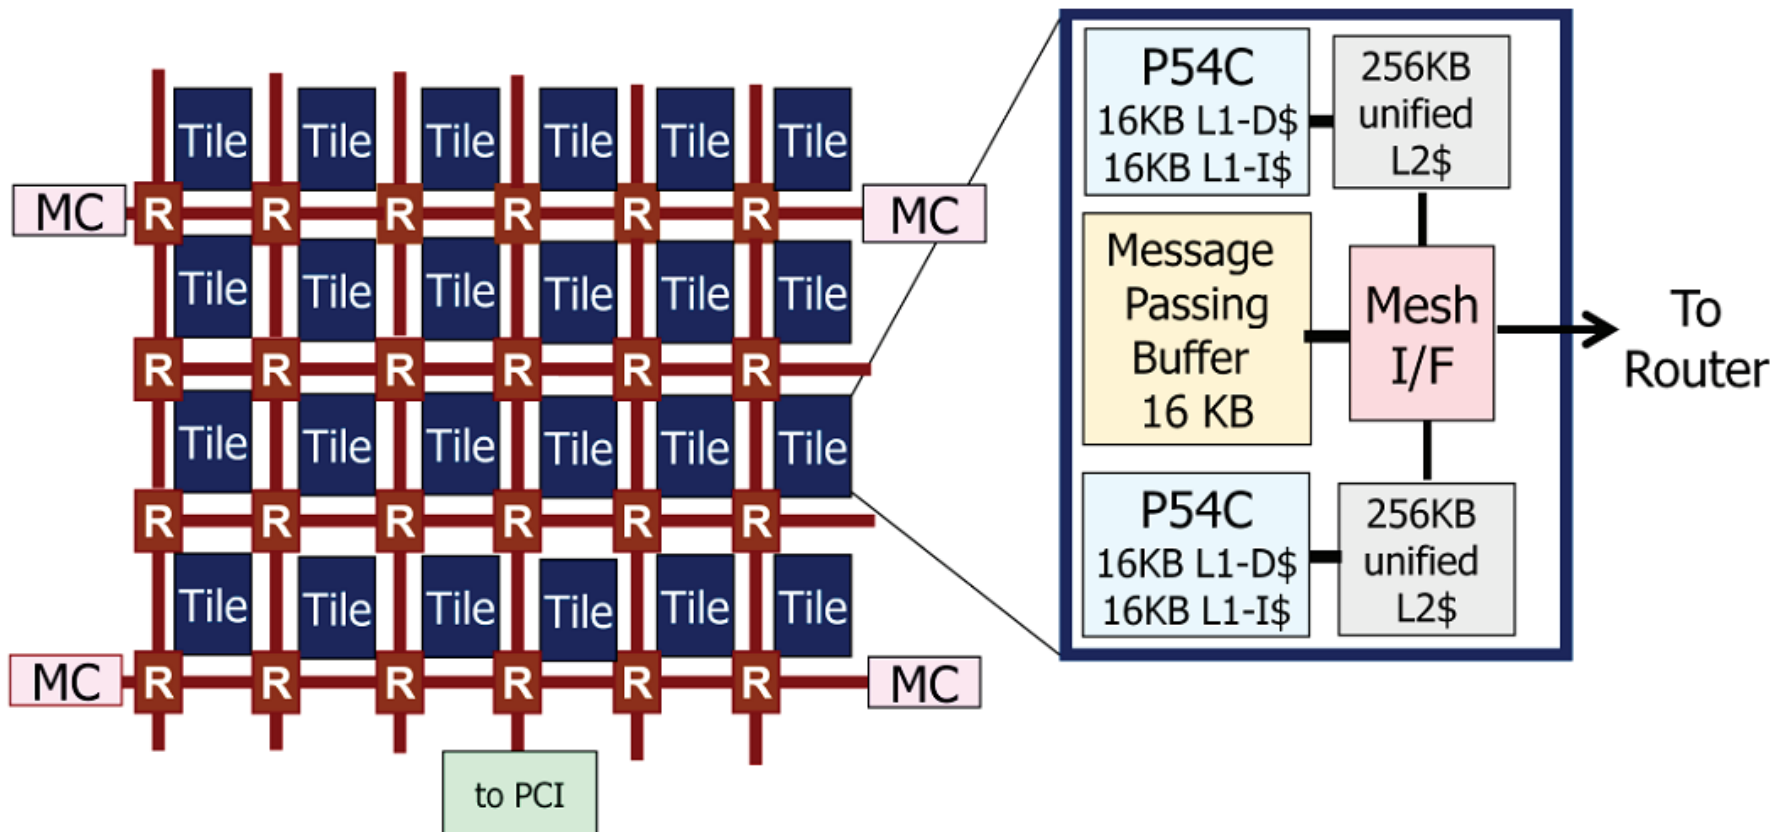
No caching

Caching

Caching +
private memory

Scalability?

Intel's Single-Chip Cloud Computer (SCC)



Parallel computer architectures summary

❑ Parallel computing motivation

- Clock rate improvement is too expensive (heat dissemination)


❑ Parallel computing architectures?

- Flynn's taxonomy: SISD, SIMD, MISD, MIMD
- MIMD: HW Multithreading, Homogeneous & Heterogeneous Multiprocessors
- Homogeneous Multiprocessors: UMA vs. NUMA
Most common implementation: shared bus
- (Heterogeneous Multiprocessors: Co-processors, accelerators)

OS Types and Implications for Bus-based HW

- ☐ Master/Slave
- ☐ Private OS per node
- ☐ Shared OS

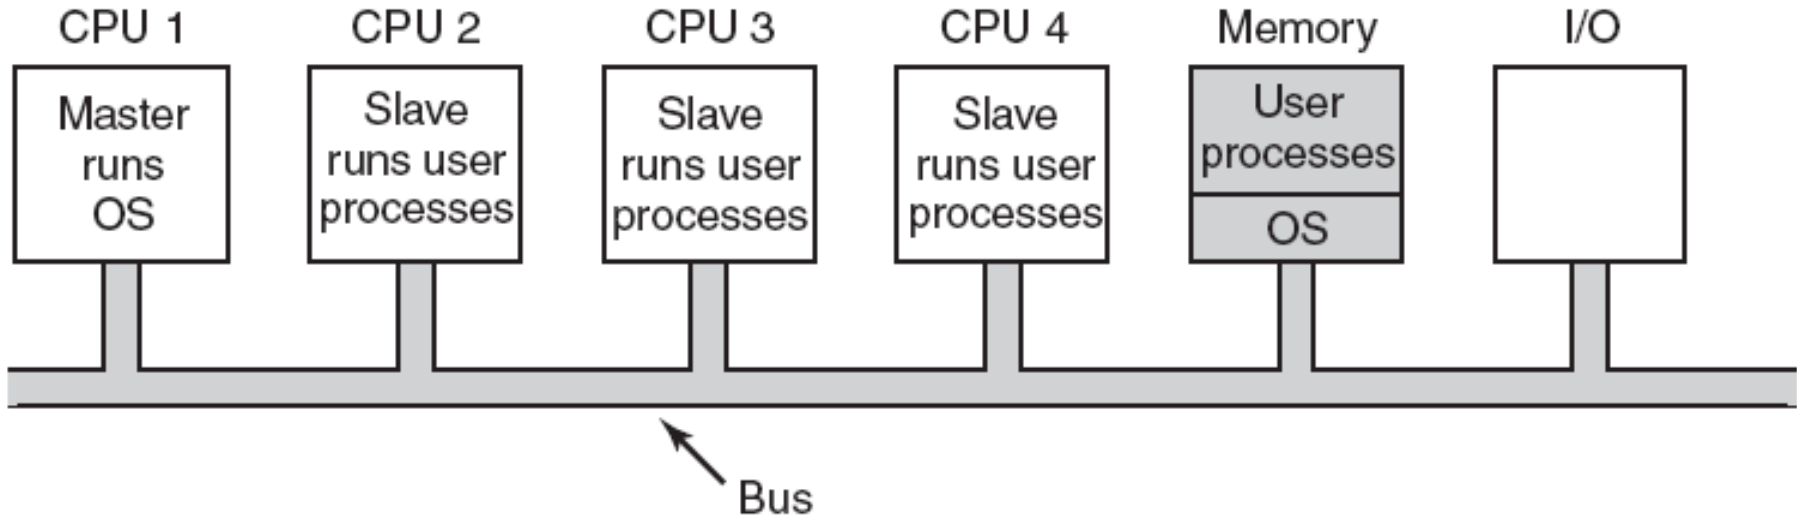
Private OS per node



© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

- ❑ Shared HW, CPUs & memory partitioned *statically*
- ❑ Soft Resources? Processes, pages, disk...

Master/Slave




© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

Central coordinating instance

Pros?

Cons?

Shared OS (Symmetric Multiprocessing - SMP)



© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

Now every user process can invoke the shared OS *locally*

→ No master CPU bottleneck, yet resource sharing

→ Other problems?

Multiprocessor OS issues

❑ Synchronization

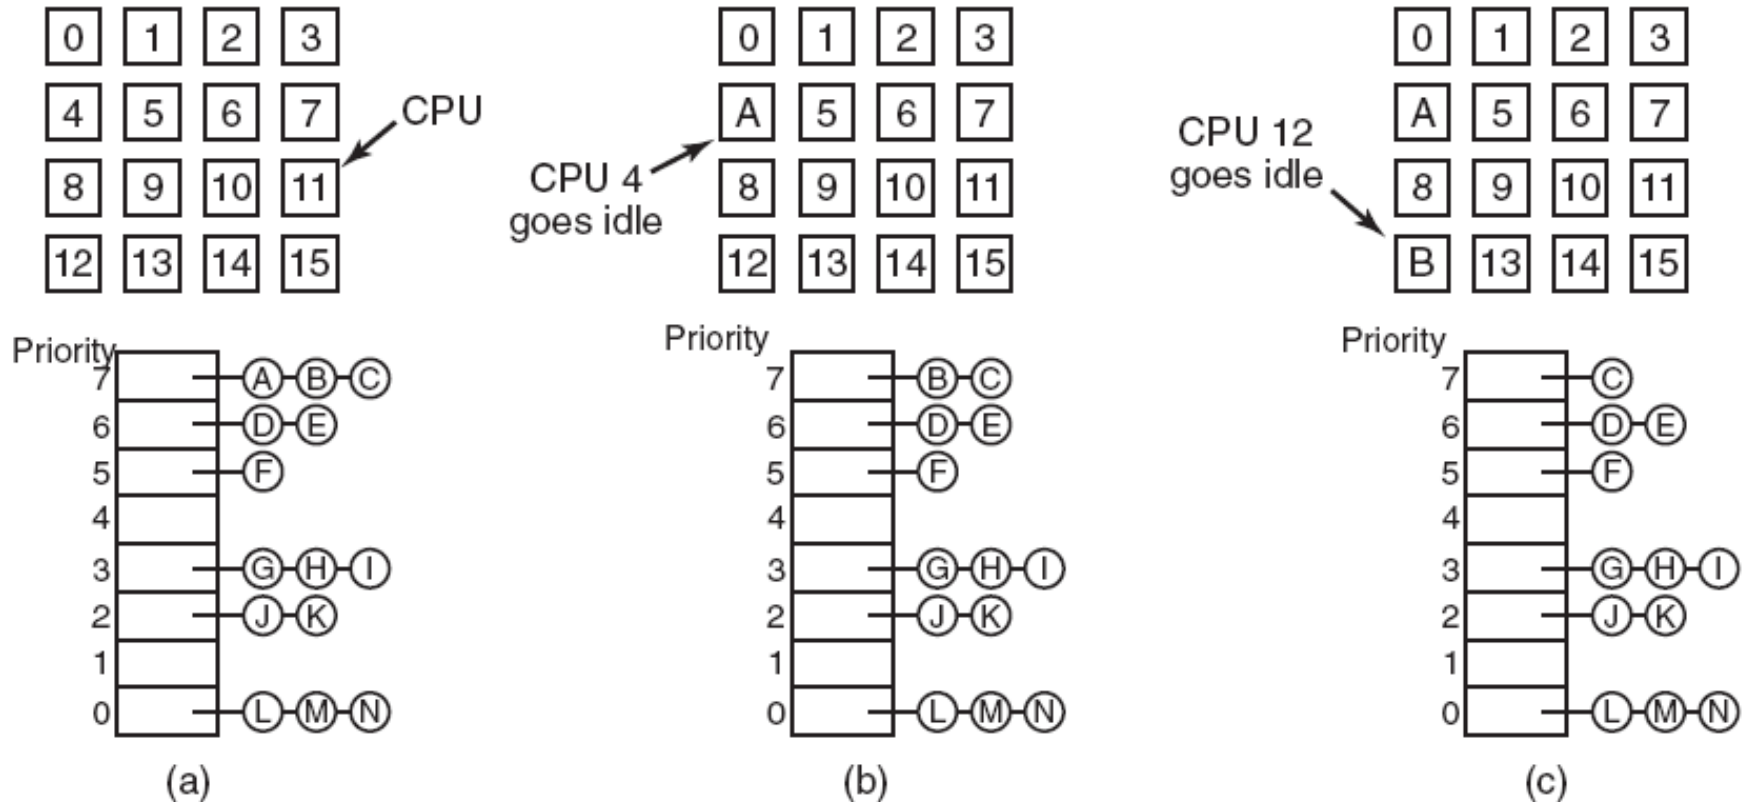
- Switching off interrupts?
- TSL?
- Covered in the DS lecture next week

❑ Scheduling

- What to run?
- Where to run?
- Which combinations?

Multiprocessor Scheduling: Time Sharing

© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

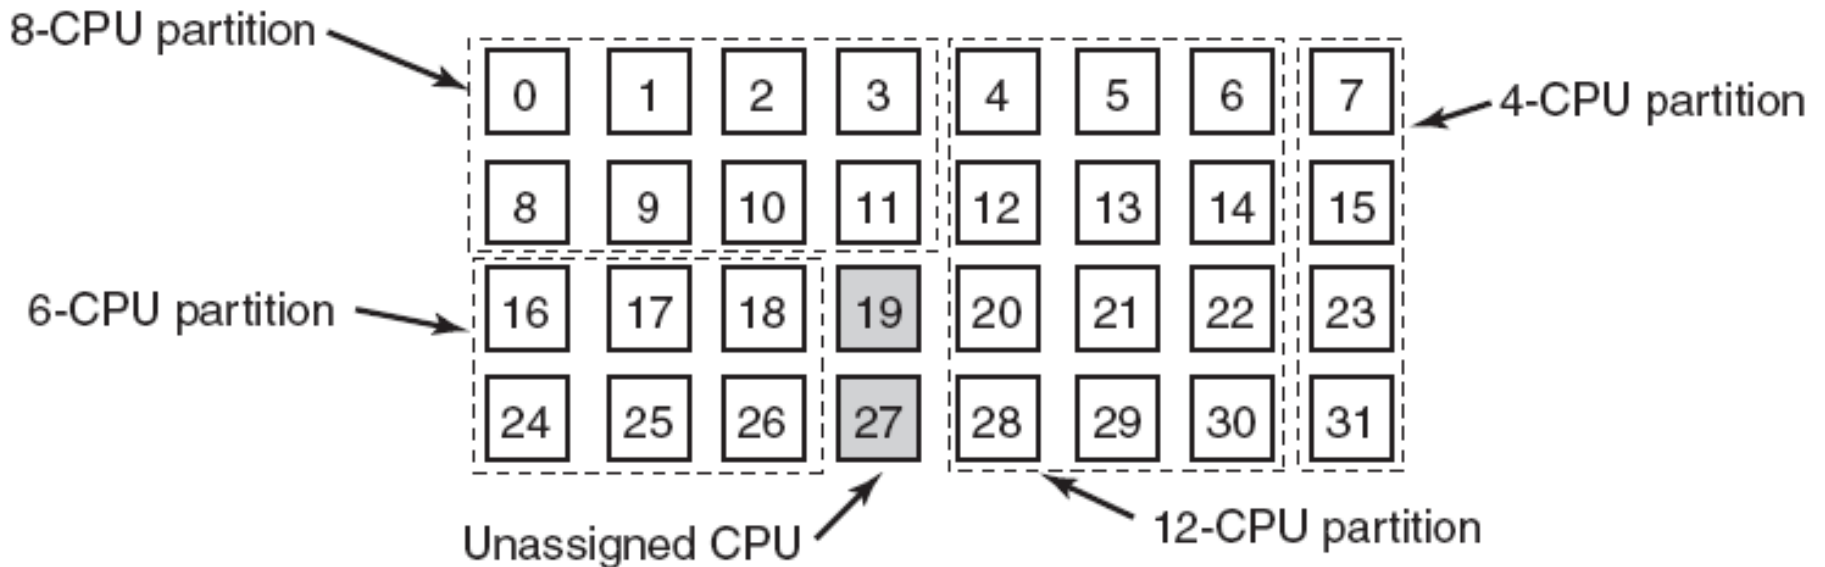


- a) All CPUs busy
- b) 4 gets idle, picks highest priority thread A
- c) 12 gets idle, picks highest priority thread B

Works well for independent threads...

Scalability? Thread interactions?

Multiprocessor Scheduling: Space Sharing




© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

- ☐ For each process the number of created threads is known
- ☐ Scheduler knows the number of free CPUs
- ☐ If enough CPUs are available, all threads of a process are scheduled simultaneously, else another process is scheduled
- ☐ All threads run to completion

Problems?

Multiprocessor Scheduling: Gang Scheduling



© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

- ❑ A_0/B_0 scheduled on CPU 0, A_1/B_1 on CPU 1
- ❑ A_0 and A_1 cooperate
- ❑ Alternating schedules: Increased communication latency!

Multiprocessor Scheduling: Gang Scheduling

		CPU					
		0	1	2	3	4	5
Time slot	0	A ₀	A ₁	A ₂	A ₃	A ₄	A ₅
	1	B ₀	B ₁	B ₂	C ₀	C ₁	C ₂
	2	D ₀	D ₁	D ₂	D ₃	D ₄	E ₀
	3	E ₁	E ₂	E ₃	E ₄	E ₅	E ₆
	4	A ₀	A ₁	A ₂	A ₃	A ₄	A ₅
	5	B ₀	B ₁	B ₂	C ₀	C ₁	C ₂
	6	D ₀	D ₁	D ₂	D ₃	D ₄	E ₀
	7	E ₁	E ₂	E ₃	E ₄	E ₅	E ₆

© A.S.Tanenbaum: Modern Operating Systems, 3rd ed., Pearson/Prentice Hall

- ☐ Interacting threads form a *Gang*
- ☐ Gangs are scheduled simultaneously
 - Requires synchronized CPU/core schedules

Playing with multiprocessor systems

❑ Bochs x86 emulator

- Supports multicore processor emulation and HW hyperthreading
- BUT: many cores \rightarrow slooowww

❑ QEMU

- Supports multicore and multiprocessor SMP emulation
- Faster than Bochs: No processor-level emulation
- (Still limited scale)

❑ HP COTSon

- Simulates multiprocessors with x86 CPUs
- Scales up to ≥ 1000 processors
- Configurable accuracy/performance trade-off

❑ MIT Graphite

- Simulation of multiprocessors on commodity Linux clusters
- Scales up to ≥ 1000 processors