



Large-Scale Parallel Computing

Aamer Shah

shah@cs.tu-darmstadt.de

EXERCISE 3

Exercise 3

- Four programs that perform matrix multiplication
 - Sequential program (base case)
 - Three parallel MPI versions
- Had to be compiled and executed on Lichtenberg-cluster
 - Measure the execution time
 - And time spent in other program sections
- How many succeeded in running it?
- How many can access the cluster?


Matrix multiplication

- Matrix M1 size [M, N]
- Matrix M2 size [N, K]

M1 x M2

Columns of M1 must be equal to rows of M2

$$\text{Prod} = P_{(i,j)} : \sum (M1_{(i,t)} \times M2_{(t,j)}) \text{ where } 1 \leq t \leq N$$



Sequential program

Output: (time in nanoseconds)

Starting app at time: 1447778337607808256

Time spent in allocation: 333312



Time spent in initialization: 110592512

Time spent in multiplication: 74643288576


Total time spent in application: 74754214400

Execution time of sequential program

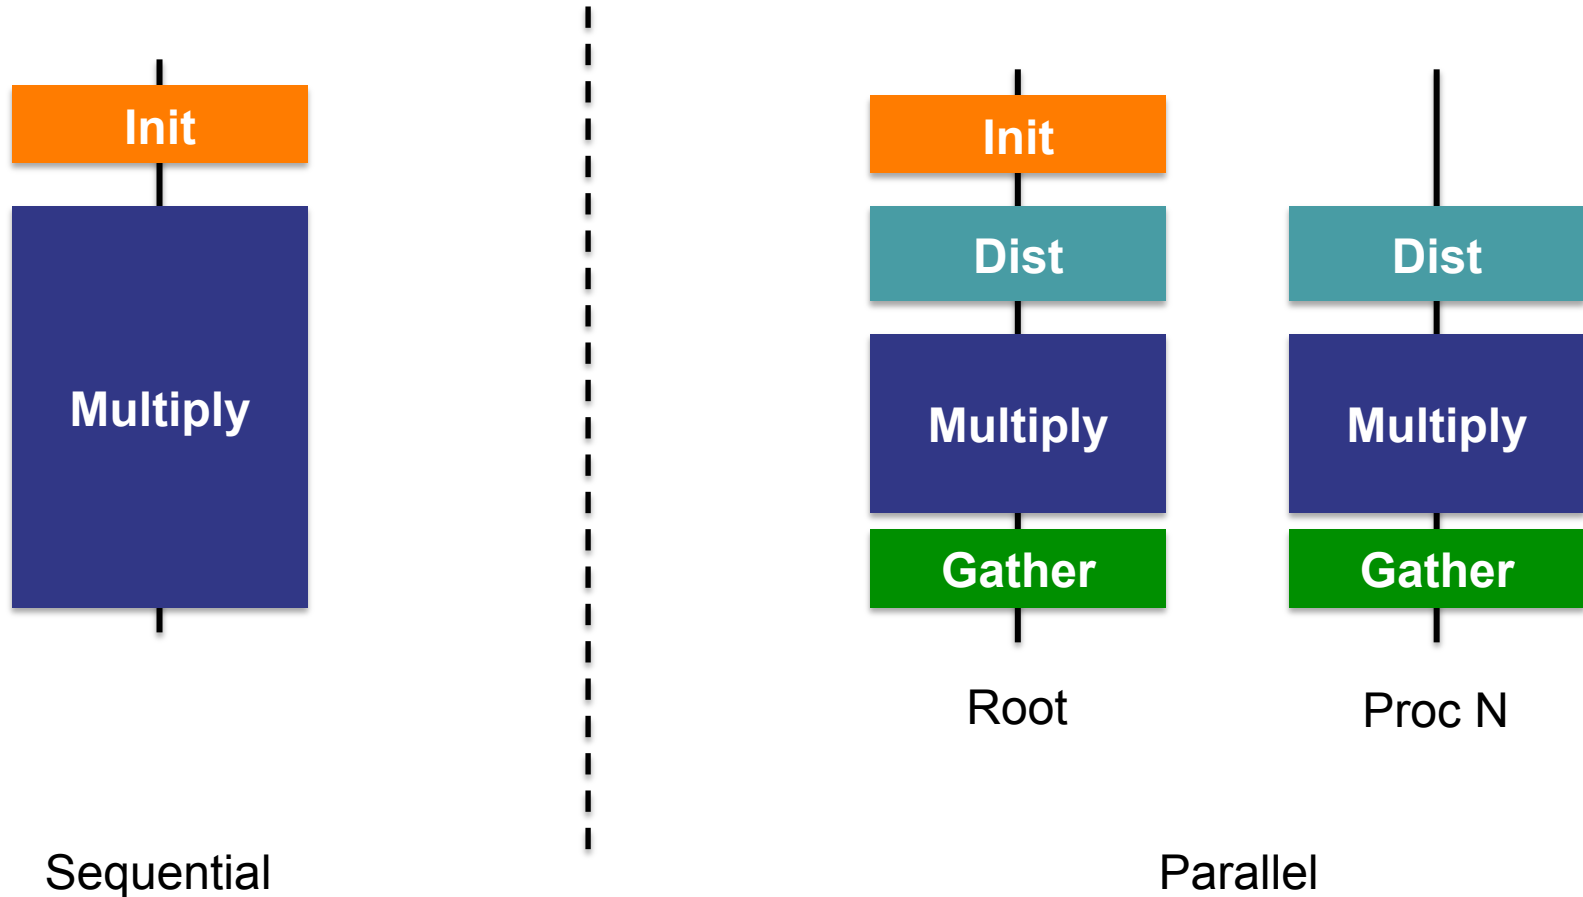
- Time distribution




Amdahl's law (3)



Process tasks – sequential vs parallel

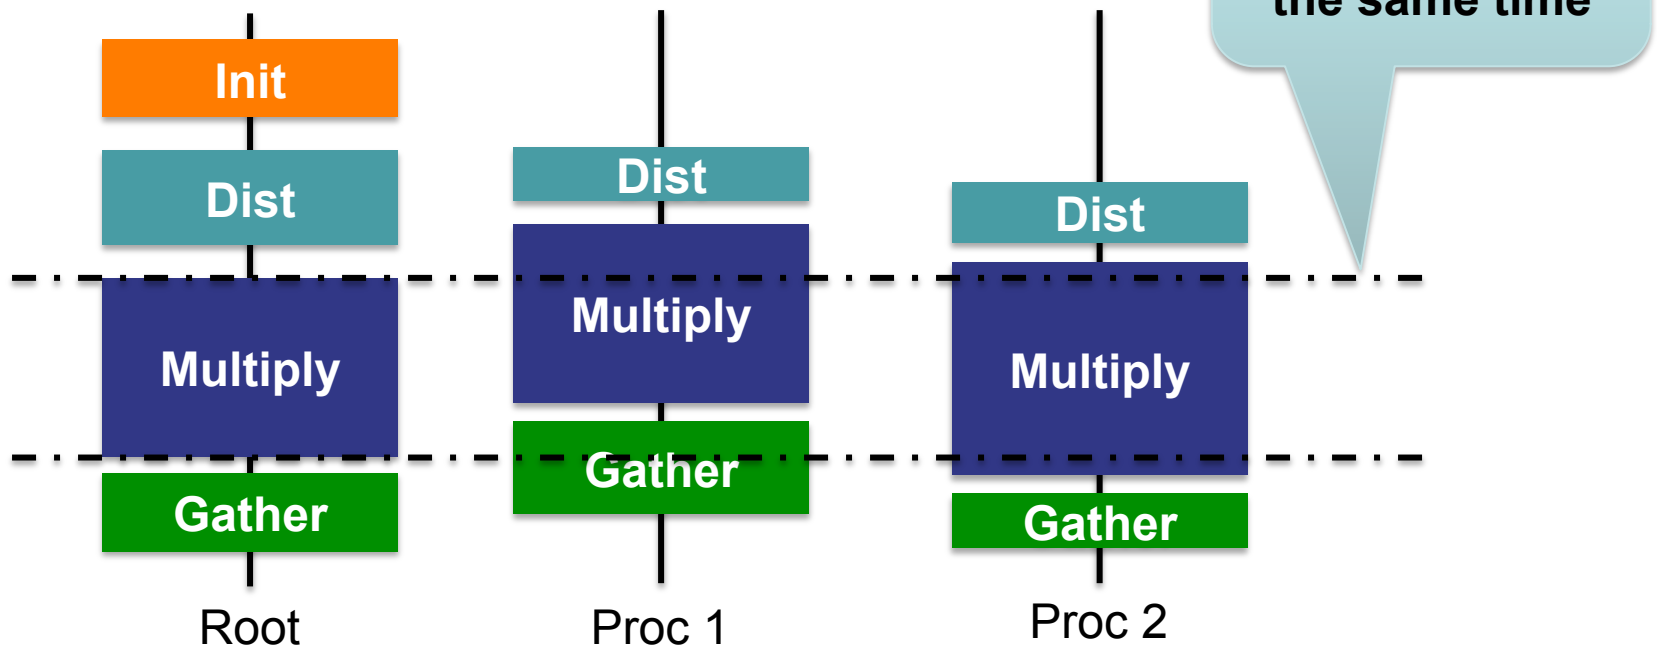


Application scaling





Reporting time

- Multiple independent processes, with synchronization in between
 - Who reports the time?




Reporting time

- Events order
 - Imbalance in computation can falsely be reported as time spent in data gathering




Data distribution

- Tile-wise distribution of data among processes

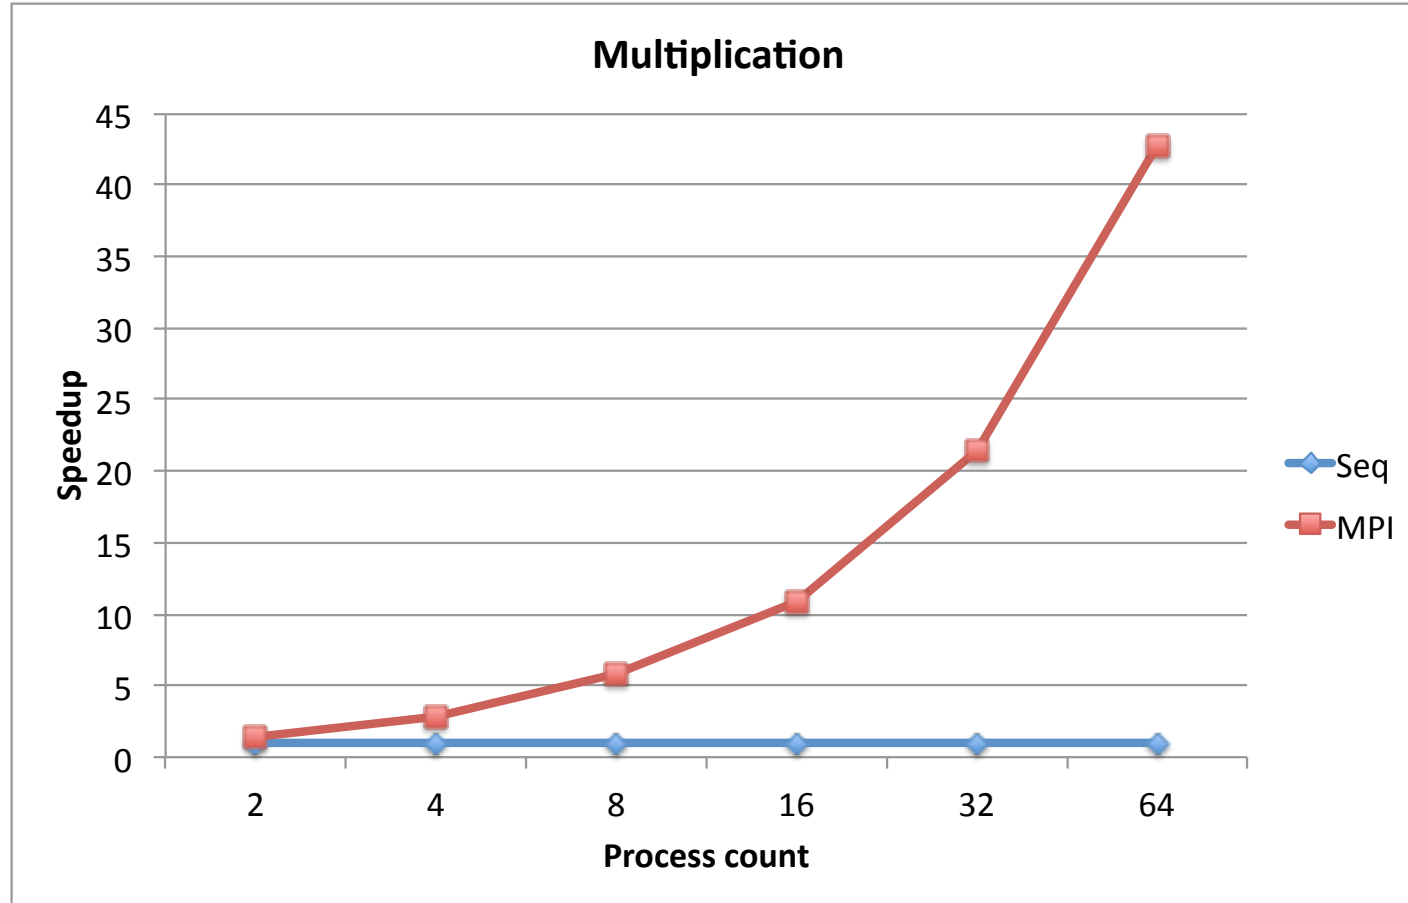


MPI - scaling


$$\text{Speedup} = \frac{\text{time parallel program}}{\text{time sequential program}}$$



MPI - scaling

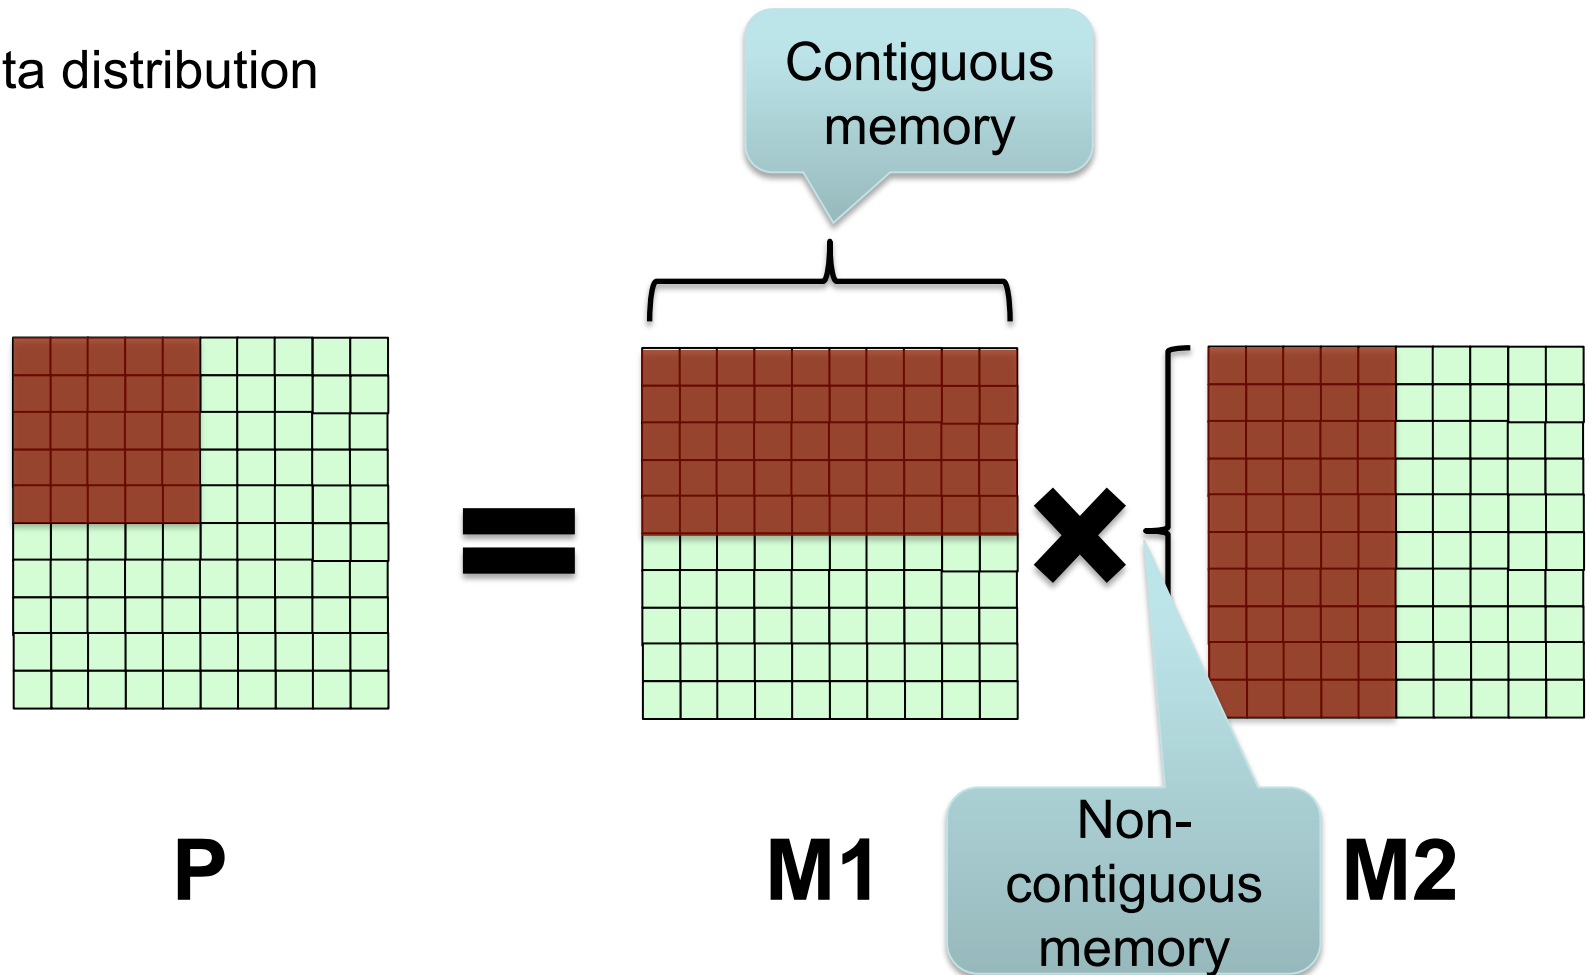


MPI – scalability graph



MPI – scalability problems

- Data distribution



MPI – scalability problem

```
int MPI_Send(const void *buf, int count, MPI_Datatype datatype, int dest, int tag,  
             MPI_Comm comm)
```

```
for(j = 0; j < cols_group_size; j++)  
{  
    int k;  
    for(k = 0; k < DIM_LEN; k++)  
        MPI_Send(&mat[k][col_disp * cols_group_size + j], 1, MPI_CHAR, i, 0, MPI_COMM_WORLD);  
}
```


Columns

```
for(j = 0; j < row_group_size; j++)  
{  
    int row_ind = row_disp * row_group_size + j;  
    MPI_Send(lmat[row_ind], DIM_LEN, MPI_CHAR, i, 0, MPI_COMM_WORLD);  
}
```


Rows

MPI – scalability problem


- Matrix multiplication
 - Non-contiguous memory access
 - Cache misses on M2



MPI – scalability problem



$$P(1,1) = M1(1,1) \times M2(1,1) + M1(1,2) \times M2(2,1) + \dots$$



MPI – scalability problem

Solution?


MPI – scalability problem - transpose

- Matrix M1 size [M, N]
- Matrix M2 size [N, K] \rightarrow $M2^T$ size [K, N]

M1 x M2


Columns of M1 must be equal to rows of M2

Prod = $P_{(i,j)} : M1_{(i,t)} \times M2^T_{(j,t)}$ where $1 \leq t \leq N$



MPI – scalability problem

- Use transpose of M2
 - Solves data distribution problem
 - Solves cache misses problem



Trans – data distribution

```
int MPI_Send(const void *buf, int count, MPI_Datatype datatype, int dest, int tag,  
             MPI_Comm comm)
```


```
int col_disp = (i % col_div);  
for(j = 0; j < cols_group_size; j++)  
{  
    MPI_Send(mat[col_disp * cols_group_size + j], DIM_LEN, MPI_CHAR, i, 0,  
             MPI_COMM_WORLD);  
}
```

Columns


```
for(j = 0; j < row_group_size; j++)  
{  
    int row_ind = row_disp * row_group_size + j;  
    MPI_Send(lmat[row_ind], DIM_LEN, MPI_CHAR, i, 0, MPI_COMM_WORLD);  
}
```

Rows


Trans – cache misses



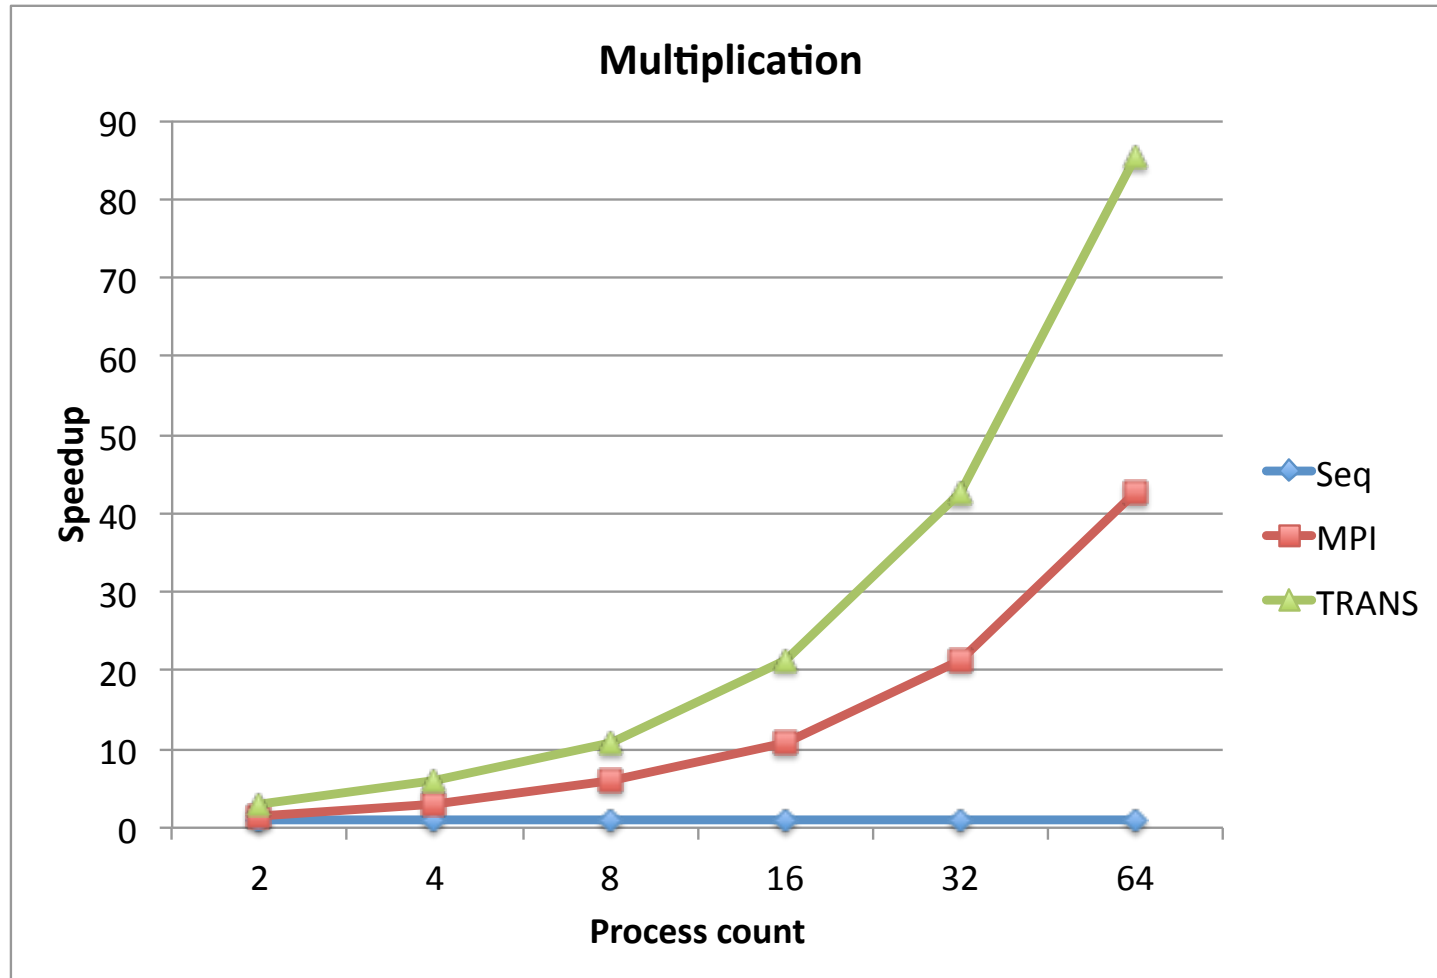
$$P(1,1) = M1(1,1) \times M2(1,1) + M1(1,2) \times \textcolor{red}{M2(1,2)} + \dots$$



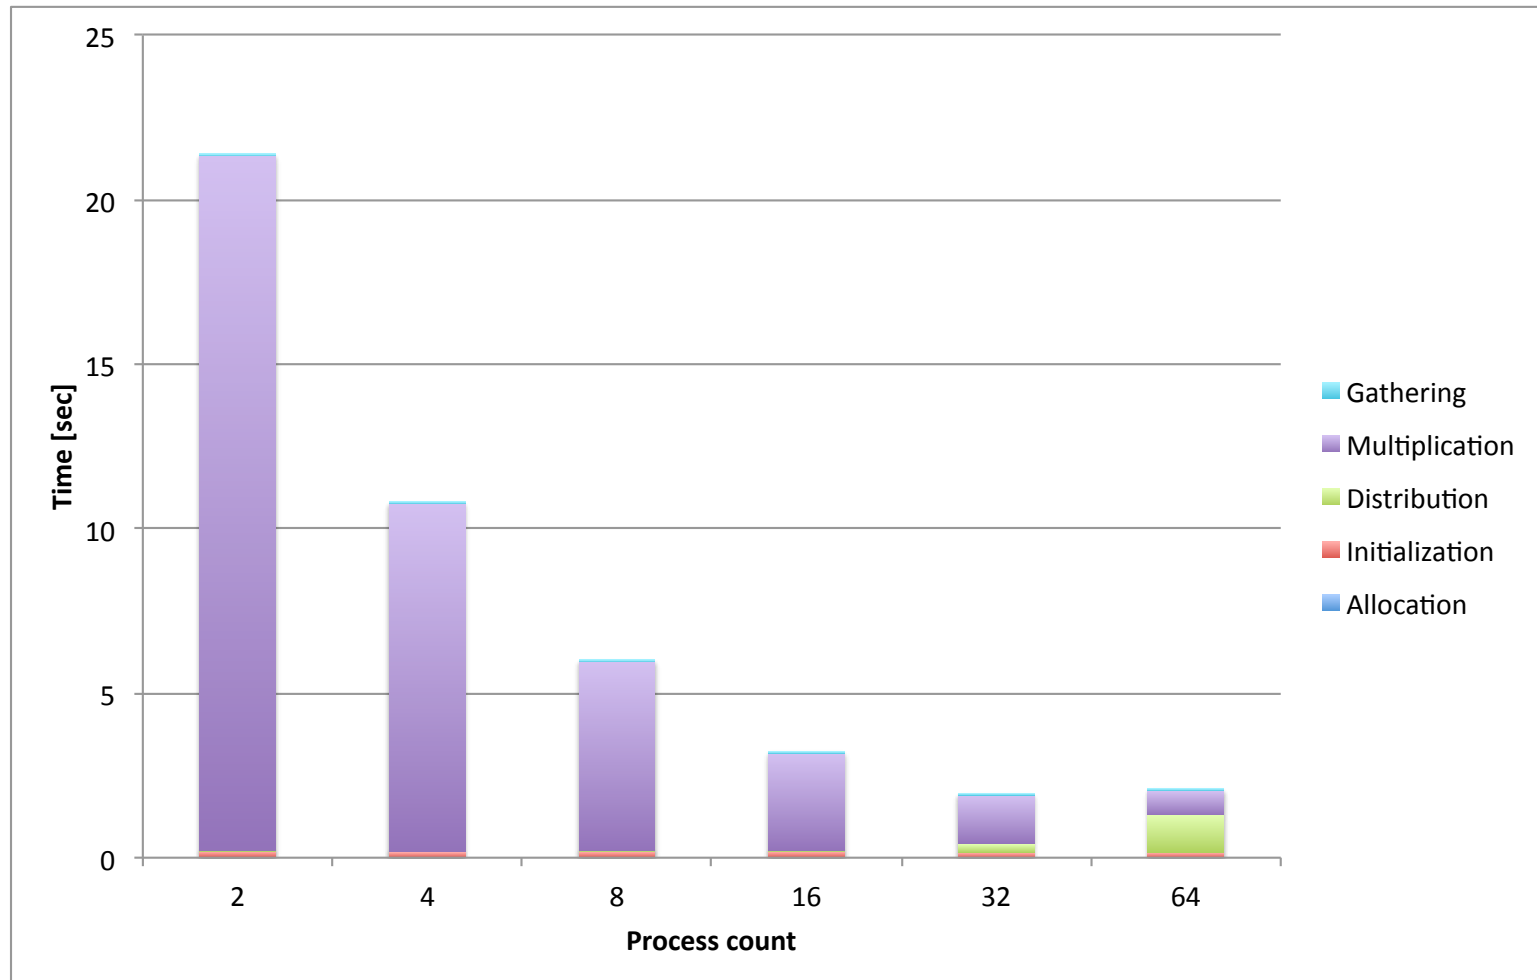
Trans - scaling



Trans - scaling

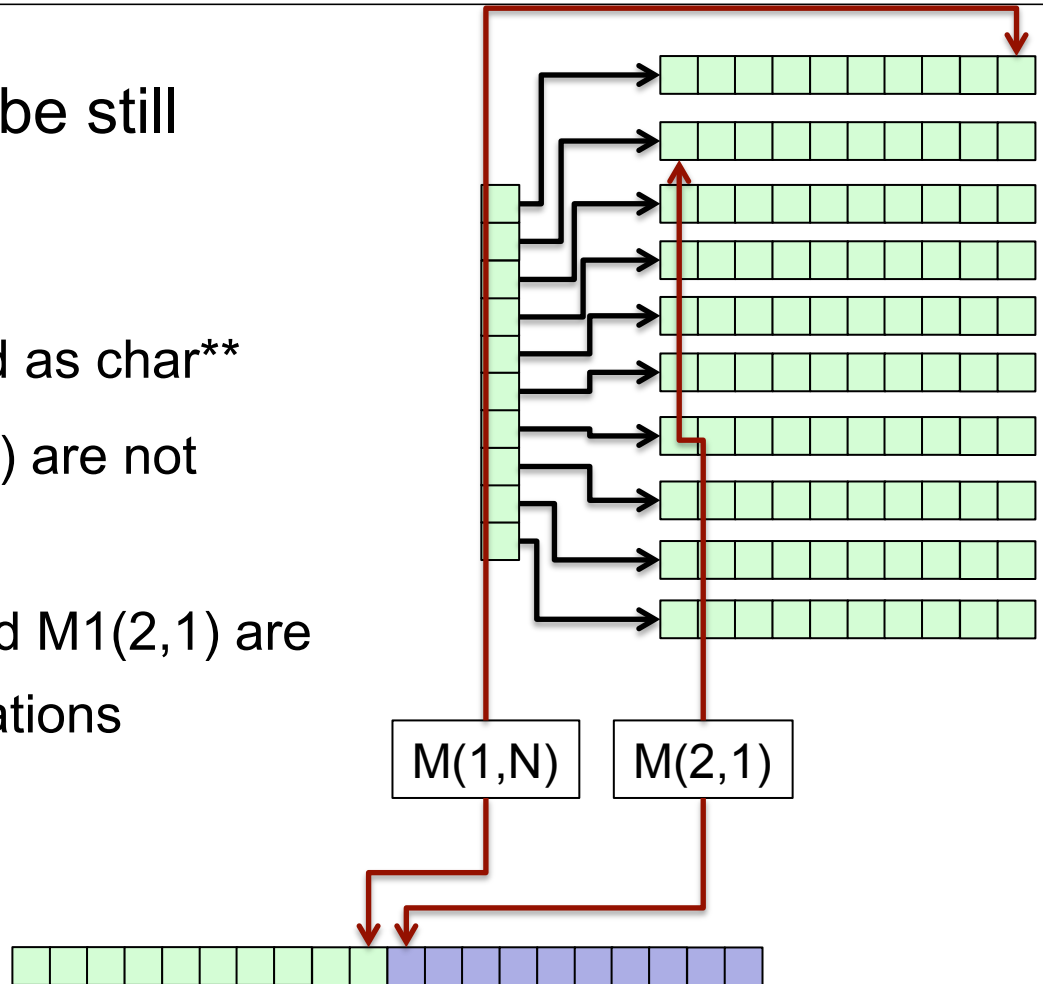


Trans - scalability




Any further improvements?

- Can the application be still improved?
- M1 and M2 are declared as `char**`
 - M1(1, N) and M1(2,1) are not contiguous
- In 2D array, M1(1,N) and M1(2,1) are contiguous memory locations




Potential optimization



$$P(1,1) = M1(1,1) \times M2(1,1) + M1(1,2) \times M2(2,1) + \dots$$

$$P(2,1) = M1(2,1) \times M2(2,1) + M1(2,2) \times M2(2,2) + \dots$$


- Cache miss if non-contiguous memory
- Cache hit if contiguous memory



Potential optimizations


- For data distribution, MPI provides utility functions that have better performance

```
int MPI_Scatter(const void *sendbuf, int sendcount, MPI_Datatype  
sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype, int  
root, MPI_Comm comm)
```

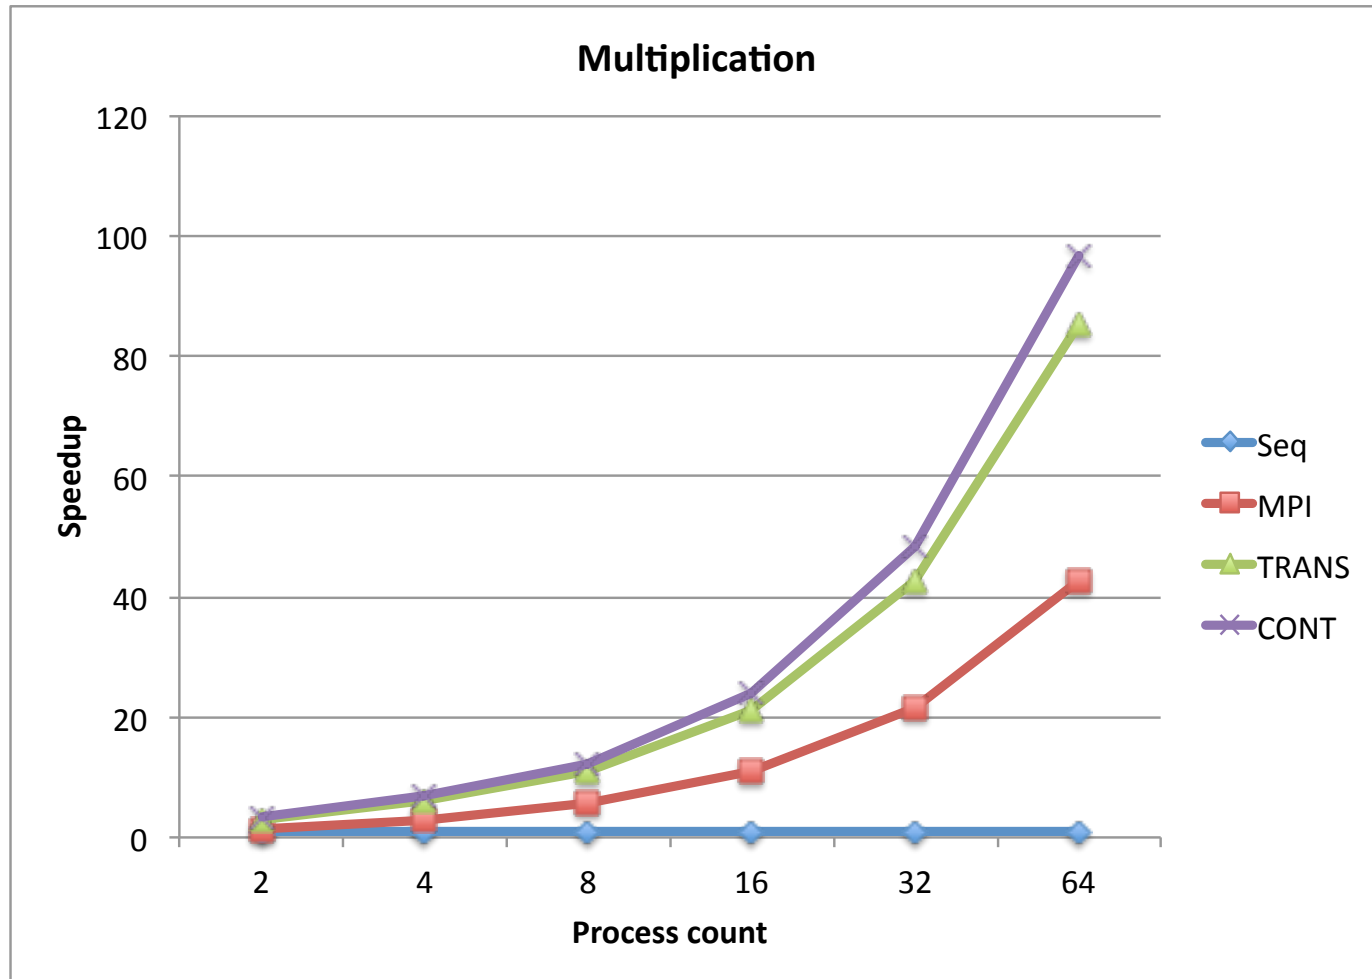


```
MPI_Scatterv(mat, scount, displacement, MPI_CHAR, mat,  
cols_group_size, MPI_CHAR, 0, MPI_COMM_WORLD);
```

Cont - scaling



Cont - scaling



Cont - scalability

