



Large-Scale Parallel Computing

Aamer Shah

shah@cs.tu-darmstadt.de


EXERCISE 6

Hands-on session

- Hands-on session
 - Students will develop the solution during the exercise session
-
1. Login to Lichtenberg cluster (with `-Y` option)
 2. Copy `ex06.tgz` from `/home/as65huly/public`
 3. Implement Task 1 in `matmul_dist.c`
 - Derived data types
 4. Implement Task 2 in `matmul_nbc.c`
 - Non-blocking communication

Derived data types

- What is the problem? Matrix multiplication as an example
 - How to distribute matrix columns among processes?
 - Column is in non-contiguous memory location
 - Send element-by-element?
 - Tedious and takes time
 - Is there a way to tell MPI what a column is and then send it?



Type contiguous

- Different ways to define different types of derived data types
 - Create data types for arrays / rows of a matrix

```
int MPI_Type_contiguous(int count, MPI_Datatype  
oldtype, MPI_Datatype *newtype)
```

```
MPI_Type_contiguous(10, MPI_INT, &mat_row)
```

1	2	3	4	5	6	7	8	9	0
2									
3									
4									
5									
6									
7									
8									
9									
0									

Type contiguous

- `MPI_Type_contiguous()` can be used for a matrix row
- A single instance can cover only one row or multiple rows

```
MPI_Type_contiguous(50, MPI_INT, &mat_rows)
```

```
MPI_Send(mat, 1, mat_rows, dest, tag, MPI_COMM_WORLD)
```

1	2	3	4	5	6	7	8	9	0
2									
3									
4									
5									
6									
7									
8									
9									
0									

Type vector

- How to transfer columns?
 - Non contiguous in memory, BUT equal distance between each element

```
int MPI_Type_vector(int count, int blocklength, int stride,  
MPI_Datatype oldtype, MPI_Datatype *newtype)
```

```
MPI_Type_vector(10, 1, 10, MPI_INT, &mat_col)
```

**blocklength = 1,
only 1 contiguous
element**

**count = 10,
10 elements in
1 column**

1	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	

**stride = 10,
Two consecutive
elements in a
column have a
distance of 10
elements between
their starts**

Type vector

- A single instance can cover either one column or multiple columns

```
MPI_Type_vector(10, 4, 10, MPI_INT, &mat_cols)
```

```
MPI_Send(mat, 1, mat_cols, dest, tag, MPI_COMM_WORLD)
```

**blocklength = 4,
4 contiguous
element**

**count = 10,
10 elements in
1 column**

1	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	

**stride = 10,
Two consecutive
elements in a
column have a
distance of 10
elements between
their starts**

Type vector

- How to make a data type for a tile?

```
MPI_Type_vector(5, 5, 10, MPI_INT, &mat_tile)
```

```
MPI_Send(mat, 1, mat_tile, dest, tag, MPI_COMM_WORLD);
```

**blocklength = 5,
5 contiguous
element**

**count = 5,
5 elements in
1 column**

1	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	

**stride = 10,
Two consecutive
elements in a
column have a
distance of 10
elements between
their starts**

Size and extent?

- Whats happens with this?
 - What is the share of the second process?

```
MPI_Type_vector(5, 5, 10, MPI_INT, &mat_tile)
```

```
MPI_Scatter(mat, 1, mat_tile, proc_mat, 1, mat_tile, 0, MPI_COMM_WORLD);
```

Rank 0 share starts here:

- Size is the amount of space needed to store the data

1	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	

- Extent includes the holes in between data types
- MPI communication functions use extent of data types

Rank 1 share starts here:



How to distribute columns/tiles?

```
MPI_Send(mat, 1, mat_tile, 0, tag, MPI_COMM_WORLD); //Rank 0
```

```
MPI_Send(&mat[tile_row_size], 1, mat_tile, 1, tag, MPI_COMM_WORLD); //Rank 1
```

```
MPI_Send(&mat[mat_row_size * tile_column_size], 1, mat_tile, 2, tag,  
MPI_COMM_WORLD); //Rank 2
```

```
MPI_Send(&mat[mat_row_size * tile_column_size + tile_row_size], 1,  
mat_tile, 2, tag, MPI_COMM_WORLD); //Rank 3
```

Rank 0 share starts
here:

Rank 1 share starts
here:

Rank 2 share starts
here:

Rank 3 share starts
here:


1	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	

Task 1

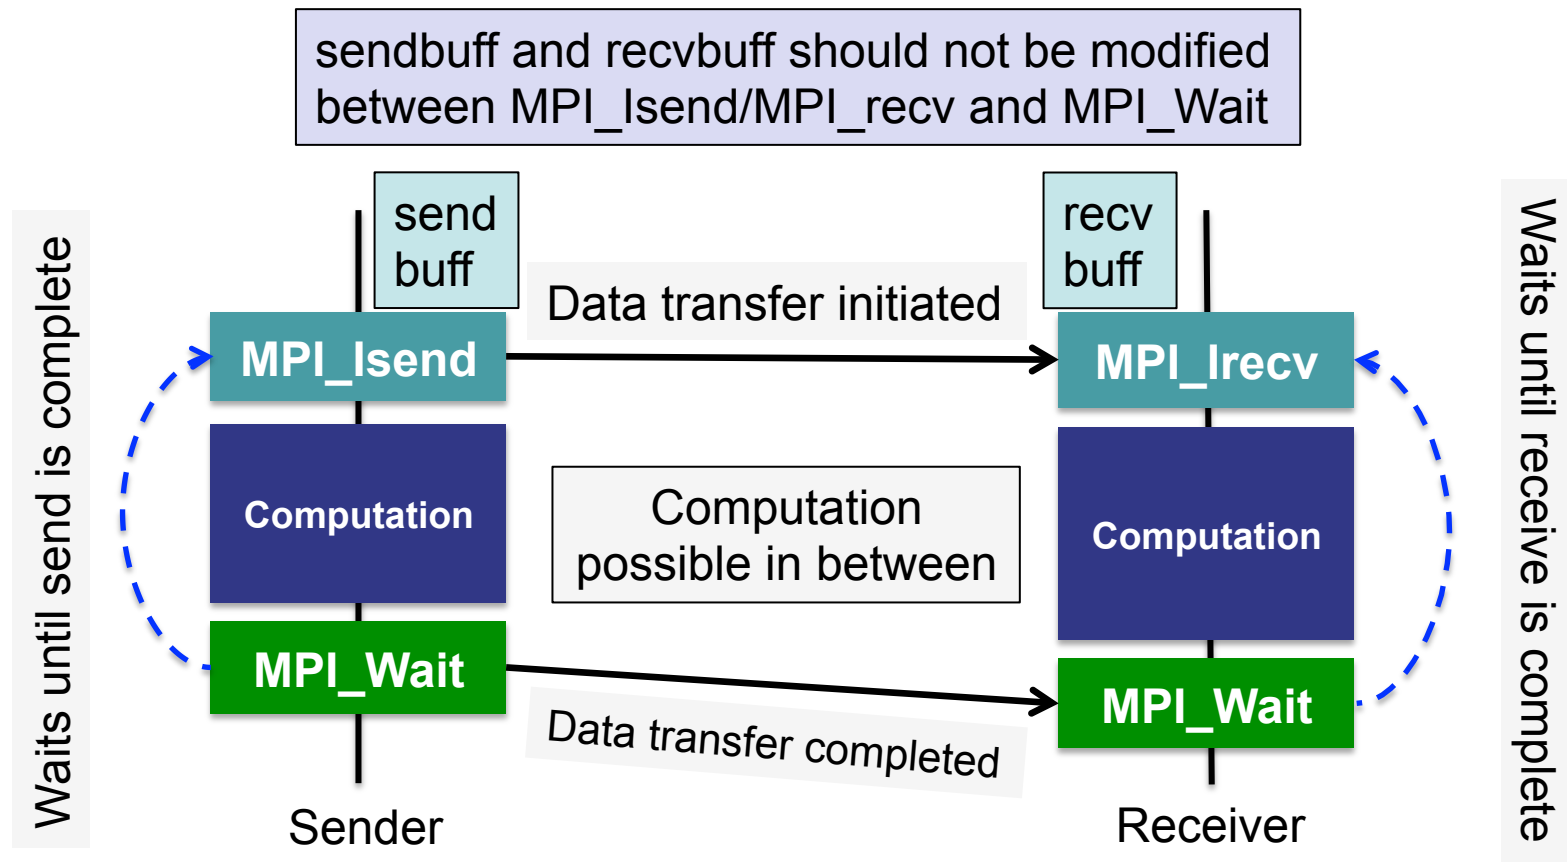
- Matrix multiplication where each process calculates a product tile
- Use derived data types to distribute rows, columns and tiles
 - Can rows/columns be distributed by MPI_Scatter?
 - Can the tiles be collected by MPI_Gather?
- Steps
 1. Create data types for rows/columns/tiles
 2. Distribute rows
 3. Distribute columns
 4. Multiply
 5. Collect tiles

Non-blocking communication

- Distributing data among processes takes time
- Not all data is needed for each process to **start** multiplication
- If computation can be started when enough data is transferred to a process, the data distribution time can be reduced
- Non-blocking communication : overlap computation and communication



Non-blocking communication



Non-blocking communication

```
int MPI_Isend(void *buf, int count, MPI_Datatype datatype, int dest,  
int tag, MPI_Comm comm, MPI_Request *request)
```

```
int MPI_Irecv (void *buf, int count, MPI_Datatype datatype, int source,  
int tag, MPI_Comm comm, MPI_Request *request)
```

```
int MPI_Wait(MPI_Request *request, MPI_Status *status)
```


```
int MPI_Waitall(int count, MPI_Request *array_of_requests, MPI_Status  
*array_of_statuses)
```

```
int MPI_Waitany(int count, MPI_Request *array_of_requests, int *index,  
MPI_Status *status)
```


```
int MPI_Waitsome(int count, MPI_Request *array_of_requests, int  
*numcompl, int *indices, MPI_Status *statuses)
```

Matrix multiplication



- Transfer rows and columns one-by-one



1. Receive a row and multiply it with already transferred columns



2. Receive a column and multiply it with already transferred rows



Task 2

- Use non-blocking communication to transfer rows and columns
- The data type should be defined such that
 - Row type is two rows of the actual matrix
 - Column type is two columns of actual matrix
- Use element-wise multiplication algorithm to find product
- Collect product at the end using the same method as Task 1