

# Software Defined Networking

Past, Current, and Future




TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

David Hausheer

Department of Electrical Engineering  
and Information Technology  
Technische Universität Darmstadt

E-Mail: [hausheer@ps.tu-darmstadt.de](mailto:hausheer@ps.tu-darmstadt.de)  
<http://www.ps.tu-darmstadt.de/teaching/sdn>



\*Original slides for this lecture provided by Scott Shenker (UC Berkeley) and Bernhard Plattner (ETH Zürich)

# Lecture Overview

---

- ❖ A Brief History of the Internet
- ❖ The Networking World Prior to SDN
- ❖ The Future of Networking, and the Past of Protocols



# A Brief History of the Internet

Bernhard Plattner

# A Brief History of the Internet

- ❖ 1962: ARPANET – network of computers
- ❖ 1973: Internet – network of networks
  - Gradually, ARPANET becomes core network
  - Applications: e-mail, interactive host access, simple c/s apps (e.g. gopher, 1991)
- ❖ 1984: First international e-mail in Switzerland (UUCP, X.400)
- ❖ 1986: NSFNet becomes core network
- ❖ 1987: The Internet starts in Europe
- ❖ 1993: Classless inter-domain routing
- ❖ 1993: IP Next Generation white paper solicitation
- ❖ 1994: Internet discovered by the public
  - The web
  - The era of global client/server communication starts

# A Brief History of the Internet

- ❖ 1998: First P2P App, Napster
  - Age of overlay networks starts
  - Google founded as a private company (IPO 2004)
- ❖ Late 90-ties: Mobile ad-hoc networks
- ❖ 2001: „Structured“ P2P systems using DHT
- ❖ 2002: SIP-based Internet telephony
- ❖ 2002: Interstellar Internet => DTN => Opportunistic Networks
- ❖ 2003: Skype
- ❖ 2004: Facebook
  - Social network era starts

# A Brief History of the Internet

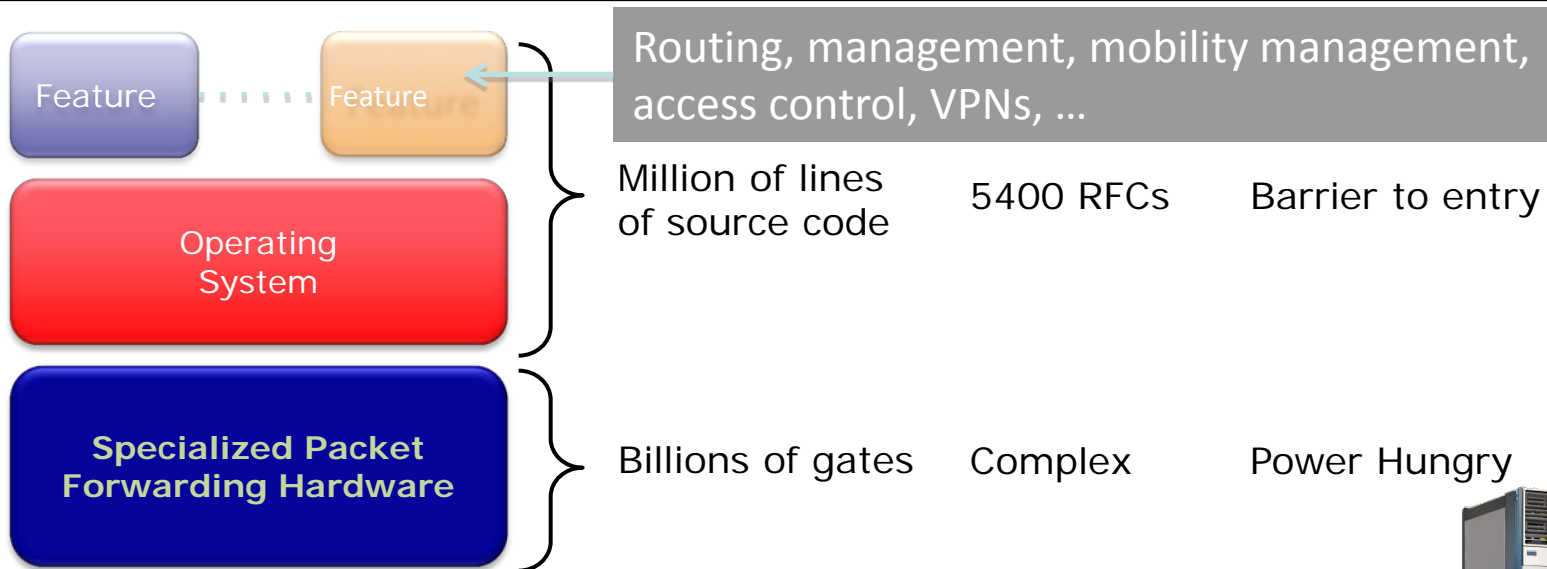
- ❖ 2005: YouTube
  - TV distribution with P2P technology (Zattoo)
- ❖ 2006: Research towards the Future Internet: US NSF
  - Find Program, FP6 Situated and Autonomic Communication, EU FIRE Program
- ❖ 2006: Cloud computing (by Amazon)
- ❖ 2007: iPhone
- ❖ 2009: SDN, currently one of the hottest topics
- ❖ 2010: Start of EU Project OFELIA: OpenFlow/SDN testbed



# The Networking World Prior to SDN

Bernhard Plattner

# Reminder: The Networking Industry (2007)



Closed, vertically integrated, boated, complex, proprietary

Many complex functions baked into the infrastructure

*OSPF, BGP, multicast, differentiated services,  
Traffic Engineering, NAT, firewalls, MPLS, redundant layers, ...*

Little ability for non-telco network operators to get what they want

Functionality defined by standards, put in hardware, deployed on nodes





# Research Stagnation: Faster networks but not *better* networks

- ❖ Lots of *deployed* innovation in other areas
  - Operating Systems: various filesystems, schedulers, virtualization
  - Distributed Systems: DHTs, CDNs, MapReduce
  - Compilers: JITs, new language paradigms
- ❖ Networks are largely the same as years ago
  - Ethernet, IPv4, WiFi – IPv6 took ages
- ❖ Rate of change in networking seems slow in comparison
  - Need better tools and abstractions to demonstrate and deploy
- ❖ Researchers need flexible and changeable platforms for experimentation

# Another problem: Closed Systems (Vendor Hardware)

---

- ❖ Can't extend – vertically integrated
- ❖ Stuck with interfaces (CLI, SNMP, etc)
- ❖ Hard to meaningfully extend
- ❖ Hard to meaningfully collaborate

- ❖ What should be the architecture of a future Internet?
  - Beyond IPv4 and IPv6
- ❖ Future Internet Architecture (FIA) project by the US National Science Foundation (2010)
- ❖ Global Environment for Network Innovations (GENI) 2006
- ❖ Future Internet and Experimentation (FIRE) Program by the European Commission's 7<sup>th</sup> Framework Program(2006)
- ❖ Clean Slate Approach

# The networking researcher's playground

Approach	Example	Performance Fidelity	Scale	Real User Traffic?	Complexity	Open
<b>“ideal system”</b>	n/a	<b>High</b>	<b>High</b>	<b>Yes</b>	<b>Low</b>	<b>Yes</b>
Simulation	NS-2, NS-3, Opnet	medium	medium	<b>no</b>	medium	yes
Emulation	Mininet	medium	<b>low</b>	<b>no</b>	medium	yes
<a href="#">Software Switches</a>	Linux box, Click router	<b>poor</b>	<b>low</b>	yes	medium	yes
HW impl.	<a href="#">NetFPGA</a>	high	<b>low</b>	yes	<b>high</b>	yes
<a href="#">Network Processors</a>	Intel IXP family	high	medium	yes	<b>high</b>	yes
<b>Vendor Switches</b>		<b>high</b>	<b>high</b>	<b>yes</b>	<b>low</b>	<b>No</b>

So let's open up the vendor switches!



# The Future of Networking, and the Past of Protocols

Scott Shenker  
with Martín Casado, Teemu Koponen, Nick McKeown  
(and many others....)

# Key to Internet Success: Layers

Applications

...built on...

Reliable (or unreliable) transport

...built on...


Best-effort global packet delivery

...built on...

Best-effort local packet delivery

...built on...

Physical transfer of bits



# Why Is Layering So Important?

- ❖ Decomposed delivery into fundamental components
- ❖ Independent but compatible **innovation** at each layer
- ❖ A practical success of unprecedented proportions...
- ❖ ...**but an academic failure**

# Built an Artifact, Not a Discipline

## ❖ Other fields in “systems”: OS, DB, DS, etc.

- Teach basic principles
- Are easily managed
- Continue to evolve

## ❖ Networking:

- Teach big bag of protocols
- Notoriously difficult to manage
- Evolves very slowly



# Why Does Networking Lag Behind?

- ❖ Networks used to be simple: Ethernet, IP, TCP....
- ❖ New **control** requirements led to great complexity
  - Isolation → VLANs, ACLs
  - Traffic engineering → MPLS, ECMP, Weights
  - Packet processing → Firewalls, NATs, middleboxes
  - Payload analysis → Deep packet inspection (DPI)
- ❖ Mechanisms designed and deployed independently
  - Complicated “control plane” design, primitive functionality
  - Stark contrast to the elegantly modular “data plane”

# Two Key Definitions

## ❖ **Data Plane:** processing and delivery of packets

- Based on state in routers and endpoints
- E.g., IP, TCP, Ethernet, etc.
- Fast timescales (per-packet)

## ❖ **Control Plane:** establishing the state in routers


- Determines how and where packets are forwarded
- Routing, traffic engineering, firewall state, ...
- Slow time-scales (per control event)

# Infrastructure Still Works!

- ❖ **Only** because of “our” ability to master complexity
- ❖ This ability to master complexity is both a blessing...
  - **...and a curse!**

# A Simple Story About Complexity

- ❖ ~1985: Don Norman visits Xerox PARC
  - Talks about user interfaces and stick shifts



# What Was His Point?

- ❖ The ability to ***master complexity*** is not the same as the ability to ***extract simplicity***
- ❖ When first getting systems to work....
  - Focus on mastering complexity
- ❖ When making system easy to use and understand
  - Focus on extracting simplicity
- ❖ **You will never succeed in extracting simplicity**
  - If don't recognize it is different from mastering complexity

# Scott Shenker's Point

- ❖ Networking still focused on mastering complexity
  - Little emphasis on extracting simplicity from control plane
  - No recognition that there's a difference....
  
- ❖ Extracting simplicity builds intellectual foundations
  - **Necessary for creating a discipline....**
  - **That's why networking lags behind**

# A Better Example: Programming

- ❖ Machine languages: no abstractions
  - Mastering complexity was crucial
- ❖ Higher-level languages: OS and other abstractions
  - File system, virtual memory, abstract data types, ...
- ❖ Modern languages: even more abstractions
  - Object orientation, garbage collection,...

## Abstractions key to extracting simplicity

# "The Power of Abstraction"

**"Modularity based on abstraction  
is the way things get done"**

Barbara Liskov

**Abstractions → Interfaces → Modularity**

*What abstractions do we have in networking?*



# Layers are Great Abstractions

---

- ❖ Layers only deal with the **data plane**
- ❖ We have no powerful ***control plane*** abstractions!
- ❖ How do we find those control plane abstractions?
- ❖ Two steps: ***define*** problem, and then ***decompose*** it.

# The Network Control Problem

- ❖ Compute the configuration of each physical device
  - E.g., Forwarding tables, ACLs,...
- ❖ Operate without communication guarantees
- ❖ Operate within given network-level protocol

***Only people who love complexity  
would find this a reasonable request***

# Programming Analogy

- ❖ What if programmers had to:
  - Specify where each bit was stored
  - Explicitly deal with all internal communication errors
  - Within a programming language with limited expressability
  
- ❖ Programmers would redefine problem:
  - Define a higher level abstraction for memory
  - Build on reliable communication abstractions
  - Use a more general language
  
- ❖ **Abstractions** divide problem into tractable pieces
  - And make programmer's task easier

# From Requirements to Abstractions

1. Operate without communication guarantees  
Need an abstraction for **distributed state**
2. Compute the configuration of each physical device  
Need an abstraction that **simplifies configuration**
3. Operate within given network-level protocol  
Need an abstraction for general **forwarding model**

***Once these abstractions are in place,  
control mechanism has much easier job!***

# Shenker's Entire Talk in One Sentence


- ❖ SDN is defined *precisely* by these three abstractions
  - Distribution, forwarding, configuration
- ❖ SDN not just a random good idea...
  - Fundamental validity and general applicability
- ❖ SDN may help us *finally* create a discipline
  - Abstractions enable reasoning about system behavior
  - Provides environment where formalism can take hold....
- ❖ OK, but what are these abstractions?

# 1. Distributed State Abstraction


---

- ❖ Shield control mechanisms from state distribution
  - While allowing access to this state
- ❖ Natural abstraction: ***global network view***
  - Annotated network graph provided through an API
- ❖ Implemented with “Network Operating System”
- ❖ Control mechanism is now program using API
  - No longer a distributed protocol, now just a graph algorithm
  - E.g. Use Dijkstra rather than Bellman-Ford

# Network of Switches and/or Routers



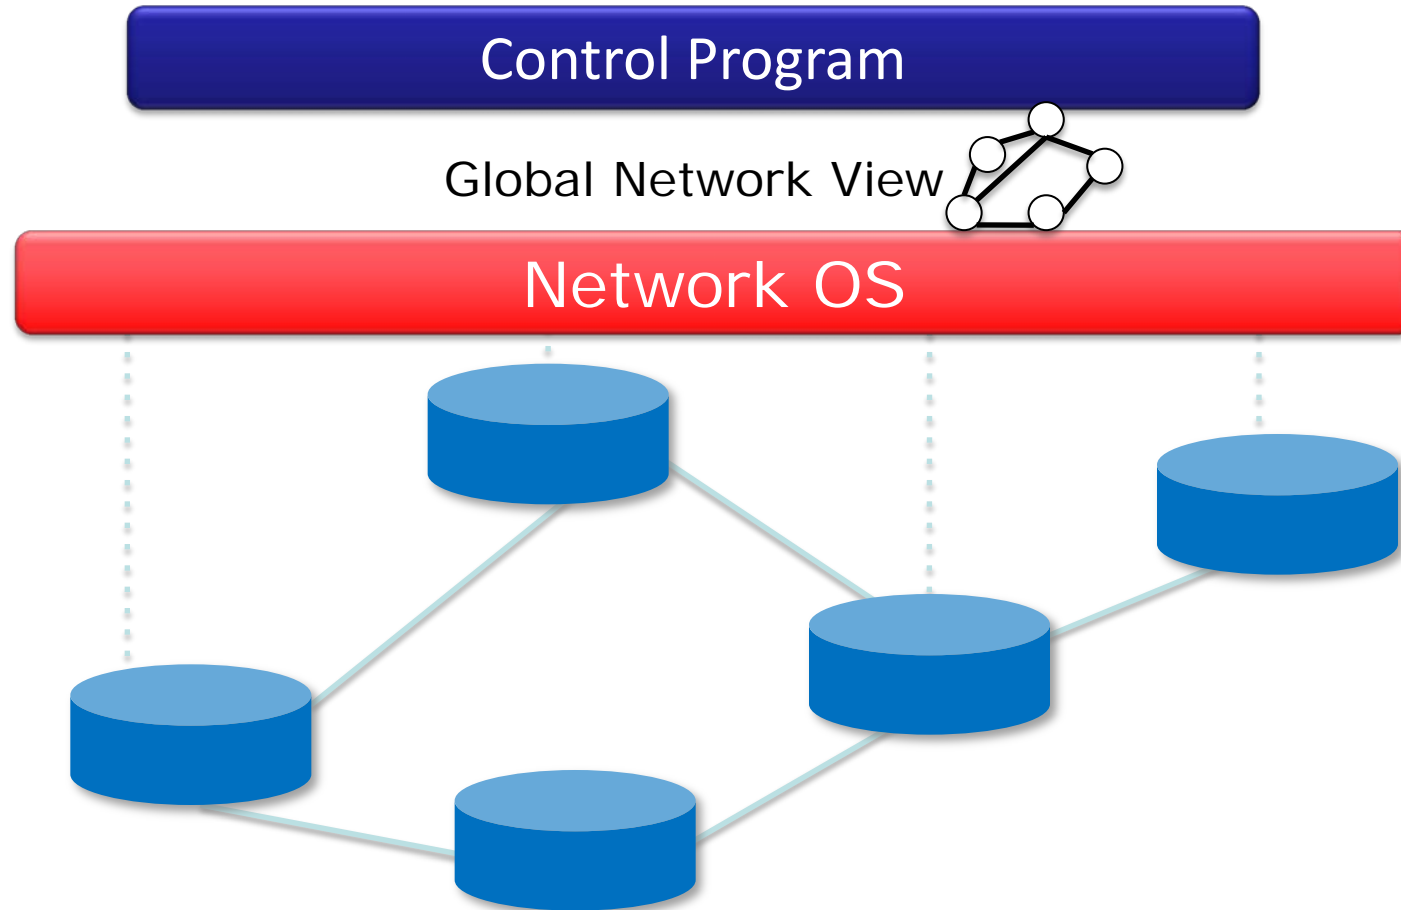
Distributed algorithm running between neighbors





# Software Defined Network (SDN)

*e.g. routing, access control*



# Major Change in Paradigm


- ❖ No longer designing distributed control protocols
  - Design one distributed system (NOS)
  - Use for all control functions
- ❖ Now just defining a centralized control *function*

**Configuration = Function(view)**

## 2. Specification Abstraction


- ❖ Control program should express desired behavior
- ❖ It should not be responsible for implementing that behavior on physical network infrastructure
- ❖ Natural abstraction: **simplified model** of network
  - Simple model with only enough detail to specify goals
- ❖ Requires a new shared control layer:
  - **Map abstract configuration to physical configuration**
- ❖ This is “network virtualization”

# Simple Example: Access Control



## What

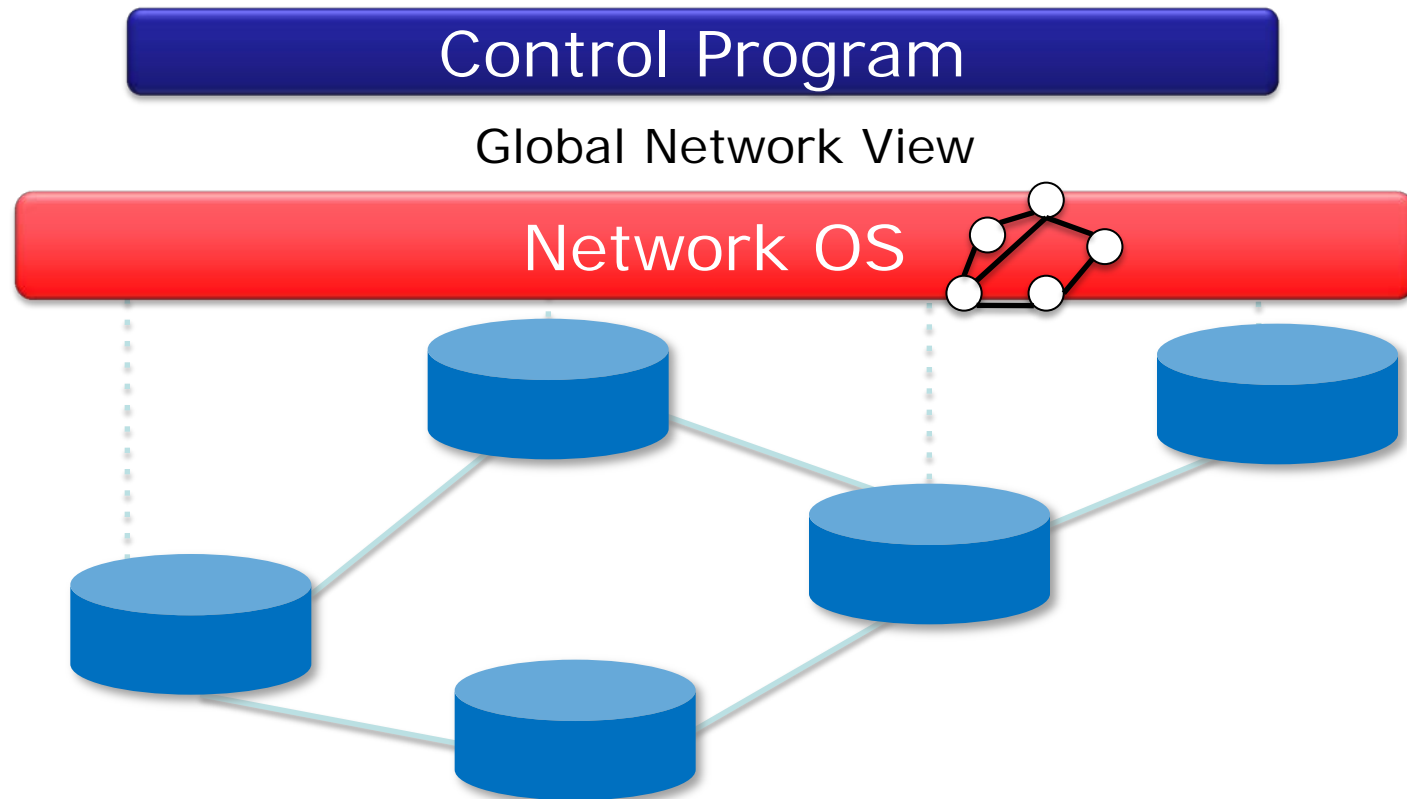
Abstract  
Network  
Model



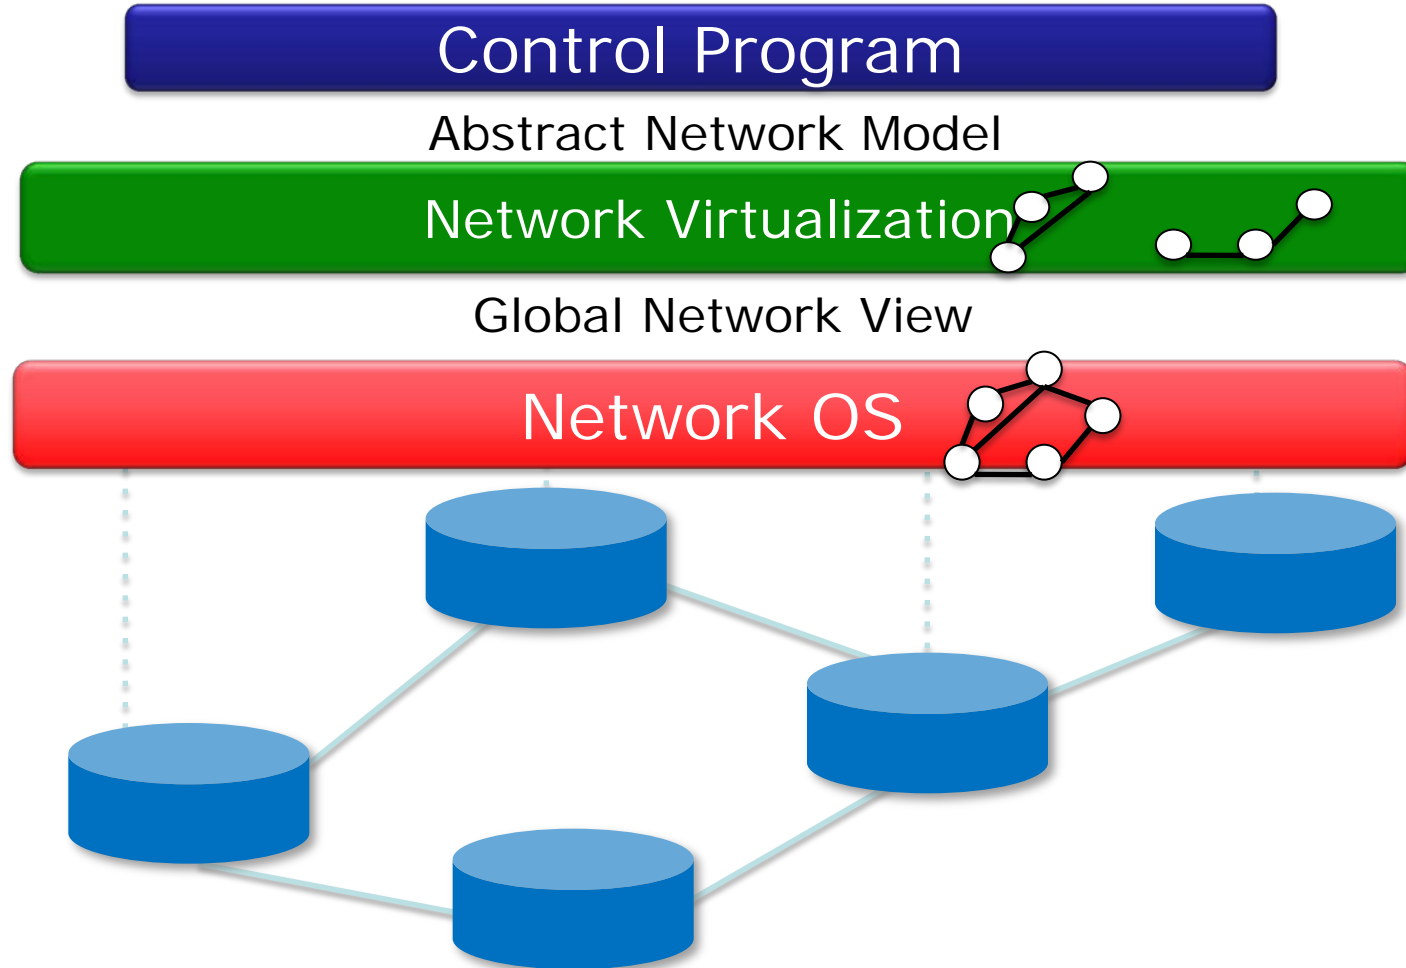
Global  
Network  
View

## How

# Software Defined Network: Take 2



# Software Defined Network: Take 2



# What Does This Picture Mean?

- ❖ Write a simple program to configure a simple model
  - Configuration merely a way to specify what you want
- ❖ Examples
  - ACLs: who can talk to who
  - Isolation: who can hear my broadcasts
  - Routing: only specify routing to the degree you care
    - Some flows over satellite, others over landline
  - TE: specify in terms of quality of service, not routes
- ❖ Virtualization layer “compiles” these requirements
  - Produces suitable configuration of actual network devices
- ❖ NOS then transmits these settings to physical boxes

# Software Defined Network: Take 2

**Specifies  
behavior**

Control Program

Abstract Network Model


**Compiles to  
topology**

Network Virtualization

Global Network View

**Transmits  
to switches**

Network OS





# Two Examples Uses

## ❖ Scale-out router:

- Abstract view is single router
- Physical network is collection of interconnected switches
- Allows routers to “scale out, not up”
- Use standard routing protocols on top

## ❖ Multi-tenant networks:

- Each tenant has control over their “private” network
- Network virtualization layer compiles all of these individual control requests into a single physical configuration

## ❖ Hard to do without SDN, easy (*in principle*) with SDN

### 3. Forwarding Abstraction

---

- ❖ Switches have two “brains”
  - Management CPU (smart but slow)
  - Forwarding ASIC (fast but dumb)
- ❖ Need a forwarding abstraction for both
  - CPU abstraction can be almost anything
- ❖ ASIC abstraction is much more subtle: **OpenFlow**
- ❖ OpenFlow:
  - Control switch by inserting <header; action> entries
  - Essentially gives NOS remote access to forwarding table
  - Instantiated in OpenvSwitch

# Does SDN Work?

- 
- |   |            |
|---|------------|
| ❖ Is it scalable?                           | <b>Yes</b> |
| ❖ Is it less responsive?                    | <b>No</b>  |
| ❖ Does it create a single point of failure? | <b>No</b>  |
| ❖ Is it inherently less secure?             | <b>No</b>  |
| ❖ Is it incrementally deployable?           | <b>Yes</b> |

# SDN: Clean Separation of Concerns

- ❖ **Control prgm: specify behavior on abstract model**
  - Driven by **Operator Requirements**
  
- ❖ **Net Virt'n: map abstract model to global view**
  - Driven by **Specification Abstraction**
  
- ❖ **NOS: map global view to physical switches**
  - API: driven by **Distributed State Abstraction**
  - Switch/fabric interface: driven by **Forwarding Abstraction**

# We Have Achieved Modularity!

- ❖ Modularity enables independent innovation
  - Gives rise to a thriving ecosystem
- ❖ Innovation is the true value proposition of SDN
  - SDN doesn't allow you to do the impossible
  - It just allows you to do the possible much more easily
- ❖ ***This is why SDN is the future of networking...***

- ❖ Open Networking Foundation is standards body
  - SDN endorsed by 49 companies
  - Almost everyone who matters.....
  
- ❖ A few products on market, many more coming
  - Some large companies using SDN internally
  
- ❖ **SDN has won the war of words, the real battle over customer adoption is just beginning....**