

Network Programmability

Research Topics in SDN & NFV

Dirk Kutscher
Fabian Schneider

NEC Laboratories Europe

NEC Laboratories Europe (NLE) - Overview

- ~100 leading researchers from all over Europe and world-wide in Heidelberg, and London (NEC E HQ)
- Close links with leading European research institutes & universities
- Collaboration with major industry in Europe,
eg. network operators, ICT vendors, automotive, utilities….

■ Research areas in NLE

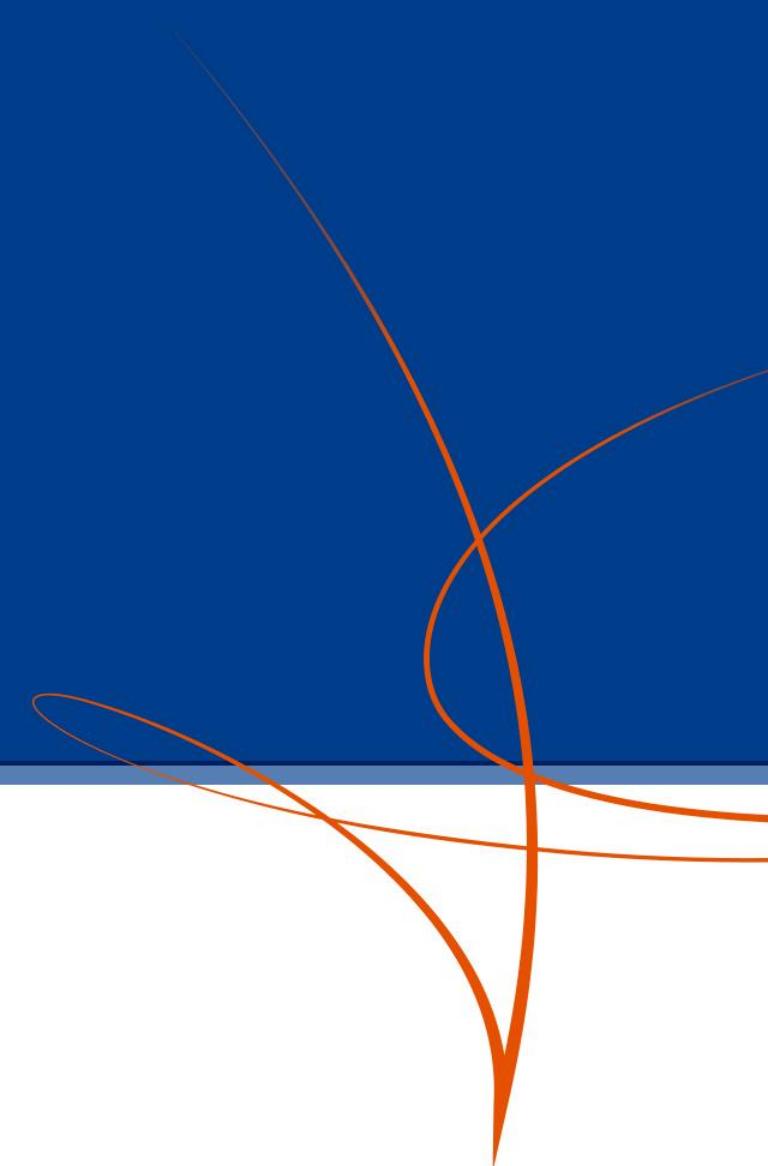
- 4G/5G, Future Internet, SDN
- Information-Centric Networking
- Cloud platform, management & services
- Security, Privacy & Performance
- Internet of Things (M2M)
- ITS and Green Telematics
- Smart Energy
- Standardization

■ NEC Laboratories Europe

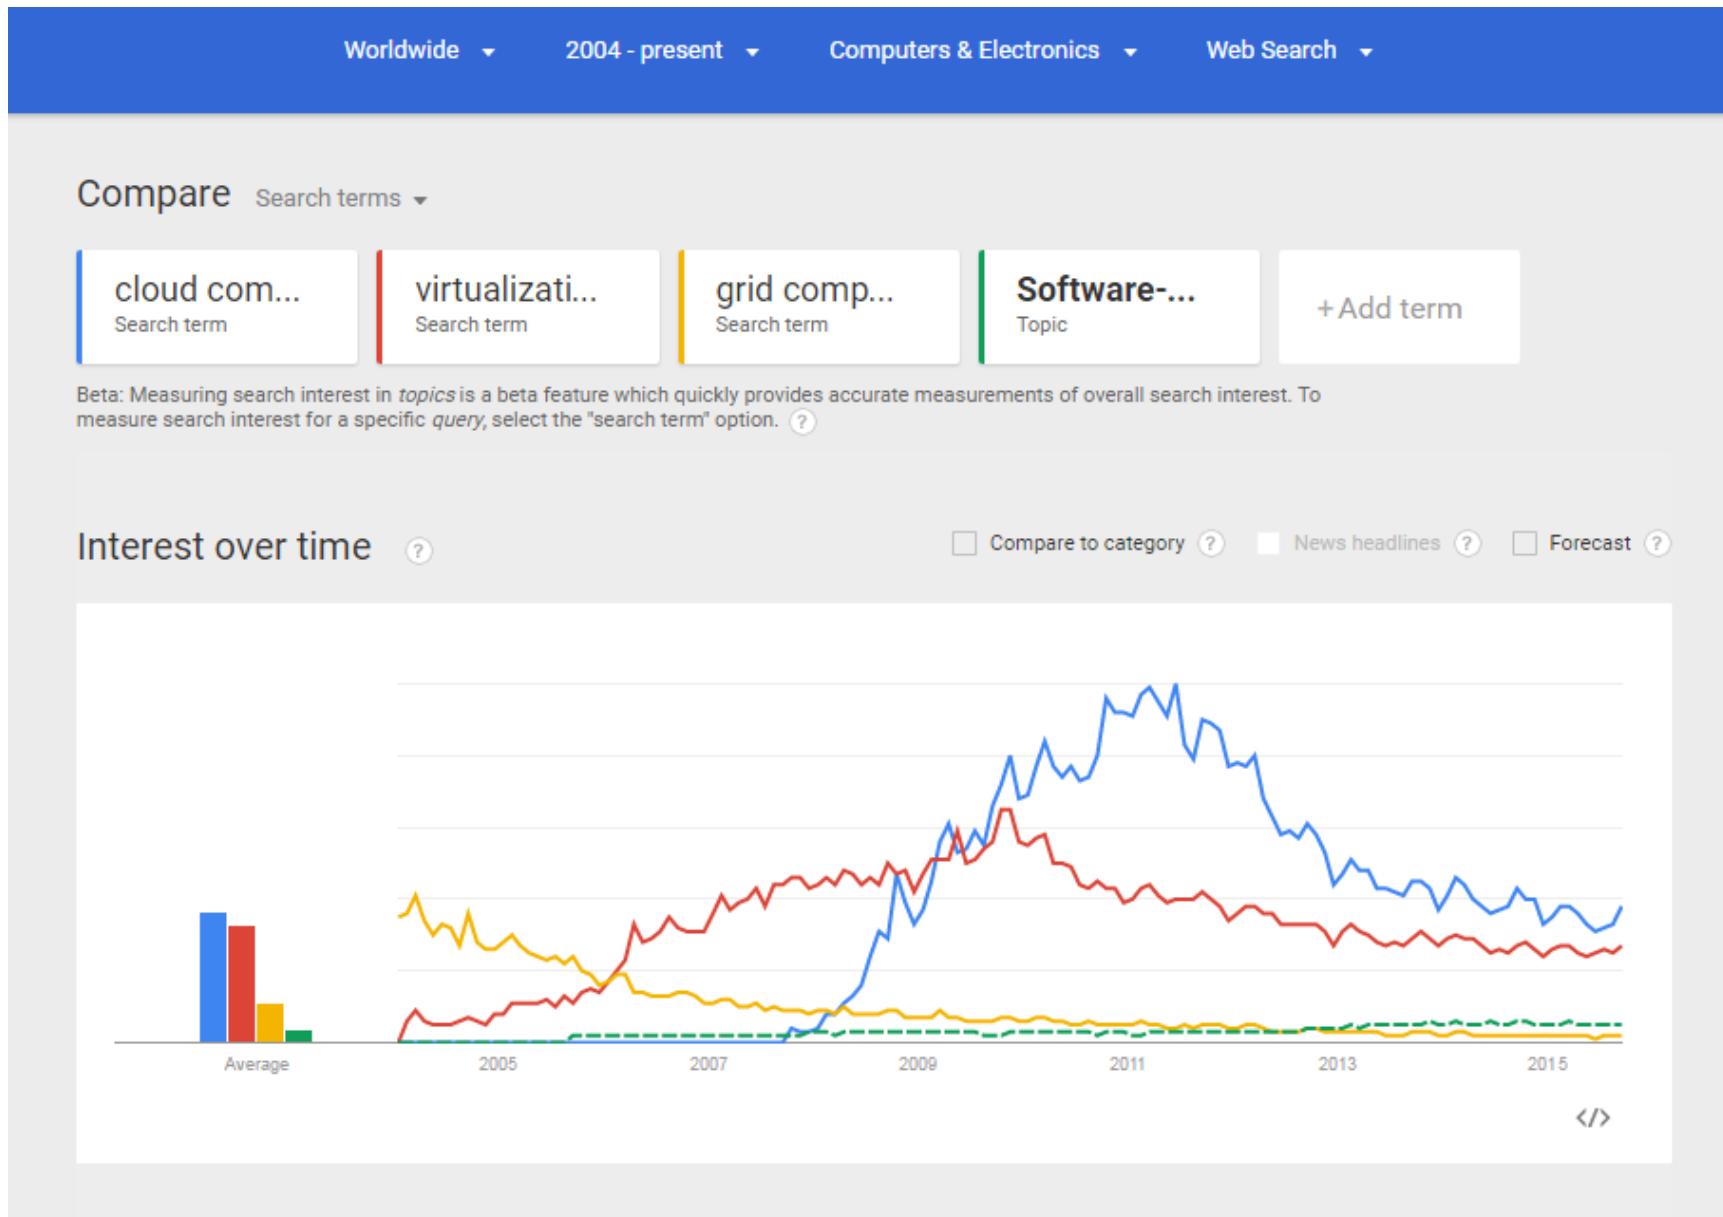
- Kurfürstenanlage 36
- D 69115 Heidelberg
- www.neclab.eu



Programmability Intro



Google Zeitgeist on Cloud Computing Keywords



Technology Trends

■ **Softwarization and programmability**

- Move more and more functions from HW to SW to increase agility…

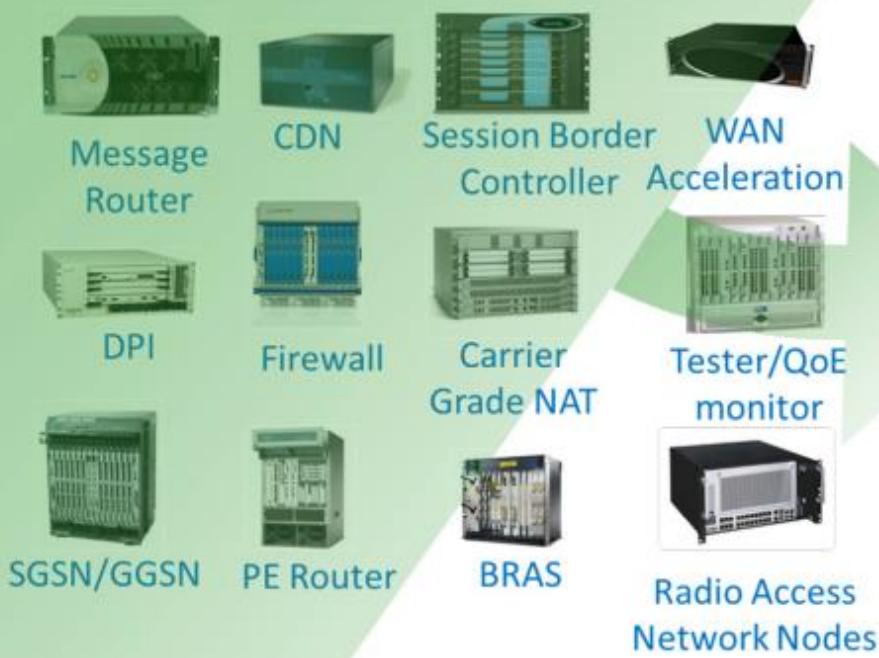
■ **Commoditization**

- initially: servers
- currently: L2/L3 switches
- next: base stations?
→ Whiteboxes (“bare metal switches”) based purely on merchant silicon

■ **Cloudification and automation**

- Cloud is not about virtualization. It is about hyper-scale, elasticity, resource pooling, … but most of all AUTOMATION

Classical Network Appliance Approach

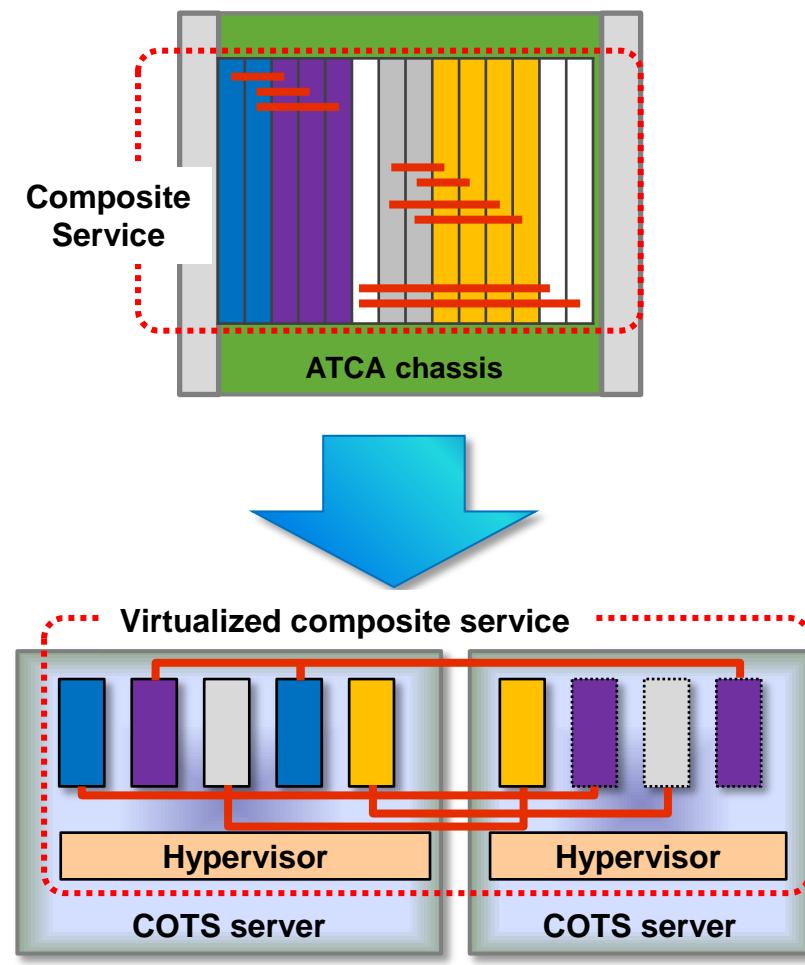


- Fragmented non-commodity hardware.
- Physical install per appliance per site.
- Hardware development large barrier to entry for new vendors, constraining innovation & competition.



Network Virtualisation Approach

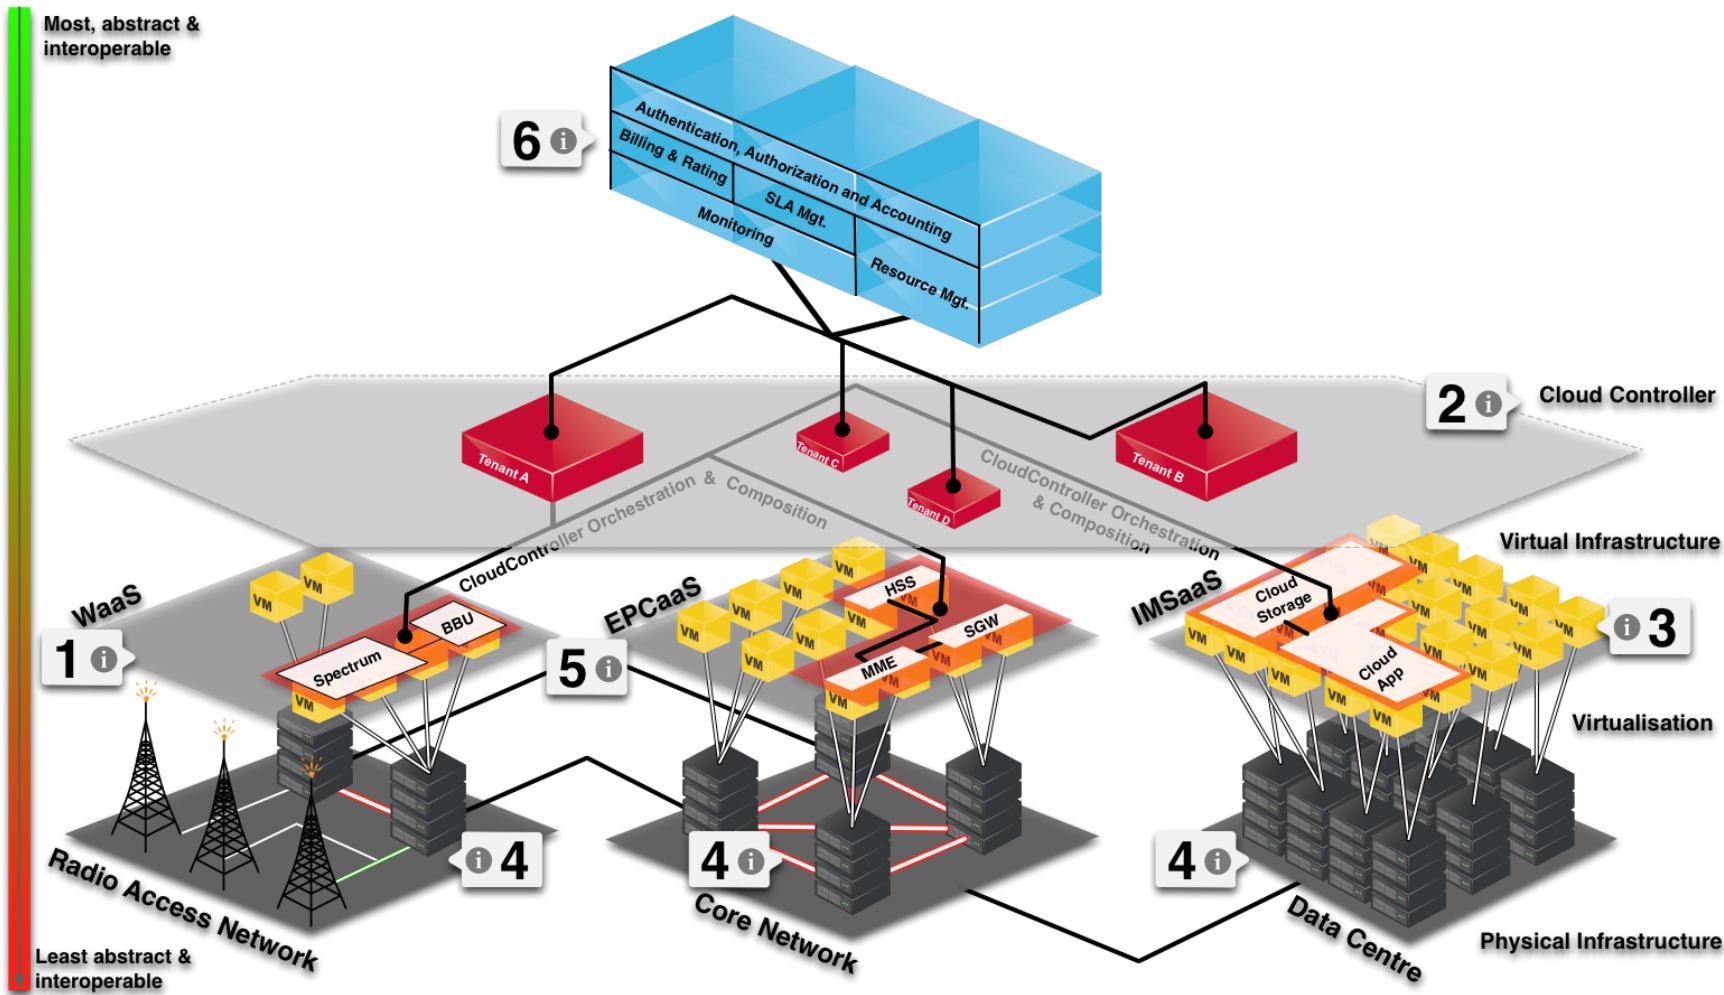
Telco Infrastructure Softwarization



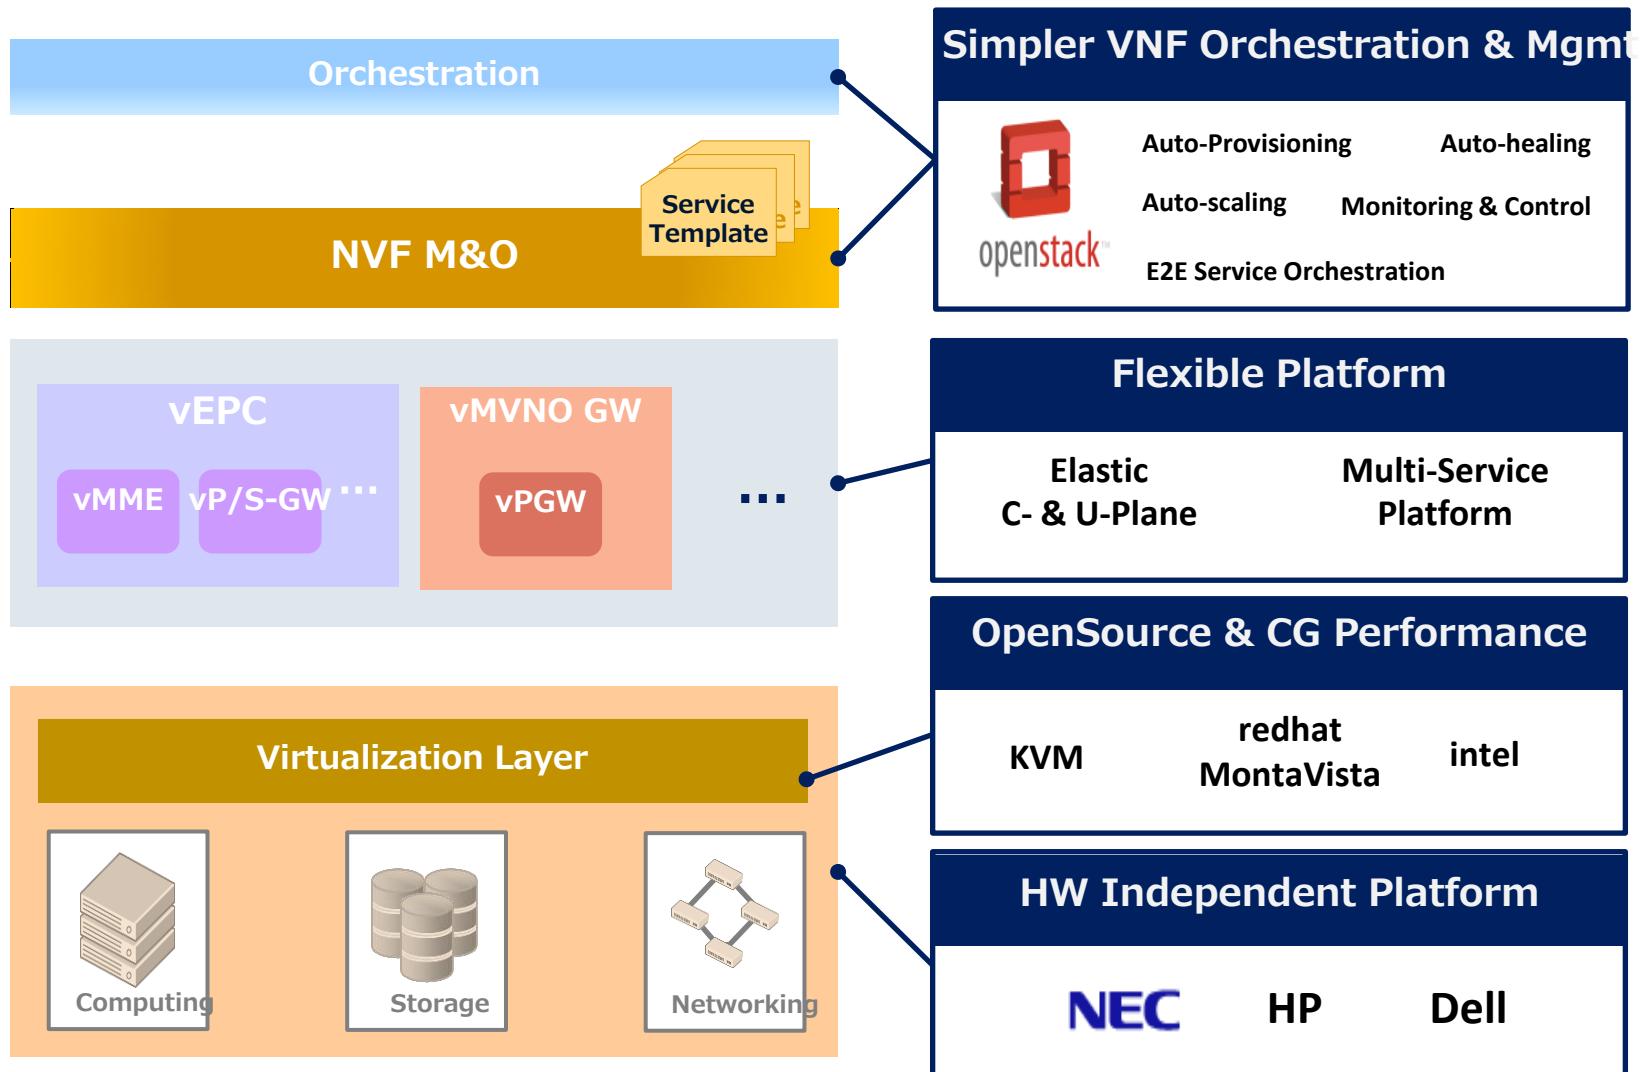
Transformation from ATCA to VNF

Example

MobileCloud Architectural Overview

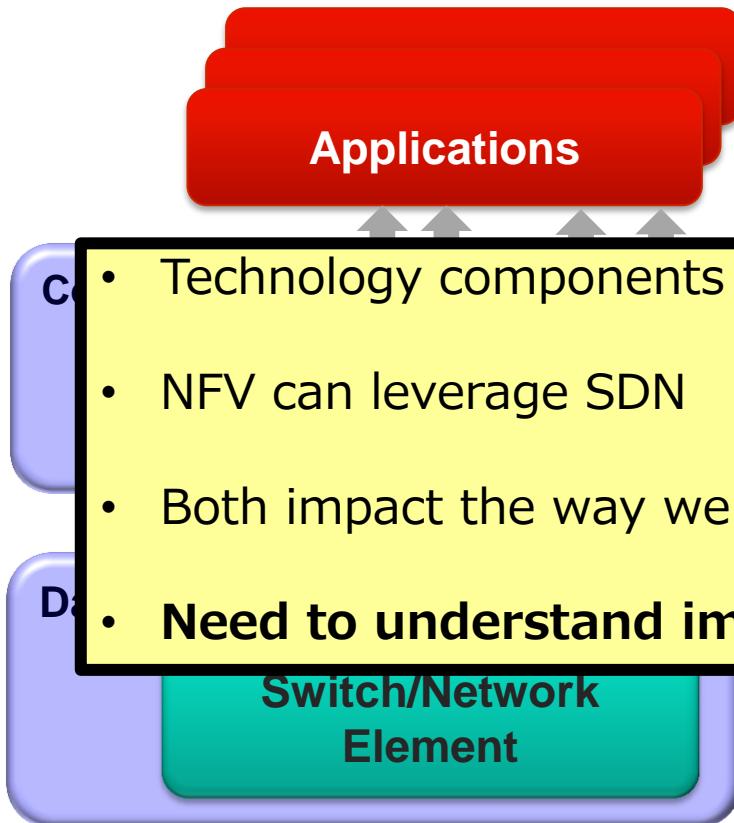


Implementation: NEC SDN/NFV

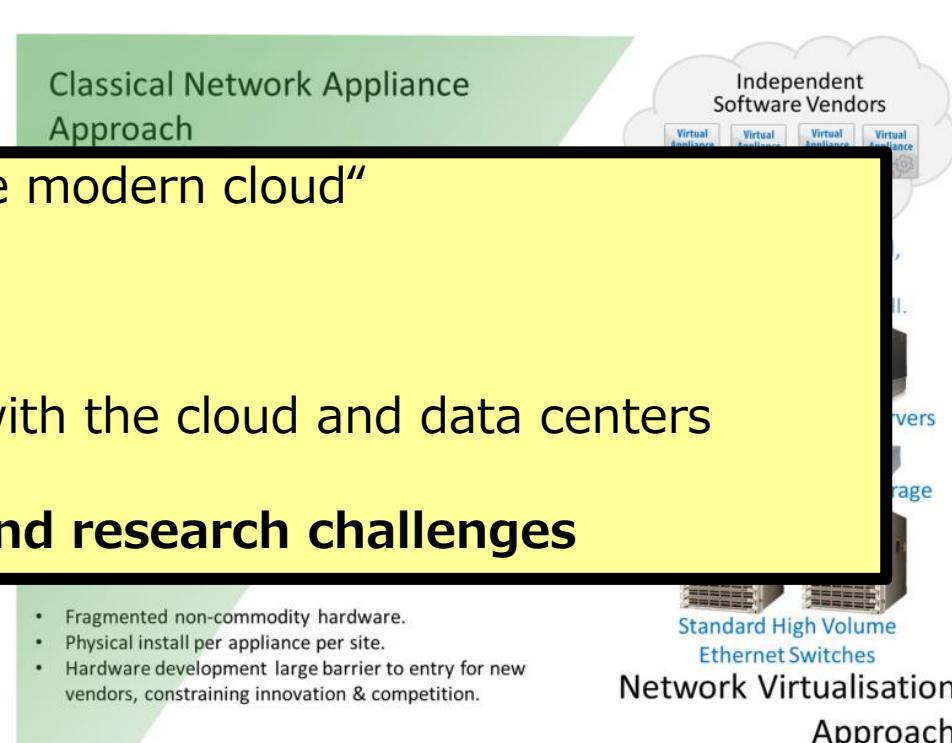


Two Main Technology Components

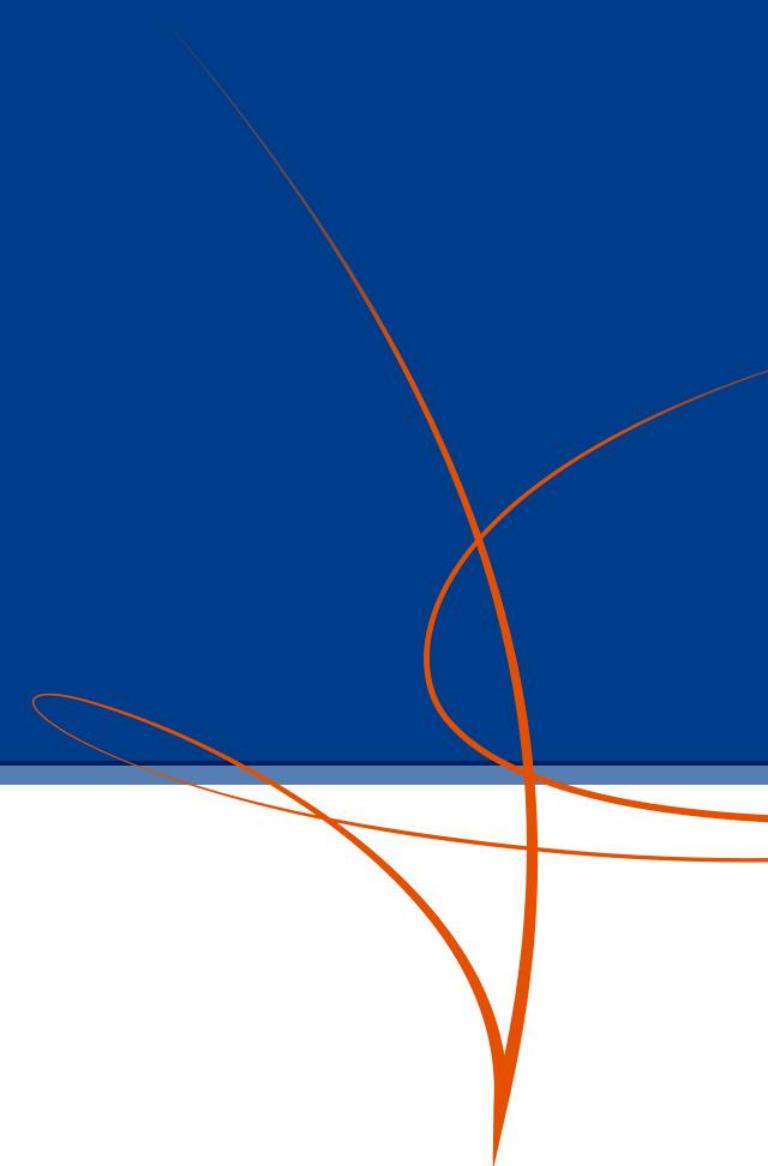
Software-Defined Networking



Network Functions Virtualization



Software-Defined Networking



Nuts and Bolts of the Internet

End-to-end principle

Best effort packet transfers

Best effort:

No guarantees for delivery, order, speed,

...

Programs exchange data ...



... using hosts ...

... connected via the Internet

Internet



Packets not circuits:

No reservations, no setup, fire and forget,

...

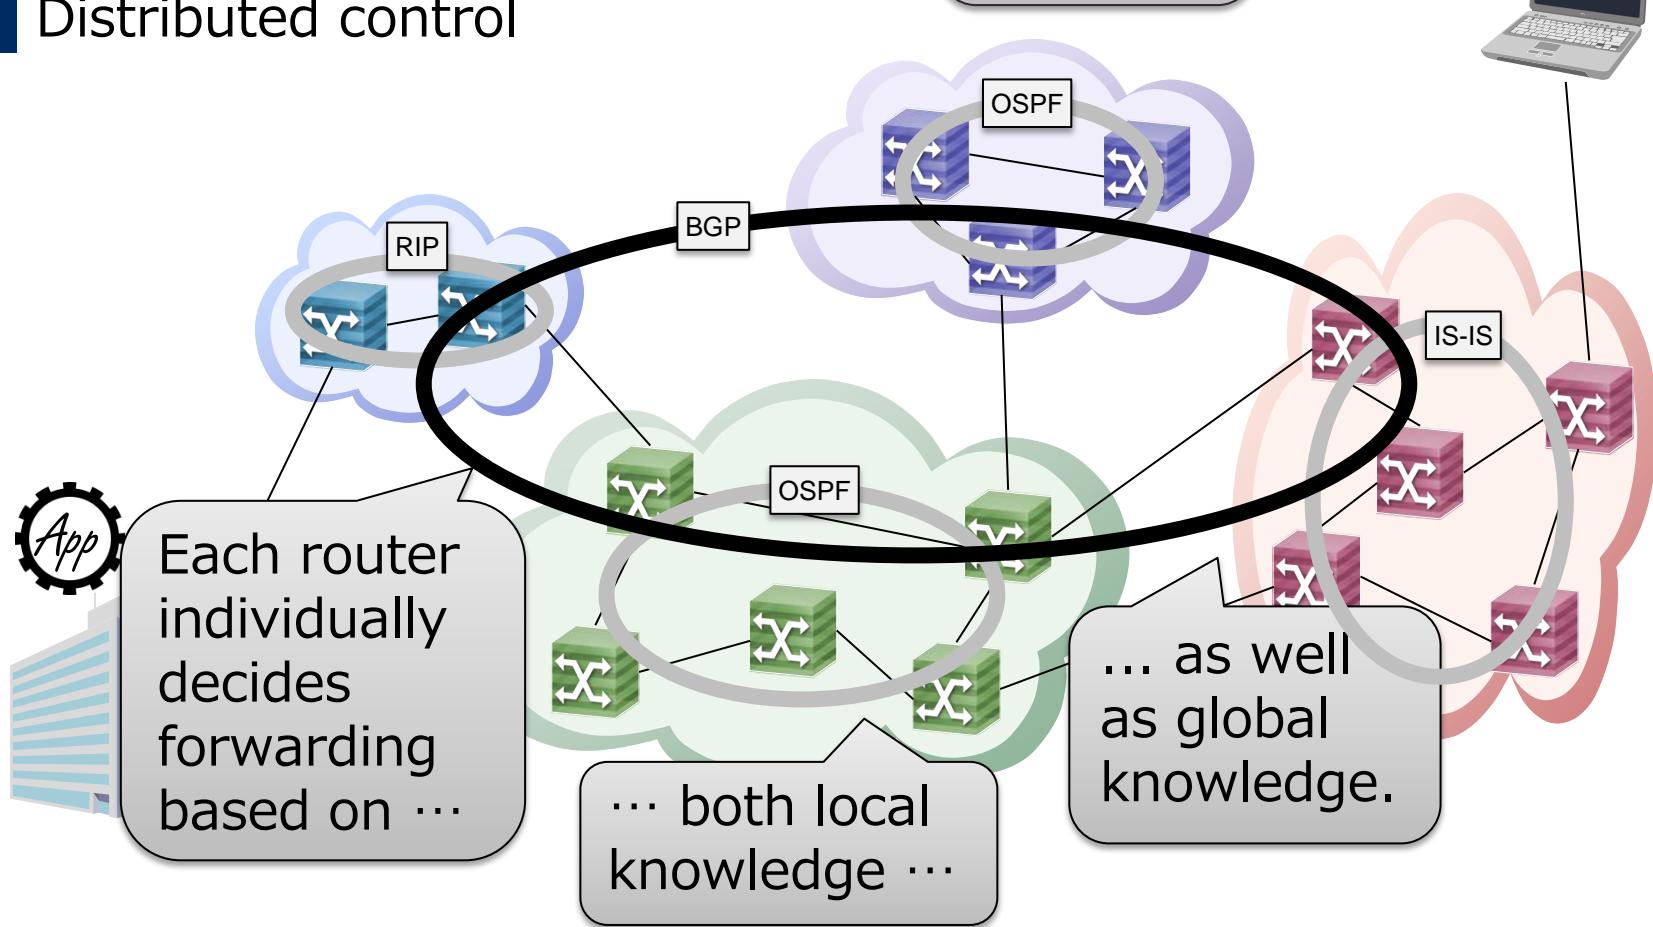
⇒ Layered protocol stack:

- Local & global addressing
- Reliability on top
- Modularity allows innovation
- Abstractions for ease of use

Nuts and Bolts of the Internet

- End-to-end principle
- Best effort transport
- Network of networks
- Distributed control

Different owners & operators



Data plane and hosts

- Built on under-specification
 - **Best-effort packet** delivery
 - OSes provide open API to network stack

Fast innovation & adoption

- Fast innovation & adoption
 - Web, P2P, VoIP, OSN, ...
 - Ethernet, optical, WiFi, 3G-4G-5G

Huge commercial success

- Huge commercial success
 - Internet started as research experiment, now:
 - Yahoo, Google, Facebook, Apple, Microsoft, Amazon, ...

Inside the network

Distributed self-adapting network control plane

- Closed switches/routers
 - Bundled with control SW
 - Proprietary interfaces

- Huge amounts of protocols
 - Slow to standardize

Slow innovation

- Slow innovation
 - Vendors write the code
 - New features need to yield short-term business

Time Warp to 2006 @ Stanford University



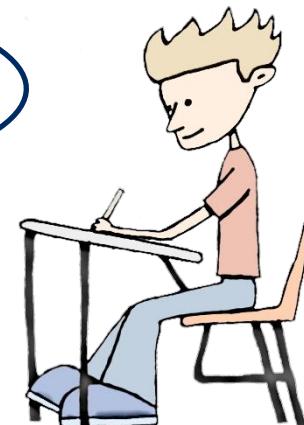
Nick McKeown
CS Professor

“We need a security architecture
for enterprise networks”

“Ok, ... so ... What is the
programming interface to
the network?”

“Well, you can change
parameters in network
control protocols!”

“Let’s build a programming
interface to forwarding.”



Martin Casado
CS Student &
SW Engineer

SANE
Usenix Sec '06

Ethane
SIGCOMM '07

OpenFlow
CCR '08

NOX
CCR '08

OVS
HotNets '09

Key to Internet Success

Only data plane protocols
Let's add the control plane

Applications

DNS SIP SDP H.323

Reliable

Layers are Great Abstractions

RSVP

Best-effort

... but ...

Best-effort

- Layers only deal with the **data plane**

- We have no powerful **control plane**

...built on...

Physical transfer of bits

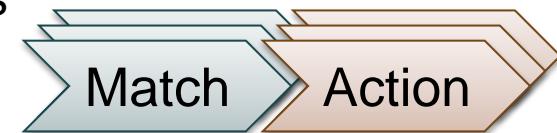
"The infrastructure still works...
Only because of your great ability
to master complexity"
Need to extract simplicity, by
finding abstractions for the control
plane!

Scott Shenker

Borrowed from Scott Shenkers talk "The Future of Networking, and the Past of Protocols" @ ONS'11
<http://opennetsummit.org/archives/oct11/shenker-tue.pdf>

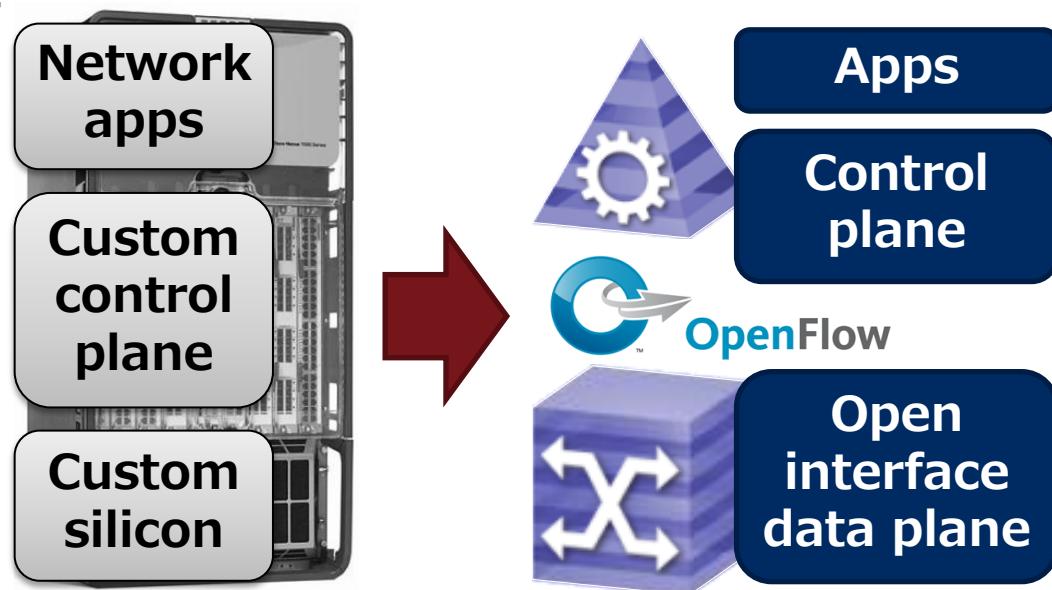
OpenFlow in a Nutshell

- Provide open protocol/interface to separate control from data plane
- Rules (Match/Action tuples + priority + counters)
 - **Match:** Select traffic based on packet header fields (e.g. L2 MACs, IP, TCP/UDP, VLAN, MPLS, …)
 - **Actions:** Define what to do with selected traffic (e.g. drop, forward to port, rewrite header, to controller)



- Support for quick & dynamic network configuration

- Latest specs:
 - OF 1.0.2, OF 1.3.5, OF 1.5.1
 - OF-Config 1.2.0
 - OF Test 1.0.1



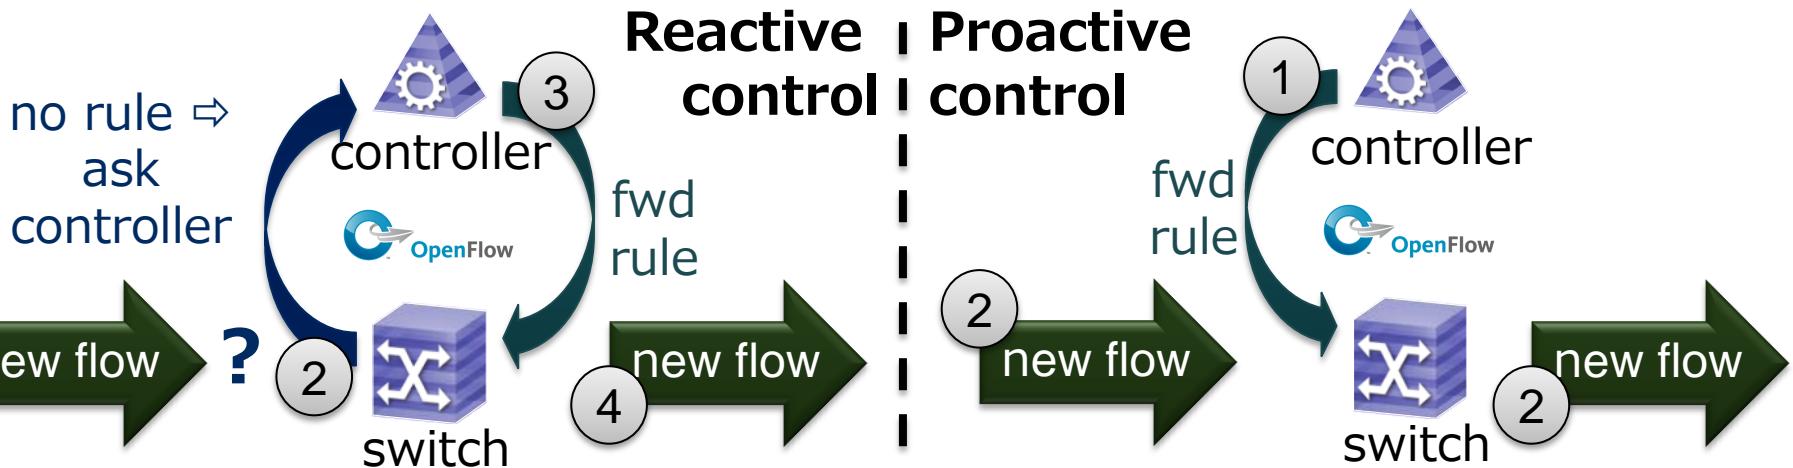
Clarifications on OpenFlow

#1: Flow granularity is not predefined!

- **No restriction** to often mentioned
 - 3-tuple (IP addresses and transport protocol) or
 - 5-tuple (adding in port numbers)
- Allows a **choice of granularity**
 - Matching on any selection of supported match fields possible

#2: **No limitation** to “reactive” control;

“Proactive” equally possible



SDN Principle #1: Decouple Control and Data Plane

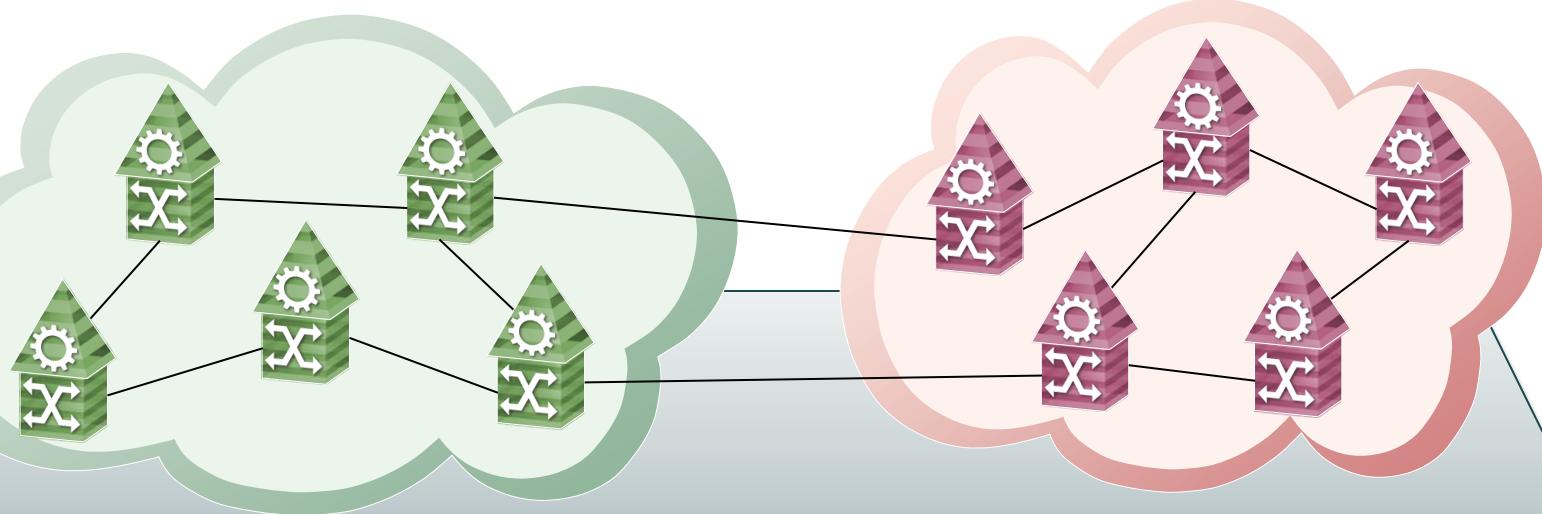


Control logic, routing, algorithms

Packet forwarding & processing

Control
Plane

Data
Plane

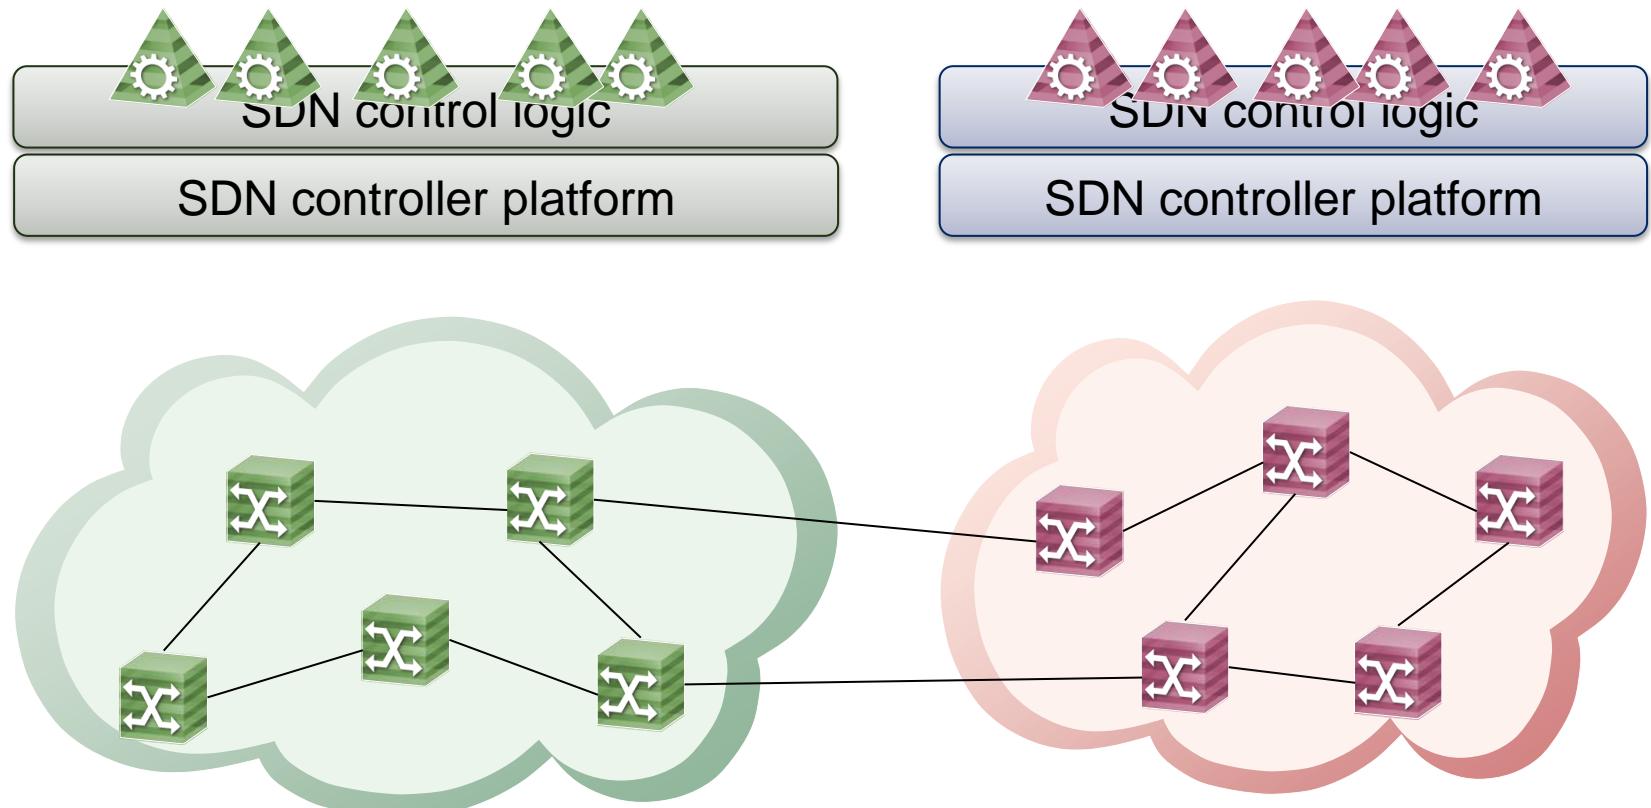


SDN Principle #2: Logically Centralized Controller

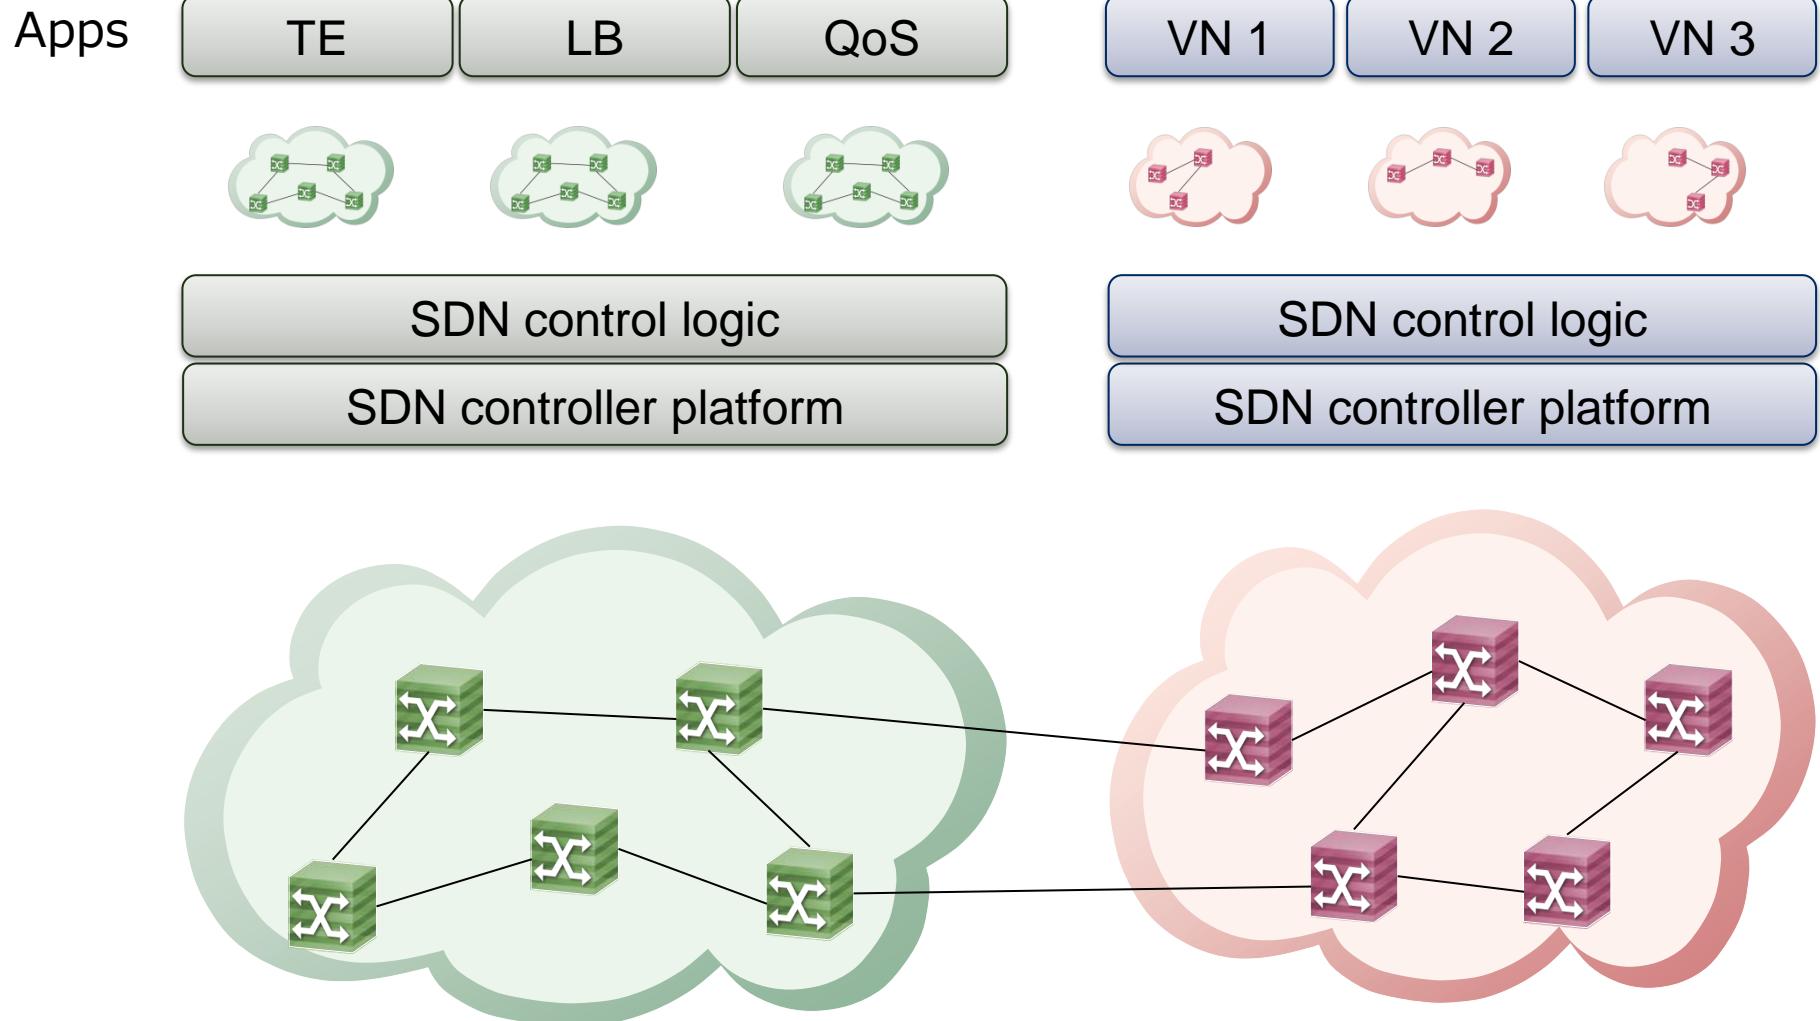


Control logic, routing, algorithms

Packet forwarding & processing



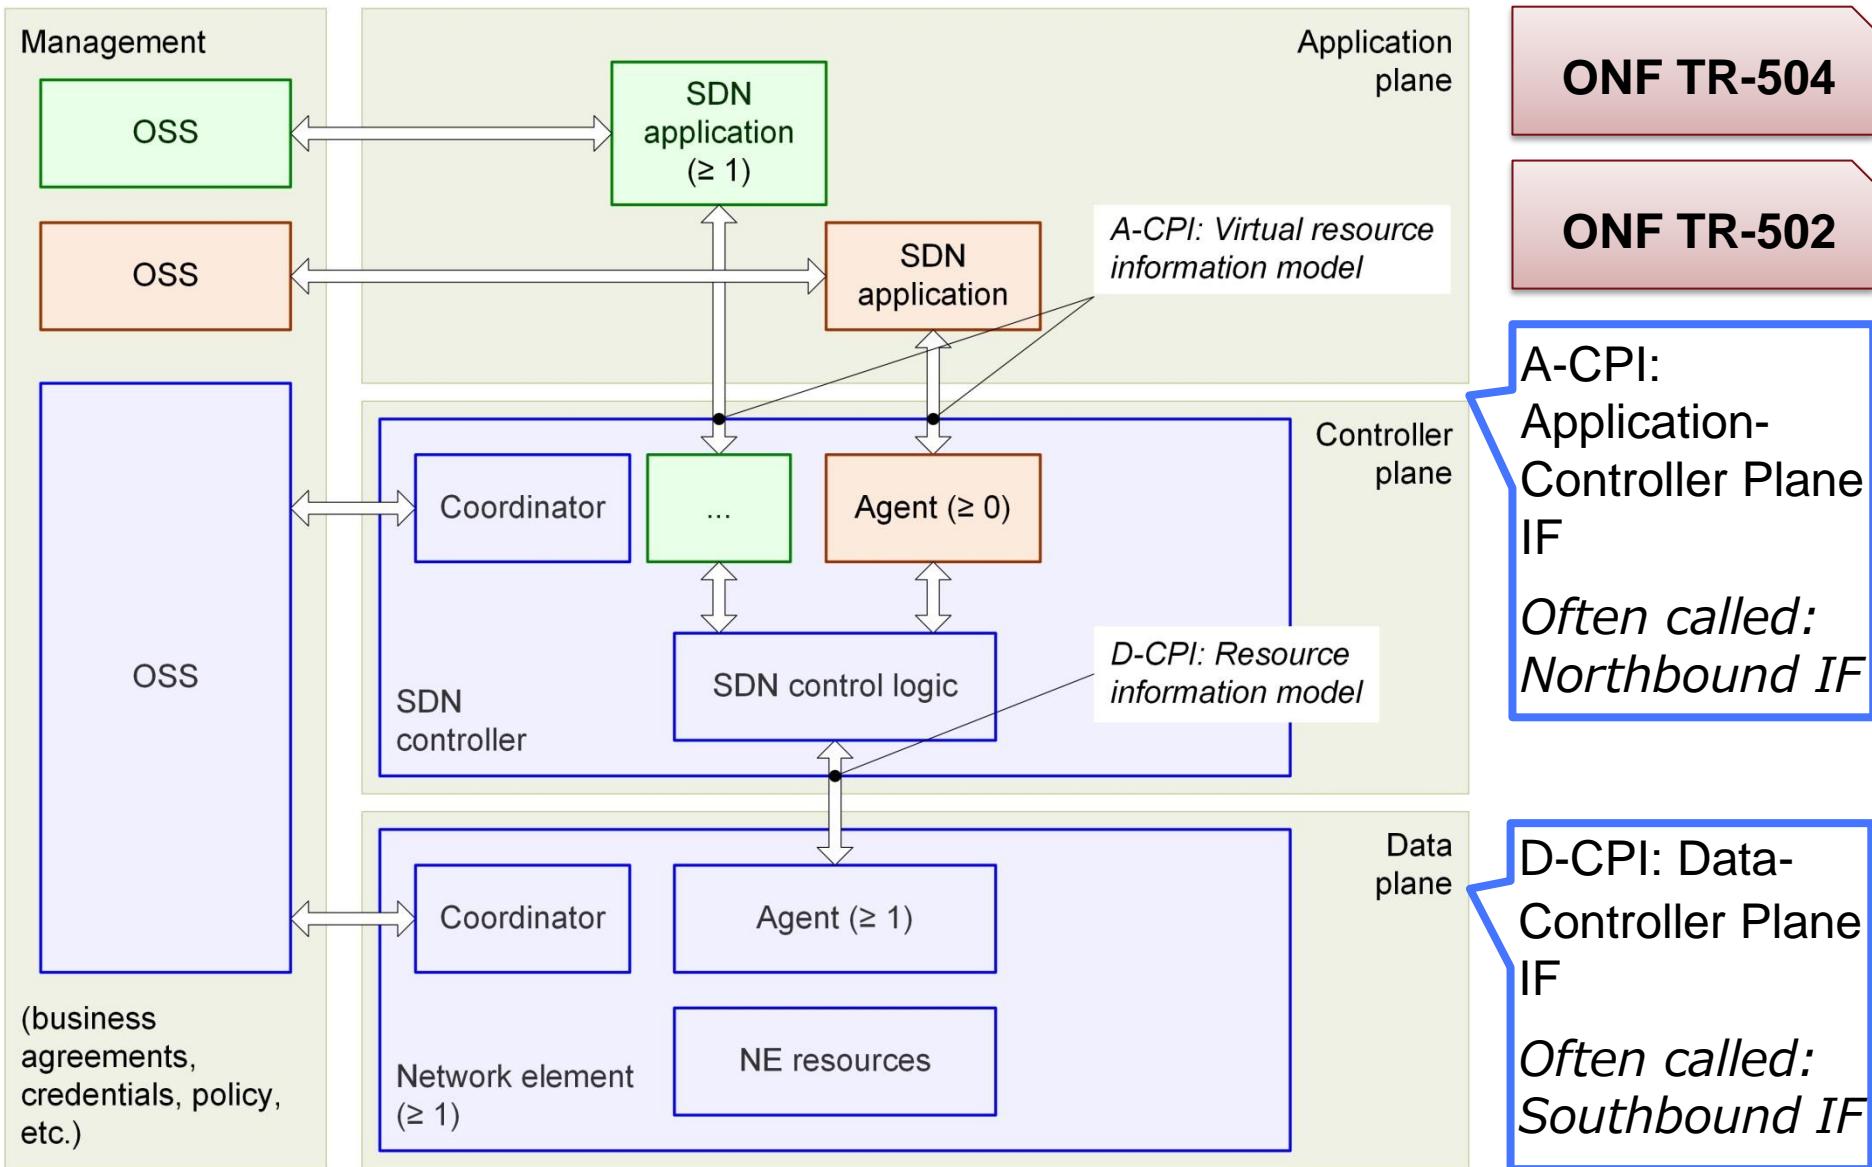
SDN Principle #3: Expose Network to Applications



ONF TR-504

1. Decoupling of controller and data plane
2. Logically centralized control
3. Expose abstract network resources to external applications

SDN Architecture



Important Properties of SDN Architecture

Apps are network aware: **SDN-enabled Apps**

- Communicate their requirements/polices to the network
- Can monitor network state and adapt accordingly

Logically centralized network control: **SDN Network Controller**

- Controller translates from app requirements to low-level rules
- Controller summarizes the network state for apps

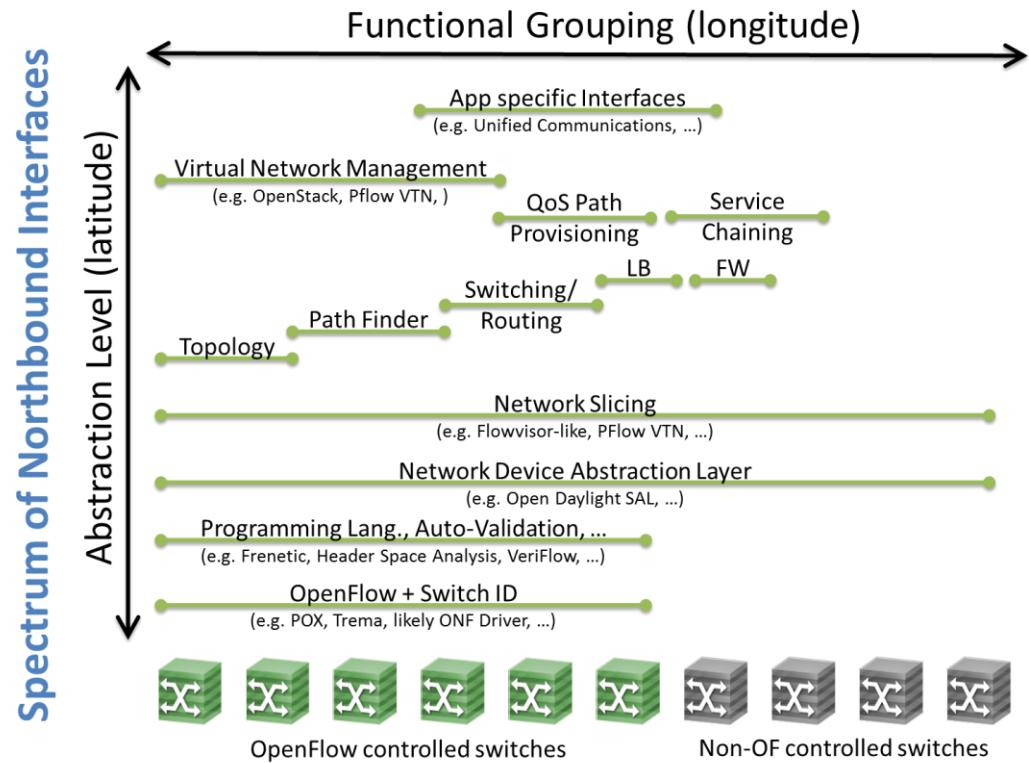
Well-understood driver-like model for devices: **SDN Datapath**

- Programmatic low-level control of all forwarding and configuration
- API for capabilities advertisement and publishing stats
- **No resource contention** with other entities
 - Controller “owns” this device, subject to capabilities advertisement

Defining Northbound Interfaces (NBIs)

Not an easy task

- Level of **abstraction unclear**
 - Bottom: OF+SwitchIDs (e.g. Trema, NOX/POX)
 - Middle: network prog. lang (e.g. Frenetic)
 - Top: Neutron/Quantum level
- **Scope unclear**
 - Single NBI to rule them all
 - Or one per operation call



Current trend: Intent

- „Finance staff needs priority access to SAP servers“
- Finance = Set of ports; SAP server = Set of IPs; Priority = reserved BW
- Include conflict detection and **conflict resolution** in design

OpenFlow ++ ?

Today:

Limited to existing
data plane protocols

Today:

Limited to typical
packet forwarding

Is
Match-
Action
enough
?

Match

Action

Next: PIF

protocol independent
forwarding

User-defined
protocol
headers

**Next:
OpenState**

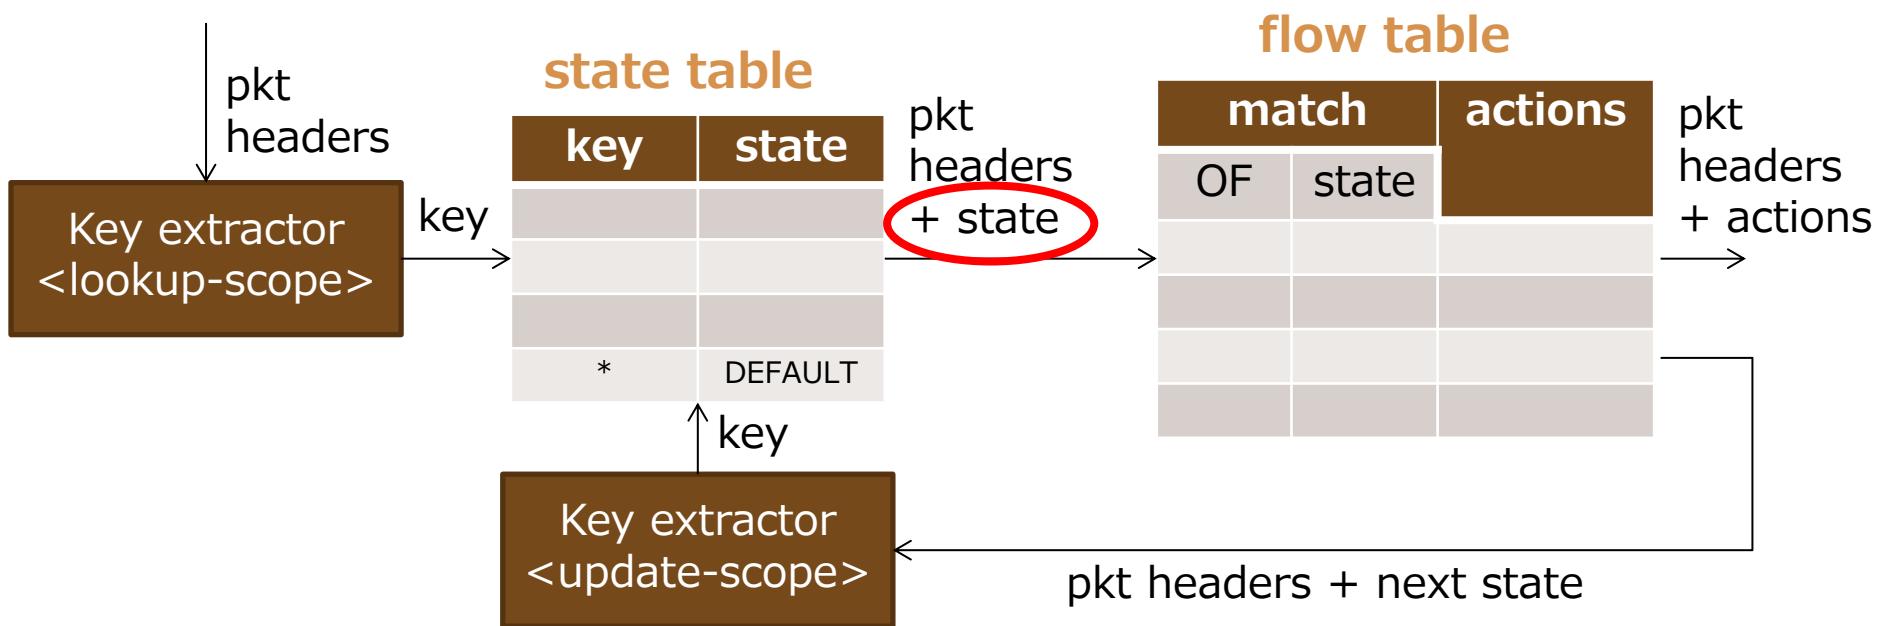
Stateful
matching

**Next: In-switch
processing**

Autonomous
Functions,
Reply packets

OpenState: Stateful flow processing

- New “state table” type: exact match, DEFAULT if miss
- Move state in packet metadata to flow table
- Key extractor: full headers \Rightarrow selected bit sequence
- Actions can include state updates
- Update scope can be different from lookup scope



PacketTemplates: In-switch packet generation

New “packet template table” type:

- Template has byte array of packet contents and metadata

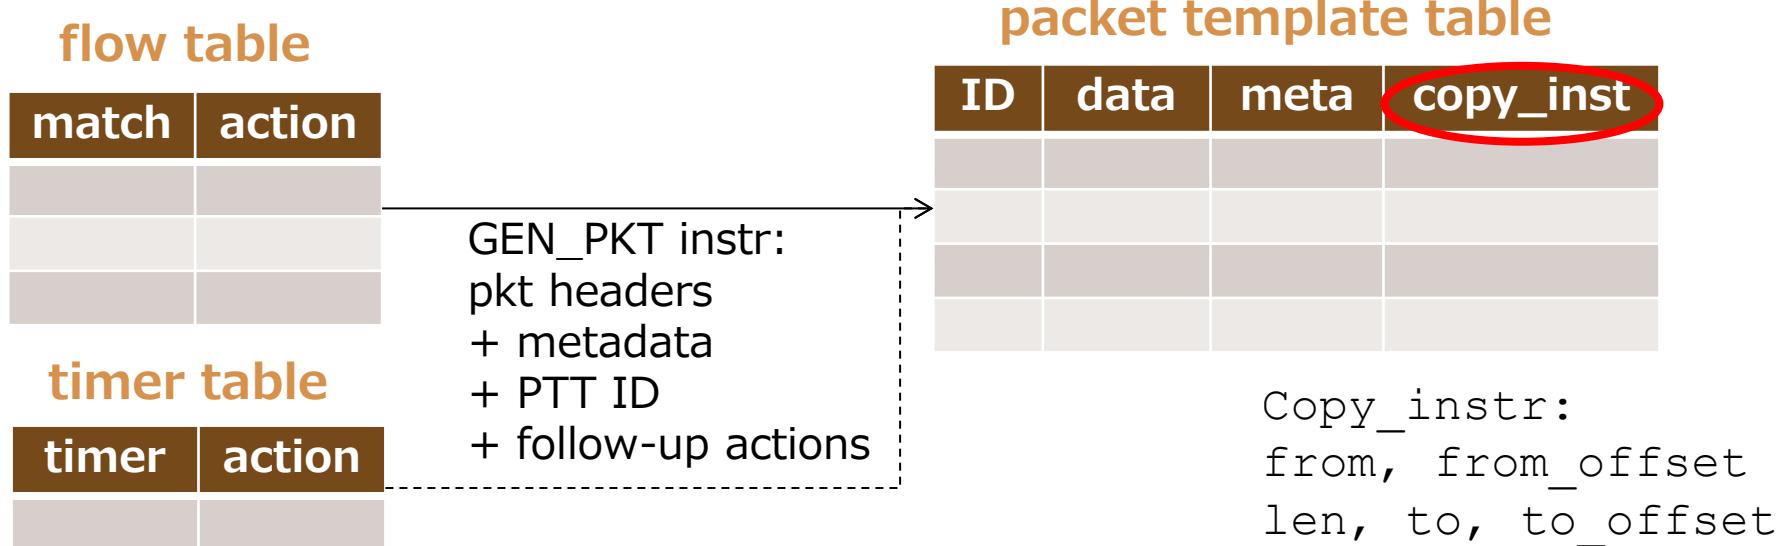
Copy instructions: Program how to change template

- Data source can be a received packet or a table (e.g. state table)

New instruction to refer to packet from flow table

- Additional action specify what to do with the generated packet

Option to use timers as triggers



Initial goals

- Make network forwarding „**programmable**“
- Find good **abstractions** for network control

SDN principles

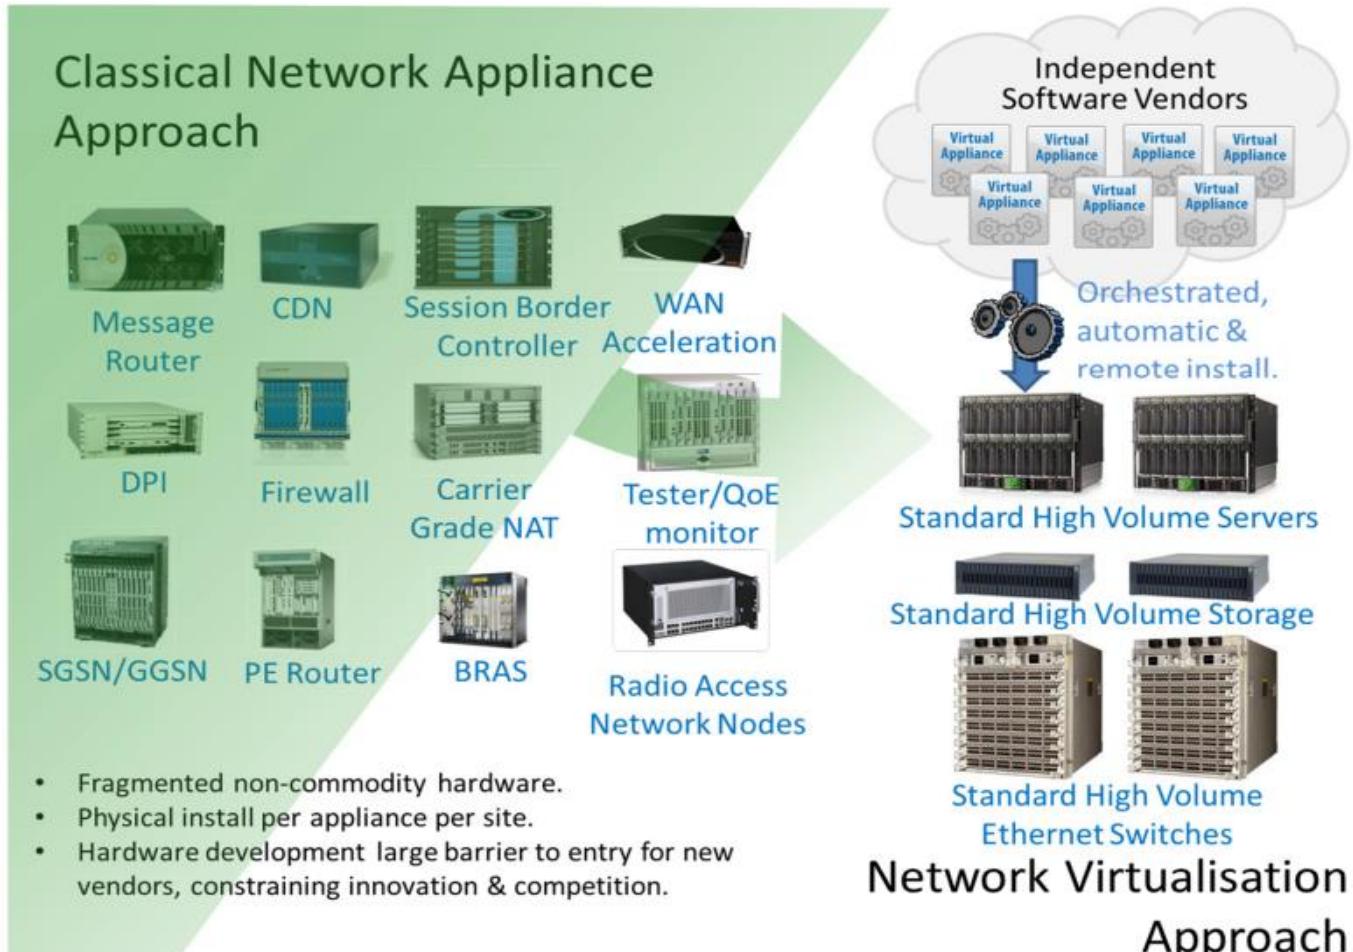
- **Decouple** data and control plane
- Logically centralize control
- Expose network **to applications**

Current Hot Topics

- Develop **higher-level abstractions** than OpenFlow
- **Extend programmability** beyond legacy forwarding
- Better approaches to **integrate network and VM control**

Network Functions Virtualization

NFV = Network Functions Virtualization



NFV = Network Functions Virtualization

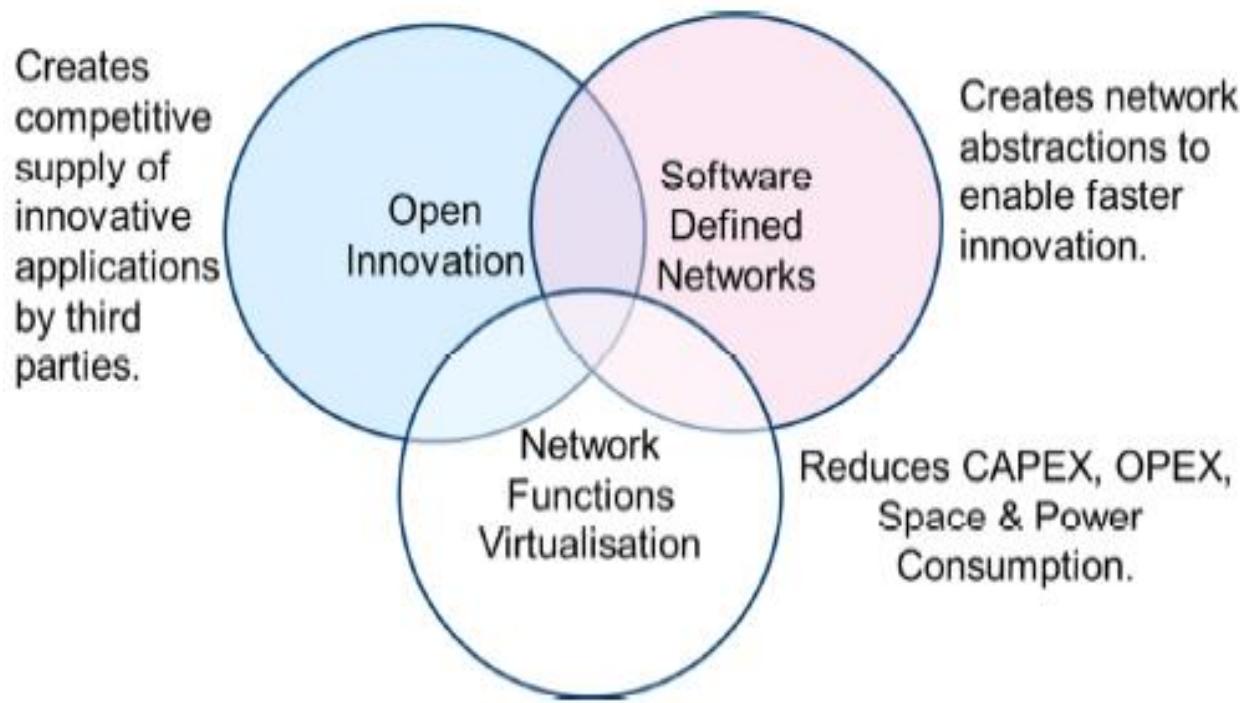
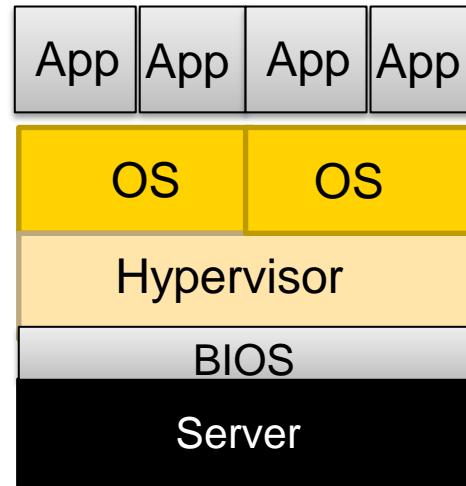
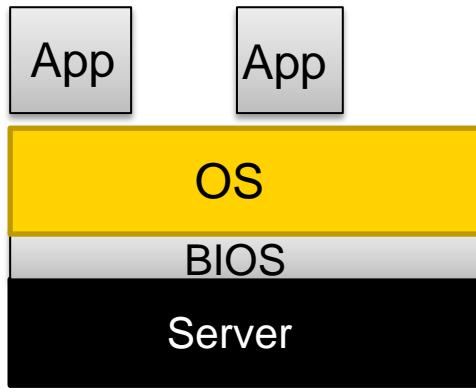


Figure 2: Network Functions Virtualisation Relationship with SDN

Source: ETSI ISG NFV White Paper; http://portal.etsi.org/NFV/NFV_White_Paper.pdf

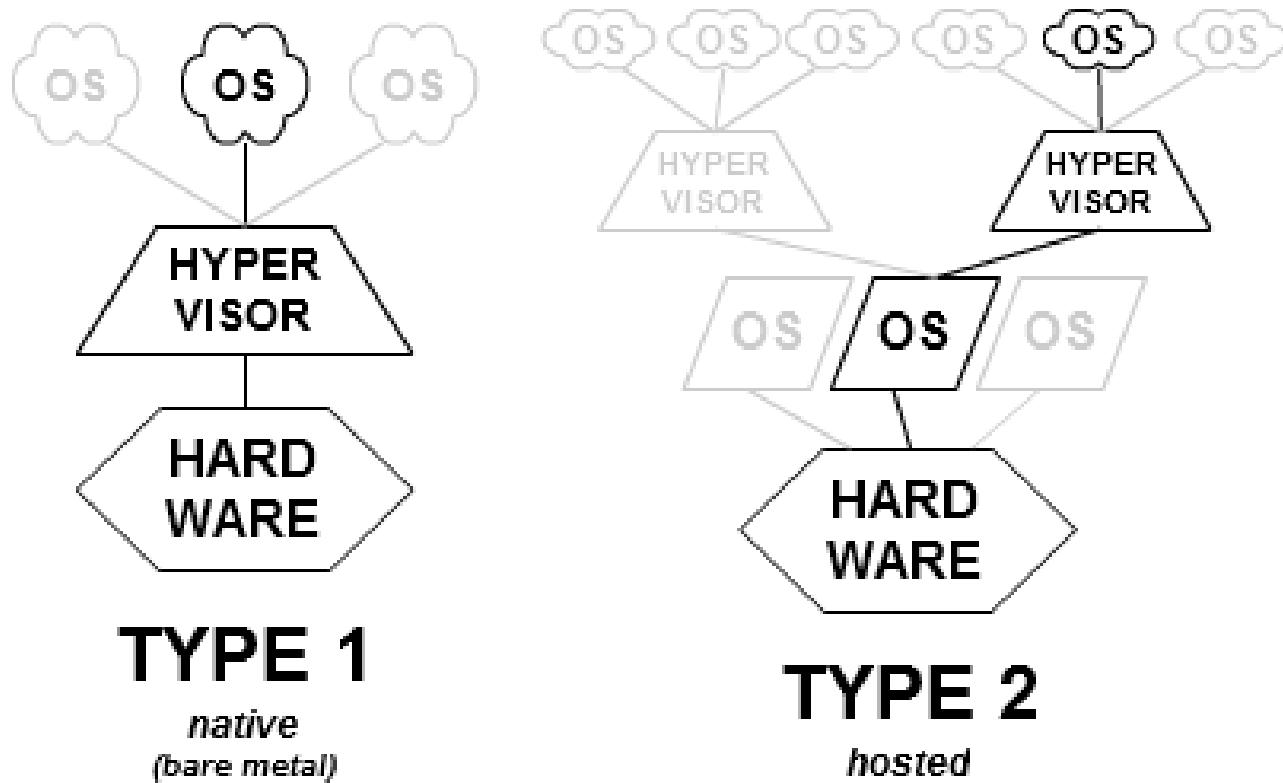
Software on top of COTS Hardware: two approaches

COTS=commercial off-the shelf

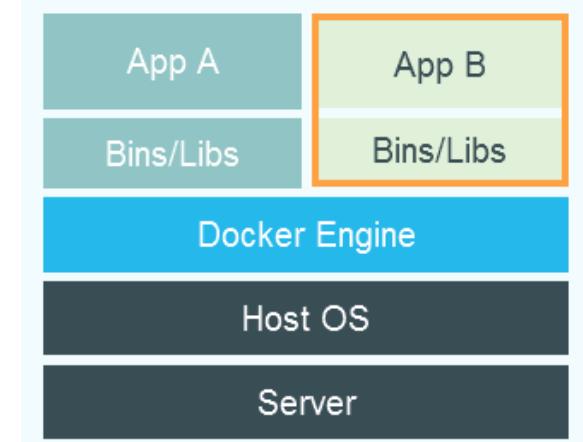
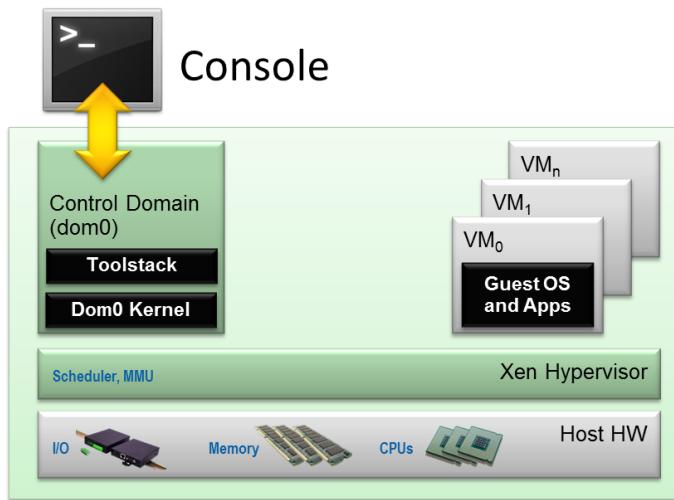
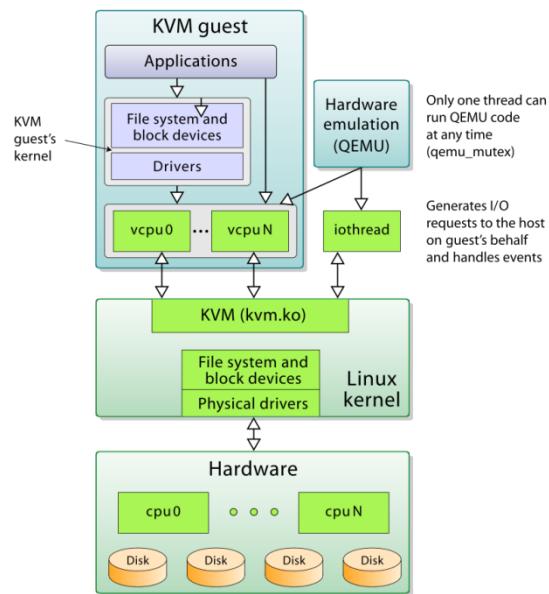


- Traditional Approach
 - Abstraction:
 - Multiple OS / Apps on one server
 - Mobility:
 - VM migration
 - Cloning:
 - Scale in/out
 - Resource management
 - Real Isolation

Virtualization Approaches



Virtualization Approaches



Type 2: Kernel VMs

- Linux Kernel as hypervisor
- „Fat“ hypervisor
- VM as user space process
- (still priv. Mode)
- Shared resources (drivers)

Type 1: HV on bare metal

- Shared domain concept
- Pure hypervisor w/o linux kernel/drivers
- Drivers in one (or more) VMs (driver domain)
- Split driver model

Process virtualization

- Linux Containers
- Leaner, but less isolation capabilities
- Own namespace/libs per app but
- all apps use same Kernel

Pictures taken from projects' websites

Networking in the virtualization era: Some Facts

- **Every server** has a **SOFTWARE VSWITCH** inside
 - to switch packets across multiple VMs
- **Network functions** (inside VMs) are **I/O intensive**
 - **IT functions** are **computing intensive**
- **Consolidation of VMs** on a single server brings **CAPEX reduction**
 - trend towards **many VMs on a PC**



Performance Requirements

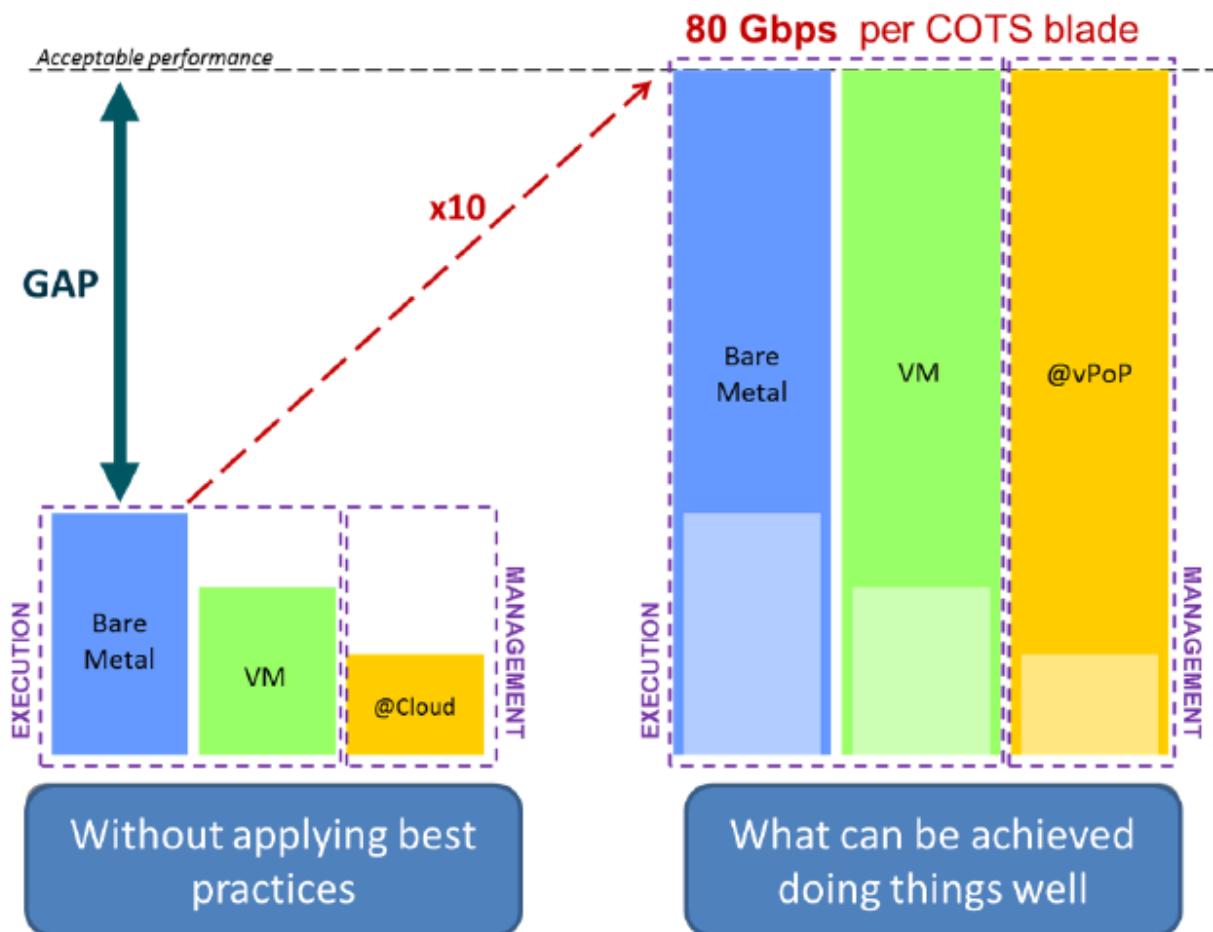
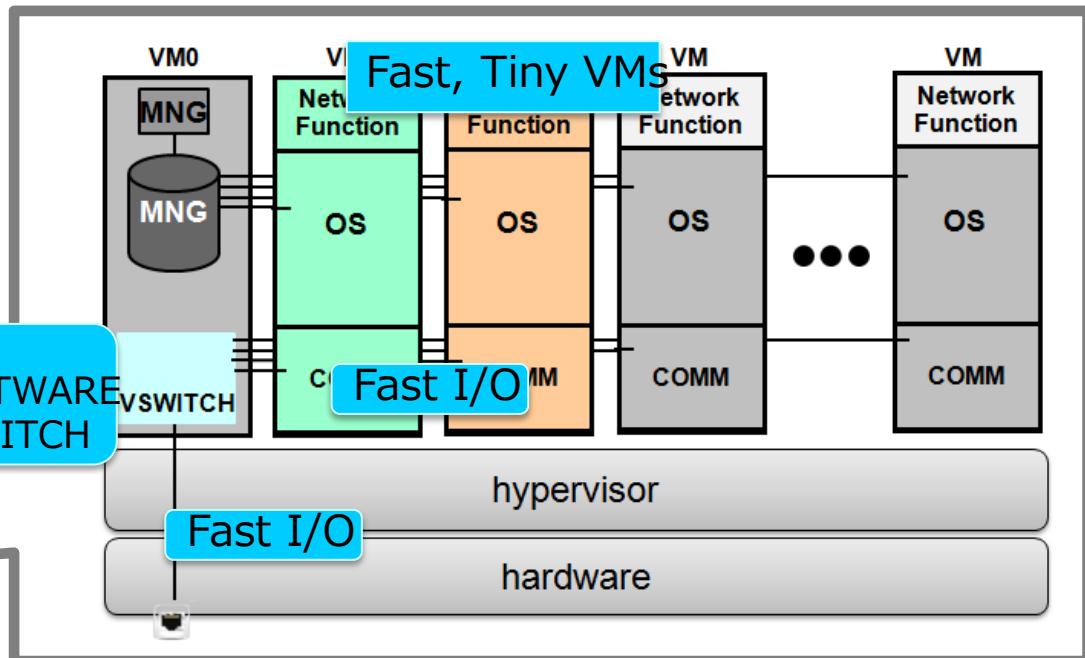


Figure 3: Performance Illustration

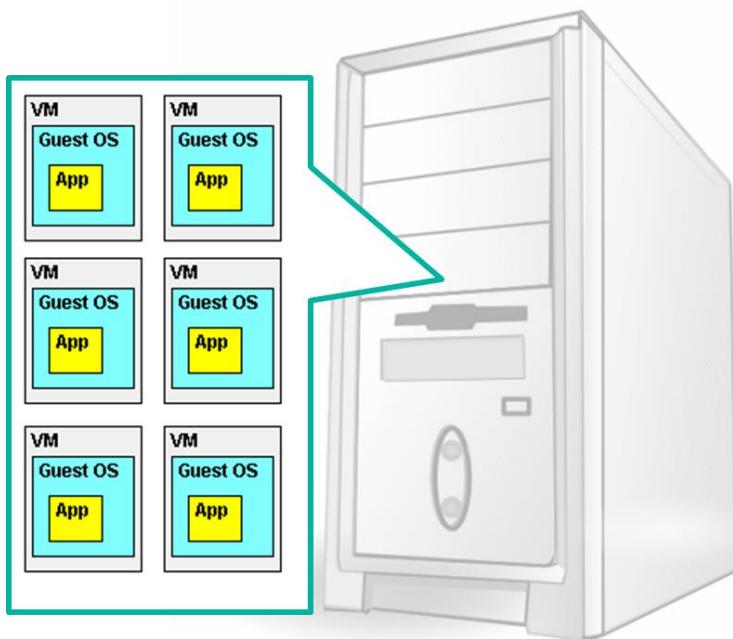


1. fast I/O and switching among fat VMs
mSwitch: a modular “10Gb-and-beyond” software switch
2. longer term perspective for tiny, fast VMs
ClickOS: fast and dynamic virtualized OS for network functions

Network Functions Virtualization: Smaller VMs

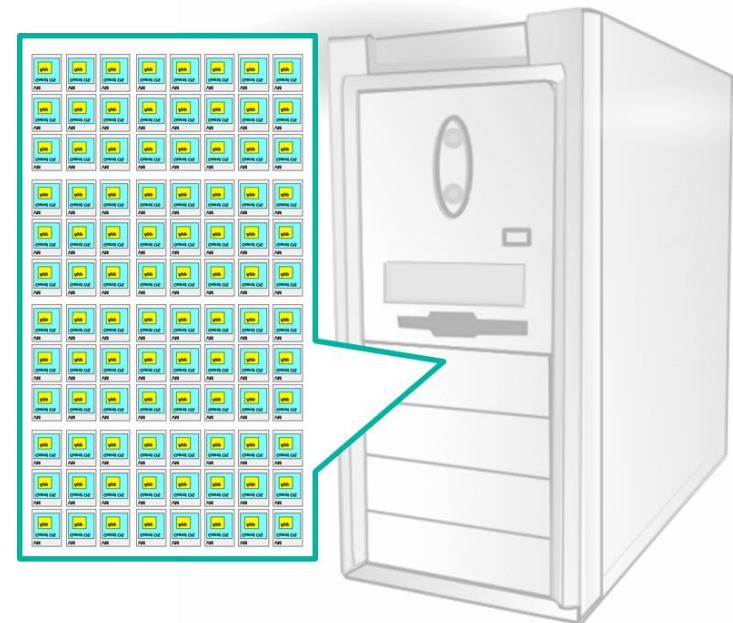
Classical virtualization of nodes:

- Fat, big and slow



New virtualization technology

- Tiny & tailored
- Fast (10Gbps on a PC & small boot times)
- Natively support NW functions



For details, see EWSDN 2012

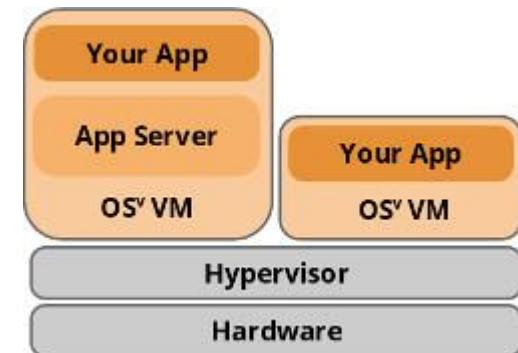
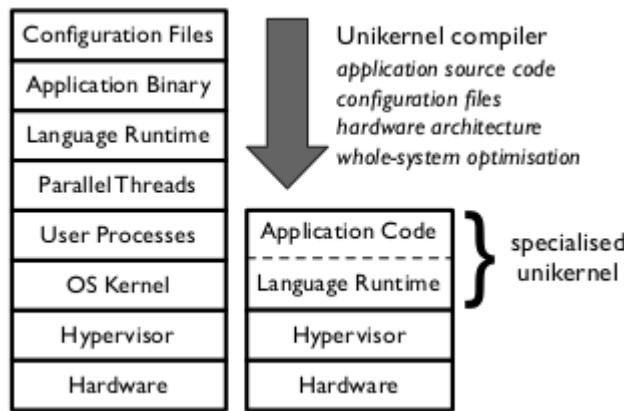
S. Niccolini: Free your middlebox functions down to the data plane with tiny, fast network VMs

<http://www.ewsdn.eu/presentations/NetworkFunctionsVirtualization-EWSDN-v0.3-public.pdf>

Uni-Kernels, Cloud Operating Systems

Specialized OS kernels

- Linked with application code
- More efficient than traditional layering
- For example: many system calls can be transformed to library calls
- Requires pre-compilation of code, configuration into specialized uni-kernel



Mirage OS

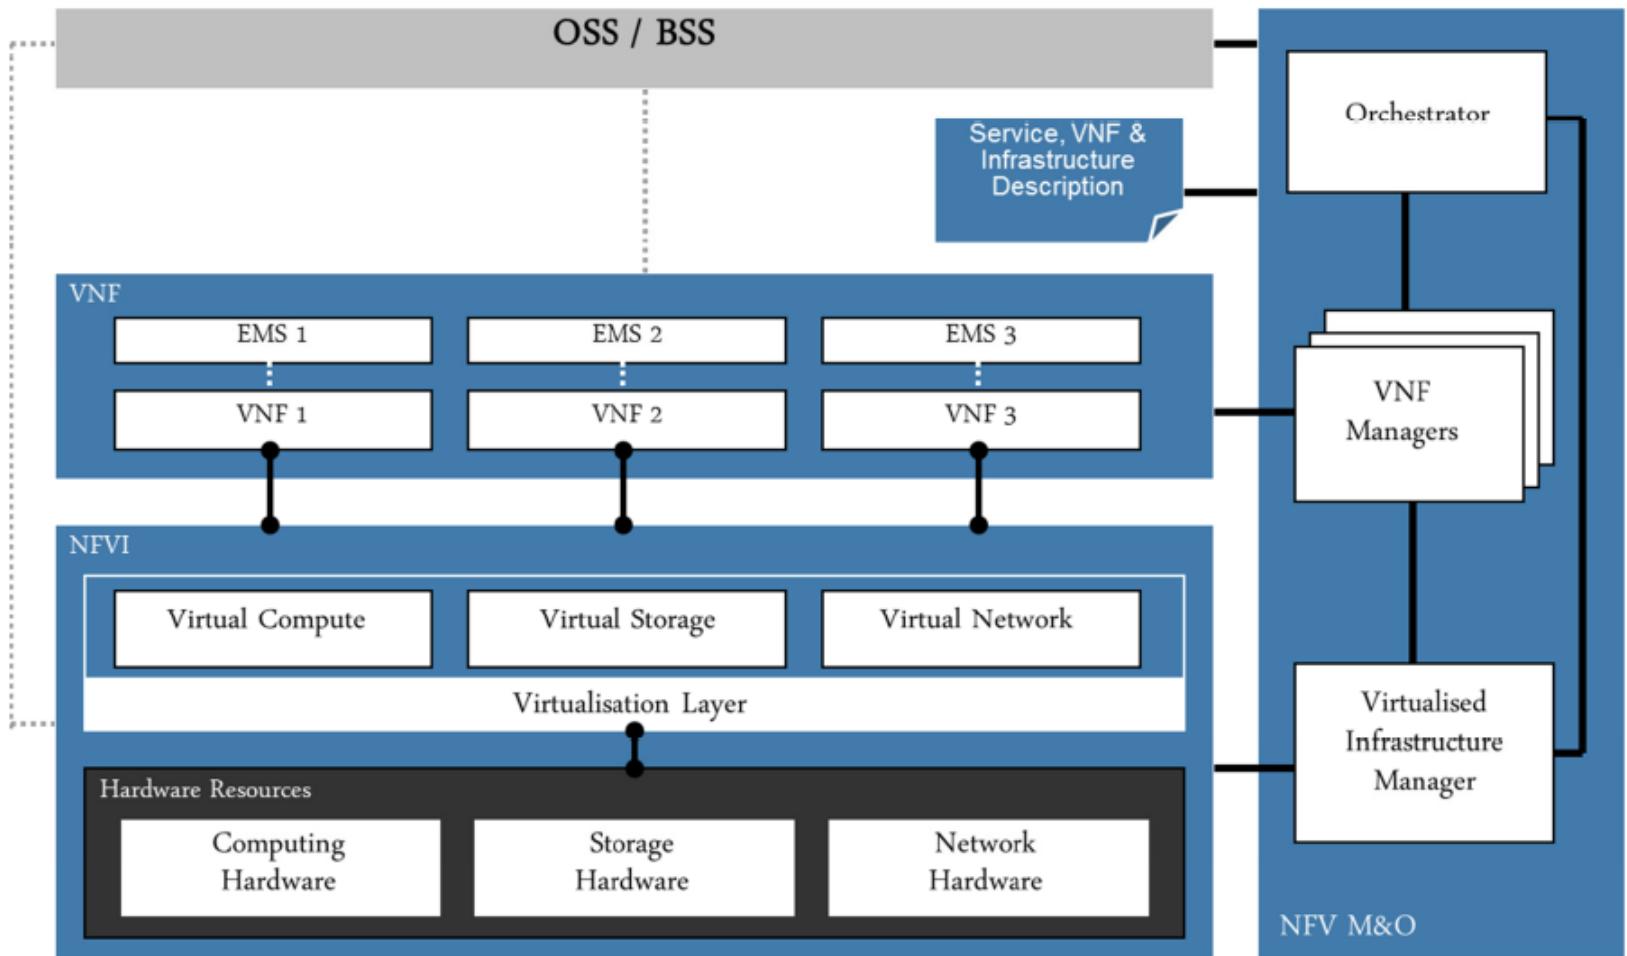
- <https://mirage.io>
- Ocaml-based runtime and app development environment

OsV

- <http://osv.io/>
- Optimized OS
- Applications as single processes, no protection between user space and kernel
- Optimizations: zero-copy, lockless network APIs, enhanced JVM etc.
- Linux-compatible

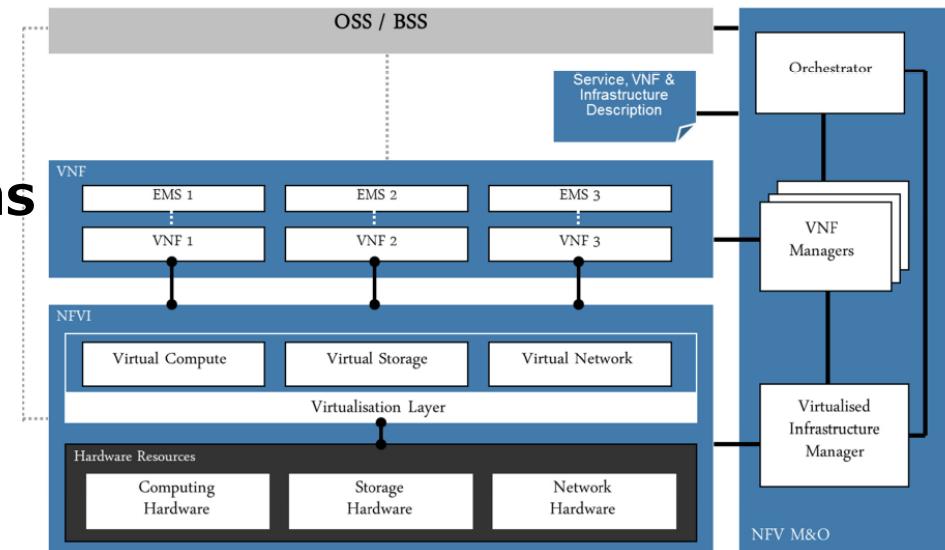
- | **Virtualization Performance**
- | **More fine-granular virtual machines, appliances**
- | **Higher dynamicity in instantiating, destroying VMs**
- | **DC workload assessment, prediction**
- | **New functions based on library OSs**

NFV Architecture (ETSI NFV ISG)



Infrastructure (NFVI)

- Provides virtual resources to support execution of VNFs
- COTS hardware, accelerator components
- Software-layer to virtualize and abstract from underlying hardware



Virtualized Network Functions (VNFs)

- Software implementations of network functions
- Running on the NFVI
- Can come together with an Element Management System (EMS)

Management and Orchestration (MANO)

- Orchestration and life-cycle management of physical/software resources
- Life-cycle management of VNFs
- Examples: on-boarding, failure recovery, scale-out
- Interacts with Operations/Business Support Systems (OSS/BSS)

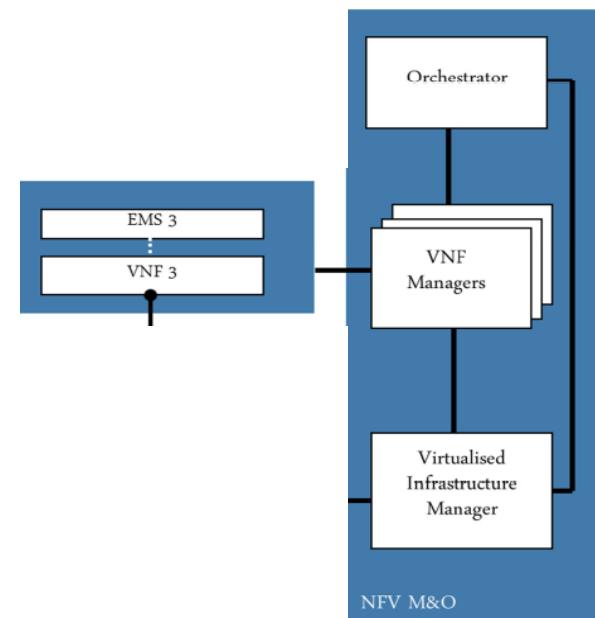
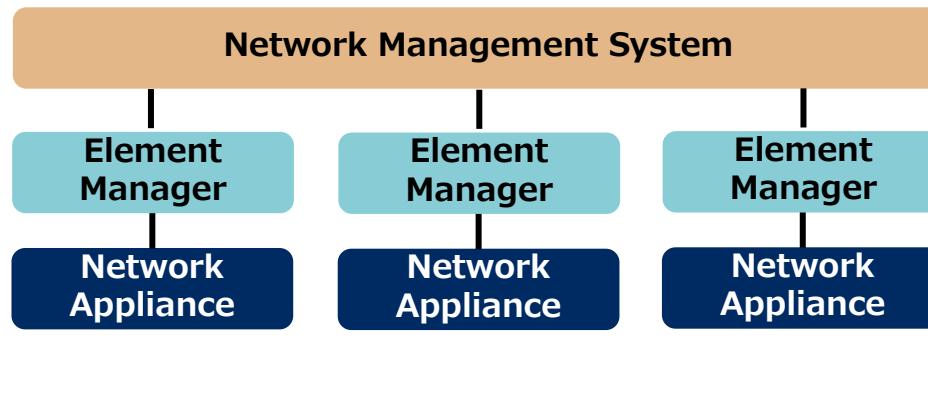
Automation in NFV

Two main motivations for moving telco networks into the cloud

1. **Resource efficiency** (consolidation of functions onto fewer servers)
2. **Automation** (fully automated management of resources and functions)

Traditional telco management approach

- Specific (vendor-provided) element manager and network management systems
- Specific management approaches, fault recovery mechanisms etc.

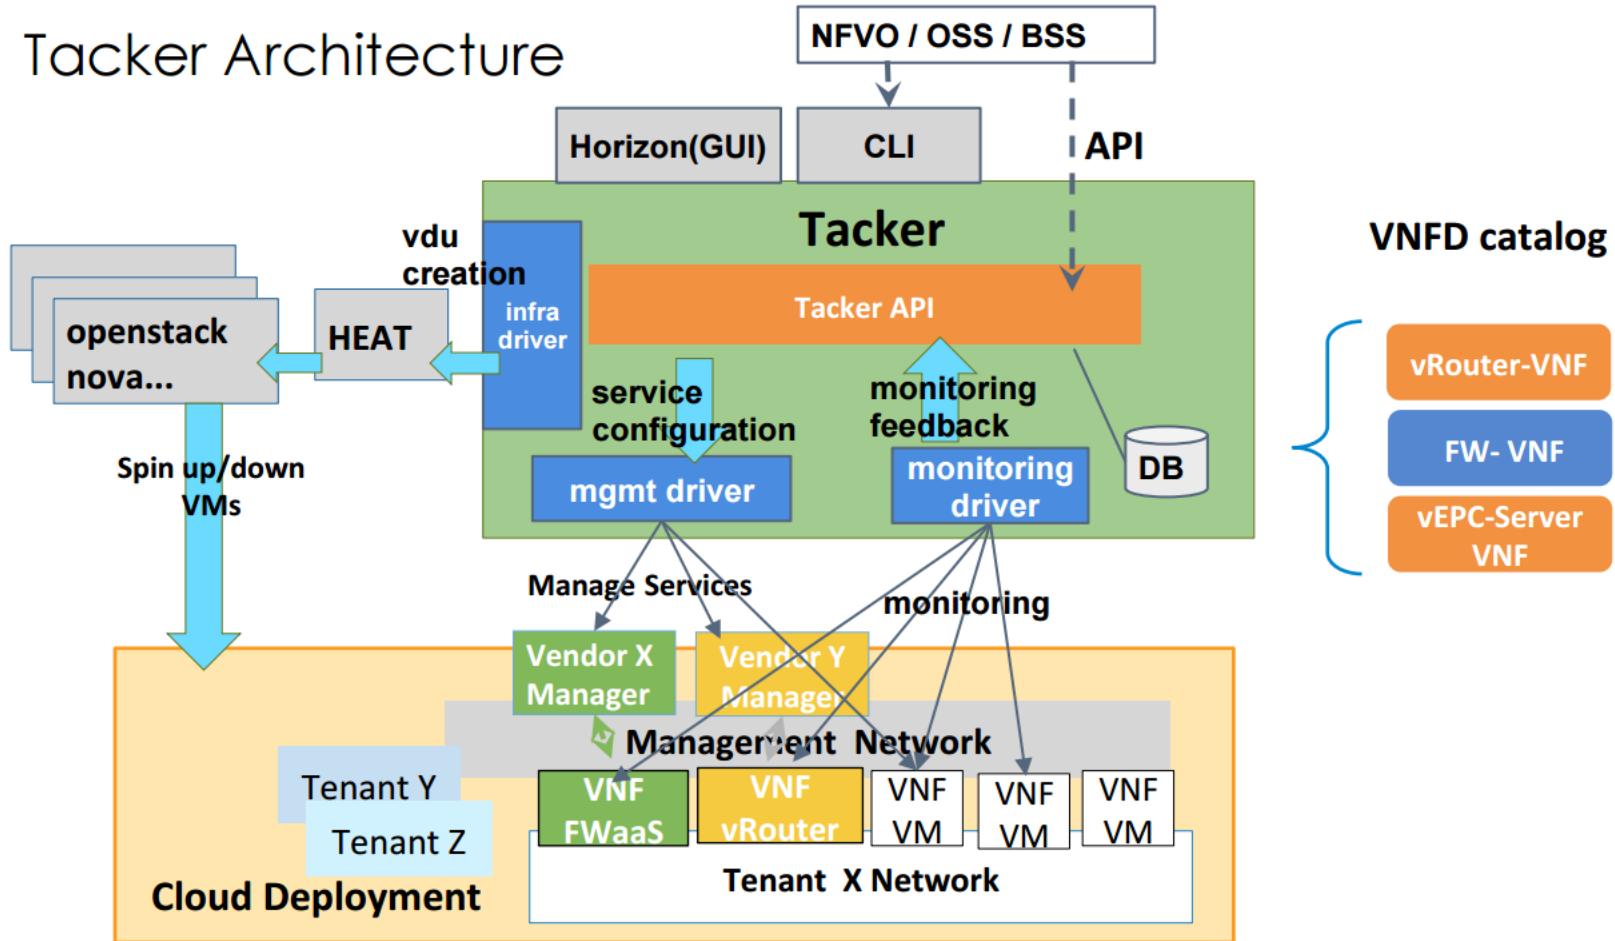


Cloud management approach

- Generalized VNF management
- Part of NFV management and orchestration

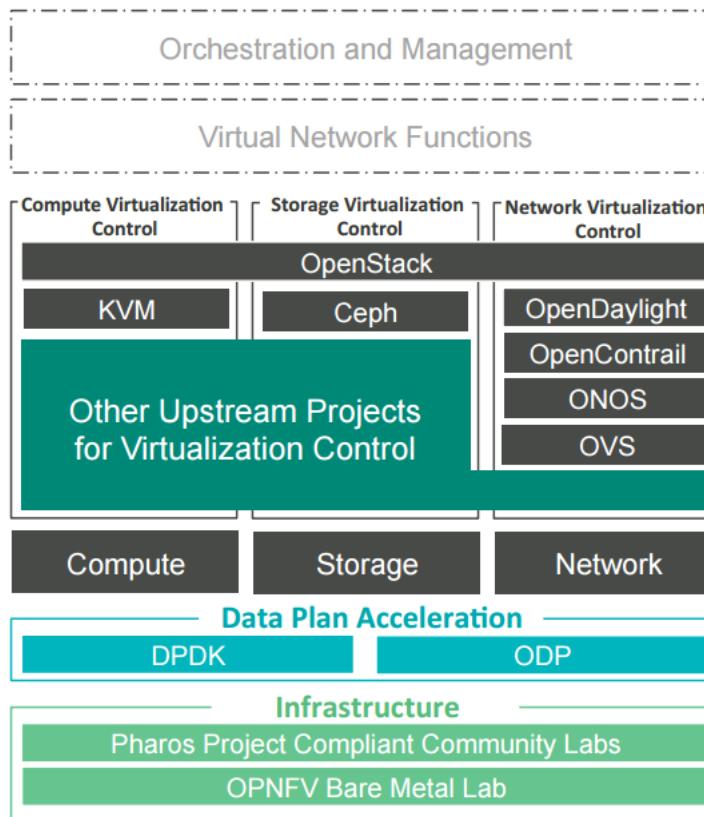
OpenStack Tacker

Tacker Architecture



<https://www.ietf.org/proceedings/93/slides/slides-93-nfvrg-25.pdf>

Open Source NFV



www.opnfv.org

Next steps for auto management

- Managing larger sets of smaller VMs
- Light-weight management
- Predicting load, failures
- Optimizing instantiation, VM positioning etc.

Telco networks

- Operator services
- Optimizers
- Service differentiation
- Policy enforcement for QoS, resource management etc.

Traffic passes through different network functions

- Depending on user context, flow specifics, network weather condition etc.

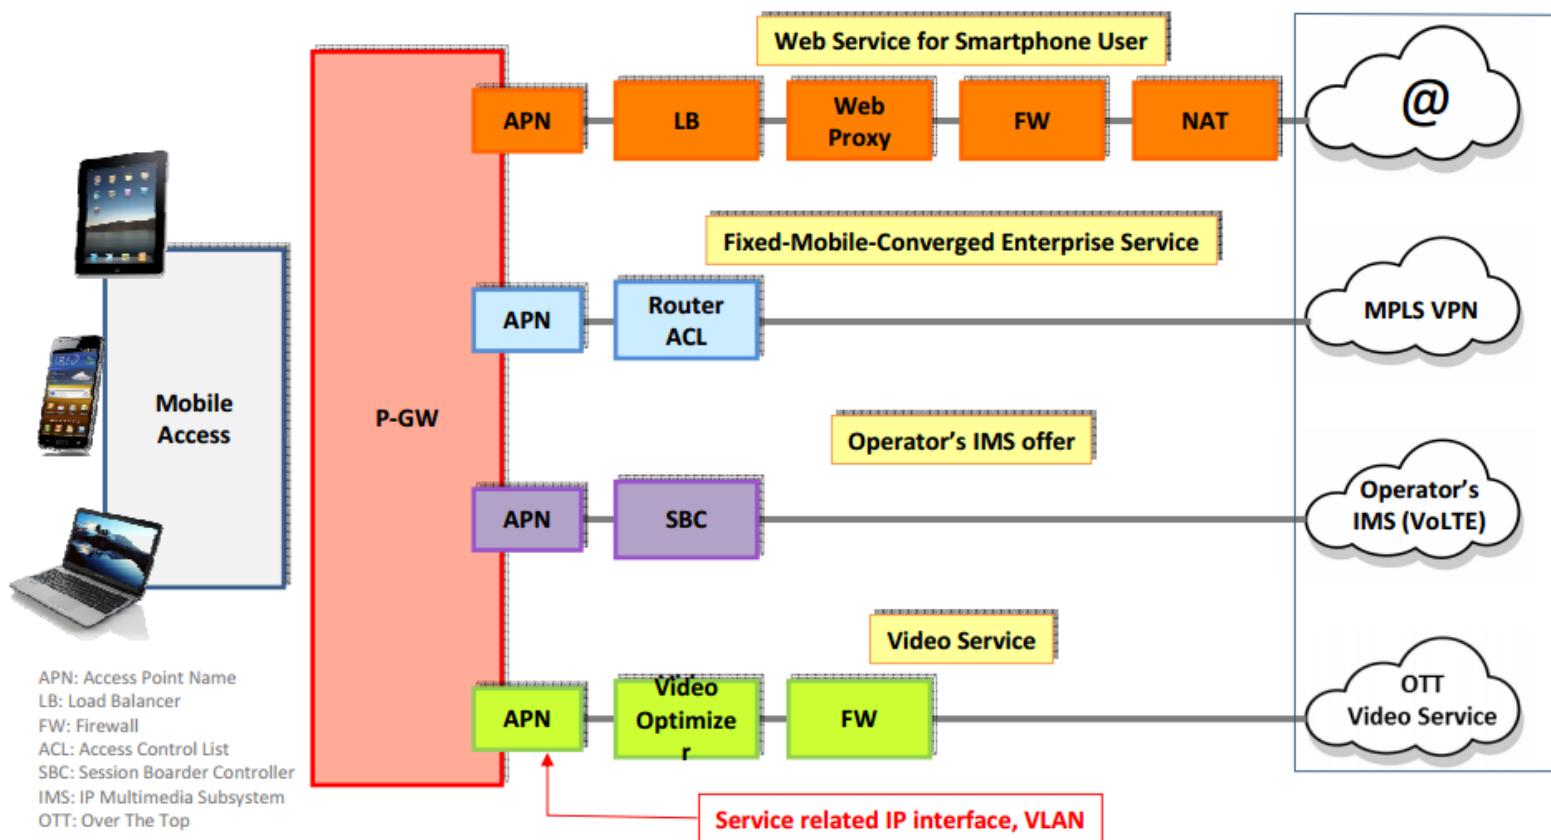
Different types of functions

- Network layer translation: CG-NAT/PAT
- Security related: intrusion-detection, firewalls, access control, en/decryption etc.
- Qos related: video optimizers, TCP optimizers, policers
- User identification/tracking: HTTP header enrichment
- Other „value-added“ services: parental control, malware protection

Vodafone Example

Service Function Chaining

Current Common Approach – Logical View on Typical Use Cases



Service Chaining Requirements

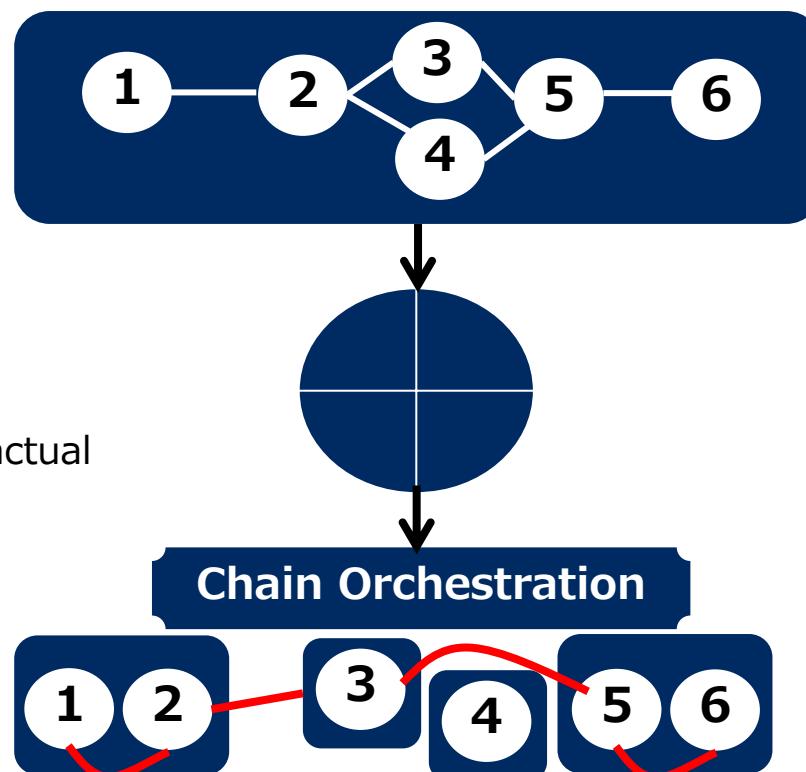
Legacy service chaining

- Static chains, hard-wired on network layer – through manual configuration
- Not very flexible – difficult to change chain configurations
- Connectivity graphs dependent on actual network graph
- Too complicated to allow for per-user chains
- Not compatible to cloud networking (automatic scale out etc.)

Ideally:

Abstract specification

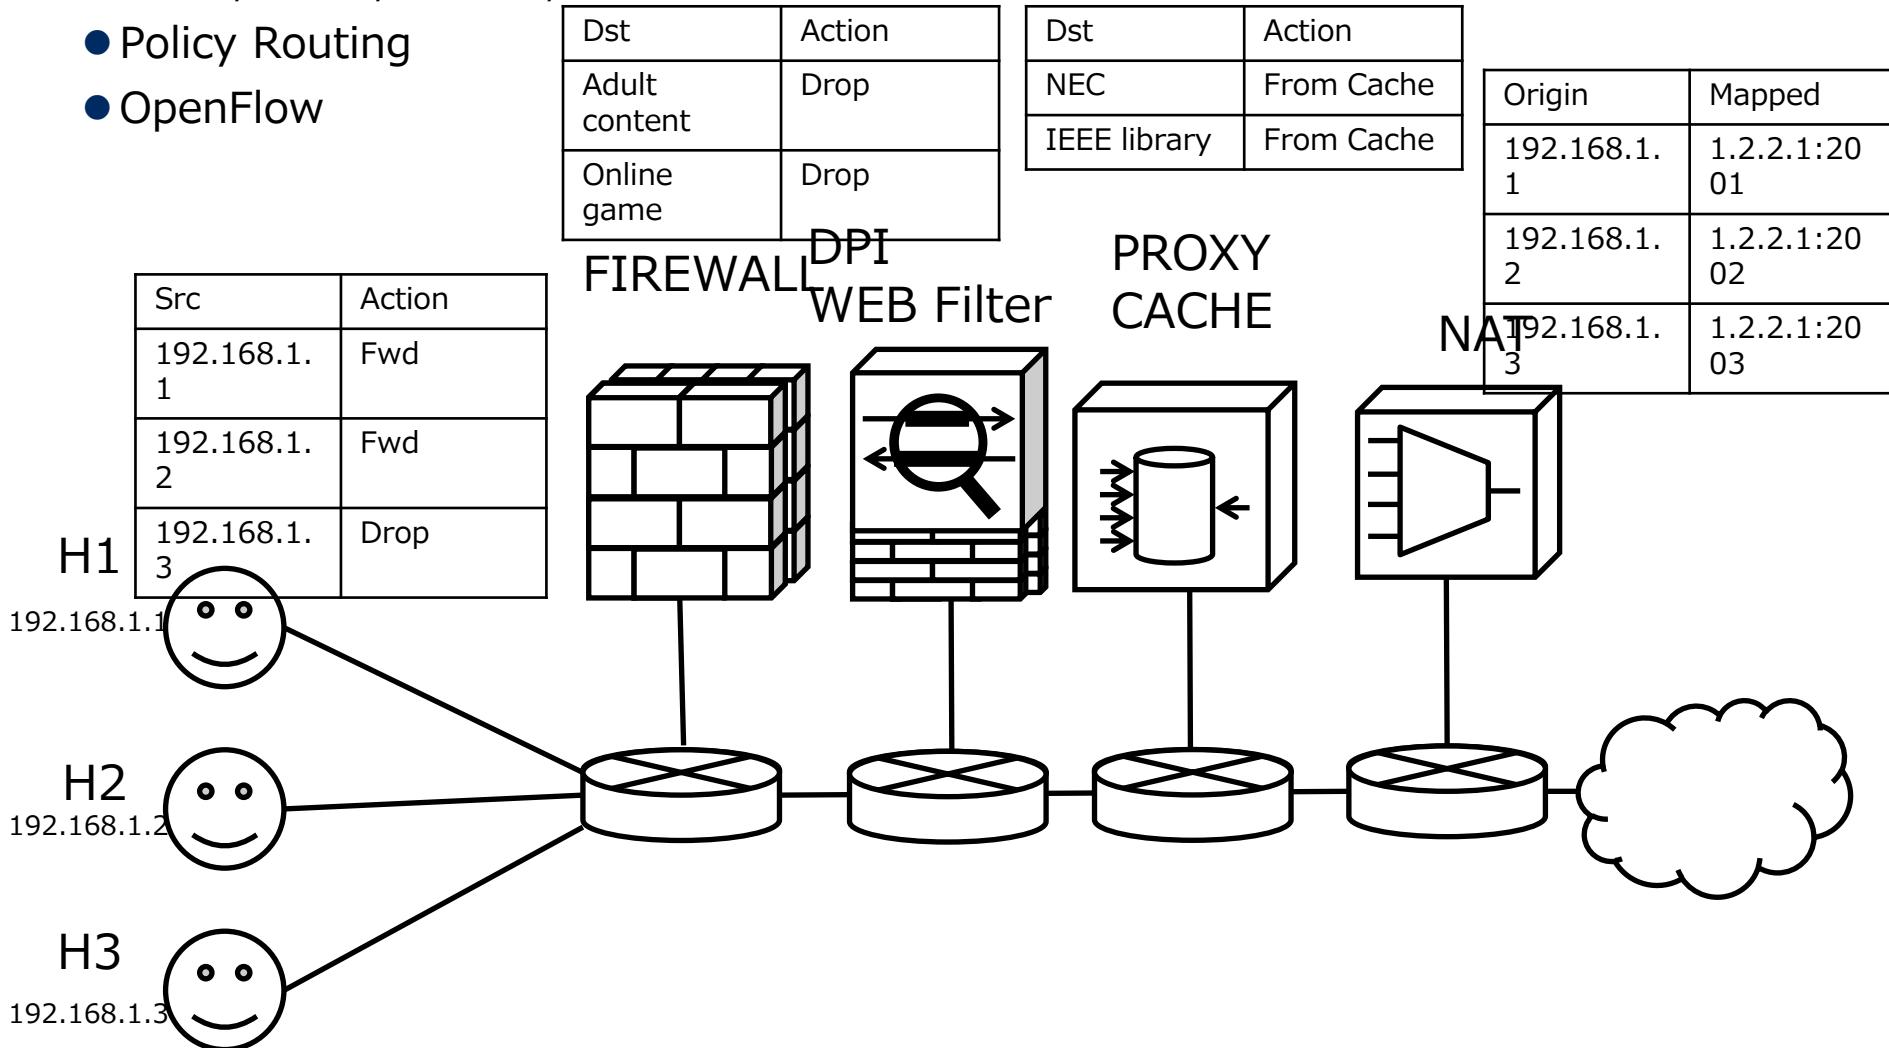
- Conditions
 - Scaling requirements
-
- Service chain compiler
 - Creating initial embedding on actual network graph
 - Management and orchestration
 - Actual instantiation
 - Monitoring
 - Scaling, graph adaptation etc.



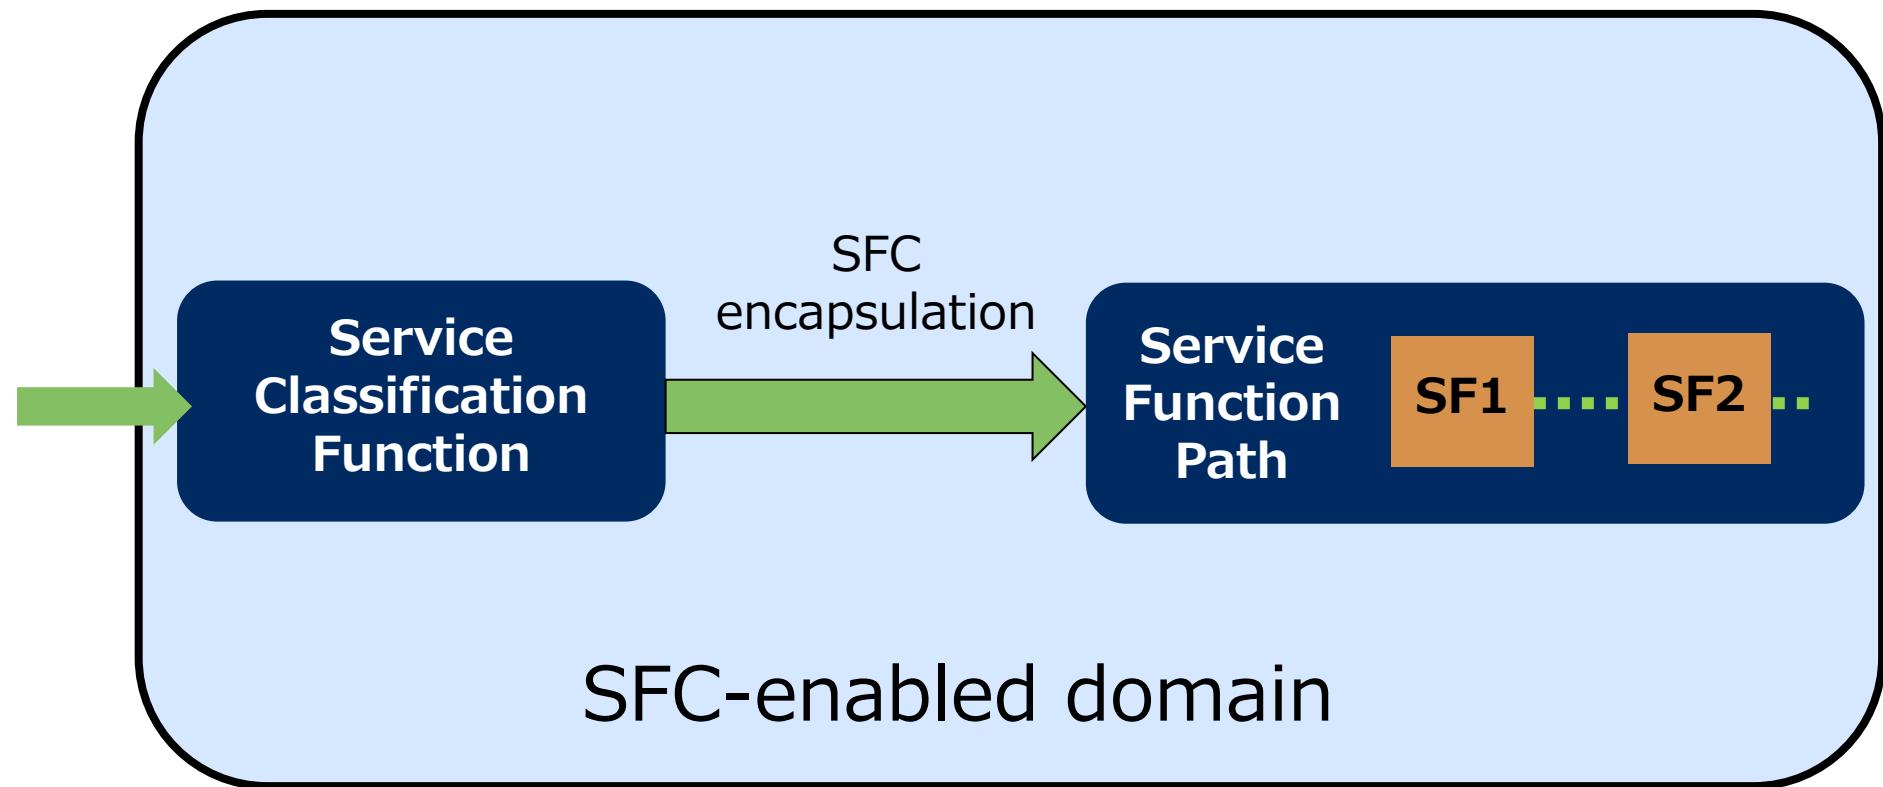
Implementing Service Chaining

Flow Steering – different approaches

- VLAN, MPLS, VXLAN, GRE
- Policy Routing
- OpenFlow

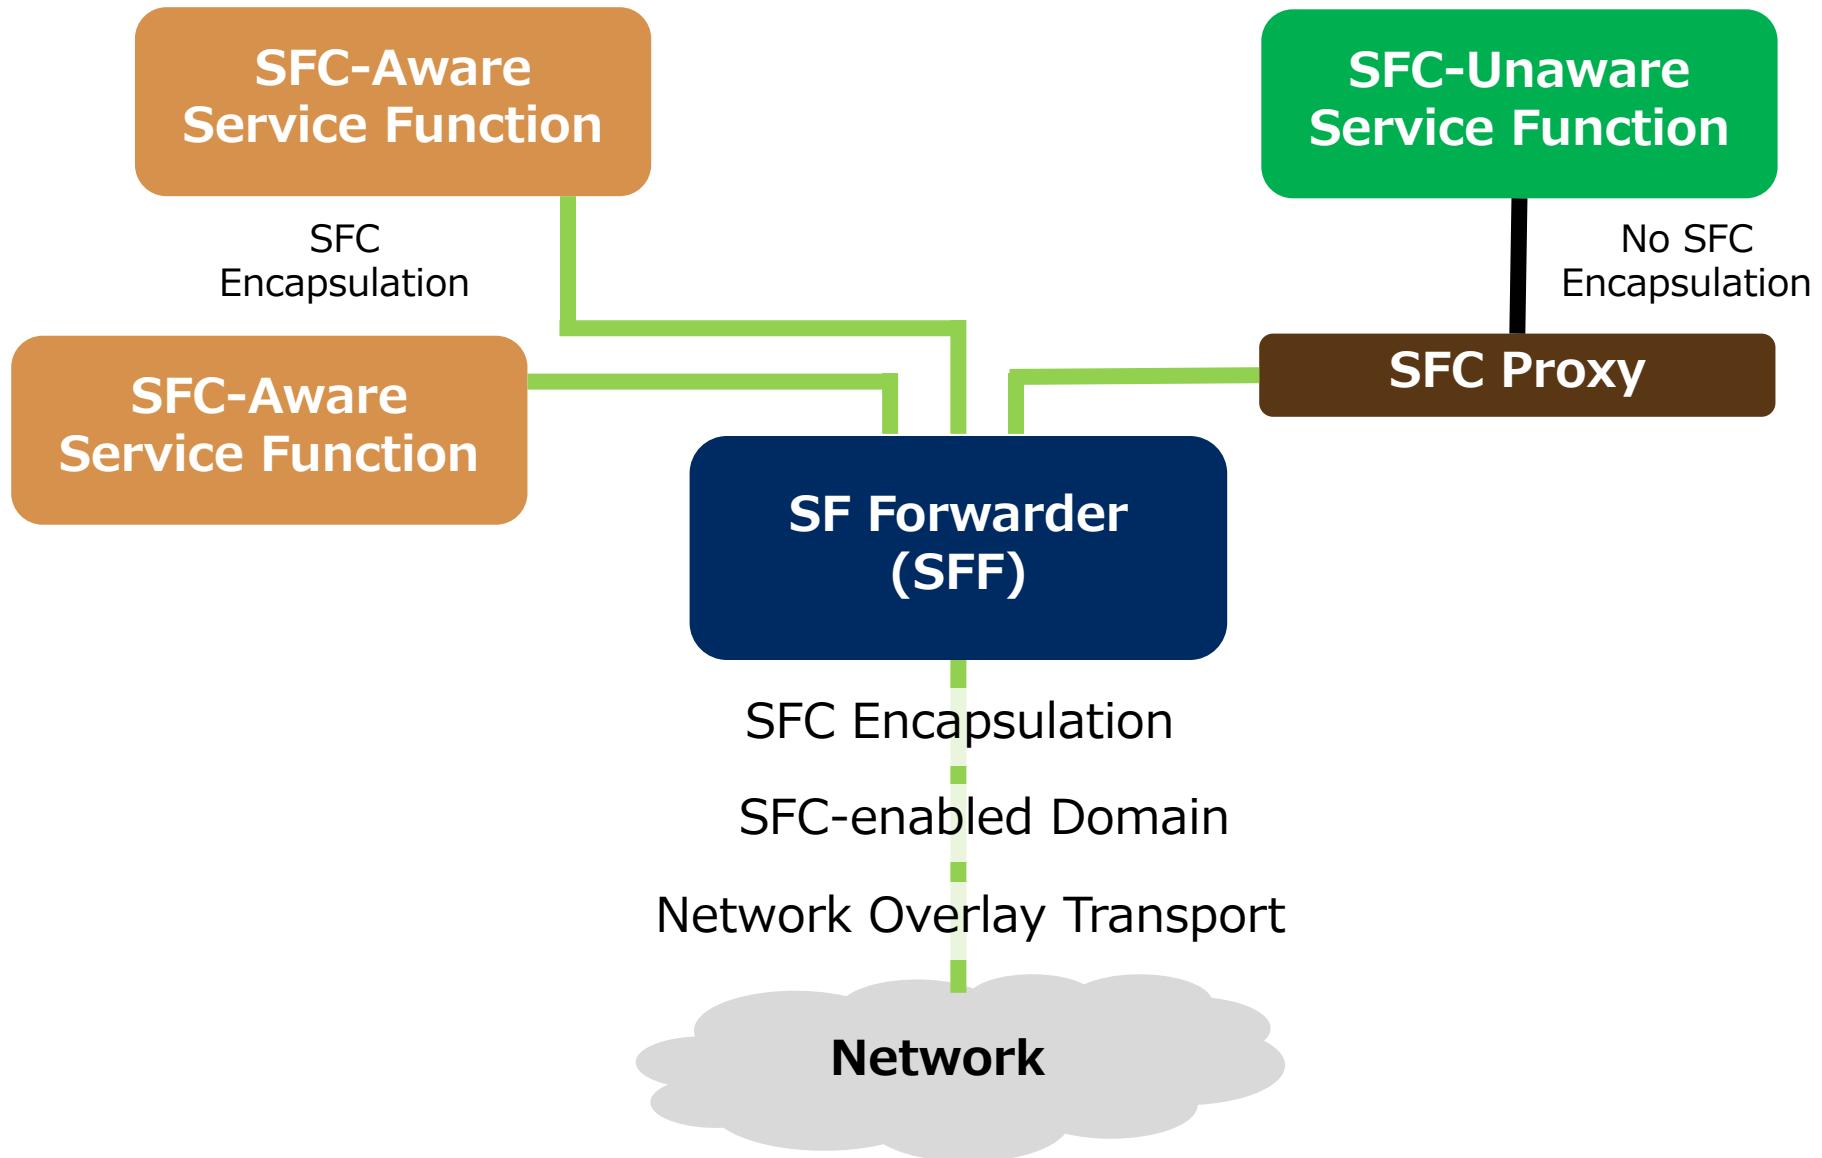


IETF SFC Architecture



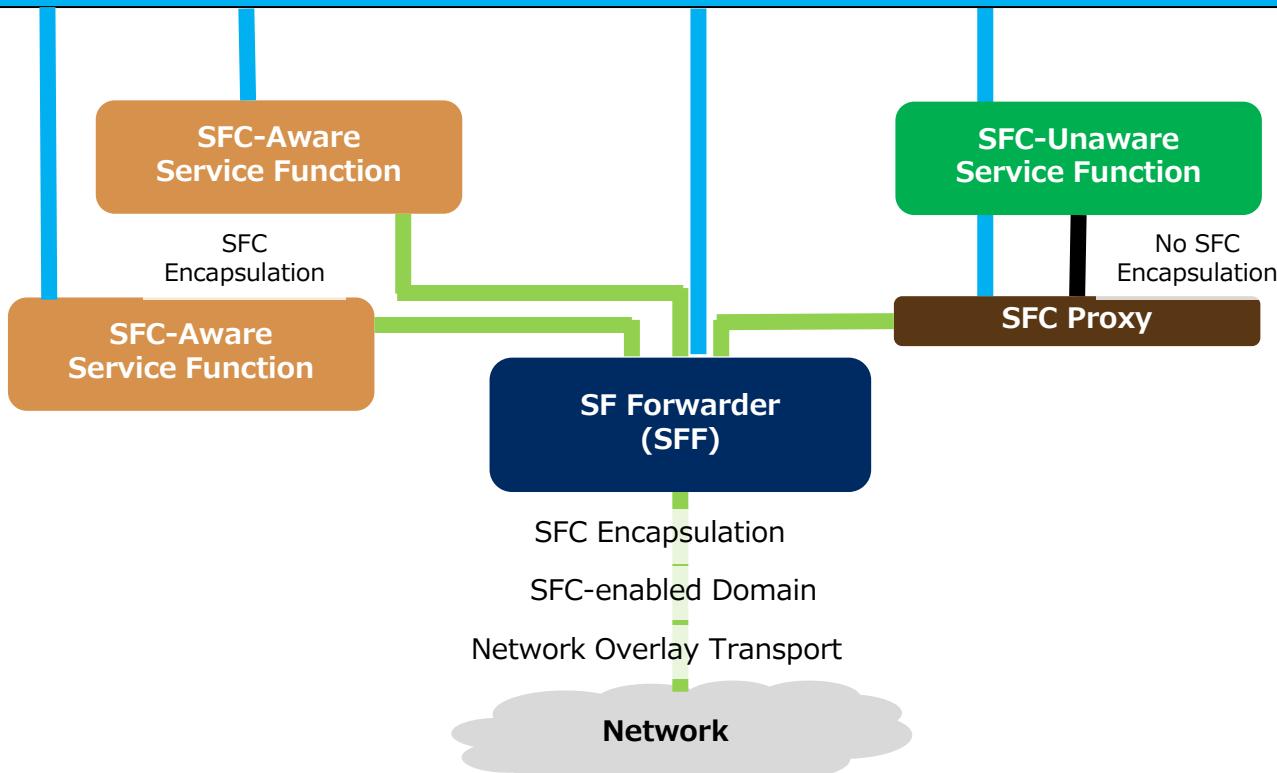
- Encapsulation-based approach
- Intended to support SFC-unaware service functions

IETF SFC Architecture



SFC Control Plane

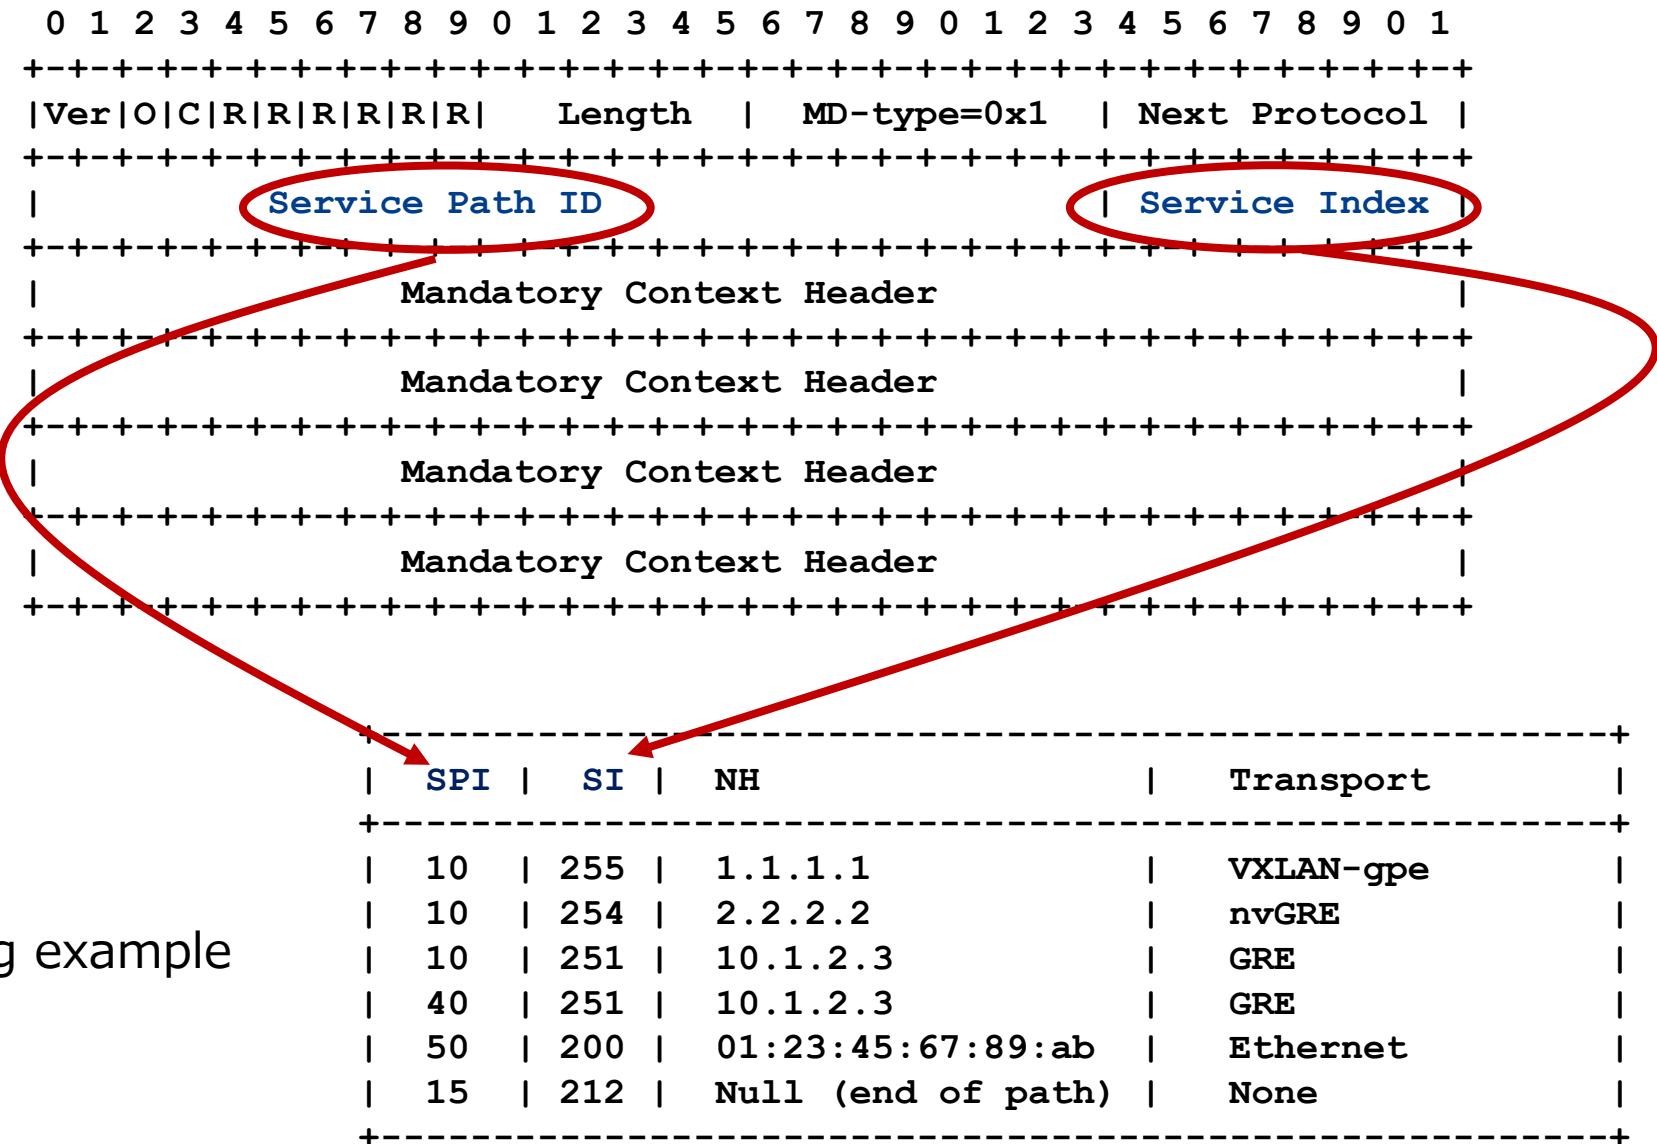
- Domain-wide view of SF resources and their locators
- Uses SFC Policy to construct service function chains and actual paths
- Selection of SFs for SFCs (statically or dynamically)
- Provide SFC data plane to components
- Provide metadata and usage info to classifiers
- Provide policy information to SFC elements to interpret metadata



NSH Header Type 1

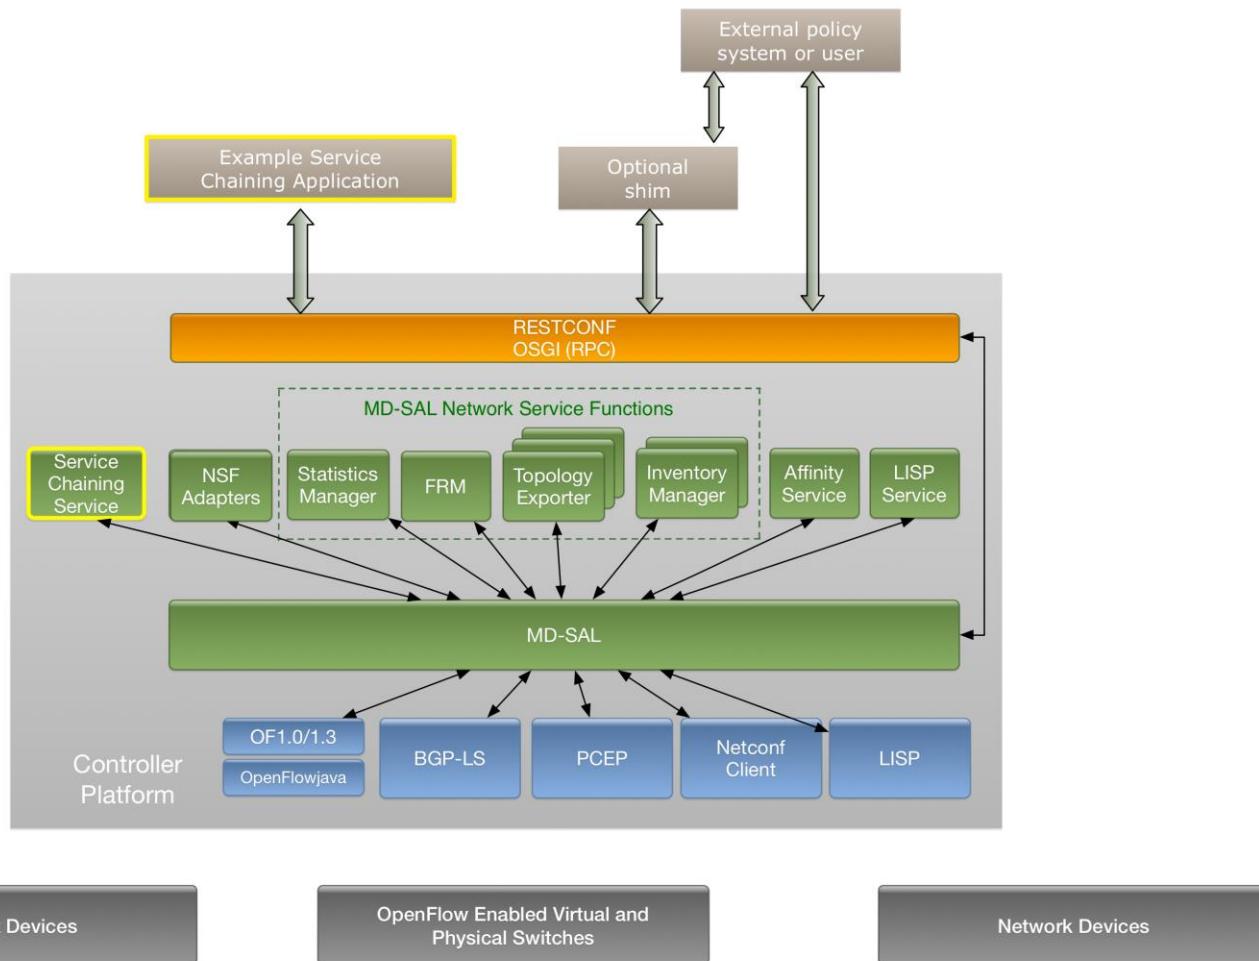
0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	
	Ver		O		C		R		R		R		R		R		Length		MD-type=0x1		Next Protocol	
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+
	Service Path ID													Service Index								
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+
	Mandatory Context Header																					
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+
	Mandatory Context Header																					
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+
	Mandatory Context Header																					
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+
	Mandatory Context Header																					
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+

NSH Header Type 1



Maping example

OpenDaylight SFC Project



https://wiki.opendaylight.org/view/Service_Function_Chaining:Main

Service Chaining Research Topics

Transport Performance

- Impact on TCP performance
- Placement optimizations

Security

- Service chaining in the presence of ubiquitous encryption
- „Cooperative Traffic Management“

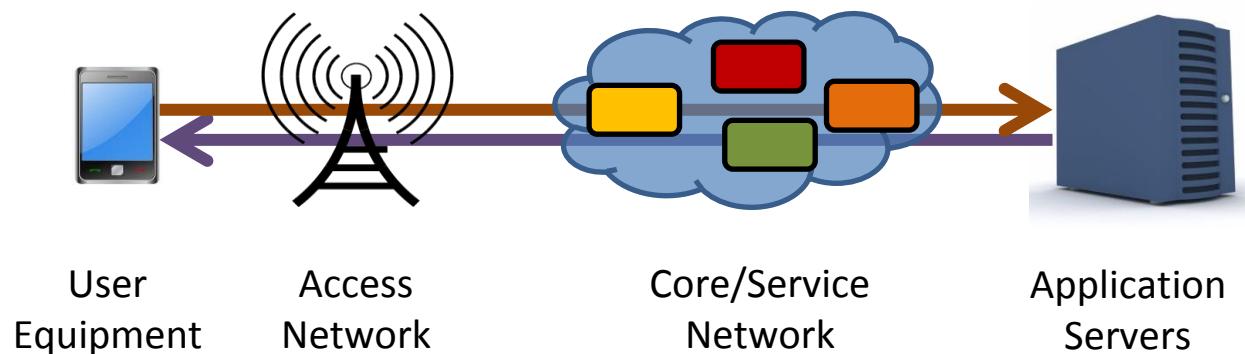
Micro service chains

- Fine-granular service chaining structure
- Compose functions from building blocks
- Reason about optimal ordering, parallelism, placement

New approaches to chaining in-network functions

- Arumaithurai et al.; Exploiting ICN for Flexible Management of Software-Defined Networks
- Integrating distributed computing and storage: Named Function Networking

Applying SDN/NFV



- Softwarization previously in data centers only (cloud)
- Now moving towards core and access networks

- New options for enhanced functionality, new features, news architectures

Information-Centric Networking

Accessing Named Data Objects (NDOs) in the network

- ADUs, chunks, fragments

Data-centric security approach

- Disentangled means for name-content binding validation, publisher authentication, confidentiality

Name-Content binding validation:

hash-based schemes

Publisher authentication

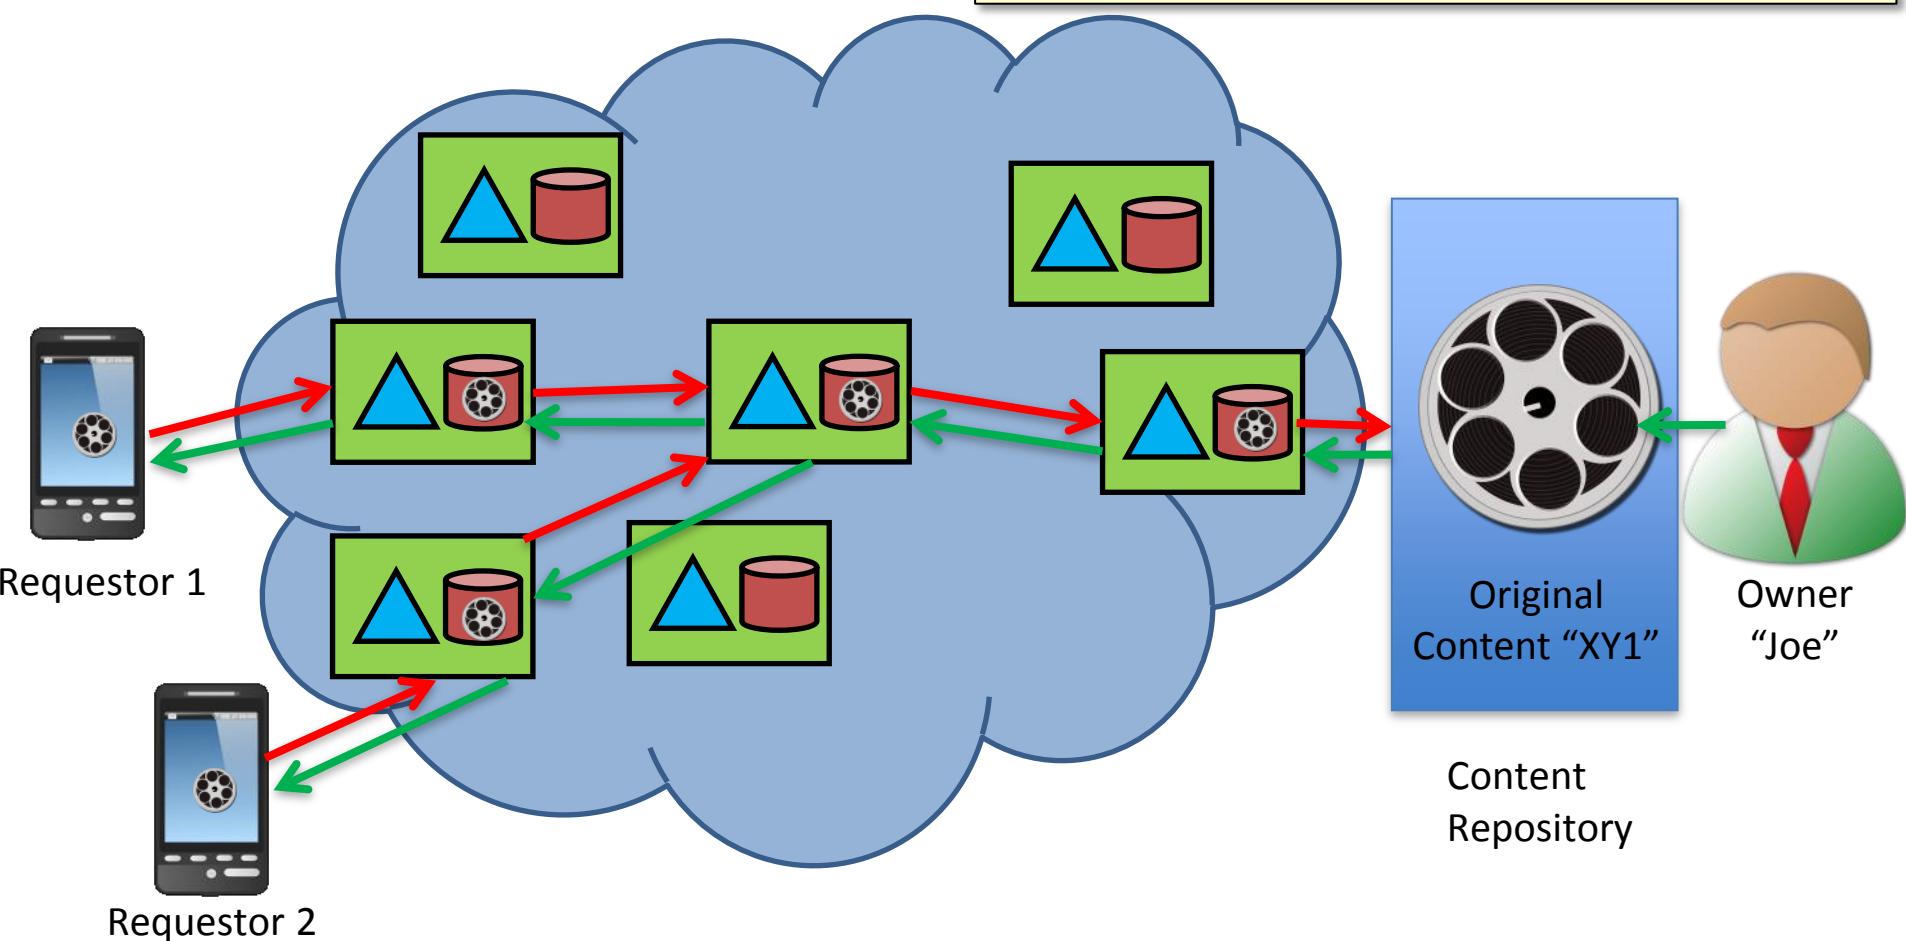
- One approach: publishers to sign NDOs, signature part of NDO meta data; trust model a la PKI

Confidentiality

- Encrypting payload with publisher key

ICN Overview

- Request Response
- Pending Interesting Tables
- Receiver-driven
- Route-by-name (prefix)
- Ubiquitous caching



ICN Traffic and Resource Management

Key ICN properties

- Requesting individual Named Data Objects
- Ubiquitous Caching

Implicit caching

- Every router can store NDO – depending on configuration, policy etc.
- Even with encrypted traffic, caching can help with local retransmissions, media replay etc.

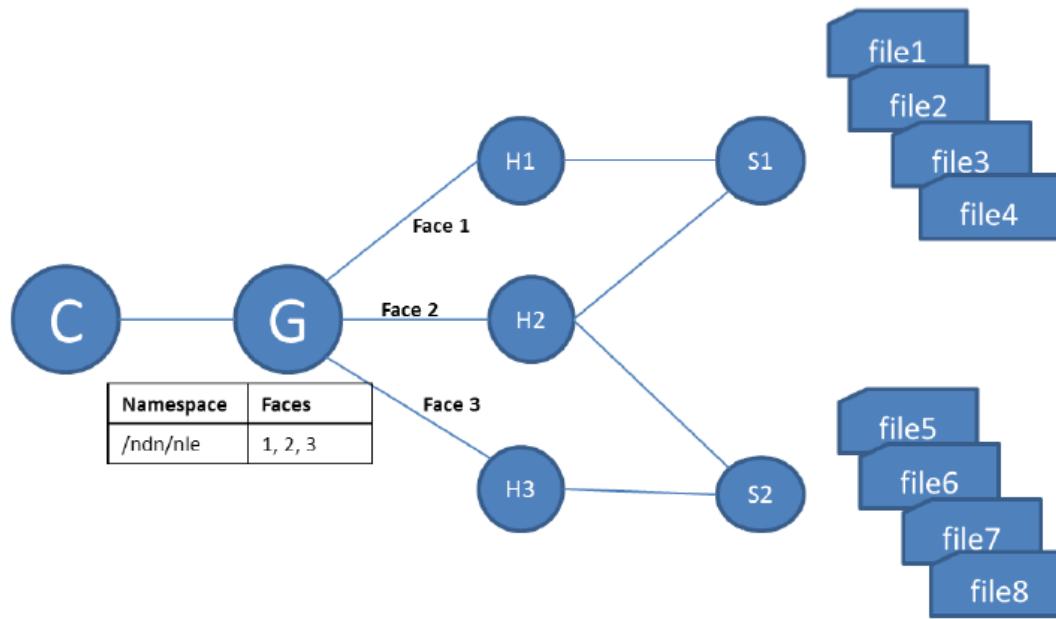
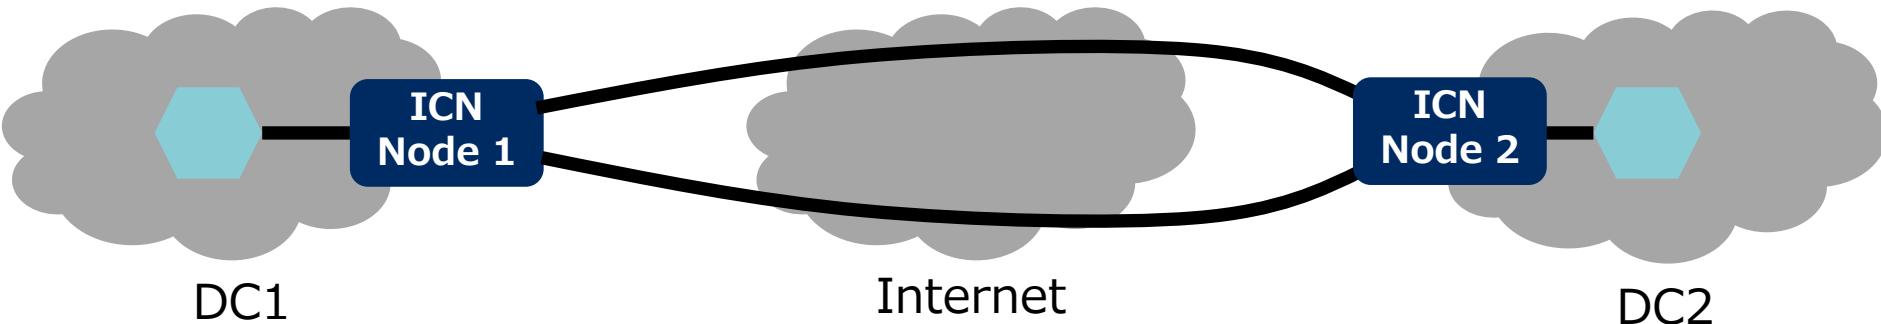
Flexible multipath communication

- Every router can make forwarding decisions depending on strategy, network characteristics, name prefix, policy

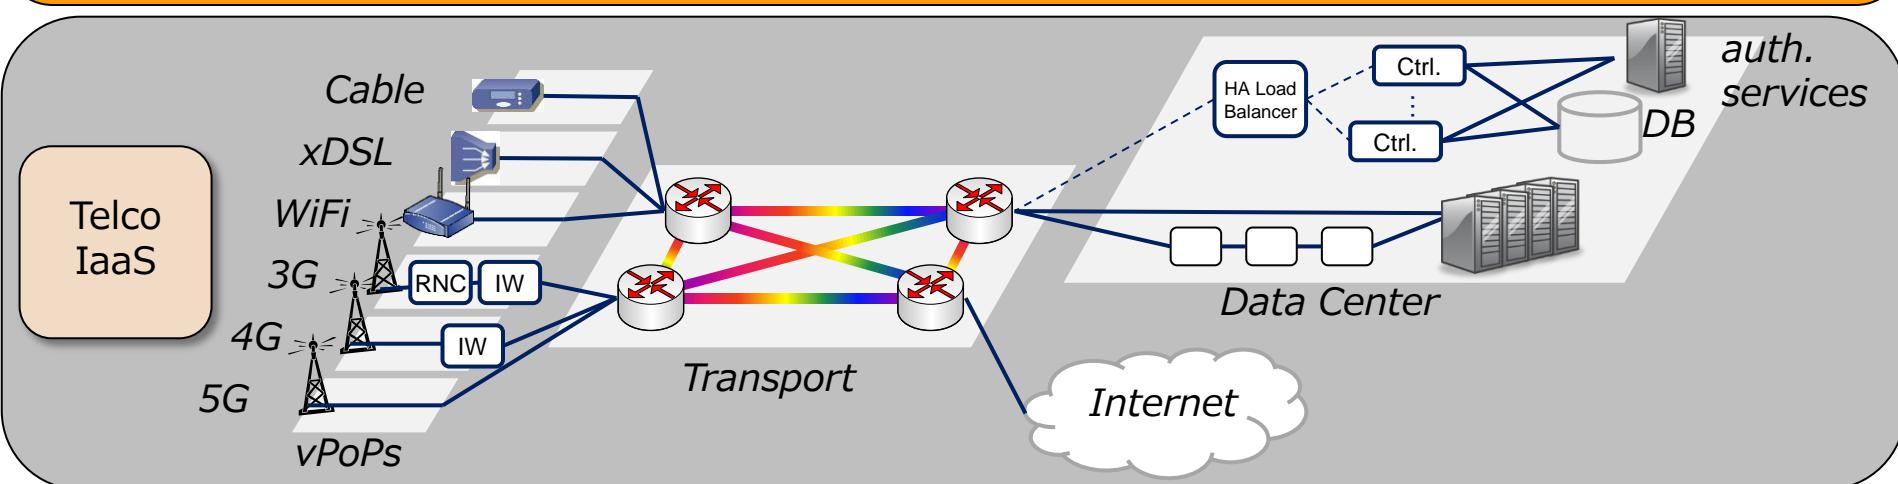
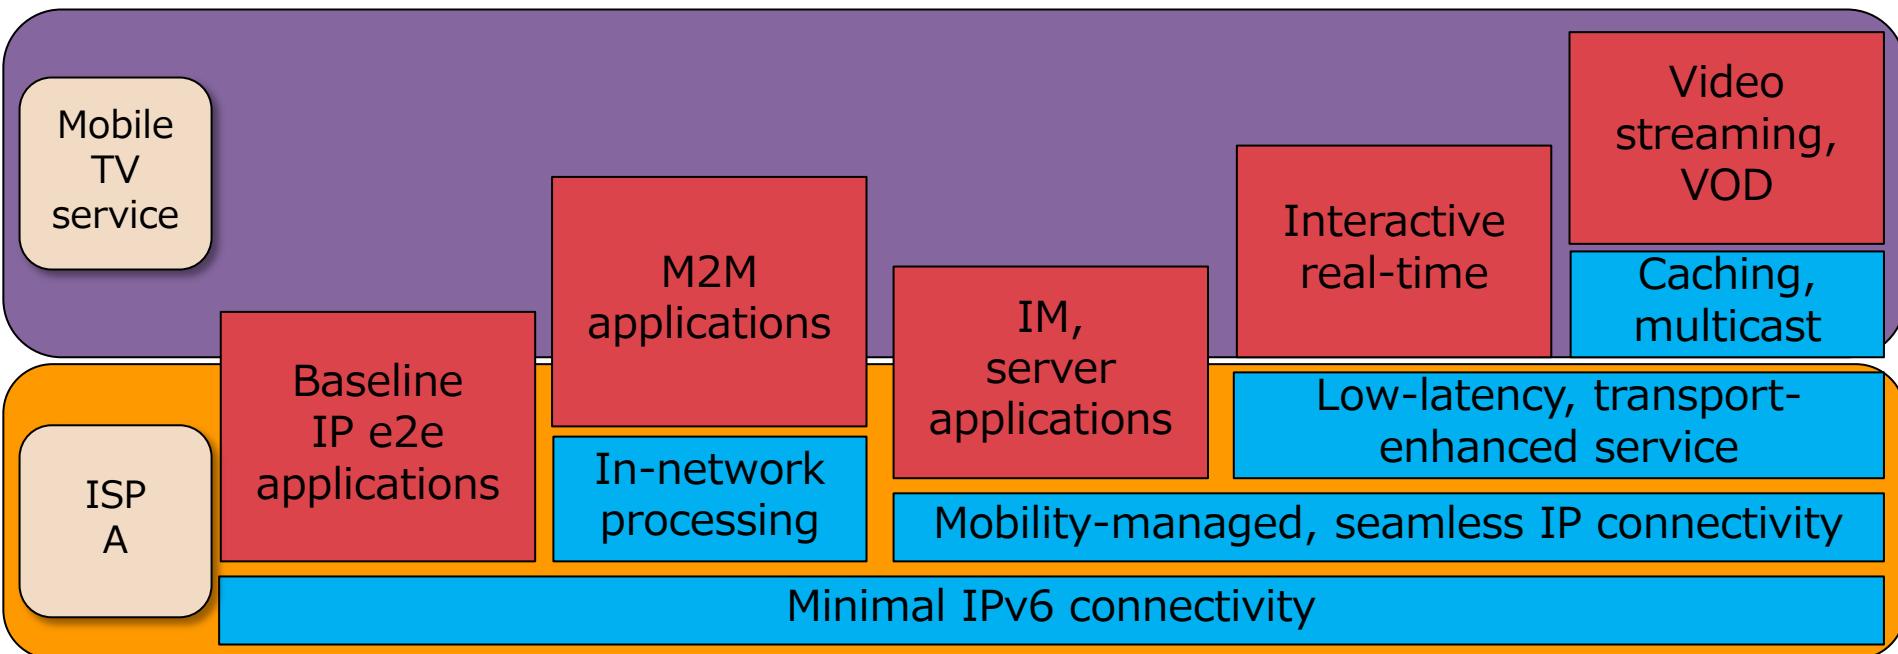
Easy policing and filtering

- Requestors, publishers and requestors see ICN requests and responses
- Policing without DPI
- Enabling other optimizations: in-network pre-fetching etc.

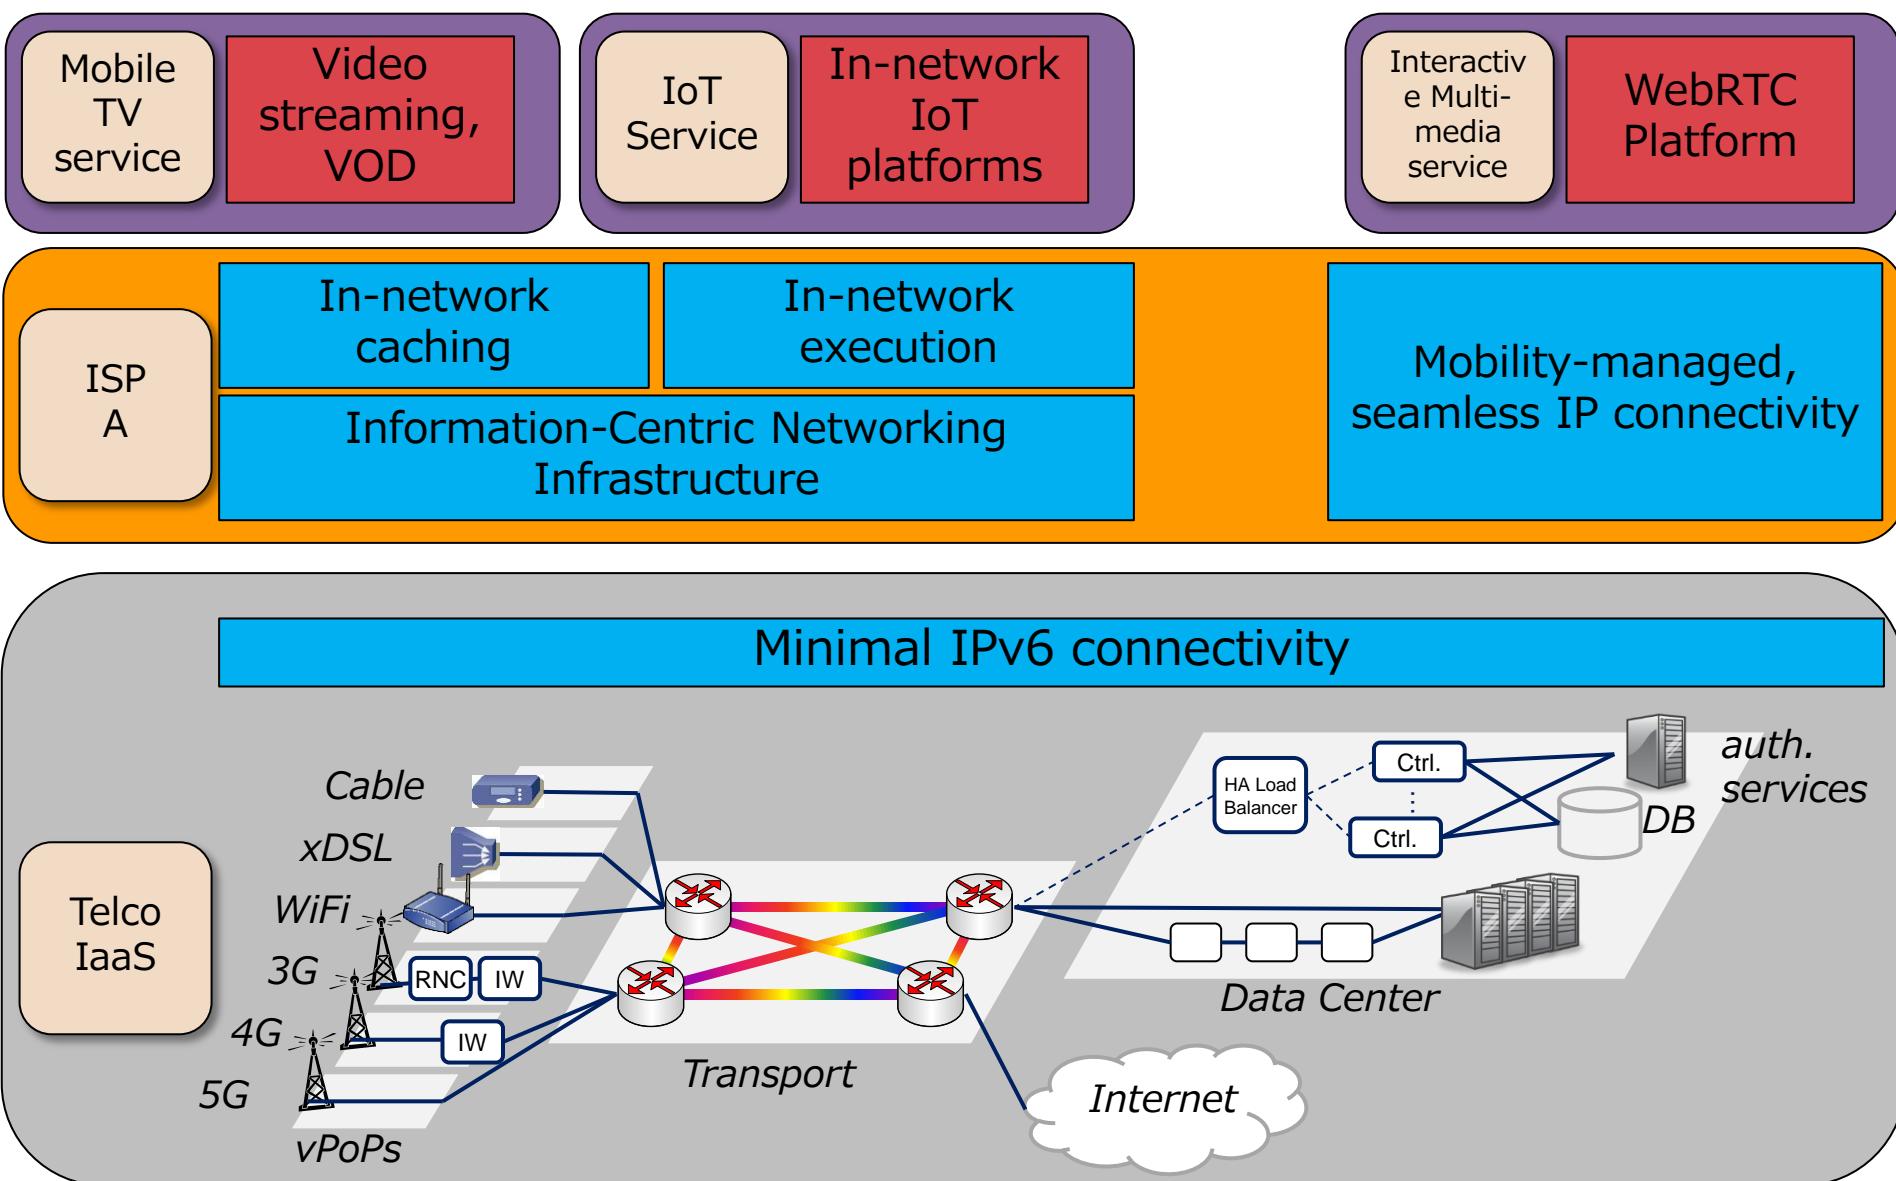
Better Multipath-Support through ICN



ICN in 5G Blueprint – Organizational Layers



5G Blueprint Reshuffled and Modernized



SDN References

- SANE: <http://dl.acm.org/citation.cfm?id=1267346>
- Ethane: <http://dl.acm.org/citation.cfm?id=1282382>
- OpenFlow: <http://dl.acm.org/citation.cfm?id=1355746>
- NOX: <http://dl.acm.org/citation.cfm?id=1384625>
- OVS: <http://openvswitch.github.io/papers/hotnets2009.pdf>
- ONF SDN Arch: [TR-502](#), Overview [TR-504](#)
- Further ONF specifications: www.opennetworking.org
- More scientific papers @
<https://sites.google.com/site/sdnreadinglist/>

NFV References

Overview

- [https://portal.etsi.org/nfv/nfv white paper.pdf](https://portal.etsi.org/nfv/nfv_white_paper.pdf)
- [https://portal.etsi.org/nfv/nfv white paper2.pdf](https://portal.etsi.org/nfv/nfv_white_paper2.pdf)
- [https://portal.etsi.org/Portals/0/TBpages/NFV/Docs/NFV White Paper3.pdf](https://portal.etsi.org/Portals/0/TBpages/NFV/Docs/NFV_White_Paper3.pdf)

Open Source

- www.opnfv.org
- [https://wiki.opendaylight.org/view/Service Function Chaining:Main](https://wiki.opendaylight.org/view/Service_Function_Chaining:Main)

IRTF NFVRG

- <https://irtf.org/nfvrg>

IETF Service Chaining

- <https://tools.ietf.org/wg/sfc/>

Interested in a Master Thesis project in our lab?

Dirk.Kutscher@neclab.eu

www.neclab.eu

Orchestrating a brighter world

未来に向かい、人が生きる、豊かに生きるために欠かせないもの。

それは「安全」「安心」「効率」「公平」という価値が実現された社会です。

NECは、ネットワーク技術とコンピューティング技術をあわせ持つ

類のないインテグレーターとしてリーダーシップを発揮し、

卓越した技術とさまざまな知見やアイデアを融合することで、

世界の国々や地域の人々と協奏しながら、

明るく希望に満ちた暮らしと社会を実現し、未来につなげていきます。

\Orchestrating a brighter world

NEC