

Software Defined Networking

SDN Mobile Wireless




TECHNISCHE
UNIVERSITÄT
DARMSTADT

David Hausheer

Department of Electrical Engineering
and Information Technology
Technische Universität Darmstadt

E-Mail: hausheer@ps.tu-darmstadt.de
<http://www.ps.tu-darmstadt.de/teaching/sdn>



*Based on original slides by Giuseppe Bianchi (University of Roma Tor Vergata)



Wireless MAC layer Reconfigurability *from an SDN perspective*

Giuseppe Bianchi, University of Roma Tor Vergata

Credits to: I. Tinnirello, P. Gallo, D. Garlisi, F. Giuliano, F.
Gringoli

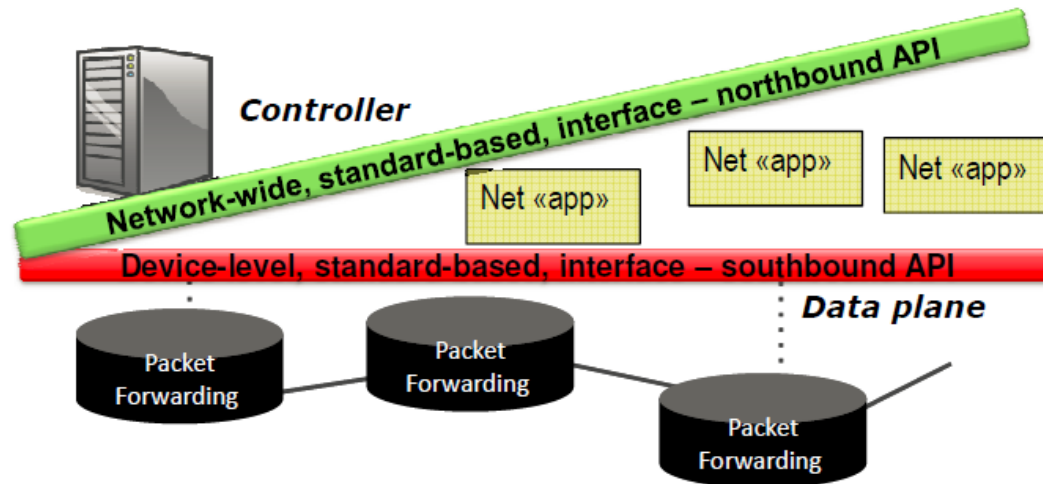
Software-Defined Radio (from wikipedia)

- ❖ A software-defined radio system, or SDR, is a radio communication system where **components** that have been typically implemented in hardware (e.g. mixers, filters, amplifiers, modulators/demodulators, detectors, etc.) are instead **implemented by means of software** on a personal computer or embedded system.
- ❖ **20+ years long research path**
 - AirBlue, CalRadio, GNURadio, RUNIC, SORA, USRP, WARP,...
- ❖ **Niche commercial exploitation**
 - Military, etc

Software-Defined Networking (from wikipedia)

- ❖ Software defined networking (SDN) is an approach to building computer networks that **separates and abstracts elements of these systems** [...] SDN allows network administrators to have programmable central control of network traffic without requiring physical access to the network's hardware devices.
- ❖ **5 years long research path**
 - Pioneered by 2008 OpenFlow paper
- ❖ **Almost 2B\$ company acquisitions in 2012**
 - Mainly Nicira, but also Contrail, Big Switch, Cariden, Vyatta,...

Why SDN == \$\$\$



- ❖ Business: provisioning and control of network services
 - Fostering easy deployment => fast innovation
- ❖ Technical enabler: open configuration APIs
 - Southbound (device level): typically openflow (but not necessarily)
 - Northbound (network control): hot topic
- ❖ Today: mostly on southbound (wireless differs!)

Common misconceptions

❖ SDN = OpenFlow

- OpenFlow is just ONE possible data-plane API

❖ SDN = control/dataplane separation

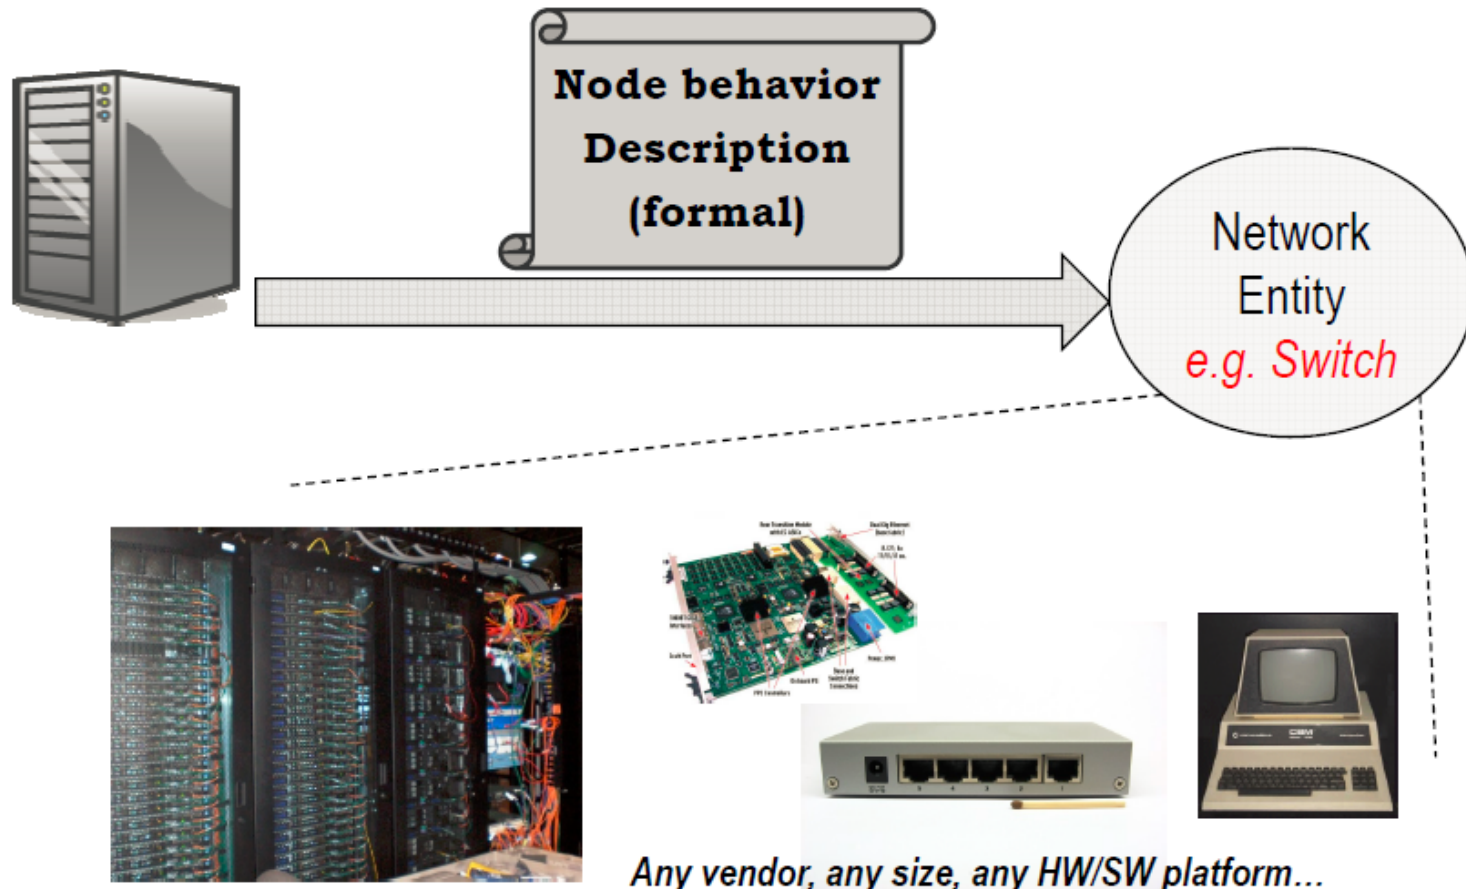
- SDN relies ALSO on control/data plane separation
- but does not reduce to just this

❖ Would not nearly be new, especially in wireless!

- Well established separation in the wireless domain
- Enterprise WLAN controllers, CAPWAP, etc: around for more than a decade

SDN: it's all about abstractions

- ❖ So far mostly dealt with in wired networks



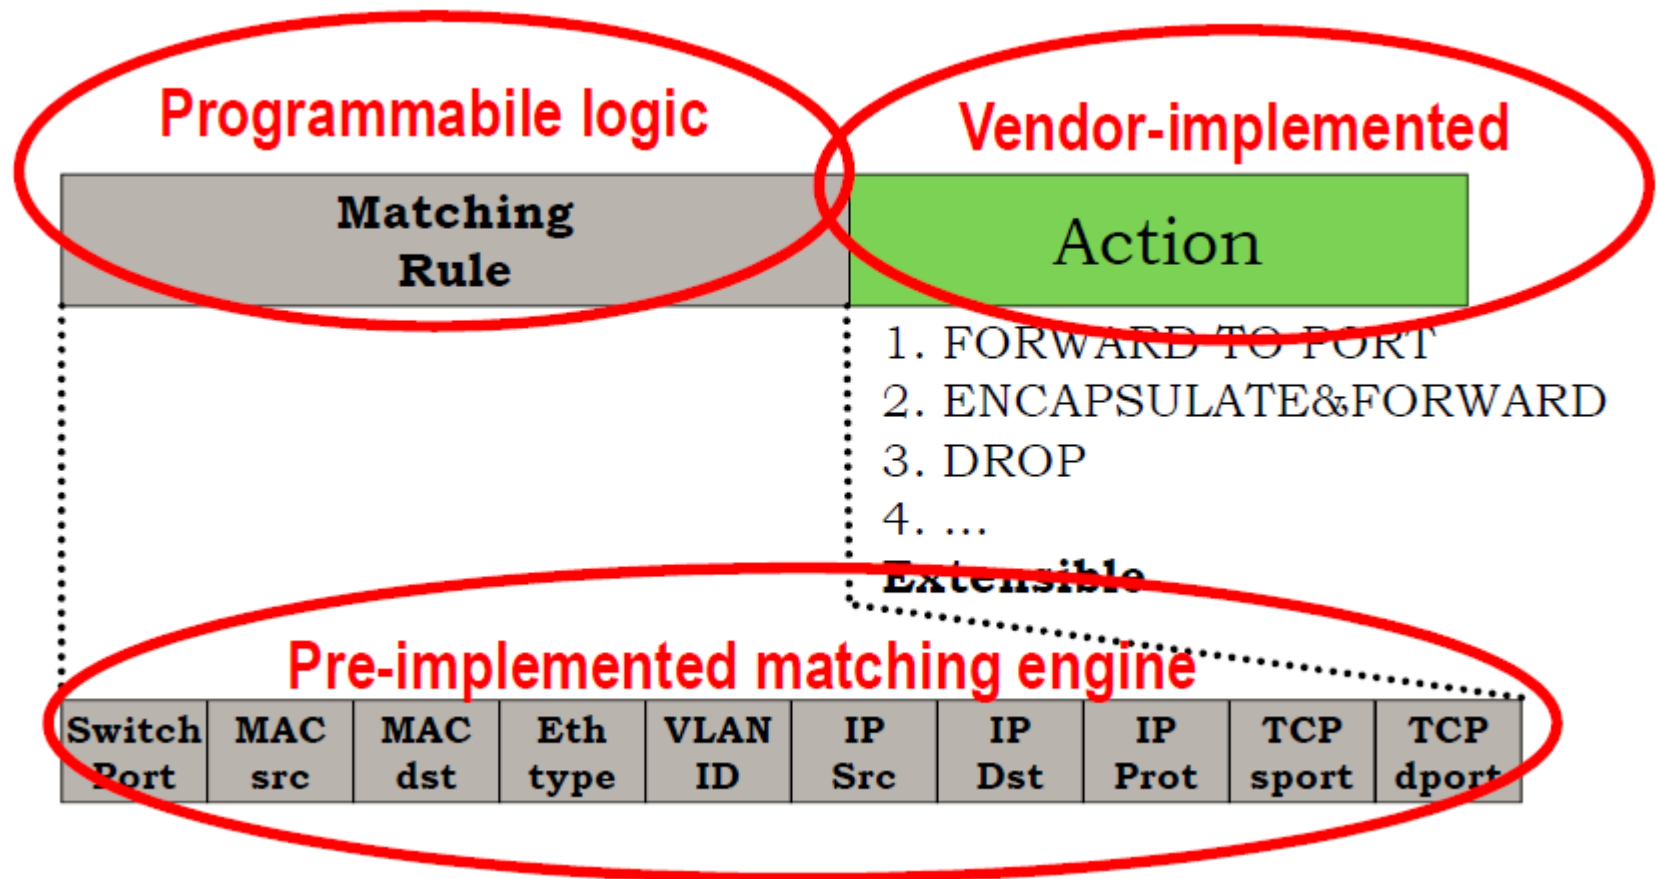
OpenFlow: a compromise

- ❖ **Best approach:** “persuade commercial name-brand equipment vendors to provide an open, programmable, virtualized platform on their switches and routers”
 - Plainly speaking: *open the box!! No way...*
- ❖ **Viable approach:** “compromise on generality and seek a degree of switch flexibility that is”
 - High performance and low cost
 - Capable of supporting a broad range of research
 - **Consistent with vendors’ need for closed platforms.**
- ❖ A successful compromise, indeed... ask Nicira ...

[original quotes: from OF 2008 paper]

OpenFlow: just *one* abstraction

❖ Good for switches, not for «all»



What about SDN in wireless?

❖ Wireless Openflow...

- Wireless specific actions: very helpful...
- ... but match/action API way too skinny
 - We all agree now: SDN >> OpenFlow

❖ Challenge: which viable programming abstractions for wireless terminals and nodes?

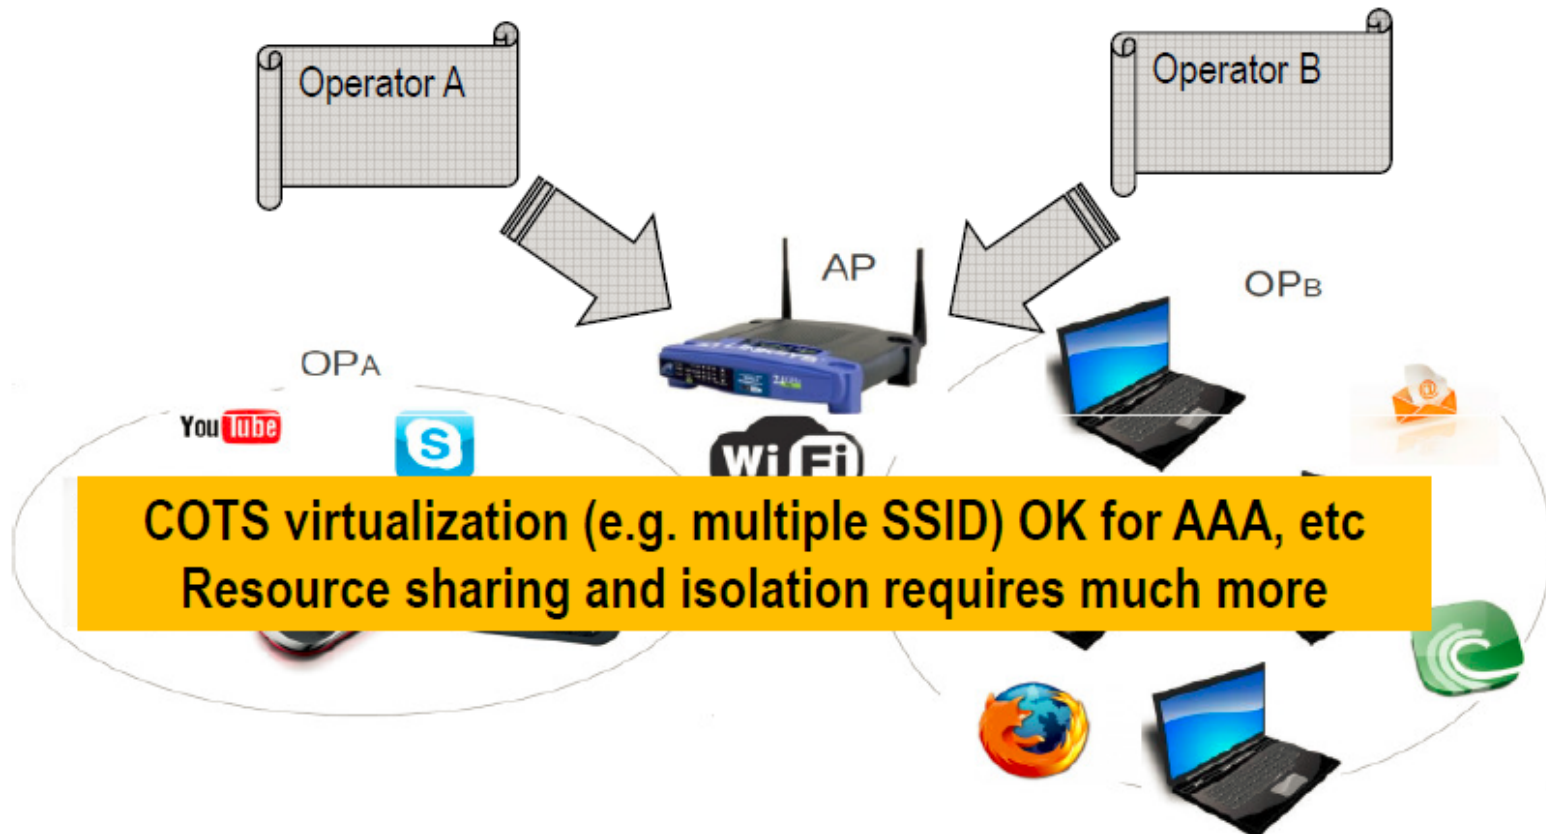
- Without requiring to «open the box»

Beneficial to multiple scenarios

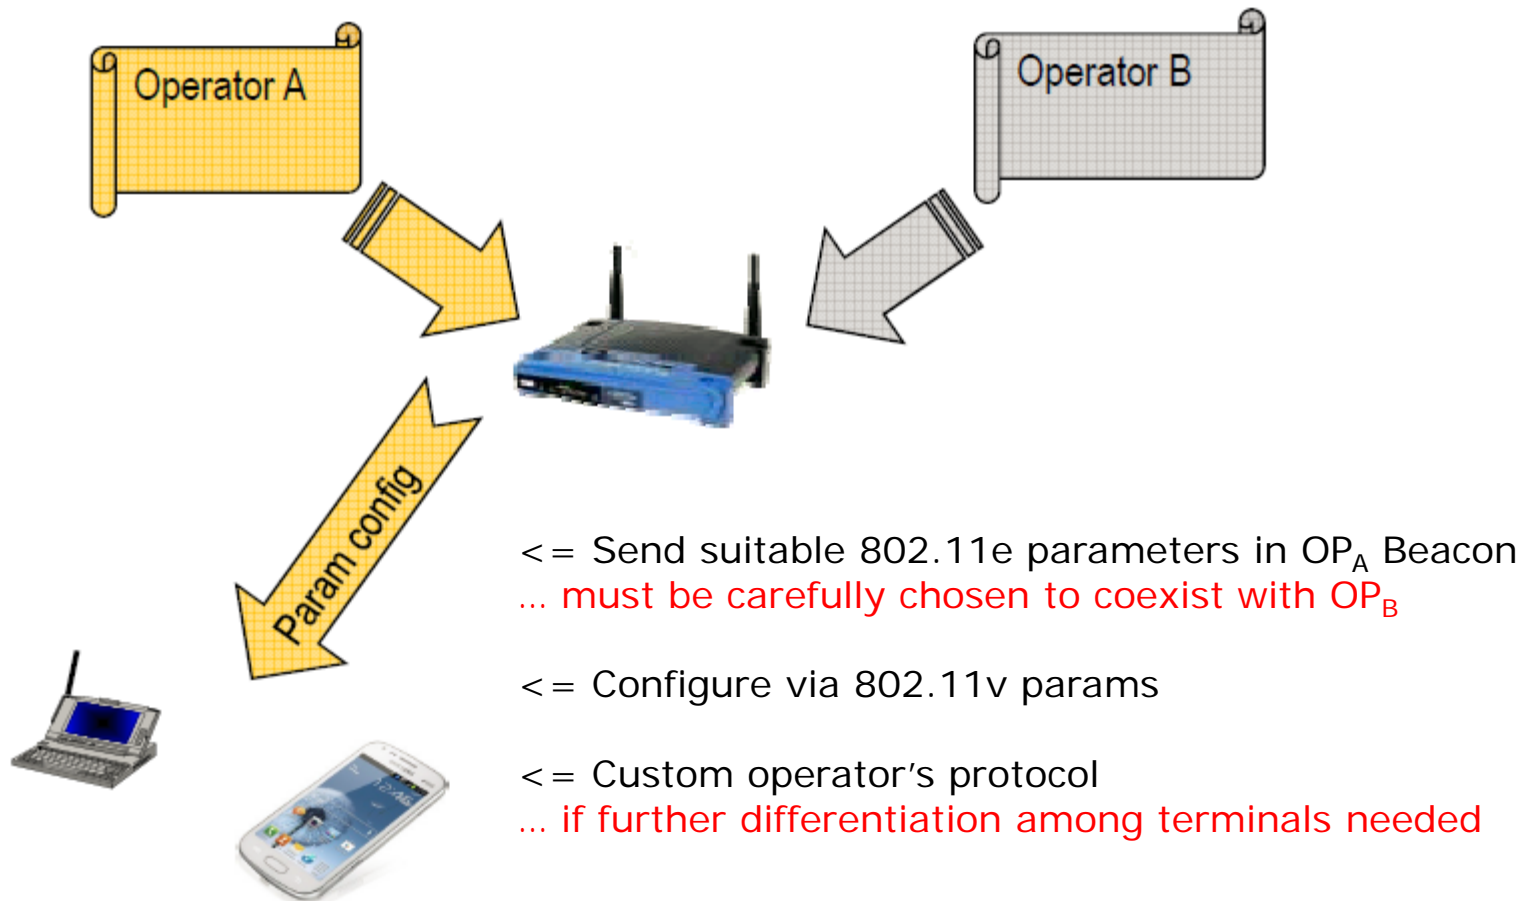
- ❖ Dynamic spectrum access
- ❖ Cognitive
- ❖ Performance optimization in niche environments
 - Home, industrial, ...
 - Adaptation to specific context or applications
- ❖ Improved support for new PHY
- ❖ Virtualization and access network sharing
- ❖ And many more...

A basic (but compelling) use case: multi-tenant WLAN sharing

- ❖ Virtual Operators over shared WLAN infrastructure
 - e.g., airport, hotel, enterprise, etc.



Well, we might «hack» this



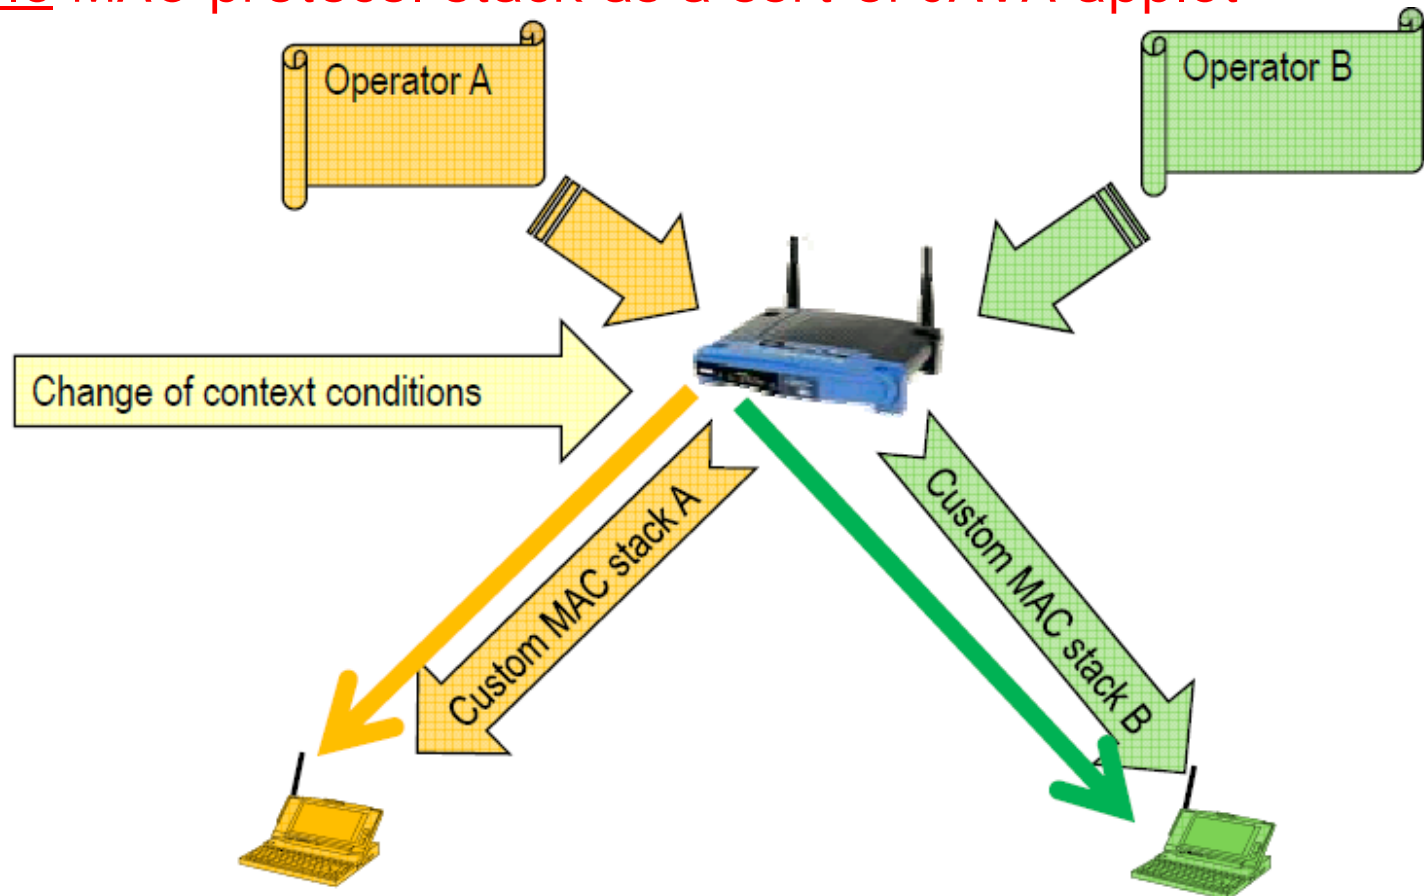
The point is another

- ❖ All-in-one MAC protocol, e.g. 802.11
 - We can probably stretch it to fit our context
 - Creative parameter configs, overlay tricks, ...
 - We are good at mastering complexity
 - and brings to accepted papers
 - When impossible? Just promote an amendment!

- ❖ But what if... we could change the MAC protocol for each and every context?
 - And we could trivially program our MAC operation?
 - Much simpler!
 - No anymore amendments, unless HW changes

Vision: Software-Defined MAC...

- ❖ Whole MAC protocol stack as a sort of JAVA applet

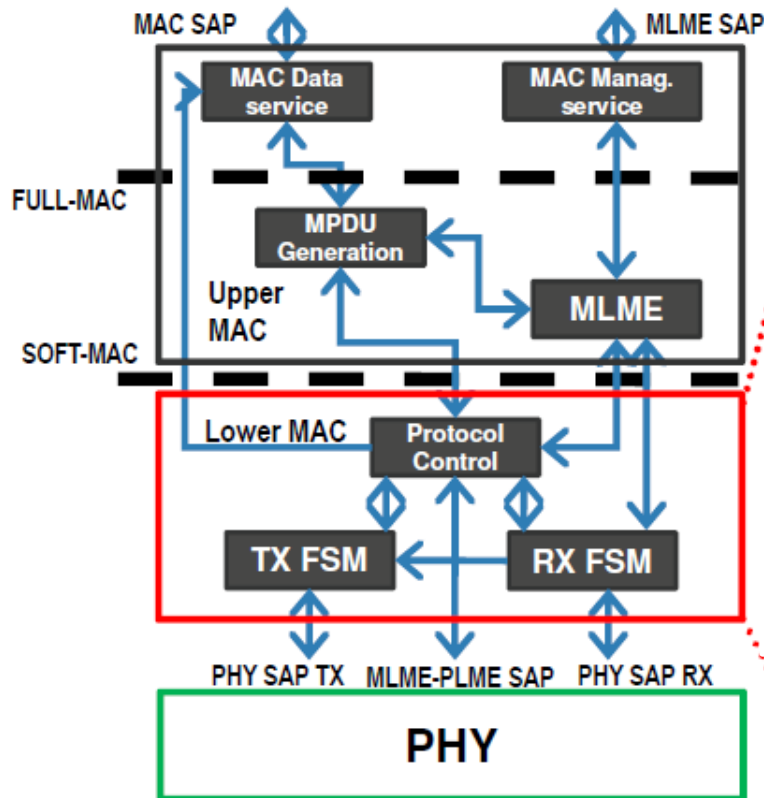


... but...

- ❖ **Best approach:** “persuade commercial name-brand equipment vendors to provide **an open programmable platform on their Wireless NICs**”
 - Plainly speaking: *let me hack your NIC!! No way...*
- ❖ **Viable approach:** “compromise on generality and seek a degree of **Wireless NIC** flexibility that is
 - High performance and low cost
 - Capable of supporting a broad range of research
 - **Consistent with vendors’ need for closed platforms.**
- ❖ **Compromise in Wireless MAC cannot be just a rule-action table!**

Current SW coding is wrong answer

❖ [even assuming Boxes are opened]



❖ DSP/FPGA SDR boards

- Cost, performance: just for research
- «open» box approach: must convince vendors

❖ Open firmware

- Probably only openFWWF, sneaked out...
- not “much” (!) vendor support

❖ BUT in both cases...

- Huge skills/experience, low level languages, inter-module dependencies
 - Assembly, VHDL, low level C, ...
- Complexity! Slow deployment time

Right answer

❖ Find the right abstractions!

- Must yield simple programming models
- Must not impair performance
- Sufficient flexibility to support most customization needs
- **Must be «vendor-friendly» 😊**

❖ Another attempt at this:

- Wireless MAC processor: Computing environment and abstractions for programming MAC protocols
- MAClets: from offline programming to online, dynamic, MAC stack injection and ultra fast reconfiguration, $< < 1$ micro second

Learn from computing systems?



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- ❖ 1: Instruction sets: perform elementary tasks on the platform
 - A-priori given by the platform
 - Can be VERY rich in special purpose computing platforms
 - Crypto accelerators, GPUs, DSPs, etc.
- ❖ 2: Programming model: actions + conditions
 - Convey desired platform algorithm
- ❖ 3: Central Processing Unit (CPU): execute program over the platform
 - Unaware of what the program specifically does
 - Fetch/invoke instructions, update registers, etc.

Let's MIMIC
all this!

Clear decoupling between:

- platform's vendor => implements (closed source!) instruction set & CPU
- programmer => produces SW code in given language

1: Which elementary MAC tasks? ("our" instruction set!)

❖ ACTIONS

➤ frame management, radio control, time scheduling

- TX frame, set PHY params, RX frame, set timer, freeze counter, build header, forge frame, switch channel, etc

❖ EVENTS

➤ available HW/SW signals/interrupts

- Busy channel signal, RX indication, inqueued frame, end timer, etc

❖ CONDITIONS

➤ boolean/arithmetic tests on available registers/info

- Frame address == X, queue length > 0,
- ACK received, power level < P, etc

Actually implemented API

- ❖ Platform: Broadcom Airforce54g commodity card

<i>events</i>	<i>actions</i>	<i>conditions</i>
CH_UP CH_DOWN	set/get(reg, value) switch RX()	dstaddr mvaddr

Just “one” possible API
convenient on our commodity platform

“others” possible as well
improved/extended
tailored to more capable radio HW

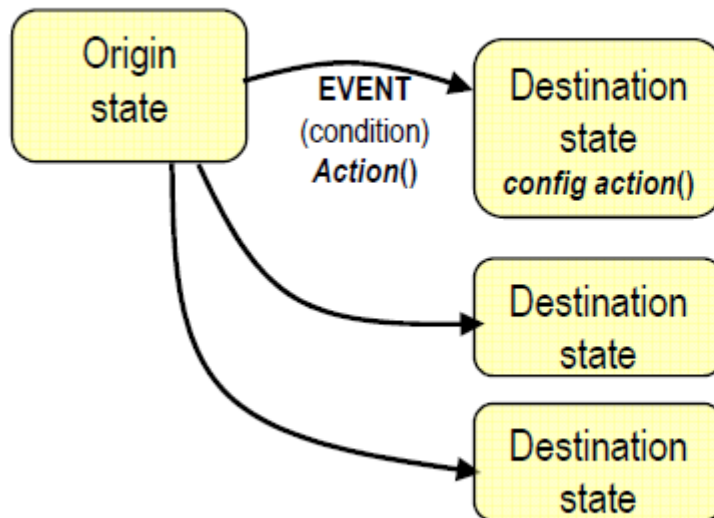
Our point:
have a specified set of actions/events/conditions,
not “which” specific one)

2: How to compose MAC tasks? ("our" programming language!)

❖ Convenient "language": XFSM

eXtended Finite State Machines

- Compact way for composing available acts/ev/cond to form a custom MAC protocol logic

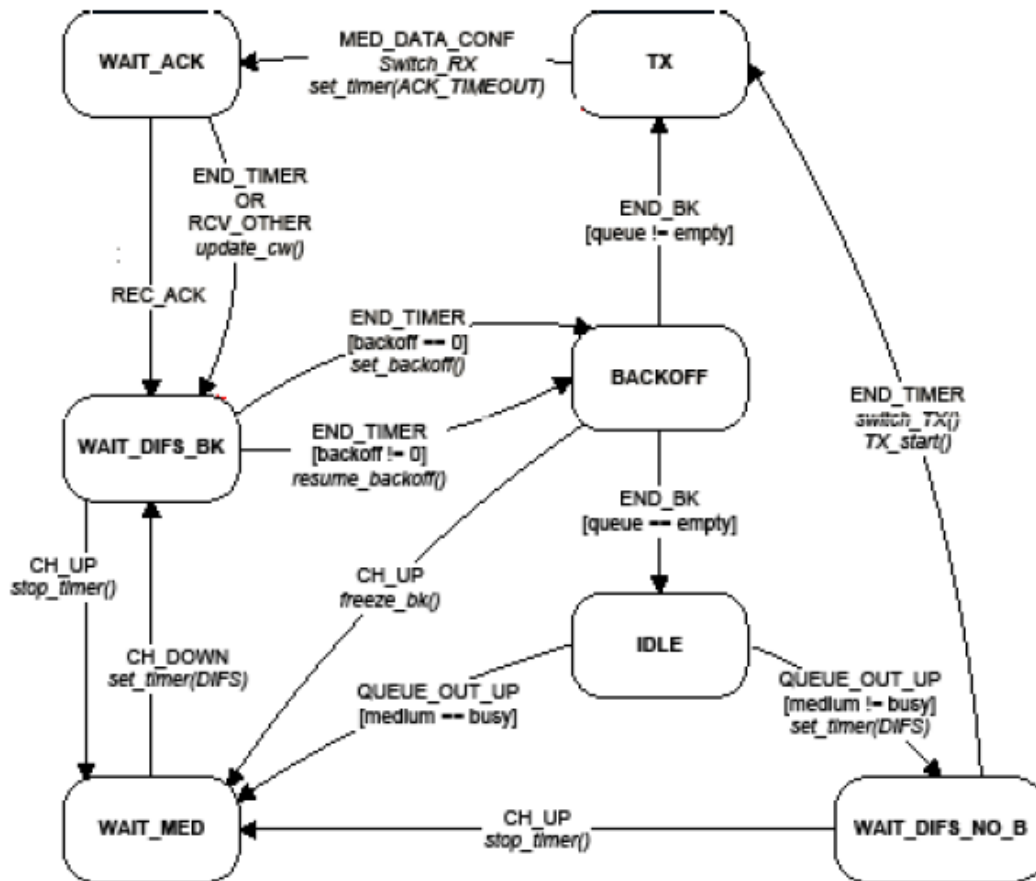


XFSM formal notation	meaning
S	symbolic states MAC protocol states
I	input symbols Events
O	output symbols MAC actions
D	n-dimensional linear space $D_1 \times \dots \times D_n$ all possible settings of n configuration registers
F	set of enabling functions $f_i : D \rightarrow \{0, 1\}$ Conditions to be verified on the configuration registers
U	set of update functions $u_i : D \rightarrow D$ Configuration commands, update registers' content
T	transition relation $T : S \times F \times I \rightarrow S \times U \times O$ Target state, actions and configuration commands associated to each transition

XFSM example: legacy DCF



❖ simplified for graphical convenience



Actions:

set_timer, stop_timer,
set_backoff,
resume_backoff,
update_cw,
switch_TX, TX_start

Events:

END_TIMER,
QUEUE_OUT_UP,
CH_DOWN, CH_UP,
END_BK,
MED_DATA_CONF

Conditions:

medium, backoff,
queue

3: How to run a MAC program? (MAC engine – XFSM onboard executor - our CPU!)



- ❖ MAC engine: specialized XFSM executor
(unaware of MAC logic)
 - Fetch state
 - Receive events
 - Verify conditions
 - Perform actions and state transition


- ❖ Once-for-all “vendor”-implemented in NIC
(no need for open source)
 - “close” to radio resources = straightforward real- time handling

MAC Programs

❖ MAC description:

➤ XFSM

❖ XFSM => tables



❖ Transitions

- «byte»-code event, condition, action
 - Portable over different vendors' devices, as long as API is the same!!
- Pack & optimize in WMP «machine-language» bytecode

A	T(A,B)
B	T(B,C)
C	T(C,A) T(C,B)

MAC protocol specification:
XFSM design
(e.g. Eclipse GMF)

Machine-readable code

Custom language compiler

**Code injection
in radio HW platform**

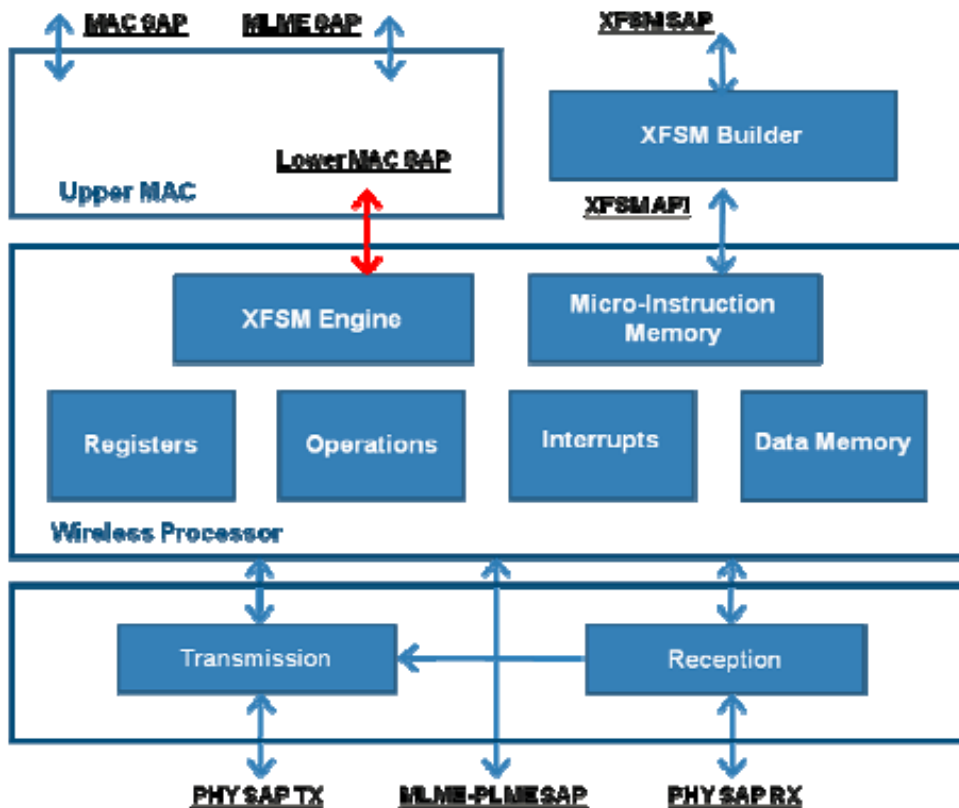
**MAC Bytecode
MAC Engine**



Machine Language Example (DCF, 544 bytes)

Memory address	Memory								Description
	<i>Initial State Descriptor</i>								
0x0BC0:	0100	FFFF	0B00	0014	A5FF	6ADA	0014	A5FF	
0x0BD0:	6ADA	6C00	80A4	FF00	FF00	3600	80EE	FF00	
0x0BE0:	0000	0000	0000	0000	0000	0000	0000	0000	
	00	01	02	03	<i>Coded state machine</i>				
0x0C00:	0100	0100	0100	0401	0108	0508	1C01	010B	Outgoing transitions for state 01 0401 0108 0508 = trans. 1 1C01 010B 010B = trans. 2 3001 010D 0200 = trans. 3 FFFF = delimiter
0x0C10:	010B	3001	010D	0200	FFFF	5101	010E	030D	
0x0C20:	0000	0100	010F	C100	0102	0602	E100	0106	
0x0C30:	0106	0401	0108	0508	1C01	010B	030B	FFFF	
0x0C40:	CD00	0104	0E0C	0000	0100	0D00	FFFF	0E01	
0x0C50:	0109	0909	1C01	010B	0D0B	FFFF	C700	0103	Transition 1 0401 = event pointer 01 = event parameter 08 = event index 05 = target state 08 = action
0x0C60:	0C03	E100	0106	0106	FFFF	6601	0110	1600	
0x0C70:	0000	0100	0100	FFFF	0E01	0109	0109	1C01	
0x0C80:	010B	010B	FFFF	5F01	010F	0A00	0000	0100	
0x0C90:	0D00	FFFF	C100	0102	0A02	C700	0103	0B03	
0x0CA0:	E100	0106	0D06	FFFF	D300	0105	0D05	E100	State 01 03 = transitions offset (9 bits) E = FFFF delimiter
0x0CB0:	0106	0D06	FFFF	D300	0105	0705	E100	0106	
0x0CC0:	0106	FFFF	6D01	0111	1800	0000	0100	0100	
0x0CD0:	0000	0100	0D10	7401	0112	1512	0000	0100	
0x0CE0:	1111	9601	0113	0513	0000	0100	0500	0000	
0x0CF0:	0100	0304	E100	0106	1206	0401	0108	0508	
0x0D00:	1C01	010B	120B	FFFF	A901	0115	0100	B401	
0x0D10:	0117	1200	0000	0100	0100	0000	0100	1214	
0x0D20:	B901	0118	0310	0000	0100	0300	0401	0108	
0x0D30:	1708	1501	010A	010A	1C01	010B	010B	C501	
0x0D40:	0119	0800	0000	0100	0500	3001	010D	0200	
0x0D50:	0401	0108	0508	1C01	010B	180B	CB01	011A	
0x0D60:	0200	0000	0100	0100	0000	0000	0000	0000	
0x0D70:	0000	0000	0000	0000	0000	0000	0000	0000	
0x0D80:	0000	0000	0000	0000	0000	0000	0000	0000	
	00	01							
0x0D90:	00F0	03FE	0DF2	13FE	20FE	27FE	2EFE	35FE	
0x0DA0:	0000	0000	0000	0000	0000	0000	0000	0000	

Wireless MAC Processor: Overall architecture



- ❖ MAC Engine: XFSM executor
- ❖ Memory blocks: data, prog
- ❖ Registers: save system state (conditions);
- ❖ Interrupts block passing HW signals to Engine (events);
- ❖ Operations invoked by the engine for driving the hardware (actions)

From MAC Programs to MAClets

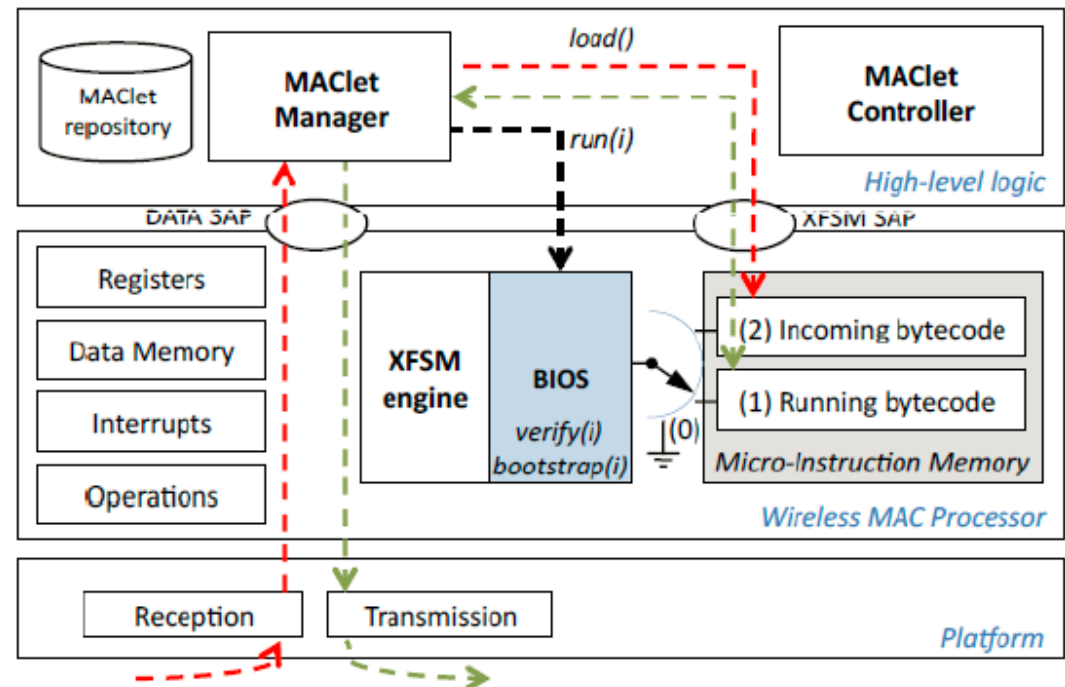
- ❖ Upload MAC program on NIC from remote
 - While another MAC is running
 - Embed code in ordinary packets

- ❖ WMP Control Primitives

- load(XFSM)
- run(XFSM)
- verify(XFSM)
- switch(XFSM1, XFSM2, ev, cond)

- ❖ Further primitives

- Synchro support for distributed start of same MAC operation
- Distribution protocol



“Bios” state machine: DEFAULT protocol (e.g. wifi) which all terminals understand

From theory to practice

- ❖ Obviously, instruction set and MAC Engine can be “easily” implemented in a software-defined radio...
 - e.g., FPGA, WARP, ...
- ❖ But... can this be done on commodity HW?
 - e.g., ultra-cheap ordinary WLAN NIC
- ❖ Yes!!!
 - Reference platform: broadcom Airforce54g 4311/4318
 - Hands-on experience on card's assembly language FW
 - general purpose processor (88 MHz), 64 registers, 4KB data memory, 32 KB code memory
 - Partly leveraging existing card HW facilities
 - HW configuration registers for radio resource and event handling
 - Frequency, power, channel sensing, frame forging facilities, etc
 - Available HW events (packet queued, plcp end, rx end, rx correct frame, crc failure, timer expiration, carrier sense, etc)

Implementation at a glance

- ❖ Delete 802.11 firmware
 - Both Broadcom and openFWWF
 - we do NOT want yet another firmware MAC to hack!
- ❖ Replace it with *[once for all developed]*:
 - Implementation of actions, events, conditions
 - in part reusing existing HW facilities
 - MAC engine: XFSM executor
- ❖ Develop “machine language” for MAC engine
 - Custom made “bytecode” specified and implemented
 - 6 bytes instructions, state transition table (sparseness exploited)
- ❖ Address several annoying technical hurdles
 - NO direct HW interrupts control available in Broadcom
 - State and state transition optimizations, ...

Public-domain

- ❖ Supported by the FLAVIA EU FP7 project
 - <http://www.ict-flavia.eu/>
- ❖ Ongoing integration in the CREW EU FP7 federated testbed
 - <http://www.ict-flavia.eu/>
- ❖ Public domain release
 - Project page: <http://wmp.tti.unipa.it>
 - Download: <https://github.com/ict-flavia/Wireless-MAC-Processor>
- ❖ Released distribution:
 - Binary image for WMP
 - You DO NOT need it open source!
 - Remember the “hard-coded” device philosophy...
 - Conveniently mounted and run on Linksis or Alix
 - Source code for everything else
 - Manual & documentation, sample programs




WMP Functional validation «static» MAC programs

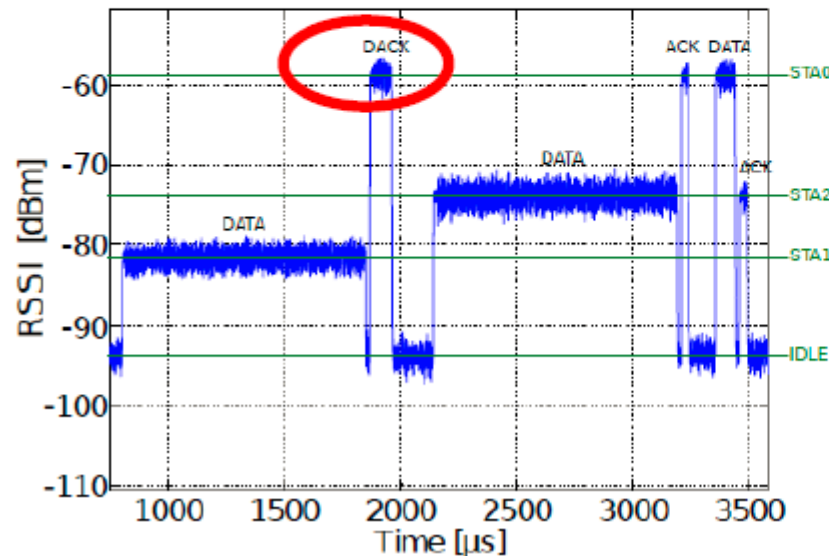
- ❖ Success IF WMP permits very easy/fast Lower MAC modifications or re-design (vs months or hands-on experience with openFWWF/assembly)
- ❖ “scientifically trivial” use cases, tackling distinct MAC aspects recurring in literature proposals
 - Piggybacked ACK
 - Programmable management of frame replies
 - Pseudo-TDMA
 - Precise scheduling of the medium access times
 - Randomized multi-channel access
 - Fine-grained radio channels control
 - Multi-tenant access network sharing, with different protocols
 - Virtualization
- ❖ Development time: $O(\text{days})$
 - Including bug fixing in engine/API, otherwise hours

Piggybacked ACK


- ❖ If available, send TCP ACK instead of MAC ACK, otherwise send normal ACK



Channel activity trace



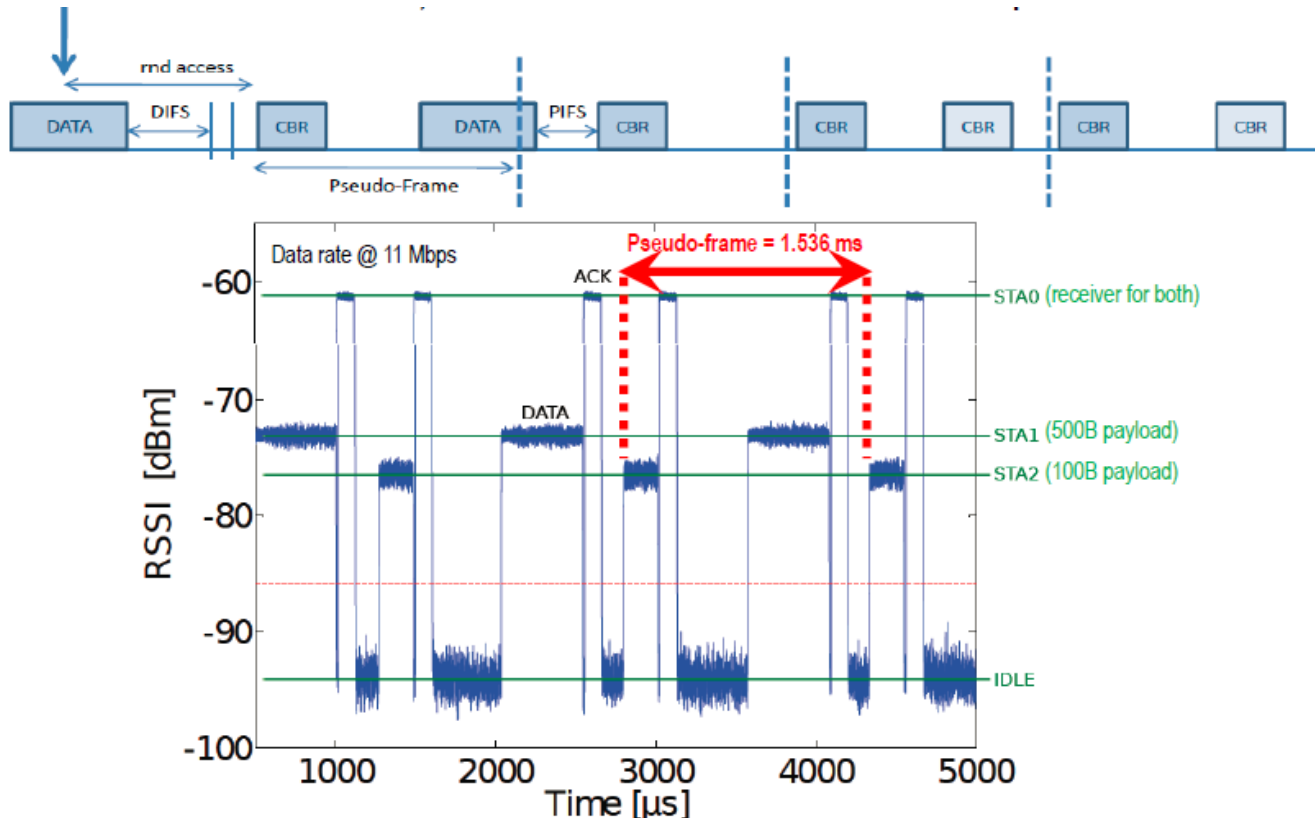
Performance gain



Pseudo-TDMA


[literature proposal]

- ❖ After first random access, schedule next transmissions at fixed temporal intervals

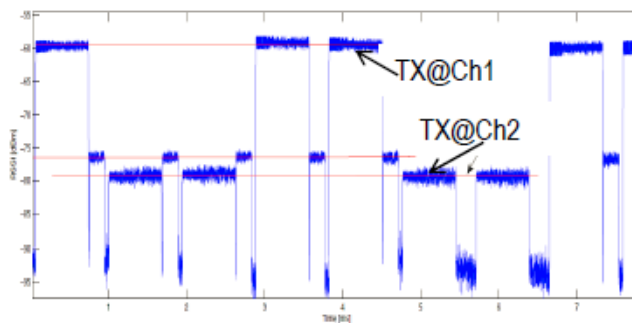


Randomized multichannel access


- ❖ Per EACH frame, randomly select backoff AND channel (switch on as little as per frame basis)



Channel activity trace

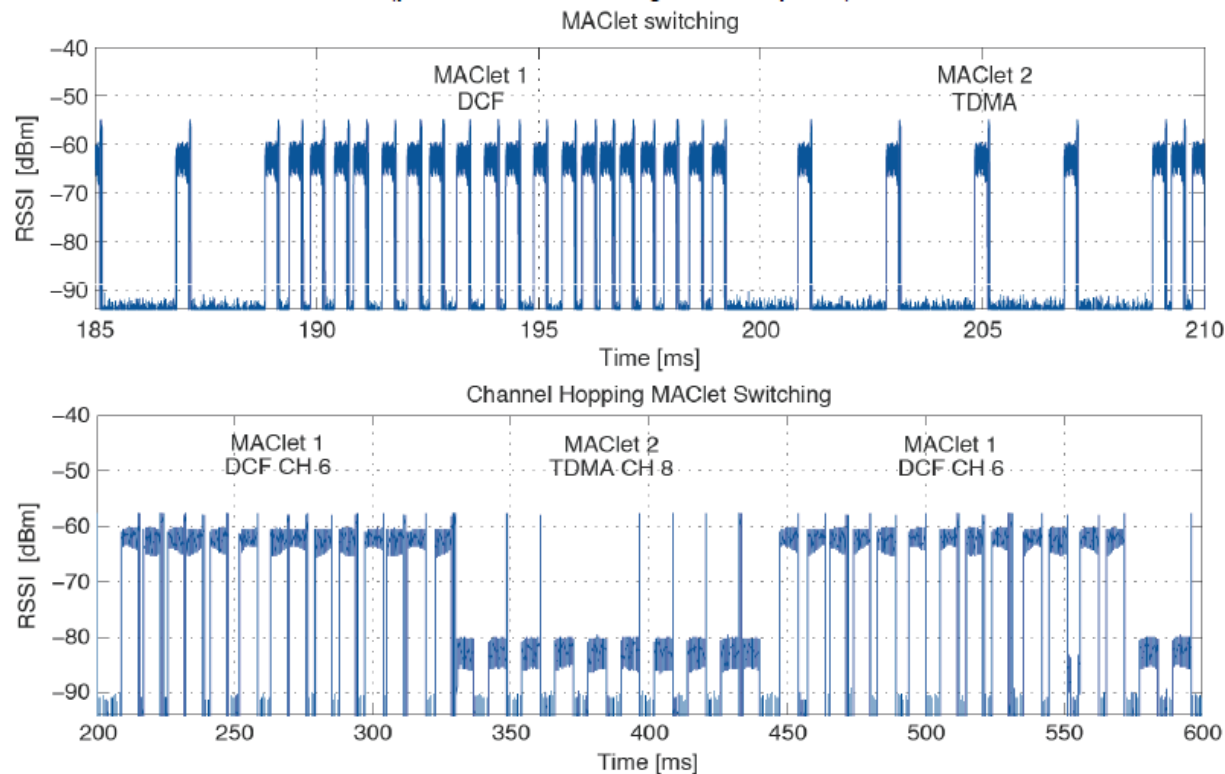


Performance gain



Multi-threaded MAC

- ❖ Success IF seamless switch from one MAC program to another in negligible time
 - Result: less than 0.2 μ s over such cheap hardware!
 - (plus channel switching time if required)



AP Virtualization with MAClets

- ❖ Two operators on same AP/infrastructure


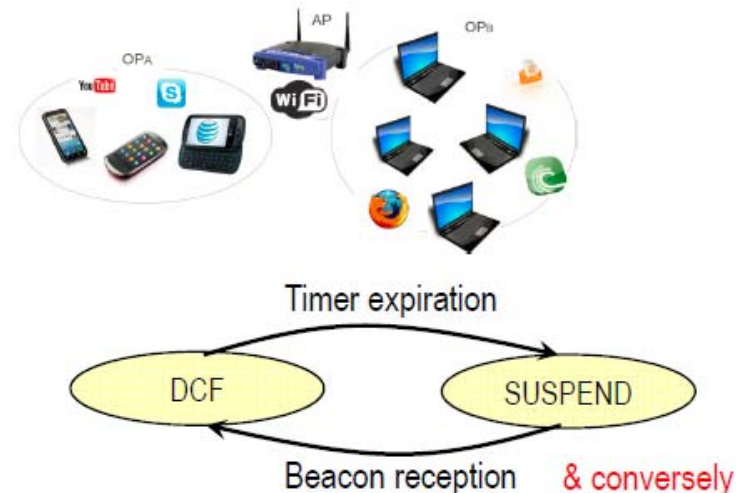
- A: wants TDM, fixed rate
- B: wants best effort DCF

- ❖ Trivial with MAClets!

- Customers of A/B download respective TDM/DCF MAClets!

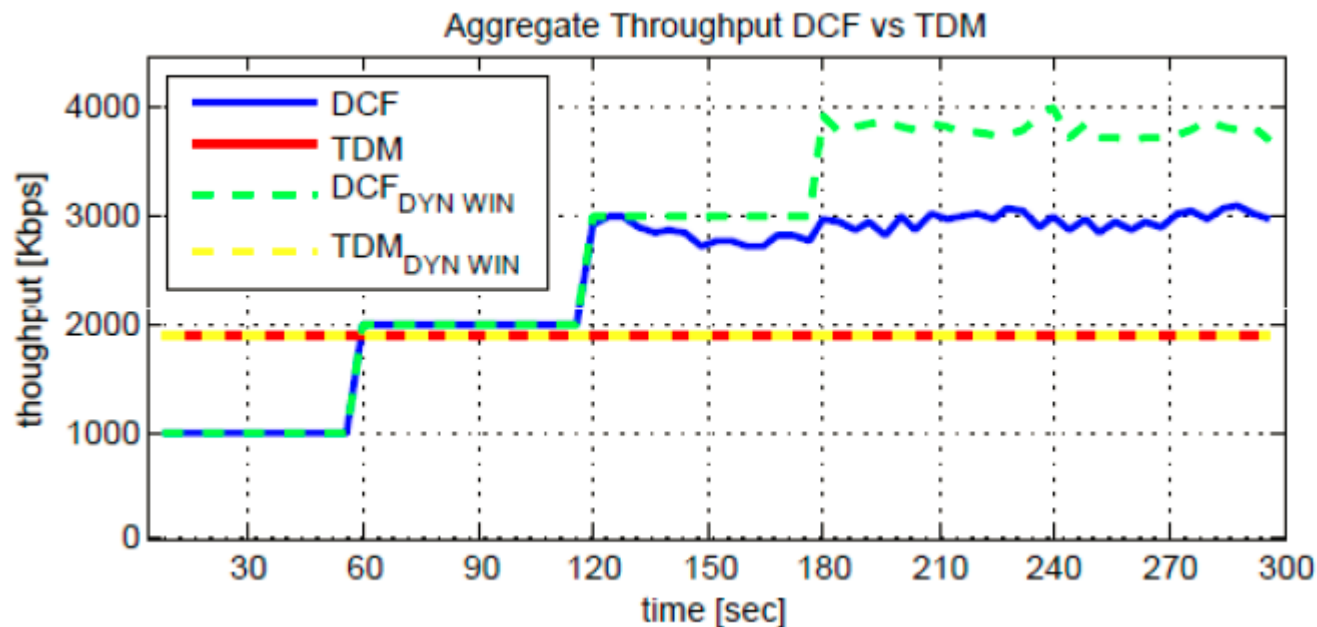
- ❖ Isolation via MAClet design

- Time slicing DESIGNED INTO the MAClets! (static or dynamic)



An Example of Throughput Performance

- ❖ 3 FIXED stations @ 0.63 Mbps vs. 5 BEST stations @ 1Mbps



What about control?

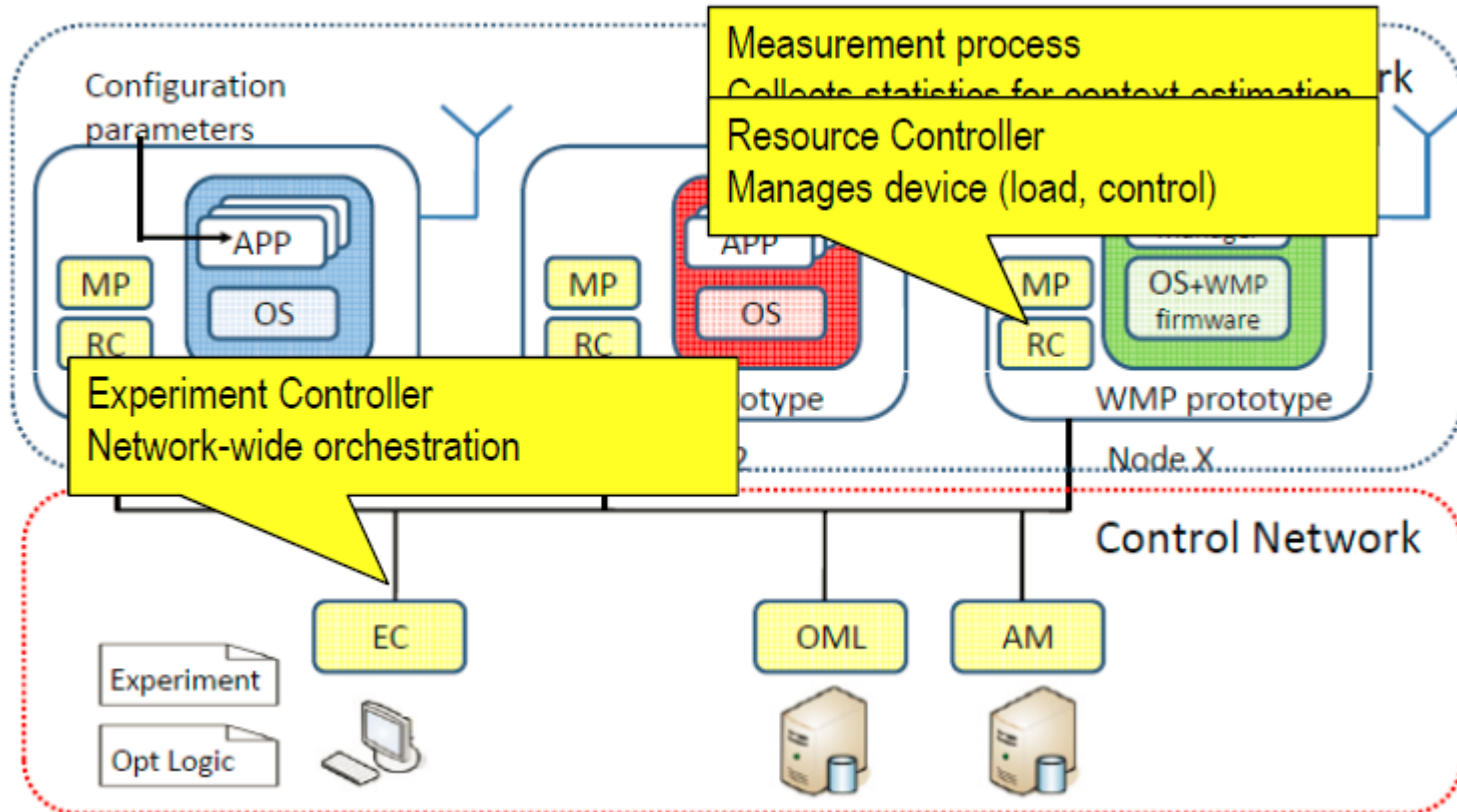
❖ Context-dependent wireless operation

- Gather context information (measurements)
- Adapt protocol logic accordingly

❖ Preliminary insights

- No need to reinvent the wheel!
- WMP ported on OMF testbed control framework
 - Different perspective, but same elementary building blocks

Deployed architecture



Control primitives

available on CREW remote testbed

- ❖ Loading messages
 - Upload bytecode programs

- ❖ Call messages
 - Trigger actions on devices

- ❖ Event messages
 - Notifications from devices to Controller

- ❖ Management messages
 - Reports on node status and configuration

«network» programs



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- ❖ Scripting language OEDL (ruby-like)
- ❖ Very simple & intuitive

Sensing Configuration

```
defGroup('N1', property.res1 ) do |g|  
  g.addApplication('WMPdump') do |app|  
    app.property('cw')  
  end  
end
```

Context Definition

```
defEvent(:HIGH_CW) do |event|  
  app_status =  
  group('N1').state("apps/app[@name='WMPdump']/io/out/line")  
  if !app_status.nil?  
    app_status.each do |element|  
      event.fire if element>31  
    end  
  end  
end
```

Reaction Mechanism

```
onEvent(:HIGH_CW) do |event|  
  group('N1').stopApplication('WMPdump')  
  group('N1').addApplication('bytecode-manager') do |app|  
    app.setProperty('run', 2)  
  end  
  group('N1').startApplication('bytecode-manager')  
end
```

Scenario

```
onEvent(:ALL_UP_AND_INSTALLED) do |event|  
  group('N3').exec('iperf -s')  
  wait 1  
  group('N2').exec('iperf -c 192.168.0.3 -t 100 -i 1')  
  group('N1').exec('iperf -c 192.168.0.3 -t 100 -i 1')  
  wait 1  
  group('N1').startApplications  
  wait 30  
  Experiment.done  
end
```

- ❖ Device-level programming abstraction
 - Consistent with vendors' need for closed platforms

- ❖ Context-specific protocol adaptation!
 - Wireless access: no more one-size-fits-all
 - Beyond the need for standards?

- ❖ From node to network
 - Preliminary control framework deployment
 - Towards cognitive reconfiguration of whole access logic!

OpenFlow versus SDN in Wireless

OpenFlow	SDN in Wireless
Switches (e.g. NEC PFLOW PF5240F)	Wireless NICs (e.g. Broadcom Airforce54g)
High performance and low cost Capable of supporting a broad range of research Consistent with vendors' need for closed platforms => Vendor-implemented actions (closed source)	
SDN Controller	XFSM Builder
Flow Entries	State Transitions
Packets	Events (e.g. busy channel, end timer)
Matching Rule (e.g. port == 80)	Conditions (e.g. frame address == X)
Action (e.g. drop packet)	Action (e.g. switch channel, set timer)
FlowVisor	MAClets

References

- ❖ Open Networking Foundation: OpenFlow-Enabled Mobile and Wireless Networks, ONF Solution Brief, September 30, 2013.
<https://www.opennetworking.org/images/stories/downloads/sdn-resources/solution-briefs/sb-wireless-mobile.pdf>