



Telecooperation Lab
Prof. Dr. Max Mühlhäuser

TK1: Distributed Systems - Programming & Algorithms

Chapter 1: Introduction

Lecturer: Dr. Benedikt Schmidt, Dr. Immanuel Schweizer,
Prof. Dr. Max Mühlhäuser

Copyrighted material – for TUD student use only



Computer Networks: Recall NCS

Computer Network := interconnected collection of autonomous computers

$$\text{CN} := \{\text{AS}\} \cup \text{CSS}$$

1. **AS:** Autonomous System (node) := pair (CPU, local memory) [+ net + 'addressability']

- 1 AS may be multiprocessor (shared memory)
- Note: AS able to run stand-alone
- Note: a shoebox may contain a distributed system!

2. **CSS:** Communication Subsystem („the network“)
:= „whatever enables the ASs to exchange messages“

- Wired/wireless, arbitrary topology (ring, star, bus, [partly] meshed,...)
- Note: AS may be part of CSS (Host & Router)



For Ubiquitous Computing:

- Consider 4-tuple from above: (address \approx identity, communication, CPU, memory)
- Ubiquitous computing: your shirt, suitcase may be an Internet node !
- But then: minimalistic node may consist of: (passive?) communication & identity \rightarrow *not* an AS, no 4-tuple
- This definition of AS is not used here any further, reserved for TK3

Computer networks are everywhere: LAN/Intranet, Internet, in-car, on-body, ...



A. “Definition” of Distributed Systems

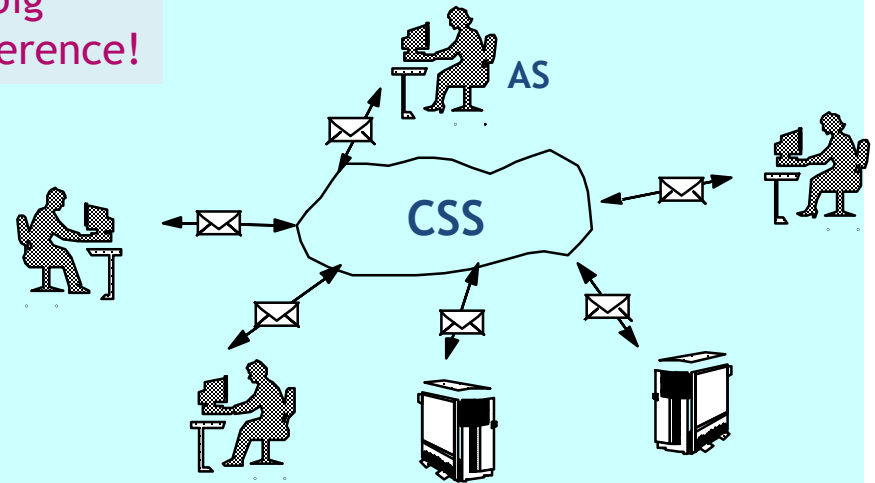
„A distributed computing system consists of multiple autonomous processors that do not share primary memory, but cooperate by *sending messages* over a communication network.“

-- H. Bal

„A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.“

-- L. Lamport

no big
difference!



Summary of many definitions: In a **Distributed System ('DistSys')**, networked computers

- communicate and coordinate their actions **only by passing messages**
- may be spatially separated by **any distance**

Definitions suck! So, what's the **distinction** *Computer Network* \leftrightarrow *Distributed System*?
well... not precisely defined either, but:

Distributed Systems* establish some level of transparency atop *Computer Networks
... of locations, distribution, concurrency, performance ... (details further below)



B. Basic Problems of DistSys



Basic problems in Distributed Systems form a dangerous *Bermuda Triangle*

Basic Problem #1: Global State Not Accessible (without unacceptable slowdown)

- no synchronized global variables, no global shared memory
- message / agent travelling $A \rightarrow B$: out-dated state of A arriving at B

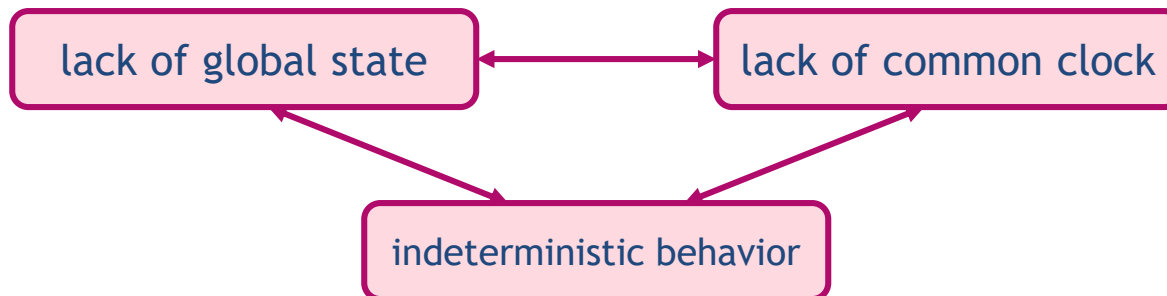
Basic Problem #2: Clocks Not 100% synchronized

Events E_A, E_B at A, B with recorded times $t(E_A) < t(E_B)$:

- May have happened at $t(E_A) > t(E_B)$!!
- When is it safe to „believe“ $t(E_A) < t(E_B)$?
- How to find out which is true? if undecidable: does distinction matter?

Basic Problem #3: Indeterminism – multiple execution of same system may yield different results

- Race conditions (messages from different senders, different threads) ...
- ... plus ‘erroneous’ underlying computer network \rightarrow „correct program“ has unpredictable result!!





C. Requirements – selected (1)

1. Heterogeneity Support

Integration of different vendors → different ...

- Networks, computer hardware, operating systems, programming languages
- Sub-requirement: **support exchange of data & code despite heterogeneity**
 - **Data:** in different HW architecture, OS, programming languages:
 - Known as „presentation“ issue
 - XML is not the solution (only syntax), may be *used* to build one
 - **Code:** mobile objects, mobile code

2. Openness

„Anyone may come along and participate“

- Remember, for computer networks: Open Systems Interconnection OSI
 - Standardized comm. Protocols → different implementations interoperate → global Internet
- For DistSys? Build application such that foreign processes may participate?
 - e.g., Reflection: ... by means of application that „adapts itself“
 - e.g., Evolution: ... supporting „new versions“ of modules/parts
 - Publish/Subscribe paradigm (see later): plug in new process at any time



C. Requirements – selected (2)

3. Scalability

- # of processes&nodes, # of users, # of transactions, ...

4. Security

- “open system” → malicious users may come in
- Authentication, authorization, trust

5. Failure tolerance

- How to handle AS or CN failures?

6. Concurrency

- Distributed applications are inherently “multithreaded”
- Concurrent access to shared resources -> ensure consistency

7. Transparency

- Abstraction of an aspect; many variants – see below




1. Client/Server

- Traditional model, easily comprehensible abstraction

- Clients request service (initiate connection)
- Servers provide service (answer requests)

- Examples:

- Web Client/Server
- Mail Client/Server (well...)
- FTP ... ; well ...
- these are all using P2P (interprocess comm.) technologies
- MAY be considered client-server IF: re-entrant behavior of server („many“ clients)
- true C/S technology: remote procedure call (RPC)






D. System Models (2)

2. Peer-to-Peer

- Brand new paradigm?
 - no, the *oldest* one
 - appears to be new due to
 - MP3 piracy → overlay network (DHTs etc.)
 - new idea: resource client := resource server
 - Consideration of entire net, not just two peers
- P2P-*new* successful since 2000's
 - First tool: Napster (file sharing)
 - route around censorship
 - Other services include
 - streaming media
 - distributed computing



Other cool P2P technologies: ;-)

- Telephone
- Usenet
- DNS
- IP Routing


Actually, P2P = original Internet model
(symmetric network - all hosts are
(ftp, telnet, ..) clients and servers)



E. Abstractions (1)


Level 1: physical configuration - seems irrelevant for DS:

- Object of SysOp people
- For DisProg people, should be abstracted from
- but:
 - ownership → cost (public net?), security, ... !
 - bandwidth etc. → performance
 - reliability?



Level 2: logical configuration - „the“ CompNet abstraction!

- CSS = „cloud“, classes of ASs; AS may be part of CSS
- sometimes, abstraction too high (see above)
- usually, abstraction too low for DistProg






E. Abstractions (2)

Level 3: process network (logical distribution): abstracts from real distribution

(\approx Distributed Program DistProg):


processes (objects, agents) exchange messages

- e.g., processes reached via mailboxes
- state distribution (no global view), see below
- no common time (no exact global “clock”), see below
- reliability, indeterminism assumed?
depends further on abstraction & underlying support



Level 4: Distributed Algorithm - abstracts from

- Target environment (\rightarrow reliability, performance, ...) not necessarily from reliability / performance issues
- Target process configuration i.e. # of processes (well, ought to...) not necessarily from interconnection / topology issues!
- Implementation language, platform, lifecycle





F. Transparency (1)

- Definition:

Transparency = „**Concealment**“ from the user and the application programmer of the **separation of components** in a distributed system, so that the **system is perceived as a whole** rather than as a collection of independent components.

- Note: „transparency“ is about „hiding something“ in English
- We introduce 8 forms of transparency below, but literature varies → there are more!

1. Access transparency

- Local & remote resources accessed using identical op's

2. Location transparency

- Resources accessed w/o knowledge of their physical/network location

3. Concurrency transparency

- Several processes operate concurrently using shared resources without interference between them



F. Transparency (2)

4. Replication transparency

- Multiple instances of resources used (→ reliability, performance) w/o knowledge of replicas by users & programmers

5. Failure transparency

- Concealment of faults, allowing users & programs to complete tasks despite HW/SW failure

6. Mobility transparency

- Movement of resources/clients w/o affecting users & programs

7. Performance transparency

- Local/remote op (exec, data access) don't differ by orders of magnitude (**most persistent problem!**)
- allows system to reconfigure to improve performance as loads vary

8. Scaling transparency

- Allows system/applications to expand in scale without change to system structure or application algorithms




G. Classification

1. According to xlxD classification, DistSys are classified as MIMD (Flynn's taxonomy)

- Like multiprocessor systems (but w/o shared memory → big difference!)
- remember: SISD = PC etc., SIMD = vector processor, MISD = <unused>

2. Federation classes: autonomous ASes access certain services (never 100% pure)

- **function** federation (use particular function of “other” system)
- **load** federation
- **data** federation
- **availability** federation: here, availability



- 1: pretty useless, 2: helpful for understanding, but: exemplars quasi nonexistent
- many more classifications exist: not very helpful either

Important lesson: **Myths & Reality & Future of DistSys**

- Dream was: DistSys for parallel processing
- Reality is (~90%): DistSys due to “locality of data / CPU”, central server
- Will central servers go away? ... no! rather, importance may decrease
- Will server farms “explode” soon? ... was envisioned since 20 years!

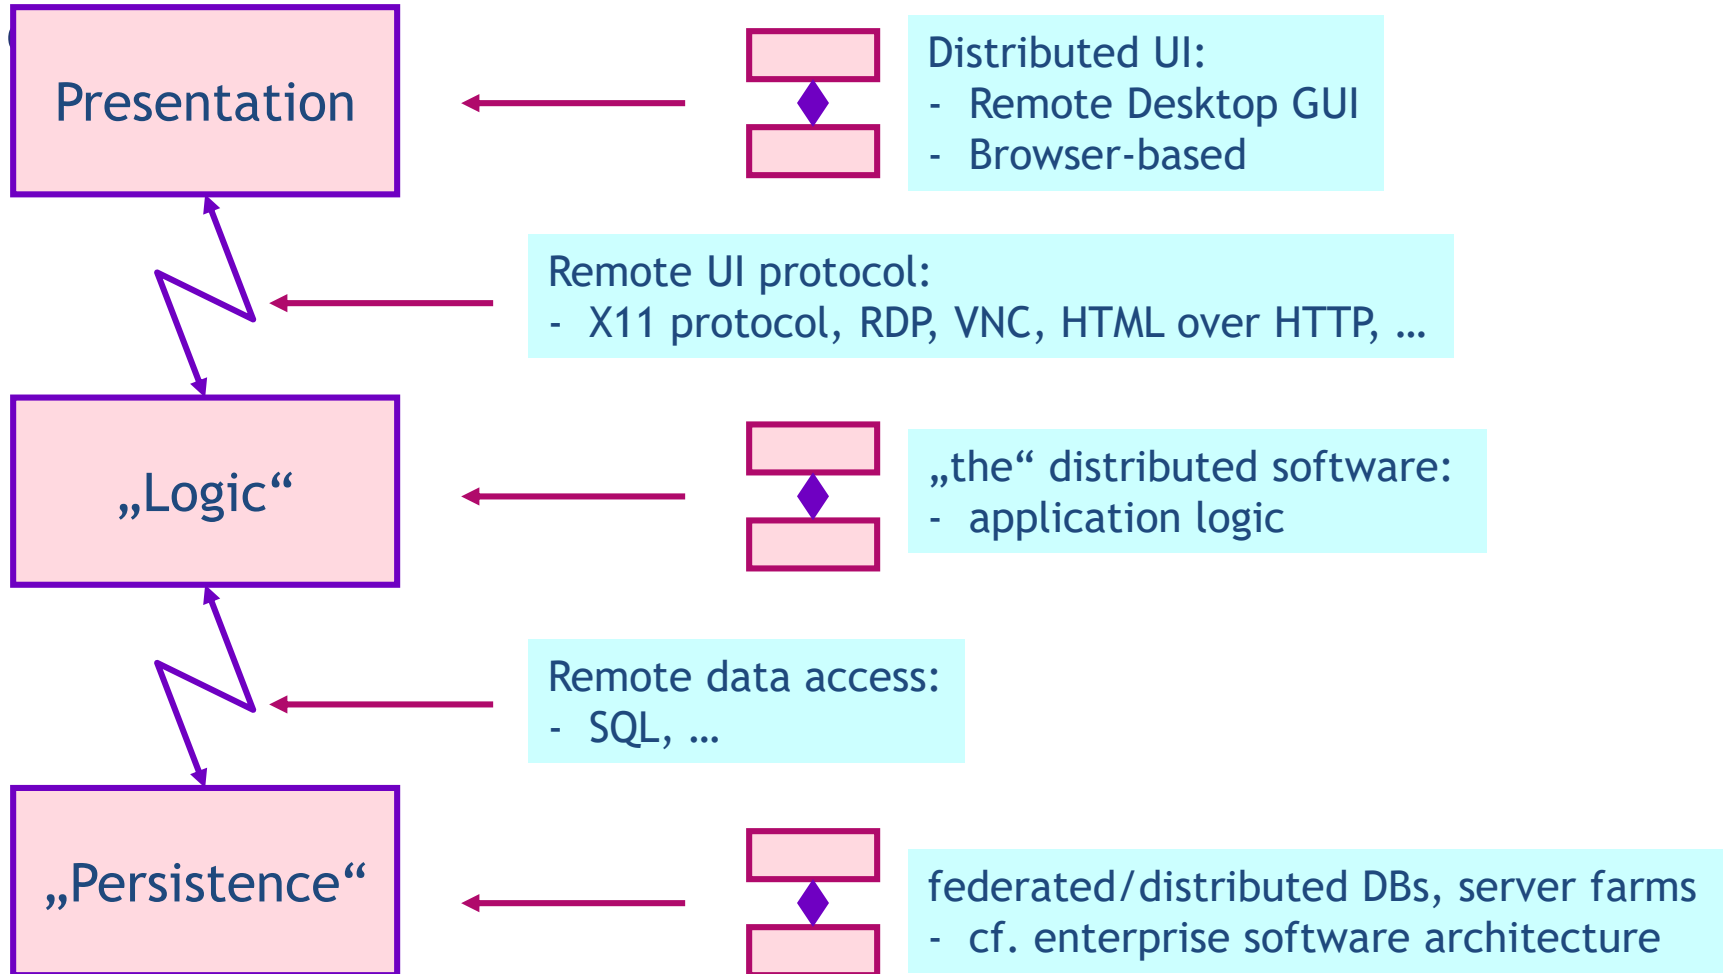


H. Tier Architectures (1)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Classical (logical) SW Architecture is 3-tiered → 5 distribution






H. Tier Architectures (2)

Note: physical tiers! (enterprise computing lingo, used by database people)

- **2-Tier:** early Client-Server

- tier 1: GUI & application
- net: file I/O, protocol (http, ...), SQL/ODBC, ...
- tier 2: DBMS, „Resources“



- **3-Tier:**


- different definitions exist, but mostly:
- Client – Application Server – Database Server

- **n-Tier:** 2nd Tier expanded

- specialized application servers (OrderProcessing, Accounting, Human..)
- cf. ERP (enterprise resource planning) companies like SAP, Oracle,...
- scalability considered to require „TP monitors“

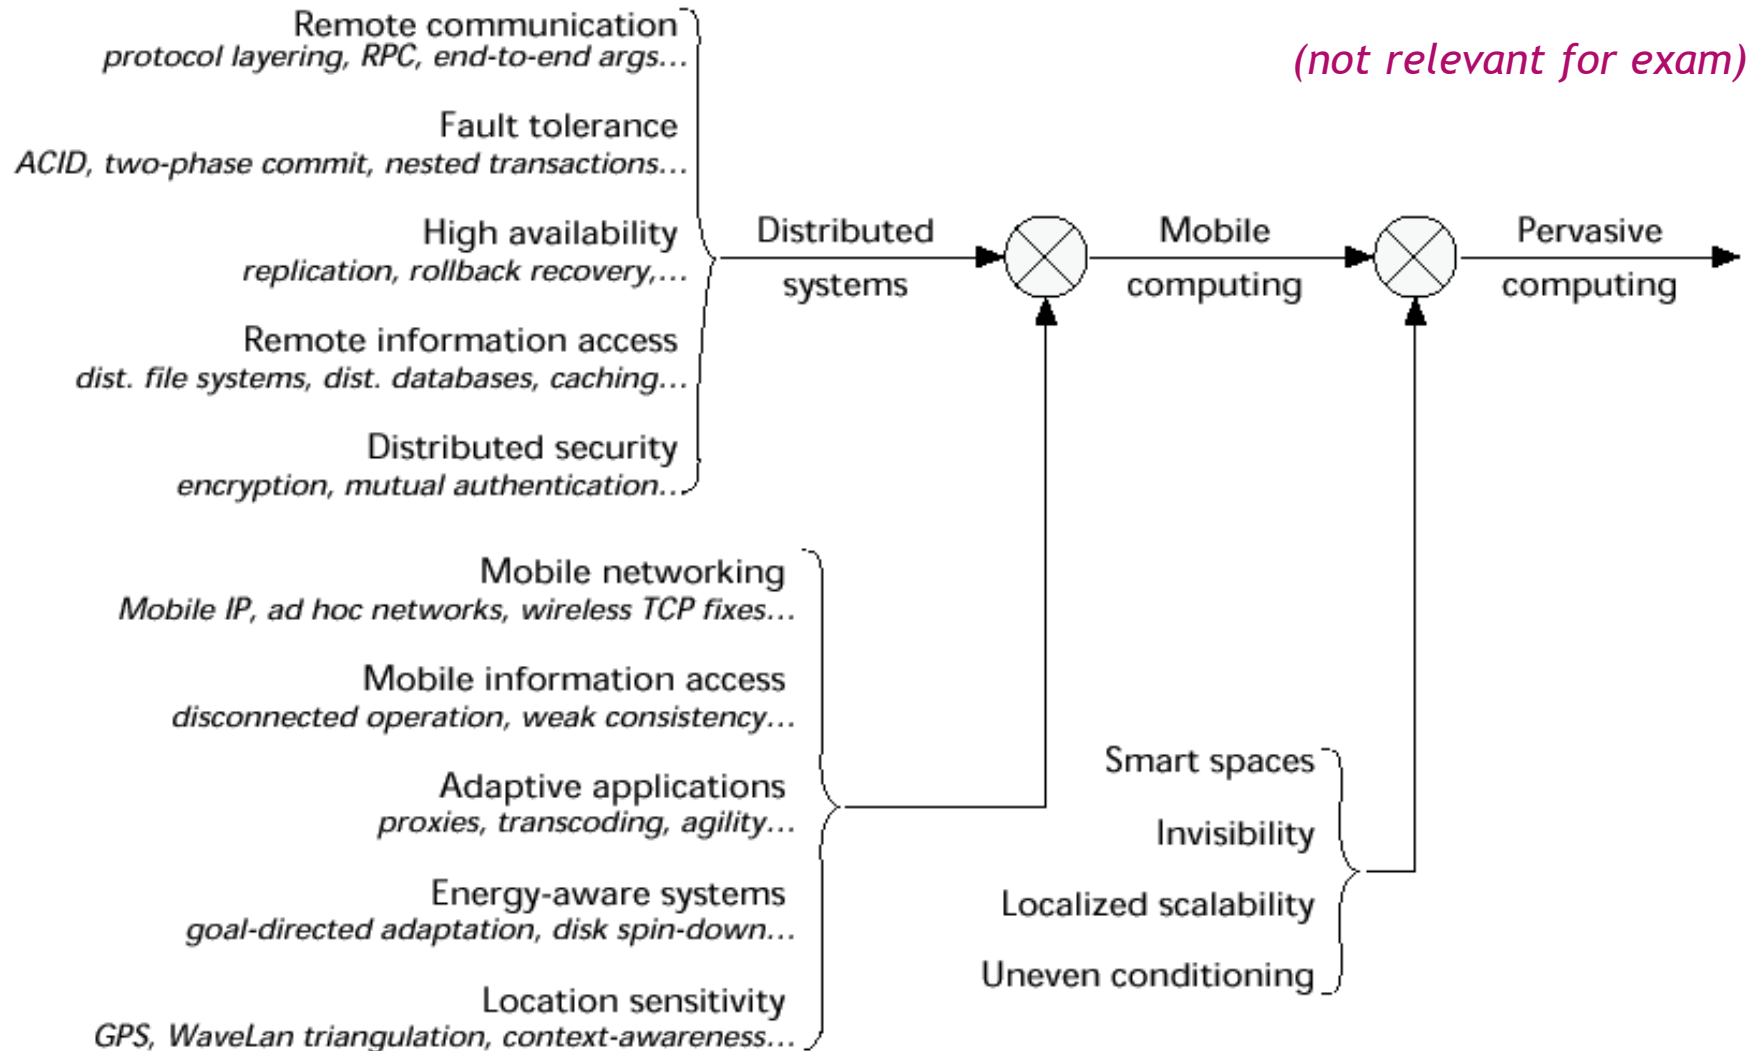
- note: each tier may be „n nodes“

- many clients anyway; server farms distribute load
- server-calls-server is rather common





I. Challenges





4 principles for abstractions wrt. Distributed Software Development

1. Distributed operating system approach

- Support for distributed programming is part of operating system
- Pro: Quite general solution, *parallel* programming paradigm
- Con: Needs wide-scale adoption of the same system
 - Large systems always heterogeneous

2. Distributed database approach

- Same as above, except OS replaced by a database system
- Pro: Allows for all DB features (semantics, ...), isolated *sequential* prog's
- Con: Independent applications with shared database
- Con: Many distributed algorithms hard to realize in this case

3. Protocol approach for dedicated purposes (Xwindow etc.)

- Standardized protocols for connecting to servers (e.g., HTTP)
- Pro: Open, global; ~ sequential prog's (+ callback threads)
- Con: Limited to standard functionalities

4. In this lecture: Distributed Programming 'Language' approach

- the only one that is wide spread – up to now!




J. SWE Approaches → Language Approach (2)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Arbitrary, heterogeneous OSES (and databases)
- Distributed programming language



Distributed
Programming
Language
Approach

- original approach (almost died out – this may change!):
 - “distributed programming language” + dist. runtime system
 - compiler can help a lot (like assembler vs. HPL in trad. systems)
- today’s approach:
 - Sequential programming language + extensions + middleware
 - Compiler will not see the distributed program

syntactic difference may be small if dist. lge = seq. lge + extensions
semantic difference is big though (compiler sees / ignores dis. prog.)



J. SWE Approaches → Language Approach (3)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- **Distributed Programming ,Language‘ Approach - Pro:**
 - DRTS intertwined with language (efficiency↑)
 - side-by-side competition of languages → competition of concepts
 - heterogeneity, migration, autonomy etc. may be addressed
 - encompassing support for DistProg (including „Platforms“) possible on a „global“ scale
- **Distributed Programming ,Language‘ Approach - Con:**
 - n languages → n DRTSes
 - loosely intertwined w/ OS (efficiency?)
 - many problems to be addressed (→ extensible solutions?)
... but that's true for the other approaches, too
 - new language constructs ↔ programmers' acceptance / learning effort
 - „compiler sees everything“ ↔ we may not *have* everything @ compiletime

Conclusion: for the rest of this lecture...

- we'll stick to the language approach
- we'll dream of dist. languages, look at lge. extensions/middleware
- E2E argument → if you know what can be done, you can BYO



J. SWE Approaches → Language Approach (4)



Language Approach: Requirements

...think of >100 processes, >10000 objects, dynamic changes (new users etc.)

Sequential language + extensions → distributed lge.

- **minimum** extensions: API + IDL (interface definition lge.); **better**: compiler support
- support for what?
 - **structuring** (hierarchy or 'network' of processes? 'ports' for different message types?)
 - **communication interplay** (,compatibility' of msges sent/received, flow of msges over time, ...)
 - **management** / administration support (installation, configuration)
 - **dynamics** (wrt. processes / connections added/removed)
 - **flexibility** wrt. the „model“ applied (see further in this chapter)
 - **safety in programming** (remember: this is why we have prog.lge's!)
 - compiler & RTS „understand“ your „entities“ & how you cope with them
 - typing of comm. partners, connections, objects interchanged, ...
 - **late distribution** (see DOC subchapter)
 - **openness** wrt. cooperation of separately developed parts
 - **integration** of existing services
 - **distribution transparency** (mimic concurrent prg.)
 - all other kinds of **transparency** (see before)



Language Approach: Support

Software production environments:

- methods & tools must be adapted; or even: need for new methods & tools?
- „integrated SPE“ wrt. UI, functional, & data layers
- integration along life cycle (today, even design → implementation ??)
- aspect oriented (e.g., performance, reliability, ...) across life cycle?
- extensible?
- distributed SPE (distributed-software eng. vs. distributed SW-eng.)

Platforms:

- middleware that includes numerous services, tools, ...
- „horizontal“: generally useful services
- „vertical“: intended for specific application domain / market

But recall End-to-End argument! ... no feature-laden middleware

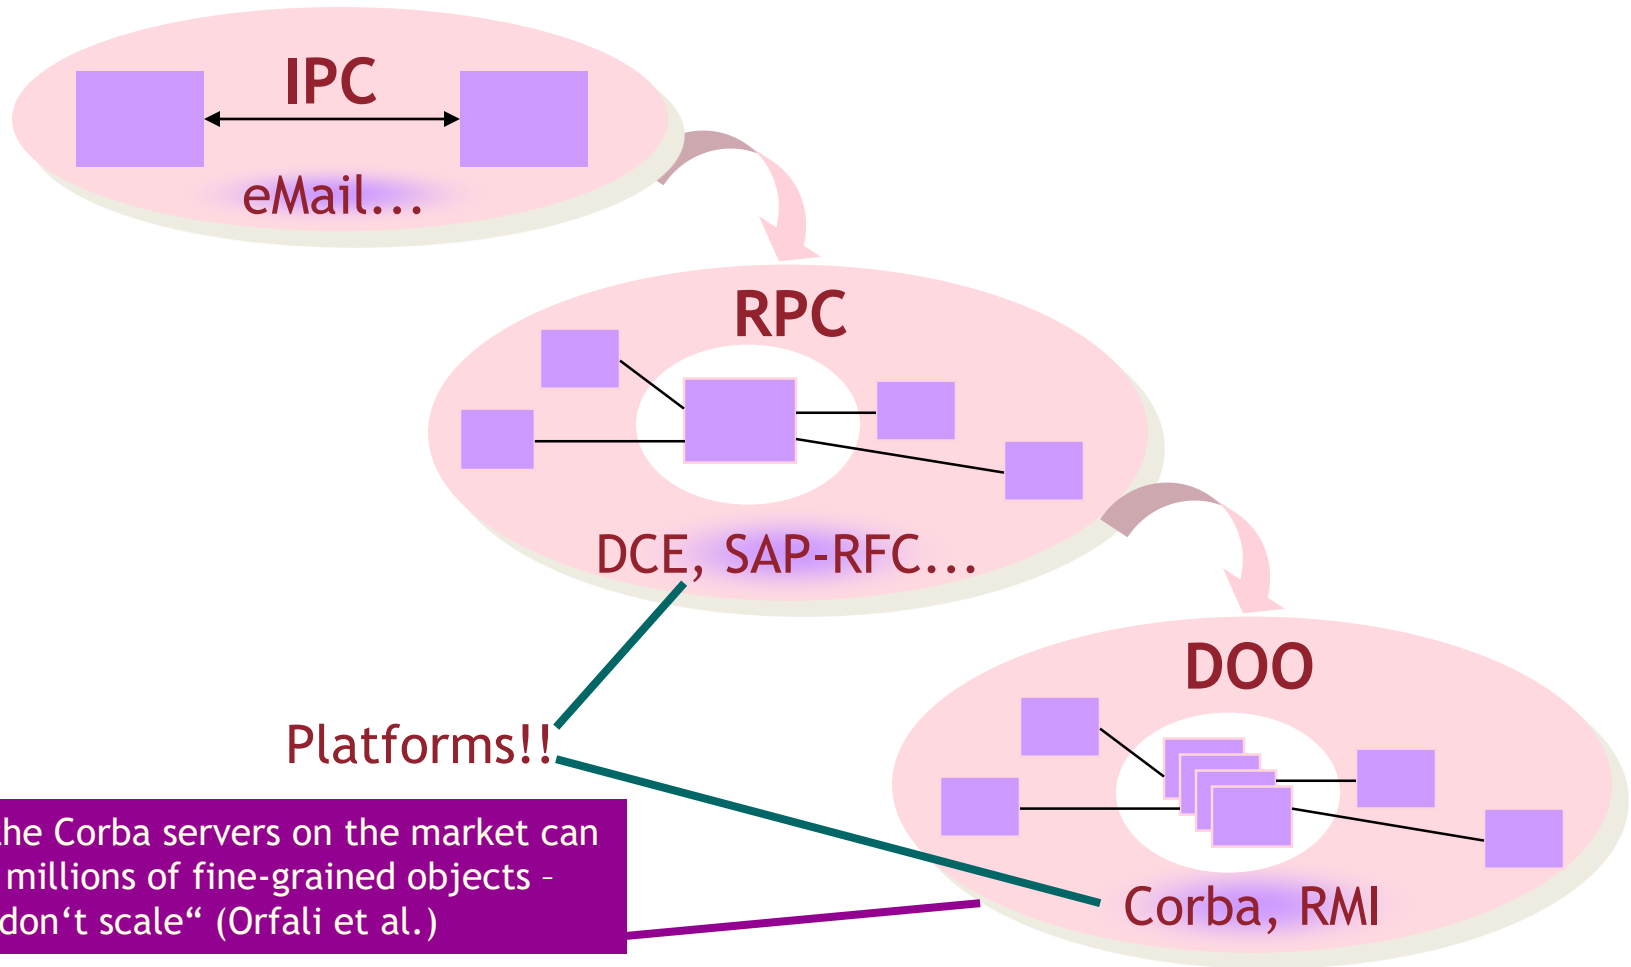
- at least, consider „plain TCP /UDP“ as an alternative! Rest: build your own BYO
(on the other hand, this is the ‚assembler approach‘ of distributed programming)
- know many concepts from literature (/lecture) !!!



K. Distributed Programming Paradigms



TECHNISCHE
UNIVERSITÄT
DARMSTADT



mainstream paradigms




K. Distributed Programming Paradigms




TECHNISCHE
UNIVERSITÄT
DARMSTADT


TupleSpace



Agents&MobileObjects



Cooperative



plus:

streaming media


components

compound docs

Message → event
based (MOM → Pub/Sub)



Cloud

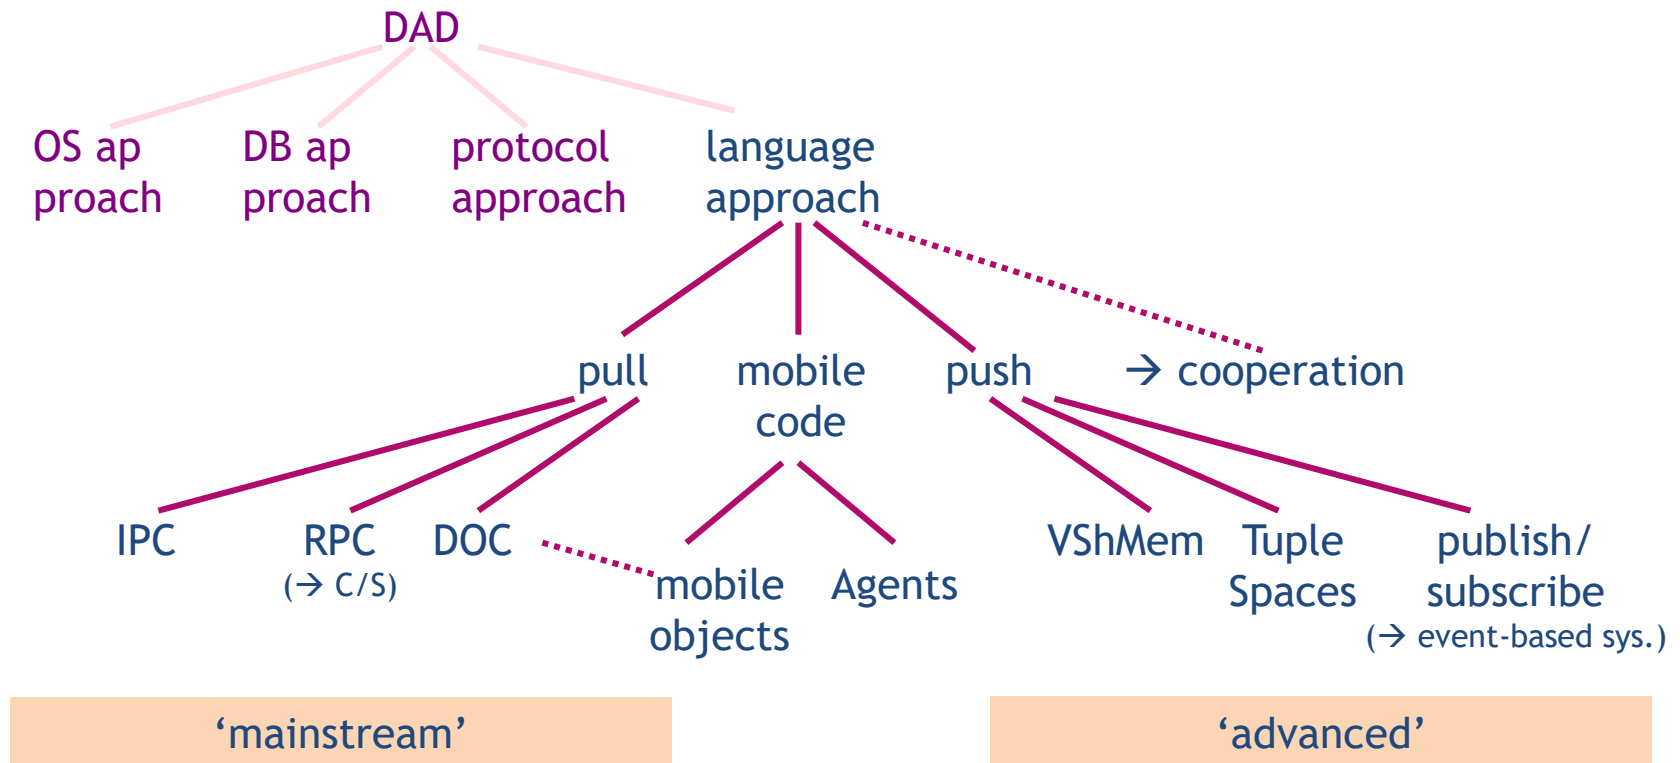


advanced paradigms



Pragmatic Taxonomy

all in one, we get the following taxonomy for distributed application development (DAD):





Famous „Religious“ Disputes



TECHNISCHE
UNIVERSITÄT
DARMSTADT

(not discussed in lecture 1, but during all of course)

- **Escape from Flatland („Flatland“: romance in 19th century)**
 - „DS speed/cost depends on interconnection fabric, need new visions!“
 - but: nothing really changed for 40 years (well, trend went from central switches to shared medium LANs back to switches)
 - will wireless tech., will multiplexing schemes bring about change? (wavelength division multiplexing, CDMA/OFDM air interfaces)
- **„End-to-End argument ...“** (J.Salzer et al., ACM ToCS 2(4), 1984)
 - „feature-rich protocols/middleware = big waste! Build-your-own (BYO)!“
 1. customized selections needed
 2. can't provide 100% reliability etc.! anyway
 - but: development cost! Will components do?
 - **truth-in-the-middle**: reentrant, highly optimized code not bad anyway
- **„A note on distributed programming“** (J. Waldo et al., Sun TR'94)
 - many systems have attempted to paper over the distinction between local and remote objects → they fail to support basic requirements of robustness and reliability!
 - **truth-in-the-middle**: consider SWE effort, new paradigms in middleware



Overview of this Lecture



TECHNISCHE
UNIVERSITÄT
DARMSTADT

1. Introduction

2. Distributed Programming

1. **Mainstream Paradigms** for Distributed Programming

- IPC: Interprocess Communication
- Inlet: Distributed Programming Languages
- RPC: Remote Procedure Call
- Inlet: Concurrency
- DOC: Distributed object-oriented computing
- Web Services

2. **Advanced Paradigms** for Distributed Programming

- Event-Based & Publish/Subscribe Communication
- Tuple Spaces
- Distributed Shared Memory Approach

3. **Mobile Objects**, Unified Objects, Mobile Agents

4. **Cloud Computing**

5. **Formal Approaches**: Process Calculi

3. Distributed Algorithms

1. **Foundation**: Motivation, Properties, Characteristics

2. **Synchronization**: Logical Clocks, Physical Clocks, Global States

3. **Coordination**: Failure Detection, Mutex, Election

4. **Cooperation**: Multicast (On Different Topologies), Consensus

5. **Local Algorithms**