



Telecooperation Lab
Prof. Dr. Max Mühlhäuser

TK1: Distributed Systems - Programming & Algorithms

Chapter 2: Distributed Programming

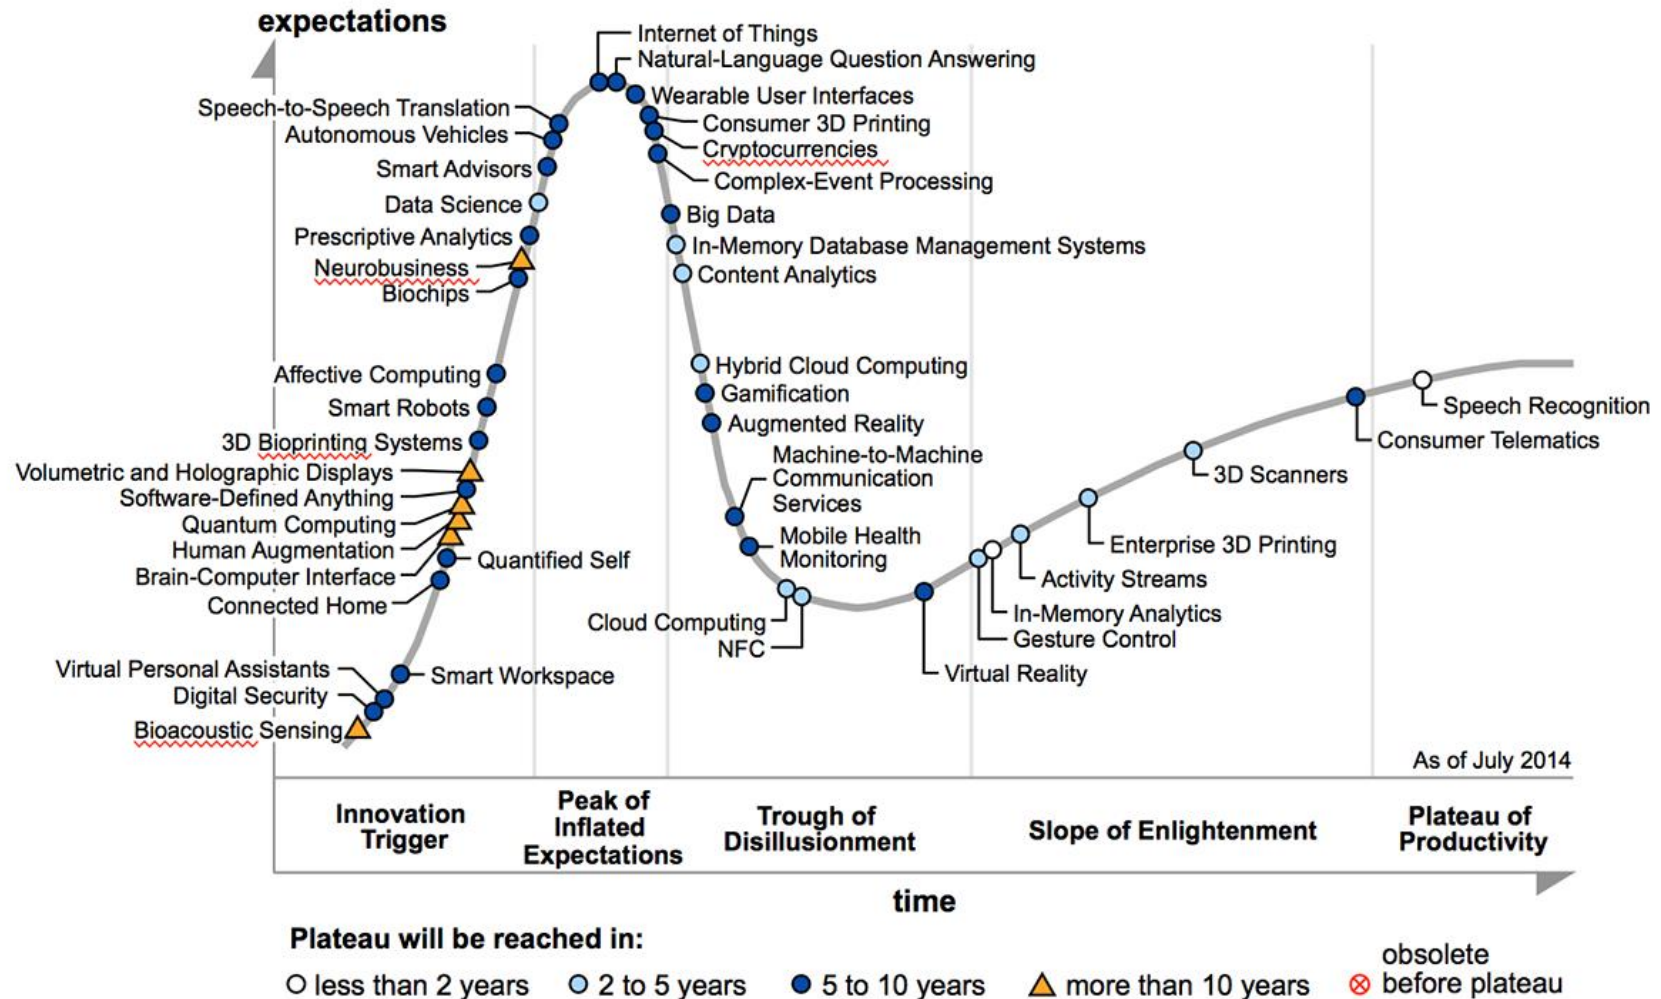
Section 3: Cloud Computing

Lecturer: Prof. Dr. Max Mühlhäuser

Copyrighted material – for TUD student use only



Cloud Computing @ “Hype Cycle”





Cloud Computing: Basic Idea



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- **Provide resources or services** over the Internet to customers with the **reliability of a data center**
- Often **virtualized infrastructure** that can be **scaled on demand** regarding storage, computation and network capacity
- Service providers/users do not need to have knowledge of maintaining the technology infrastructure in the "cloud"





Cloud Computing: NIST Definition



TECHNISCHE
UNIVERSITÄT
DARMSTADT



National Institute of
Standards and Technology
U.S. Department of Commerce

- **On-demand self-service:**

- Unilateral provision of computing capabilities (**server time and network storage**)
- Provision as needed automatically
- Requiring no **human interaction** with each service provider

- **Broad Network Access**

- Capabilities available over the network
- Capabilities accessed through standard mechanisms

- **Resource Pooling**

- Provider pools resources to serve multiple consumers using a multi-tenant model
- Demand based reassignment of physical and virtual resources

- **Rapid Elasticity**

- Capabilities can be elastically provisioned and released, in some cases automatically, to scale rapidly outward and inward commensurate with demand.
- To the consumer, the capabilities available for provisioning often appear to be unlimited and can be appropriated in any quantity at any time.

- **Measured Service**

- Cloud systems automatically control and optimize resource use by leveraging a metering capability at some level of abstraction appropriate to the type of service (e.g., storage, processing, bandwidth, and active user accounts).
- Resource usage can be monitored, controlled, and reported, providing transparency for both the provider and consumer of the utilized service.



In-House vs. External Hosting: Analogy



Cow

Big upfront investment

Produces more than you need

Uses up resources

Maintenance needed

Unpleasant waste product

Supermarket: Bottle of Milk

Small continuous expense

Buy what you need

Less resources needed

No maintenance

Waste is not my problem

Source: Donald Kossmann & Kraska 2009: What is new in the Cloud?




Analogy “translated”



TECHNISCHE
UNIVERSITÄT
DARMSTADT





Cloud Computing (vs. traditional in-house hosting)

- Sold on demand
(typically by minute or hour)
 - “pay as you go” – no need for long-standing contracts
- Elastic: user can determine amount of service
(CPU, Storage, ...) he wants at any given time
- Service is fully managed



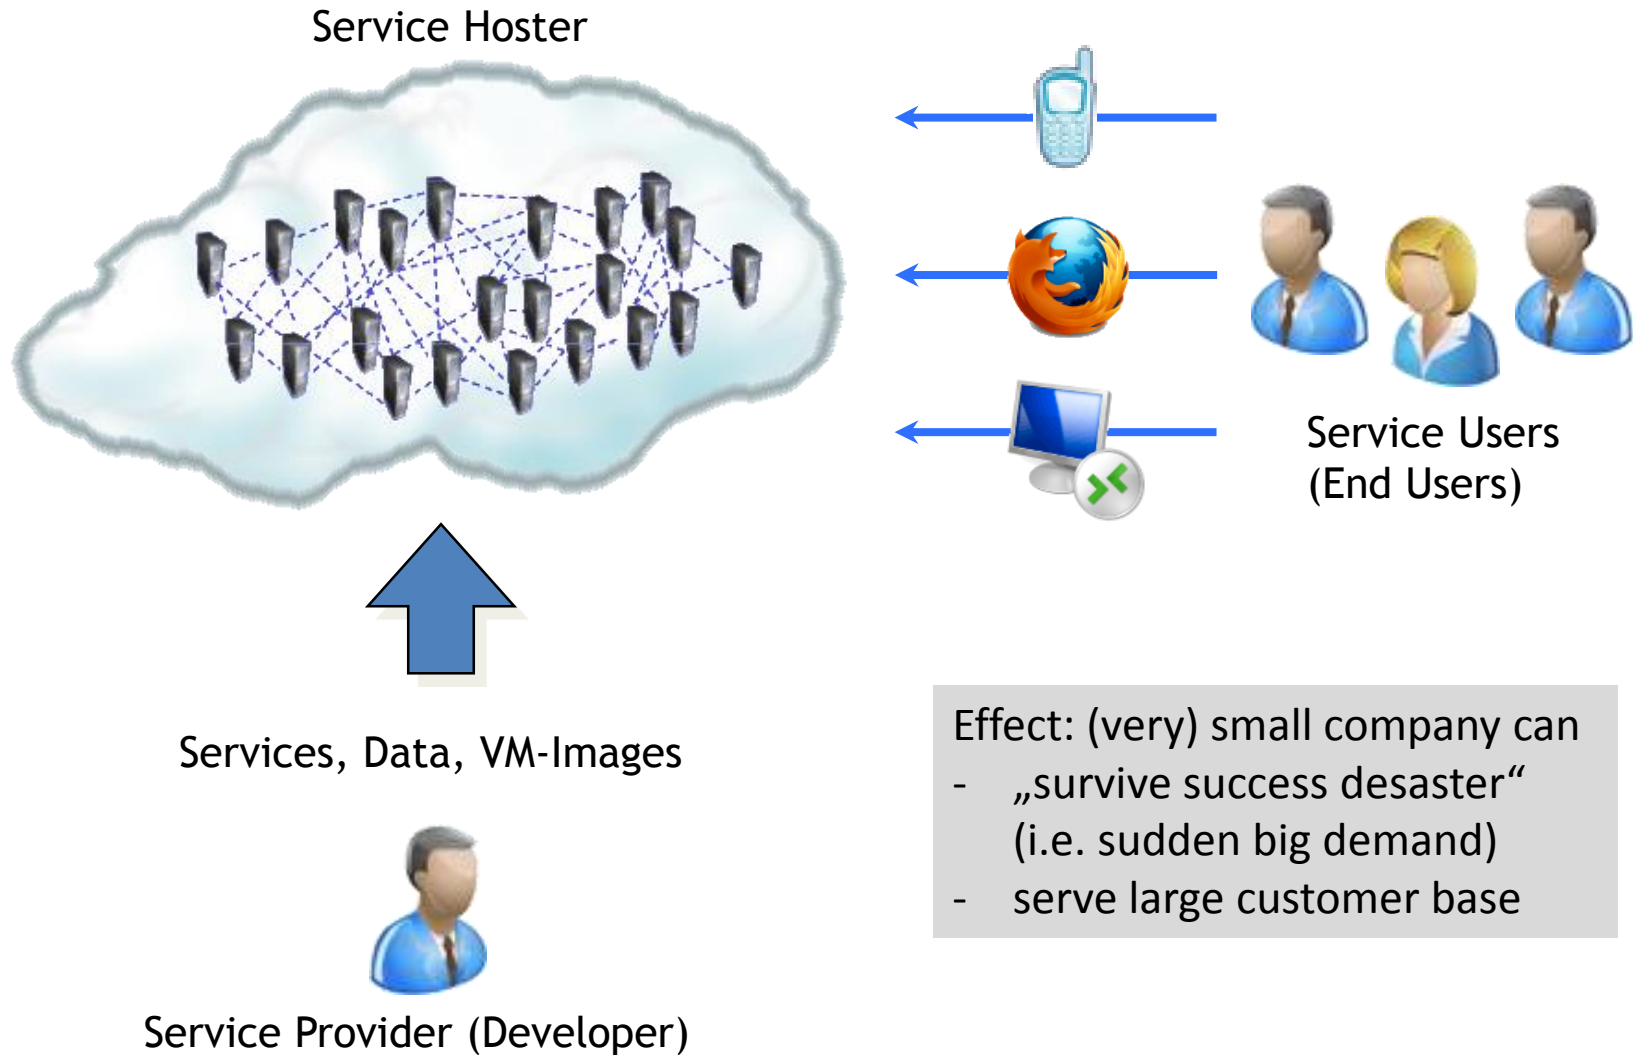


Elasticity





Cloud Computing: Roles





Cloud Types



- **Public Cloud:** sells service to anyone on the Internet
 - Observation (by IBM): 70% of IT budget is spent on average for **maintenance** versus adding new capabilities
 - Goal: Replace fixed IT-costs with variable costs
 - Gartner says: “Organizations are switching from company-owned hardware and software assets to per-use service-based models”
 - Low entry barrier
 - For applications you have to do once
 - For applications you do not have the resources for
 - e.g., Amazon EC2
- **Private Cloud:** used inside a business (e.g., IBM Blue Cloud)
 - Observation: Most systems are oversized, run at 10-20% utilization most of the time
 - Goal: more efficient utilization of resources in a business
- **Community Cloud:** infrastructure shared between several organizations
 - Common concerns (security, compliance, jurisdiction, etc.)
- **Hybrid Cloud:** composition of ≥ 2 cloud infrastructures (private / community / public)
 - Clouds remain unique entities
 - Bound together to realize data or application portability
 - Usage: cloud bursting for load balancing



Key characteristics

- **„Unbounded scale“**
 - Rapid up- and downscaling possible
 - No more oversized systems that are mostly idle
 - Still, the peak load capacity can be provided
 - Assumption: data center hosts many different apps, only few need peak resources at a given time
- **Reliability**
 - Redundant computers and data centers
- More **location independence**
 - Hoster may operate data centers worldwide
- Often based on **cheap hardware**
- Compared to clusters & Grid:
Slow interconnection network
- Subscription-based or **Pay-per-use**
- **Three basic (layers) variants, see below**






Service Models



TECHNISCHE
UNIVERSITÄT
DARMSTADT





































Service Models: Selected „Products“



TECHNISCHE
UNIVERSITÄT
DARMSTADT

 AppStream	 EC2	 Redshift
 CloudFormation	 ElastiCache	 Route 53
 CloudFront	 Elastic Beanstalk	 S3
 CloudSearch	 Elastic Transcoder	 SES
 CloudTrail	 EMR	 SNS
 CloudWatch	 Glacier	 SQS
 CodeDeploy	 IAM	 Storage Gateway
 Cognito	 Kinesis	 SWF
 Config <small>PREVIEW</small>	 Lambda <small>PREVIEW</small>	 Trusted Advisor
 Data Pipeline	 Mobile Analytics	 VPC
 Direct Connect	 OpsWorks	 WorkSpaces
 Directory Service	 RDS	 Zocalo
 DynamoDB		



Infrastructure as a Service (IaaS)

■ **Computer Infrastructure** delivery model

- A platform virtualization environment
 - „virtualization taken a step further“
- Computing resources, such as processing capacity and storage
- Pay-for-what-you-use model
 - resembles electricity, fuel, water ... consumption → also called **utility computing**
- Example: you want to run a batch job but you don't have the infrastructure necessary to run it in a timely manner → use Amazon EC2

■ **Providers**

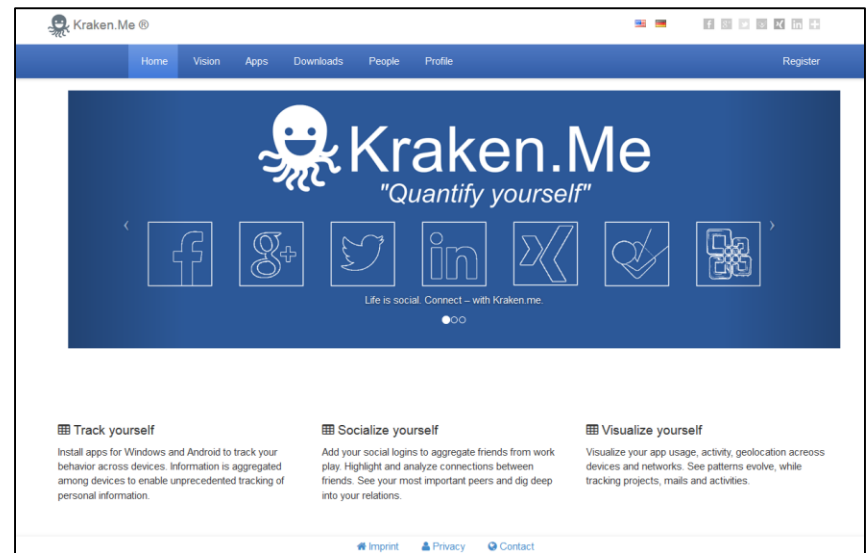
- **Amazon** EC2, S3 (Elastic Compute Cloud, Simple Storage Service)
 - Virtual server instances with unique IP addresses
 - Blocks of storage on demand
 - API to start, stop, access, configure virtual servers and storage
- **Rackspace**: cloudServers
- **FlexiScale**



Infrastructure as a Service (IaaS)

■ Example (Kraken.me)

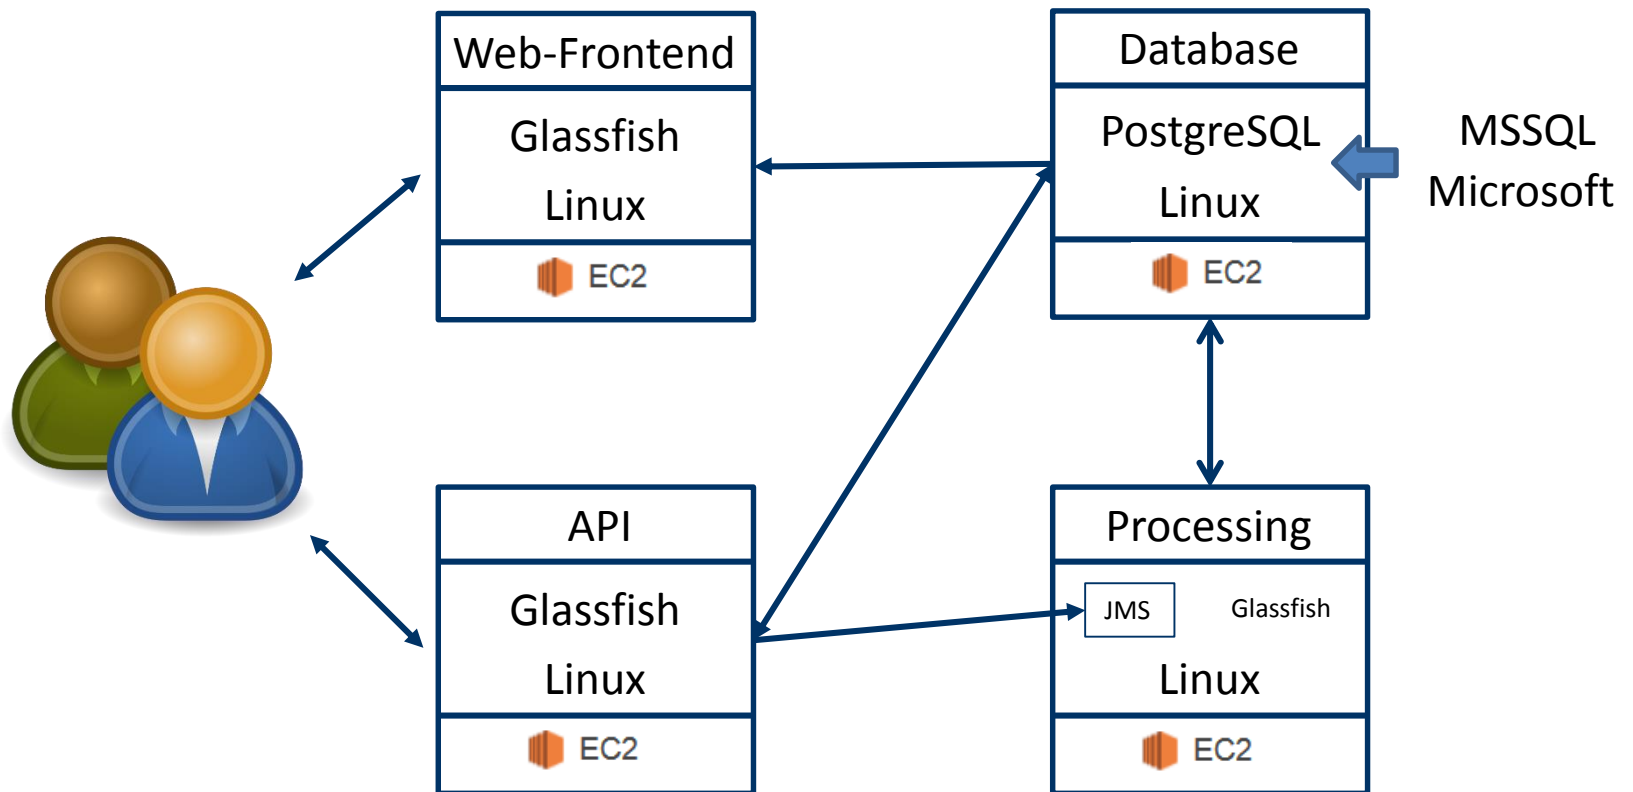
- Web Application with Real-time User Interaction
- Database for User Data
- Processing for Generating Insights
- API





Infrastructure as a Service (IaaS)

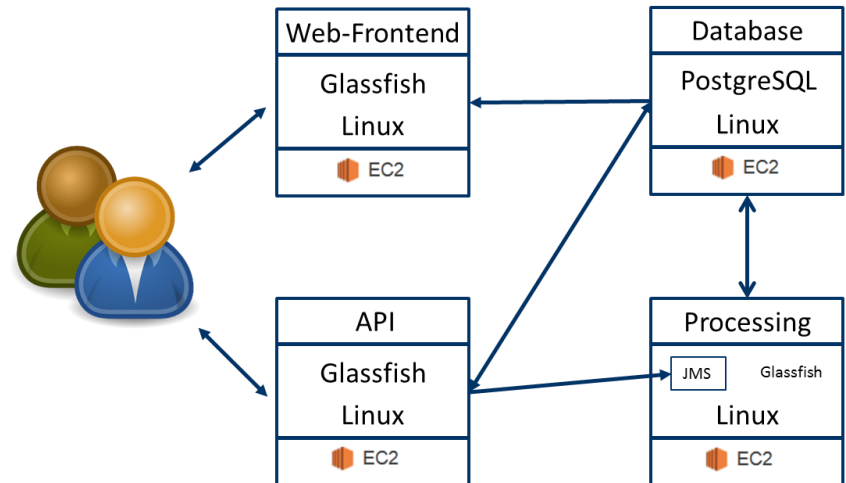
User builds / modifies Platform & Apps, Cloud „just“ provides elastic VMs / StorageSpace





Elasticity (IaaS)

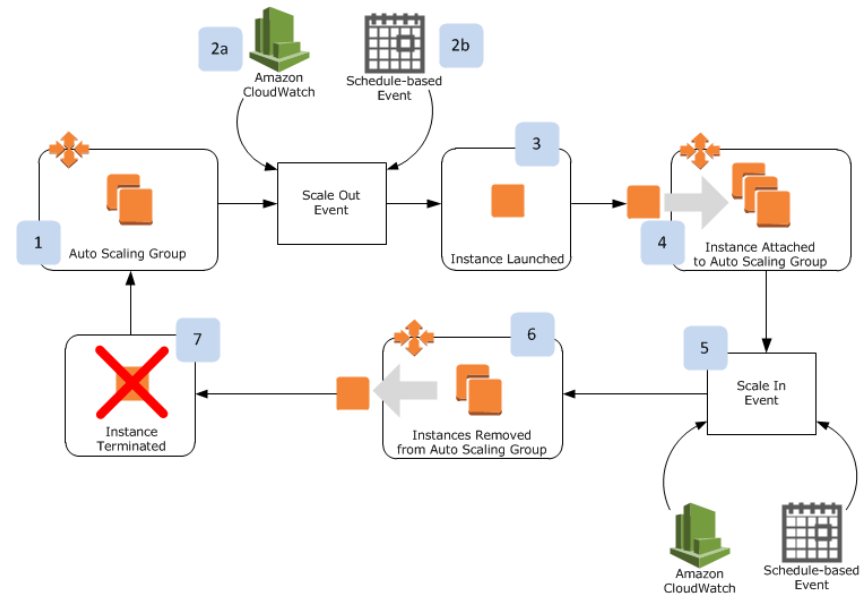
- Larger Instance per EC2
37 instance types, e.g.:
 - T1.micro (1 vCPU, 0,613 GB)
 - I2.8xlarge (32 vCPU, 244GB)
- Limitations to resizing
 - Storage of Virtual Image
 - Format of Virtual Image
 - Downtime





Elasticity (IaaS)

- Add more instances
 - Scale In / Scale Out
 - Groups of instances with same functionality
 - Granularity
 - Load Balancer / Scheduler
- Load Balance between regions etc.





Platform as a Service (PaaS)

■ **Platform** delivery model

- Platforms are built upon infrastructure
- Provider is responsible for maintenance of OS & base frameworks
- Problem: currently no standards for interoperability / data portability
- Examples
 - You want to start storage services on your network for a large number of files and you do not have the storage capacity → use Amazon S3 (IaaS)
 - You want to host a website and expect many hits, but only for a few days → use Rackspace cloudSites (PaaS)

■ **Providers:** Google App Engine, Force.com, Rackspace, ...

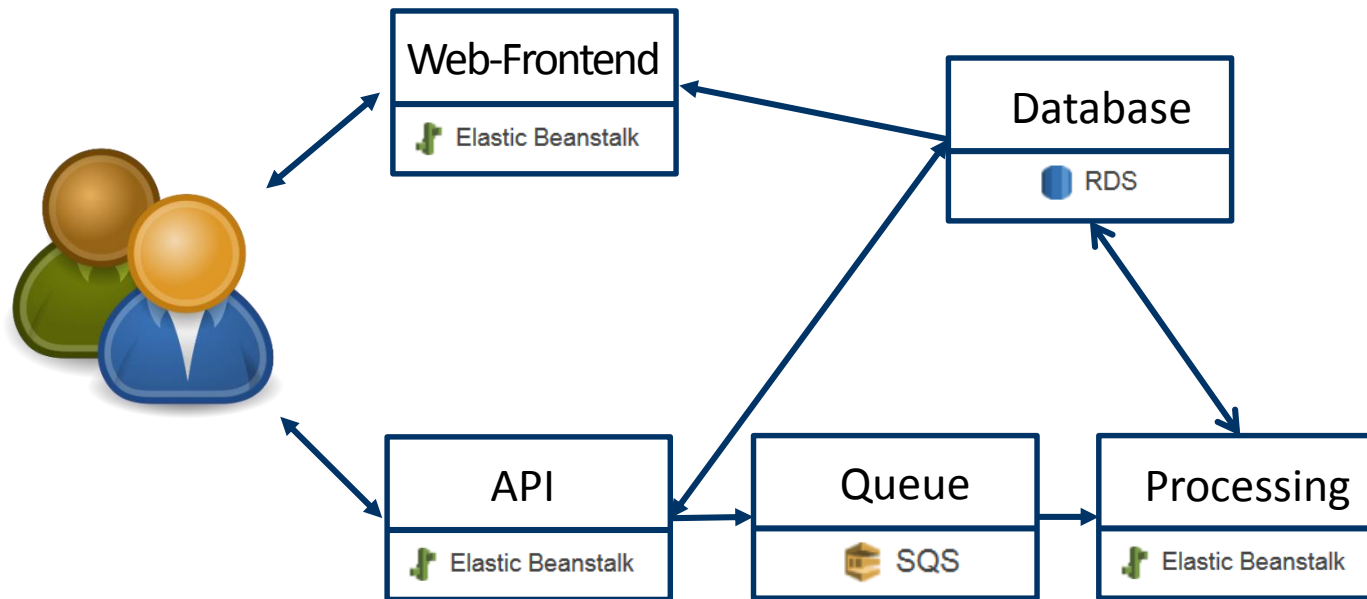
- Rackspace: cloudSites, cloudFiles
 - cloudSites: scalable website w/o need to manage dedicated server



Platform as a Service (PaaS)




TECHNISCHE
UNIVERSITÄT
DARMSTADT



- Build upon elasticity of IaaS
 - more / larger Instances
 - but *automatically*
- “just” develop application
 - built to take advantage of PaaS

compare to IaaS:





Software as a Service (SaaS)

- Software delivery model
 - No need to manage hardware or software
 - Service delivered through a browser (in most cases)
 - Provider hosts both application and data
 - End user is free to use the service from anywhere
 - Customers use the service on demand
- **Instant Scalability**
 - Examples
 - Your email is hosted on an Exchange server in your office and it is very slow
→ outsource this using Hosted Exchange
 - Your current CRM package is not managing the load or you simply don't want to host it in-house → use a SaaS provider such as Salesforce.com
 - SaaS is a very broad market: many domain-specific applications




- Human as a Service (HaaS)
 - Crowdsourcing applications consider humans as resource to execute tasks
 - Very useful for tasks, computers are not good at (e.g. prepare labeled data sets as test/training set for machine learning)
 - Example: Amazon Mechanical Turk





Cloud Computing: Paas vs. SaaS



■ PaaS for Developer

- A set of software & product development tools hosted on the provider's infrastructure
- Developers create applications on the provider's platform over the Internet

■ SaaS for End User



Short History of Cloud Computing



TECHNISCHE
UNIVERSITÄT
DARMSTADT

IBM (key contributor Jim Rymarczyk) launches CP-67 software; one of IBM's first attempts at virtualizing mainframe operating systems

Salesforce.com introduces the concept of delivering enterprise applications via a website

Amazon Web Services provide a suite of cloud-based services, including storage and computation

Cloud providers offer browser-based enterprise applications

1961 1967 1999 2002 2006 2008 2009 2010

John McCarthy envisions that "computation" might someday be organized as a public utility

The term *grid computing* is originated by Ian Foster and Carl Kesselman's work, *The Grid: Blueprint for a New Computing Infrastructure*

Amazon launches Elastic Compute (EC2) as a commercial Web service that small companies and individuals rent to run their own computer applications

Private cloud models make their appearance

Cloud 2.0 model is emerging


G. Pallis
Cloud computing: the new
frontier of internet computing
IEEE Internet Computing
(2010) [14:5: 5562494:70-73]



History: Technology Development

Grid Computing

- Solving large problems with parallel computing
- Made mainstream by Globus Alliance



Utility Computing

- Offering computing resources as a metered service
- Introduced in late 1990s




Software as a Service

- Network-based subscriptions to applications
- Gained momentum in 2001



Cloud Computing

- Next-generation Internet computing
- Next-generation Data Centers





- Grid Definition: (Ian Foster)
 - **Coordinated resource sharing and problem solving in dynamic, multi-institutional virtual organizations**
- Hence, grid is coordinated, aimed at solving problems
 - Motivation to join a grid: more resources available
 - Resources typically provided by everyone and usable by all
- Focus on large scale resource sharing, innovative applications, and high performance
 - This focus distinguishes grid from traditional distributed computing
- Main challenges: authentication, authorization, resource access, and resource discovery



Why 'Grid'? From Grid to Cloud

Why 'Grid'?

- Word has several meanings in English (map grid, starting grid, ...)
- Relationship to grid computing? (are computers arranged in a grid? 😊)
- Best analogy is **power grid**: distribution of electricity
 - In power grid, there are producers and consumers
 - Grid connects them and lets any consumer use any producer seamlessly
 - Grid computing can be defined in a similar way, with similar goals:
 - Anytime, anywhere, just plug in and you are connected

From Grid to Cloud

- Both share a lot: intention, architecture & technology; problems
- Intentions:
 - Manage large, central facilities
 - Define methods by which consumers discover, request & use resources
 - Provide means for efficient parallel execution of software



Grid vs. Cloud Computing: Differences



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Business Model / **accessibility**
 - Grids are only accessible within pre-defined multi-organizations
 - Participants share resources, i.e., most participants also offer resources
 - Each individual organization maintains full control of their resources
 - User/Hoster-Roles in Cloud Computing
 - Cloud computing hoster offers its service to anyone on the Internet
- Computing Model / **application types**
 - Grids are mostly limited to **batch processing**
 - Clouds support **interactive applications** (and also batch)
 - Good support for “web-faced” applications met demand of many e-commerce applications -> great market success
- Programming Model / **job/application submission**
 - Grid: builds mostly on **source code distribution**
 - Send archive with source code to server in Grid
 - Job will be compiled and executed there
 - Cloud: builds on **virtualization**



Grid vs. Cloud Computing



TECHNISCHE
UNIVERSITÄT
DARMSTADT

■ Interoperability

■ Grid

- Wide adoption of the OGF (Open Grid Forum) architecture/protocol standards
- Wide adoption of the Globus toolkit

■ Cloud

- Currently no interoperability
- Big players do not seem to be interested much in standards
 - Amazon, Microsoft, Salesforce, Google
- Open Cloud Consortium
 - Members: Aerospace, Cisco, Infoblox, Nasa, Yahoo, ...
 - ... but the big players are not here

→ **Vendor lock-in!**

■ Elasticity

- Cloud: You pay for the resources you use (**pay-per-use**)
- Grid: Accounting across organizations never worked



Grid vs. Cloud Computing



TECHNISCHE
UNIVERSITÄT
DARMSTADT

■ Amdahl's Law:

s ... speedup

P ... number of parallel processes

$o(P)$... communication overhead

a ... sequential fraction of program

- Max. expected speedup of overall system, assuming: *not everything* can be parallelized
- Expresses relationship btw. sequential algorithm and parallelized version
- The **sequential fraction** of a program (in terms of run time) determines the **upper bound for speedup**! - still true, but only if we look at solving one problem


■ Grids

- Parallel computers with fast interconnection networks -> low $o(P)$
- Solve often **only 1 problem**

■ Clouds

- Inexpensive hardware, „slow“ interconnects -> high $o(P)$
- Either a is really small (cf. MapReduce)
- or mostly: **solve n problems** in parallel

$$s = \frac{1}{a + o(P) + \frac{1-a}{P}} \leq \frac{1}{a}$$





The „Cloud Stack“



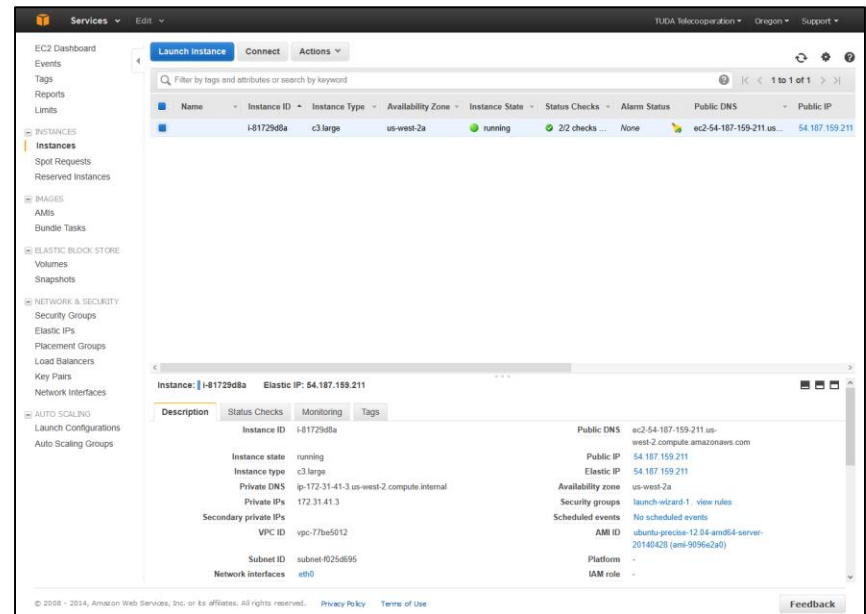
TECHNISCHE
UNIVERSITÄT
DARMSTADT

Application	Web	AppEngine (Google)	Azure Web Role (MS)		
	Services		Azure Worker Role (MS)		
Platform	Storage	Sherpa (Yahoo)	BigTable (Google)	Azure Tables (MS)	
	Batch Processing	Hadoop (Apache)	PIG (Yahoo)	MapReduce (Google)	DryadLinq (MS)
	OS				
Infrastructure	Virtualization	EC2 (Amazon)			
	Block Storage	S3 (Amazon)	GFS (Google)		



Amazon Elastic Cloud Computing

- Introduced 2006 based on internal project
 - >1.000.000 physical servers
- EC2 is not much more than “VM hosting”
 - Amazon Machine Image (AMI) = Encrypted VM image, stored in S3 repository
 - Pre-configured templates provided
- Instances: *controlled* using a Web Service interface, *accessed* via RDP or SSH





Amazon EC2: Block Store, API

Amazon EC2 API:

- Running
 - RegisterImage / DeregisterImage
 - RunInstances: launch the specified number of instances
 - TerminateInstances, RebootInstances
- Networking
 - AllocateAddress: acquire an elastic IP address; AssociateAddress: associate elastic IP address with instance
- Storage
 - CreateVolume / DeleteVolume
 - AttachVolume: mount an EBS volume
 - CreateSnapshot / DeleteSnapshot



Amazon EC2



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Step 1: Choose an Amazon Machine Image (AMI)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. You can select an AMI provided by AWS, our user community, or the AWS Marketplace, or you can select one of your own AMIs.

Quick Start

My AMIs
AWS Marketplace
Community AMIs

☐ Free tier only (1)

Amazon Linux AMI 2014.09.1 (HVM) - ami-b5a7ea85 **Select**
The Amazon Linux AMI is an EBS backed image. It includes the 3.14 kernel, Ruby 2.1, PHP 5.5, PostgreSQL 9.3, Docker 1.2, the AWS command line tools, and repository access to many other packages.
Root device type: ebs Virtualization type: hvm 64-bit

Red Hat Enterprise Linux 7.0 (HVM), SSD Volume Type - ami-99bef1a9 **Select**
Red Hat Enterprise Linux version 7.0 (HVM), EBS General Purpose (SSD) Volume Type
Root device type: ebs Virtualization type: hvm 64-bit

SUSE Linux Enterprise Server 12 (HVM), SSD Volume Type - ami-d7450be7 **Select**
SUSE Linux Enterprise Server 12 (HVM), EBS General Purpose (SSD) Volume Type. Public Cloud, Advanced Systems Management, Web and Scripting, and Legacy modules enabled.
Root device type: ebs Virtualization type: hvm 64-bit

Ubuntu Server 14.04 LTS (HVM), SSD Volume Type - ami-3d50120d **Select**
Ubuntu Server 14.04 LTS (HVM), EBS General Purpose (SSD) Volume Type. Support available from Canonical (<http://www.ubuntu.com/cloud/Services>)
Root device type: ebs Virtualization type: hvm 64-bit

Microsoft Windows Server 2012 R2 Base - ami-a13d5891 **Select**
Microsoft Windows 2012 R2 Standard edition with 64-bit architecture. (English)
Root device type: ebs Virtualization type: hvm 64-bit

Are you launching a database instance? Try Amazon RDS. **Hide**
Amazon Relational Database Service (RDS) makes it easy to set up, operate, and scale a relational database of your choice (MySQL, PostgreSQL, Oracle, SQL Server) in the cloud. It provides cost-efficient and resizable capacity while managing time-consuming database management tasks, freeing you up to focus on your applications and business. [Learn more](#)
Launch a database using RDS

© 2008 - 2014, Amazon Web Services, Inc. or its affiliates. All rights reserved. [Privacy Policy](#) [Terms of Use](#) **Feedback**

Step 2: Choose an Instance Type

Amazon EC2 provides a wide selection of instance types optimized to fit different use cases. Instances are virtual servers that can run applications. They have varying combinations of CPU, memory, storage, and networking capacity, and give you the flexibility to choose the appropriate mix of resources for your applications. [Learn more](#) about instance types and how they can meet your computing needs.

Filter by: All instance types Current generation Show/Hide Columns

Currently selected: t2.micro (Variable ECUs, 1 vCPUs, 2.5 GHz, Intel Xeon Family, 1 GiB memory, EBS only)

	Family	Type	vCPUs (1)	Memory (GiB)	Instance Storage (GiB) (1)	EBS-Optimized Available (1)	Network Performance (1)
<input checked="" type="checkbox"/>	General purpose	t2.micro Free tier eligible	1	1	EBS only	-	Low to Moderate
<input type="checkbox"/>	General purpose	t2.small	1	2	EBS only	-	Low to Moderate
<input type="checkbox"/>	General purpose	t2.medium	2	4	EBS only	-	Low to Moderate
<input type="checkbox"/>	General purpose	m3.medium	1	3.75	1 x 4 (SSD)	-	Moderate
<input type="checkbox"/>	General purpose	m3.large	2	7.5	1 x 32 (SSD)	-	Moderate
<input type="checkbox"/>	General purpose	m3.xlarge	4	15	2 x 40 (SSD)	Yes	High
<input type="checkbox"/>	General purpose	m3.2xlarge	8	30	2 x 80 (SSD)	Yes	High
<input type="checkbox"/>	Compute optimized	c3.large	2	3.75	2 x 16 (SSD)	-	Moderate
<input type="checkbox"/>	Compute optimized	c3.xlarge	4	7.5	2 x 40 (SSD)	Yes	Moderate
<input type="checkbox"/>	Compute optimized	c3.2xlarge	8	15	2 x 80 (SSD)	Yes	High
<input type="checkbox"/>	Compute optimized	c3.4xlarge	16	30	2 x 160 (SSD)	Yes	High
<input type="checkbox"/>	Compute optimized	c3.8xlarge	32	60	2 x 320 (SSD)	-	10 GigaBit
<input type="checkbox"/>	GPU instances	g2.2xlarge	8	15	1 x 60 (SSD)	Yes	High

Cancel **Previous** **Review and Launch** **Next: Configure Instance Details**

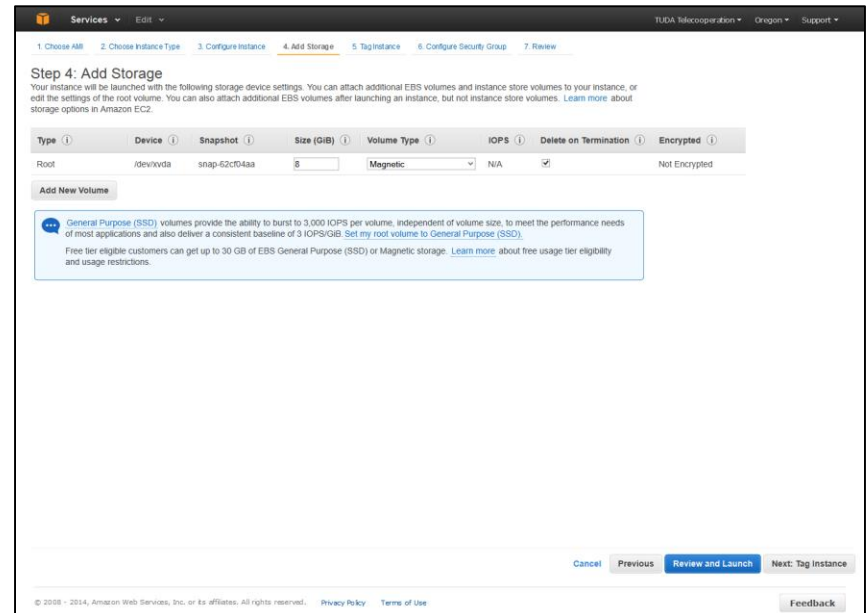
© 2008 - 2014, Amazon Web Services, Inc. or its affiliates. All rights reserved. [Privacy Policy](#) [Terms of Use](#) **Feedback**



- The local storage of instances is non-persistent
 - AMI Images are stored in Amazon S3


Amazon Elastic Block Store

- Raw block device that can be mounted (“attached”) by an instance (one instance at a time)
- Persistence beyond the lifetime of an instance, snapshots can be stored on S3
- Block store API functions (accessed via SOAP or REST):
 - CreateVolume / DeleteVolume; DescribeVolume: get info. about volume (size, source snapshot, creation time, status)
 - AttachVolume / DetachVolume
 - CreateSnapshot / DeleteSnapshot; DescribeSnapshots: get information about snapshots





- Instances are Availability Zones
 - Describe the location of a datacenter as geographic region (e.g., us-east-1a)
 - Placement of instances can be controlled
 - Distribution of instances can be used for fault tolerance





The „Cloud Stack“



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Application	Web	AppEngine (Google)	Azure Web Role (MS)		
	Services		Azure Worker Role (MS)		
Platform	Storage	Sherpa (Yahoo)	BigTable (Google)	Azure Tables (MS)	
	Batch Processing	Hadoop (Apache)	PIG (Yahoo)	MapReduce (Google)	DryadLinq (MS)
	OS				
Infrastructure	Virtualization	EC2 (Amazon)			
	Block Storage	S3 (Amazon)	GFS (Google)		



- Amazon Simple Store Service (S3)
 - Based on REST or SOAP APIs, also browser-based management tools available
 - Organized in buckets
 - bucket = “drive” in which files can be stored
 - ACLs (access control lists) → share parts of the bucket (in group, in public)
 - Other services build on S3, such as EC2
- Example: Creating a Bucket

```
PUT /onjava HTTP/1.1
Content-Length: 0
User-Agent: jClientUpload
Host: s3.amazonaws.com
Date: Sun, 05 Aug 2007 15:33:59 GMT
Authorization: AWS 15B4D3461F177624206A:YFhSWKDG3qDnGbV7JCnkfdz/IHY=
```

- Authorization: AWS <AWSAccessKeyID>:<Signature>
 - AWSAccessKeyID: customer ID
 - Signature: SHA1 hash of message, encrypted with customer’s private key



The „Cloud Stack“



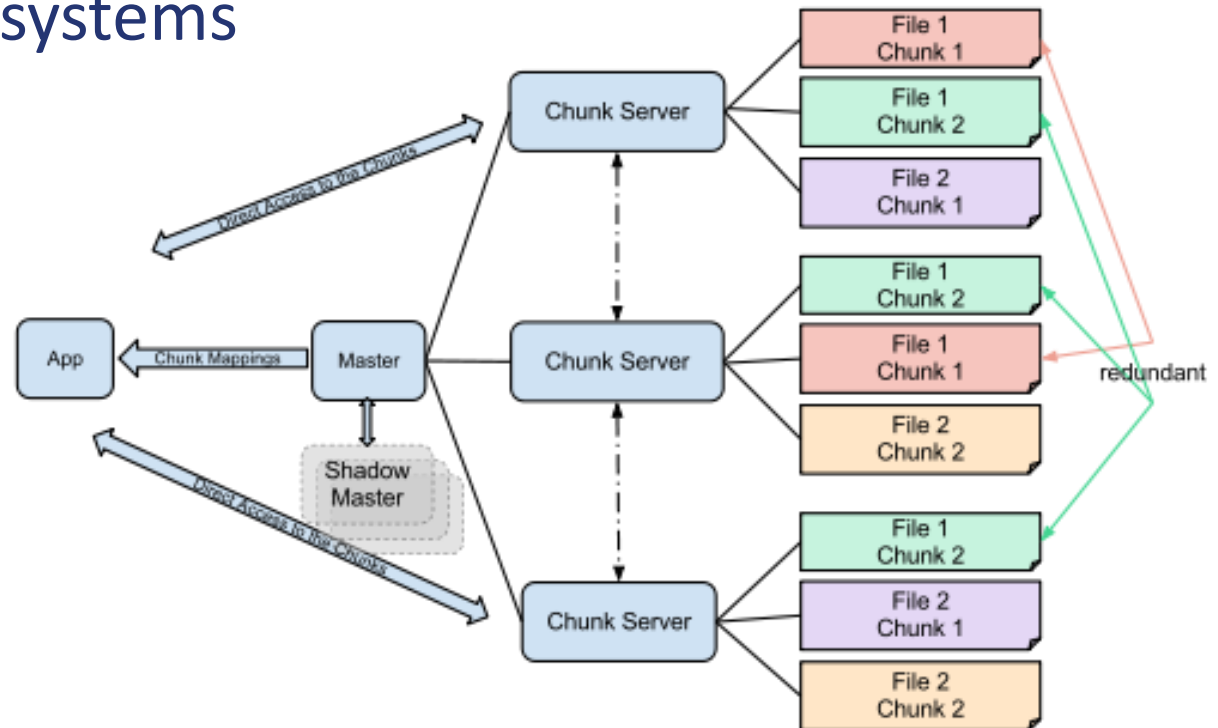
TECHNISCHE
UNIVERSITÄT
DARMSTADT

Application	Web	AppEngine (Google)	Azure Web Role (MS)		
	Services	Azure Worker Role (MS)			
Platform	Storage	Sherpa (Yahoo)	BigTable (Google)	Azure Tables (MS)	
	Batch Processing	Hadoop (Apache)	PIG (Yahoo)	MapReduce (Google)	DryadLinq (MS)
	OS				
Infrastructure	Virtualization	EC2 (Amazon)			
	Block Storage	S3 (Amazon)	GFS (Google)		



Google Filesystem (GFS)

- Shares many of the same goals as previous distributed file systems
 - Performance
 - Scalability
 - Reliability
 - Availability





Google Filesystem (GFS)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

1. Component failures are the norm rather than the exception

- Quantity and quality of components virtually guarantees failures
- Constant monitoring, error detection, fault tolerance, and automatic recovery are integral to the system

2. Huge files (by traditional standards)

- Multi GB files are common
- I/O operations and blocks sizes must be revisited
- Number of files is limited

3. Most files are mutated by appending new data

- Focus for performance optimization and atomicity guarantees
- Concurrent Appends possible
- Random file modifications are supported but inefficient

4. Co-designing the applications and APIs benefits overall system by increasing flexibility

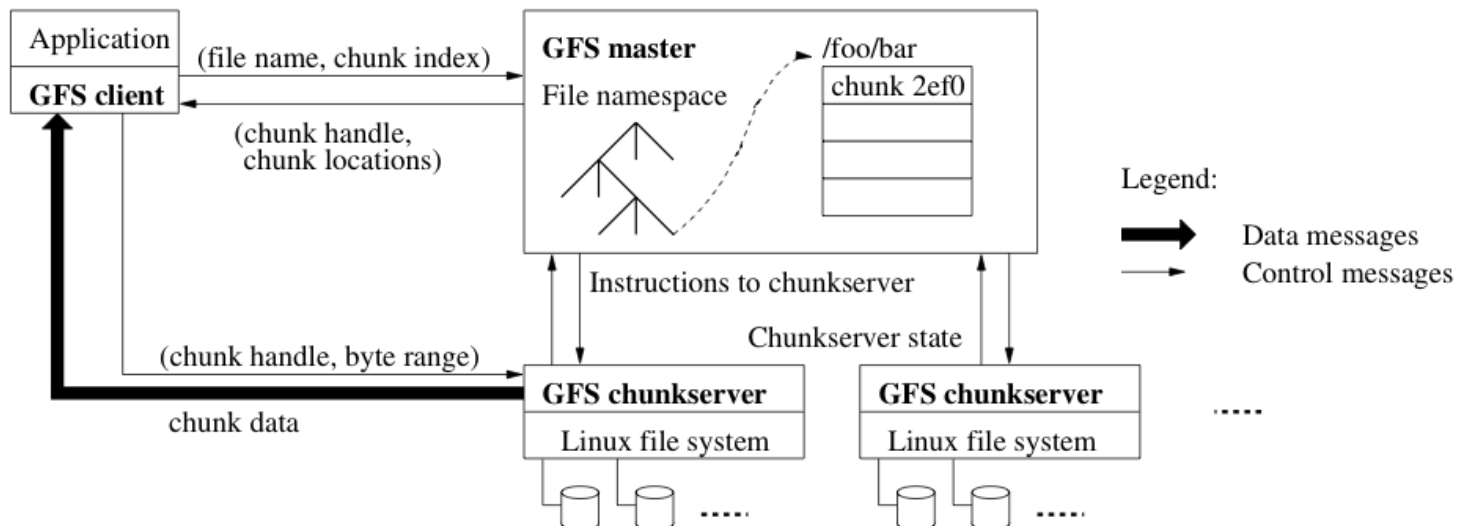
5. Two main types of read: large streaming reads and small random reads

6. High sustained bandwidth more important than low latency



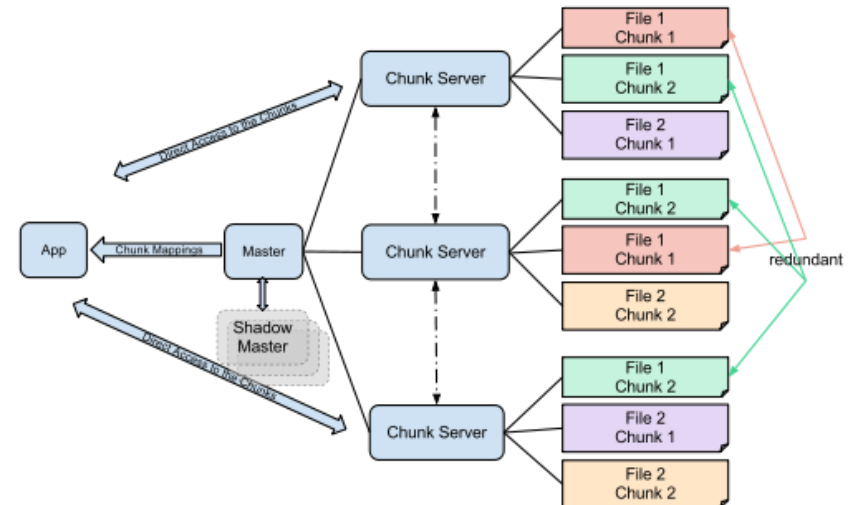
GFS: Architecture

- GFS cluster consists of
 - A single **master**
 - Multiple **chunkserver**s
 - Is accessed by multiple **clients**
- Files organized hierarchically in directories & identified by pathnames






- Maintains metadata
 - names space, access control info, file to chunk mappings, chunk locations, etc.
 - Kept in memory
- Client contacts Master to get chunk locations, then deals directly w/ chunkservers
 - Master is not a bottleneck for reads/writes
- Single master
 - Simplifies design
 - Helps make sophisticated chunk placement & replication decisions (global knowledge)
 - 'Shadow masters' keep replica of master state & log, details later





GFS: Chunks

- Files are divided into fixed-size chunks
 - 64 MB default
- Each chunk has a immutable globally unique 64-bit chunk-handle
 - Handle is assigned by the master at chunk creation
- Advantages
 - Reduced client / master interaction
 - Cache all chunk locations for multi-TB working sets
 - Reduced size of metadata
- But: Hotspots possible





- GFS does not implement a standard API such as POSIX
- Linked as library into apps
- Communicates with master and chunkservers for reading and writing
 - Master interactions only for metadata
 - Chunkserver interactions for data
- Cache metadata
 - Do not cache file data (no coherence issues)



GFS: Operations



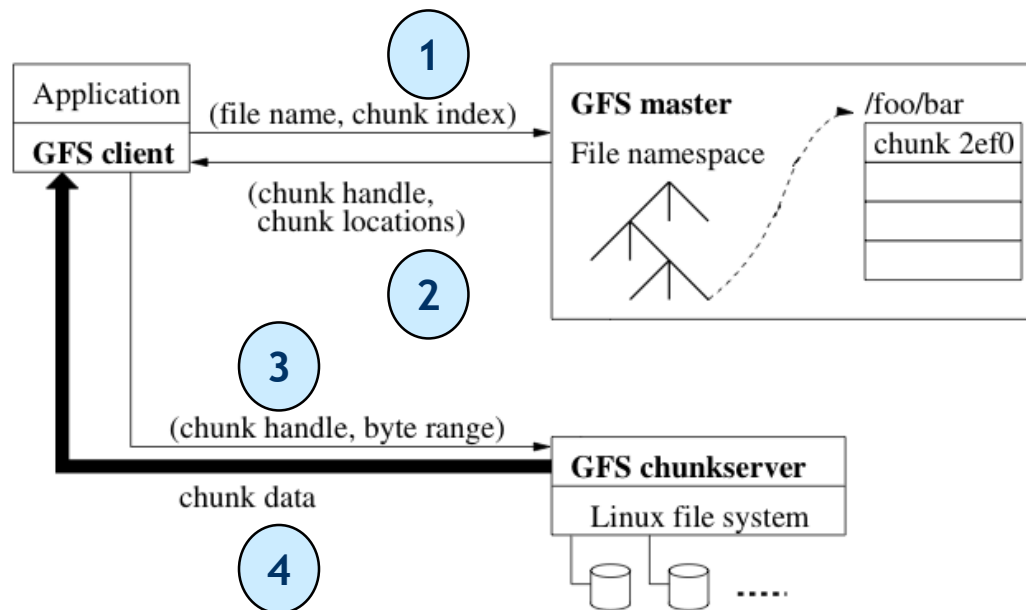
TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Standard: create, delete, open, close, read, and write files
- **Snapshot**
 - Creates a copy of a file or a directory tree at low cost
- **Record Append**
 - Allows multiple clients to append data to the same file concurrently
 - Guaranteeing the atomicity of each individual client's append
 - Useful for implementing multi-way merge results
 - Useful for producer-consumer queues that many clients can simultaneously append to without additional locking



GFS: Read

1. Client sends Master a request containing file name and chunk index (= byte offset / 64MB)
2. Master replies with **chunk handle** and **locations of replicas**
3. Client sends request to one replica (closest one)
4. Client reads data from chunkserver

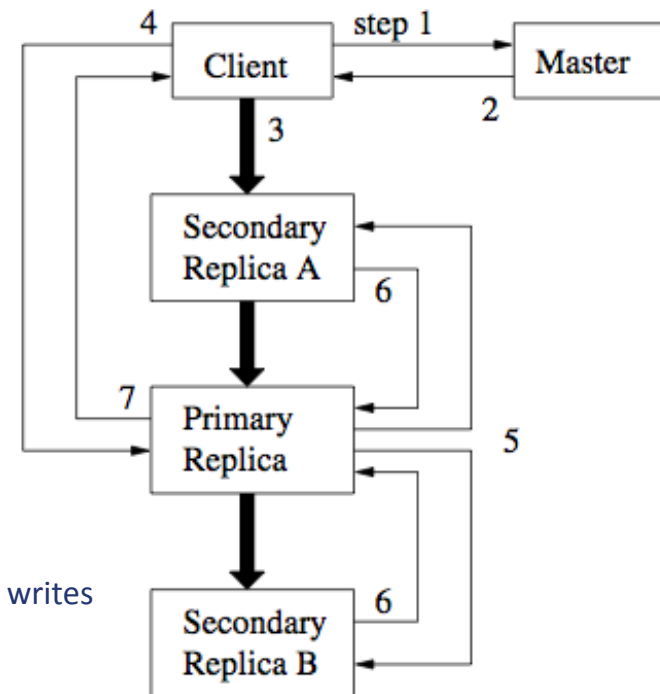




GFS: Write



1. *Client* asks *master* for *lease holder* (see below) of chunk to write to
2. Reply: identity of the *primary* and locations of *secondary* replicas
3. *Client* pushes data to all replicas on *chunkservers* (directly, not as depicted)
4. Once all the *chunkservers* have acknowledged receiving the data, client sends a write request to the *primary*
 - Request identifies the other replicas
 - Primary assigns consecutive serial numbers → serialization of concurrent writes (an atomic change to the file is called a **mutation**)
 - Data in local write buffer is ordered according to serial no. order
5. *Primary* forwards write request to all *secondary* replicas
 - Data is ordered according to serial number order
6. *Secondaries* all reply to *primary* indicating completion
7. *Primary* replies to *client*





GFS: Atomic Record Append

■ Problem

- In a traditional write, the client specifies offset at which data is to be written
- Concurrent writes to the same region are not serializable: the region may end up containing data fragments from multiple clients

■ Solution: **Atomic Record Append**

- Client specifies only the data to write
- GFS chooses and returns the offset it writes to
- No need for a distributed lock manager
 - GFS chooses the offset, not the client
- Follows similar control flow as Write
 - Primary tells secondary replicas to append at the same offset
 - Data must be written at the same offset for all chunk replicas for success to be reported



■ Purpose

- Maintain consistent *mutation order* across replicas by identifying primary copy
- De-allocation of resources in case of a failure

■ **Chunkserver: Failure Recovery**

- Write, 1st step: Client asks Master for Lease Holder
 - If no primary assigned for specified chunk → Master issues new **Lease**
 - Lease has an initial timeout of 60 seconds
 - As long as mutation goes on, primary periodically requests extensions from Master
- Now, if the master loses communication with a primary:
 - Master can safely grant a new lease to another replica after the old lease has expired
- Now, if the faulty chunk server comes back:
 - Master maintains version number for each chunk
 - Allows detecting stale replicas



GFS Fault Tolerance: Master

■ Operation Log

- Record of all critical metadata changes
- Stored on Master and replicated on other machines
- Defines order of concurrent operations
- Changes not visible to clients until metadata changes are persistent
- Master checkpoints its state whenever the log grows beyond a certain size
 - Checkpoint is a compact B-tree (used for namespace lookup)

■ Master Recovery

- Read latest complete checkpoint
- Replay operation log

■ Replication: Shadow Masters

- Help providing read-only access while primary master is up
- ... and even when primary master is down

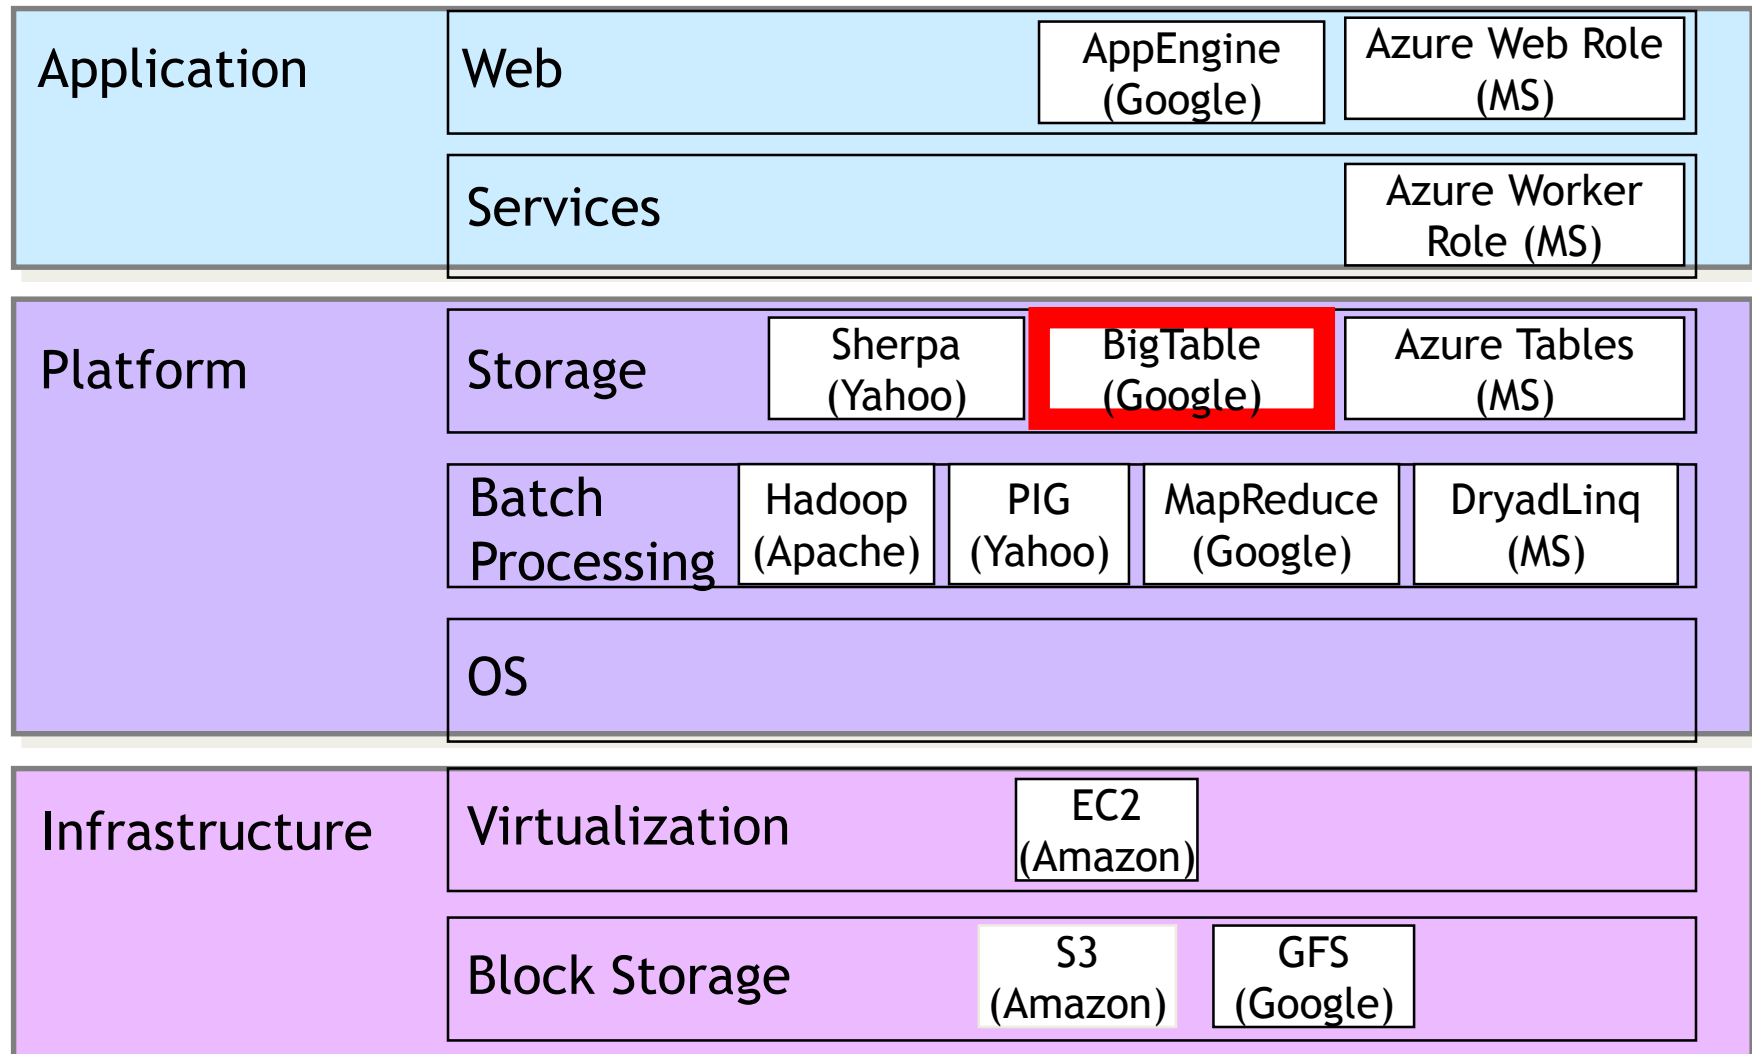
<http://static.googleusercontent.com/media/research.google.com/en//archive/gfs-sosp2003.pdf>



The „Cloud Stack“



TECHNISCHE
UNIVERSITÄT
DARMSTADT





BigTable (Google)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

BigTable: distributed storage system for managing **structured data**

- Lots of (semi-)structured data at Google
 - **URLs:** Contents, crawl metadata, links, anchors, pagerank, ...
 - **Per-user data:** User preference settings, recent queries/search results, ...
 - **Geo-locations:** places (shop, restaurant ...), roads, satellite img. data, user annotations ...
- Designed to scale to a **very large size**
 - Petabytes of data across thousands of servers
 - Billions of URLs, many versions/page (~20K/version)
 - Hundreds of millions of users, thousands or q/sec
 - Many TB of satellite image data
- Scale is too large for most SQL databases
 - Even if it weren't, licensing costs would be very high
 - Low-level storage optimizations help performance significantly
 - Much harder to do when running on top of a database layer

<http://www.cs.cornell.edu/courses/cs6464/2009sp/lectures/17-bigtable.pdf>

<http://static.googleusercontent.com/media/research.google.com/en//archive/bigtable-osdi06.pdf>



BigTable: Goals

- Want asynchronous processes to be continuously updating different pieces of data
 - Want access to most current data at any time
- Need to support ...:
 - Very high read/write rates (millions of op's per second)
 - Efficient scans over all or interesting subsets of data
 - Efficient joins of large one-to-one and one-to-many datasets
- Often want to examine data changes over time
 - E.g. Contents of a web page over multiple crawls
- Fault-tolerant, persistent
- Self-managing
 - Servers can be added/removed dynamically
 - Servers adjust to load imbalance



BigTable: Dat Model



TECHNISCHE
UNIVERSITÄT
DARMSTADT

(row, column, time) → string

- String: Array of uninterpreted bytes
 - Whatever data
 - No size limit
- Row keys: uninterpreted bytes
 - Up to 64KB
 - Data sorted by row key
 - Atomic changes to multiple columns with the same row key
- Column: uninterpreted bytes
 - Sorted into column families
 - Basic unit of access control / only a few hundred
 - Lone columns are trivial families
- Time
 - Used for versioning
 - 64bit Integers
 - Timestamps per column
 - Current time in milliseconds or application timestamp



BigTable: Applications

- Used in many Google projects (data from 2006)

Project name	Table size (TB)	# Cells (billions)	# Column Families	% in memory	Latency-sensitive?
<i>Crawl</i>	800	1000	16	0%	No
<i>Crawl</i>	50	200	2	0%	No
<i>Google Analytics</i>	20	10	1	0%	Yes
<i>Google Analytics</i>	200	80	1	0%	Yes
<i>Google Base</i>	2	10	29	15%	Yes
<i>Google Earth</i>	0.5	8	7	33%	Yes
<i>Google Earth</i>	70	9	8	0%	No
<i>Orkut</i>	9	0.9	8	1%	Yes
<i>Personalized Search</i>	4	6	93	5%	Yes



DBMS vs. Key/Value Store

Database Structure



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Relational Database

- Tables w/ columns and rows, rows contain column values, rows in table all have same schema
- Data model is well-defined in advance, schema is typed, has constraints & relationships to enforce integrity
- Data model based on „natural“ representation of data, independent of app data structures
- Data model is normalized to remove duplication. Relationships associate between multiple tables
- No good support for large amounts of binary data

Key/Value Store

- Table is a bucket where key/value pairs are put into.
- Different rows can have different „schemas“
- Records are identified by keys. A dynamic set of attributes can be associated with a key
- Sometimes attributes are all of a string type only
- No relationships are explicitly defined. No normalization models
- Handles binary data well



DBMS vs. Key/Value Store

Data Access



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Relational Database

- Data is created, updated, and retrieved using SQL
- SQL queries can act on single tables, but also on multiple tables through table joins
- SQL provides functions for aggregation & complex filtering
- „Active Database“ with triggers, stored procedures, functions

Key/Value Store

- Data manipulated through proprietary APIs
- Sometimes SQL-like syntax without join or update functionality (e.g., GQL)
- Only basic filter predicates ($=$, \neq , $<$, $>$). Often: processing not considered to be part of DB (\rightarrow MapReduce, Hadoop)
- App & data integrity logic all contained in app code



DBMS vs. Key/Value Store

Application Interface

Relational Database

- De-facto standards for SQL database drivers (e.g., ODBC, JDBC)
- Data is stored in „natural“ representation. Often mapping between app data structures and relational structure used: Object-Relational Mapping (ORM)

Key/Value Store

- Proprietary APIs, sometimes SOAP/REST
- Direct serialization of app objects into the value part of records



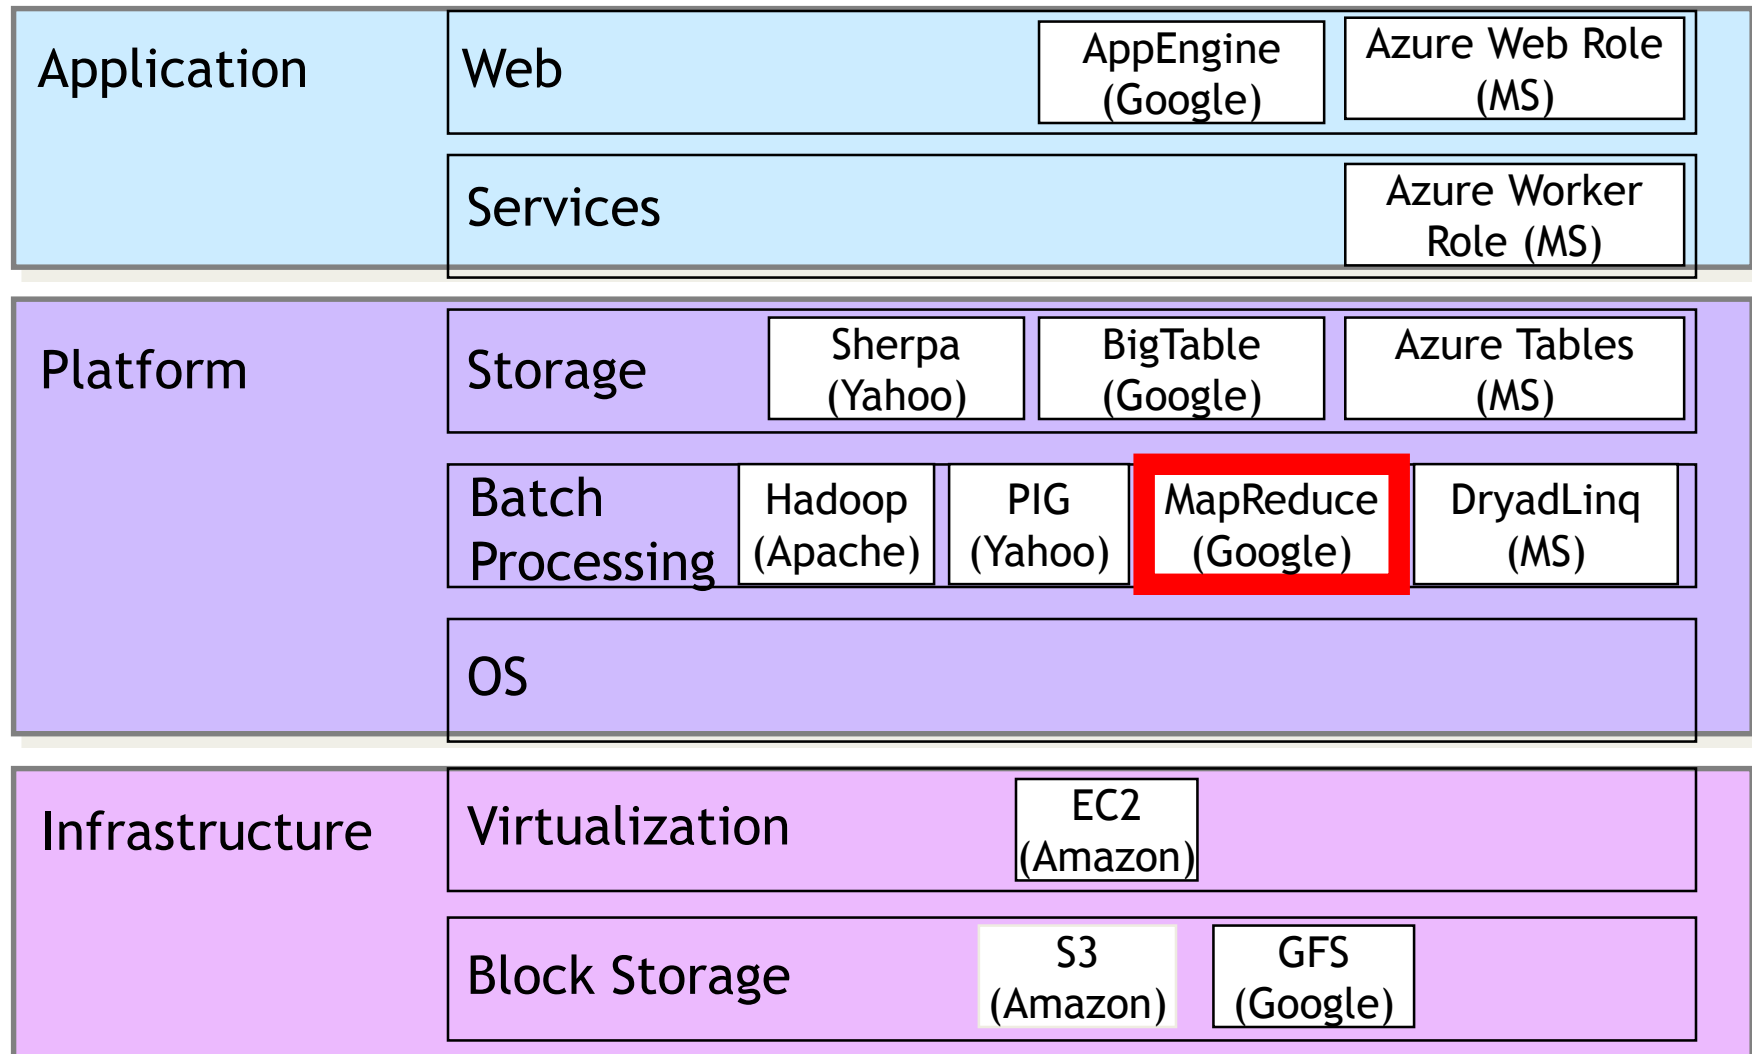
DBMS vs. Key/Value Store



- Applications for Key/Value Stores
 - Data is heavily document-oriented
→ more natural fit with the key/value model than the relational model
 - Large amounts of binary data
 - Application builds heavily on object-orientation → key/value store minimizes need for „plumbing“ code
 - Data store is cheap & integrates easily with cloud platform
 - On-demand, vast scalability is of vital importance
 - Thousands of servers
 - Terabytes of in-memory data
 - Petabyte of disk-based data
 - Millions of reads/writes per second, efficient scans



The „Cloud Stack“





Consider problem: count number of occurrences of each word in large document collection (in how many doc's does it appear?)

■ MapReduce Programming Model

- Inspired from **map** and **reduce** operations commonly used in functional programming languages like Lisp

1.map: (key1, val1) \rightarrow [(key2, val2)]

2.reduce: (key2, [val2]) \rightarrow [val3]

■ Map

- is a pure function (always evals to same result; no side effects)
- processes input key/value pair
- produces *set of* key/value pairs (intermediate results)
- e.g.: (document_name, document_content) \rightarrow [(word₁, 1), ..., (word_n, 1)]

■ Reduce

- Combines all intermediate values for a particular key
- Produces a set of merged output values (usually just one)
- e.g.: (word, [1, ..., 1]) \rightarrow [n]



- Example: count word occurrences

```
map (String input_key, String input_value):
```

```
    // input_key: document name
```

```
    // input_value: document contents
```

```
    for each word w in input_value:
```

```
        EmitIntermediate(w, "1");
```

```
reduce(String output_key, Iterator intermediate_values):
```

```
    // output_key: a word
```

```
    // output_values: a list of counts
```

```
    int result = 0;
```

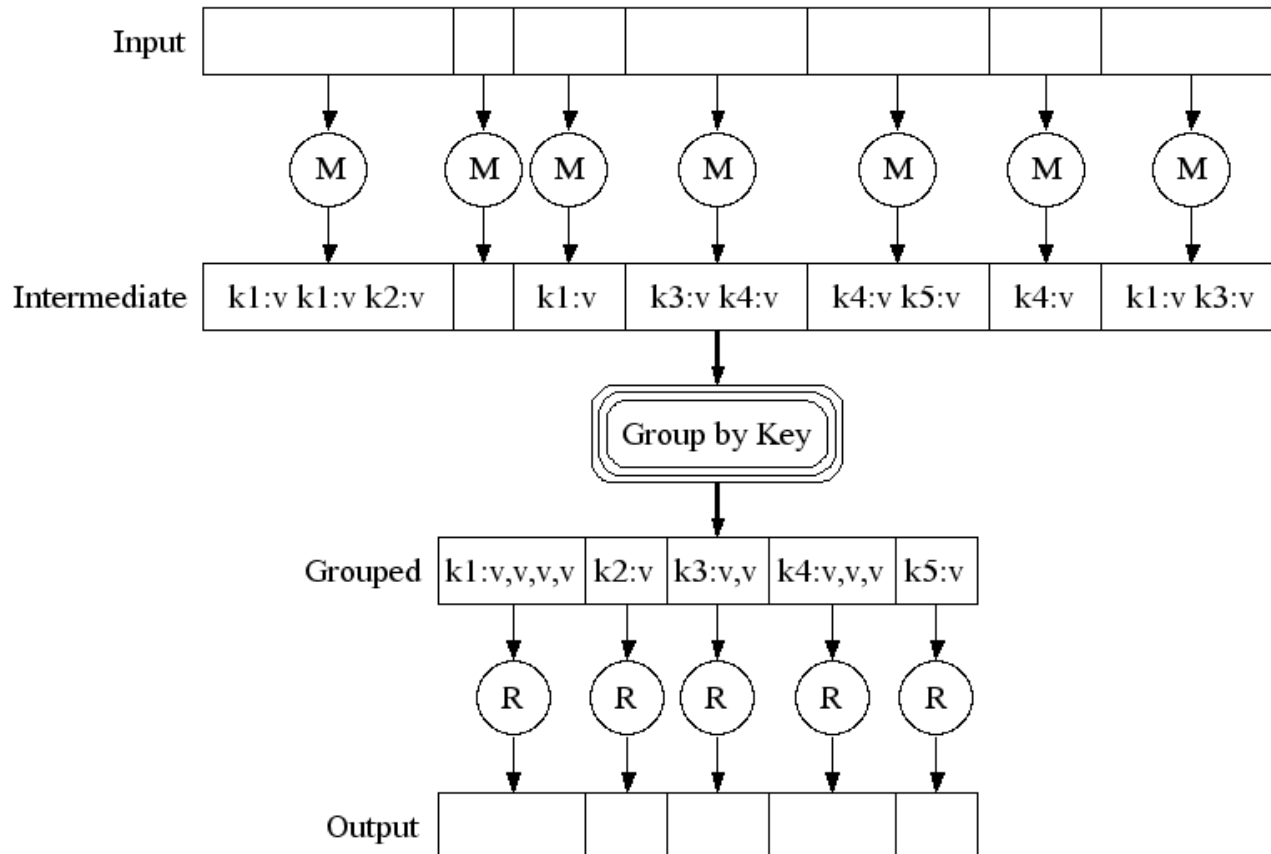
```
    for each v in intermediate_values:
```

```
        result += ParseInt(v);
```

```
    Emit(AsString(result));
```



MapReduce: Execution

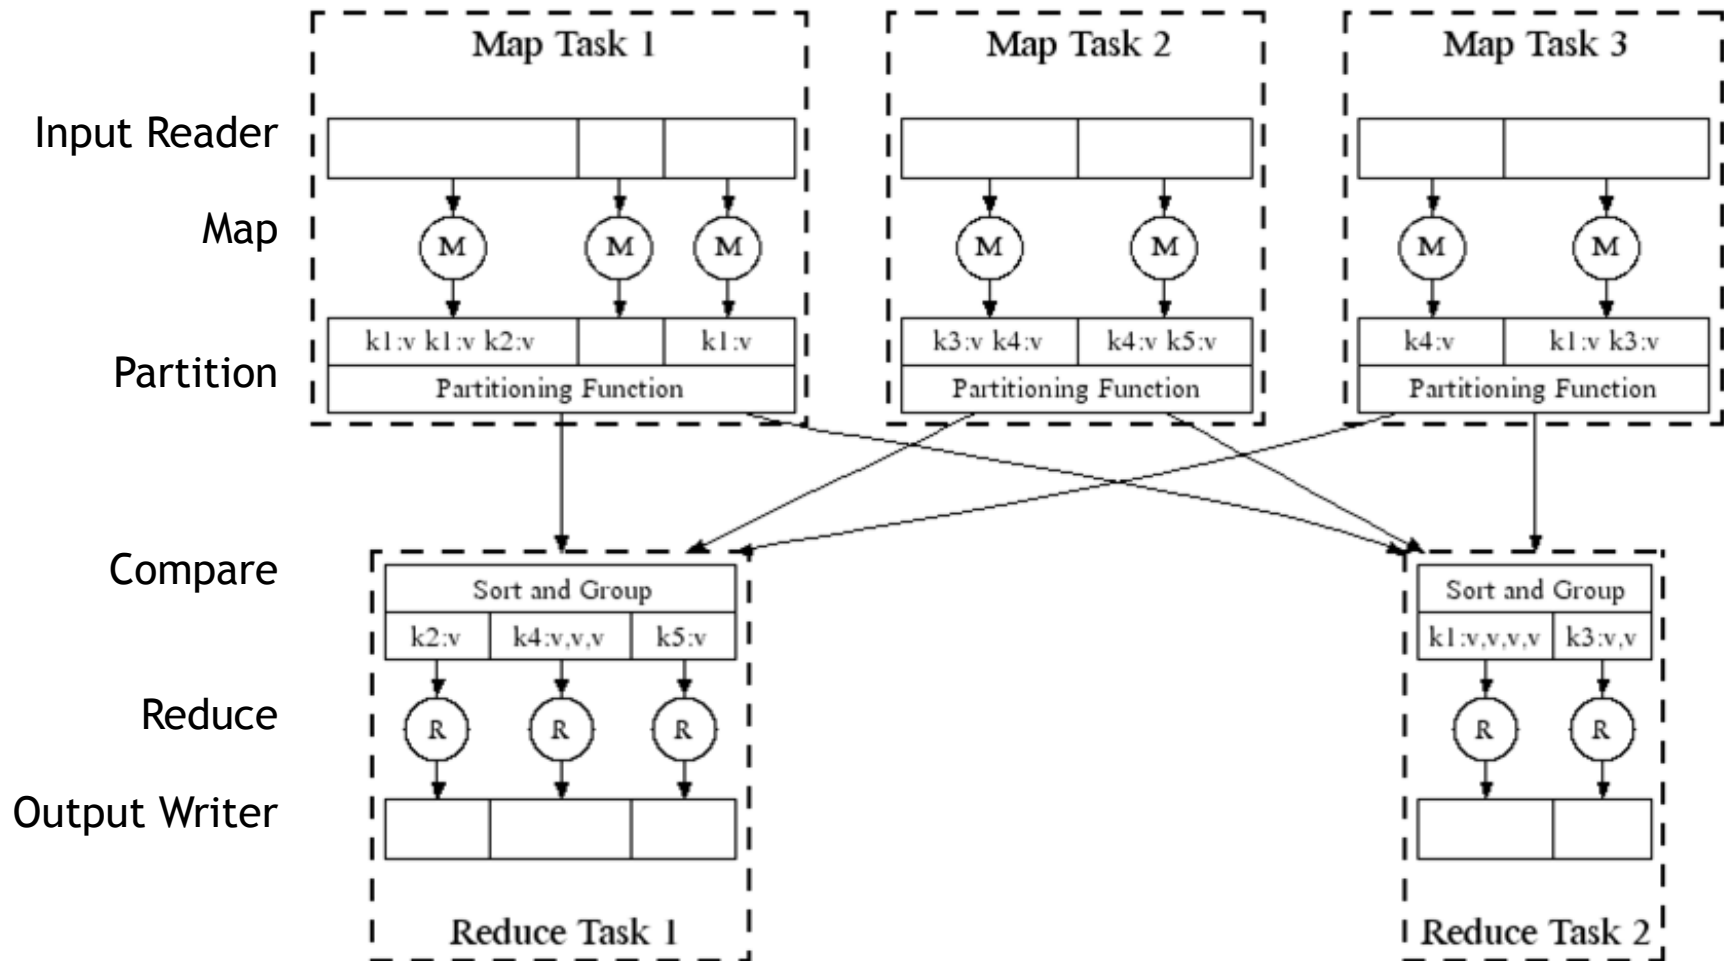




MapReduce: Distributed Execution



TECHNISCHE
UNIVERSITÄT
DARMSTADT





MapReduce: Data Flow



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Static part of MapReduce: “large distributed sort”

Variable part, defined by applications:

1. Input Reader

- Reads data from (distributed) file system, generates key/value pairs
- E.g., read directory full of text files & return each line as record

2. Map Function

3. Partition Function

- Assign output of a map to a reducer (on all map nodes, for all map functions):
`reducerIndex = partitionFunction(key, numberOfReducers)`
 - Note: all identical keys must be processed by same reducer
 - E.g., `reducerIndex = hash(key) % numberOfReducers`

4. Compare Function

- Defines which keys to be considered identical, groups input for reduce step

5. Reduce Function

6. Output Writer

- Writes output to (distributed) filesystem



The „Cloud Stack“



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Application	Web	AppEngine (Google)	Azure Web Role (MS)		
	Services		Azure Worker Role (MS)		
Platform	Storage	Sherpa (Yahoo)	BigTable (Google)	Azure Tables (MS)	
	Batch Processing	Hadoop (Apache)	PIG (Yahoo)	MapReduce (Google)	DryadLinq (MS)
	OS				
Infrastructure	Virtualization	EC2 (Amazon)			
	Block Storage	S3 (Amazon)	GFS (Google)		



Main purpose: run your web applications on Google's infrastructure

- **Python, Java, PHP, Go**
- **Sandbox** with several limitations
 - App code can only run as result of HTTP(S) request, may not spawn sub-processes
 - Cannot write to file system
 - App can only fetch URLs or use email services for communication
- **Datastore**
 - Distributed storage service with query engine and transactions
 - Stores data objects of mixed types, each object consists of typed key/value pairs
 - Query with GQL, a limited SQL variant
- **Google Cloud SQL**
 - MySQL in the cloud
- **Google accounts**
 - App can access information from Google accounts



The „Cloud Stack“

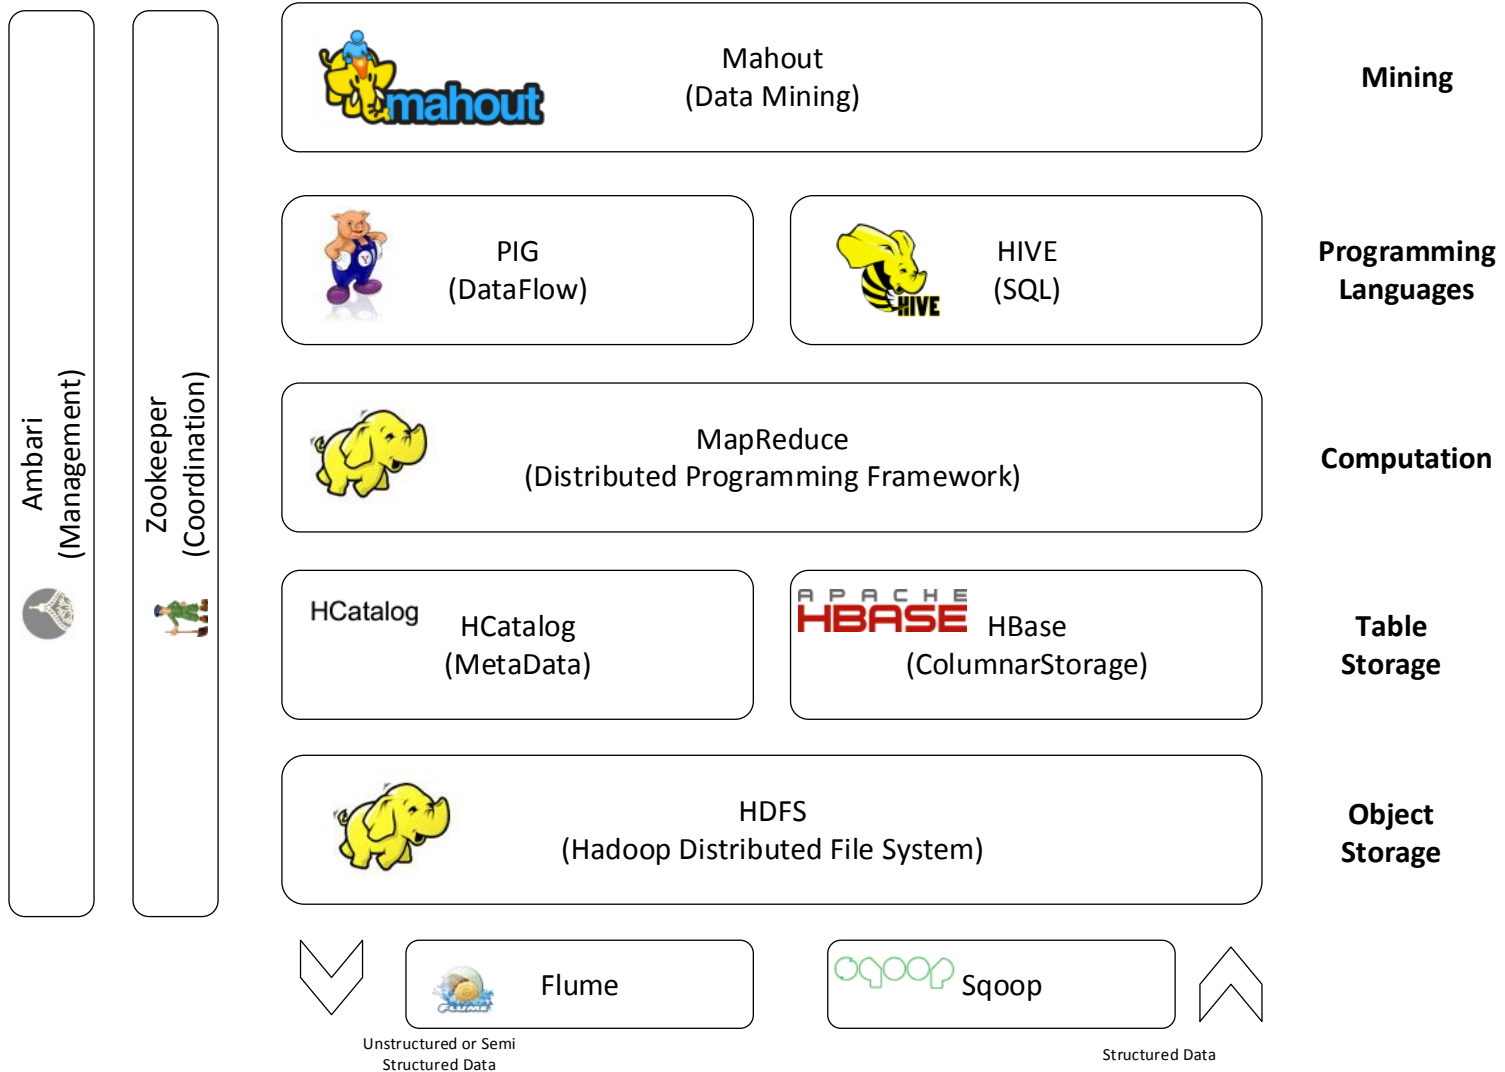


TECHNISCHE
UNIVERSITÄT
DARMSTADT

Application	Web	AppEngine (Google)	Azure Web Role (MS)		
	Services		Azure Worker Role (MS)		
Platform	Storage	Sherpa (Yahoo)	BigTable (Google)	Azure Tables (MS)	
	Batch Processing	Hadoop (Apache)	PIG (Yahoo)	MapReduce (Google)	DryadLinq (MS)
	OS				
Infrastructure	Virtualization	EC2 (Amazon)			
	Block Storage	S3 (Amazon)	GFS (Google)		



Hadoop Stack





Hadoop Stack

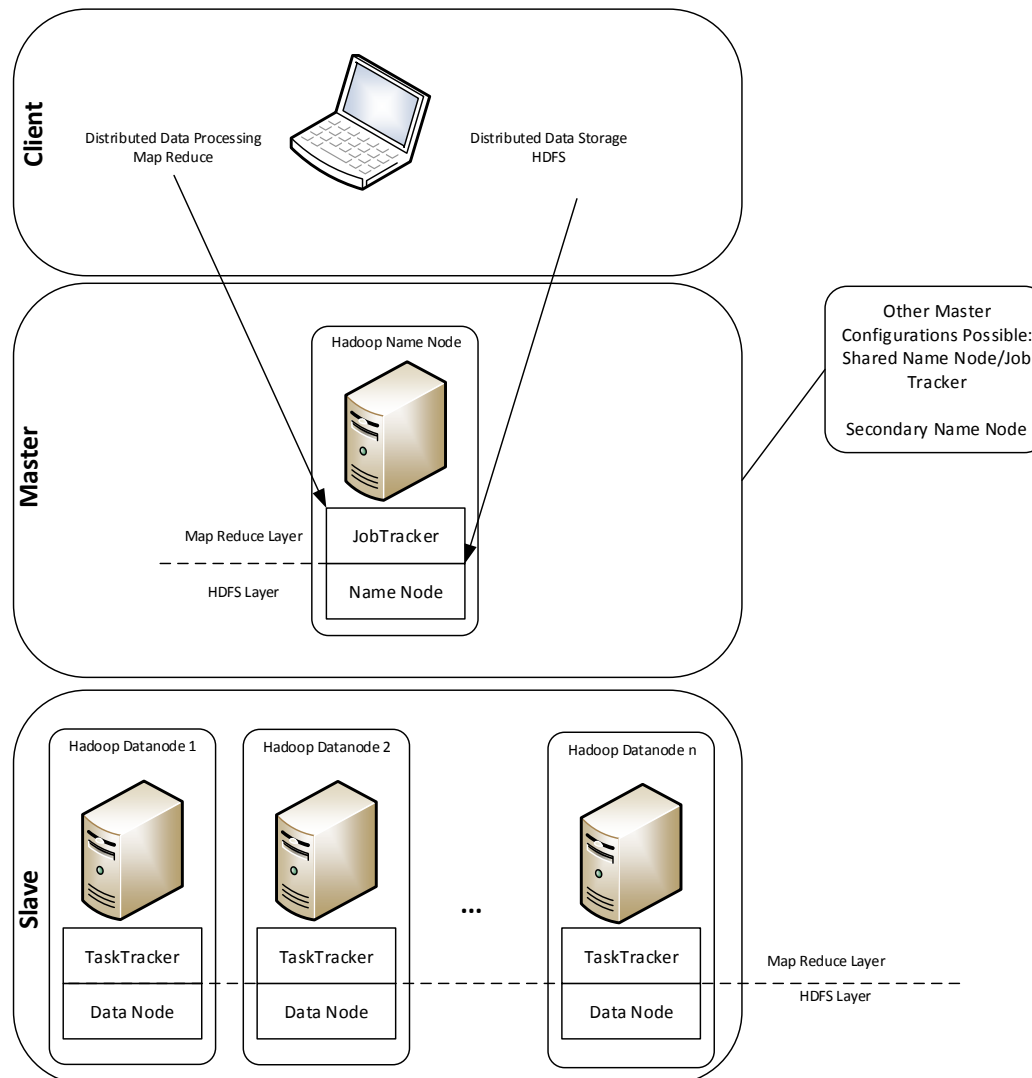
- Mahout: data mining capabilities
- Pig / Hive: High level code transformed into MapReduce Tasks
- Hadoop: MapReduce and Distributed File System
- Hbase: Key-Value Store, a persistent hash-map
- HCatalog: centralized meta-data management for Hadoop
- Scoop: exchange between Hadoop and relational database
- Flume: efficiently collecting, aggregating, and moving large amounts of log data
- Ambari/Zookeeper: Manage cluster: configuration, synchronization, etc.



Hadoop Overview



TECHNISCHE
UNIVERSITÄT
DARMSTADT





Hadoop HDFS Main Components

- Name Node (Master Node)
 - Centerpiece of an HDFS file system
 - Maintains and manages the blocks which are present on the Data Nodes
 - Keeps the directory tree of all files in the file system, and tracks where across the cluster the file data is kept
 - It is responsible for maintaining namespace image and edit log files
 - Any Changes to the file system namespace or its properties is recorded by Name Node
 - Mapping of file blocks to Data Nodes (Physical location of Data)
 - Aware of the data nodes for a particular file
 - Important files for Name Node
 - Image: File consists of meta data information
 - Check Point: Persistent record of image stored in the native file system
 - Journal: Modification log of image stored in local file system



Hadoop HDFS Main Components

- Data Node (Slave Node)
 - Slaves which are deployed on each machine and provide the actual storage
 - Responsible for serving read and write requests for the clients
 - Store and retrieve blocks when they are told to
 - Report back to Name Node periodically with the list of blocks they have
 - Data Nodes are the work horses of the Hadoop file system
 - Block replica is represented by two files
 - File that stores the data itself
 - File that stores the block meta data which includes checksum for the block and block's generation time stamp



Files in Hadoop



TECHNISCHE
UNIVERSITÄT
DARMSTADT

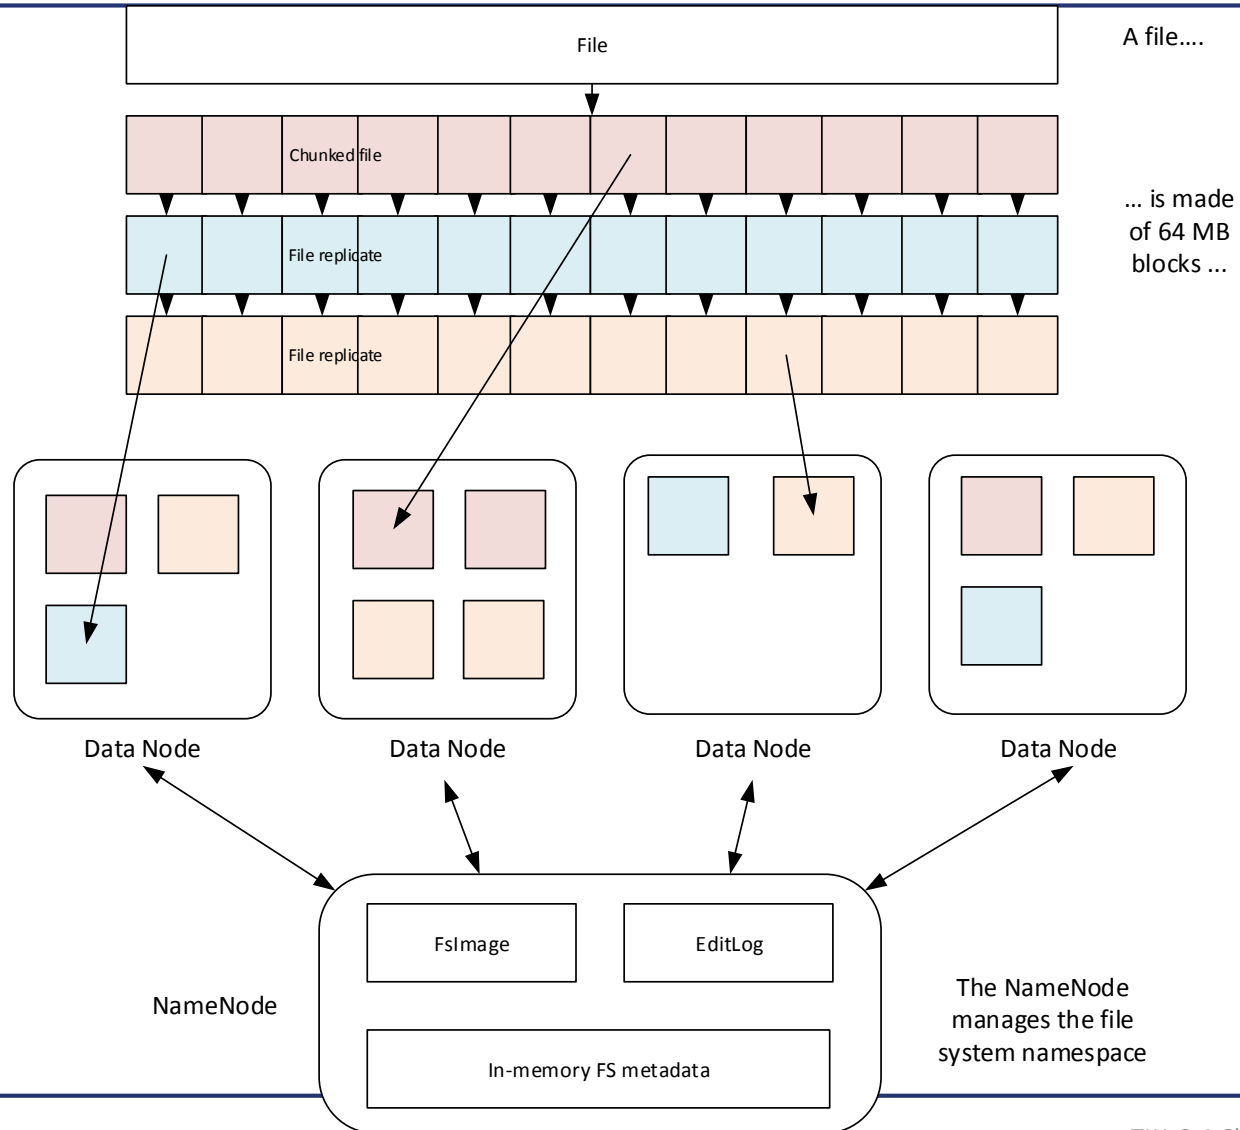
- Hadoop stores large files
- Files are broken into big chunks
 - Default 64MB
- These blocks are distributed among different machines
- Blocks are replicated for redundancy
 - Default: 3



Hadoop File Handlung

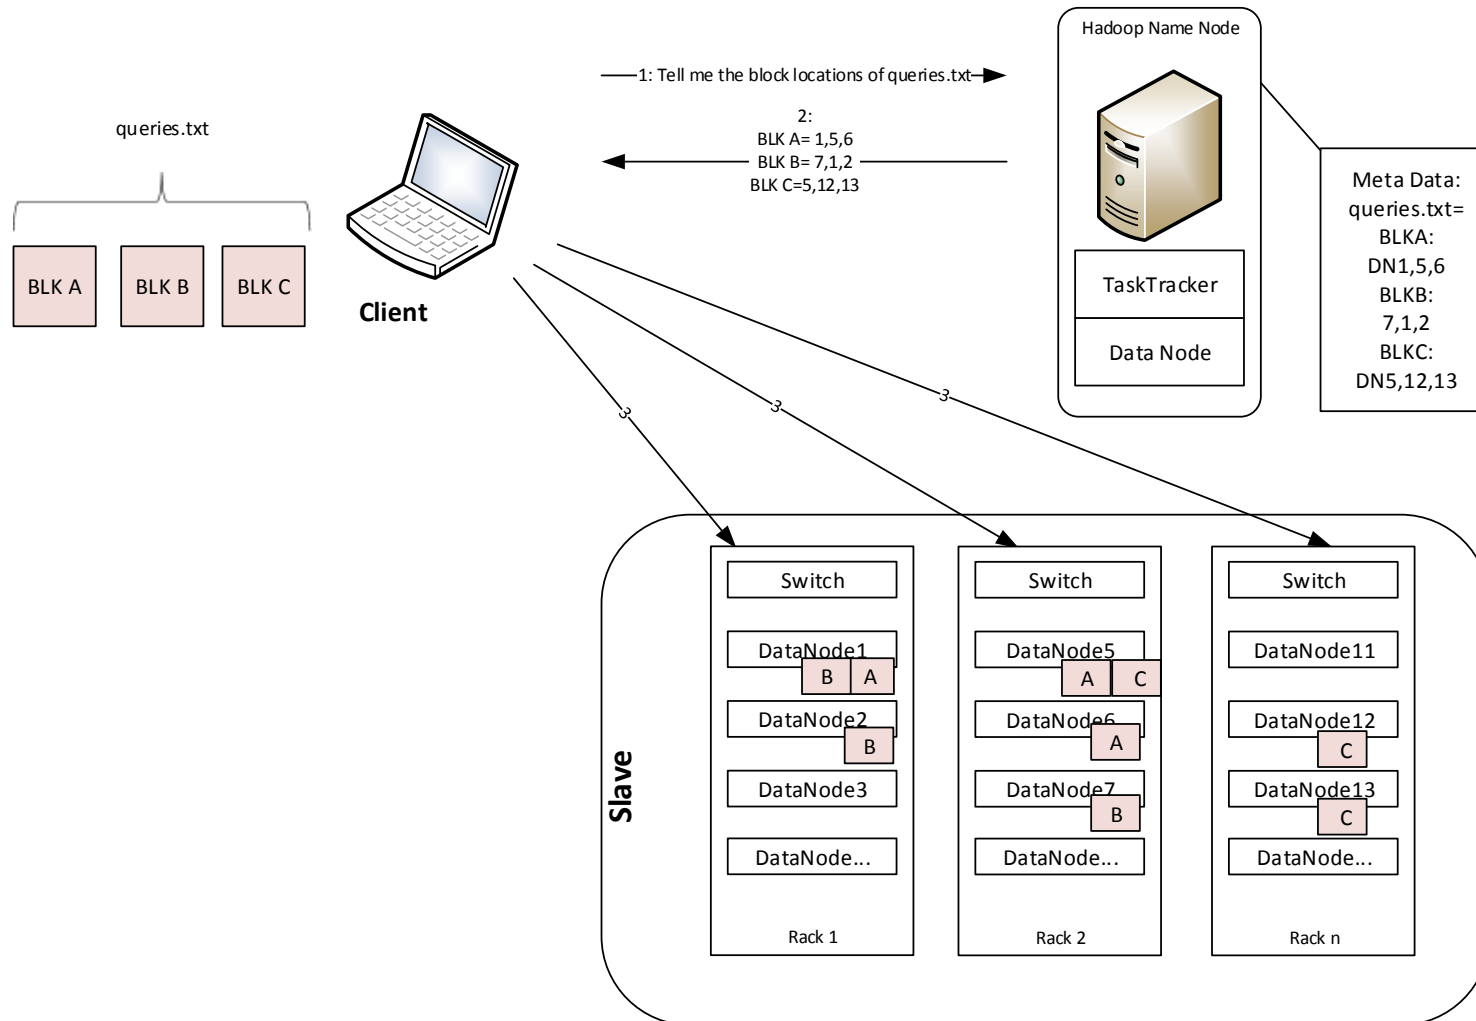


TECHNISCHE
UNIVERSITÄT
DARMSTADT



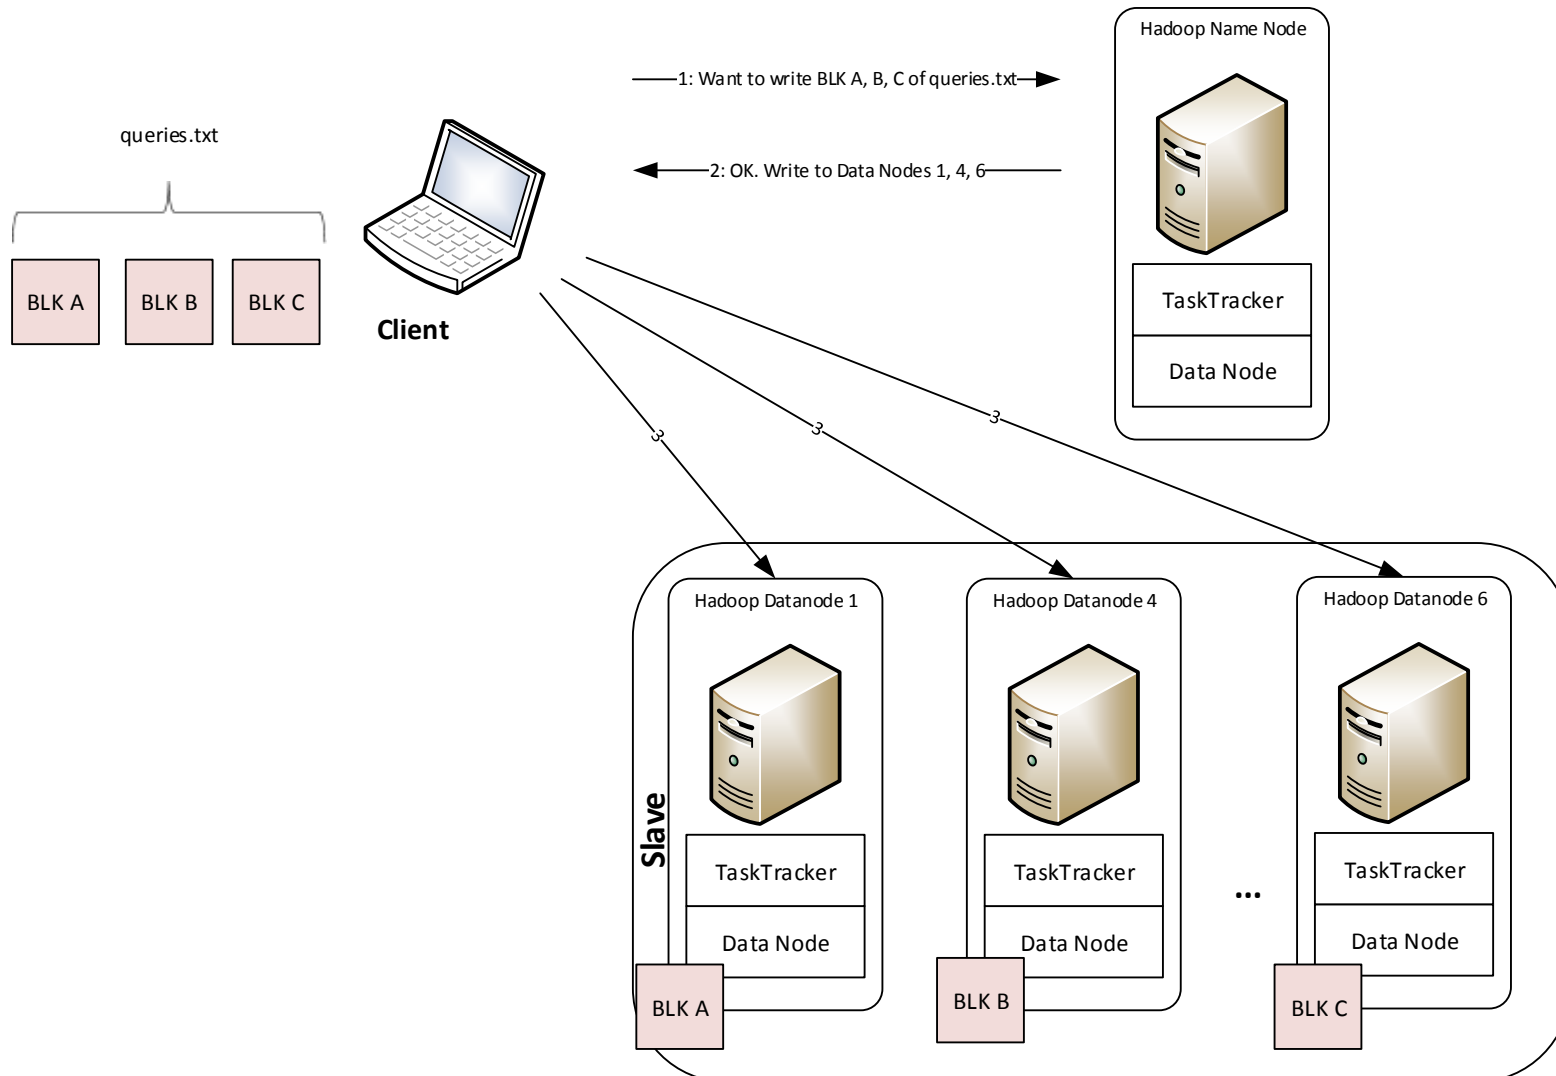


Read Process





Write Process





Node Replication



- Why replicate data?
 - Fault tolerance
 - Read latency improvement
 - Configure: Replication factor -> how many replicates?
- Replica placement policy
 - 1st replica on-rack: replicate on local rack
 - 2nd replica off-rack: random node in random, different rack
 - 3rd replica: different node, same rack as 2nd replica
 - Further replica: random, avoid too many on same rack



Hadoop and Map Reduce



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Client applications submit map reduce tasks to job tracker
- Job Tracker
 - Runs on master node
 - Schedules tasks
 - Allocates tasks to slave task trackers
 - Manages data nodes
 - Keeps track of resources, already running tasks
 - Single point of failure
- Task Tracker
 - Executes and manages the tasks received from job tracker
 - Handles data motion between map and reduce phase
 - Communicates status of task to job tracker (send heartbeats)



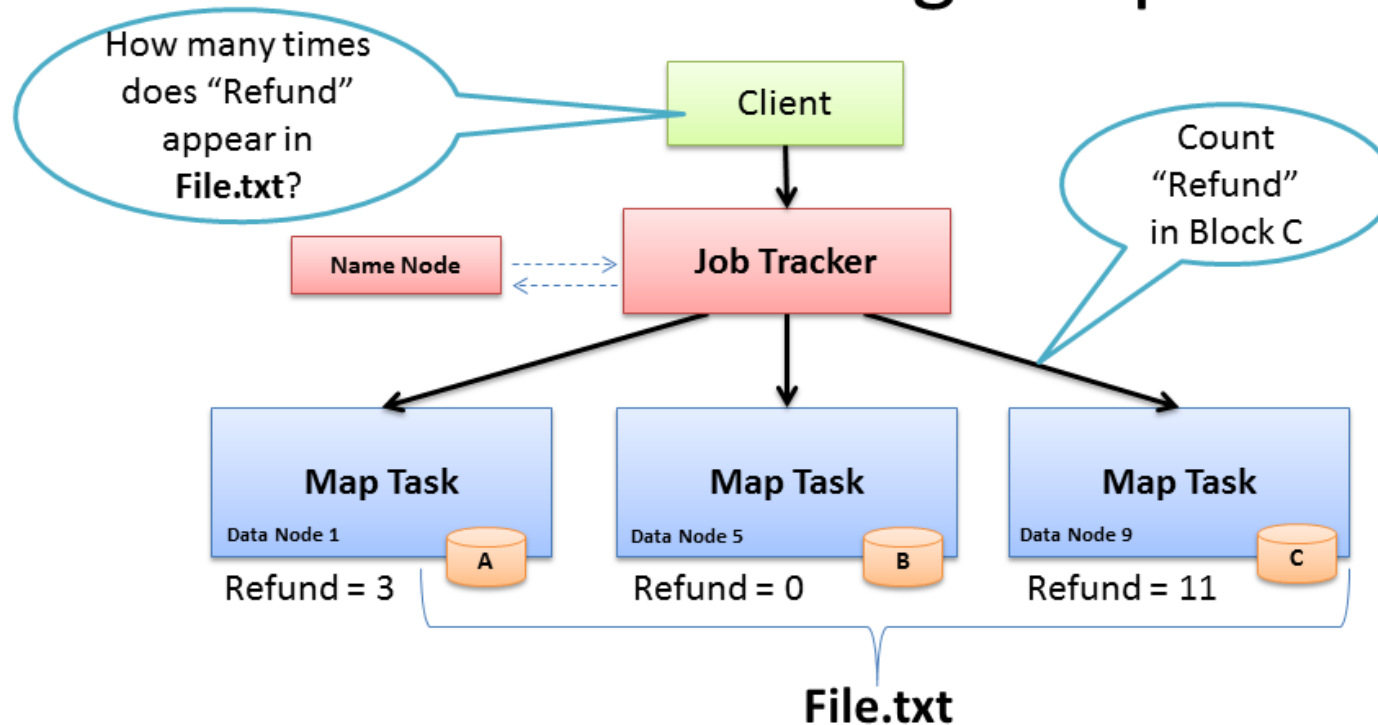
Map Reduce Process 1/4



- Client submits job to job tracker
- Job tracker gets location of data from name node
- Job tracker identifies task tracker nodes w. available slots close to data
- Job tracker submits tasks to task tracker nodes
- Job tracker monitors task trackers based on their heartbeat. No heartbeat=assume crash
- If task tracker fails with a task the job tracker is notified. Job tracker decides what to do (resubmit, mark data as to-be-avoided)
- Job tracker offers status information to client systems



Data Processing: Map



- **Map:** “Run this computation on your local data”
- Job Tracker delivers Java code to Nodes with local data

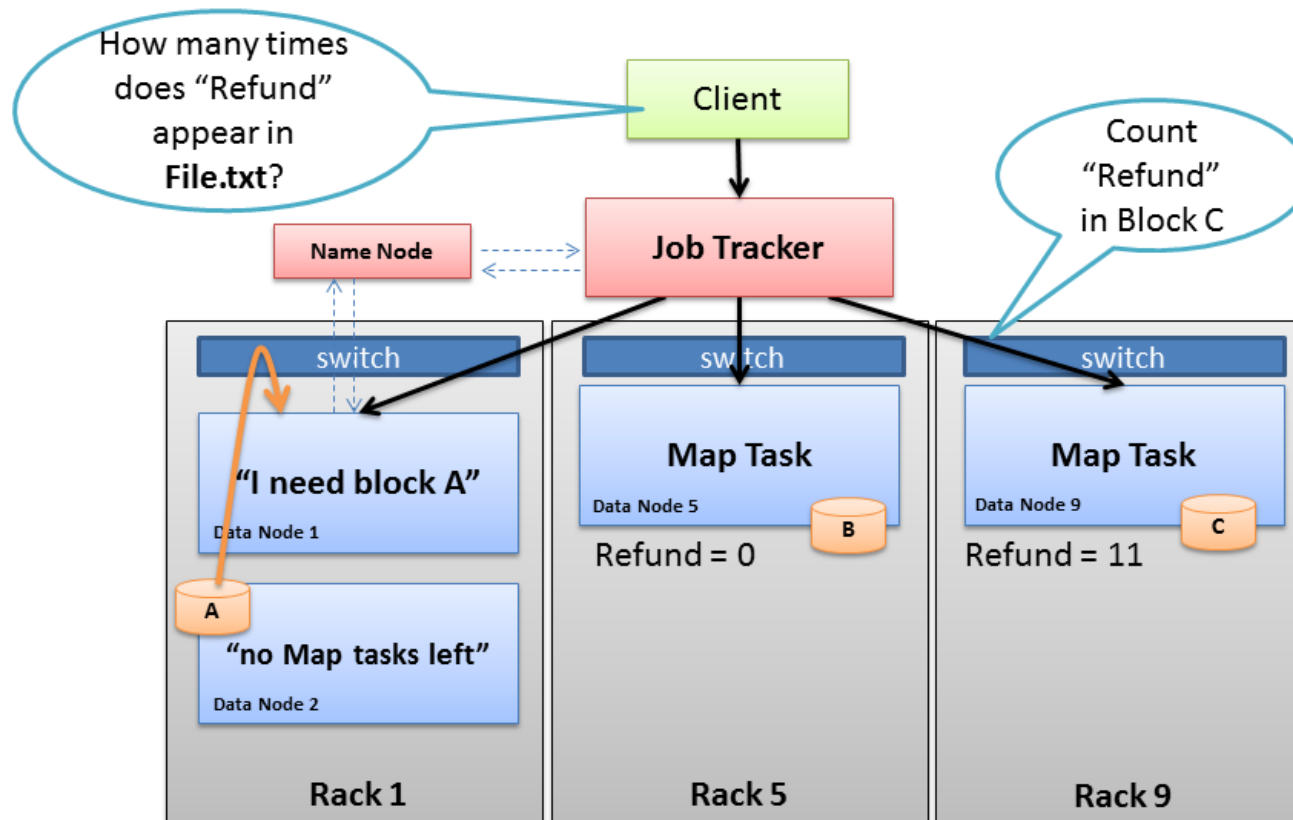


Map Reduce Process 2/4



TECHNISCHE
UNIVERSITÄT
DARMSTADT

What if data isn't local?



- Job Tracker tries to select Node in same rack as data
- Name Node rack awareness

BRAD HEDLUND .com

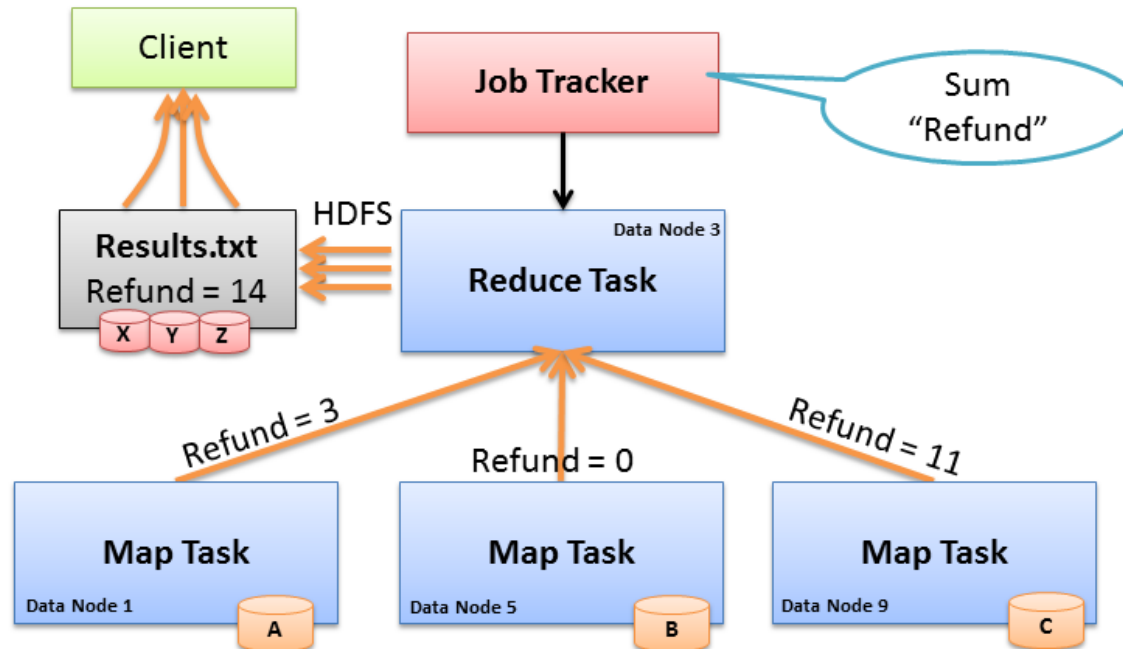


Map Reduce Process 2/4



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Data Processing: Reduce



- **Reduce:** “Run this computation across Map results”
- Map Tasks send output data to Reducer over the network
- Reduce Task data output written to and read from HDFS

BRAD HEDLUND .com



Hadoop Modes

- **Single node setup (standalone setup)** : Hadoop as a single Java process.
No daemons running and everything runs in a single JVM instance.
HDFS is not used.
- **Pseudo-distributed mode** : Simulating a cluster on a small scale.
Hadoop daemons run on a local machine. Different Hadoop daemons run in different JVM instances, but on a single machine.
HDFS is used instead of local FS.
- **Fully Distributed Mode**: Hadoop on a cluster of machines



One usage of Clouds: Big Data




TECHNISCHE
UNIVERSITÄT
DARMSTADT

- There is big data
 - Airline Jet: 10 TB of data for 30 min flight
 - NYSE: 1 TB of trade data for one trading day
 - Facebook: 1 billion active users, 3.2 billion likes per day
 - Twitter: 6000 Tweets per minute



- Hadoop provides an infrastructure built from relatively cheap commodity hardware
- What about the really fast machines and the really fast databases? E.g., your in-memory database (e.g., Hana) and the high-end server?
- Use Hadoop for the „much data“ things... raw logs, backups, analytic tasks for huge data sets
- Use the fast machines for „selected“ tasks -> use sqoop to transfer selected data (might be results from a map reduce job)





Existing Open Source Cloud Frameworks



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Eucalyptus
- Open Nebula
- Nimbus
- Cloud Stack
- Open Stack
- App Scale
- Typhoon AE
- Apache Hadoop



Summary



- Cloud Computing objectives
- Types of Cloud Service: SaaS, PaaS, IaaS
- Cloud vs. Grid Computing, Scalability
- Architecture of a Cloud Platform, “The Cloud Stack”
 - Block Storage: Amazon S3, Google File System
 - Virtualization: Amazon EC2
 - Structured Distributed Storage: BigTable
 - Batch Processing: MapReduce, Hadoop
 - Google App Engine



- Sanjay Ghemawat, Howard Gobioff, Shun-Tak Leung. "The Google File System". in *Symposium on Operating Systems Principles (SOSP)*, 2003.
- Jeffrey Dean, Sanjay Ghemawat. "MapReduce: Simplified Data Processing on Large Clusters". in *Operating Systems Design and Implementation (OSDI)*, 2004.
- Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Mike Burrows, Tushar Chandra, Andrew Fikes, Robert E. Gruber. "Bigtable: A Distributed Storage System for Structured Data". in *Operating Systems Design and Implementation (OSDI)*, 2006.
- Tony Bain. Is the Relational Database Doomed?. Retrieved Nov 25, 2015, from <http://www.readwriteweb.com/enterprise/2009/02/is-the-relational-database-doomed.php>, 2009.
- Donald Kossmann, Tim Kraska, and Simon Loesing: An Evaluation of Alternative Architectures for Transaction Processing in the Cloud, in: ACM SIGMOD'10, 2010.
- David Chappell: Introducing the Windows Azure Platform, retrieved Dec 07, 2010, from <http://www.microsoft.com/windowsazure/>
- Simon Munro: The Trouble With Sharding, retrieved Nov 25, 2015, from <http://simonmunro.com/2009/09/10/the-trouble-with-sharding/>
- Keith Brown, Sesha Mani: Microsoft Windows Identity Foundation (WIF) Whitepaper for Developers, 2009.